

**UCL**

Université  
catholique  
de Louvain



# Computer Networking : Principles, Protocols and Practice

## Part 4 : Network Layer

Olivier Bonaventure  
<http://inl.info.ucl.ac.be/>

# Network layer

---

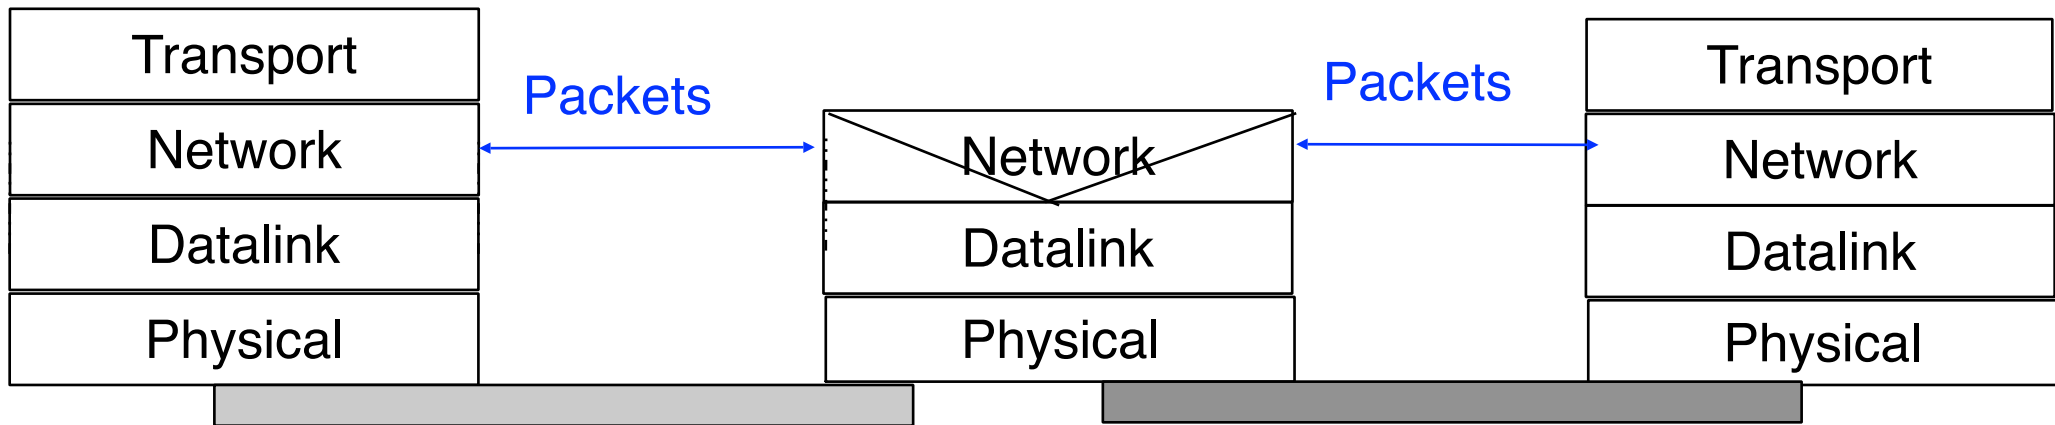
→ Basics  
Datagram mode  
Virtual circuits

Routing

IP : Internet Protocol

Routing in IP networks

# The network layer



## Goal

Allow packets to be forwarded from any source to any destination through heterogeneous networks and routers

## Services

Unreliable connectionless service

Reliable connection-oriented service

# Two types of datalink layers

---

# Two types of datalink layers

---

## WAN type datalink layer

PPP, HDLC

Reliable exchange of frames between two hosts  
attached to the same “link”

Mainly used by wide area networks

# Two types of datalink layers


---

## WAN type datalink layer

PPP, HDLC

Reliable exchange of frames between two hosts  
attached to the same “link”

Mainly used by wide area networks



# Two types of datalink layers


---

## WAN type datalink layer

PPP, HDLC

Reliable exchange of frames between two hosts attached to the same “link”

Mainly used by wide area networks




## LAN type datalink layer


Ethernet, Token Ring, FDDI, WiFi, Wimax,

Exchange of frames between hosts attached to the same LAN

limited geographical coverage



# The datalink service



## Service of datalink layer

### Unreliable connectionless service

Transmission of frames between hosts directly connected at the physical layer or directly attached to the same LAN

Unreliable transmission (frames can be lost but usually transmission errors are detected)

Most datalink layers have maximum frame length

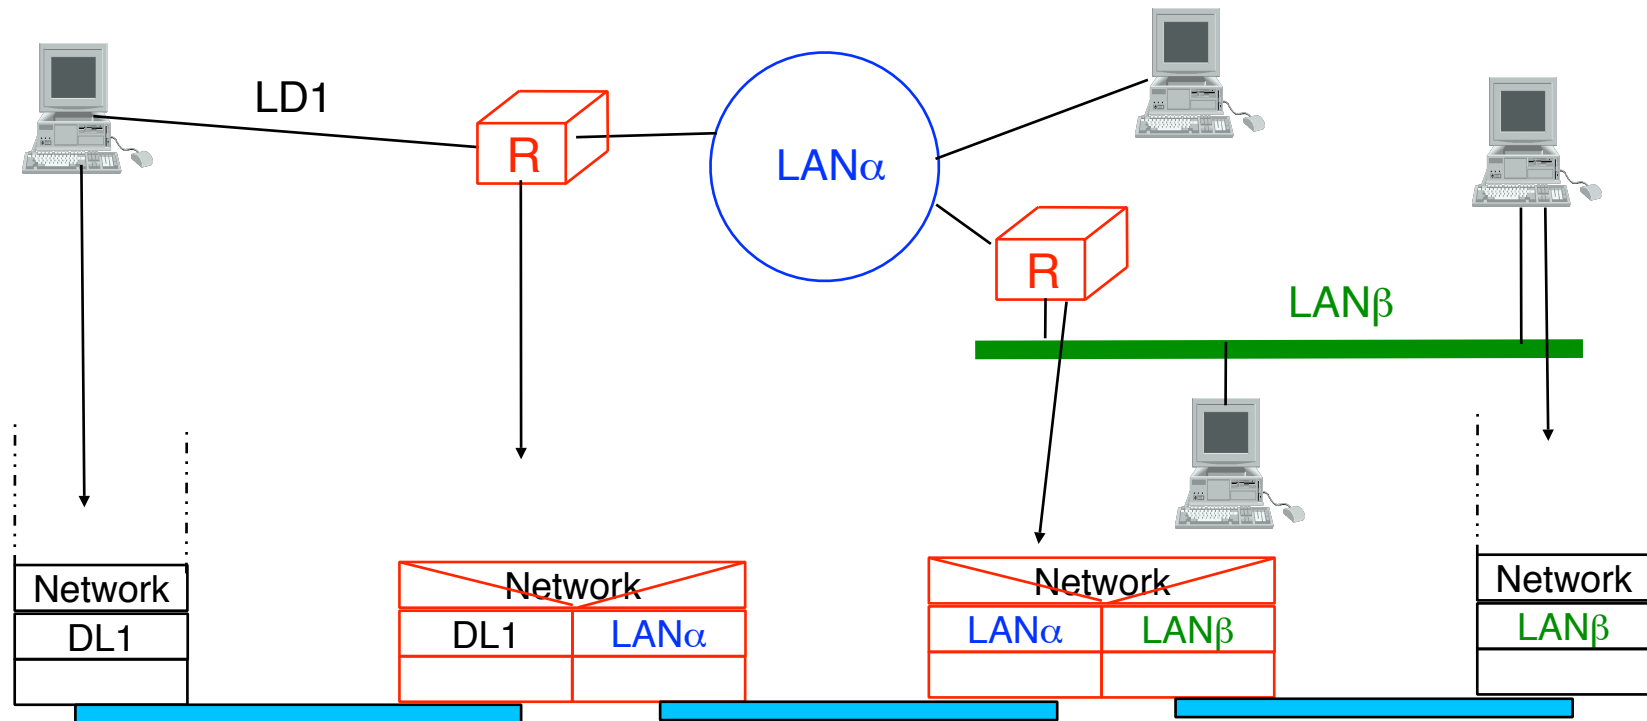
### Connection-oriented service, reliable or not

Transmission of frames between hosts directly connected at the physical layer or directly attached to the same LAN

Reliable or unreliable transmission



# Routers



**Router**

Relay within the network layer  
packet is unit of transmission

# Network layer

## Basic principles

---

# Network layer

## Basic principles

---

Each host/router must be identified by a **network layer address** which is independent from its datalink layer address

# Network layer

## Basic principles

---

Each host/router must be identified by a **network layer address** which is independent from its datalink layer address

Network layer forwards packets from source to destination through multiple routers

# Network layer

## Basic principles

---

Each host/router must be identified by a **network layer address** which is independent from its datalink layer address

Network layer forwards packets from source to destination through multiple routers

Network layer service must be completely independent from the service provided by the datalink layer

# Network layer

## Basic principles

---

- Each host/router must be identified by a **network layer address** which is independent from its datalink layer address
- Network layer forwards packets from source to destination through multiple routers
- Network layer service must be completely independent from the service provided by the datalink layer
- Network layer user should not need to know anything about the internal structure of the network layer to be able to send packets

# Internal organisation of the network layer

---

# Internal organisation of the network layer

---

Two possible organisations

datagrams

virtual circuits



# Internal organisation of the network layer

---

## Two possible organisations

datagrams

virtual circuits

The internal organisation of the network is orthogonal to the service provided, but often datagram mode is used to provide a connectionless service  
virtual circuits are used to provide a connection-oriented service

# Datagram transmission mode

---

# Datagram transmission mode

---

## Basics

Each route/host is identified by an **address**

Information is divided in packets

Each packet contains

- Source address

- Destination address

- Payload

## Router behavior

Upon packet arrival look at destination address and routing table to decide where the packet should be forwarded

- hop-by-hop forwarding, each routers takes a forwarding decision

# Datagram transmission mode

---

## Basics

- Each route/host is identified by an **address**

- Information is divided in packets

- Each packet contains

  - Source address

  - Destination address

  - Payload

- Router behavior

  - Upon packet arrival look at destination address and routing table to decide where the packet should be forwarded

    - hop-by-hop forwarding, each routers takes a forwarding decision

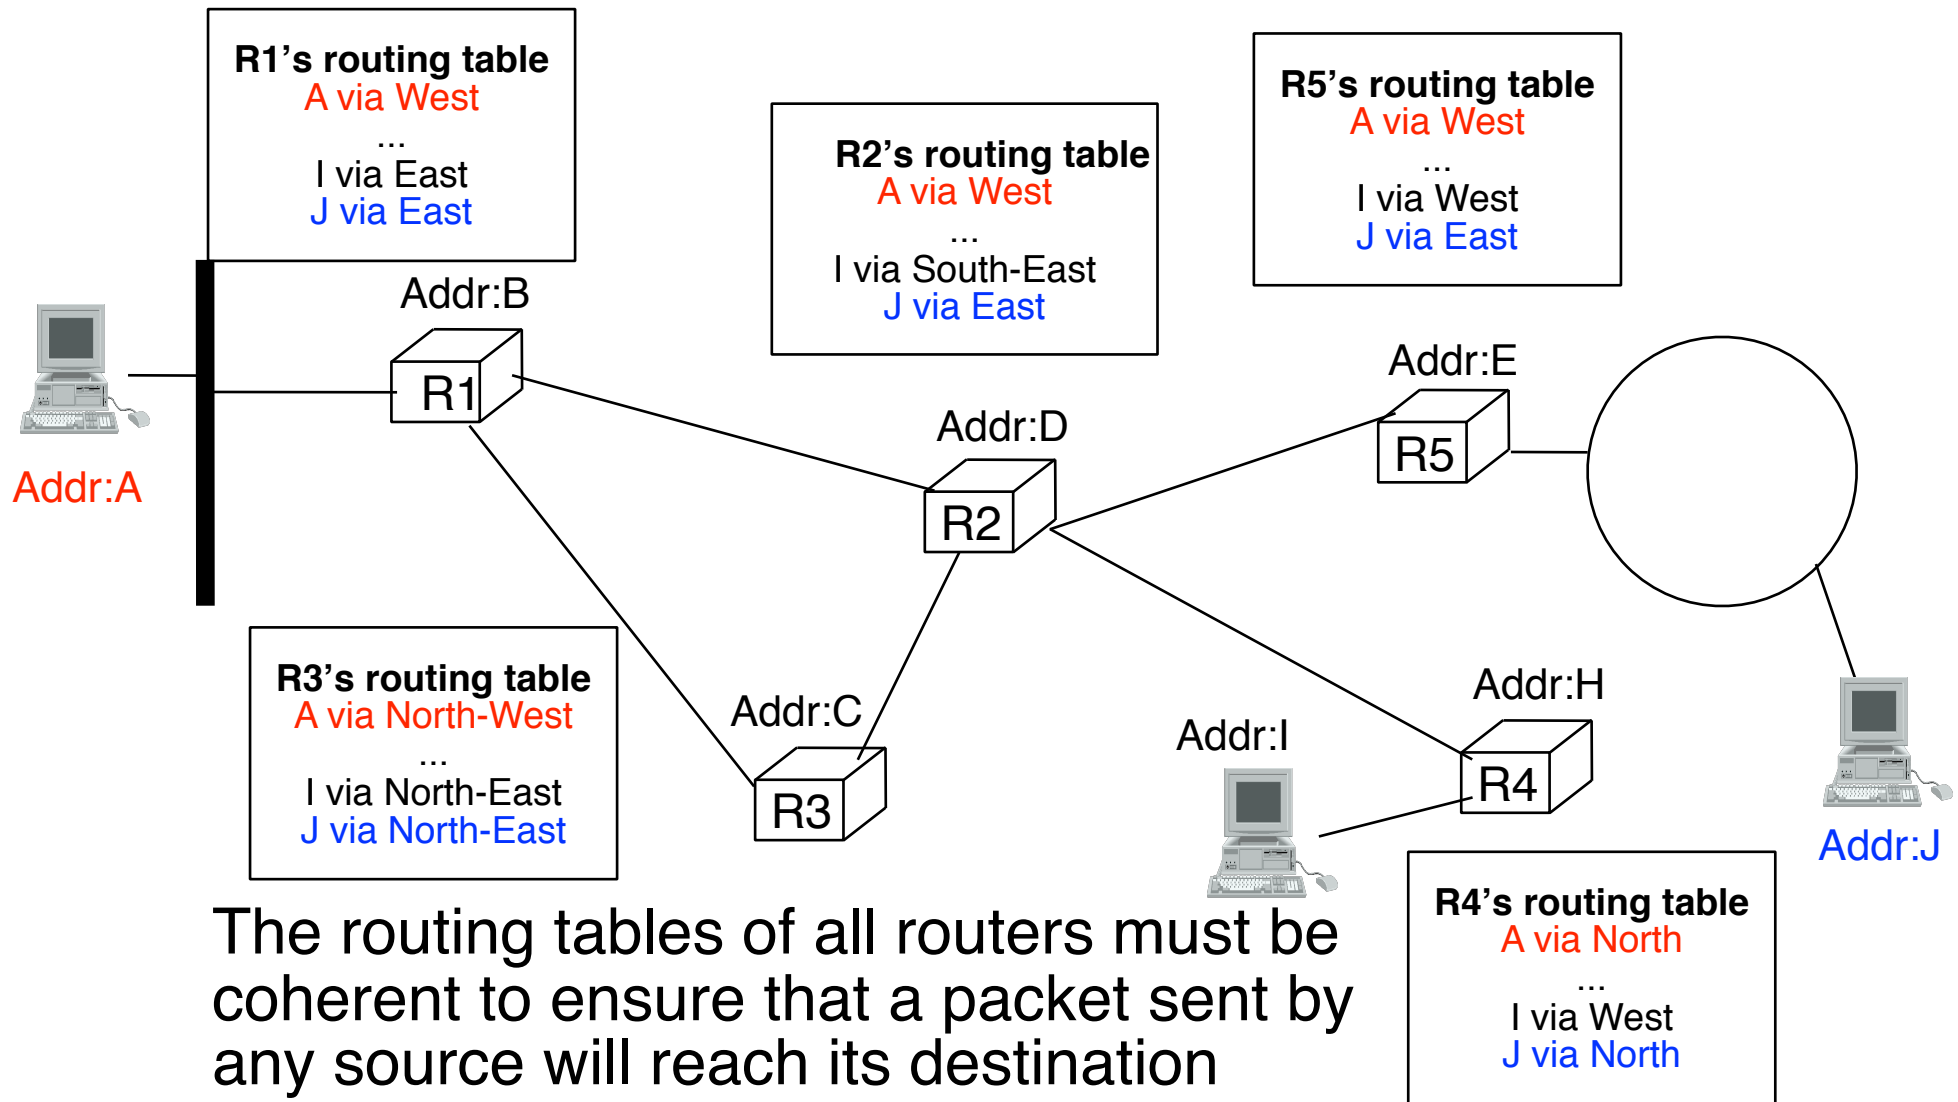
## Examples

- IP (IPv4 and IPv6)

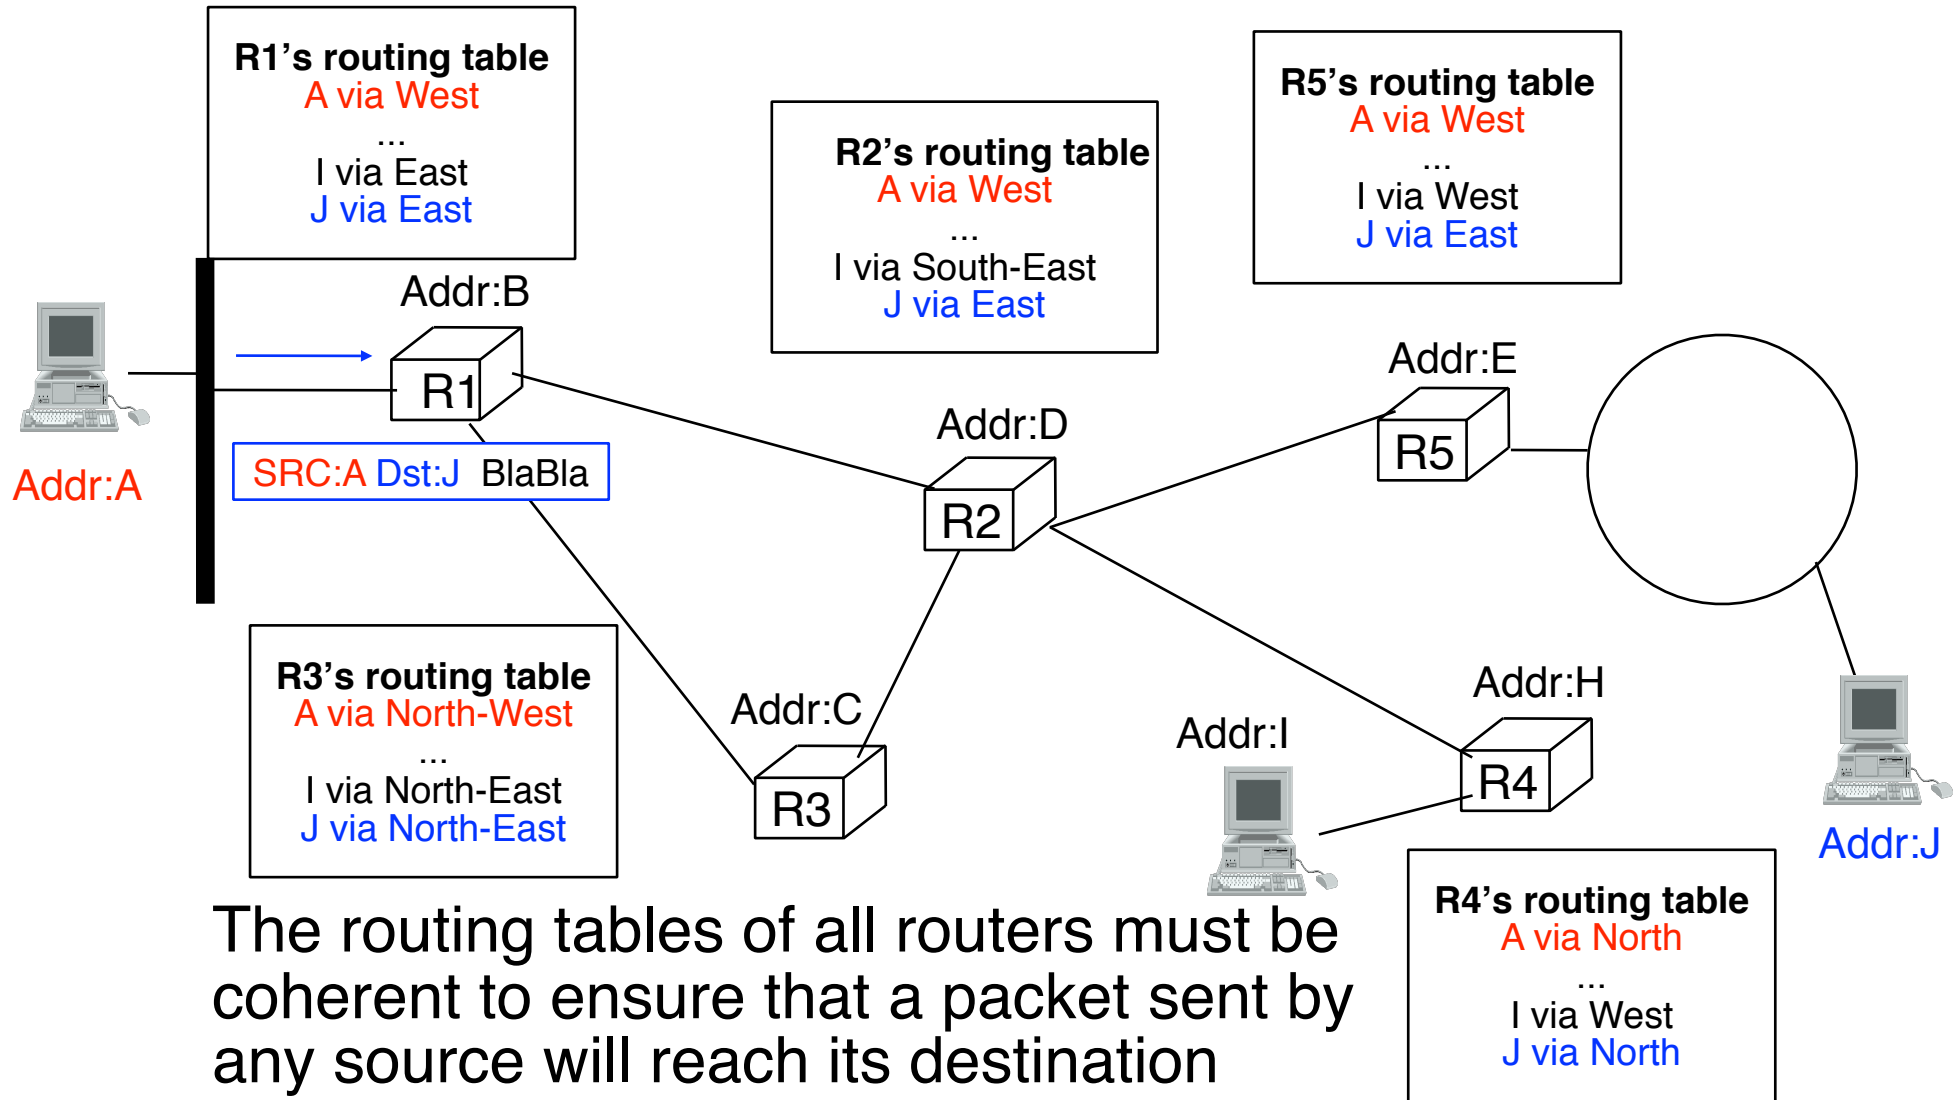
- CLNP

- IPX

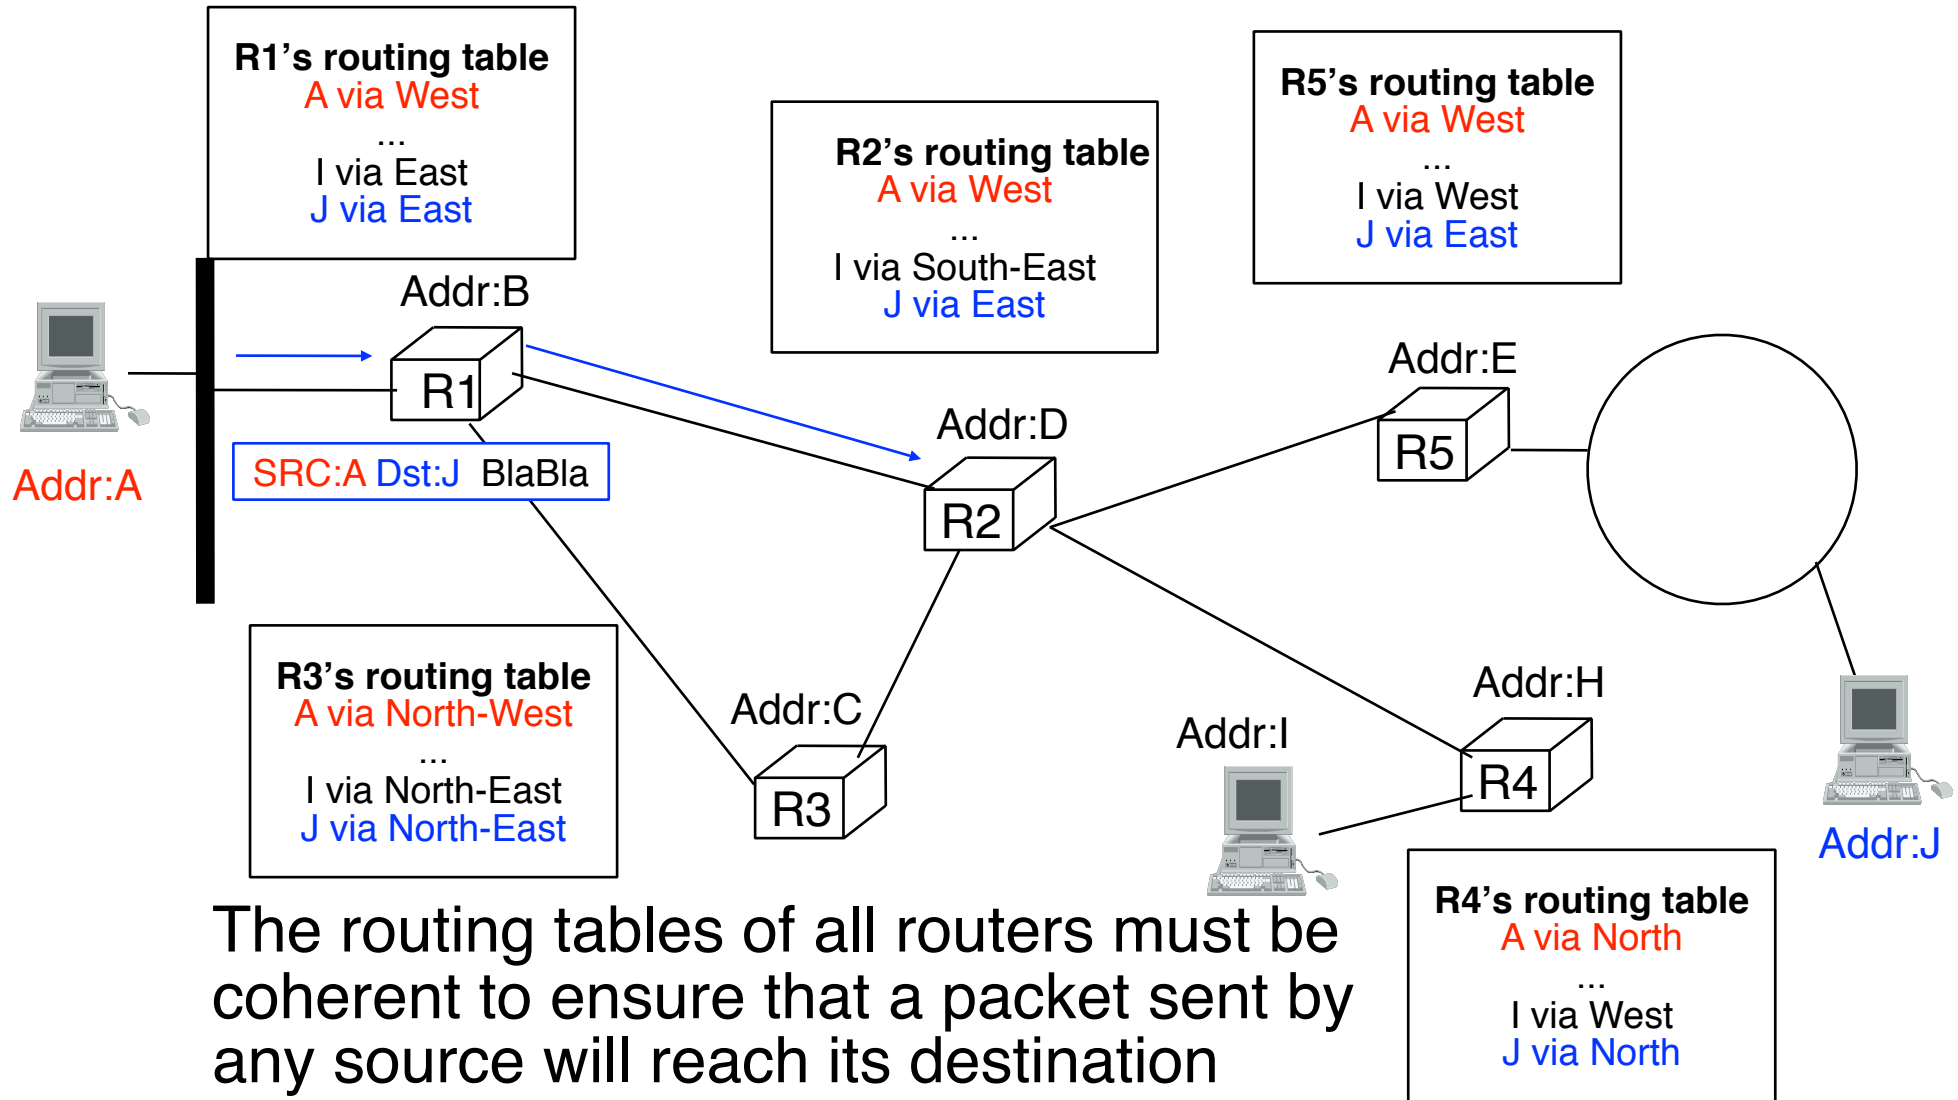
# Datagram transmission mode (2)



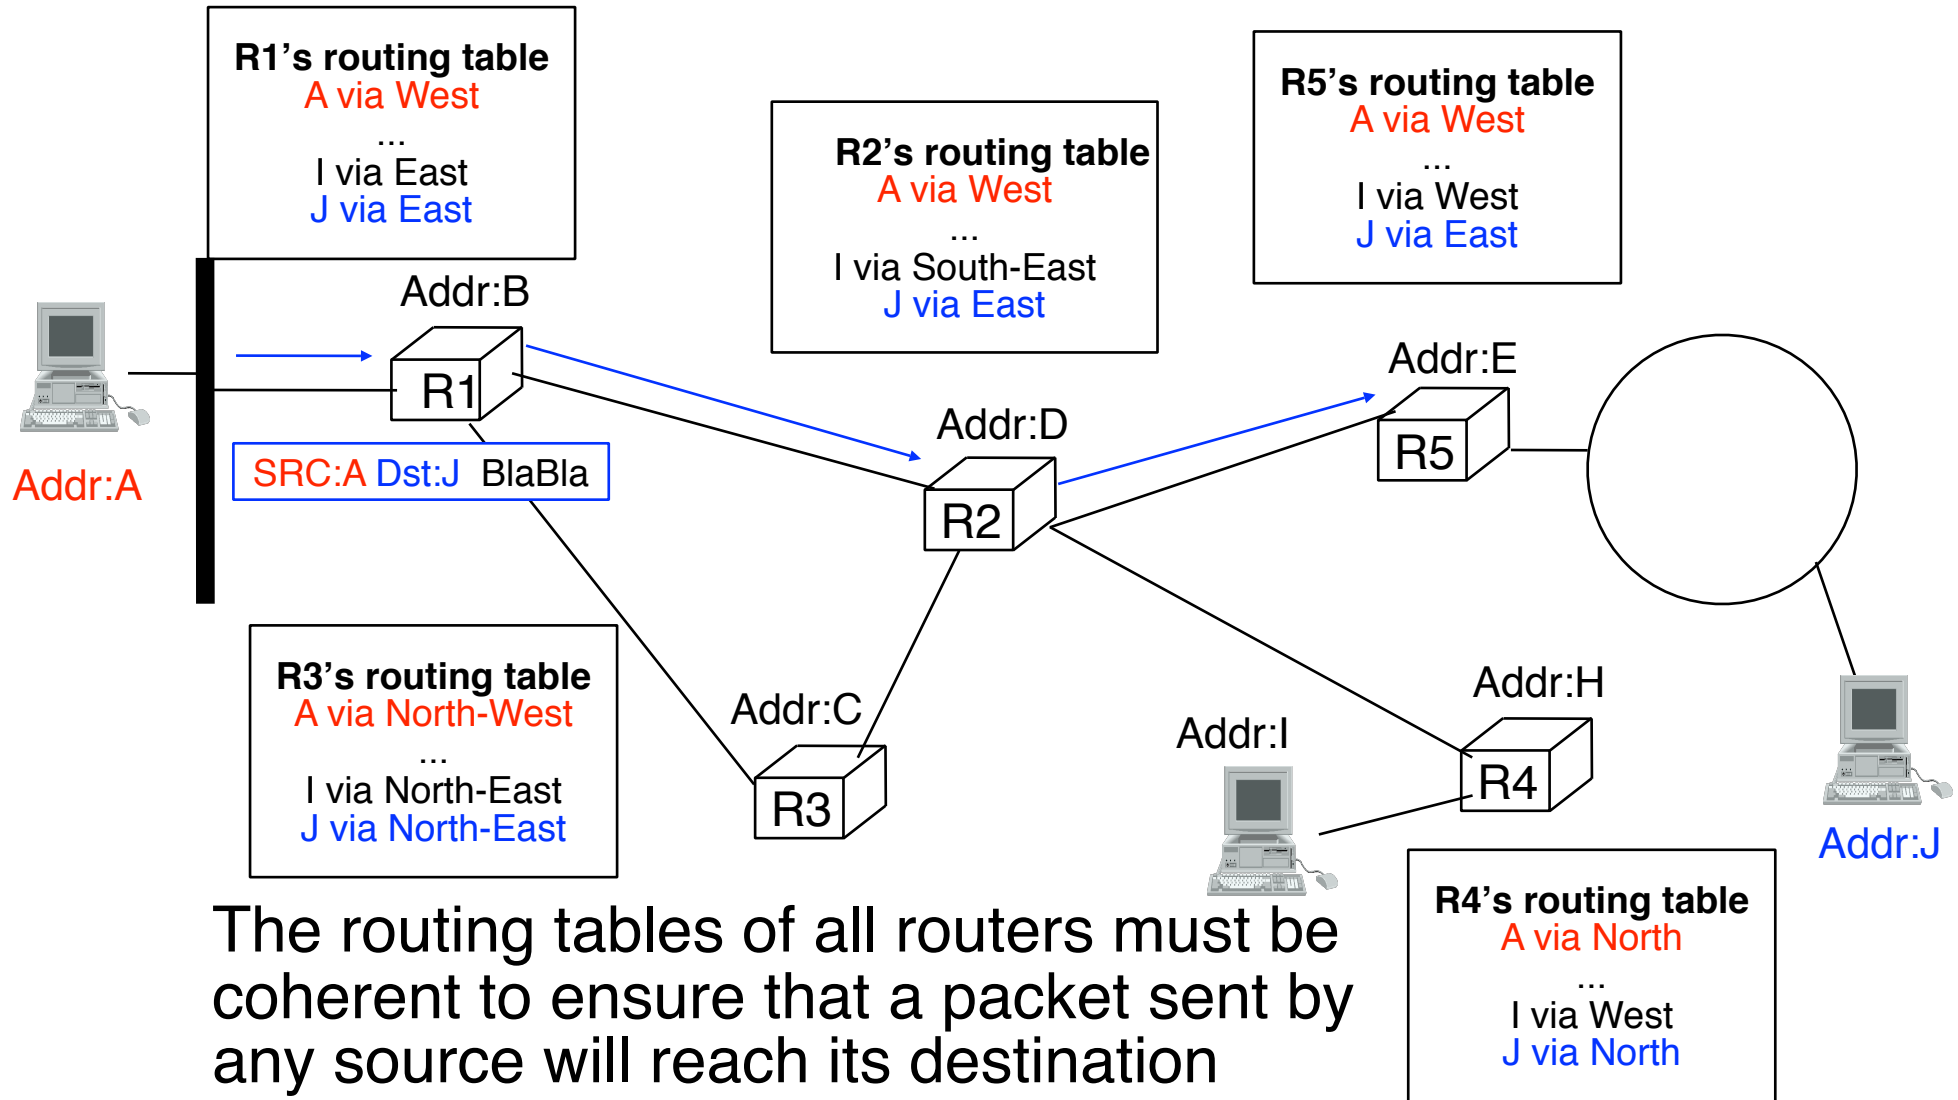
# Datagram transmission mode (2)



# Datagram transmission mode (2)

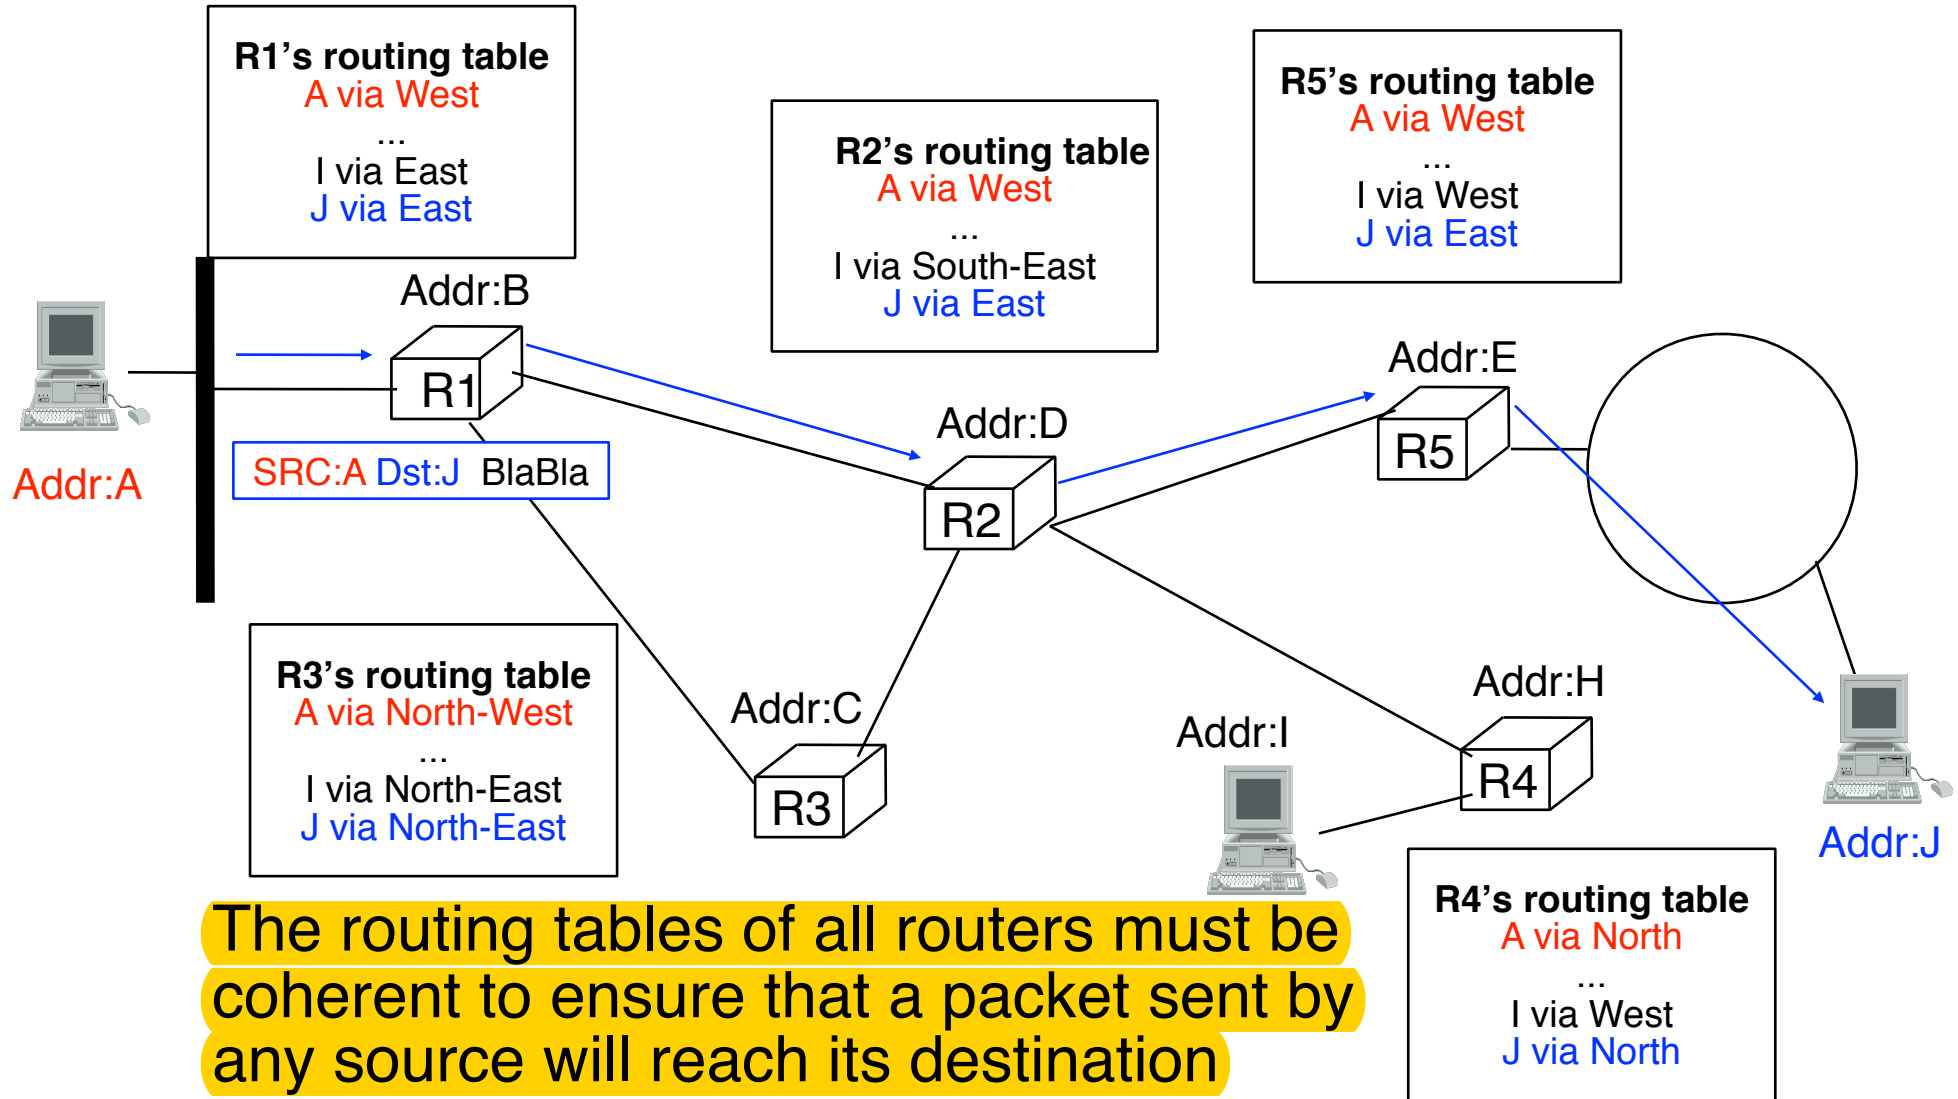


# Datagram transmission mode (2)





# Datagram transmission mode (2)



# Virtual circuit organisation

---

# Virtual circuit organisation

---

## Goals

Keep forwarding on the routers as simple as possible

consulting a routing table for each packet is costly from a performance viewpoint

# Virtual circuit organisation

---

## Goals

Keep forwarding on the routers as simple as possible

consulting a routing table for each packet is costly from a performance viewpoint

## Solution

Before transmitting packets containing data, create a virtual circuit that links source and destination through the network

During the virtual circuit establishment, efficient datastructures are updated on each transit router to simplify forwarding

Use the virtual circuits to forward the packets  
All packets will follow the same path

# Virtual circuit organisation

---

## Goals

- Keep forwarding on the routers as simple as possible

- consulting a routing table for each packet is costly from a performance viewpoint

## Solution

- Before transmitting packets containing data, create a virtual circuit that links source and destination through the network

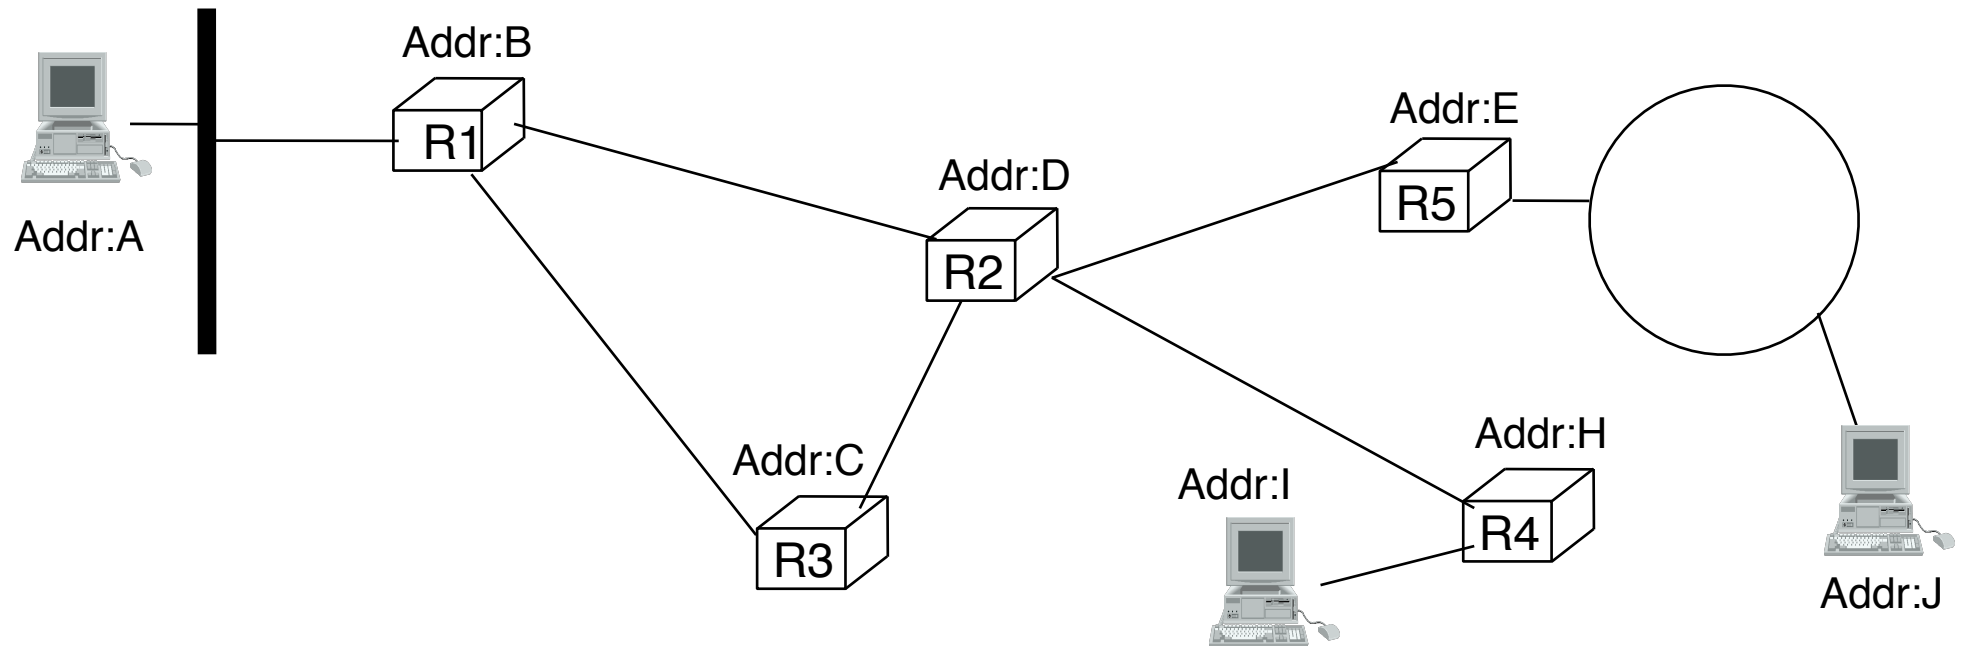
- During the virtual circuit establishment, efficient datastructures are updated on each transit router to simplify forwarding

- Use the virtual circuits to forward the packets
- All packets will follow the same path

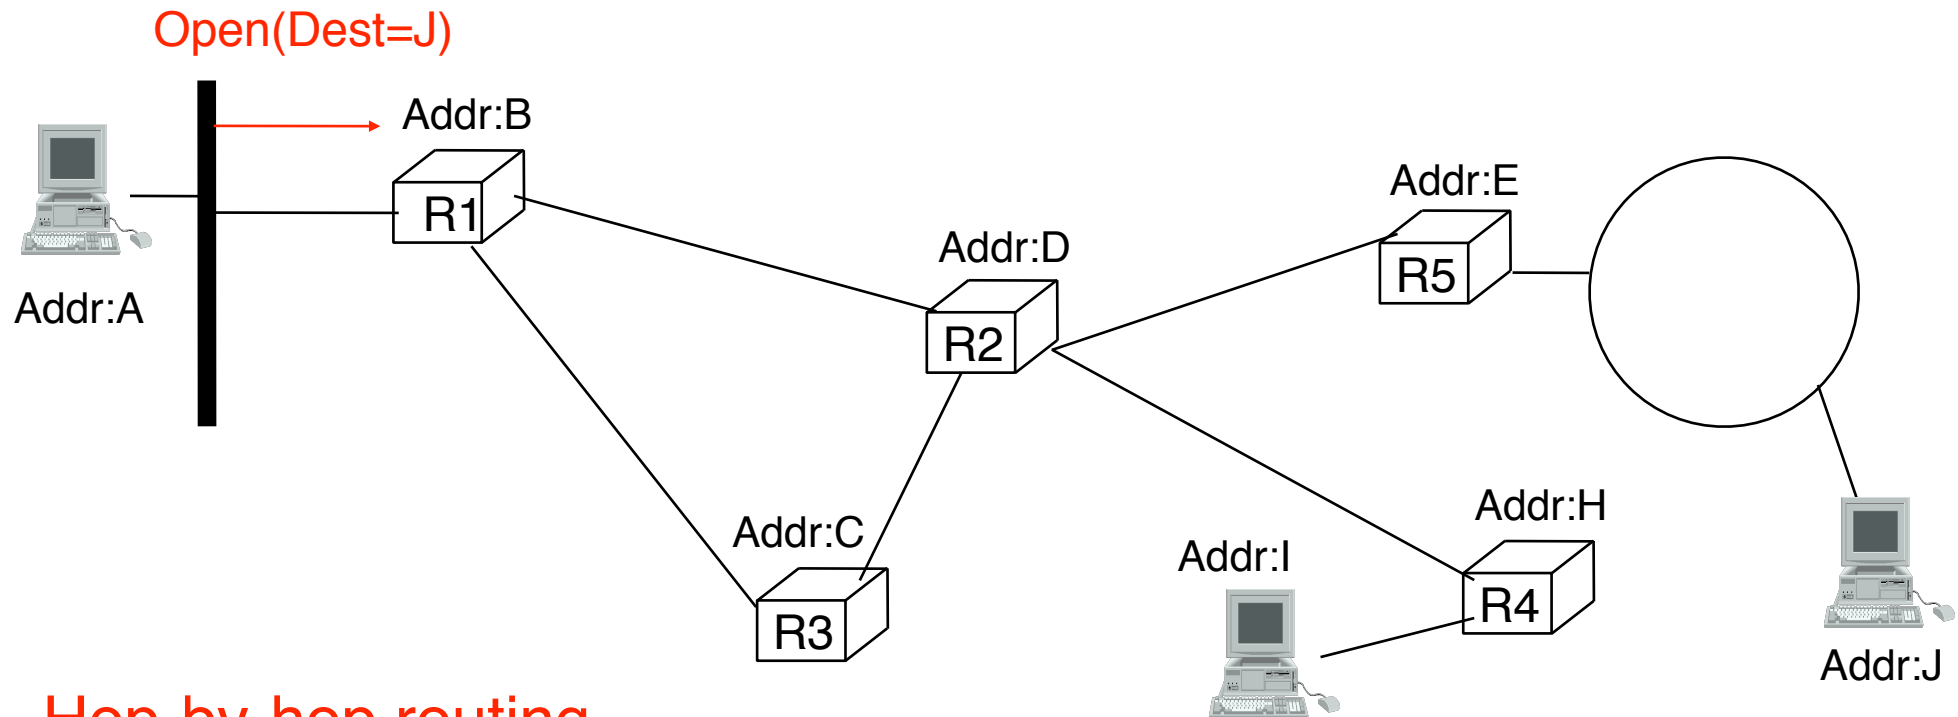
## Example

ATM, X.25, Frame Relay, MPLS, gMPLS

# Establishment of a virtual circuit



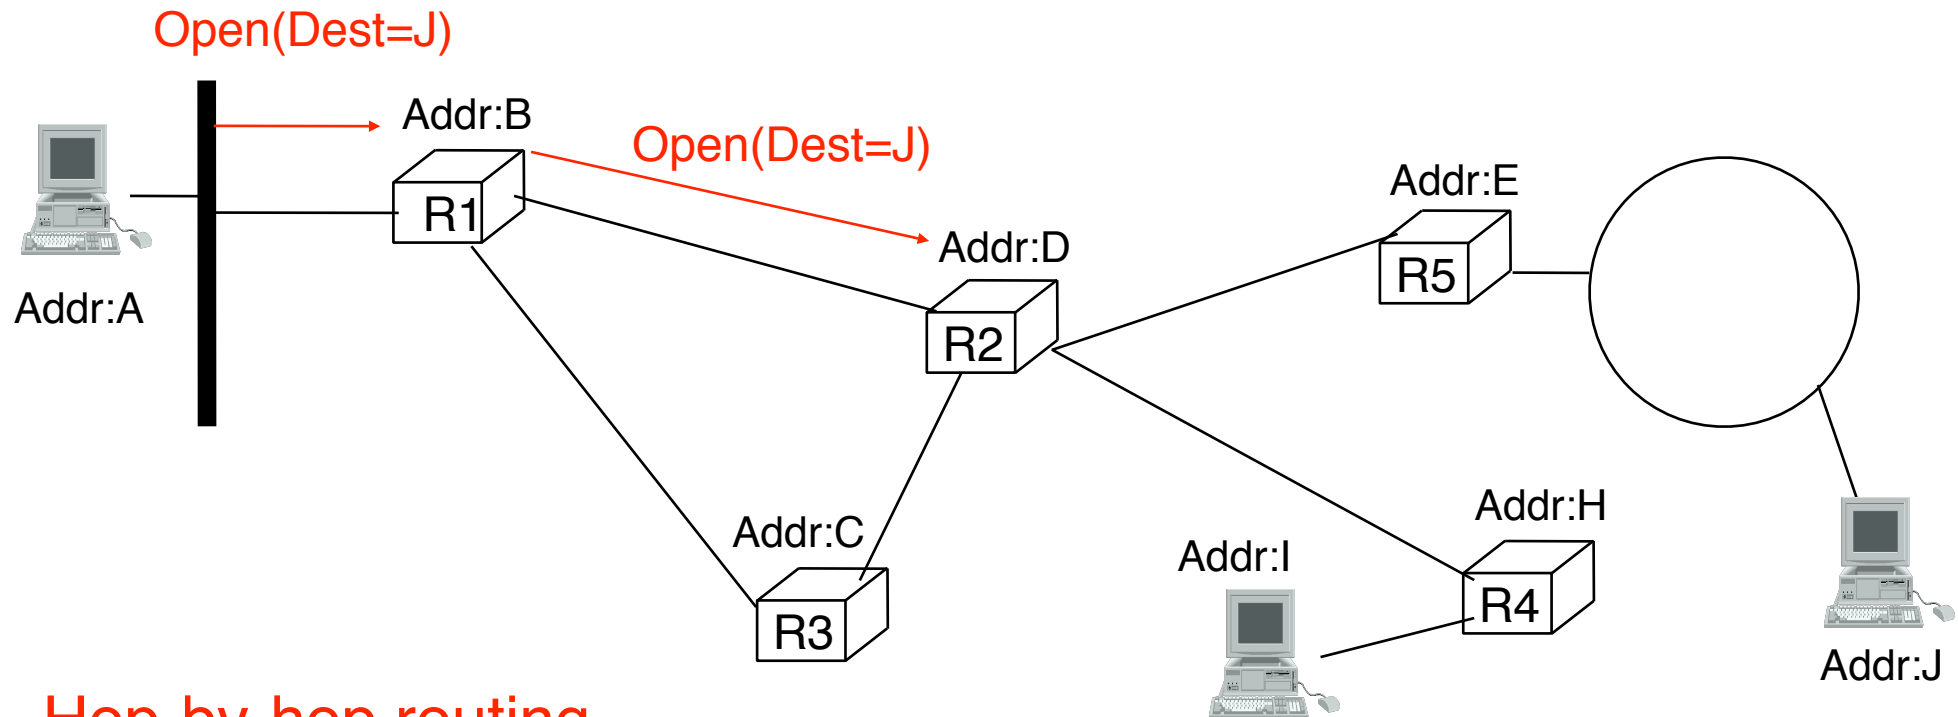
# Establishment of a virtual circuit



## Hop-by-hop routing

Each router consults its routing table to forward vc establishment

# Establishment of a virtual circuit

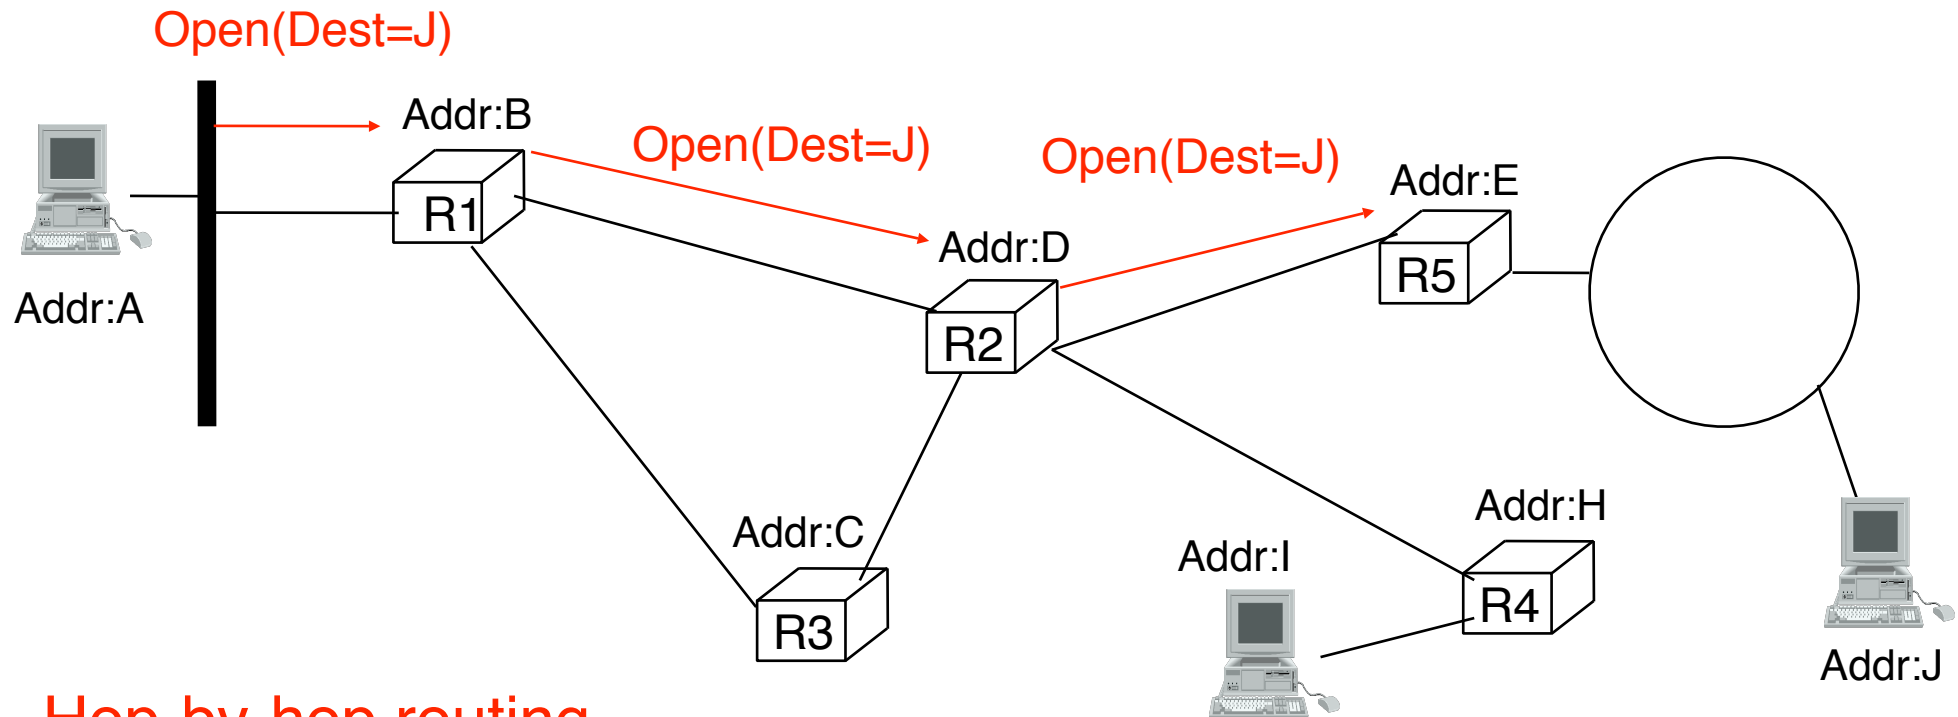


## Hop-by-hop routing

Each router consults its routing table to forward vc establishment



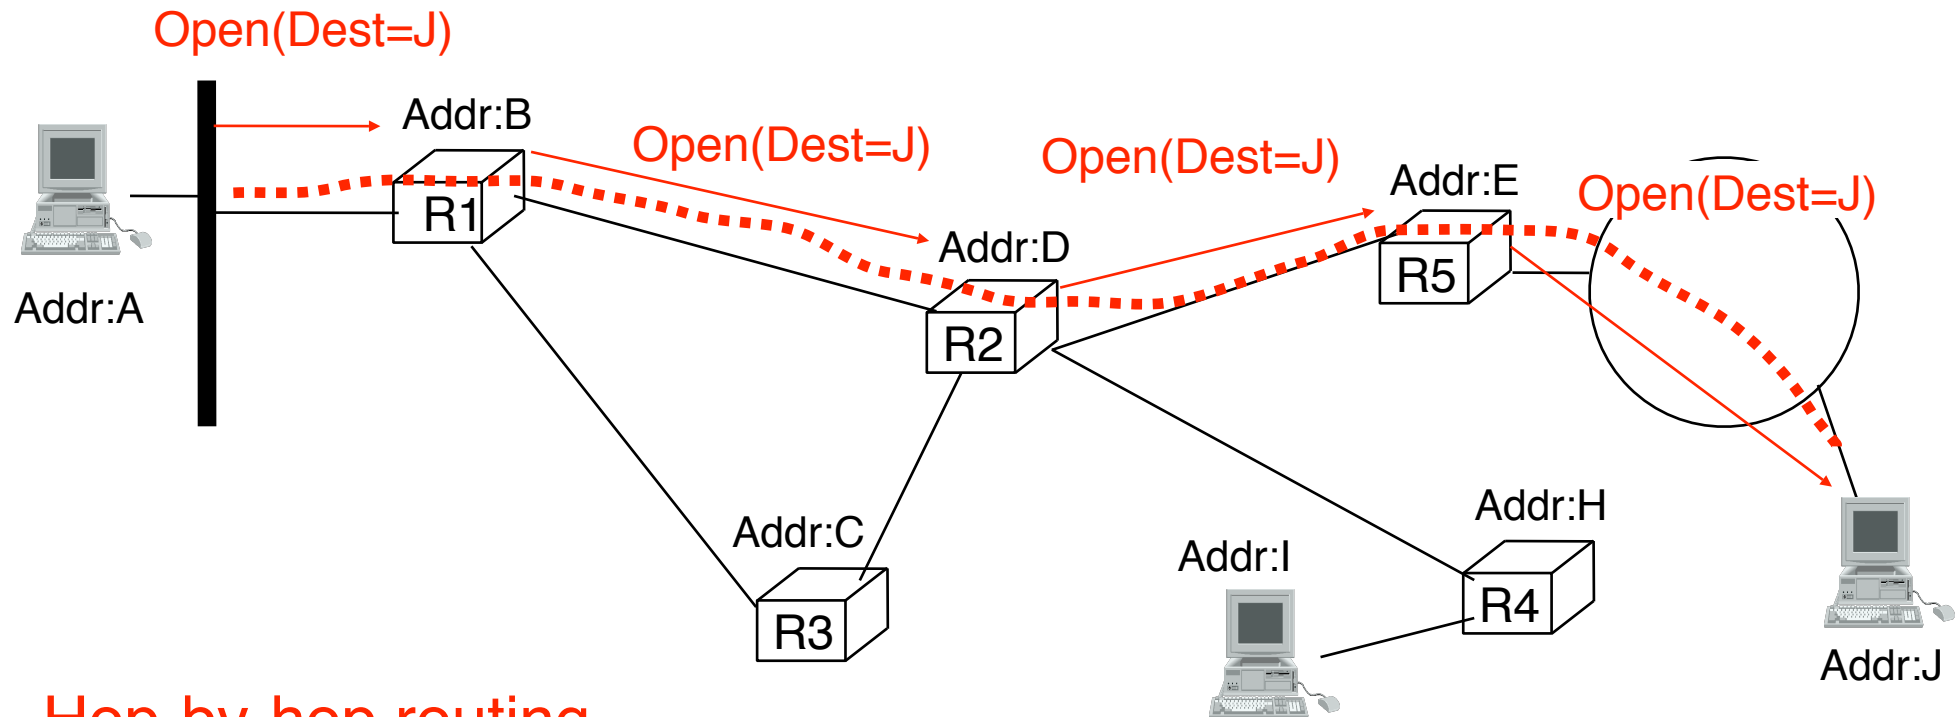
# Establishment of a virtual circuit



## Hop-by-hop routing

Each router consults its routing table to forward vc establishment

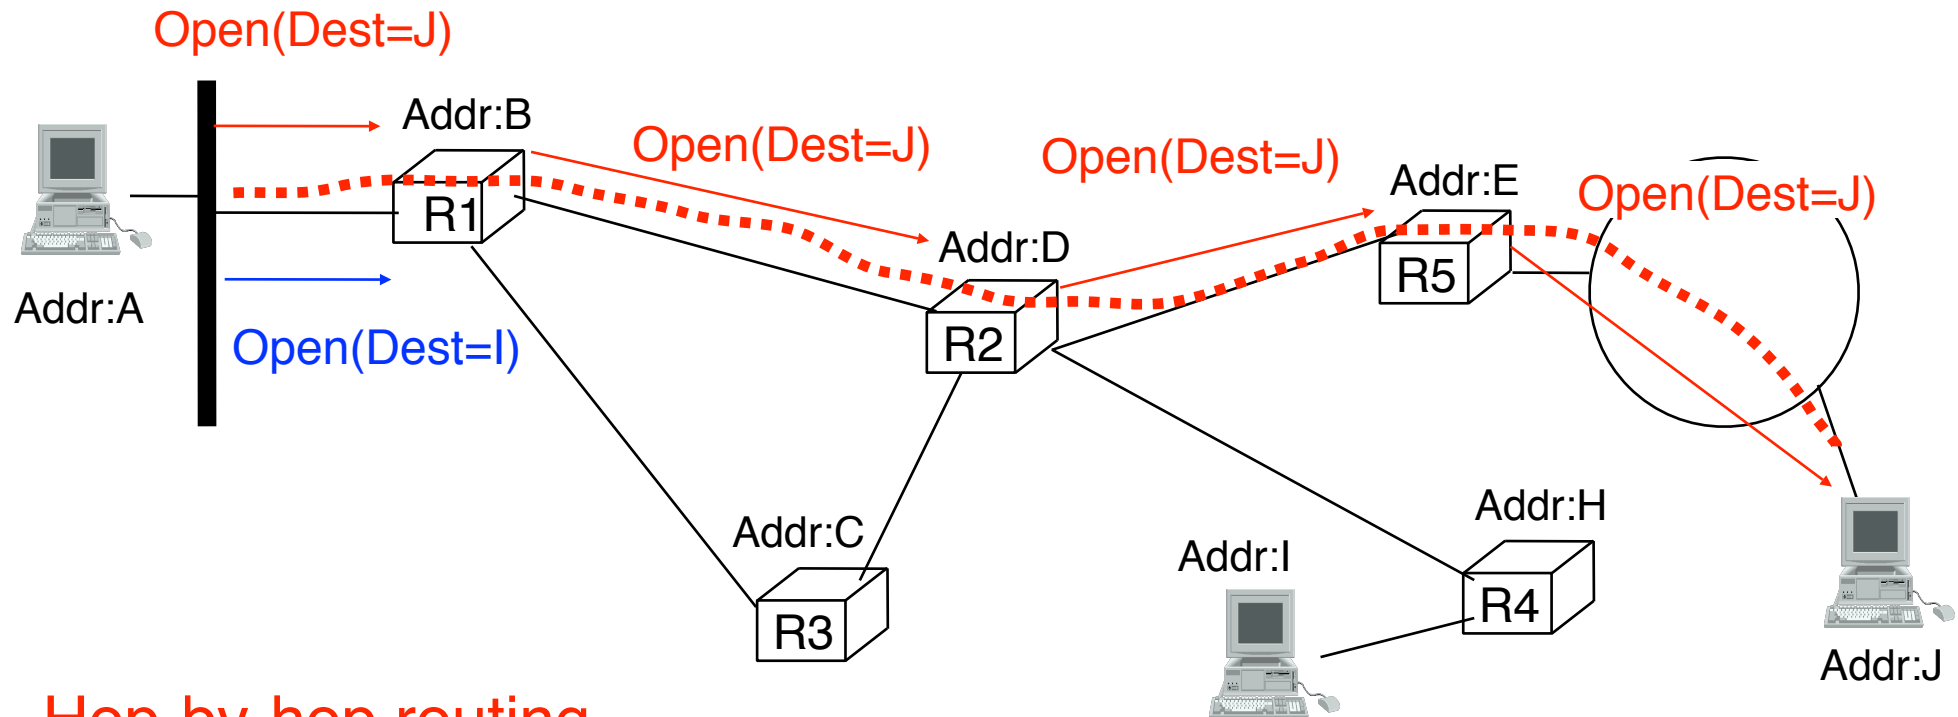
# Establishment of a virtual circuit



## Hop-by-hop routing

Each router consults its routing table to forward vc establishment

# Establishment of a virtual circuit



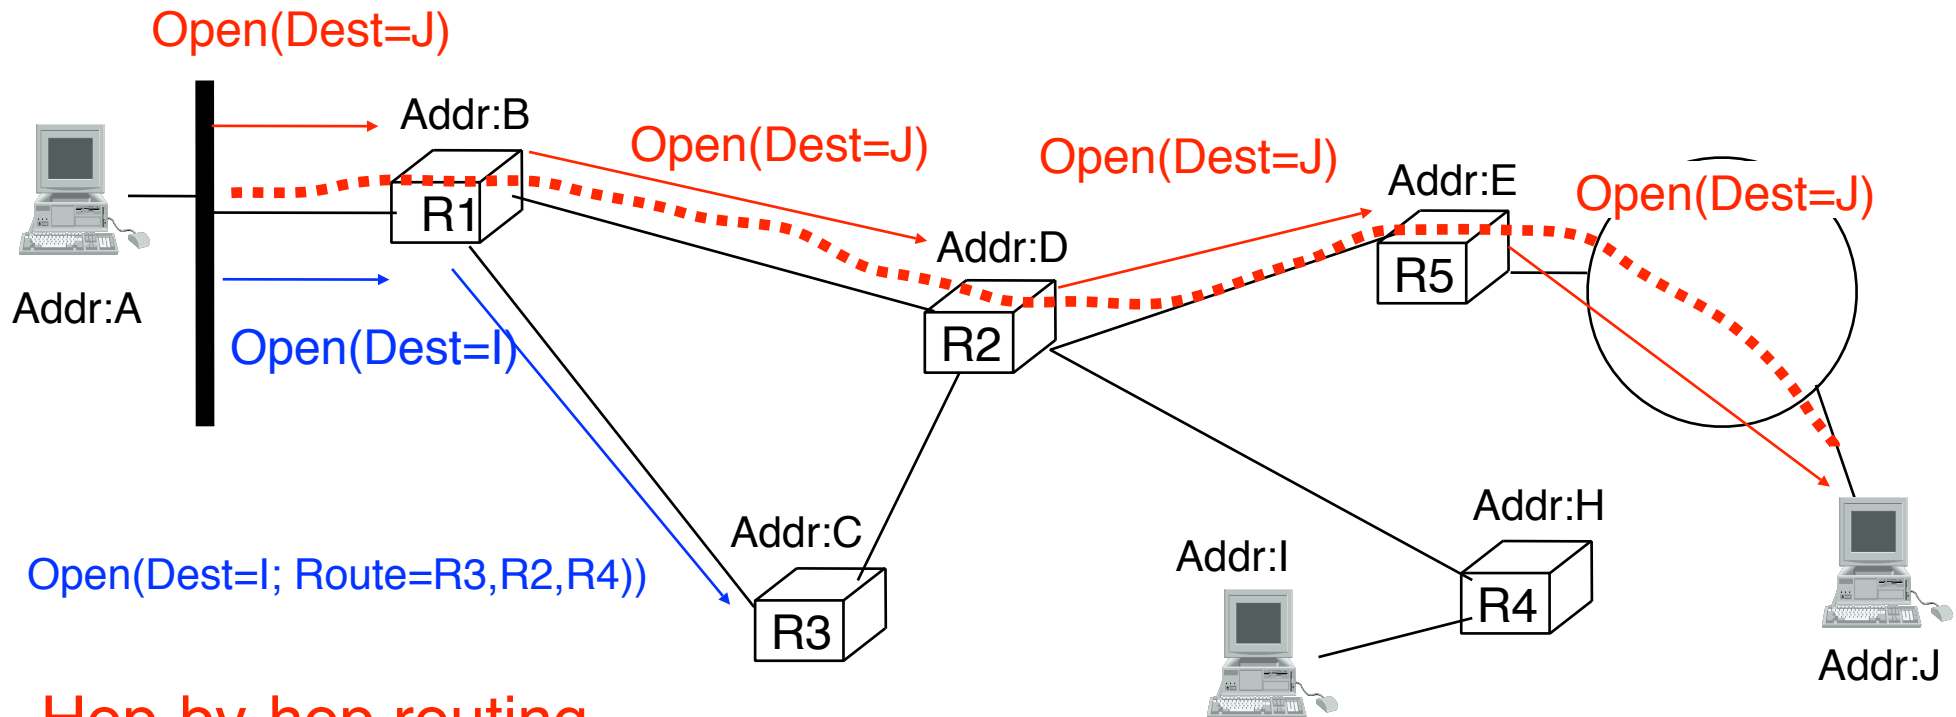
## Hop-by-hop routing

Each router consults its routing table to forward vc establishment

## Source routing/ explicit routing

Source (or first hop router) indicates in vc establishment packet the path to be followed

# Establishment of a virtual circuit



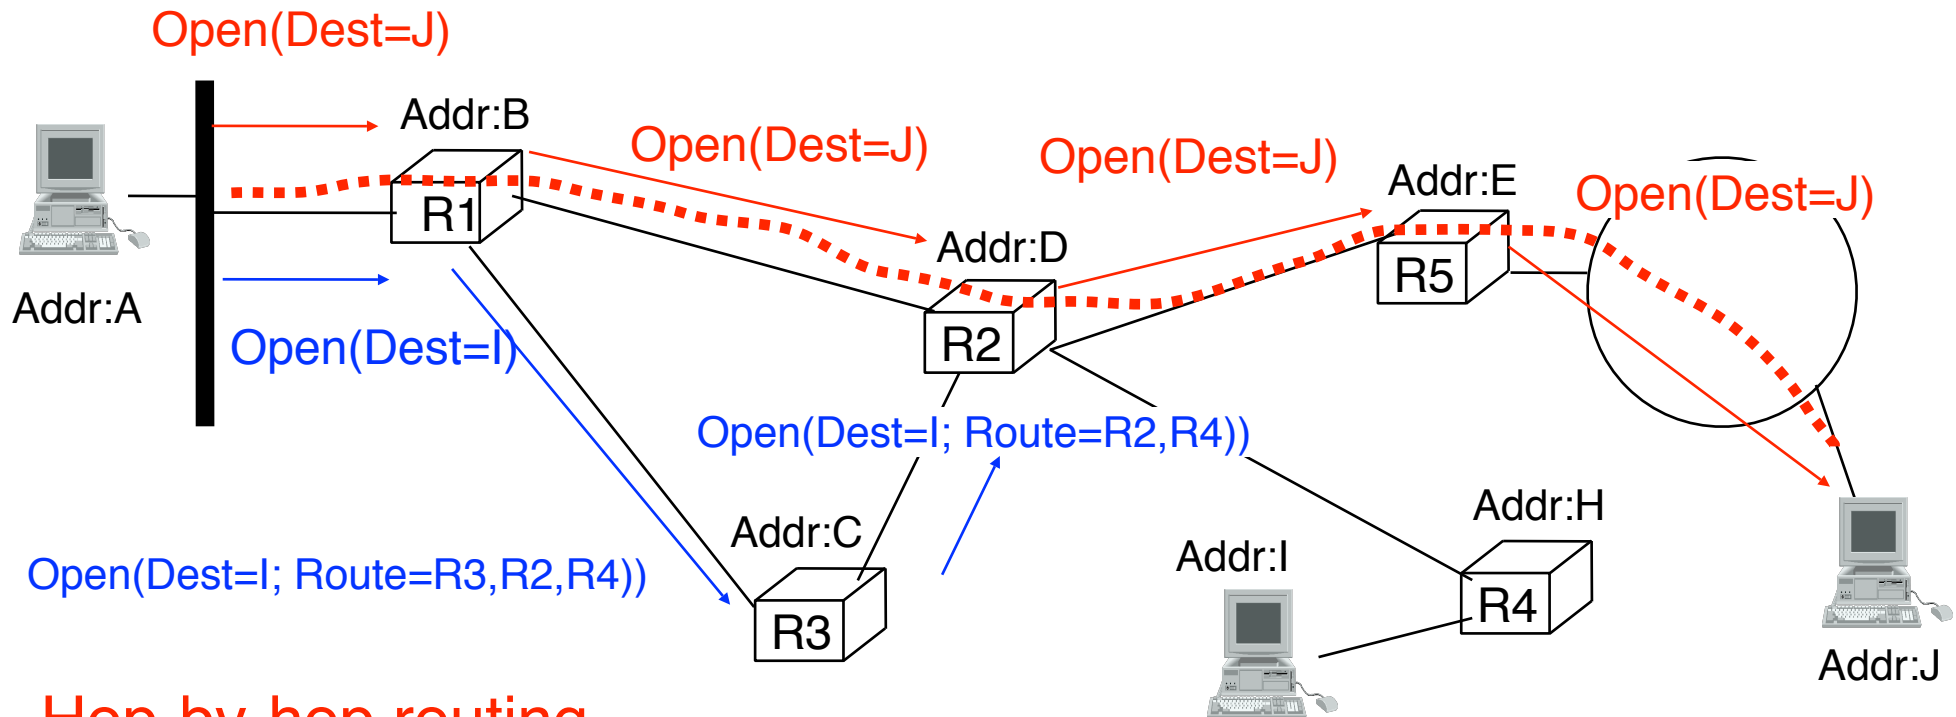
## Hop-by-hop routing

Each router consults its routing table to forward vc establishment

## Source routing/ explicit routing

Source (or first hop router) indicates in vc establishment packet the path to be followed

# Establishment of a virtual circuit



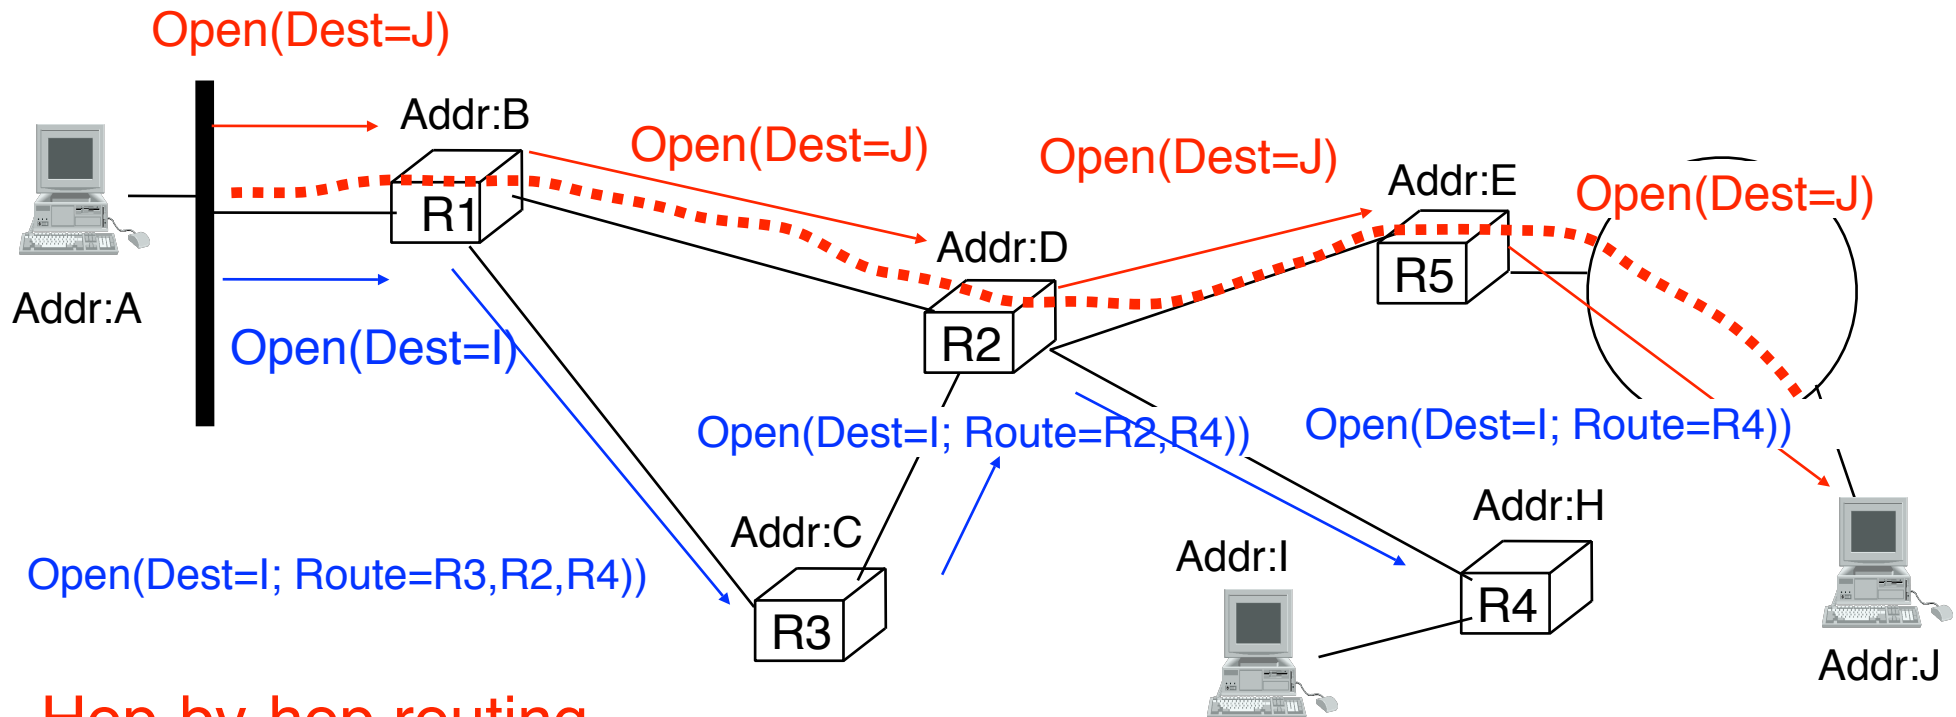
## Hop-by-hop routing

Each router consults its routing table to forward vc establishment

## Source routing/ explicit routing

Source (or first hop router) indicates in vc establishment packet the path to be followed

# Establishment of a virtual circuit



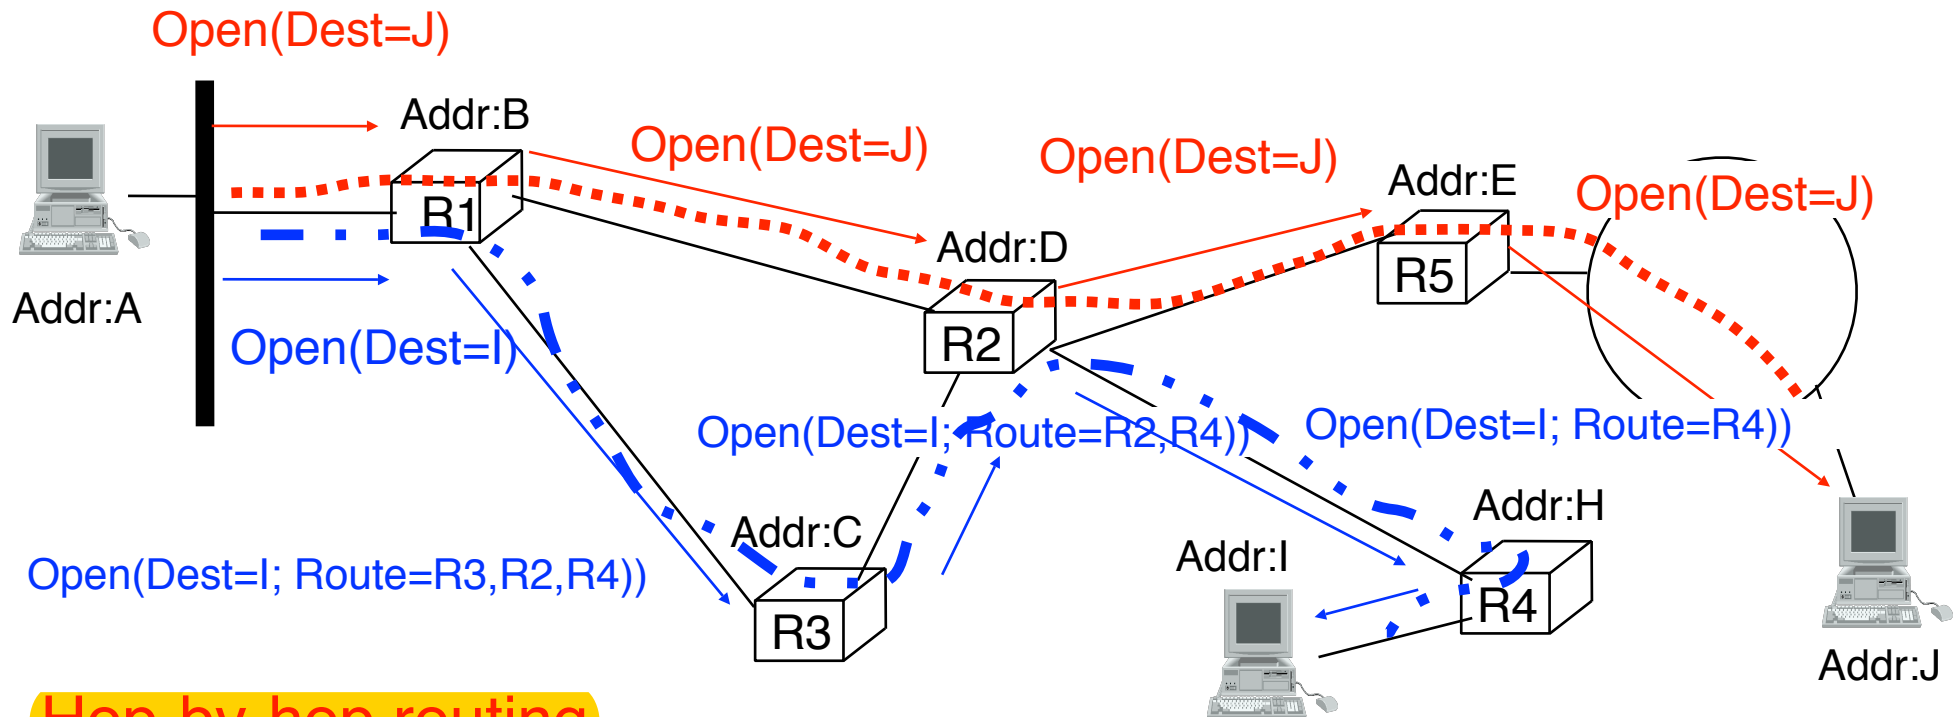
## Hop-by-hop routing

Each router consults its routing table to forward vc establishment

## Source routing/ explicit routing

Source (or first hop router) indicates in vc establishment packet the path to be followed

# Establishment of a virtual circuit



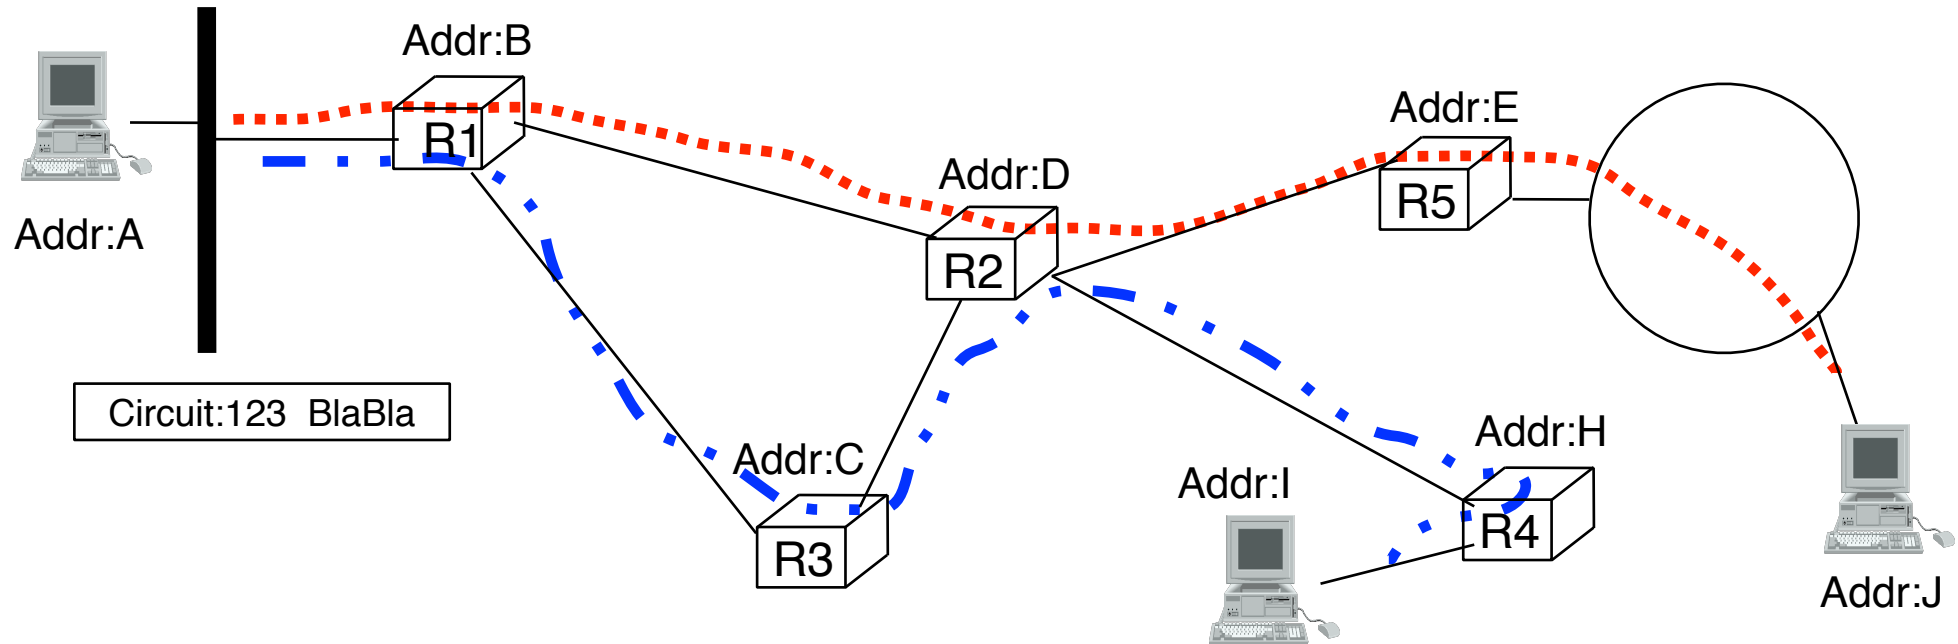
## Hop-by-hop routing

Each router consults its routing table  
to forward vc establishment

## Source routing/ explicit routing

Source (or first hop router) indicates in vc  
establishment packet the path to be followed

# Packet transmission



Packet contents  
virtual circuit identifier  
packet payload

What kind of virtual circuit identifier

Naive solution

unique identifier for all virtual circuits inside network

How to coordinate allocation of vc identifiers ?



# Packet transmission (2)

---

How unique should virtual circuits identifiers be ?

globally unique

unrealistic

unique inside a given network

then coordination among routers is necessary

unique on a given link

easier to manage, no coordination required, but

virtual circuit identifier may need to be changed from link to link

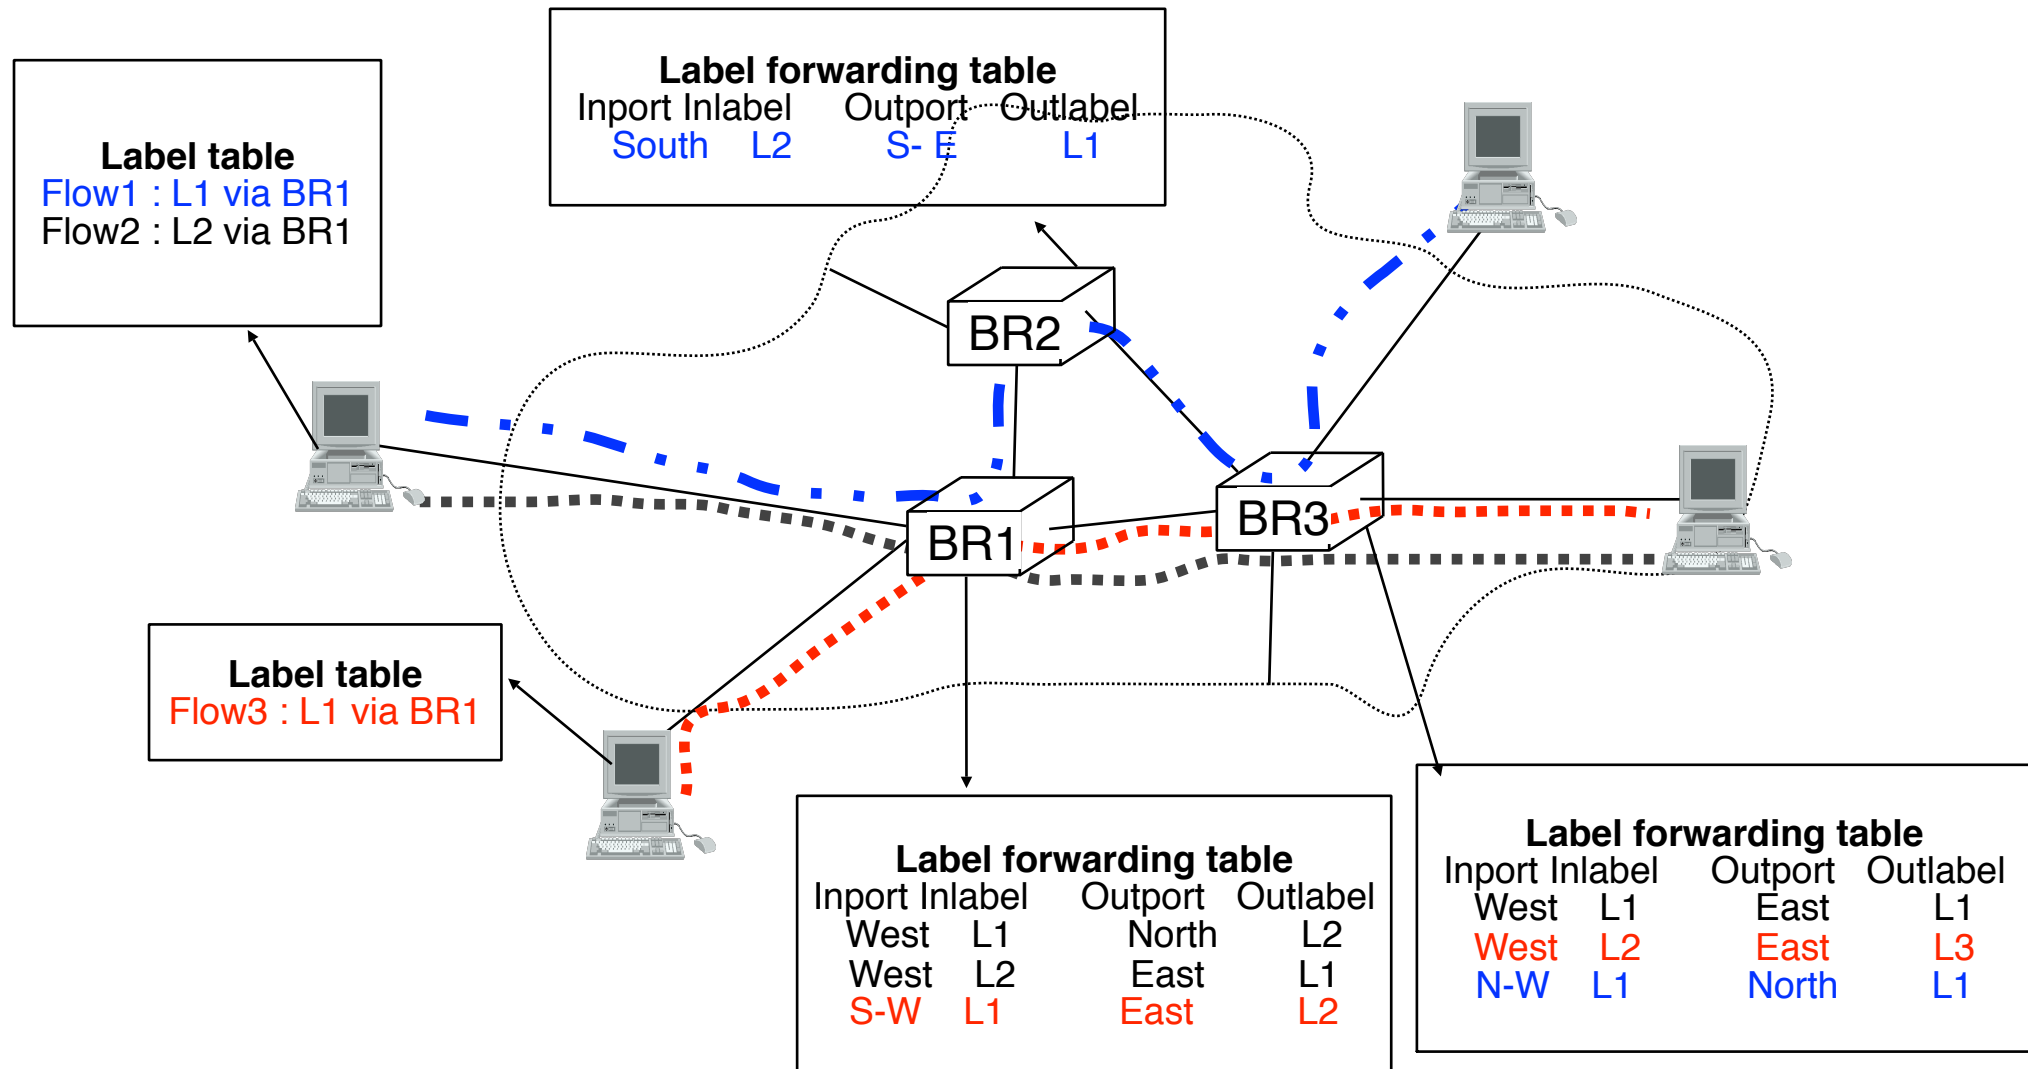
How to update the virtual circuit identifier of packets

All routers must contain a label forwarding table

this table is updated every time a virtual circuit is established

| Label forwarding table |         |         |          |
|------------------------|---------|---------|----------|
| Inport                 | Inlabel | Outport | Outlabel |
| West                   | L1      | East    | L1       |
| West                   | L2      | East    | L3       |
| N-W                    | L1      | North   | L1       |

# Virtual circuits : example



# Network layer

---

## Basics



### Routing

#### Static routing

Distance vector routing

Link state routing

## IP : Internet Protocol

## Routing in IP networks

# Routing and Forwarding

---

Main objective of network layer

transport packets from source to destination

Two mechanisms are used in network layer

forwarding

algorithm use by each router to determine on which interface each packet should be sent to reach its destination or follow its

virtual circuit


relies on the routing table maintained by each router

routing

algorithm (usually distributed) that distributes to all routers the information that allows them to build their routing tables

# Routing (2)

How to build the routing tables of each router ?



## Principle

Include in the routing table of each router the path to allow it to reach each destination

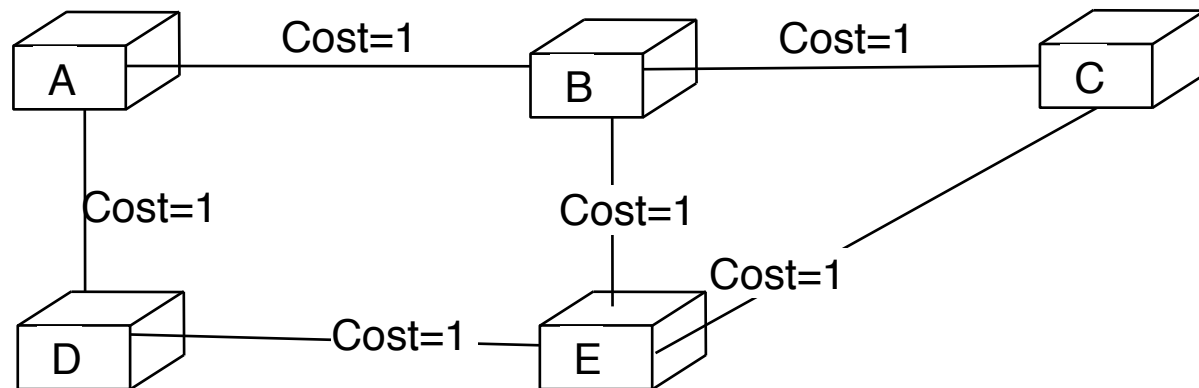
Which path to be included in the routing table

From A to C ?

From D to B

# Selection of the shortest paths

---



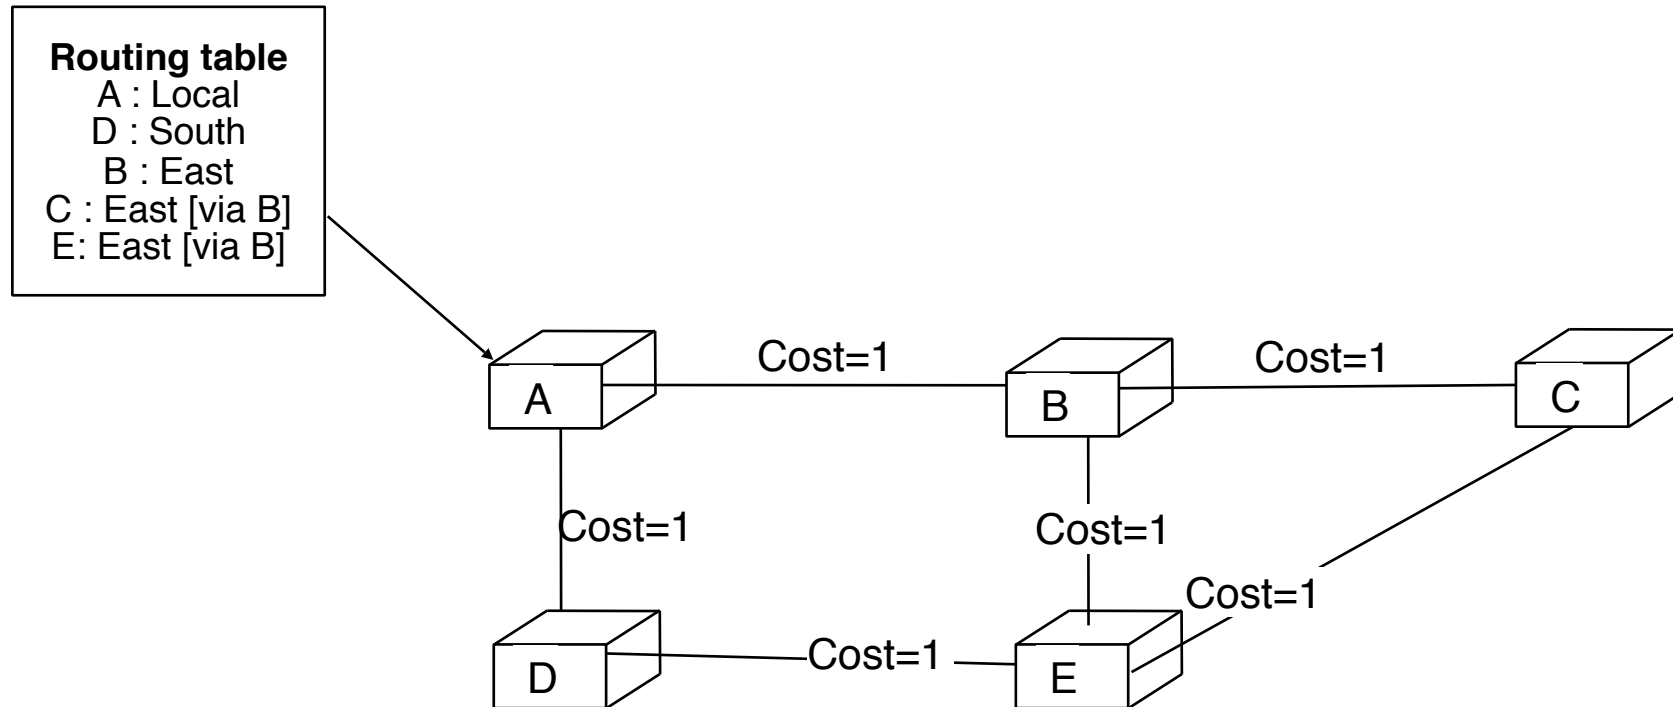
## Principle

Associate a weight/cost to each link

Each router chooses the lowest cost path

How to ensure that the routing tables of all routers are coherent ?

# Selection of the shortest paths



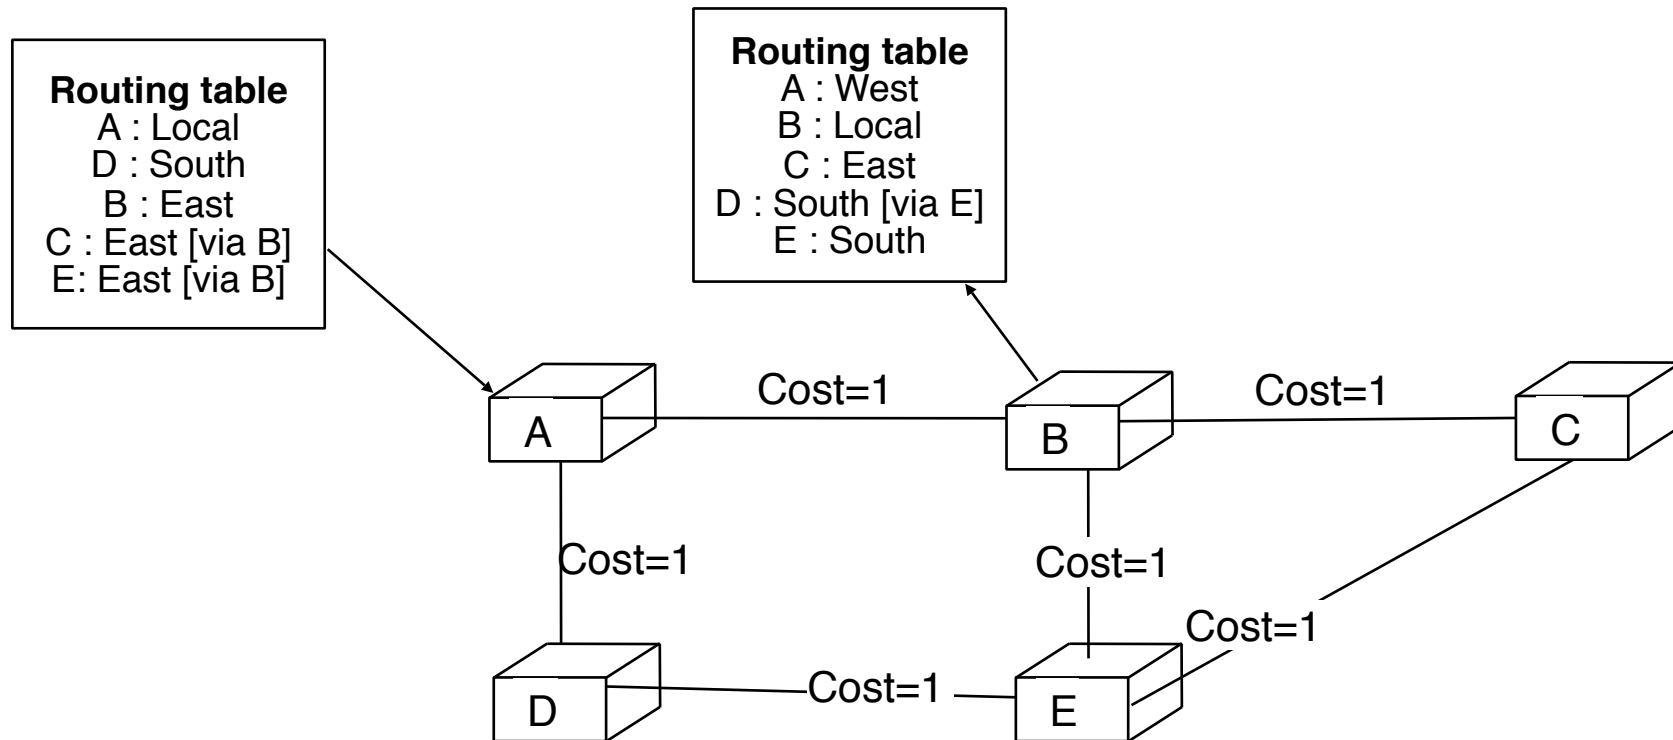
## Principle

Associate a weight/cost to each link

Each router chooses the lowest cost path

How to ensure that the routing tables of all routers are coherent ?

# Selection of the shortest paths



## Principle

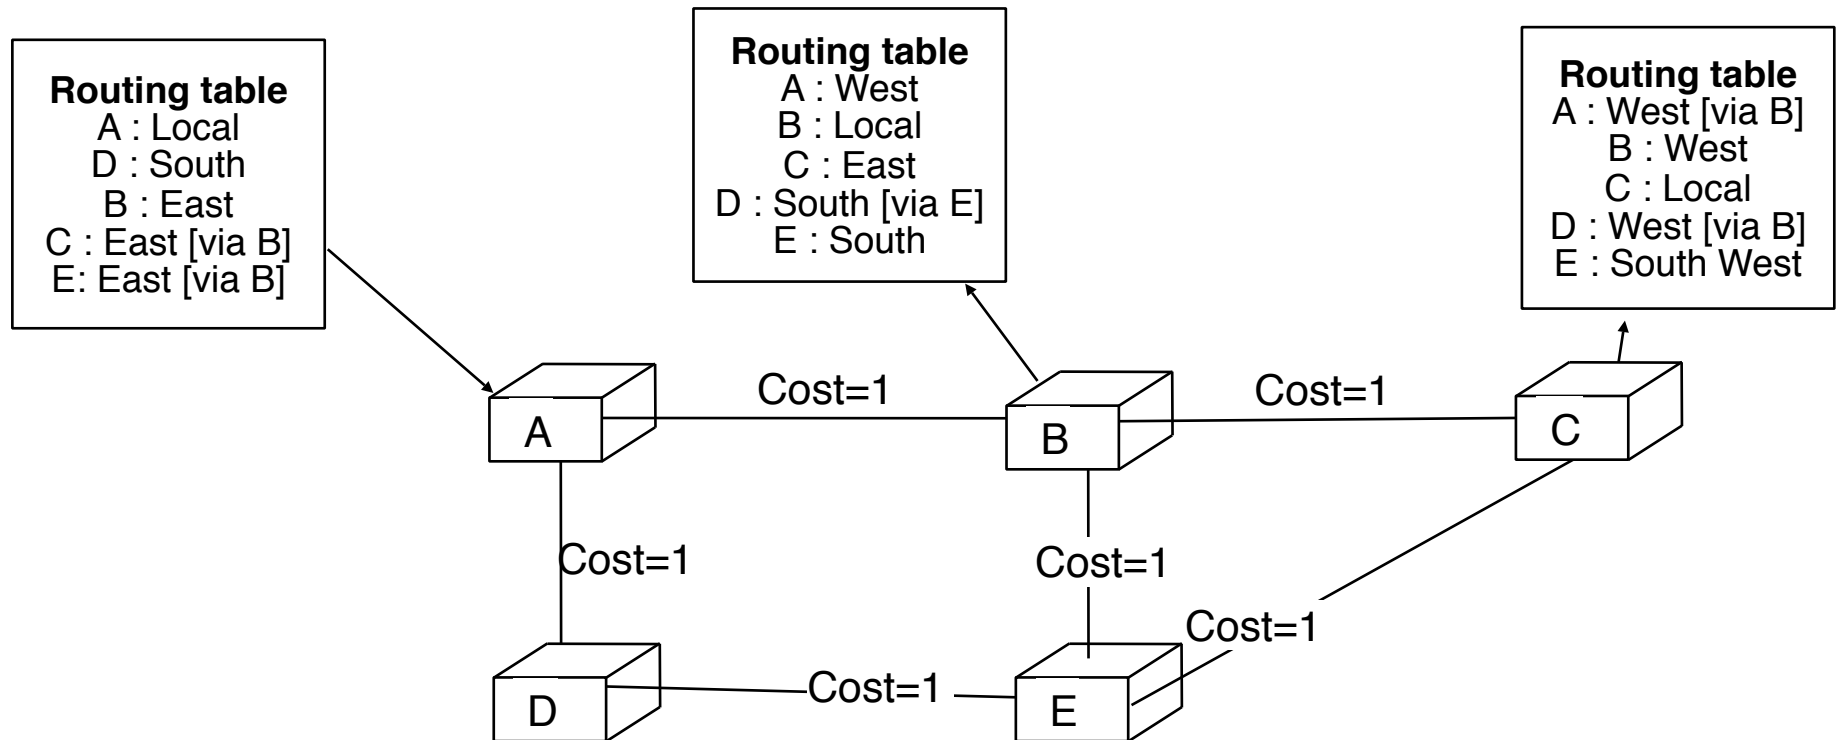
Associate a weight/cost to each link

Each router chooses the lowest cost path

How to ensure that the routing tables of all routers are coherent ?



# Selection of the shortest paths



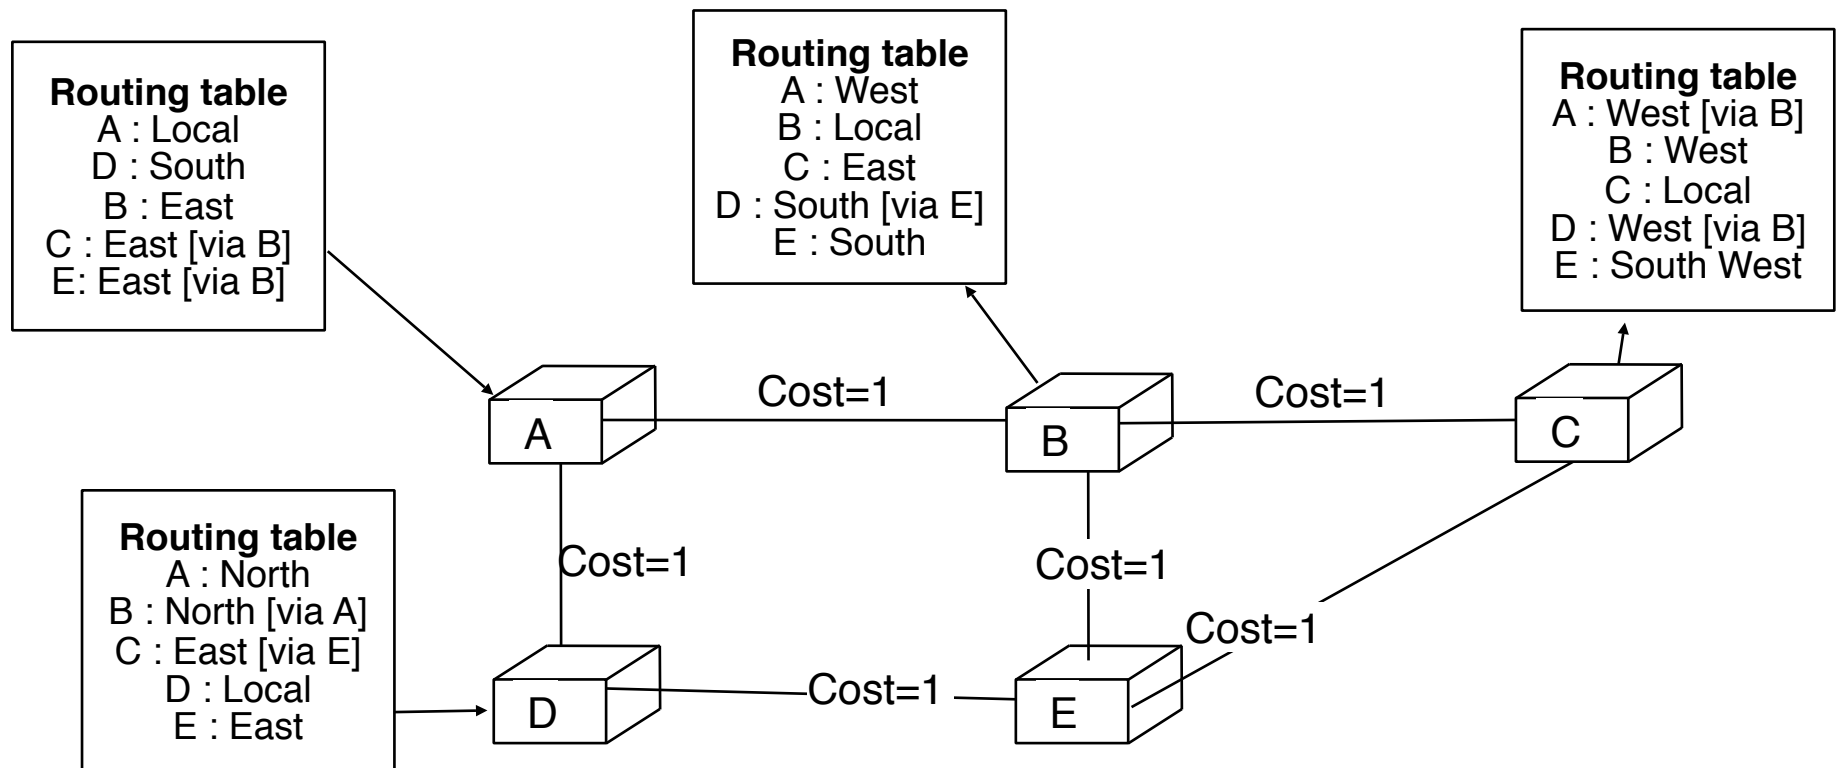
## Principle

Associate a weight/cost to each link

Each router chooses the lowest cost path

How to ensure that the routing tables of all routers are coherent ?

# Selection of the shortest paths



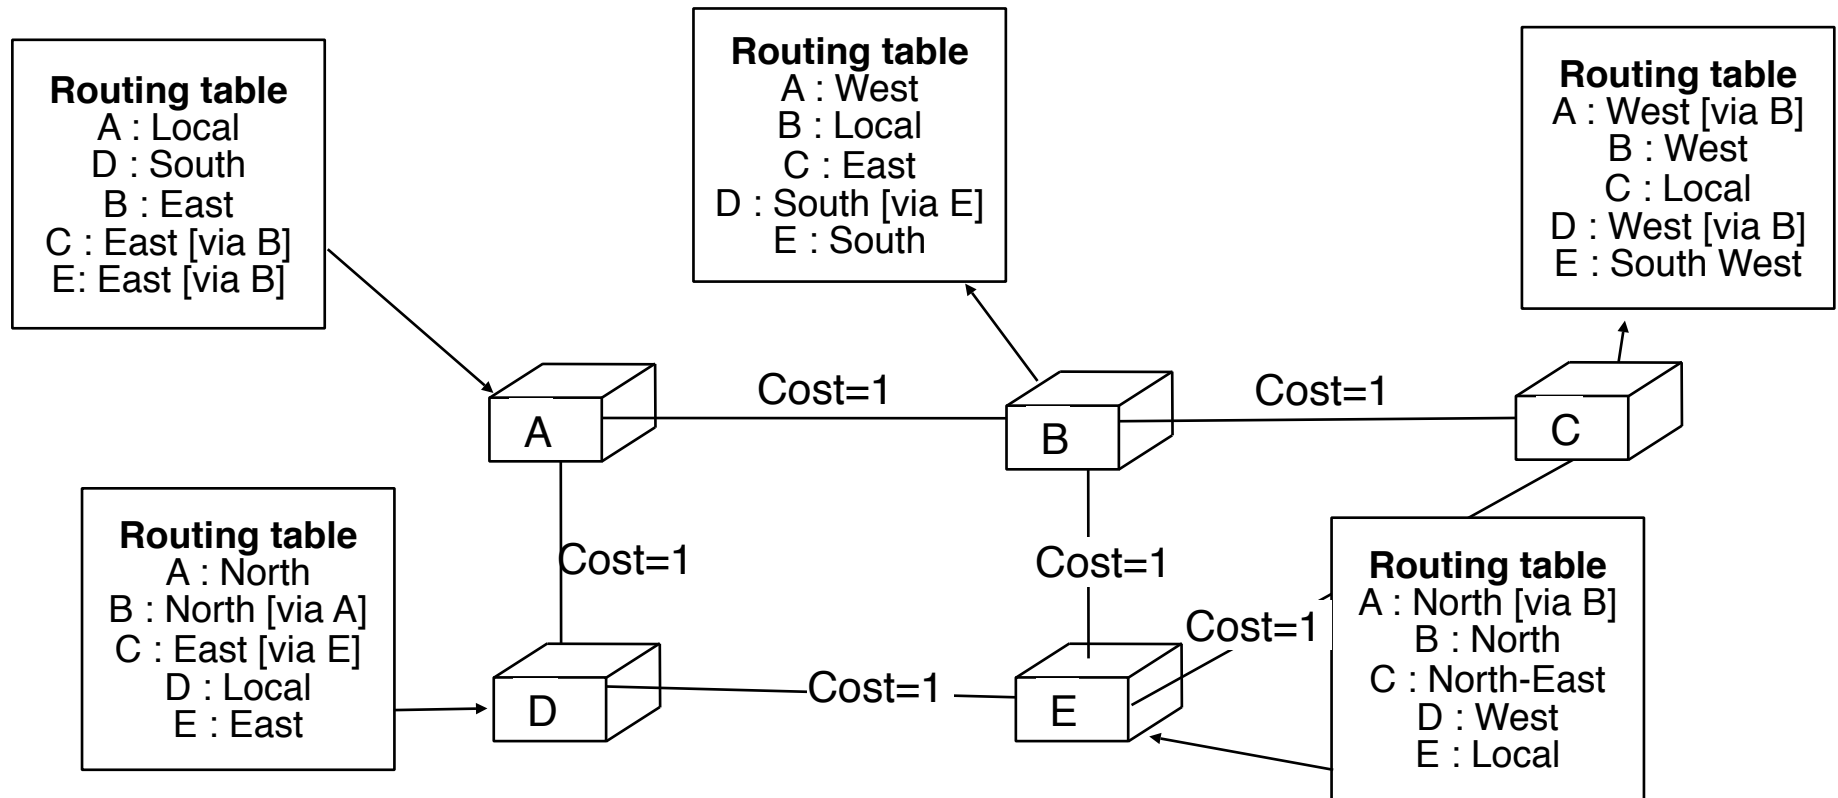
## Principle

Associate a weight/cost to each link

Each router chooses the lowest cost path

How to ensure that the routing tables of all routers are coherent ?

# Selection of the shortest paths



## Principle

Associate a weight/cost to each link

Each router chooses the lowest cost path

How to ensure that the routing tables of all routers are coherent ?

# Static Routing

---

## Principle

Network manager or network management station computes all routing tables and downloads them on all routers

How to compute routing tables ?

- shortest path algorithms

- more complex algorithms to provide load balancing or traffic engineering

## Advantages of static routing

- Easy to use in a small network
- routing tables can be optimised

## Drawbacks of static routing

- does not adapt dynamically to network load
- how to deal with link and router failures ?

# Dynamic or distributed routing

---

## Principle

routers exchange messages and use a distributed algorithm to build their routing tables  
used in almost all networks

## Advantages

can easily adapt routing tables to events

## Drawbacks

more complex to implement than static routing

## Most common distributed routing methods

Distance vector routing

Link state routing

# Network layer

---

## Basics

### Routing

Static routing

→ Distance vector routing

Link state routing

## IP : Internet Protocol


## Routing in IP networks

# Distance vector routing

## Basic principles

Configuration of each router

Cost of each link



When it boots, a router only knows itself

Each router sends periodically to all its neighbours a vector that contains for each destination that it knows

1. Destination address
2. Distance between transmitting router and destination
  - distance vector is a summary of the router's routing table

Each router will update its routing table based on the information received from its neighbours

# Distance vector routing (2)

---

## Routing table *maintained by router*

For each destination  $d$  inside routing table

$R[d].cost$  = total cost of shortest path to reach  $d$

$R[d].link$  = outgoing link used to reach  $d$  via shortest path

## Distance vector *sent to neighbours*

For each destination  $d$

$V[d].cost$  = total cost of shortest path to reach  $d$

Every  $N$  seconds:

```
Vector=null;
```

```
for each destination=d in R[]
```

```
{
```

```
  Vector=Vector+Pair(d,R[d].cost);
```

```
}
```

```
for each interface
```

```
{
```

```
  Send(Vector);
```

```
}
```



# Distance vector routing (3)


---

## Processing of received distance vectors

```
Received(Vector V[], link l)
{ /* received vector from link l */
  for each destination=d in V[]
  {
    if (d isin R[])
    { if ((V[d].cost+l.cost) < R[d].cost)
      { /* shorter path */
        R[d].cost=V[d].cost+l.cost;
        R[d].link=l;
      }
    }
    else
    { /* new route */
      R[d].cost=V[d].cost+l.cost;
      R[d].link=l;
    }
  }
}
```

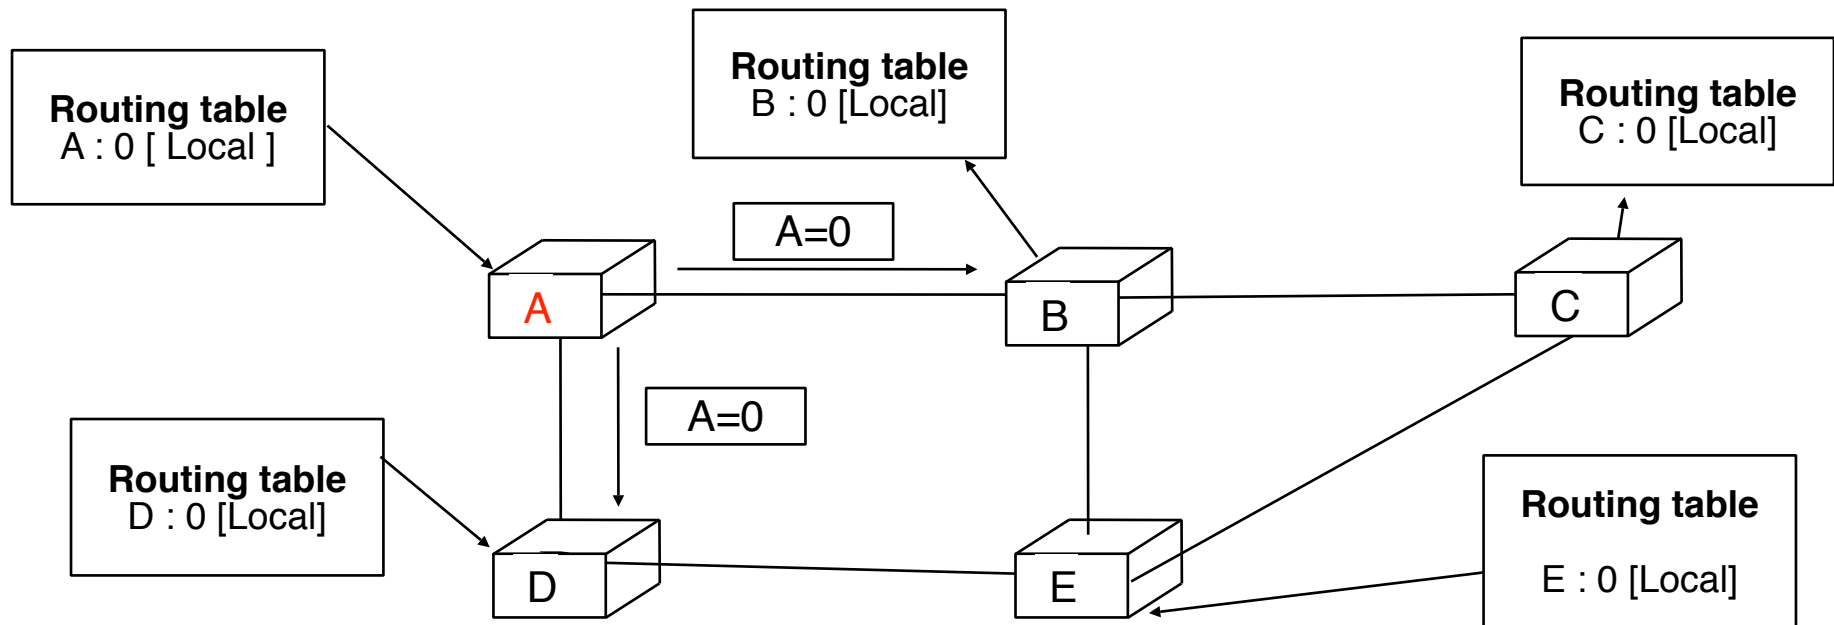
# Distance vectors example

All links have a unit cost



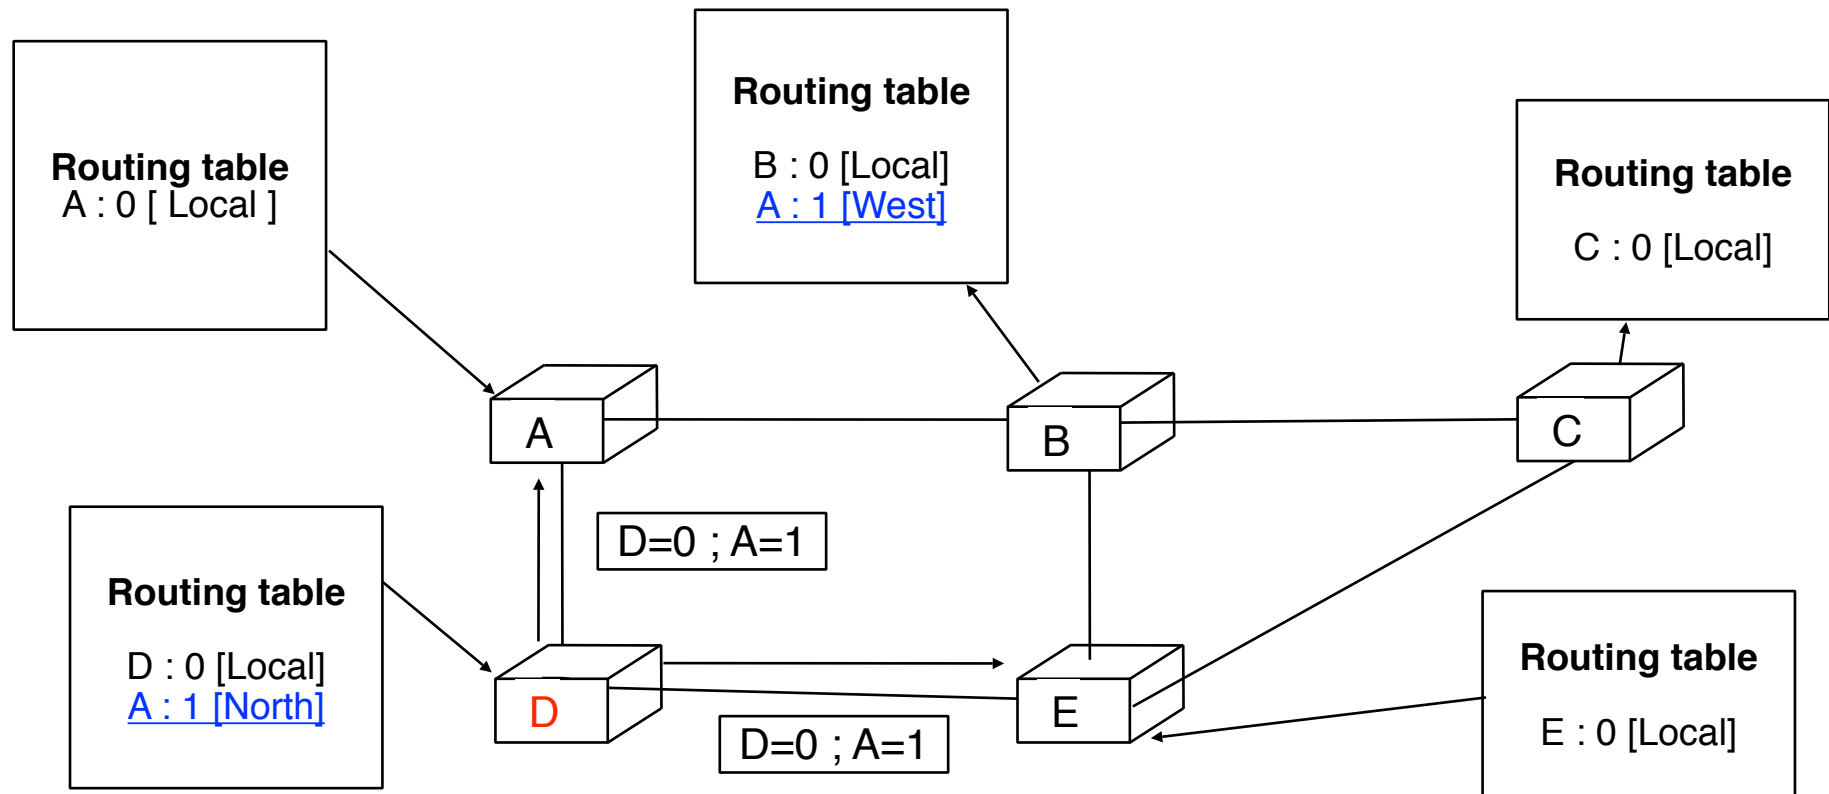
# Distance vectors example

All links have a unit cost

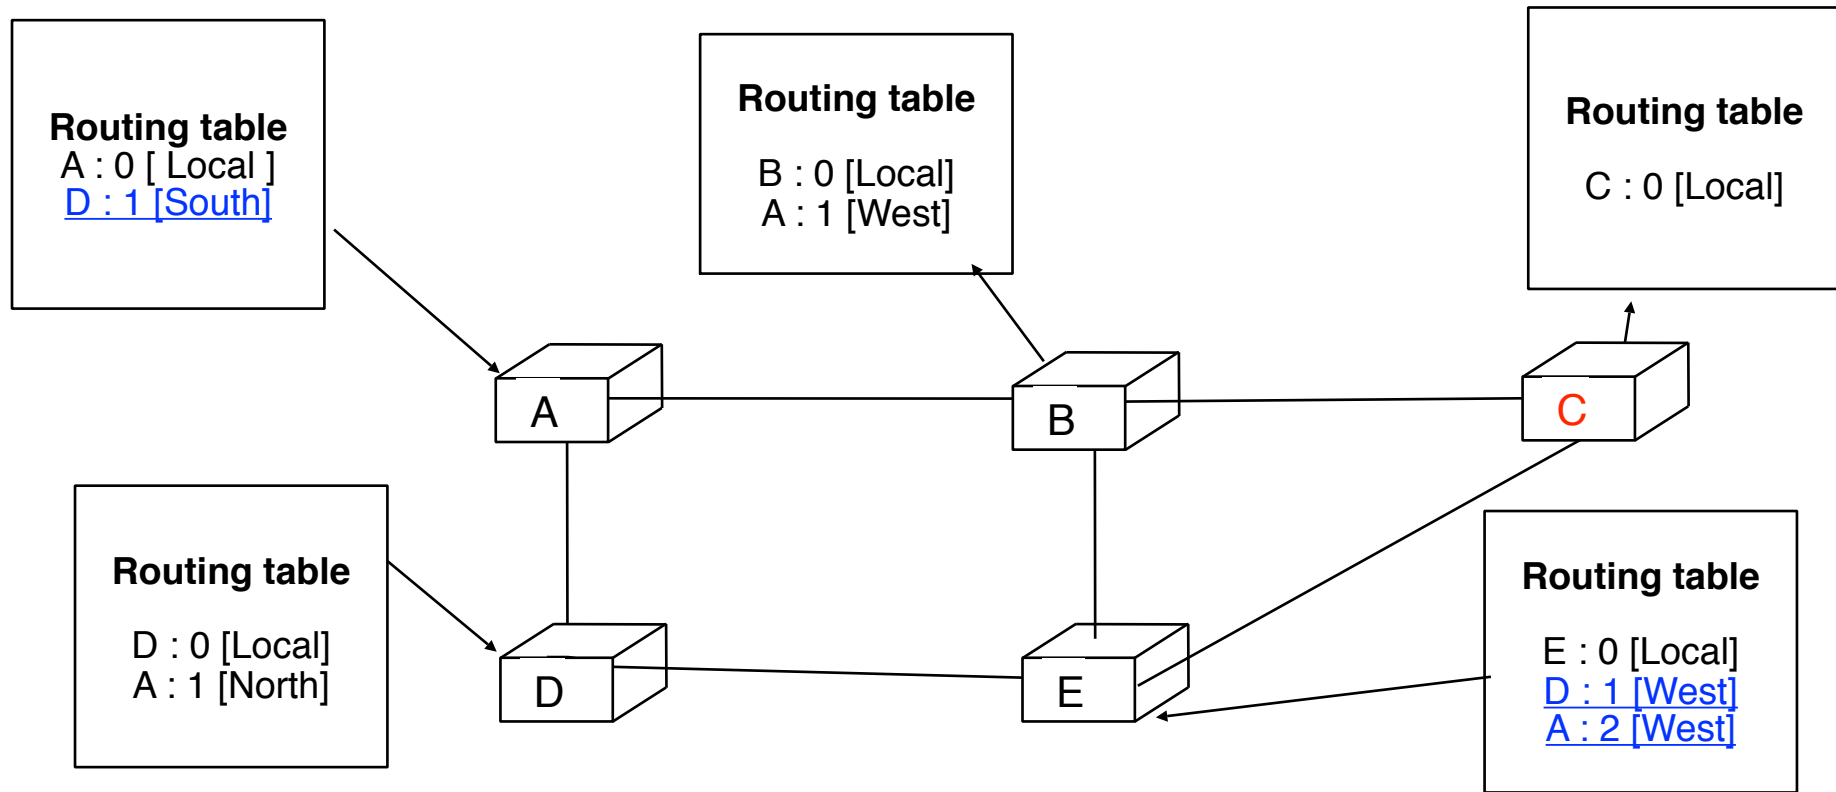


When a router boots, it only knows itself. Its distance vector thus contains only its address

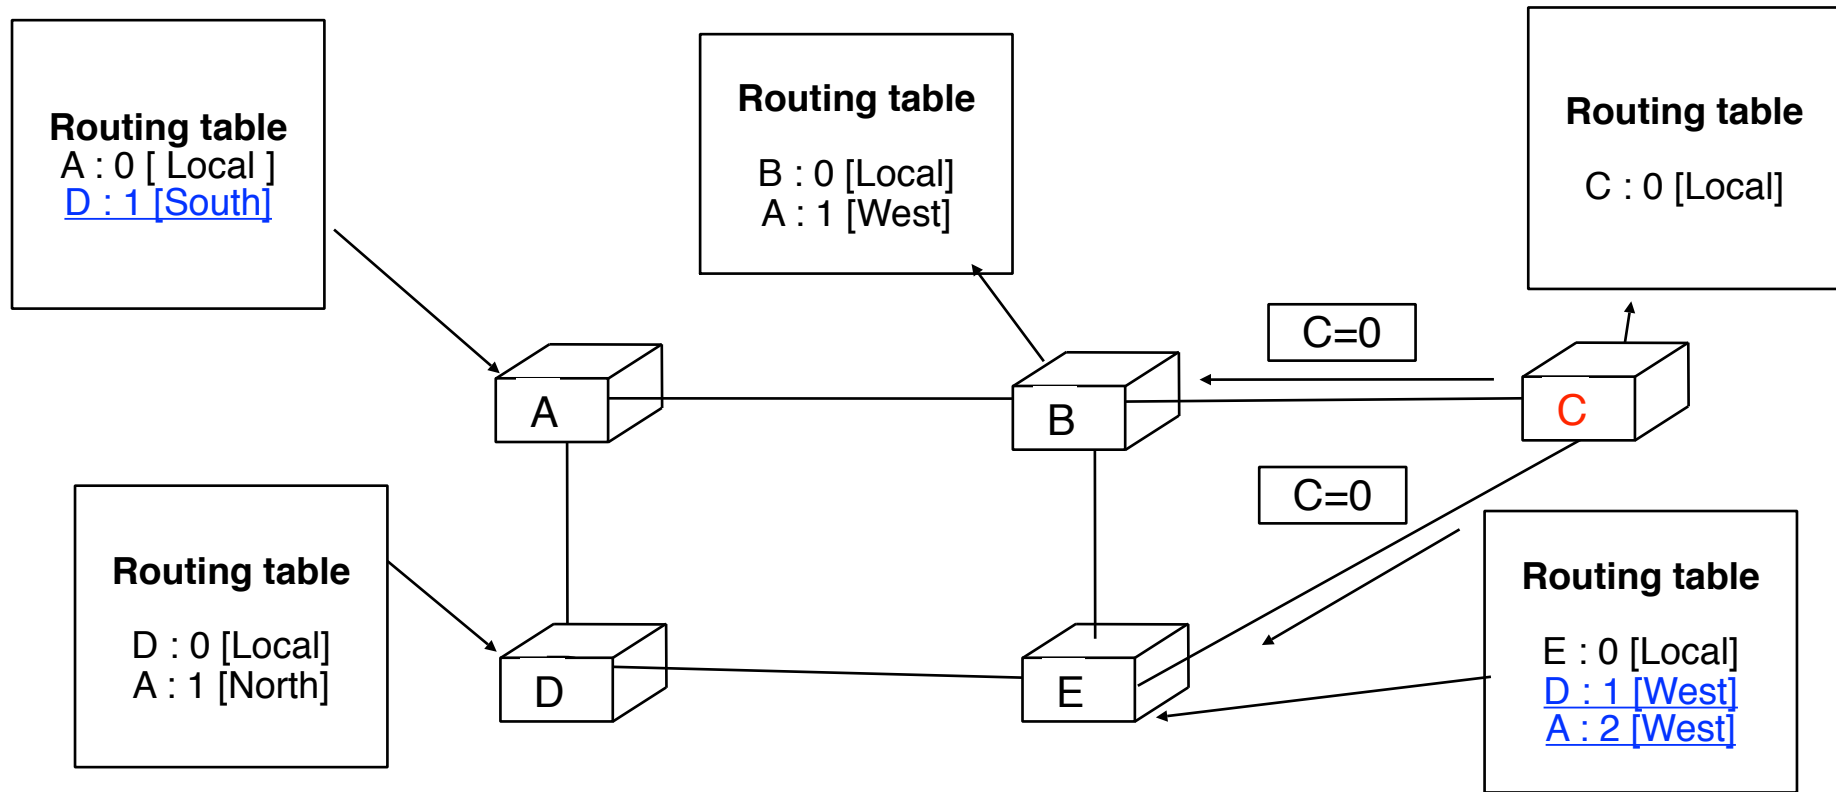
# Distance vectors example (2)



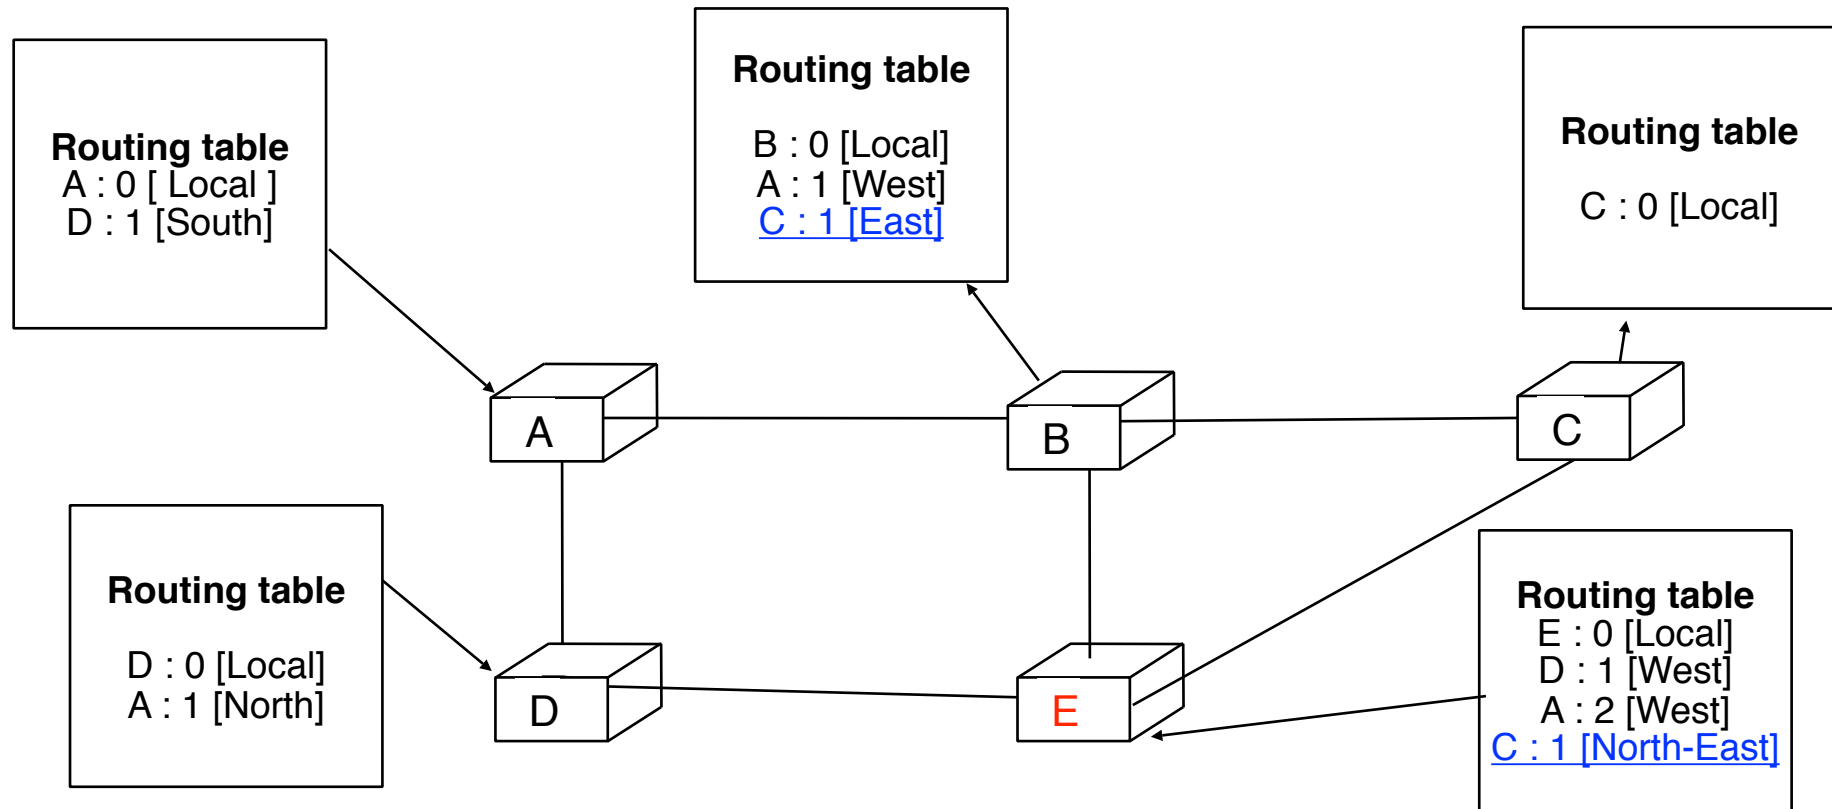
# Distance vectors example (3)



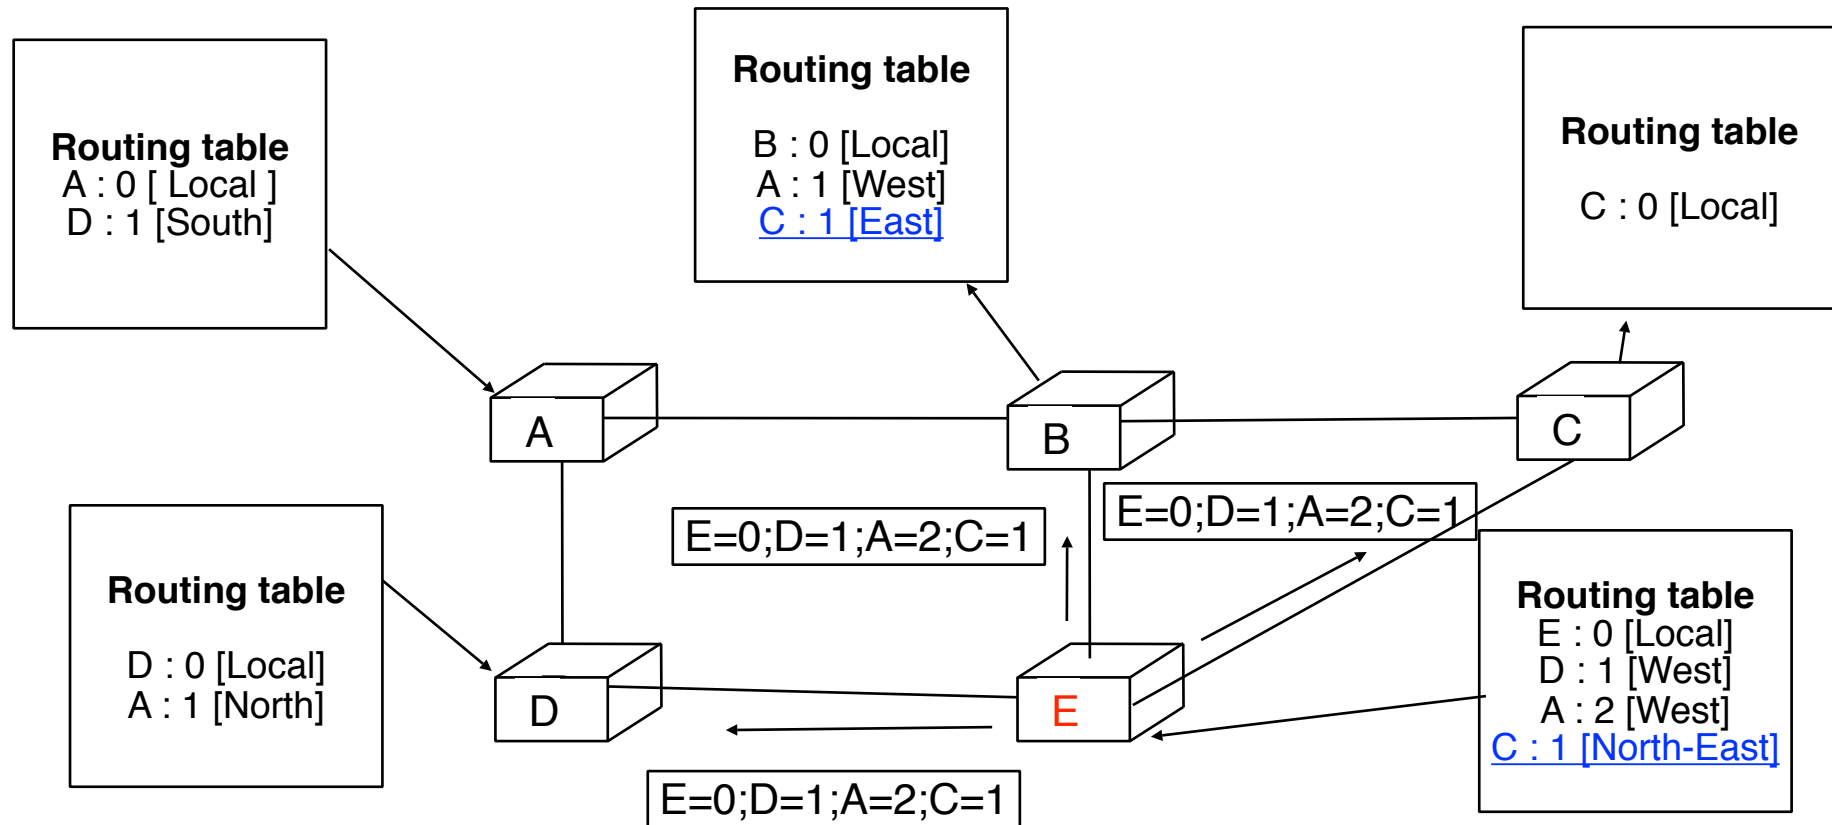
# Distance vectors example (3)



# Distance vectors example (4)

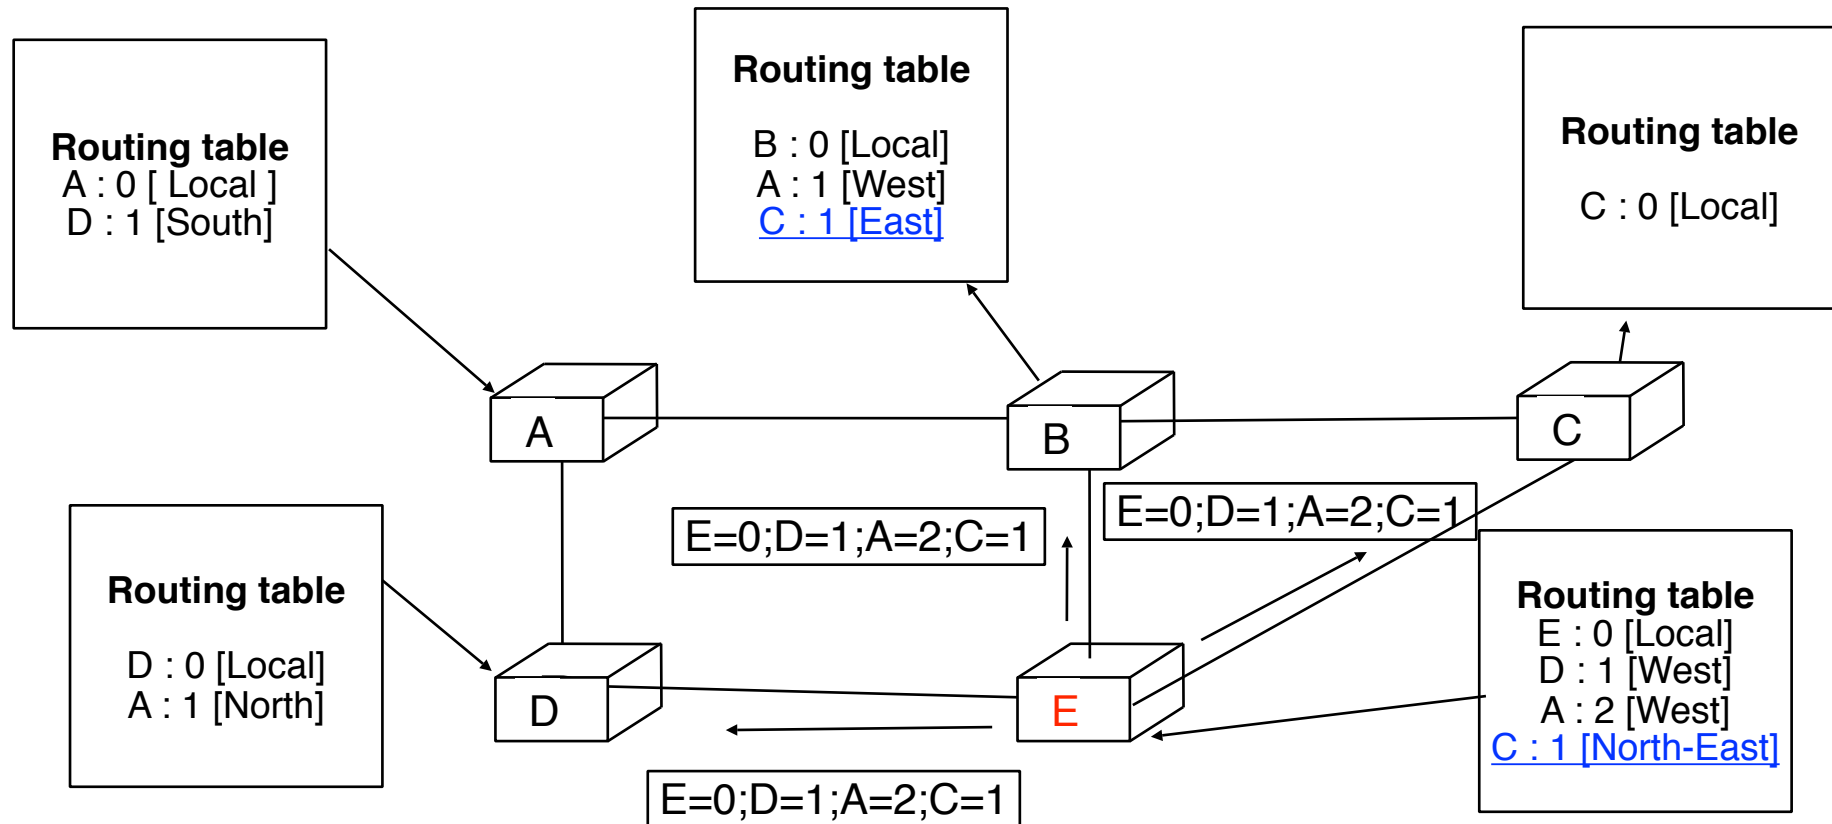


# Distance vectors example (4)





# Distance vectors example (4)



**Reception of distance vector on B**

New route to reach E and D, longer route for A

**Reception of distance vector on C**

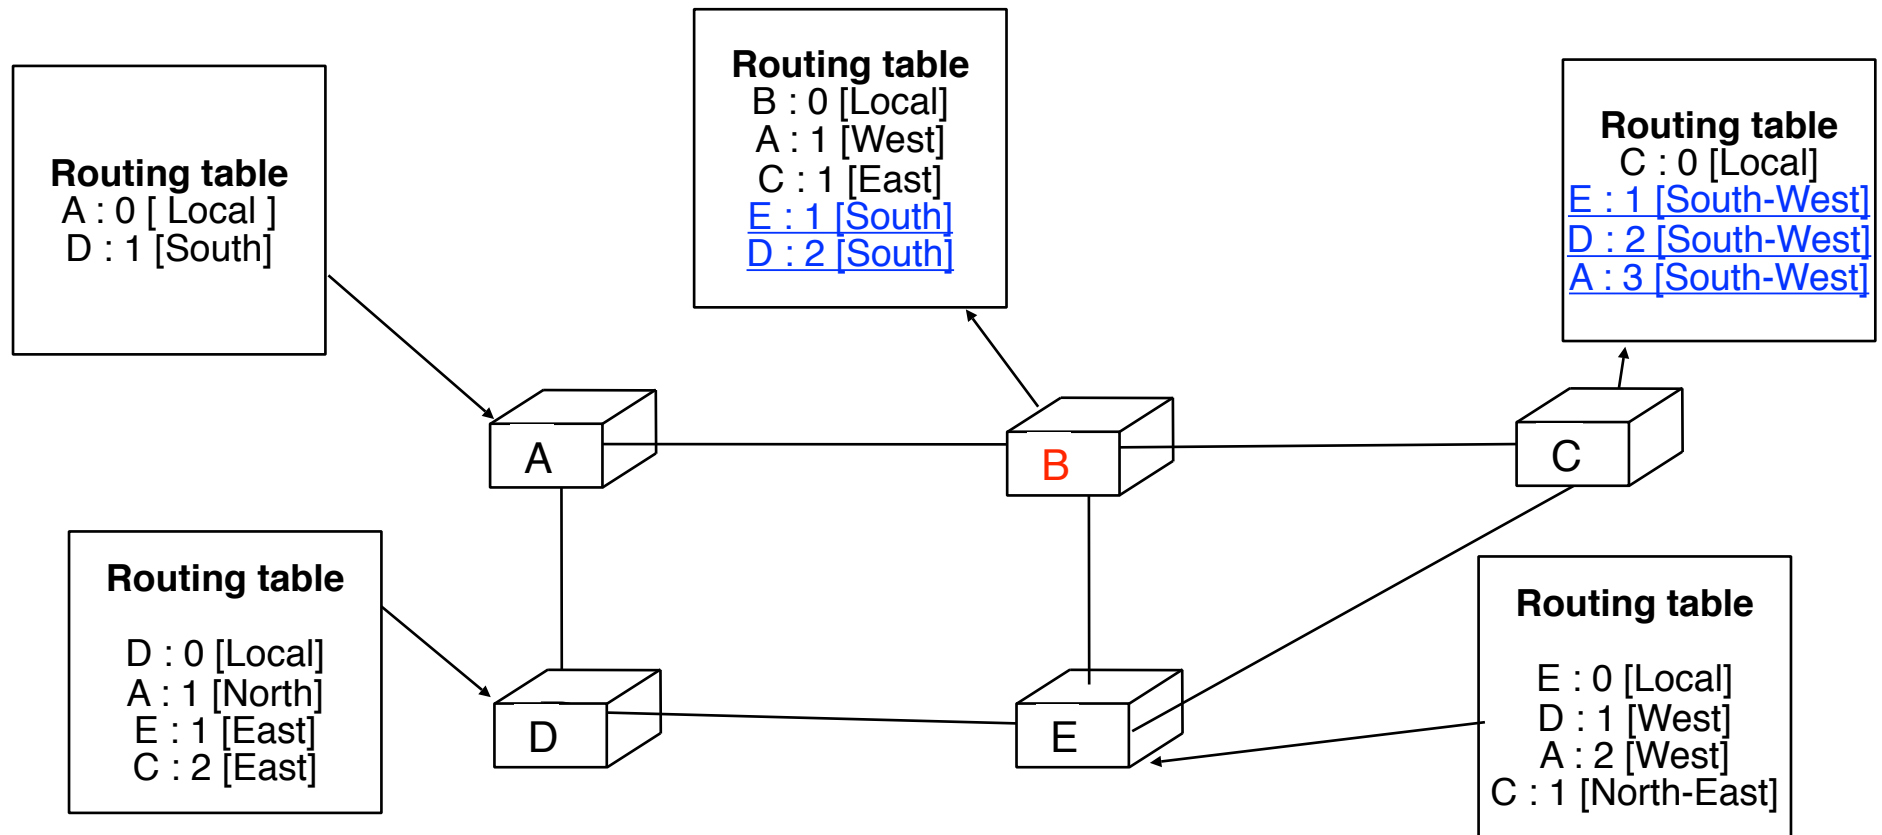
New routes to reach D, A and E

**Reception of distance vector on D**

New routes to reach E and C, longer route for A

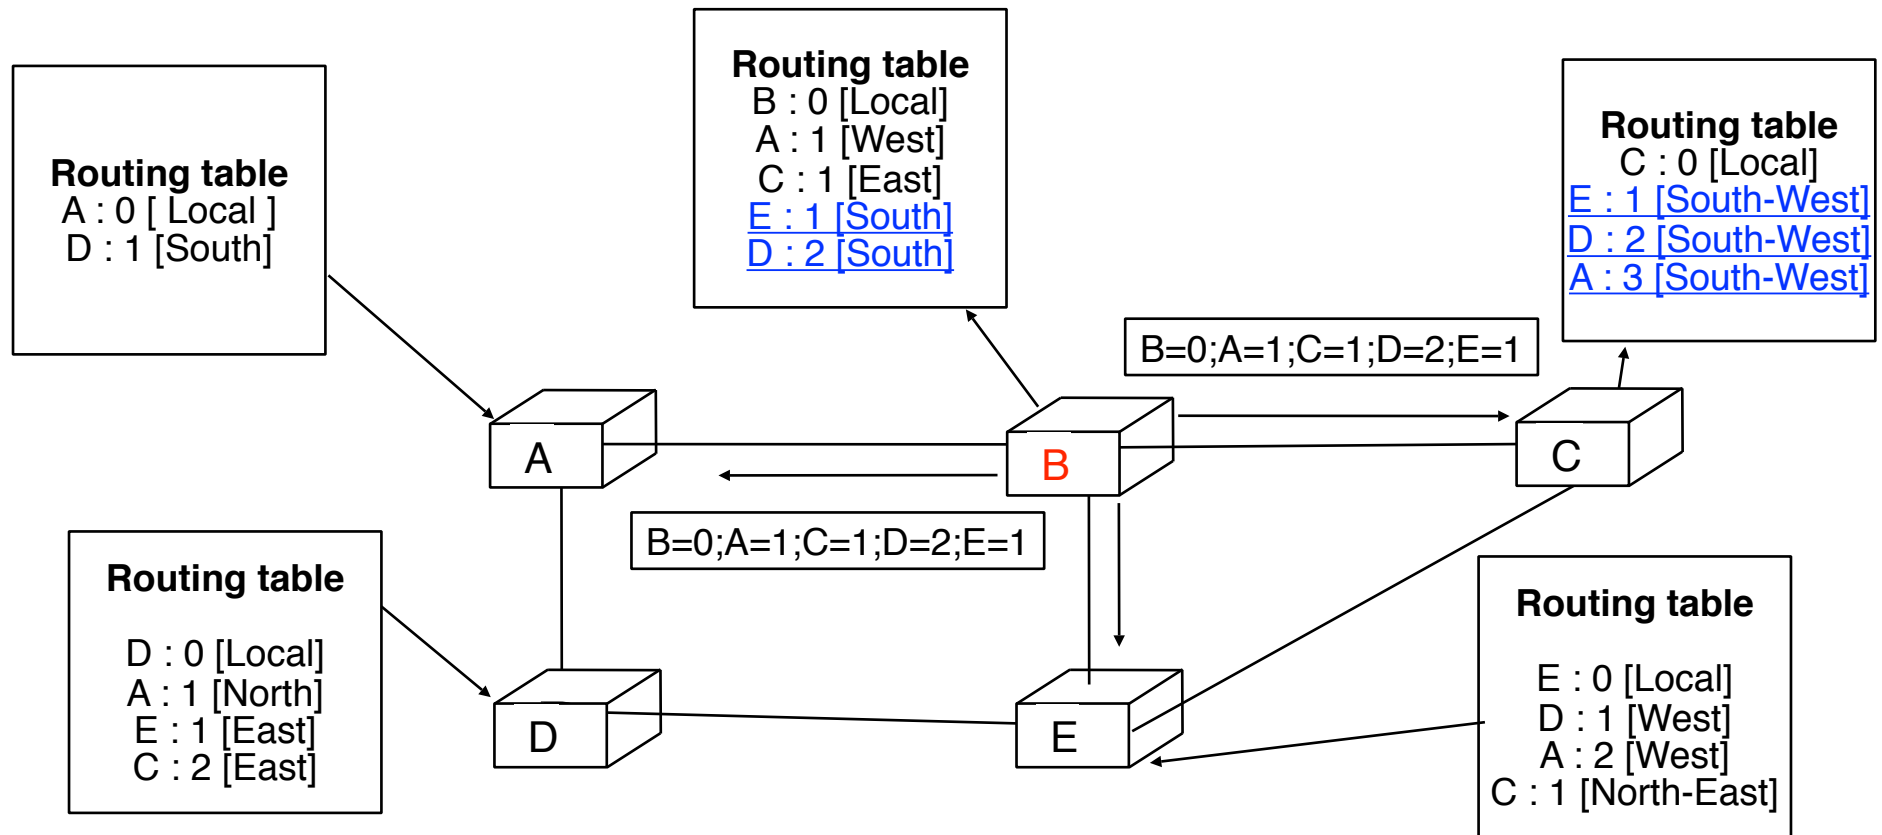
# Distance vectors example (5)

B is the first to send its vector

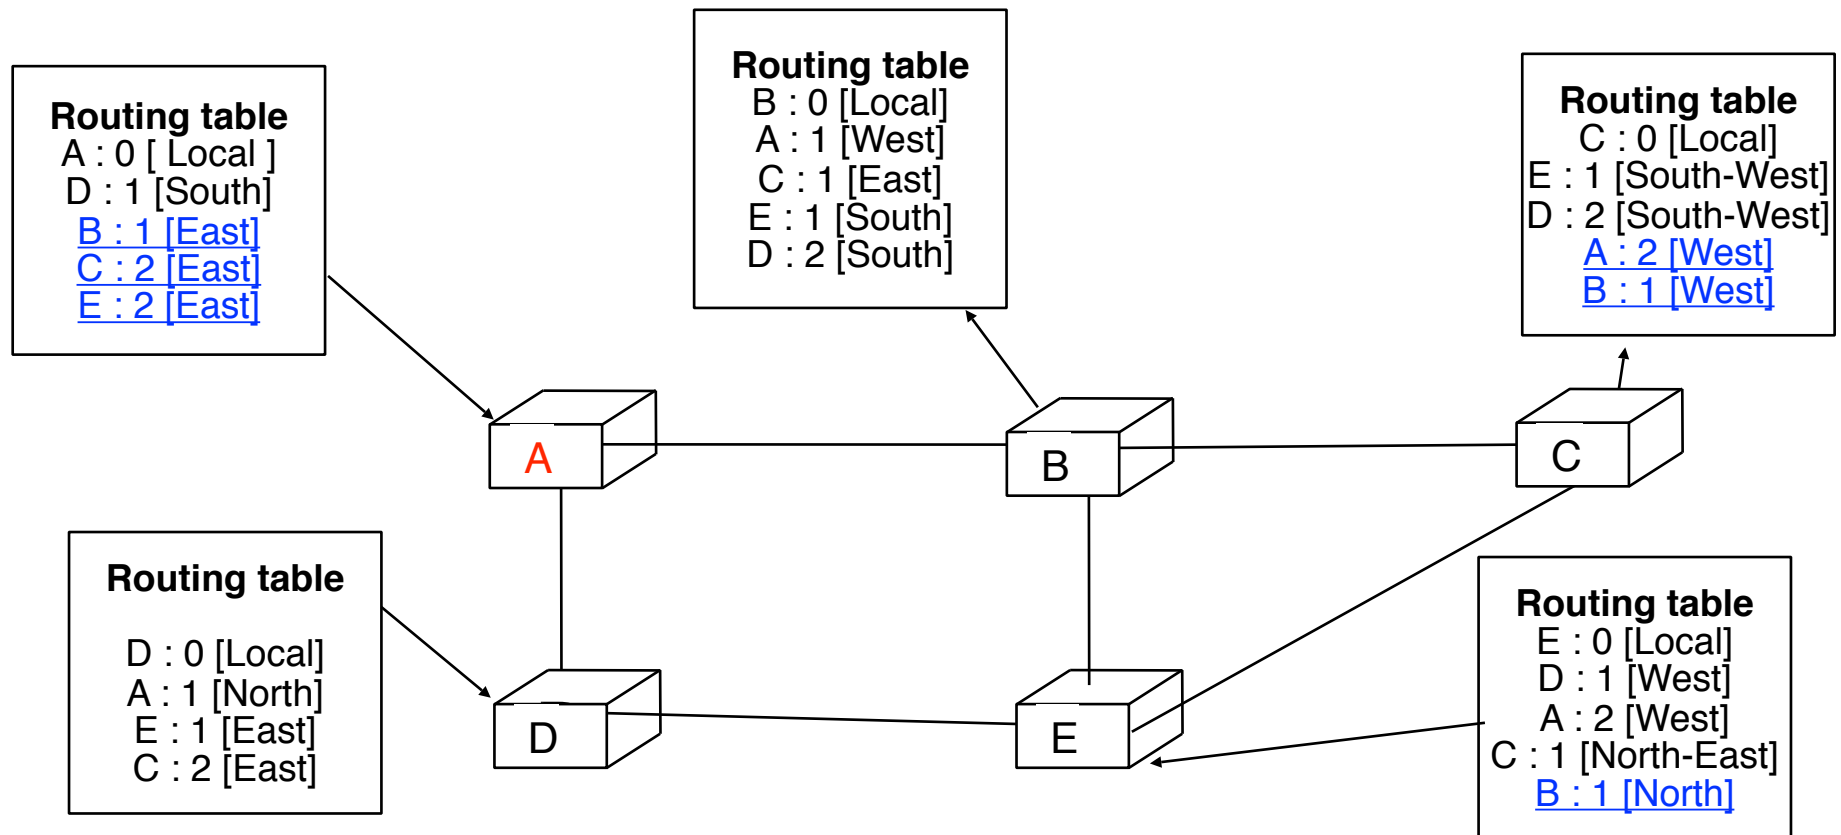


# Distance vectors example (5)

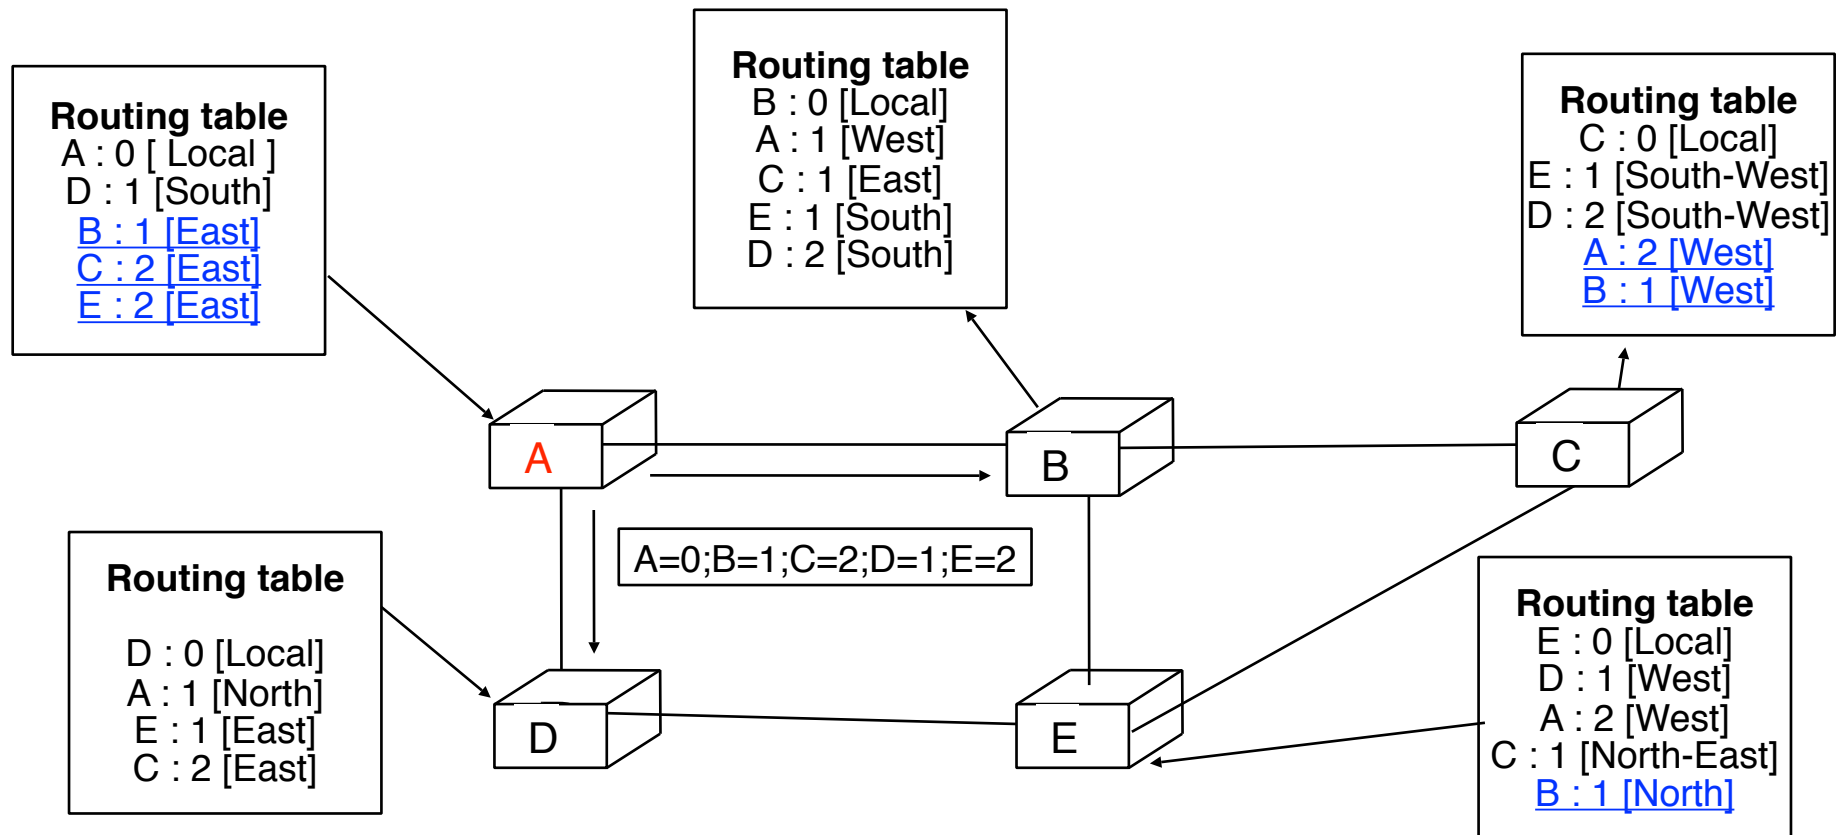
B is the first to send its vector



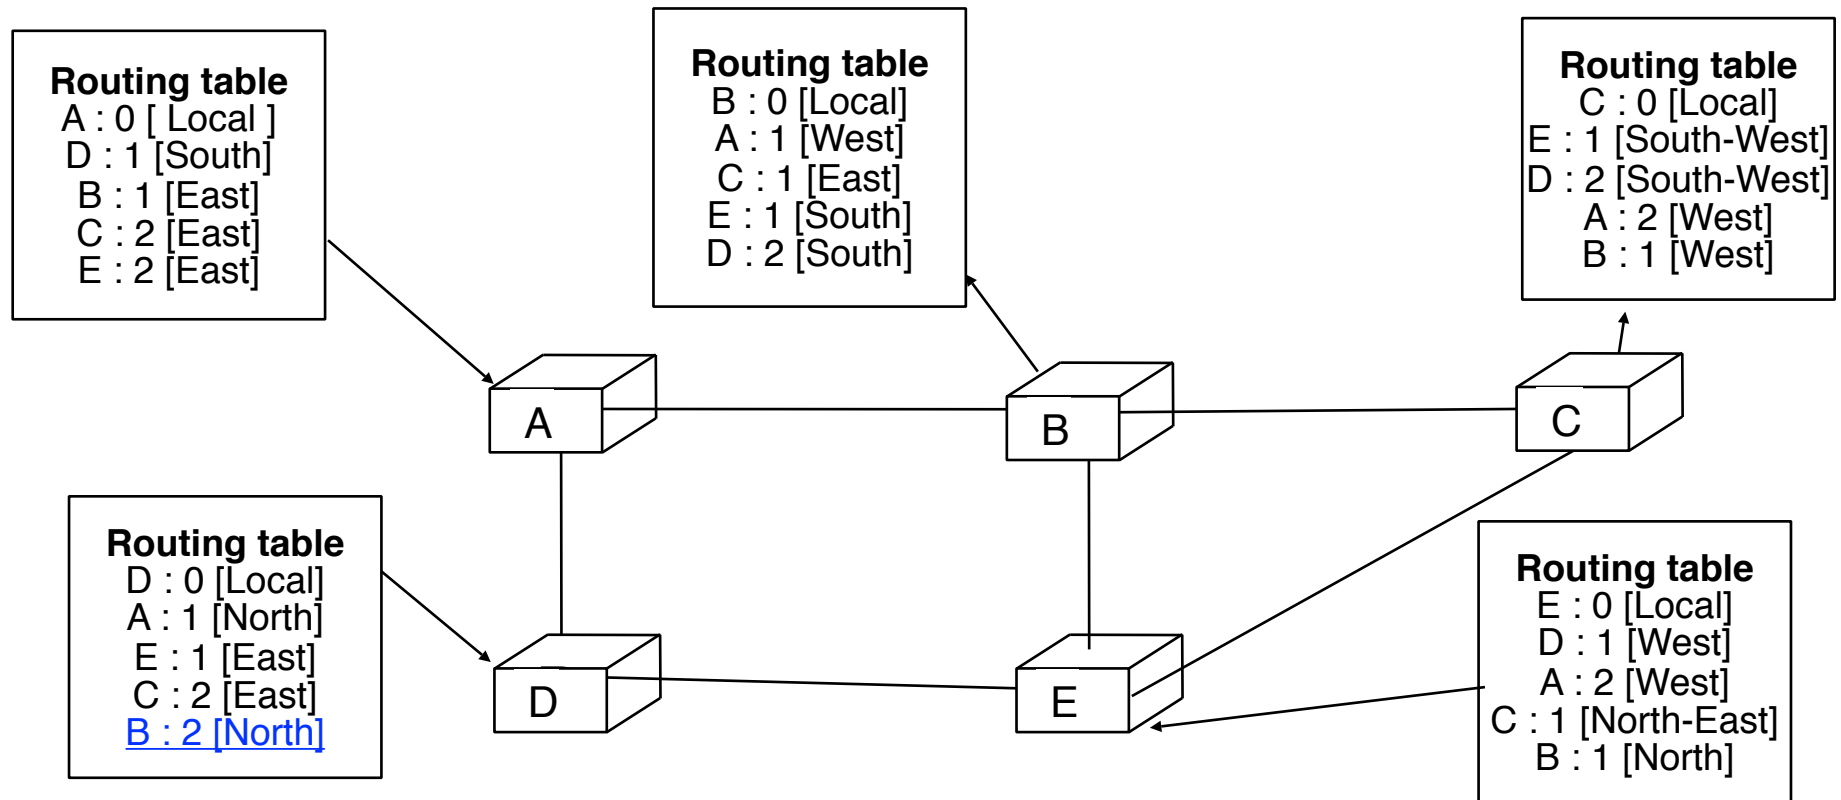
# Distance vectors example (6)



# Distance vectors example (6)



# Distance vectors example (7)



All routers know how to reach all other routers


Routing tables are stable

If a distance vector is sent by one router, it will not cause any change to the routing table of other routers in the network

# Distance vectors

## Link failures

How to deal with link failures ?



Two problems must be solved for failures

- How to detect that the link has failed ?

- How to indicate to all routers that they should update their routing table since the paths that use link A-B do not work anymore ?

# Detection of link failures

---

## Two types of solutions

- rely on failure information from datalink or physical layer

  - fast and reliable

  - unfortunately not supported by all datalink/physical layers

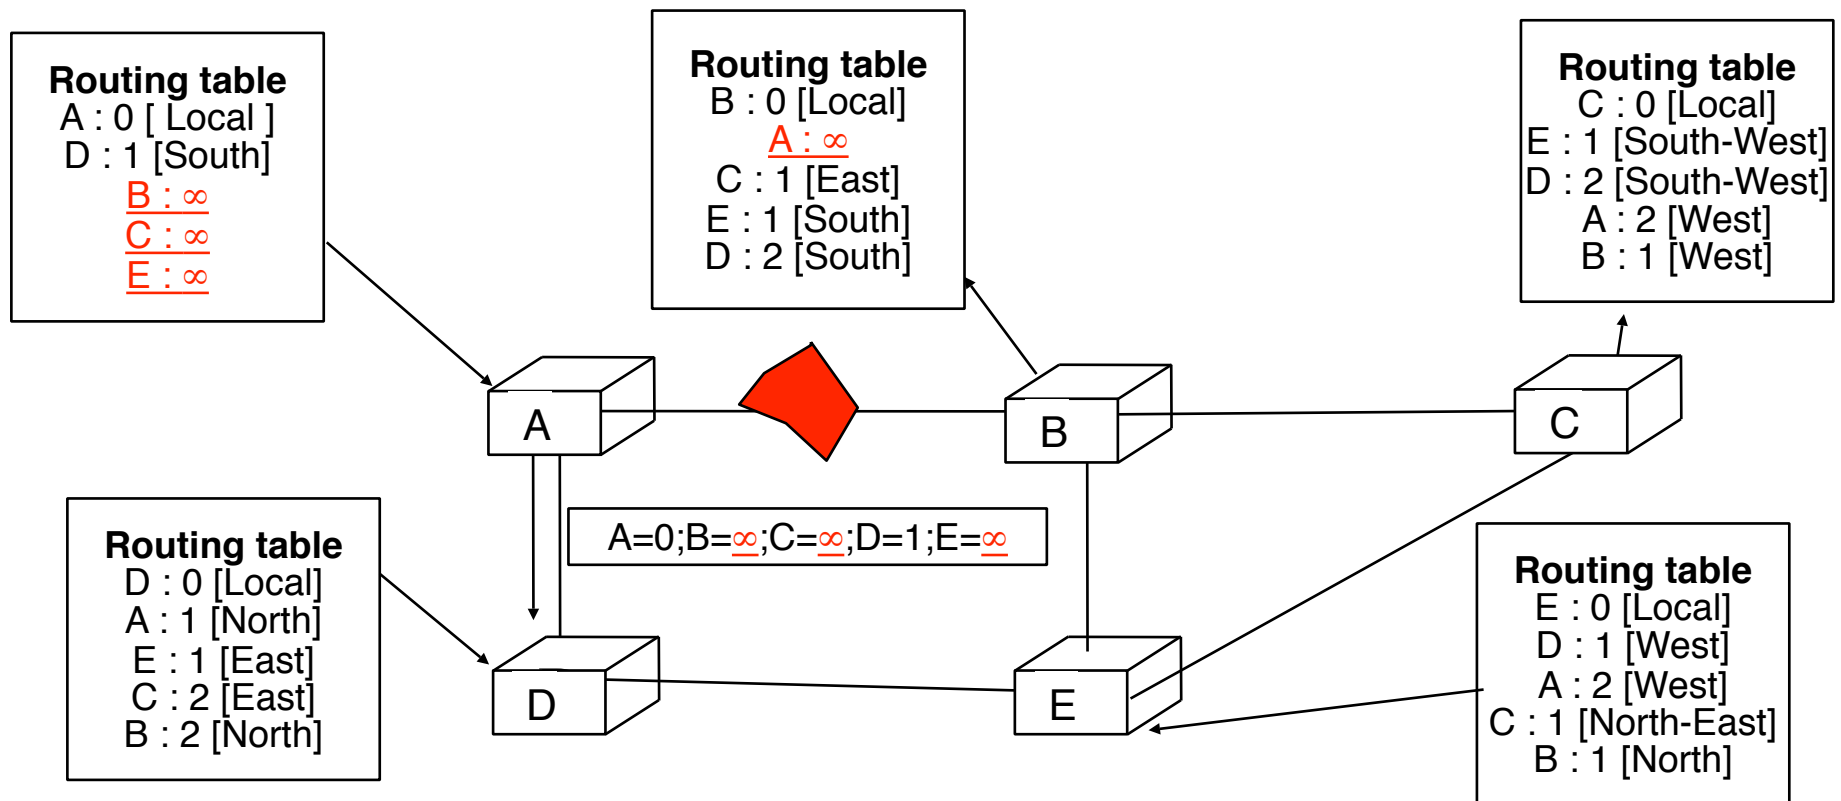
- ask each router to regularly send its distance vector (e.g. every 30 seconds)

  - If a router does not receive a refresh for a route in a distance vector from one of its neighbours during some time (e.g. 90 seconds), it assumes that the route is not available anymore



# How to update the routing table ?

All routes that use a failed link are marked with an infinite cost



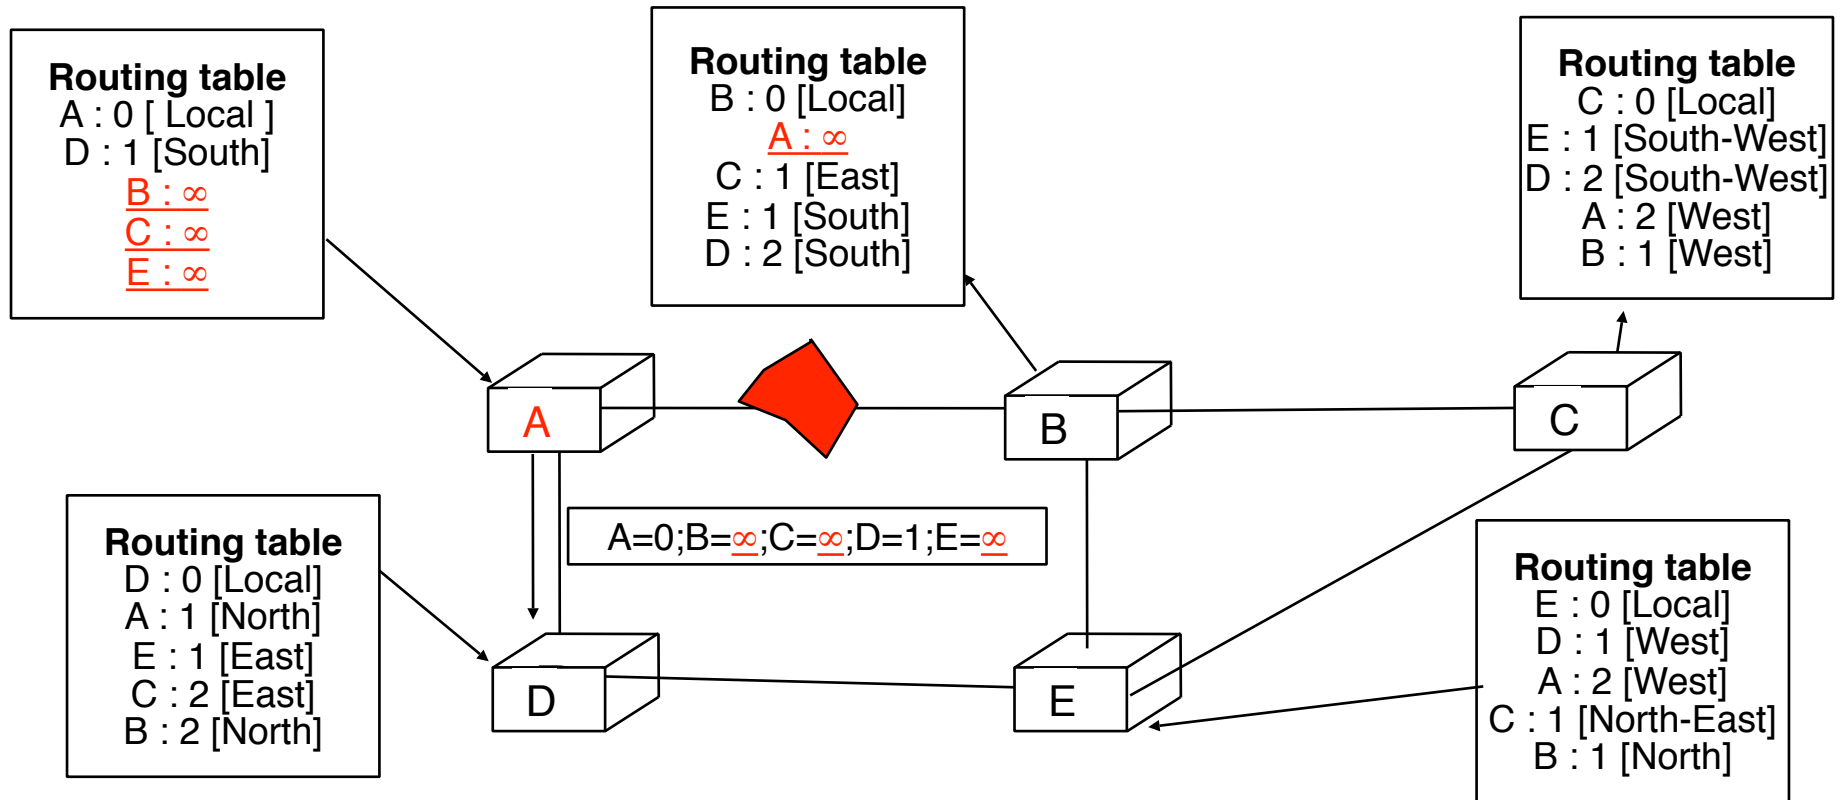
# How to update the routing table ? (2)

---


## Reception of a distance vector

```
Received(Vector V[], link l)
{ /* received vector from link l */
  for each destination=d in V[]
  {
    if (d isin R[])
    { if ( ((V[d].cost+l.cost) < R[d].cost) OR
          ( R[d].link == l) )
      { /* better route or change to current route */
        R[d].cost=V[d].cost+l.cost;
        R[d].link=l;
      }
    }
    else
    { /* new route */
      R[d].cost=V[d].cost+l.cost;
      R[d].link=l;
    }
  }
}
```

# How to update the routing table ? (3)

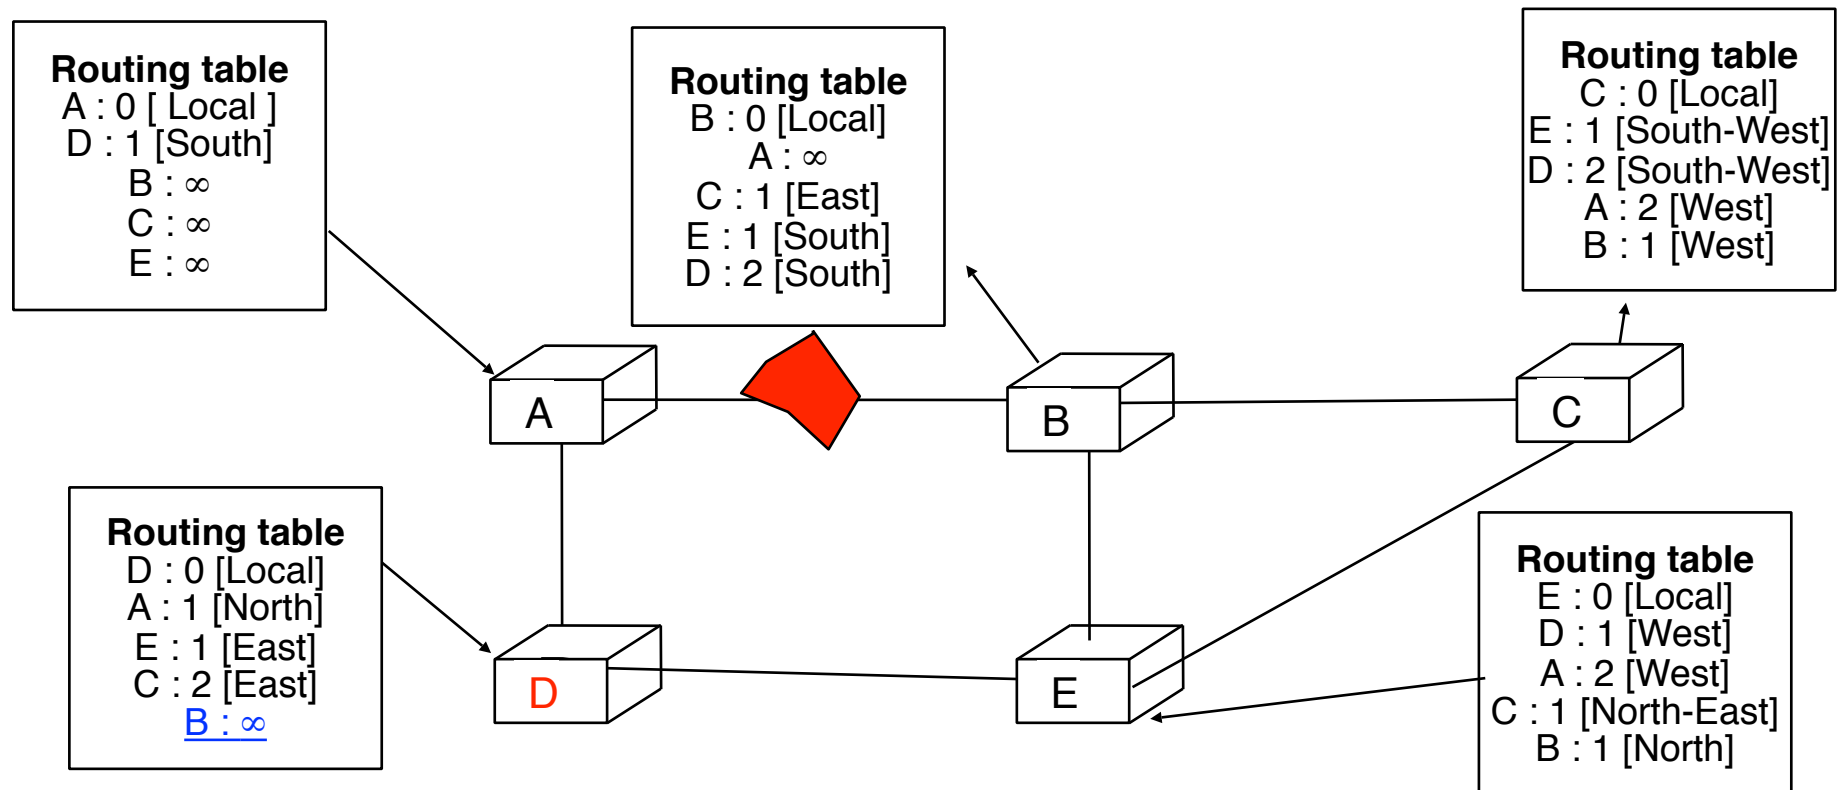


# How to update the routing table ? (3)

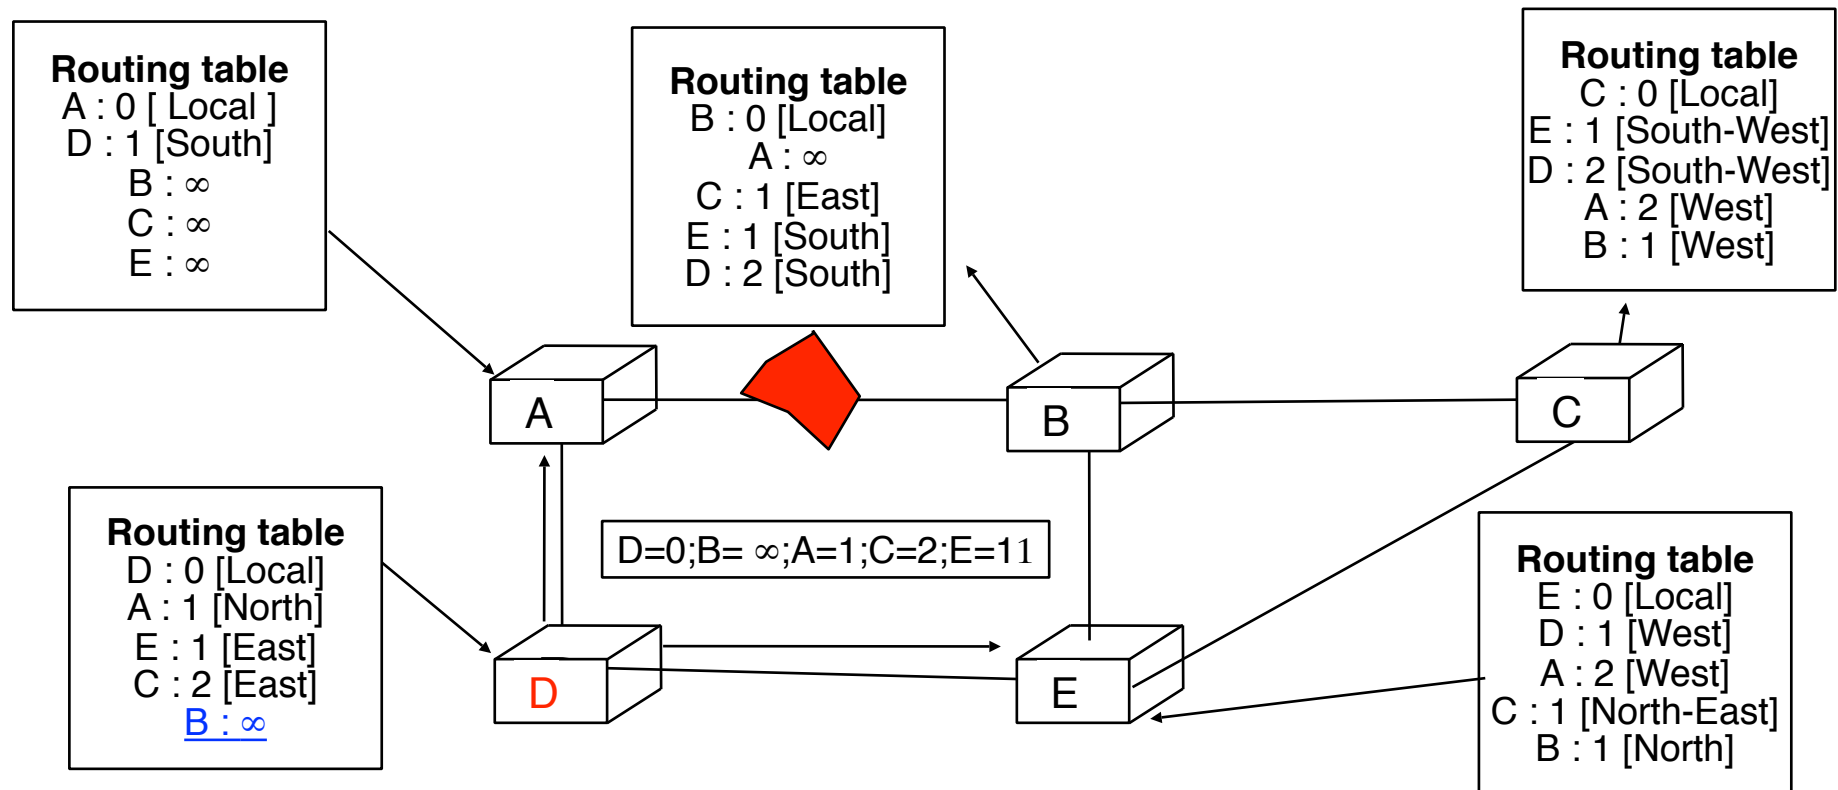


D must remove from its routing tables all the routes that it learned from its North link and are announced now with an  $\infty$  cost

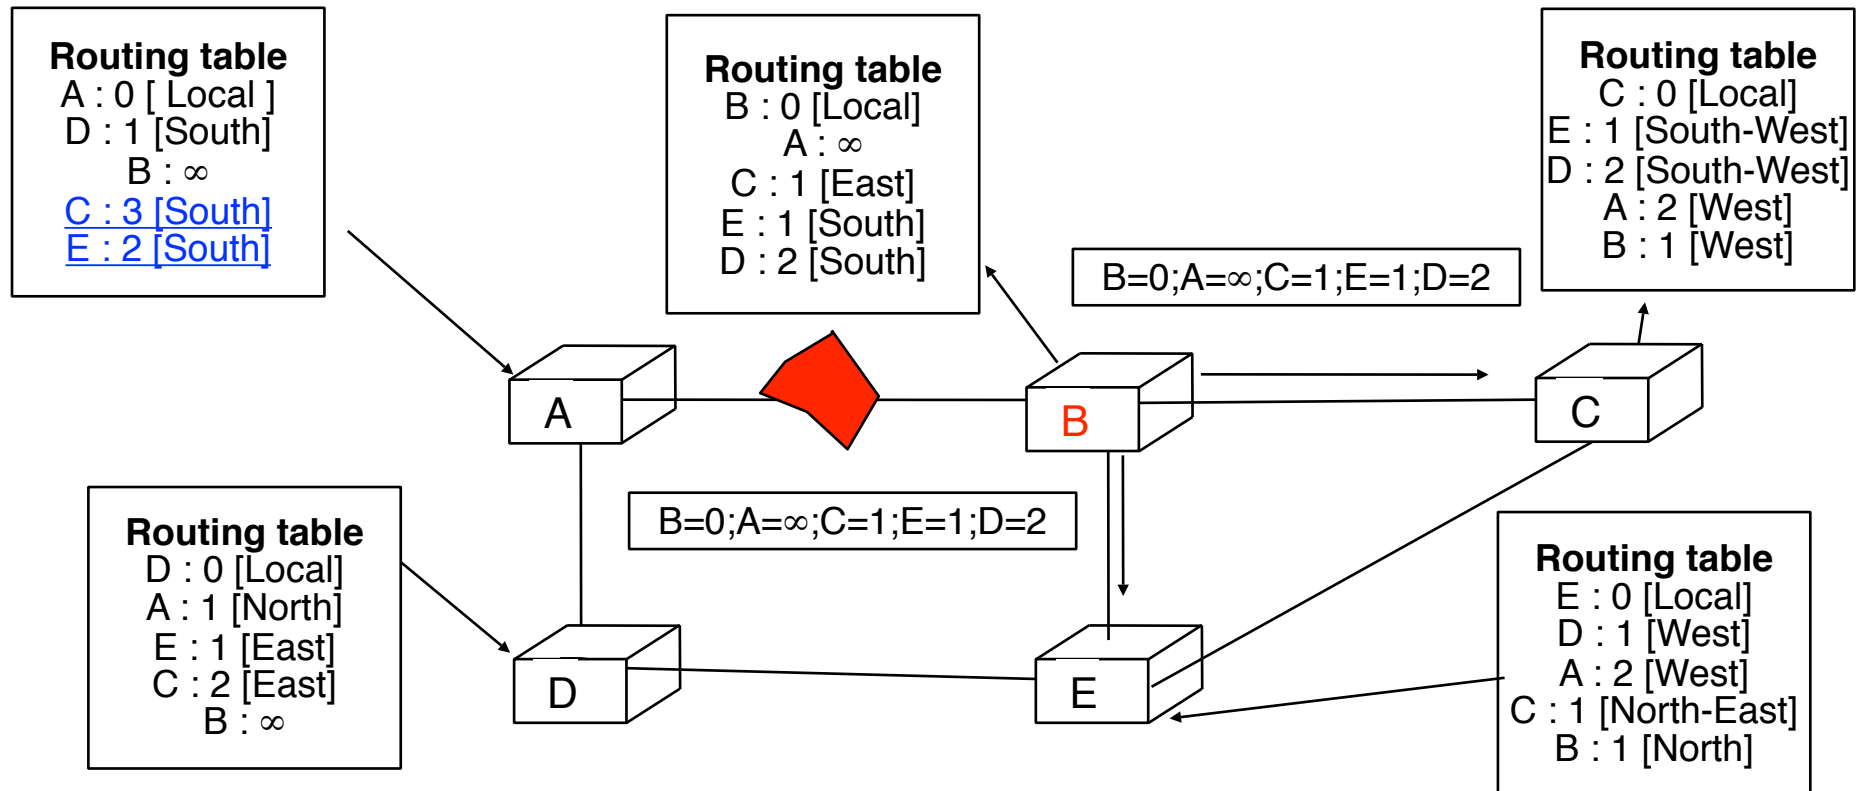
# How to update the routing table ? (4)



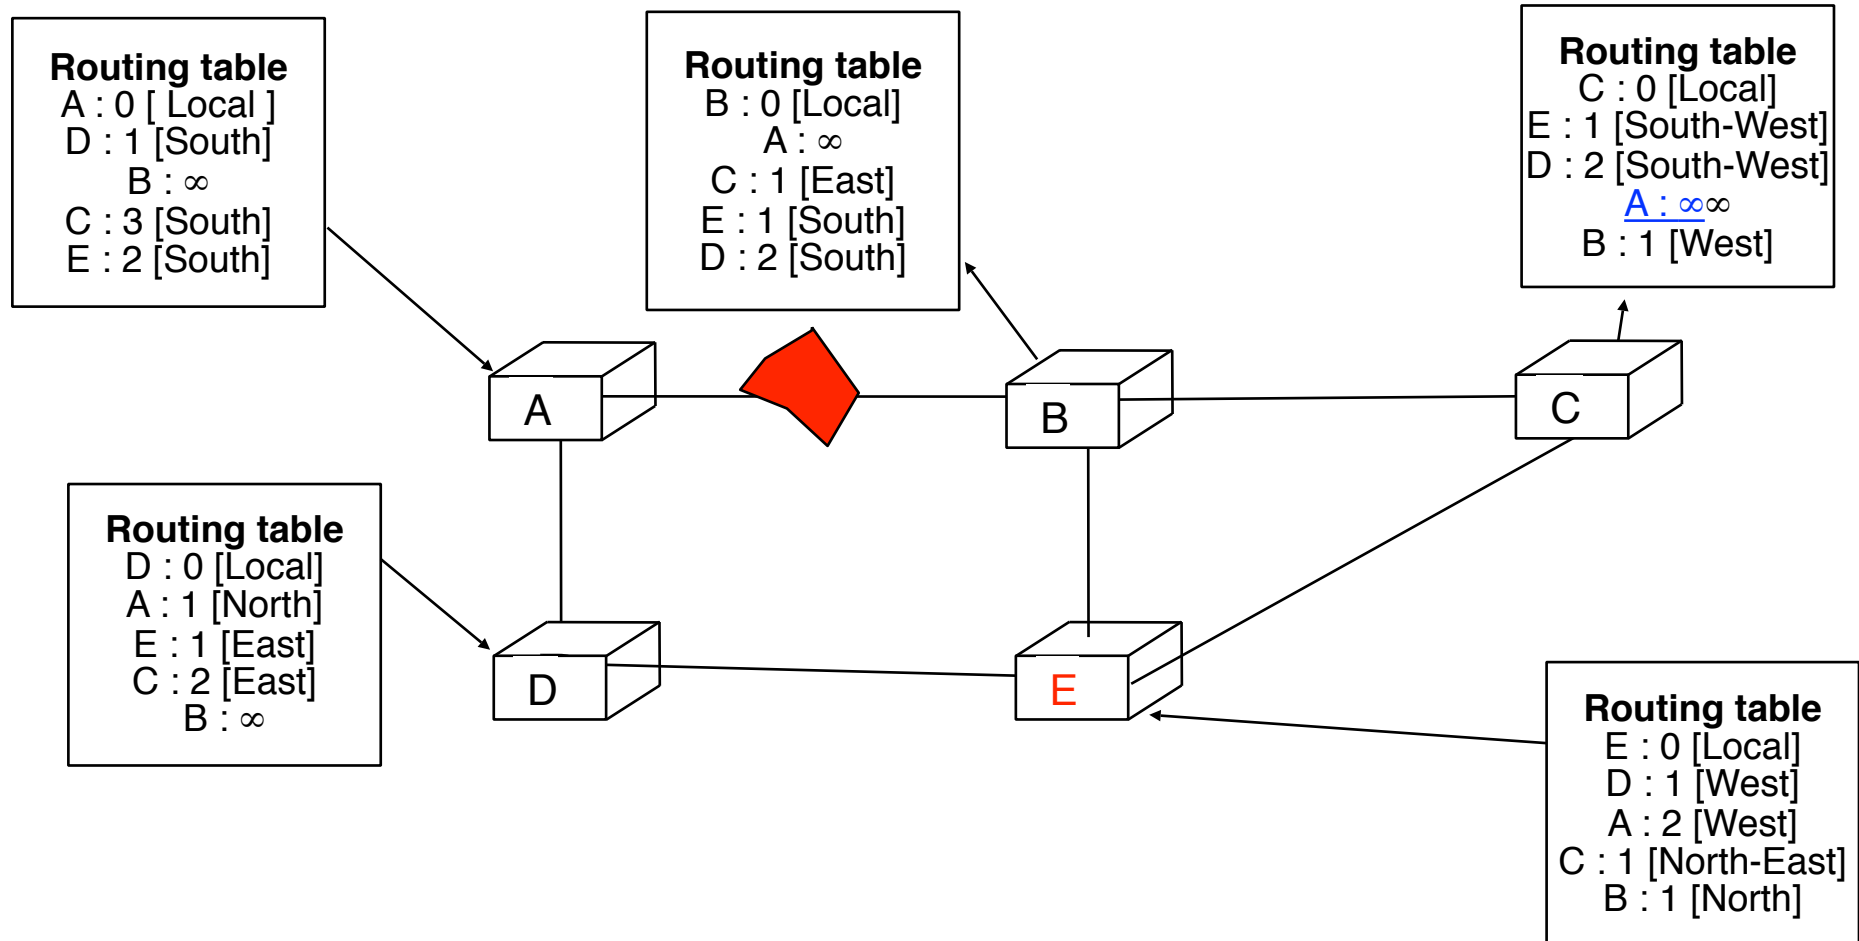
# How to update the routing table ? (4)



# How to update the routing table ? (5)

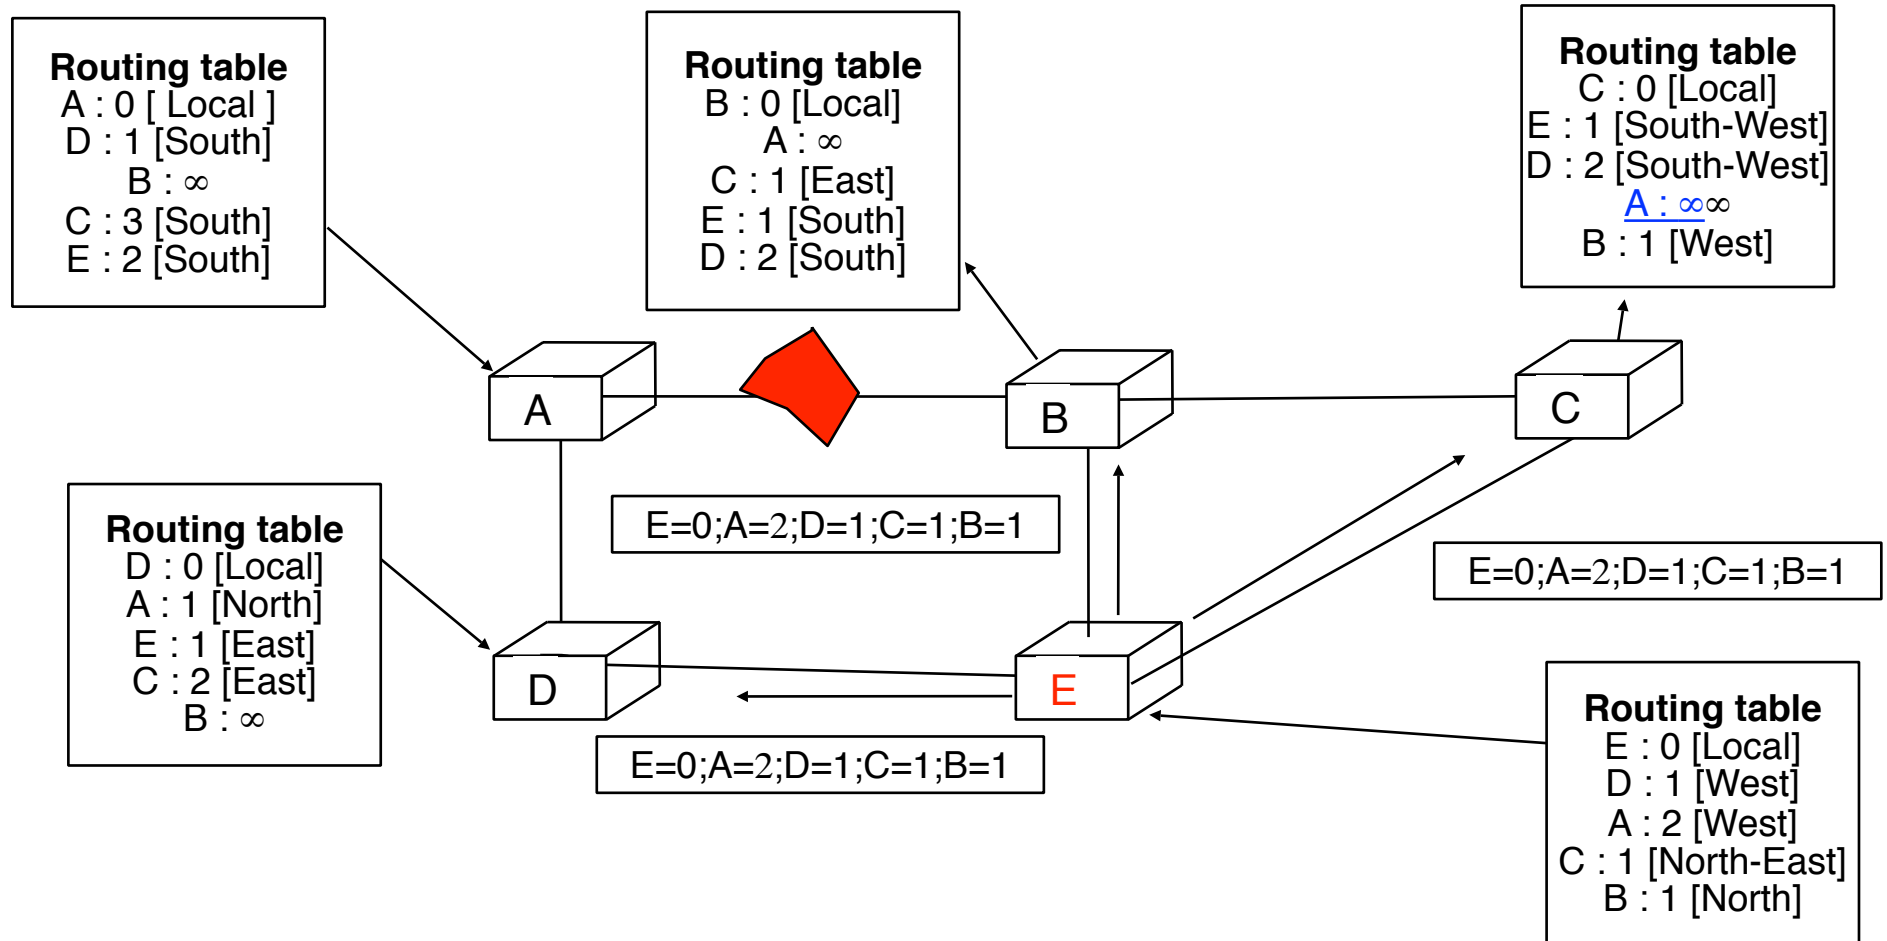


# How to update the routing table ? (6)

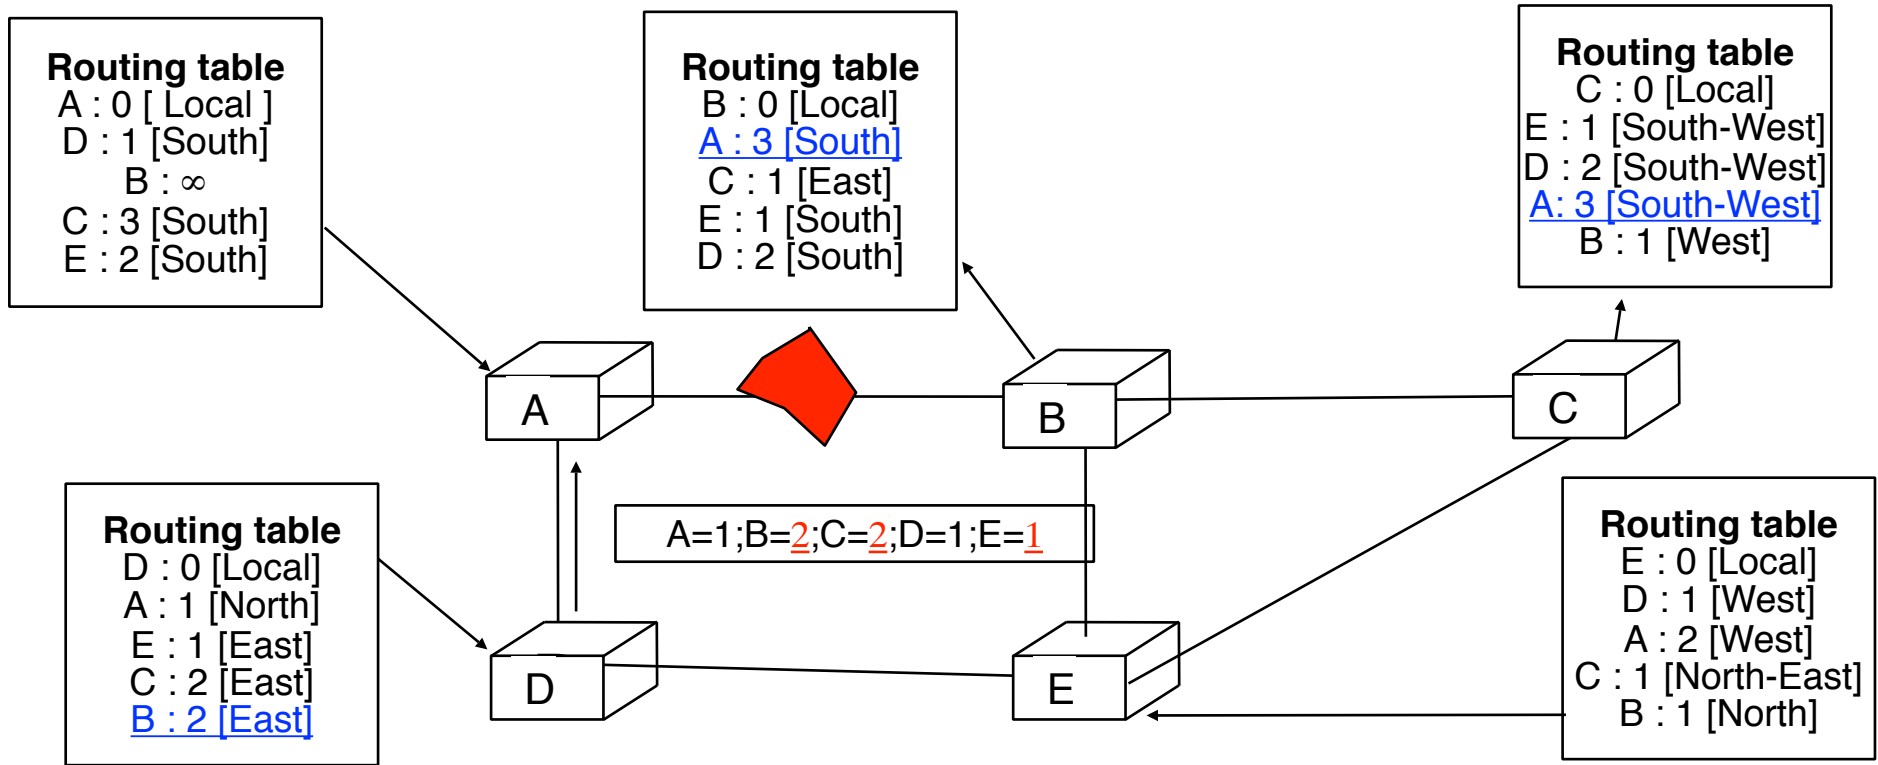





# How to update the routing table ? (6)



# How to update the routing table ? (7)

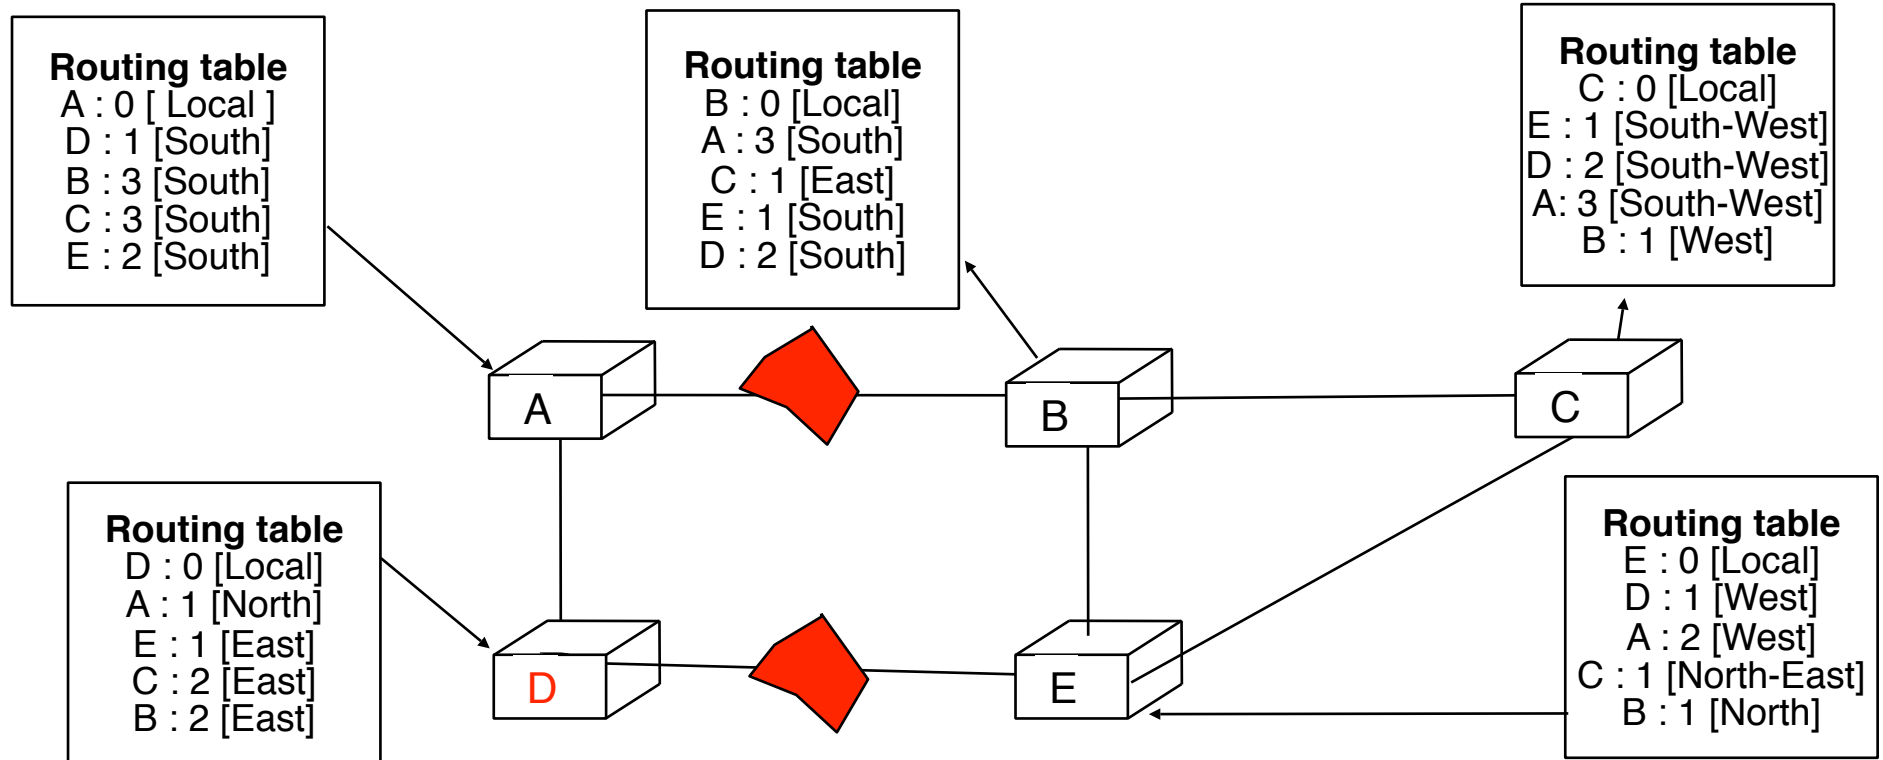


# How to update the routing table ? (7)



Failure has been recovered, all routers are now reachable again from any router

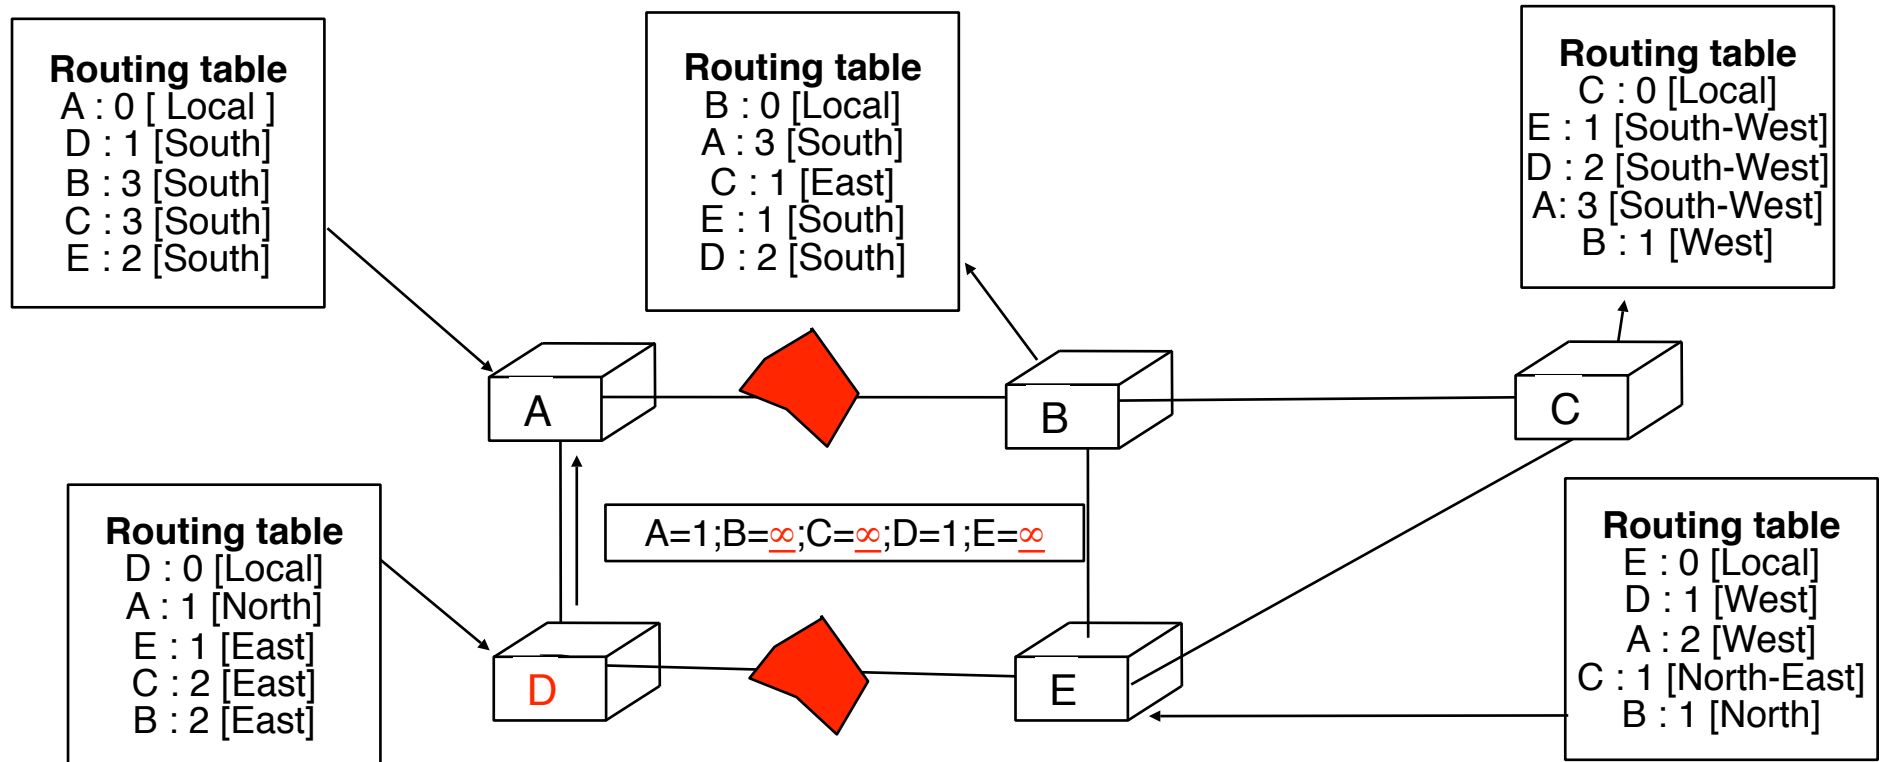
# Second link failure



**D detects the failure**

**If it is the first to send its distance vector, failure is detected and router A updates its routing table**

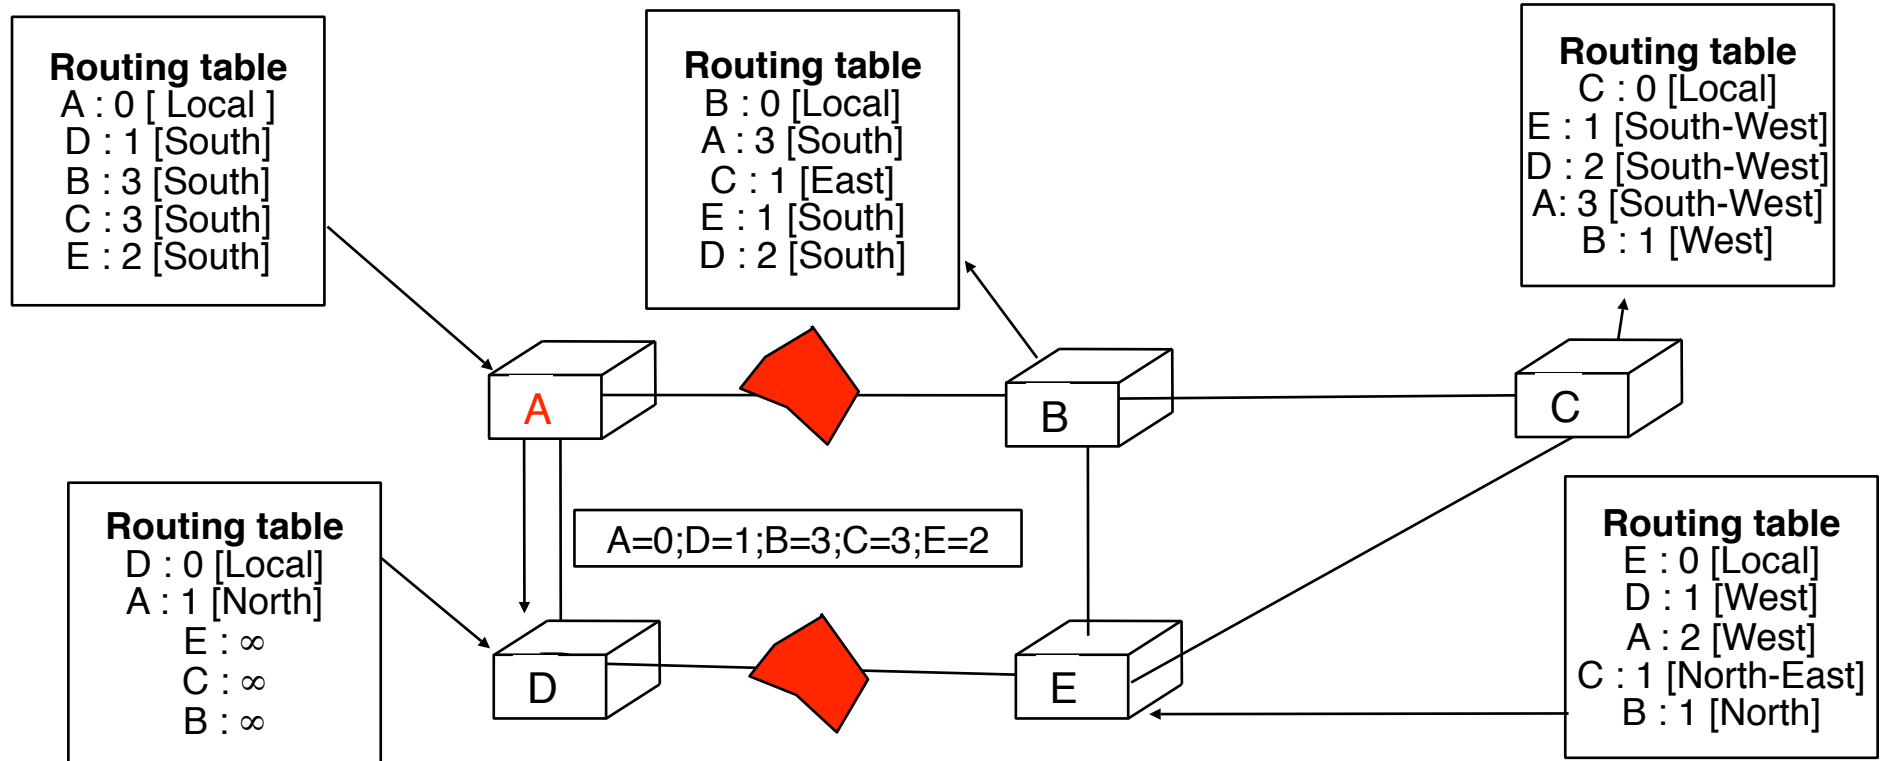
# Second link failure



**D detects the failure**

If it is the first to send its distance vector, failure is detected and router A updates its routing table

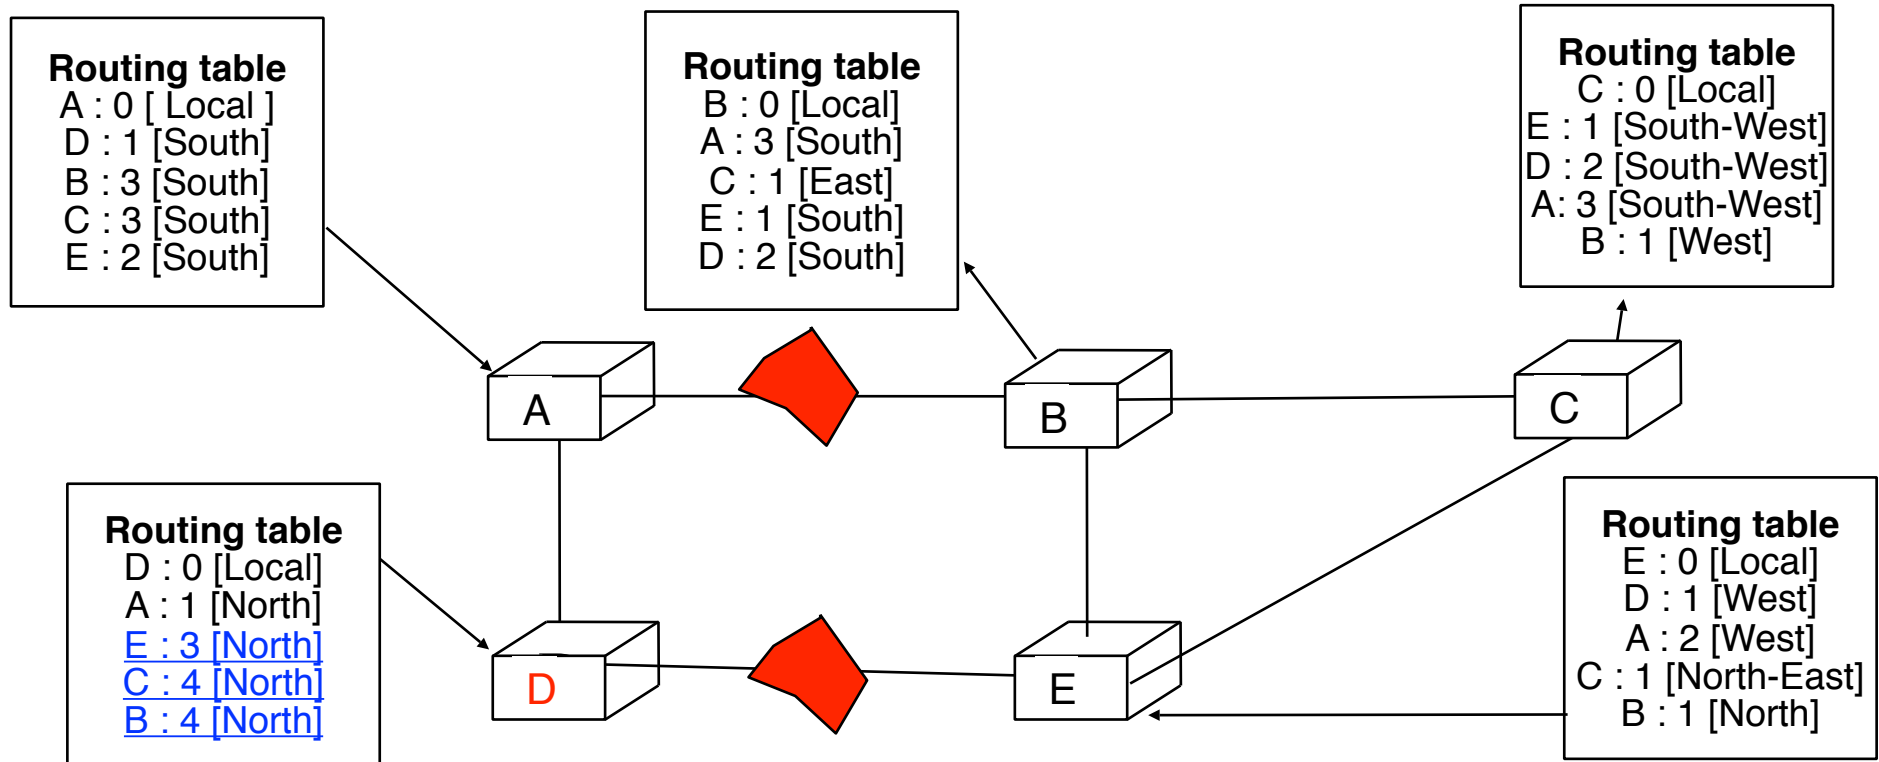
# Second link failure (2)



But if A sends its distance vector before having received or processed D's updated distance vector ...

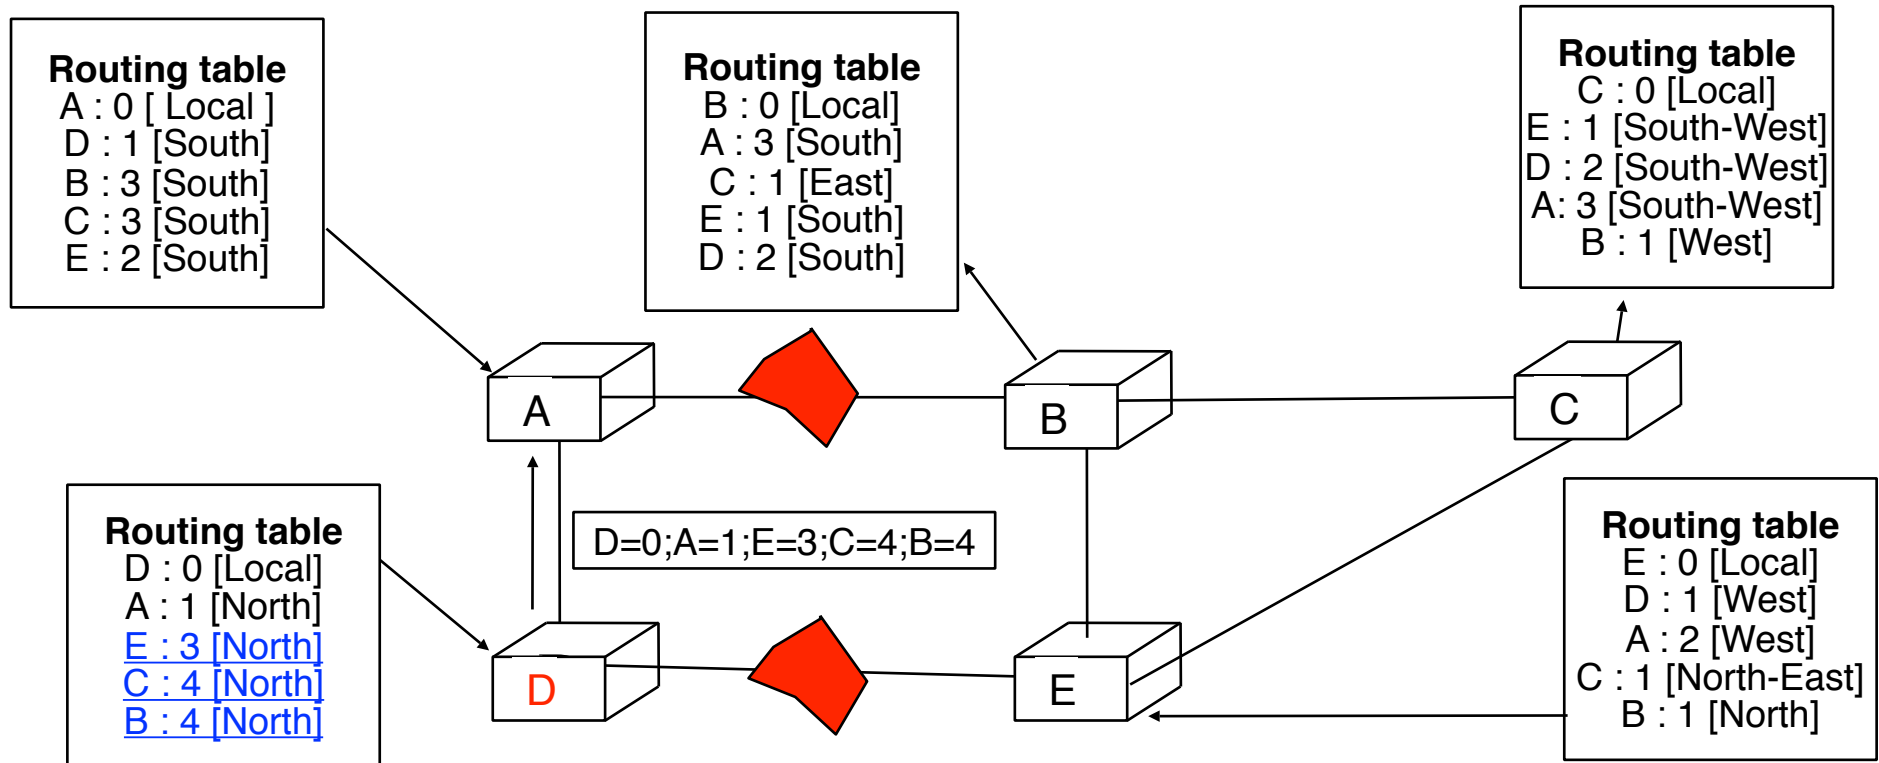
# Second link failure (3)

Upon reception of A's vector, D updates its routing table



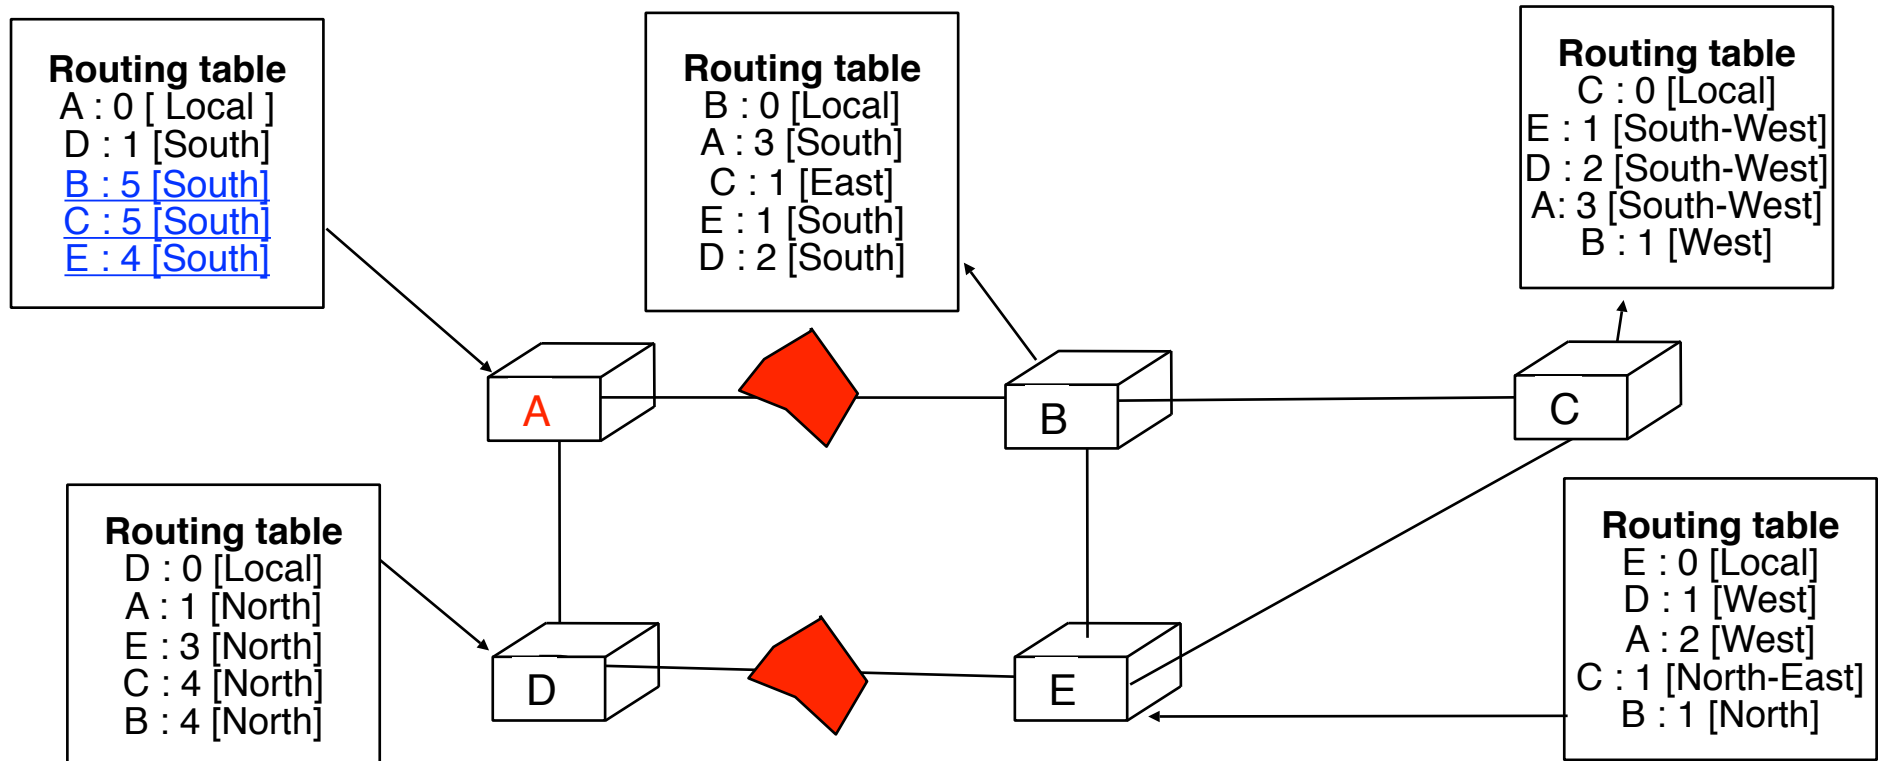
# Second link failure (3)

Upon reception of A's vector, D updates its routing table

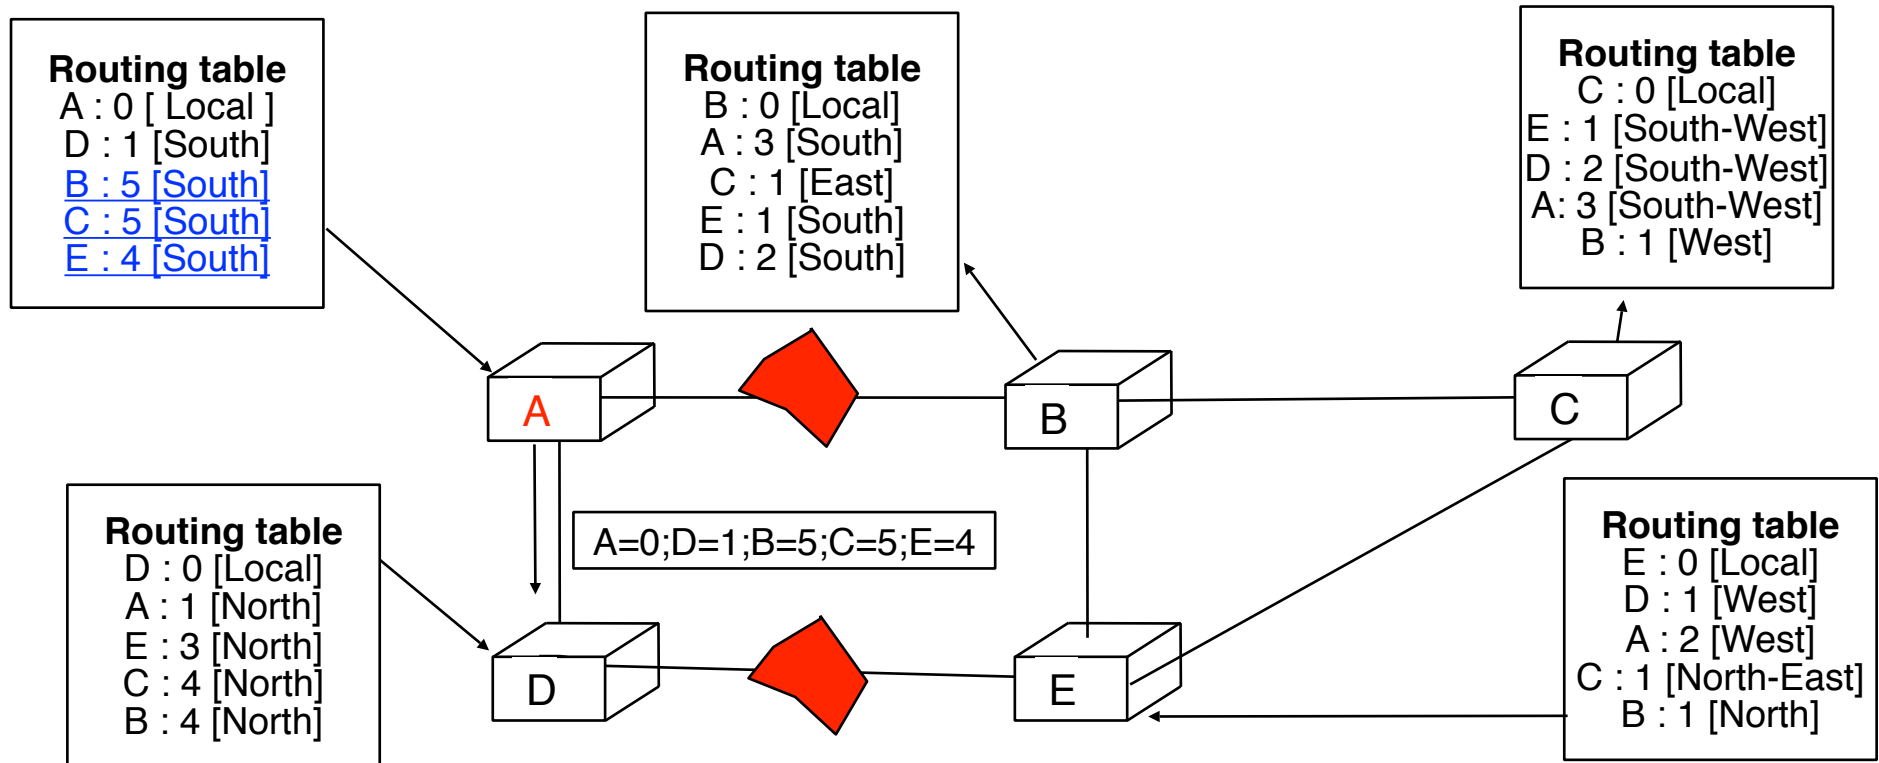




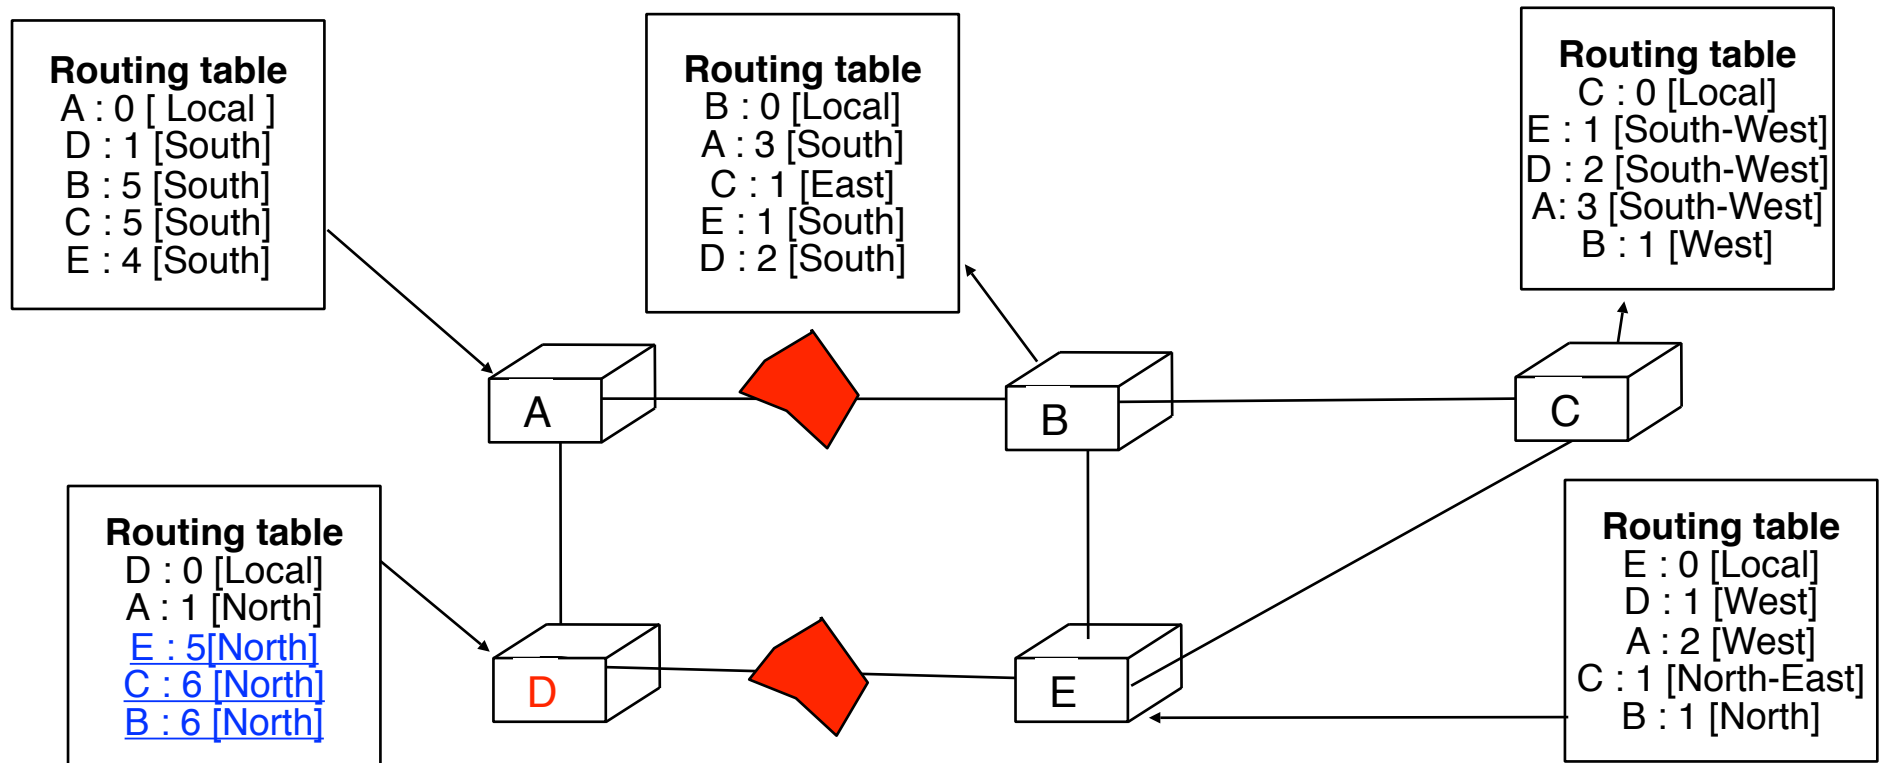
# Second link failure (4)



# Second link failure (4)



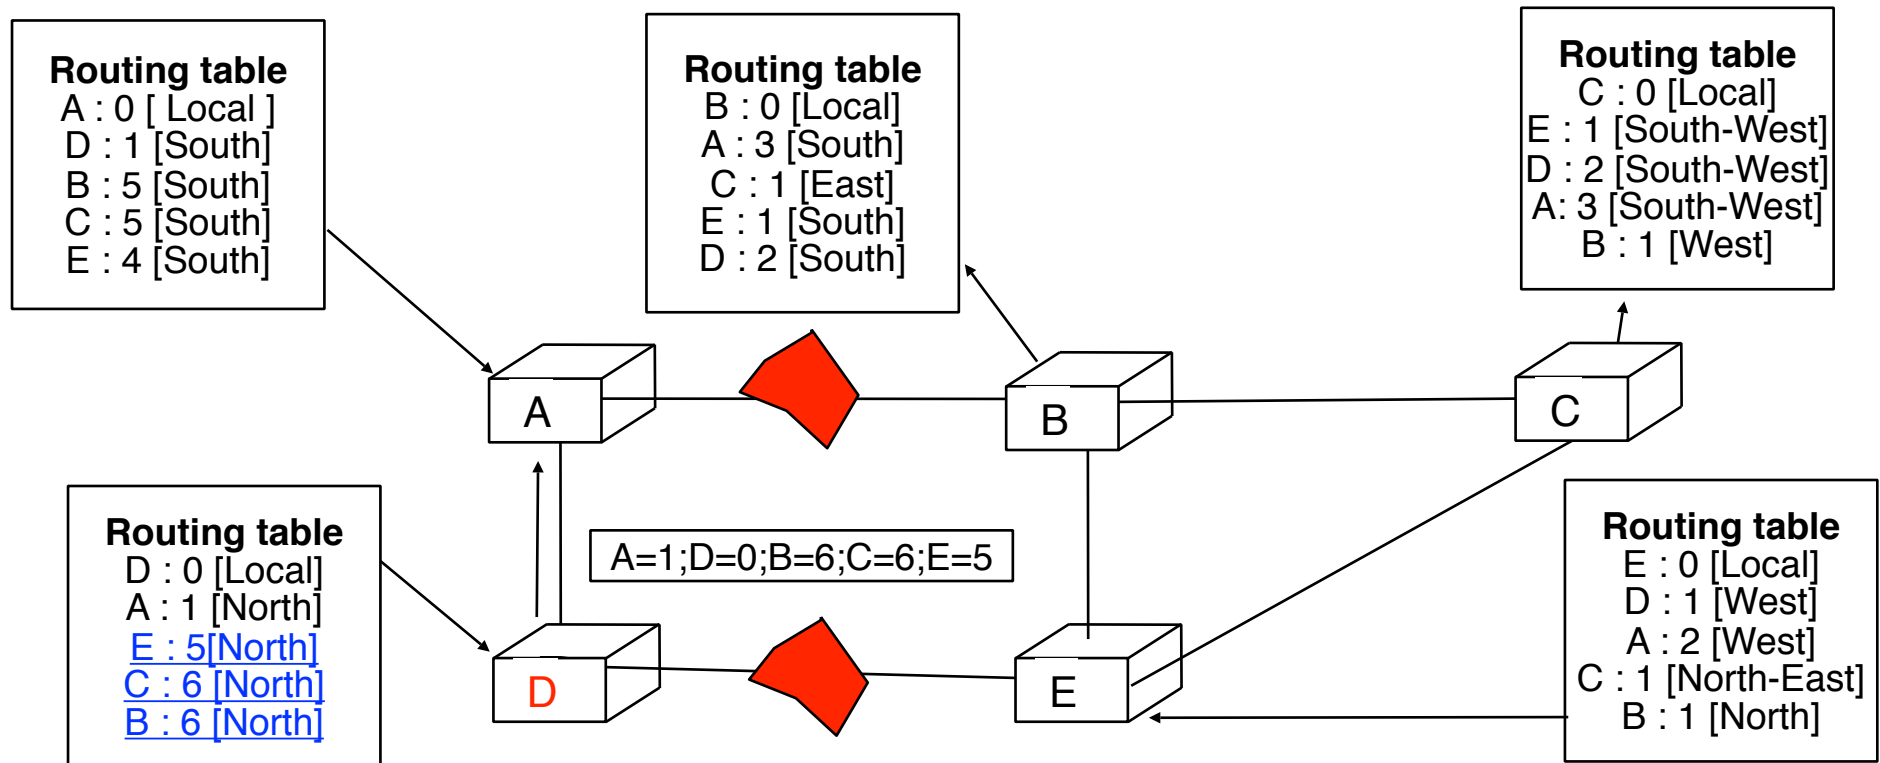
# Second link failure (5)



This problem is called counting to infinity

How can we avoid it ?

# Second link failure (5)



This problem is called counting to infinity

How can we avoid it ?

# Second link failure (6)

Where does counting to infinity comes from ?

A router announces on a link routes that it has already learned via this link

How to avoid counting to infinity ?  
split horizon

each router creates a distance vector for each link

on link  $i$ , router does not announce the routes learned over link  $i$

Pseudocode

Every  $N$  seconds:

for each link= $l$

{ /\* one different vector for each link \*/

Vector=null;

for each destination= $d$  in  $R[]$

{

if ( $R[d].link \neq l$ )

{ Vector=Vector+Pair( $d, R[d].cost$ ); }

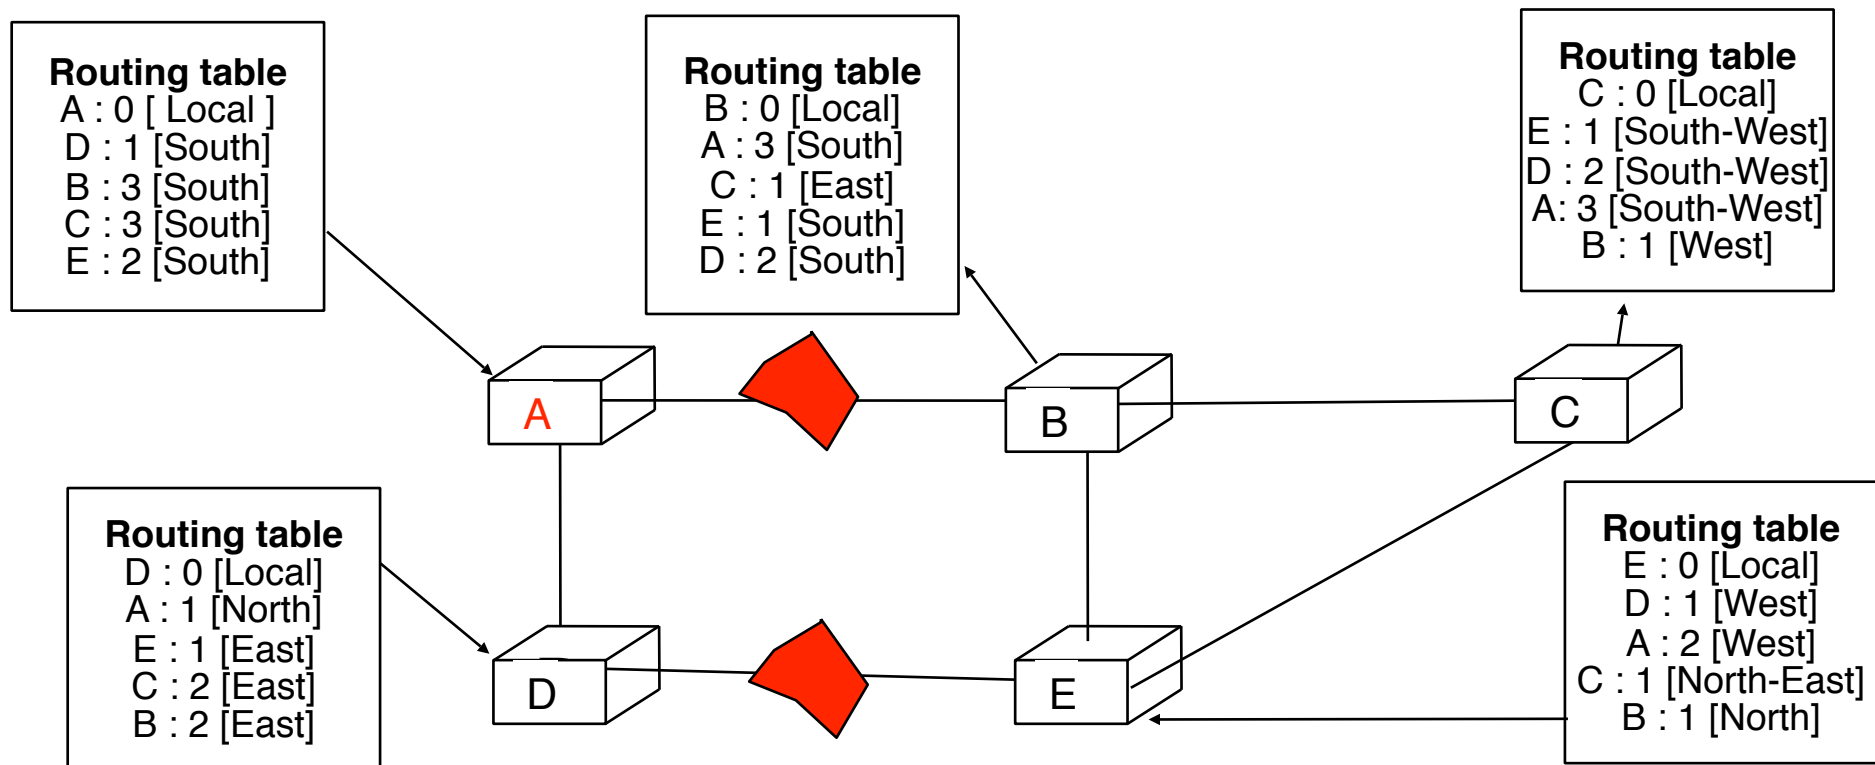
}

Send(Vector);

}

# Split horizon

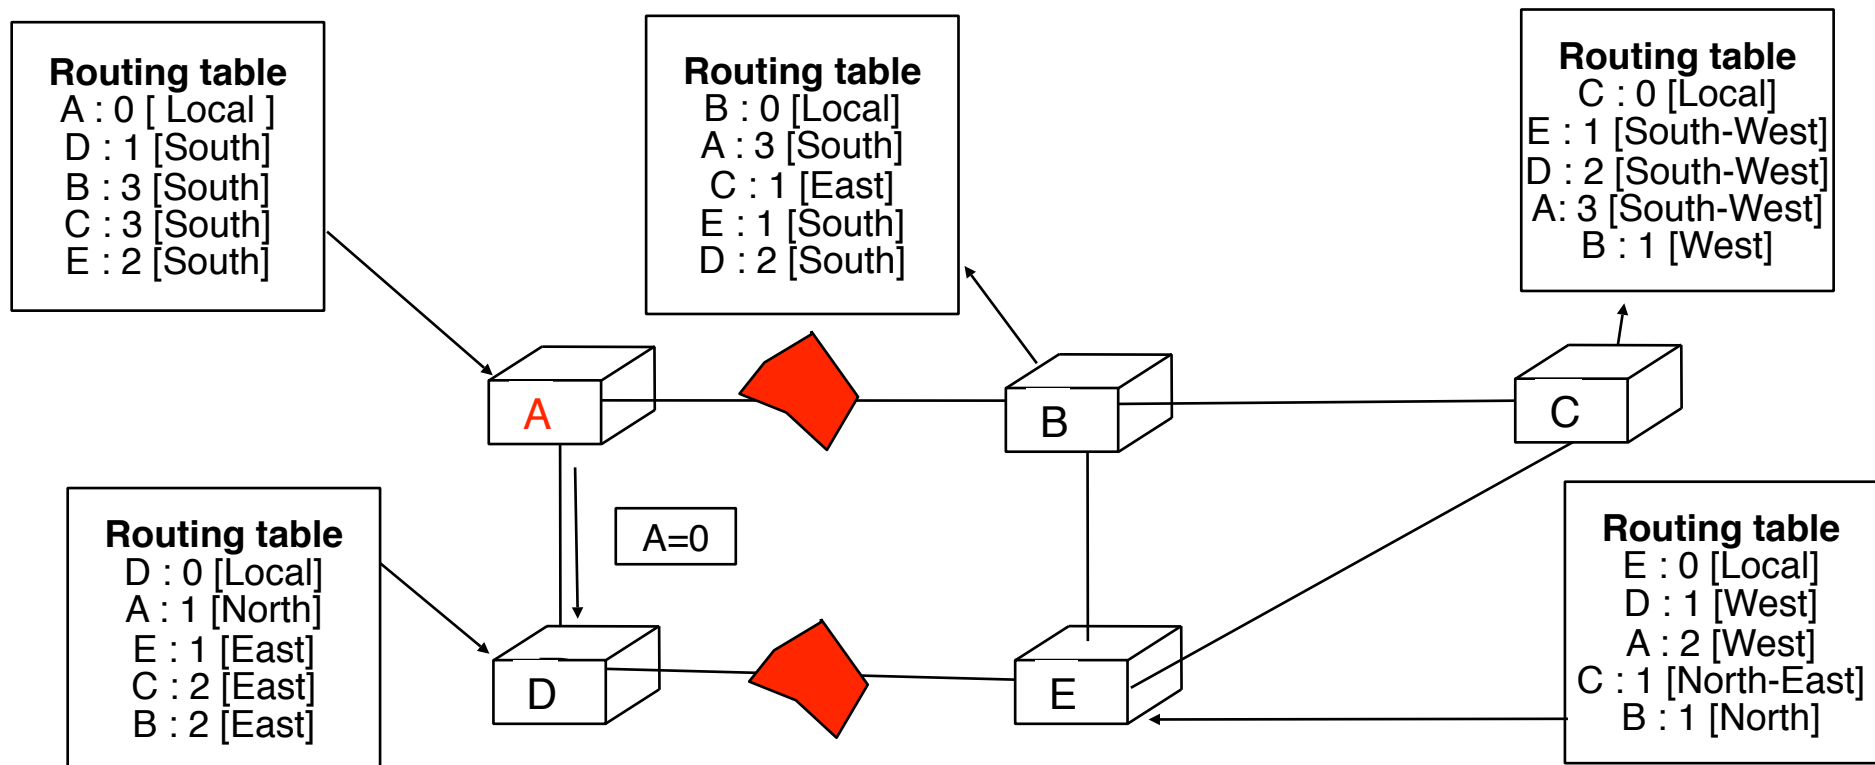
## Back to previous example



A will not pollute D's routing table with split horizon

# Split horizon

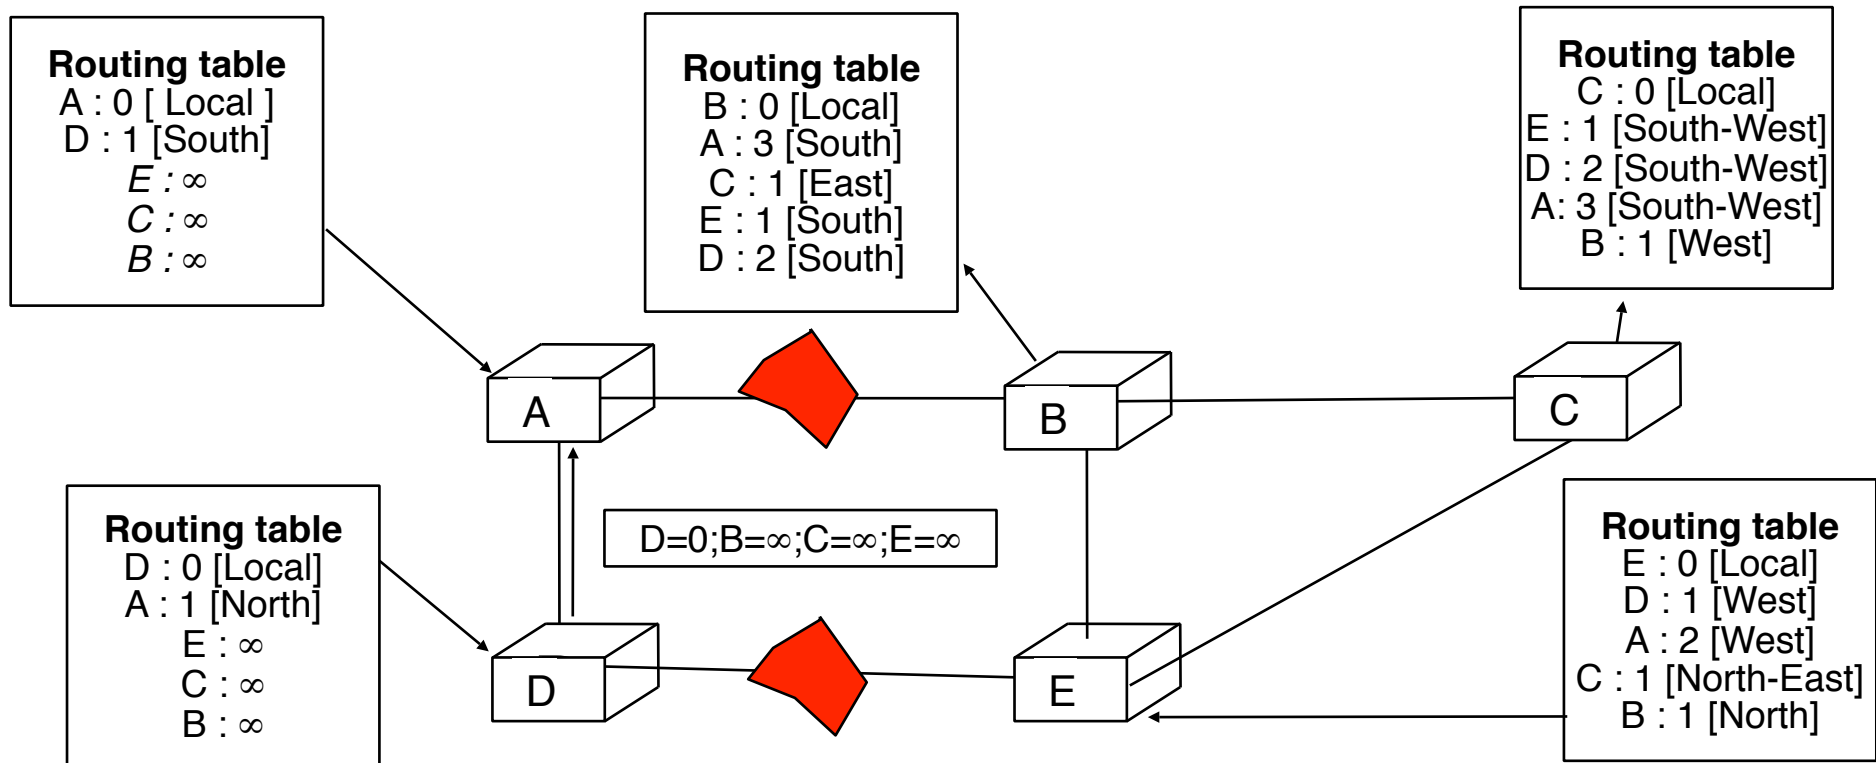
## Back to previous example



A will not pollute D's routing table with split horizon

# Split horizon (2)

D can also send its distance vector



Does split horizon allows to avoid all counting to infinity problems ?



# Split horizon with poisoning

---

## Improvement

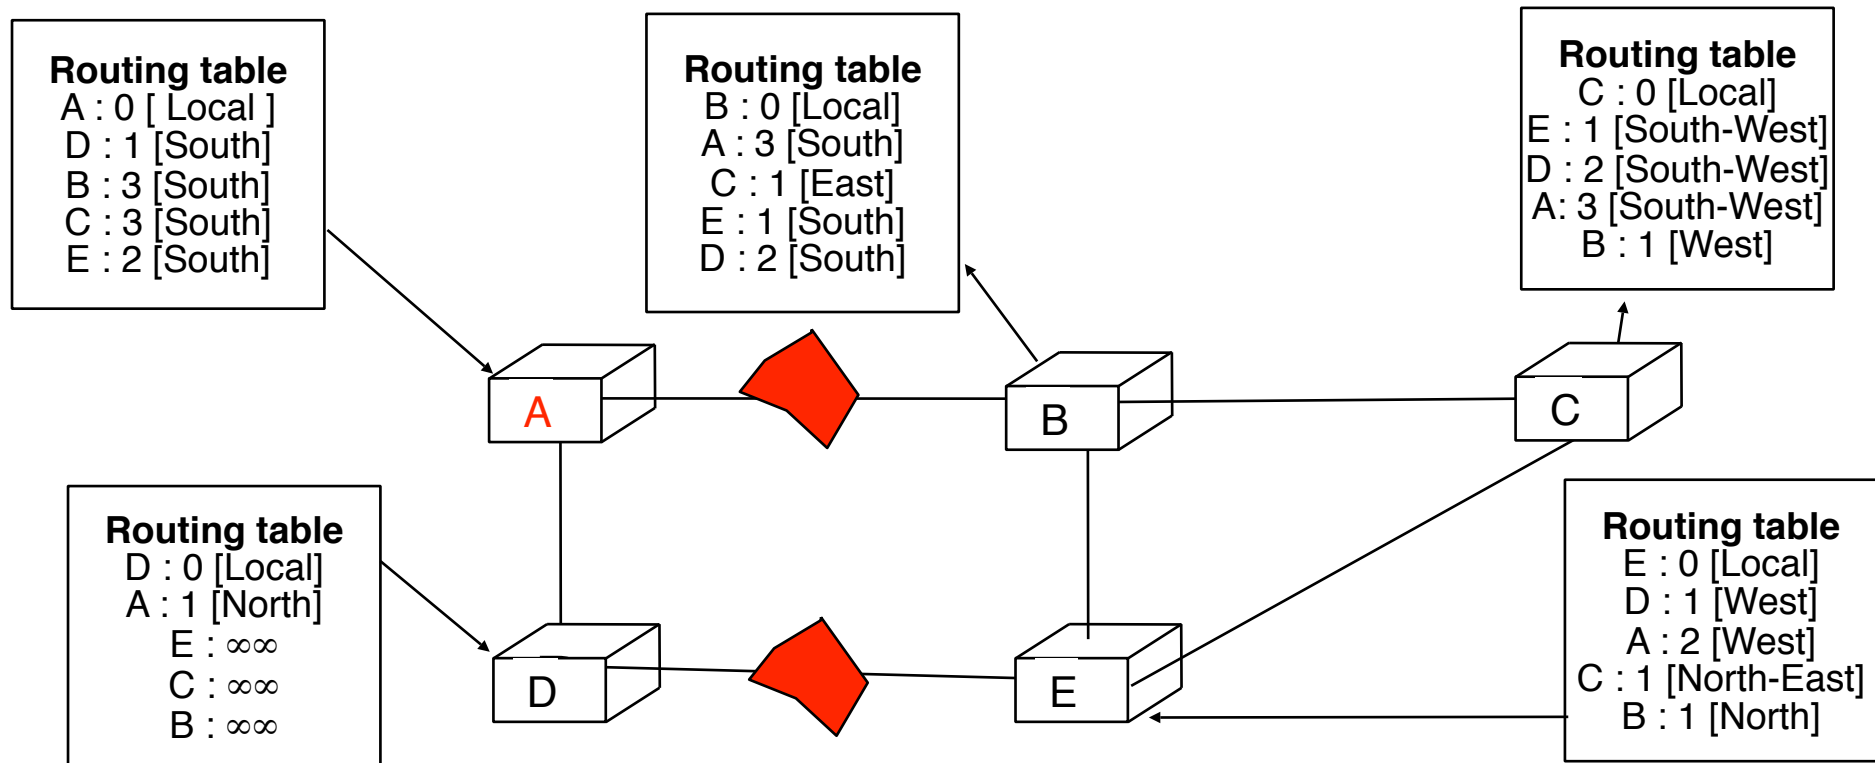
Instead of not advertising a route over the link from which it was learned, advertise it with an infinite cost

Pseudocode

```
Every N seconds:
  for each link=l
  { /* one different vector for each link */
    Vector=null;
    for each destination=d in R[]
    {
      if (R[d].link<>l)
      {
        Vector=Vector+Pair(d,R[d].cost);
      }
      else
      {
        Vector=Vector+Pair(d,∞);
      }
    }
    Send(Vector);
  }
```

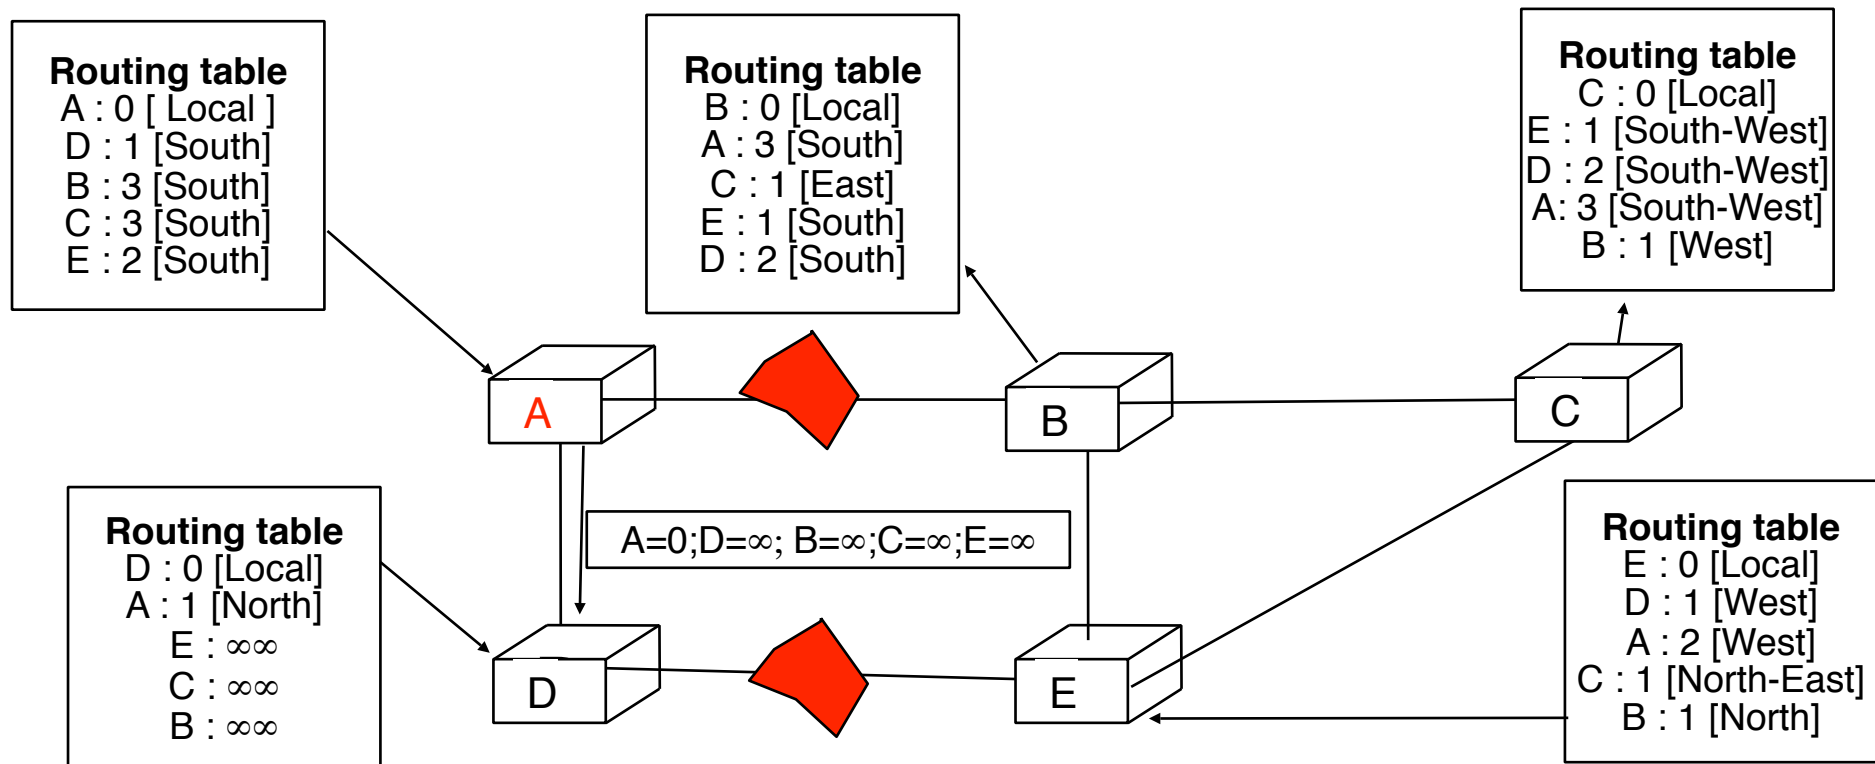
# Split horizon with poisoning (2)

Back to previous example

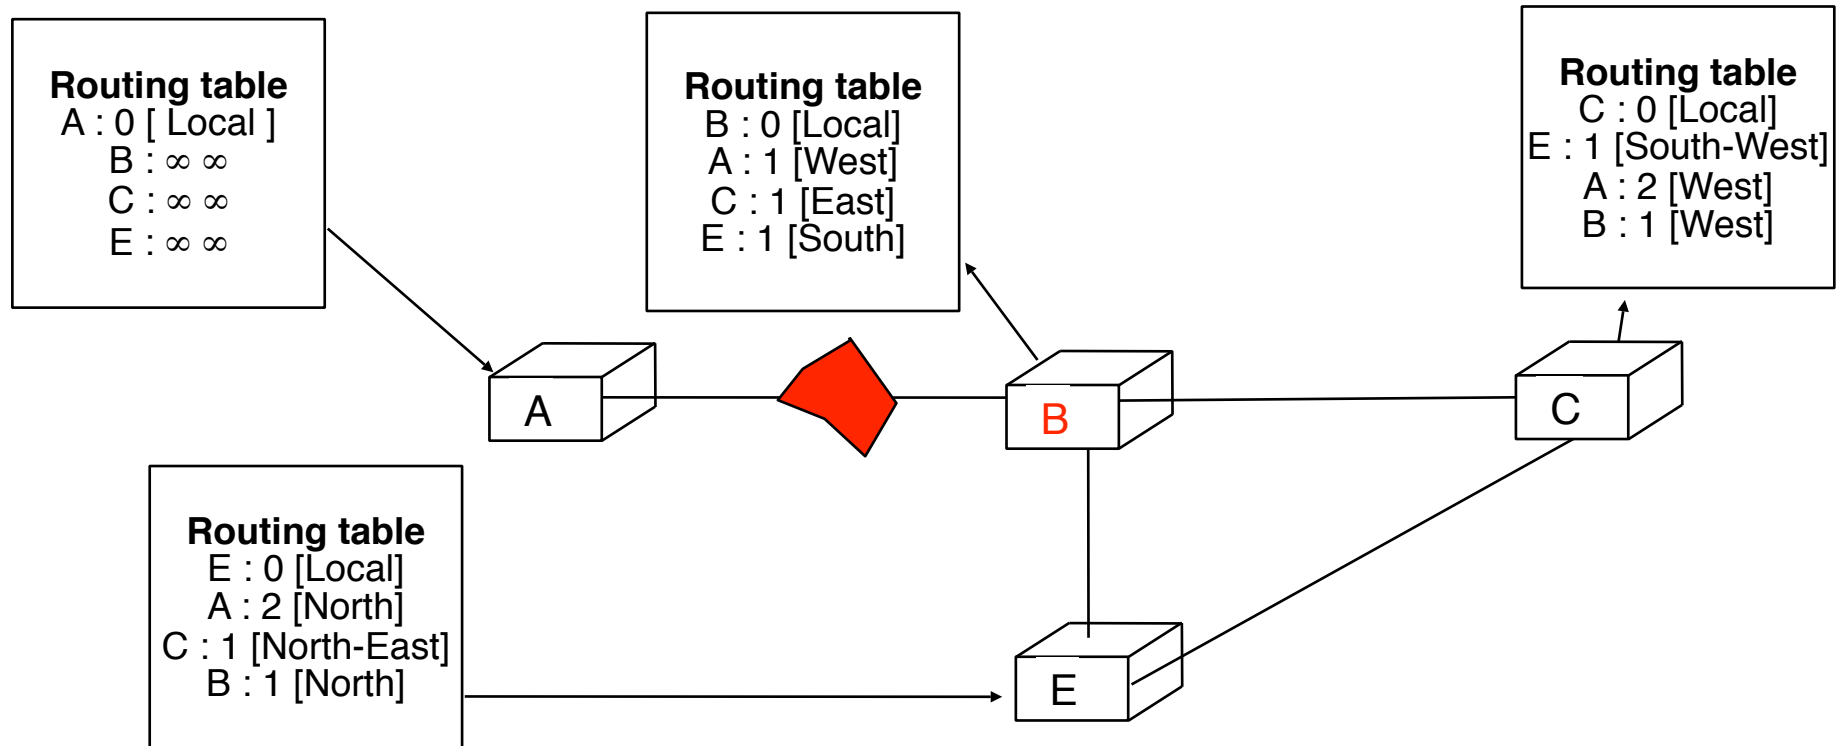


# Split horizon with poisoning (2)

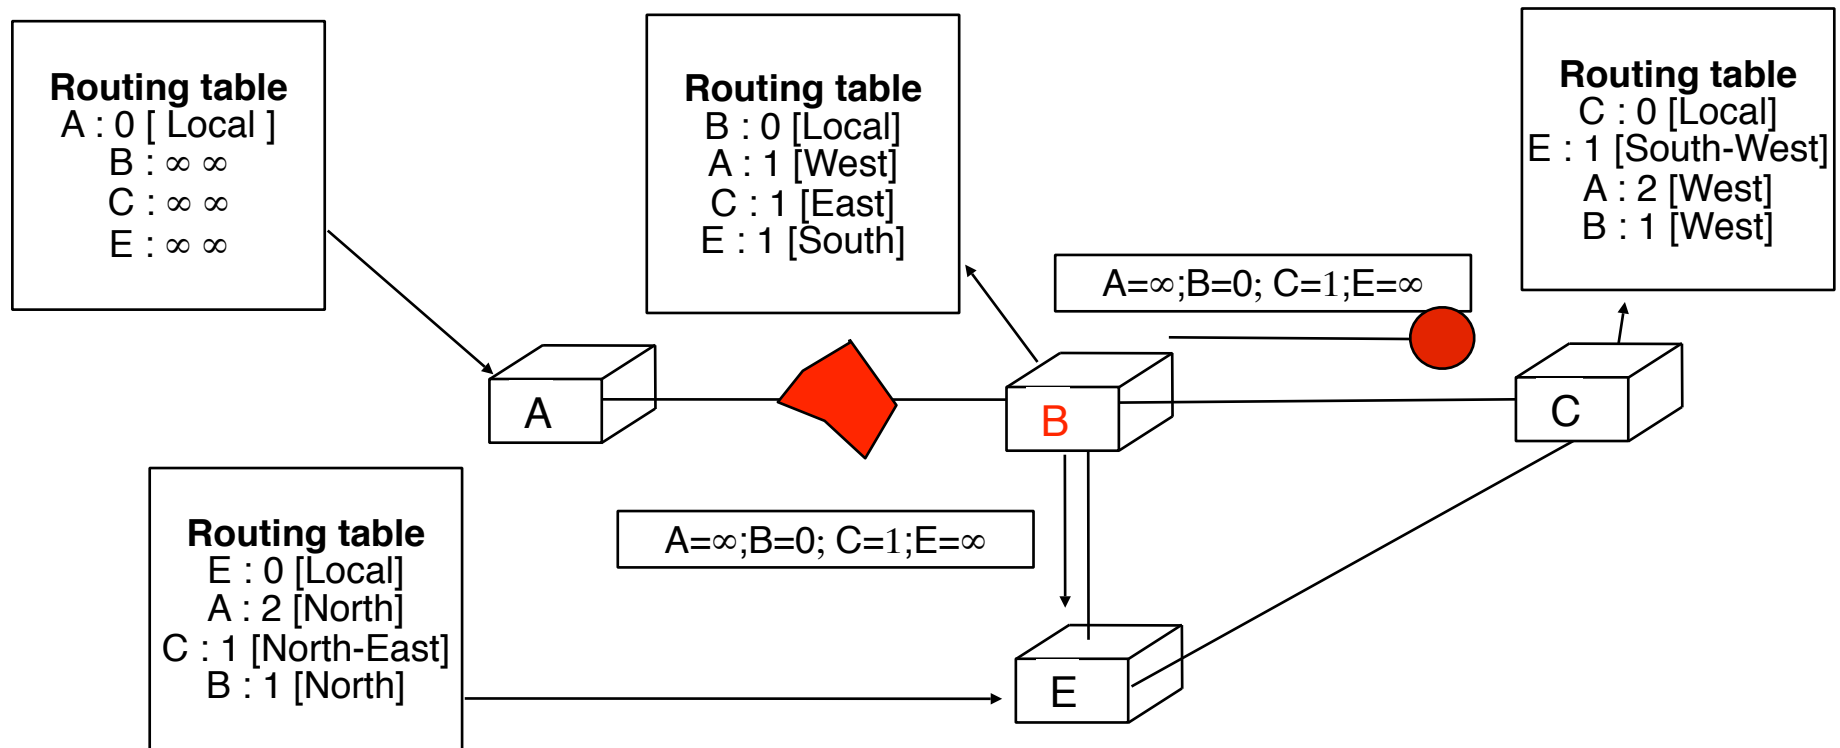
Back to previous example



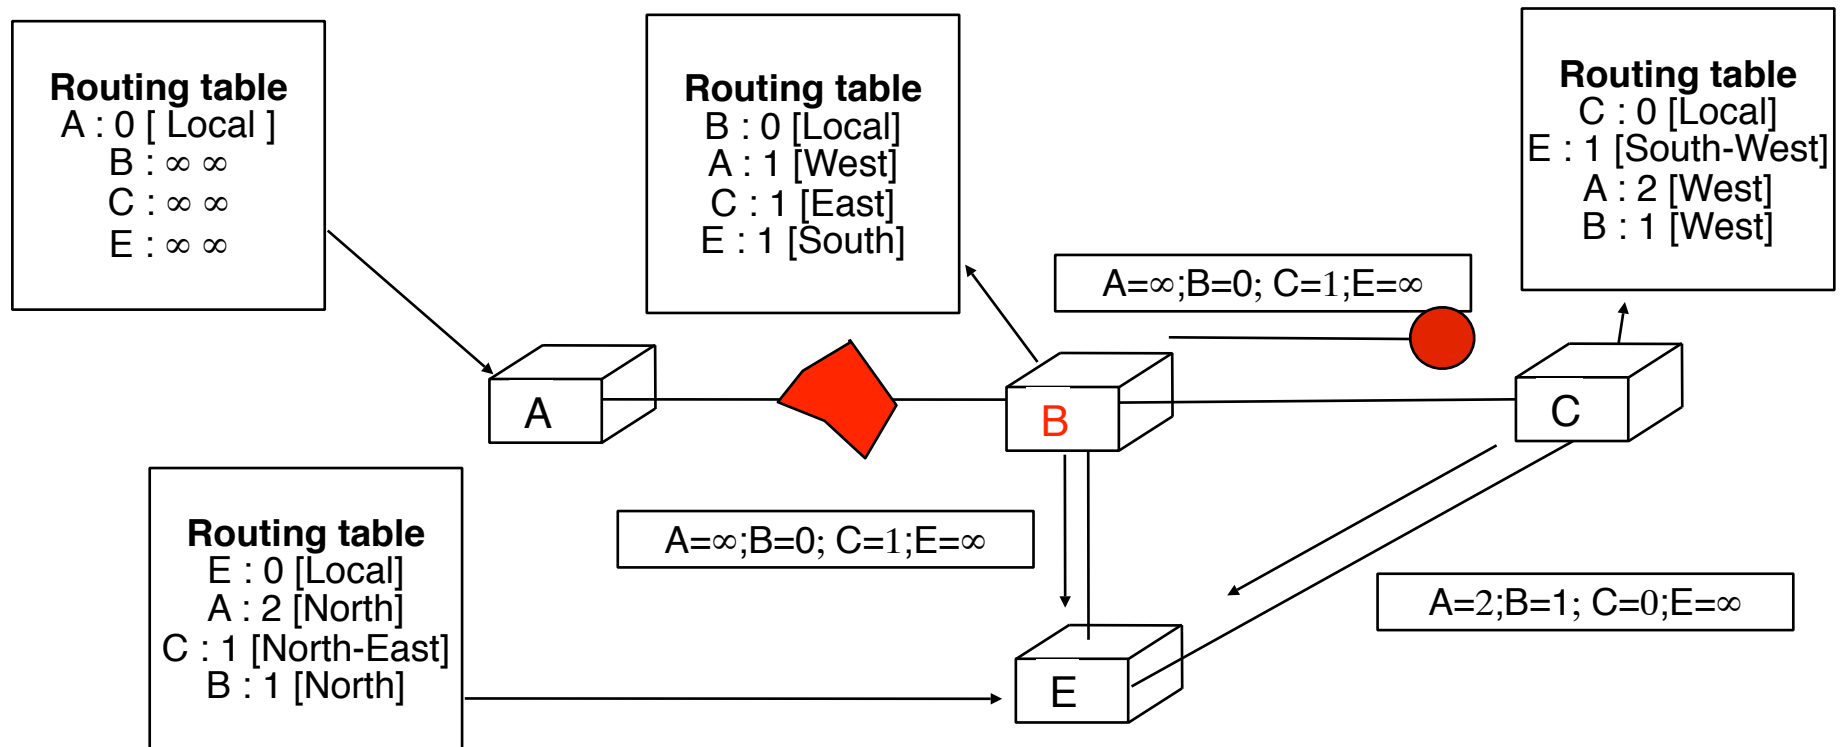
# Limitations to split horizon



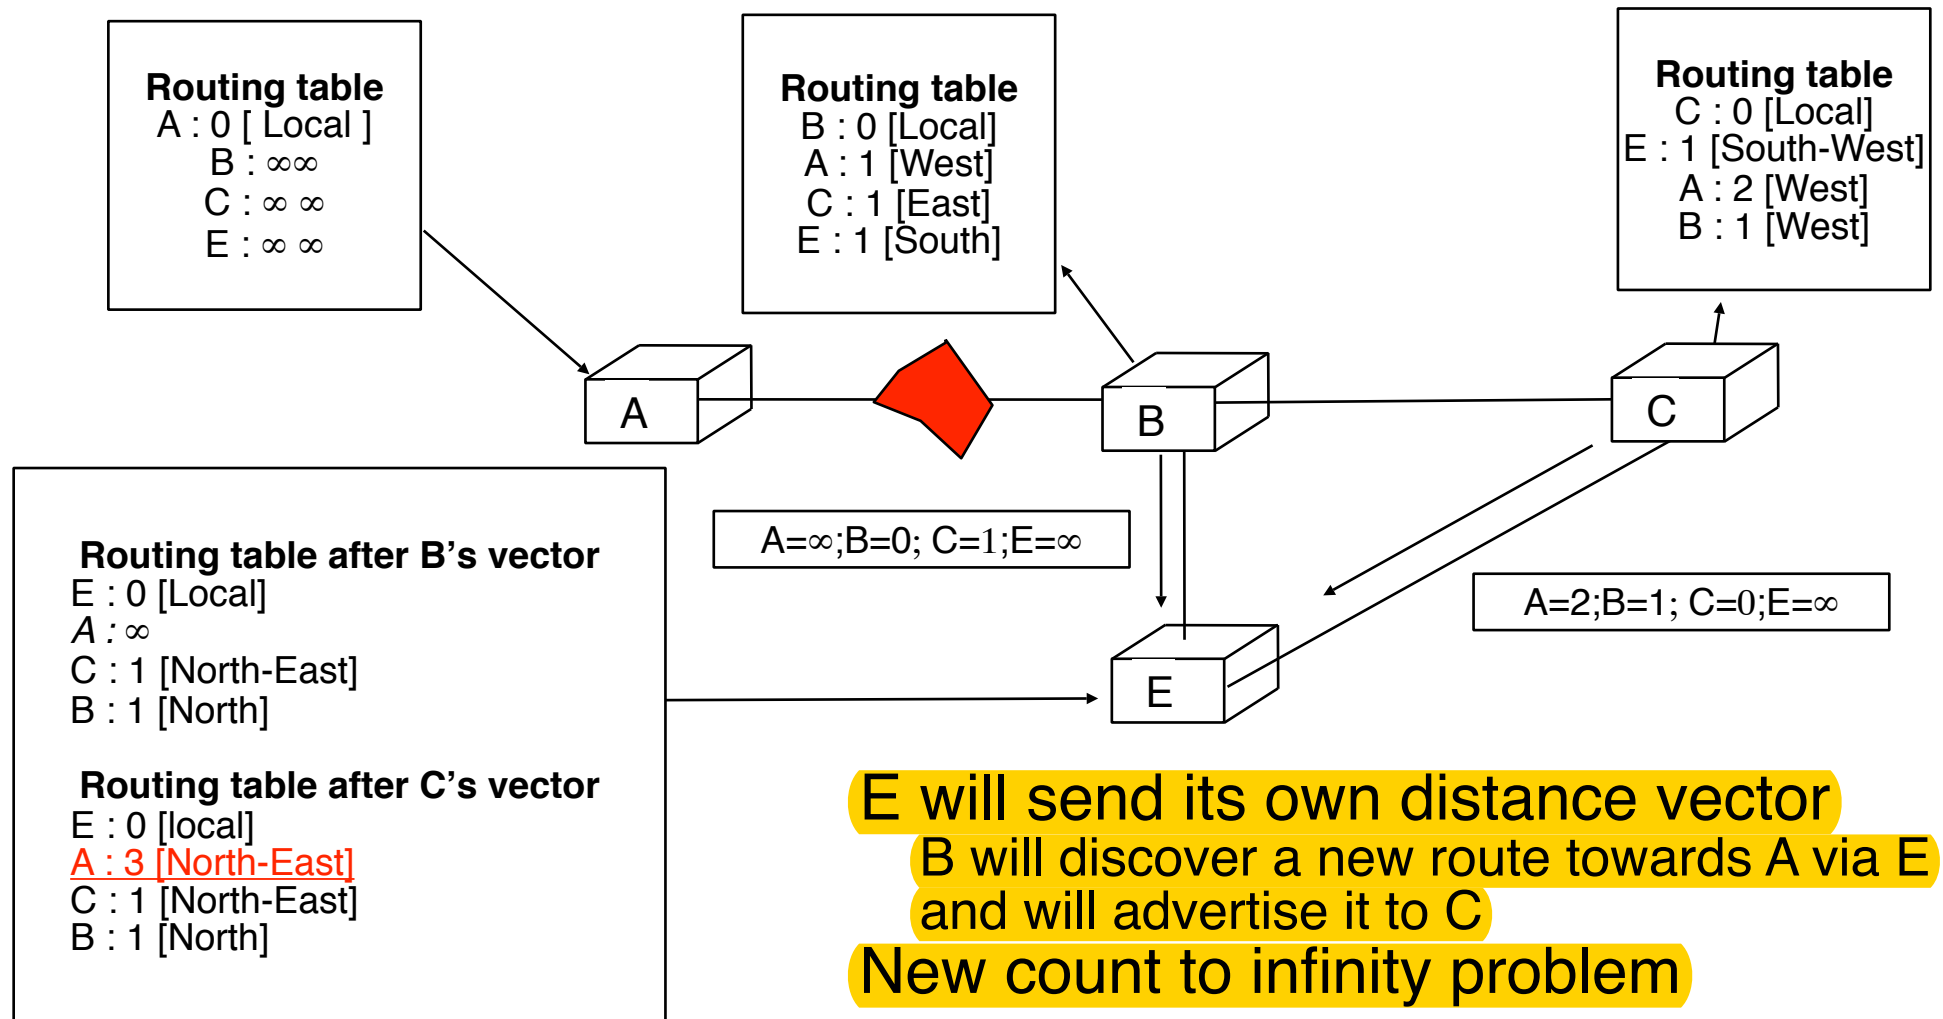
# Limitations to split horizon



# Limitations to split horizon



# Limitations to split horizon (2)



# Network layer

---

## Basics

### Routing

Static routing

Distance vector routing

→ Link state routing

## IP : Internet Protocol

## Routing in IP networks



# Link state routing

---

## Idea

Instead of distributing summaries of routing tables, wouldn't it be better to distribute network map ?

## How to build such as network map ?

Each router must discover its neighbours

It should be possible to associate a cost to each link since all links are not equal

Each router sends its local topology to all routes and assembles the information received from other routers

Routers build the network graph and used

Dijkstra's algorithm to compute shortest paths

# Neighbour discovery

How does a router discover its neighbours ?

By manual configuration


Unreliable and difficult to manage

By using HELLO packets

Every N seconds, each router sends a HELLO packet on each link with its address

Neighbours replay by sending their own address

Periodic transmission allows to verify that the link remains up and detect failures



# Neighbour discovery

How does a router discover its neighbours ?

By manual configuration


Unreliable and difficult to manage

By using HELLO packets

Every N seconds, each router sends a HELLO packet on each link with its address

Neighbours replay by sending their own address

Periodic transmission allows to verify that the link remains up and detect failures



# Neighbour discovery

How does a router discover its neighbours ?

By manual configuration


Unreliable and difficult to manage

By using HELLO packets

Every N seconds, each router sends a HELLO packet on each link with its address

Neighbours replay by sending their own address

Periodic transmission allows to verify that the link remains up and detect failures



# How to determine link costs ?

---

## Principle

- one cost is associated with each link direction

## Commonly configured link costs

### Unit cost

- simplest solution but only suitable for homogeneous networks

### Cost depends on link bandwidth

- high cost for low bandwidth links
- low cost for high bandwidth links

### Cost depends on link delays

- often used to avoid satellite links

## Cost based on measurements

### Use HELLO to measure link rtt

- allows to track link load, but be careful if the measurement is not stable enough as each delay change will cause a topology change ...

# Assembling the network topology

---

## How to assemble the network topology

- By receiving HELLOs, each routers builds its local part of the network map

- Each router summarises its local topology inside one link state packet that contains

  - router identification

  - pairs (neighbour identification, cost to reach neighbour)

## When should a router send its link state packet ?

- in case of modification to its local topology
  - allows to inform all other routers of the change

- Every N seconds

  - allows to refresh information in all routers and makes sure that if an invalid information was stored on a router due to memory errors it will not remain in the router forever

# How to distribute the link state packets ?

---

# How to distribute the link state packets ?

---

## Naive solution

Each router sends one packet to each other router in the network

This solution can only work if

- All routers know the address of all other routers in network

- All routers already have routing tables that allow them to forward packets to any destination



# How to distribute the link state packets ?

---

## Naive solution

Each router sends one packet to each other router in the network

This solution can only work if

- All routers know the address of all other routers in network

- All routers already have routing tables that allow them to forward packets to any destination

## Realistic solution

Does not rely on pre-existing routing tables

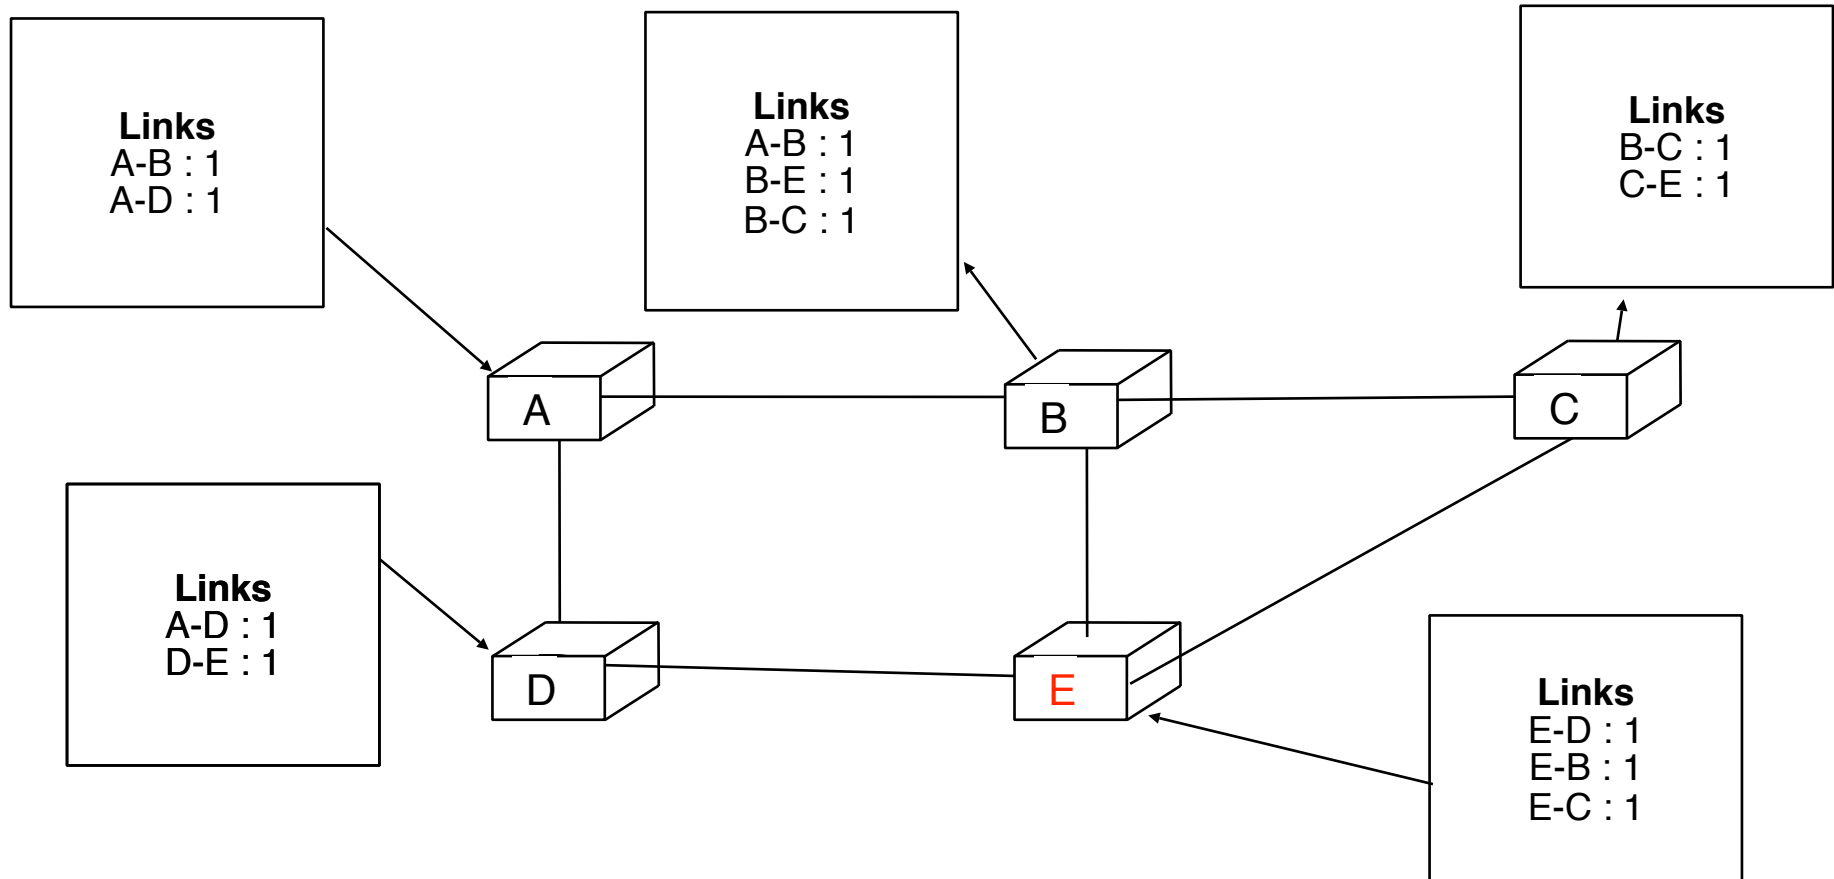
Each router must receive entire topology

### First solution

Each router sends local topology in link state packet and sends it to all its outgoing links

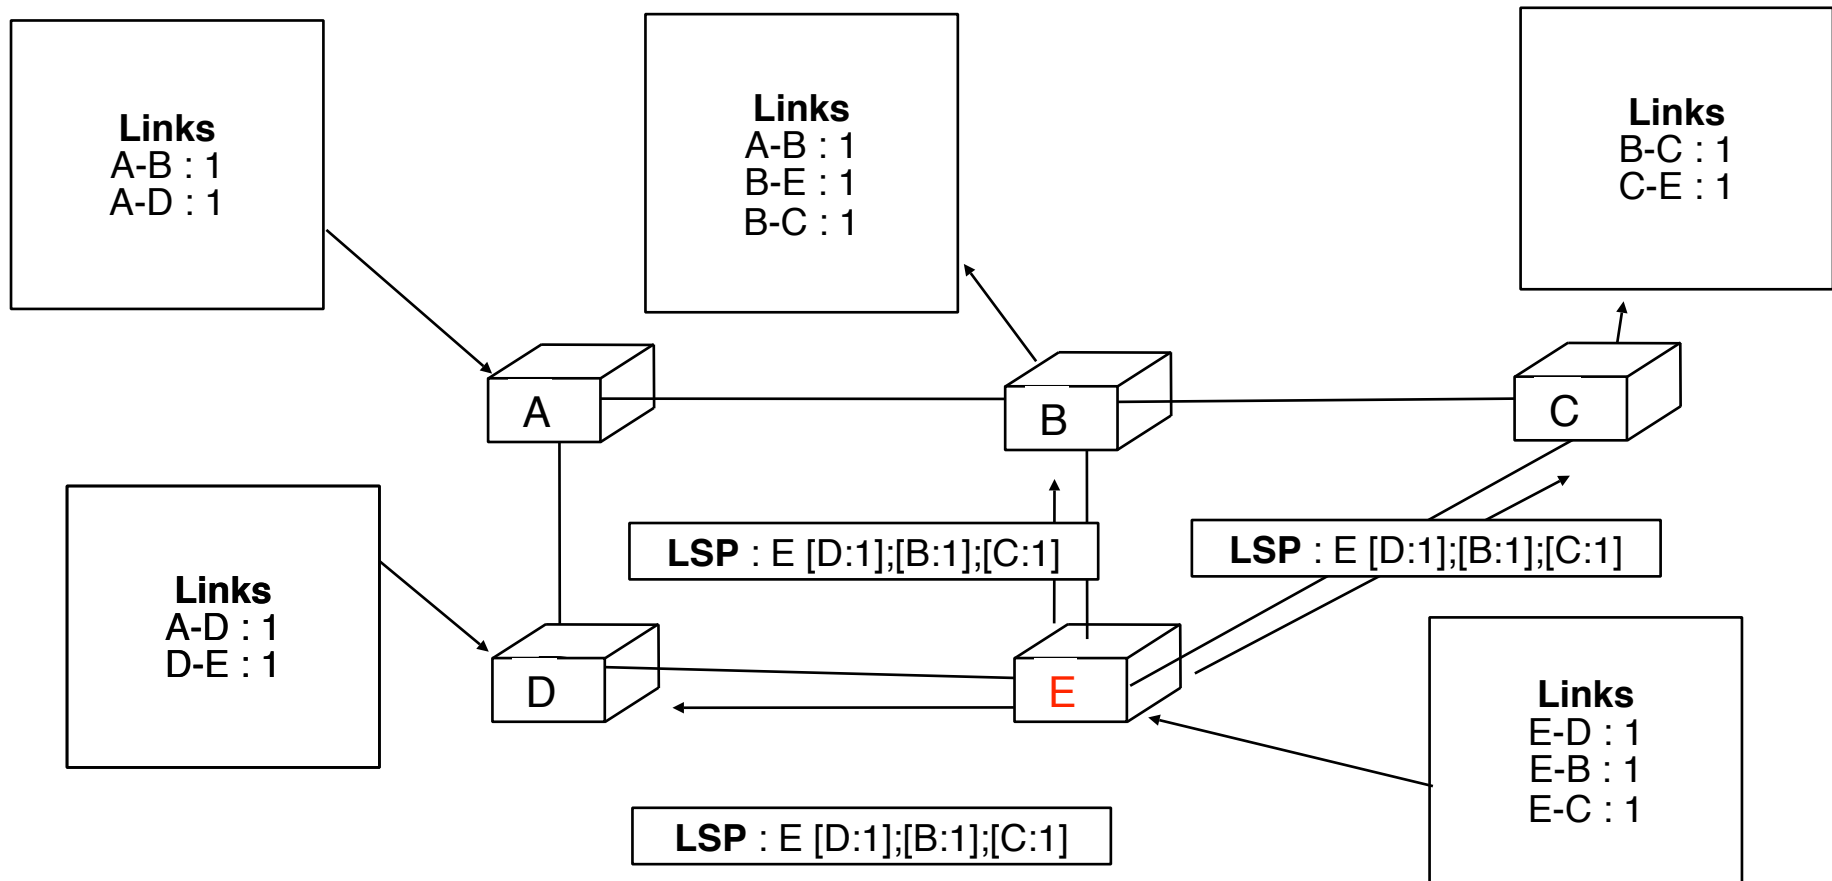
When a router receives an LSP, it forwards it to all its outgoing links except the link from which it received it

# LSP flooding




Assumes that all links have a unit cost

# LSP flooding

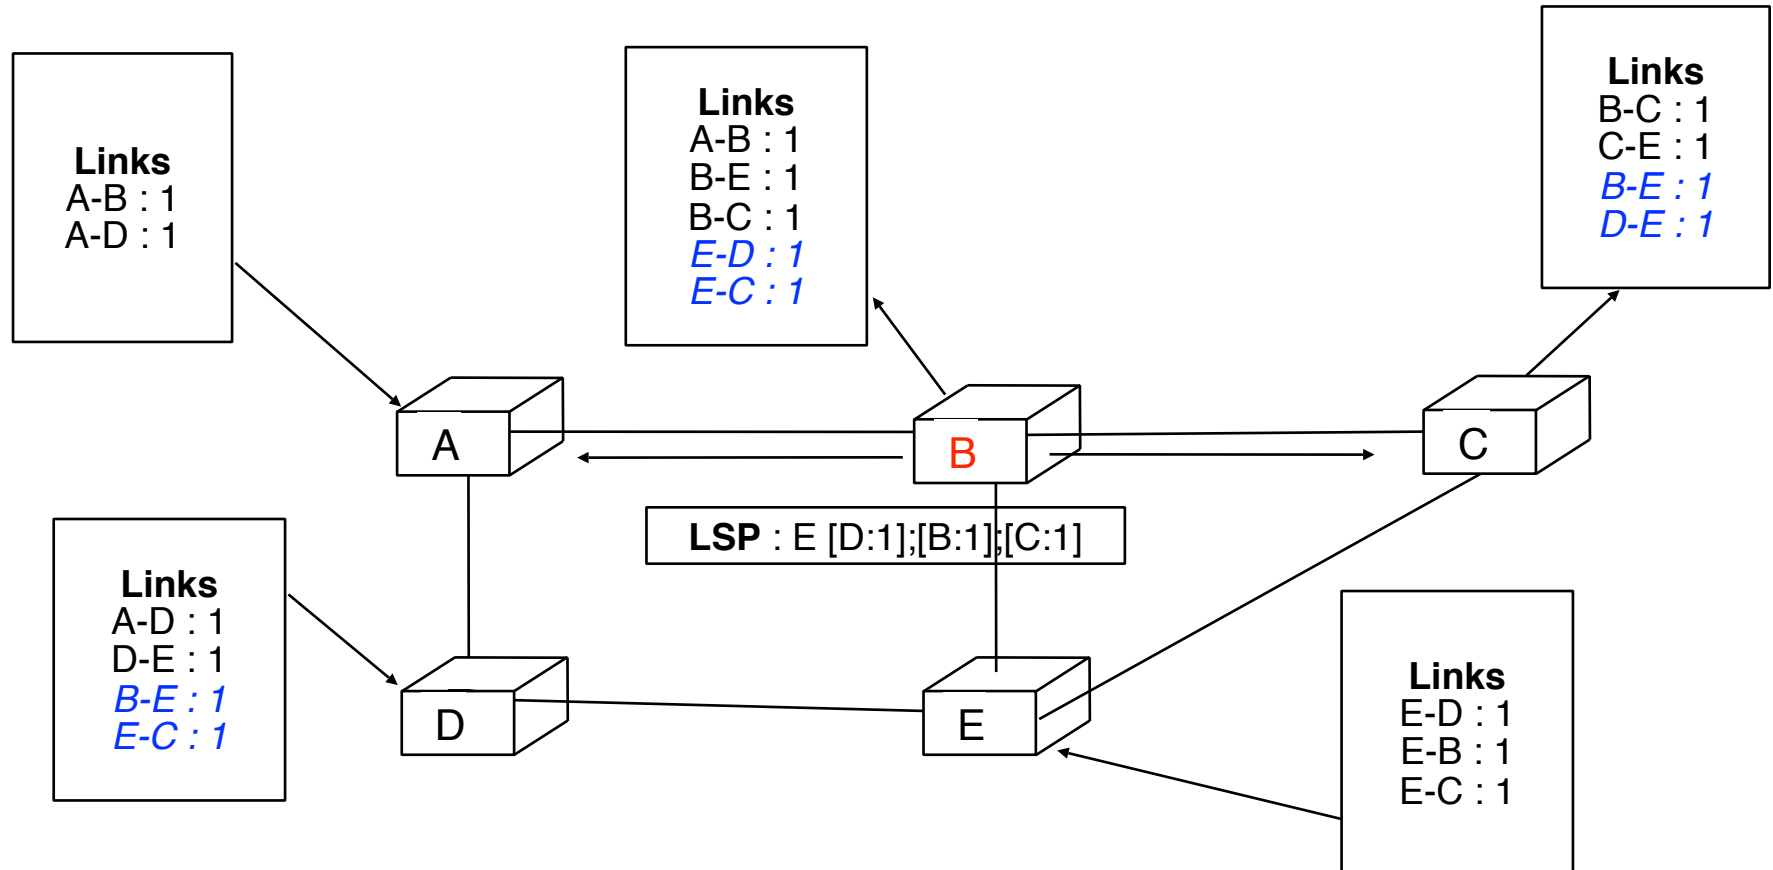


Assumes that all links have a unit cost

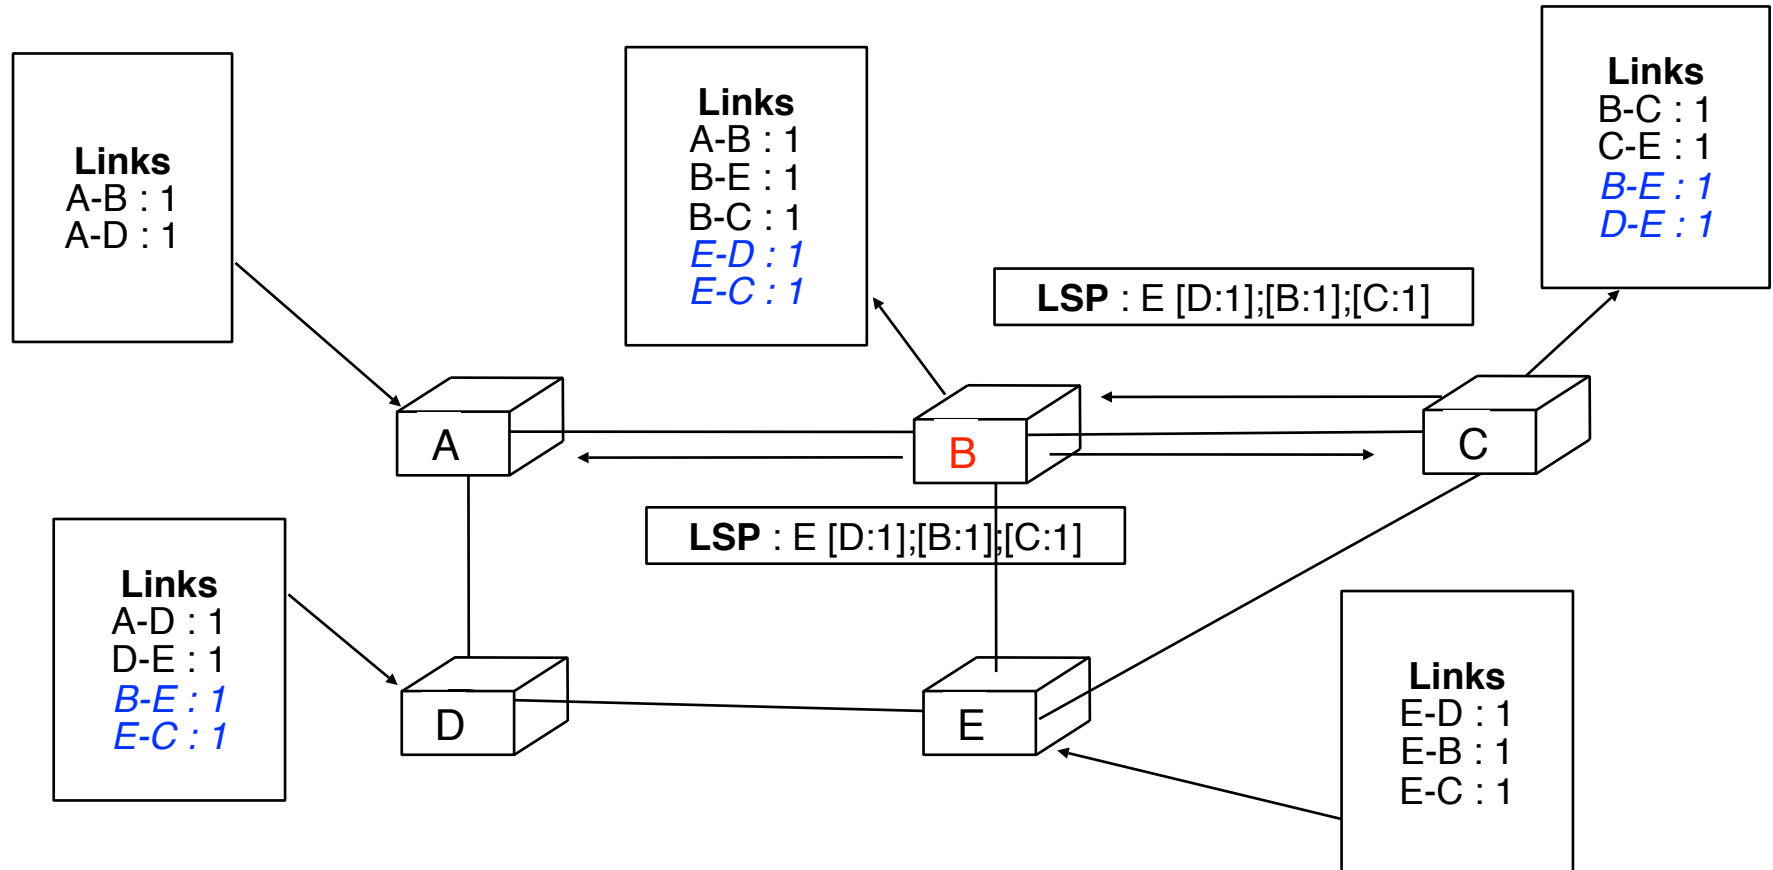
# LSP flooding (2)



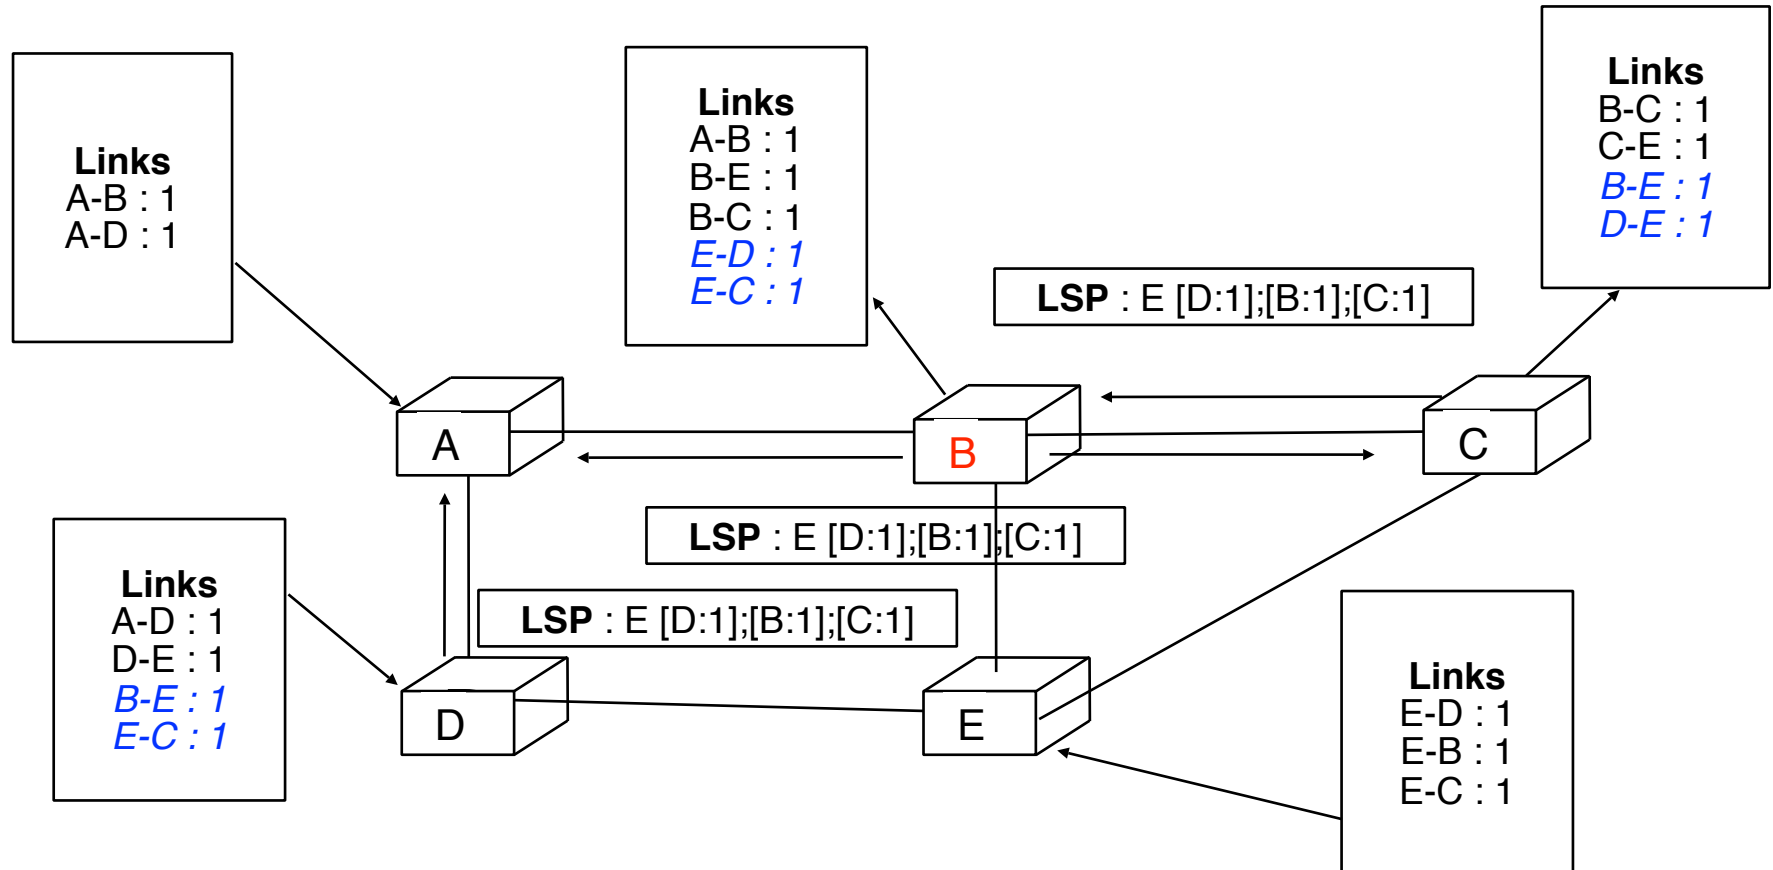
# LSP flooding (2)




# LSP flooding (2)



# LSP flooding (2)



# LSP flooding (2)



How to ensure that an LSP will not loop ?



# LSP flooding (3)

---

How to avoid LSP flooding loops ?

A router should not reflood an LSP that it has already and flooded

Solution

LSP contents

- sequence number

  - incremented every time an LSP is generated by a router

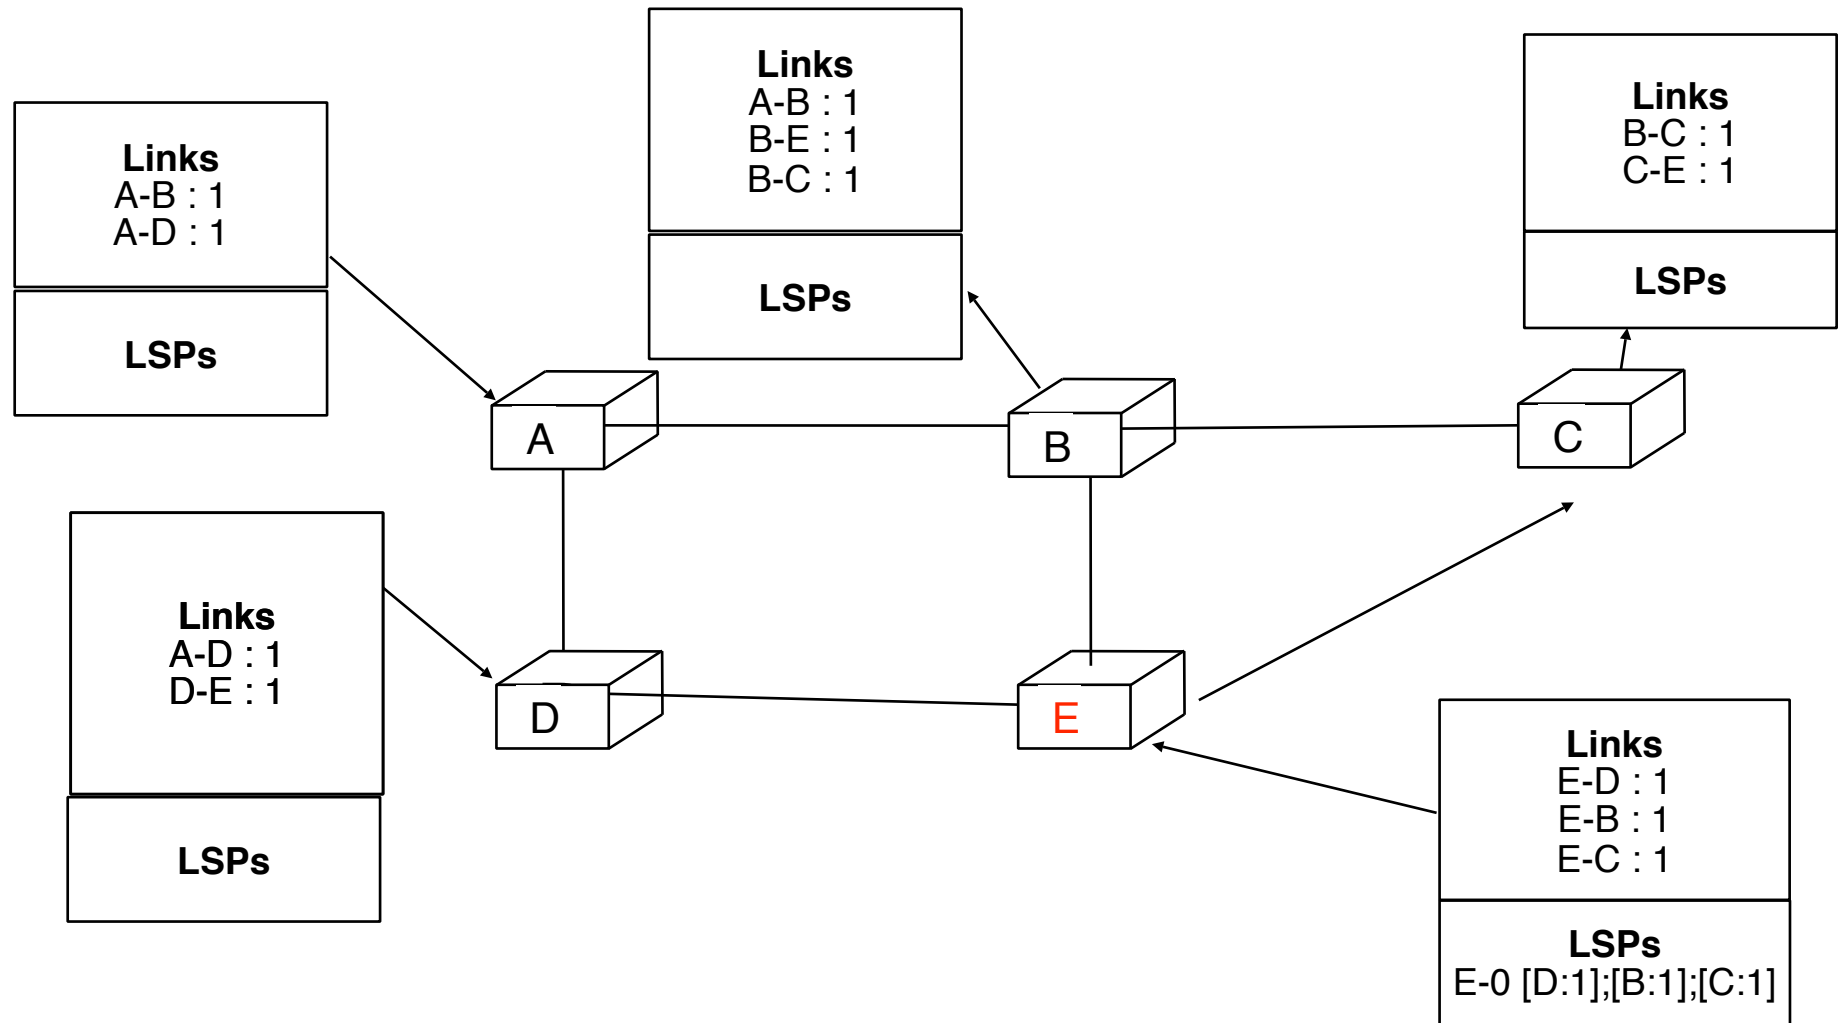
- address of LSP originator

- pairs address:distance for all neighbours of the originator

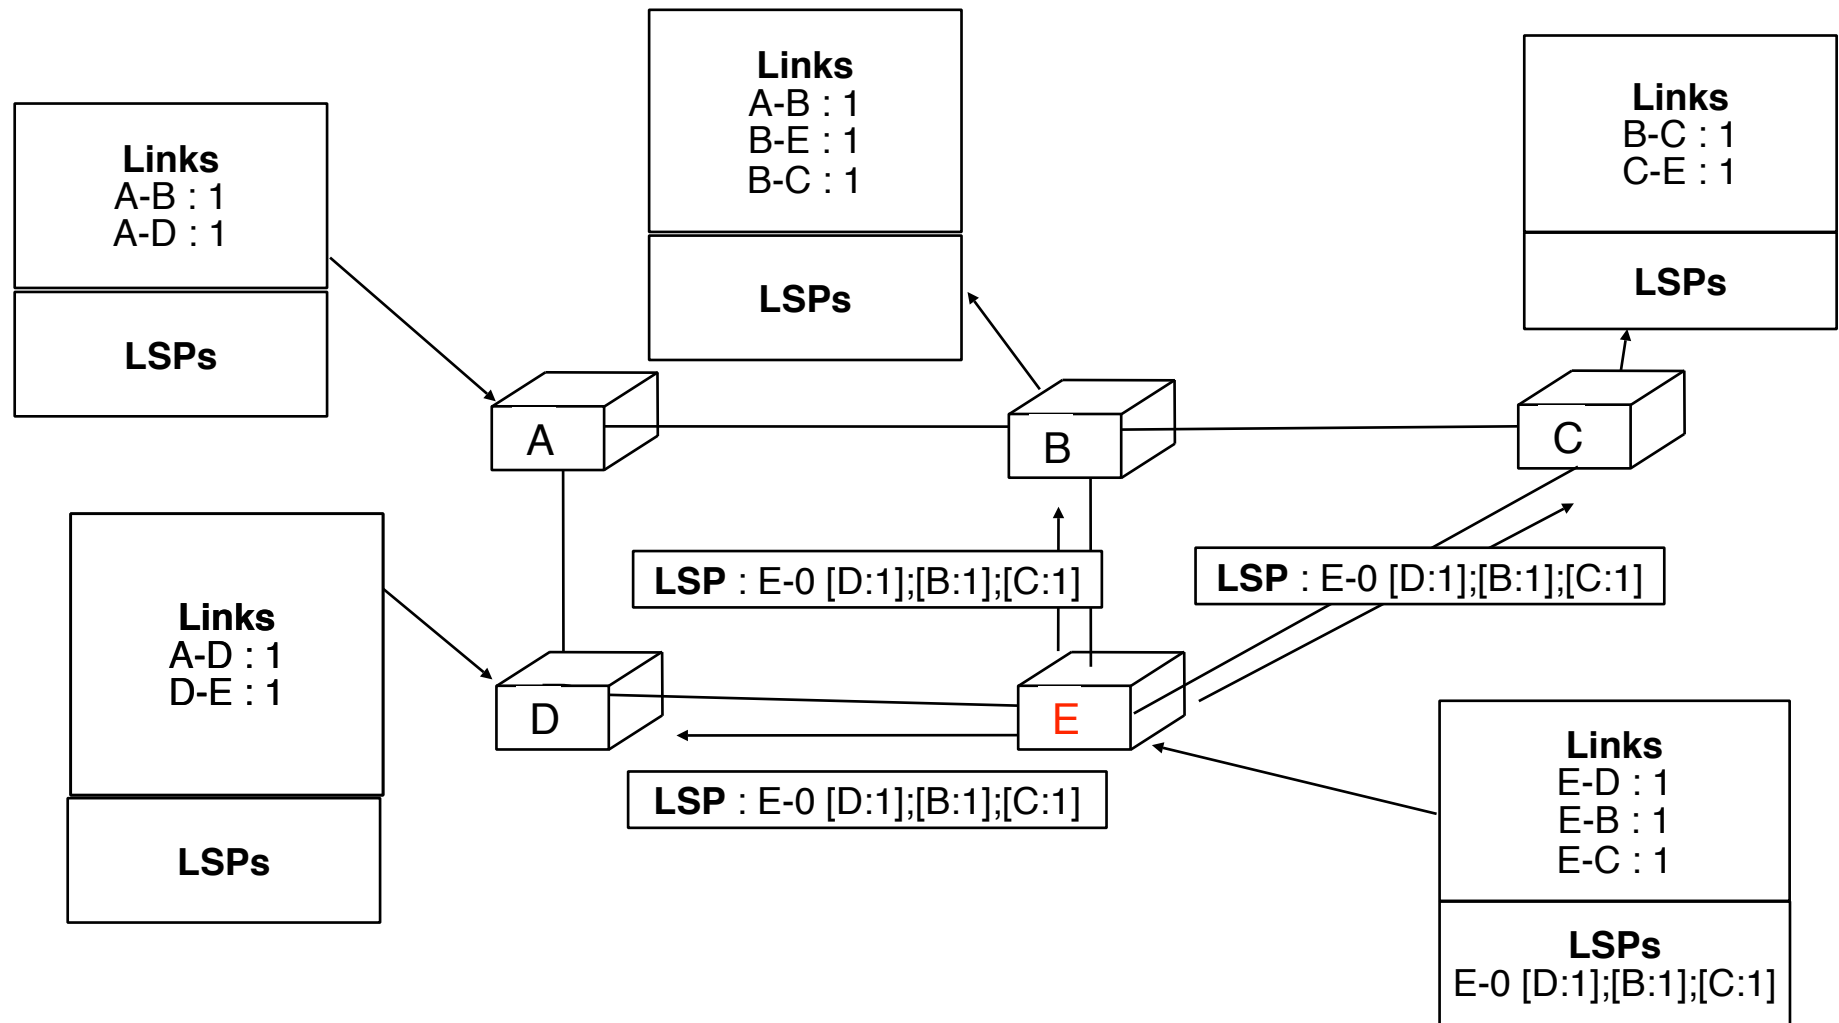
Each router must store the last LSP received from each router of the network

A received LSP is processed and flooded only if it is more recent than the LSP stored in the LSDB

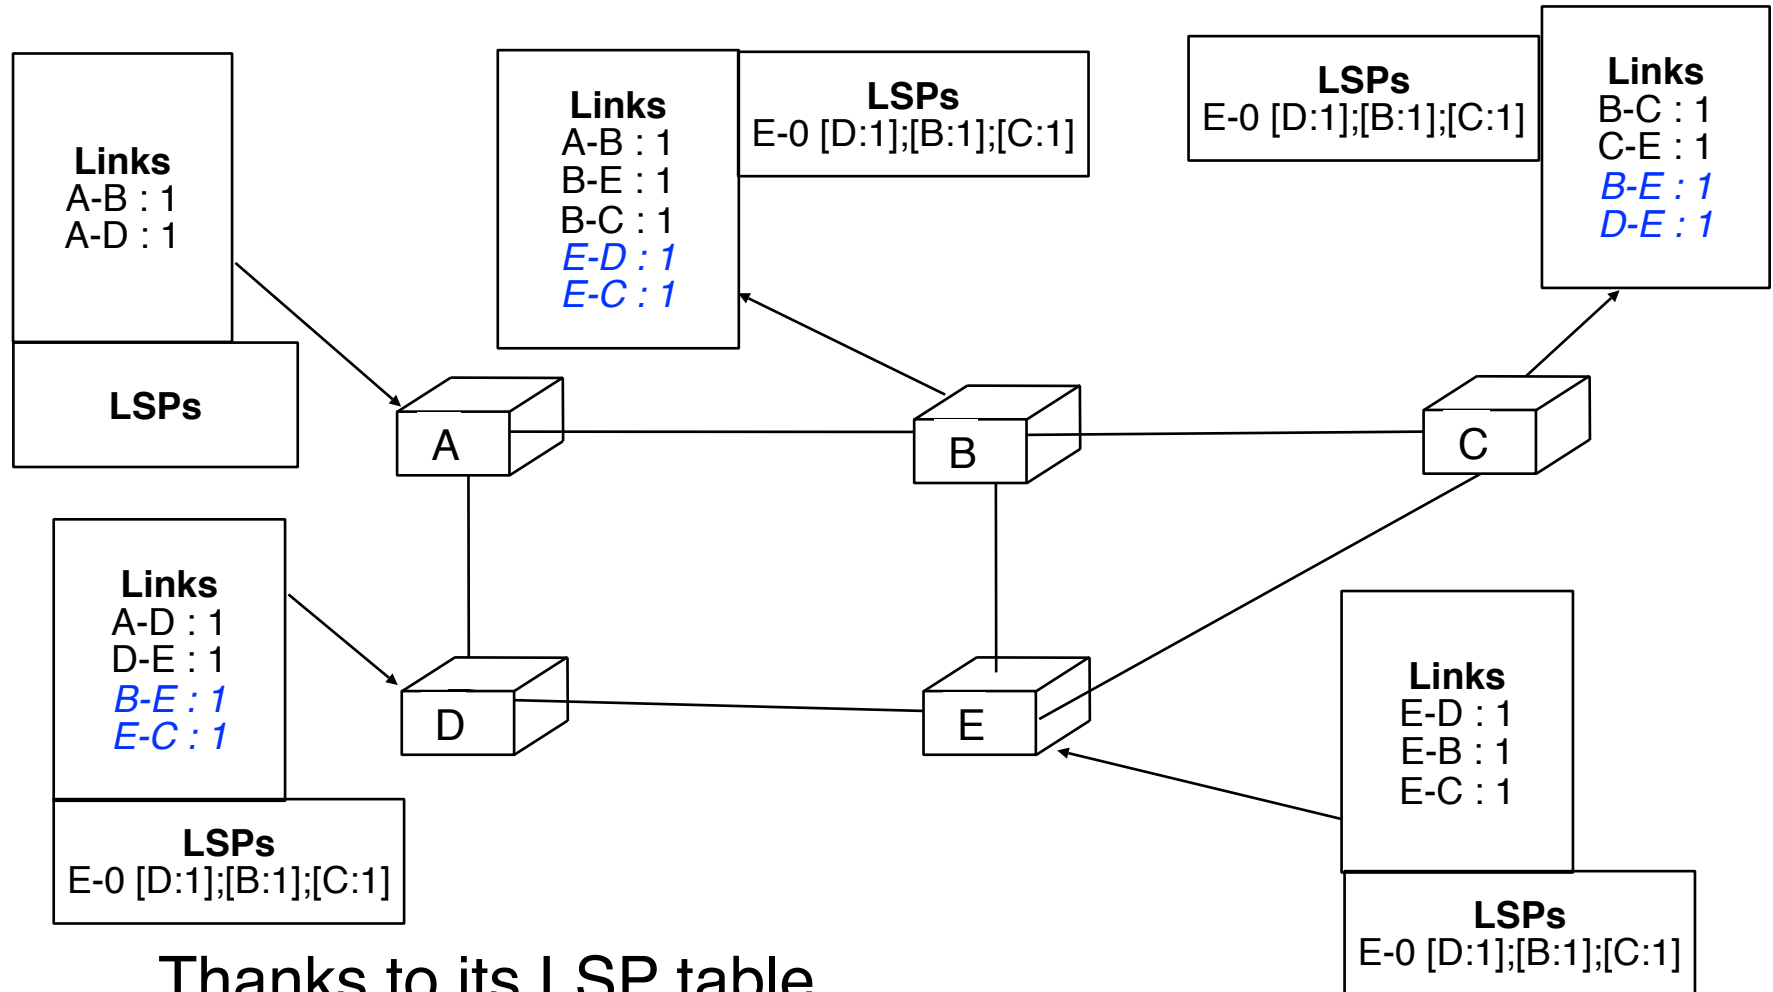
# LSP flooding (4)



# LSP flooding (4)



# LSP flooding (5)

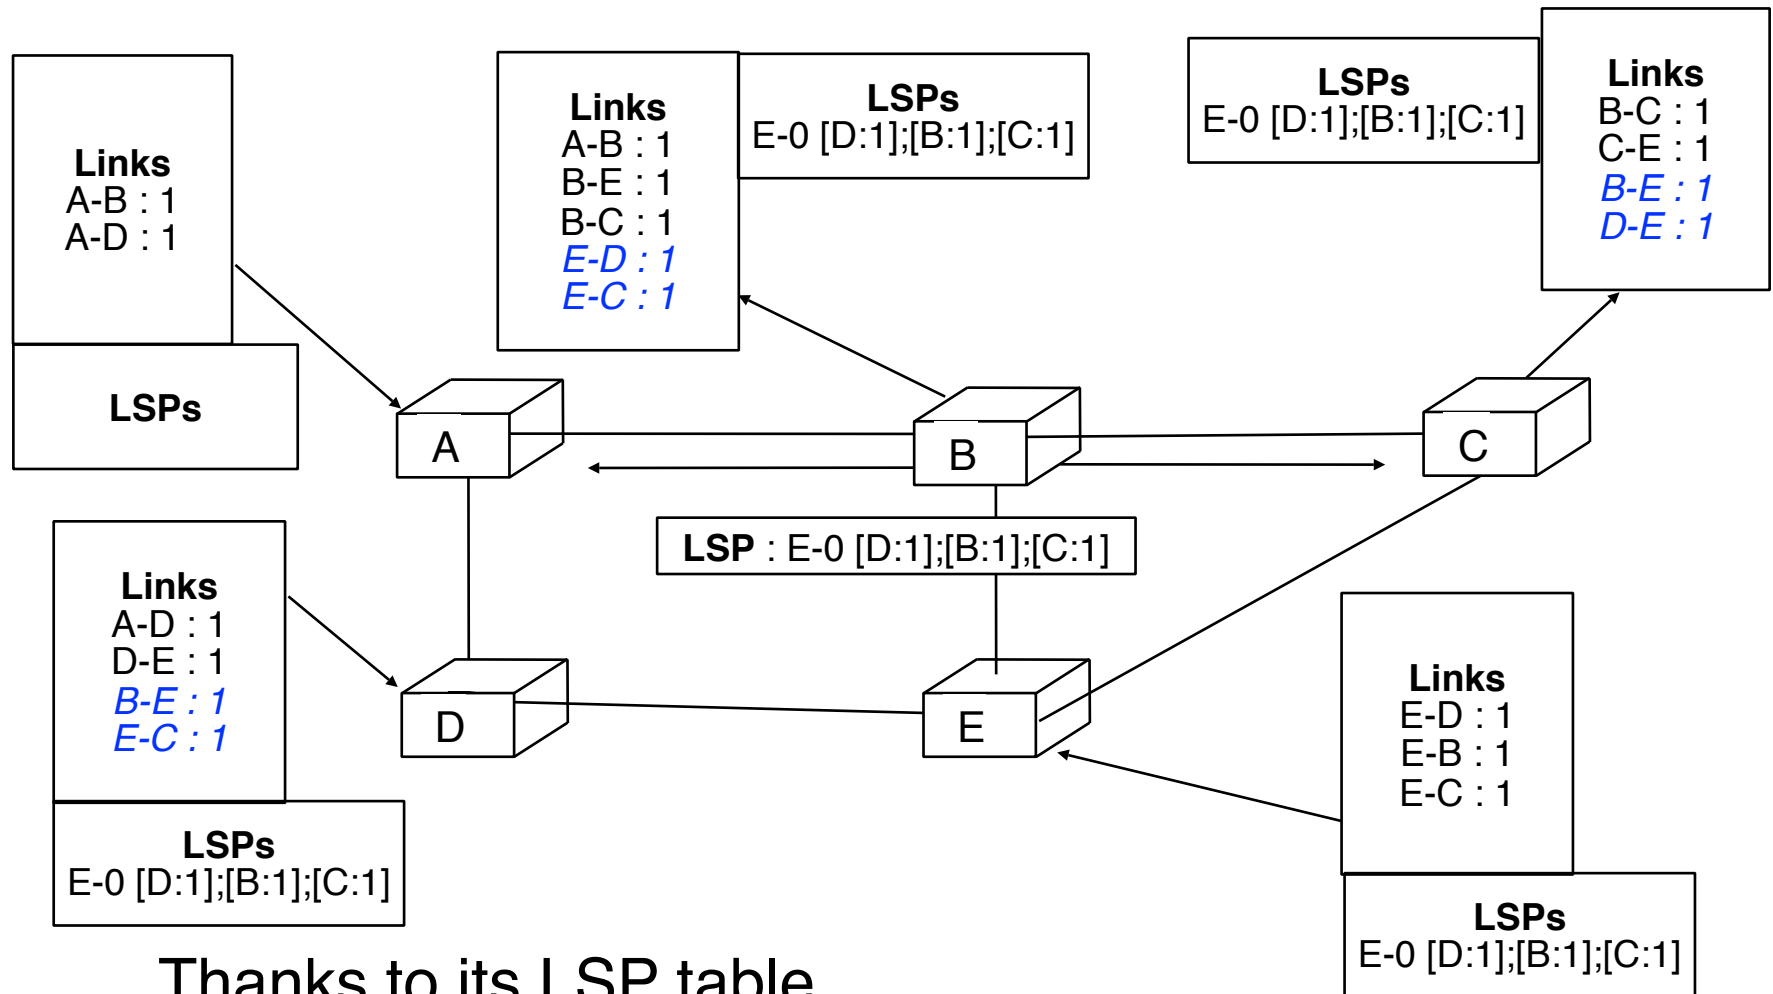


Thanks to its LSP table

A can detect that it received same LSP via B and D

C can detect that it received same LSP via B and E

# LSP flooding (5)

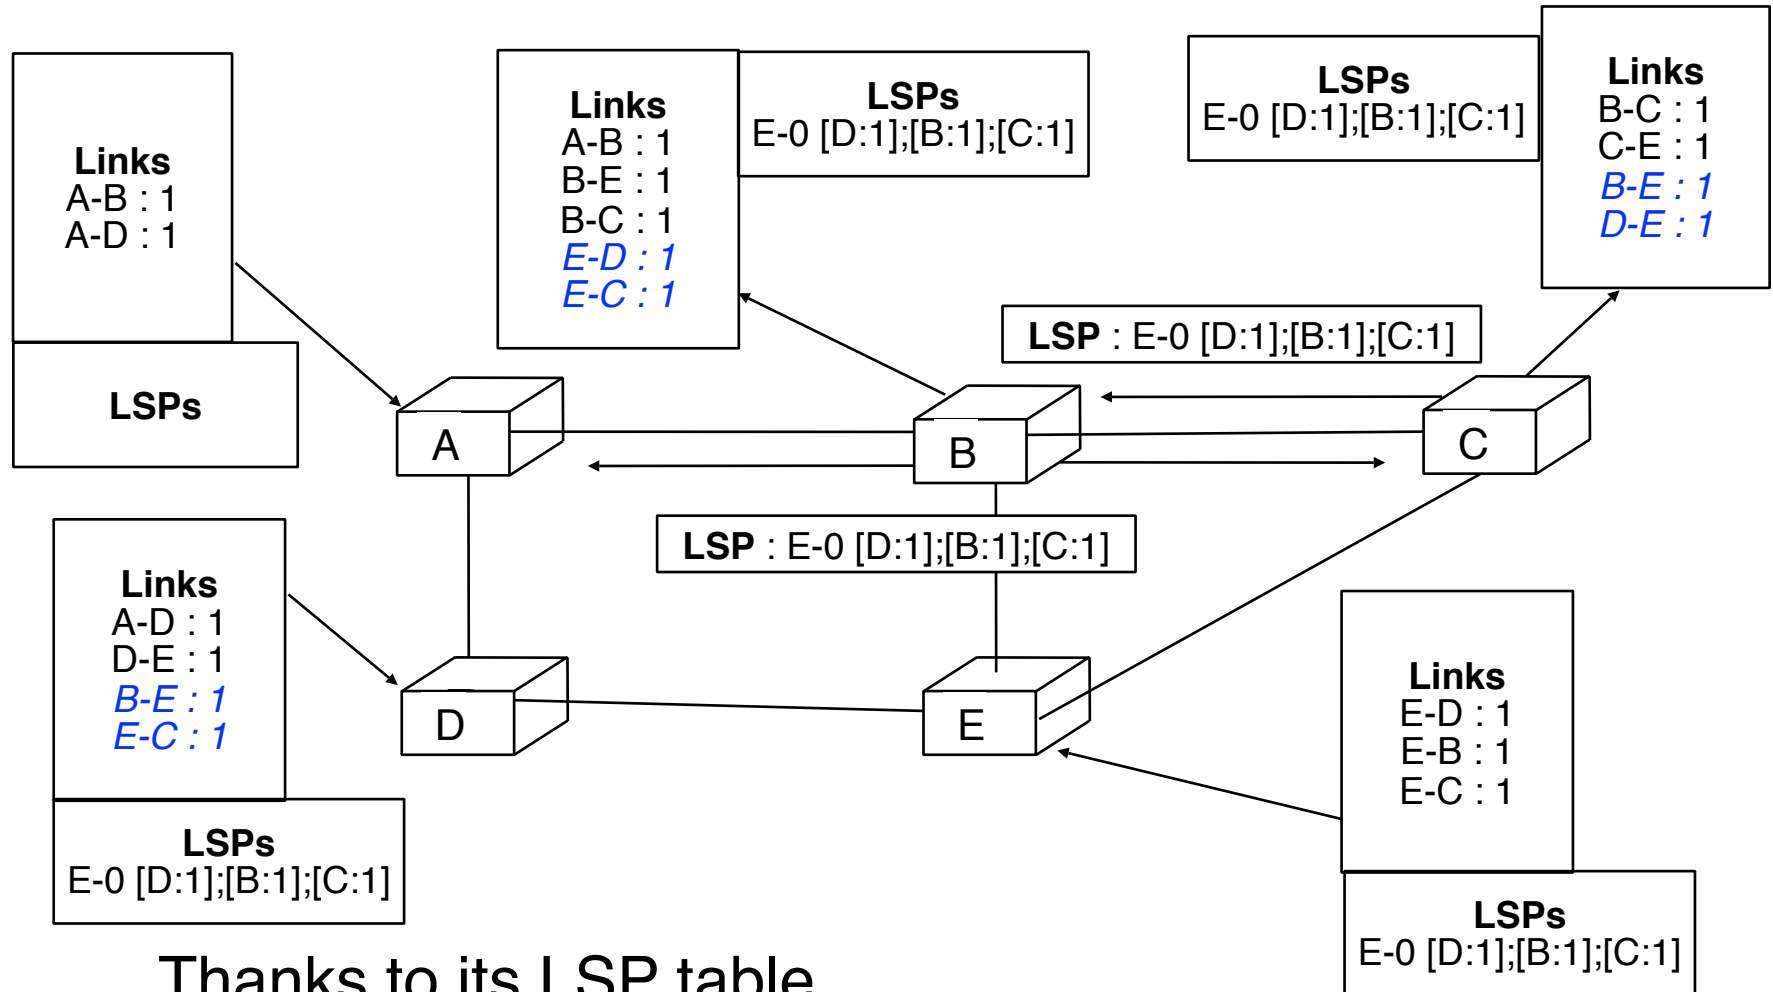


Thanks to its LSP table

A can detect that it received same LSP via B and D

C can detect that it received same LSP via B and E

# LSP flooding (5)

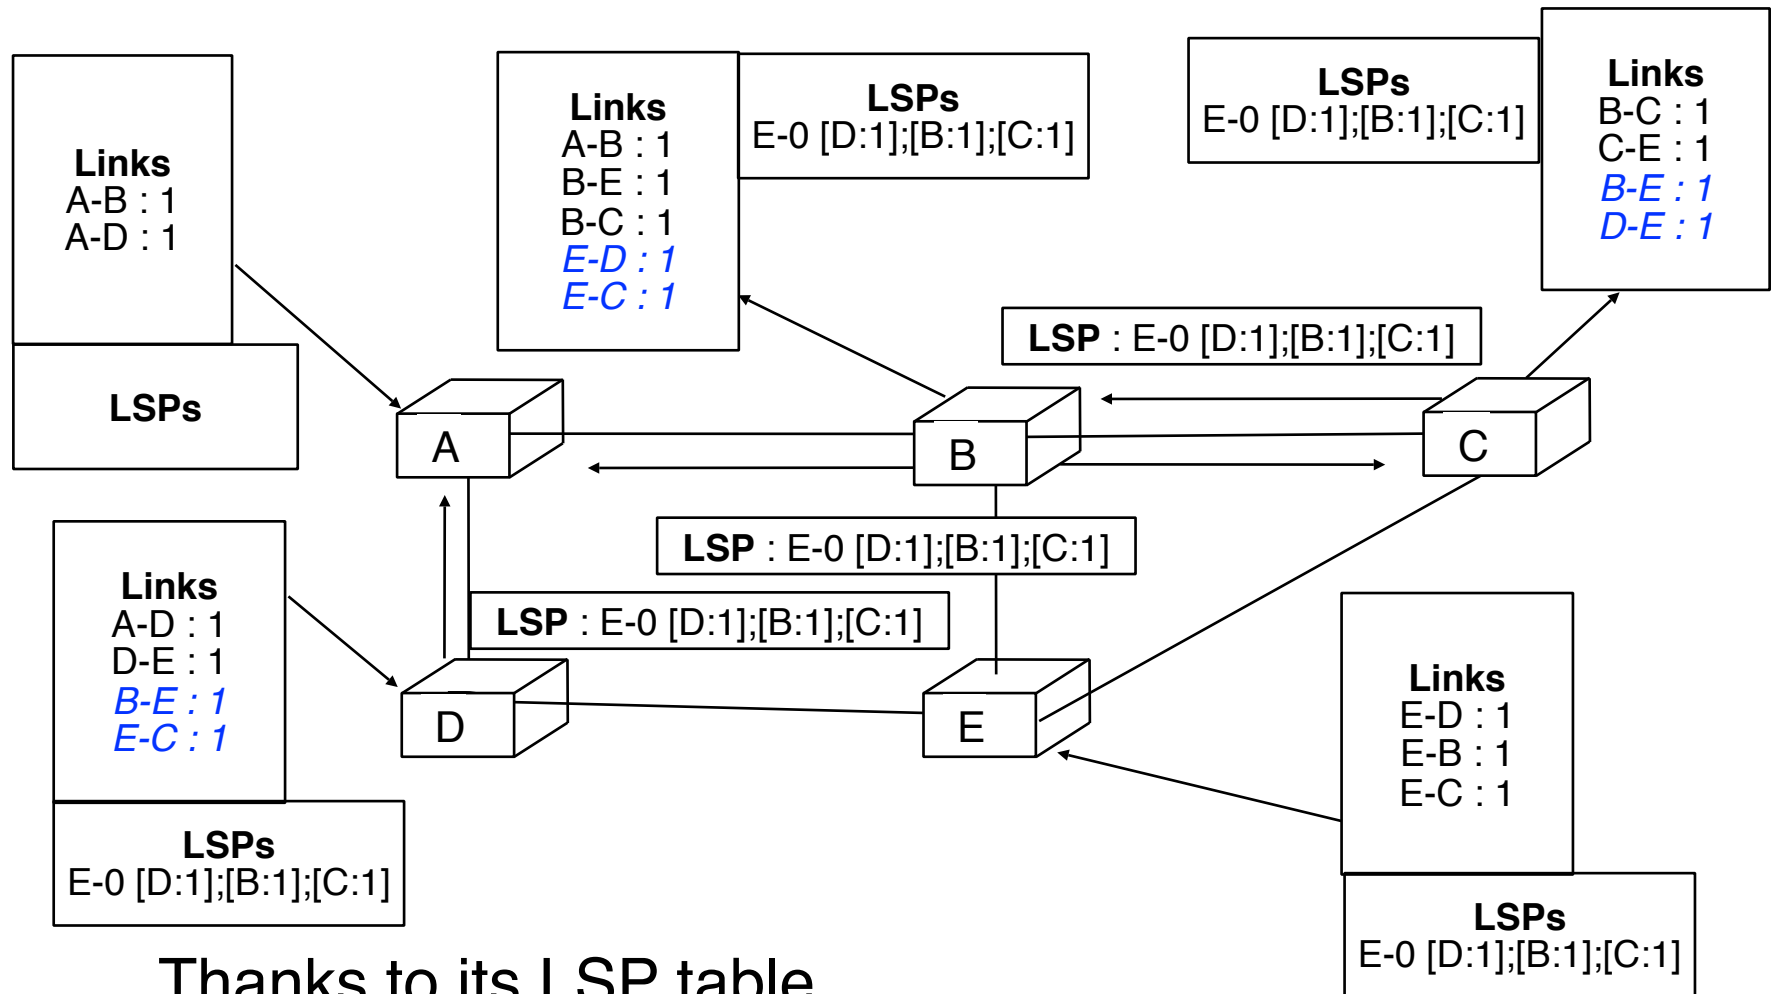


Thanks to its LSP table

A can detect that it received same LSP via B and D

C can detect that it received same LSP via B and E

# LSP flooding (5)

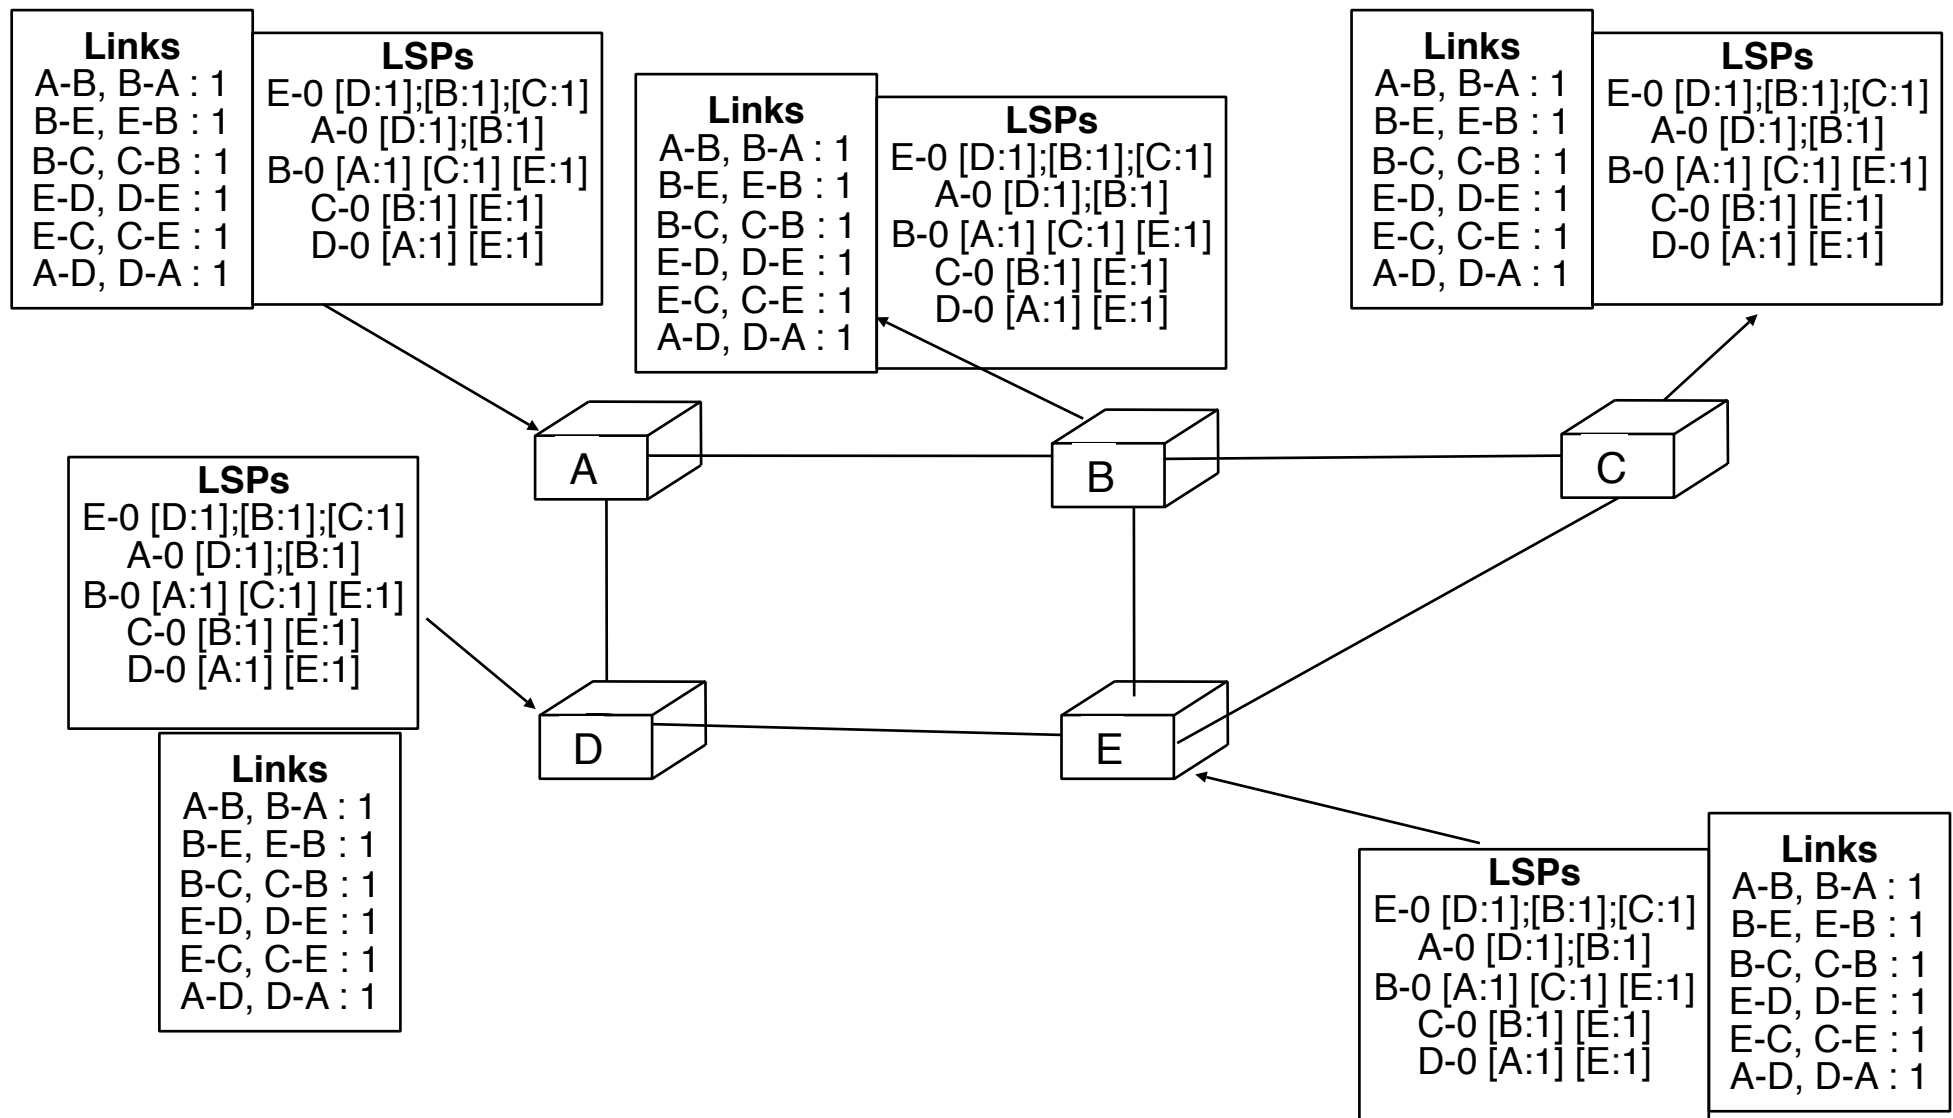


Thanks to its LSP table

A can detect that it received same LSP via B and D

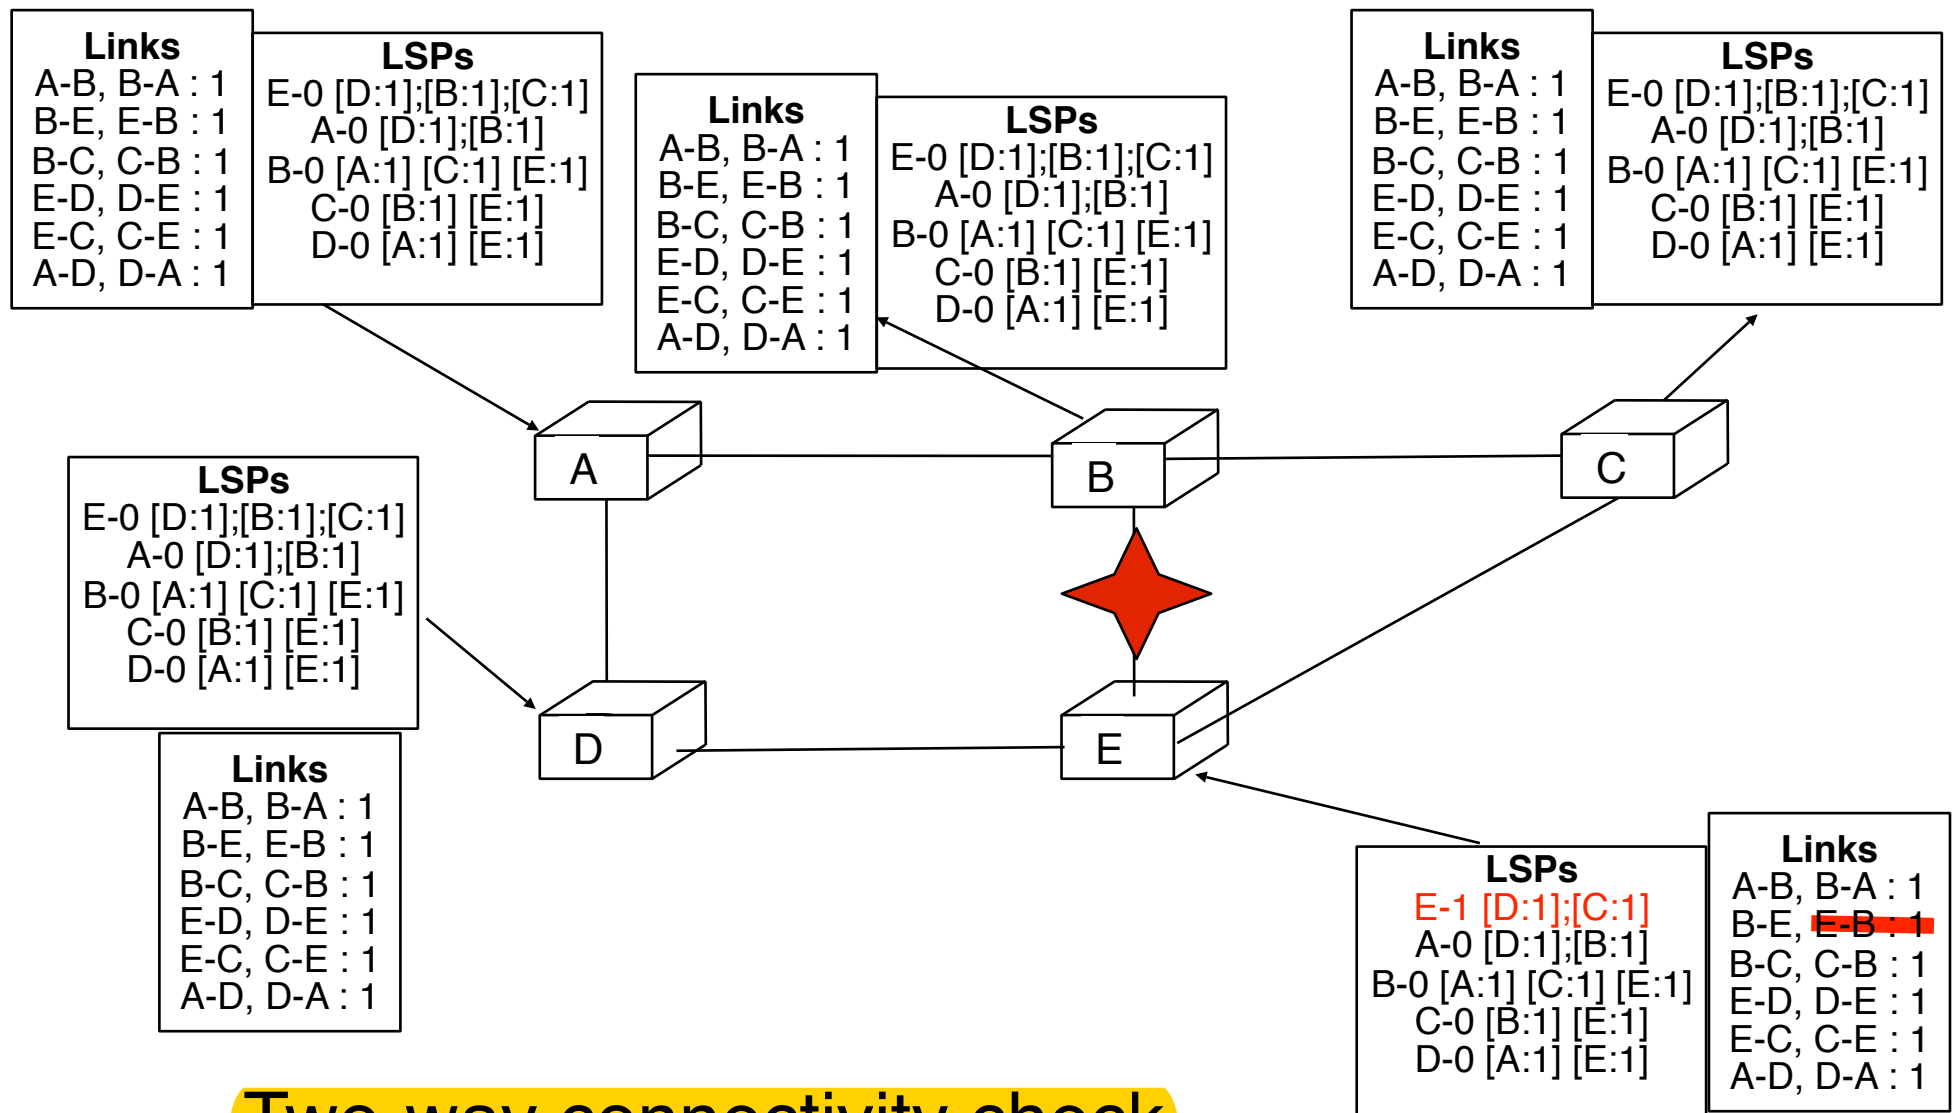
C can detect that it received same LSP via B and E

# Full topology





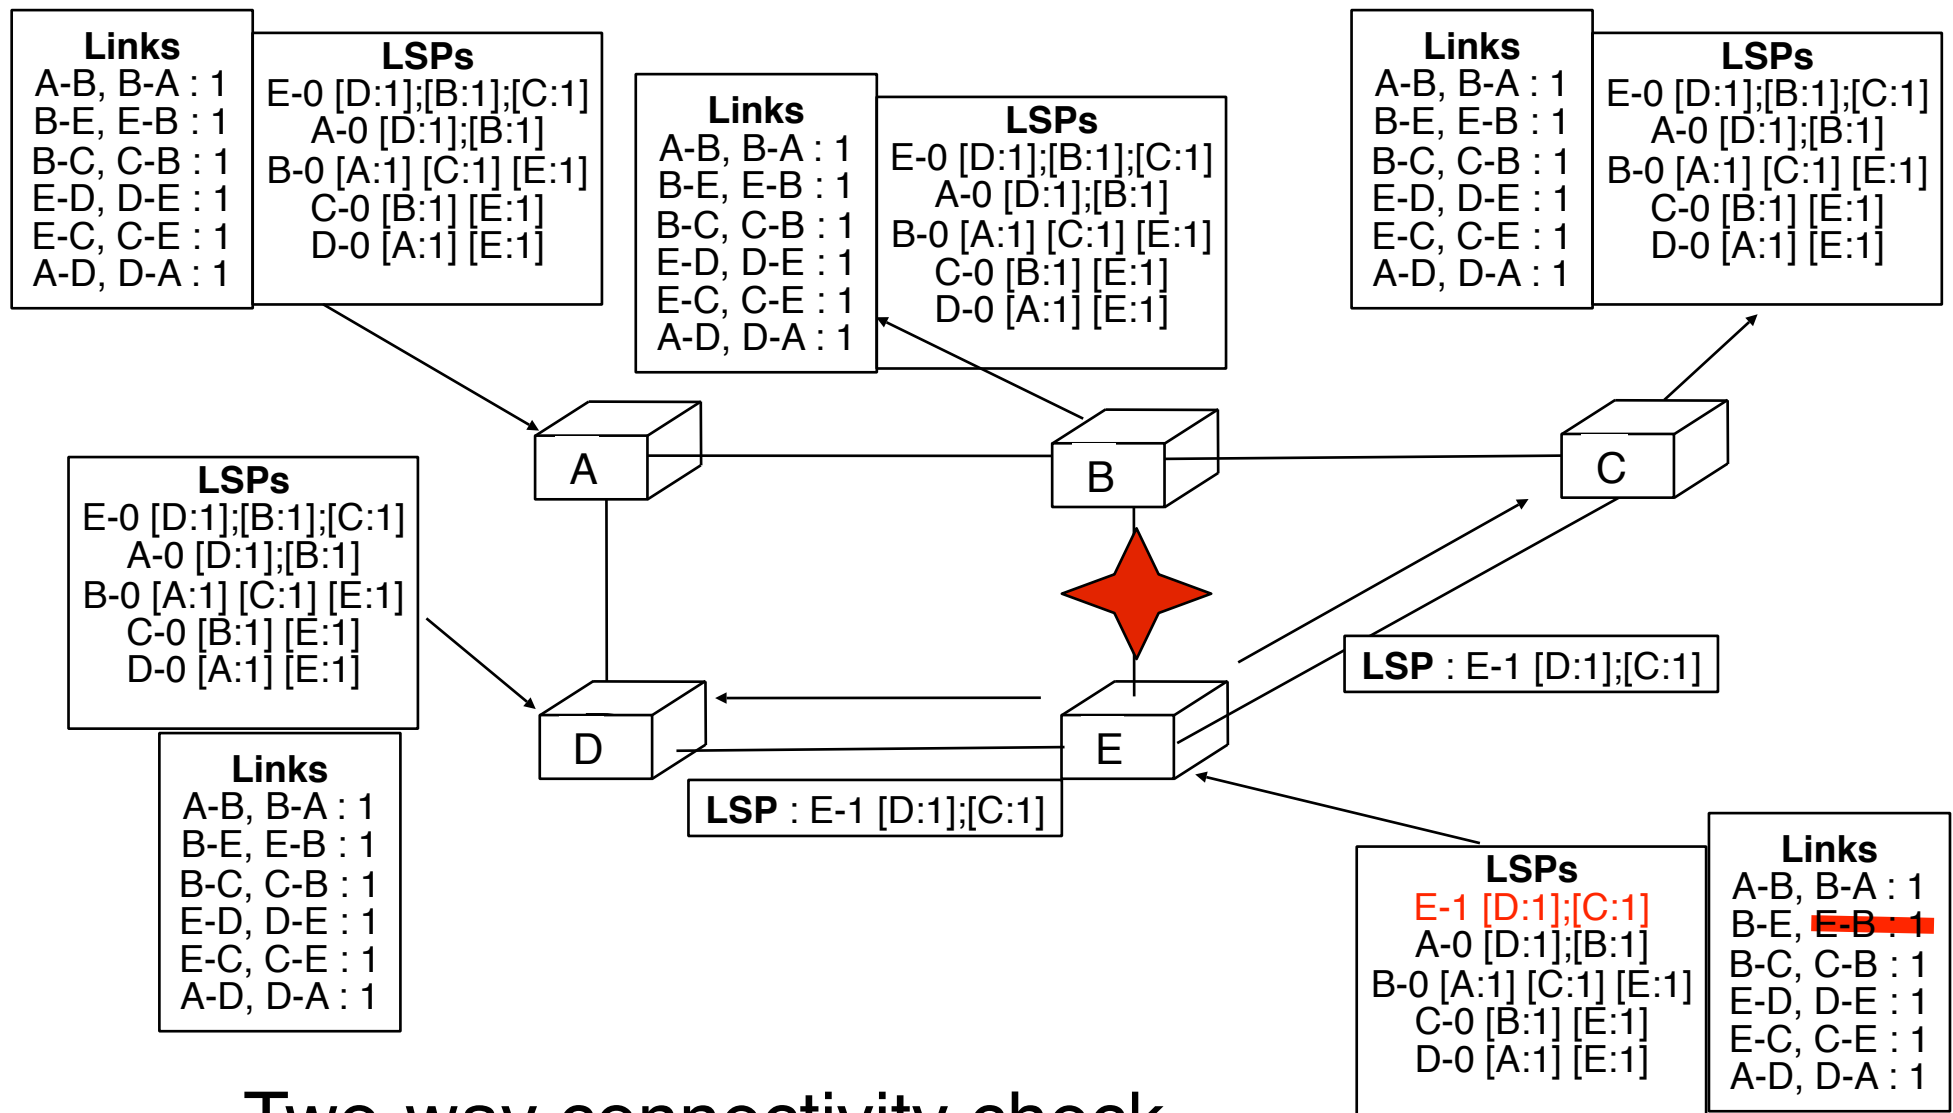
# How to deal with link failures ?



**Two-way connectivity check**

A link is only considered useable if **both** directions have been advertised

# How to deal with link failures ?



## Two-way connectivity check

A link is only considered useable if **both** directions have been advertised

# Router failures

---

# Router failures

---

What happens if a router fails ?

All its interfaces become unusable and do not reply anymore to HELLO packets

# Router failures

---

What happens if a router fails ?

All its interfaces become unusable and do not reply anymore to HELLO packets

What happens when the router reboots ?

It will send its LSP with its sequence number set to zero

If older LSPs from same router were still in network, then the new LSP will not be flooded

# Router failures

---

**What happens if a router fails ?**

All its interfaces become unusable and do not reply anymore to HELLO packets

**What happens when the router reboots ?**

It will send its LSP with its sequence number set to zero

If older LSPs from same router were still in network, then the new LSP will not be flooded

**Solution**

Add "age" field inside each LSP

Each router must decrement age regularly even for the LSPs stored in its LSDB

LSP having age=0 is too old and must be deleted

Each router must flood regularly its own LSP with age>0 to ensure that it remains inside network

# Improvements to LSP flooding

---

## Avoid sending twice same LSP on a link

When an LSP needs to be flooded on a link, wait some time to let other router flood the LSP  
reduces number of LSPs exchanged on a link but increases flooding time

## Reliable flooding

CRC inside each LSP to detect transmission errors  
Acknowledgements on each link for the LSPs exchanged on this link  
each transmission is protected by a timer

## Link state database exchange/synchronisation

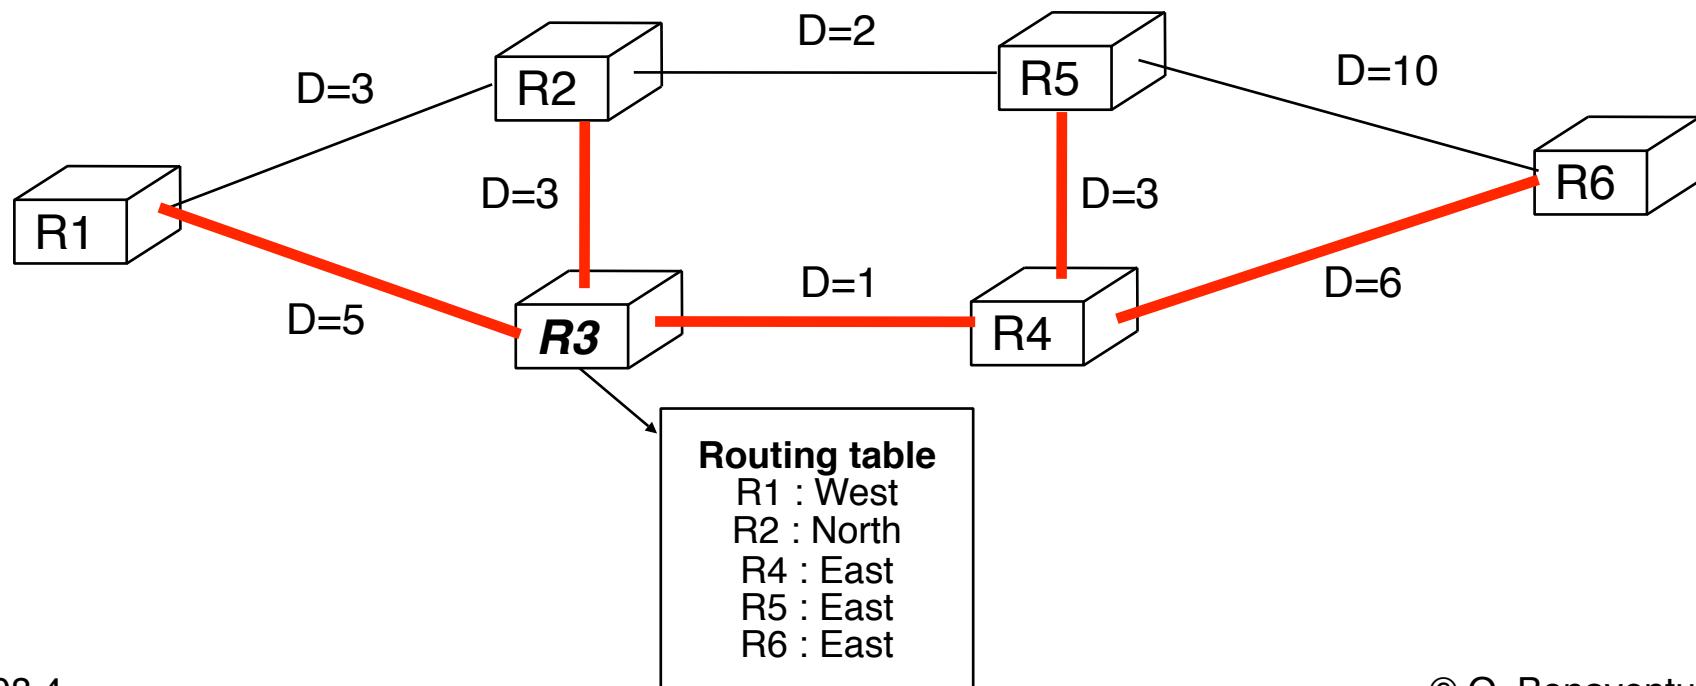
Routers can compare the content of their LSDB and exchange only missing LSPs from neighbour  
useful when the router boots and wants to receive quickly all LSPs from the network

# Computation of routing table

## Principle

Each router uses the received LSPs to build a graph and then computes the shortest spanning tree rooted on itself

From this spanning tree, it is easy to compute the routing table





# Dijkstra's shortest path

---

## Computing the shortest path tree

- At the beginning, the tree only contains the root node

- Adjacent routers are placed with the cost of their link in the candidates list

- Candidate router with lowest cost is chosen and added to the tree

- Consider the neighbours of the chosen candidate router and update the candidate router list if

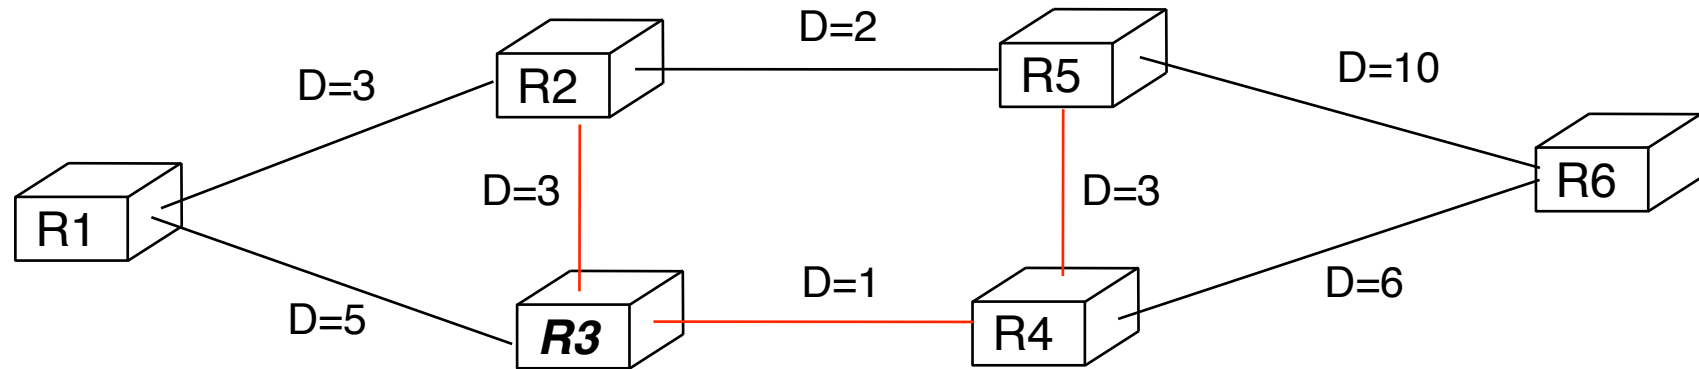
  - one of the new neighbours was not already in the candidates list

  - one of the new neighbours was already in the candidates list but with a longer path than the one in the current list

- Algorithm continues with the new candidates list and

- ends when all routers belong to shortest path tree

# Dijkstra's shortest path (2)



- 1) Routers : [R1, R2, R4, R5, R6] ; Candidates : [ - ] ; Tree : R3
- 2) Routers : [R5, R6] ; Candidates : [R1(5) ; R2(3) ; R4 (1) ]  
 selected candidate : R4  
 New tree : R3 - R4  
 New Candidates ? [R1(5) ; R2 (3) ; R5(R4-4) ; R6(R4-7) ]
- 3) Routers [ - ]  
 Selected candidate : R2 ; New tree : R2 - R3 - R4  
 New Candidates ? [R1(5) ; R5(R4-4) ; R6 (R4-7) ]
- 4) Selected candidate : R5 ; New tree : R2 - R3 - R4 - R5  
 New candidates ? [R1(5) ; R6(R4-7) ]
- 5) ...

# Network layer

---

## Basics

## Routing

- Static routing

- Distance vector routing

- Link state routing

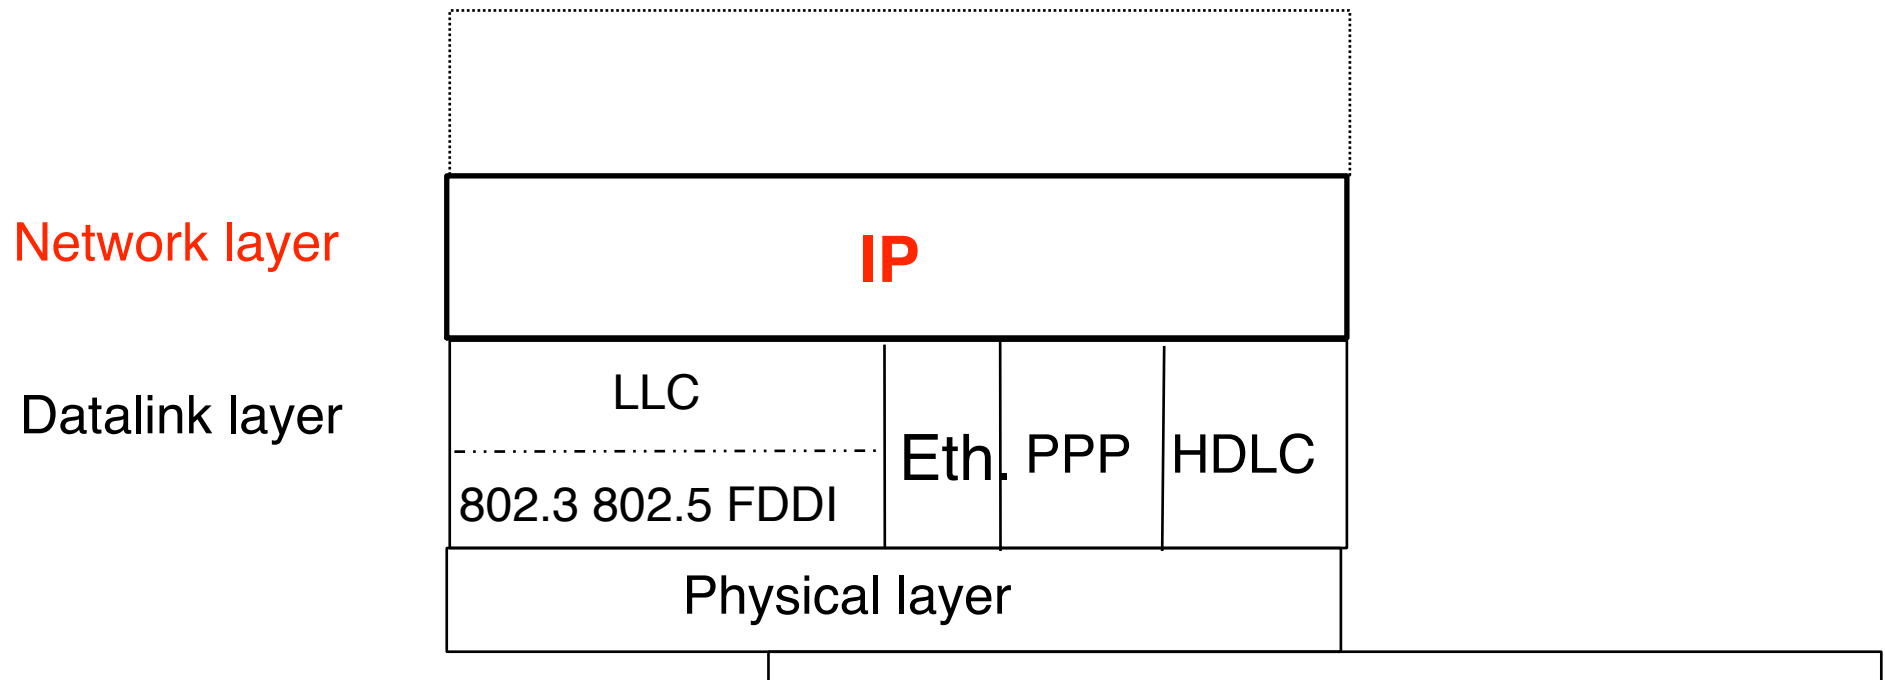
→ **IP : Internet Protocol**

- IP version 4**

- IP version 6

## Routing in IP networks

# IP : Internet Protocol



## Internet network layer

provides unreliable connectionless service

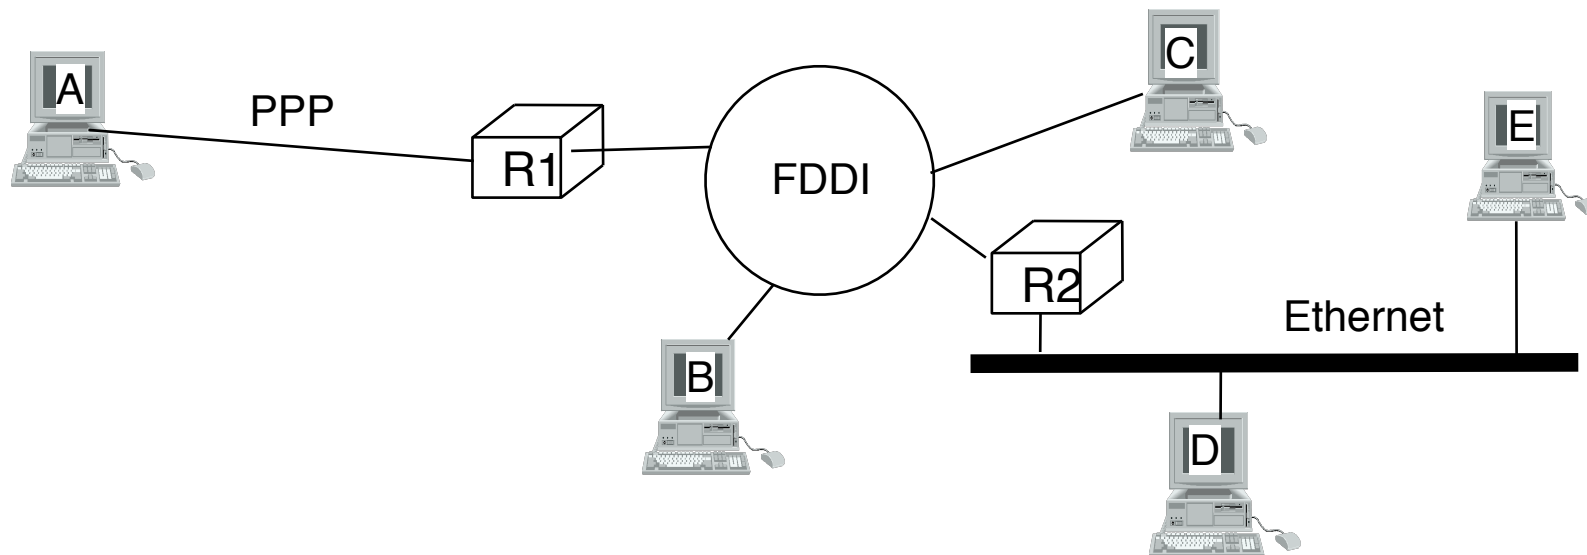
some packets can be lost

packets can suffer from transmission errors

packets can be misordered

# Basic principles

## Datagram mode




Each host is identified by one IP address  
(encoded as 32 bits number)

Each host knows how to reach at least one router

Routers know how to reach other routers

# Basic principles (2)



## Endhost

equipment able to send and receive packets originated by or destined to it

## Router

equipment able to send and receive packets originated by or destined to it

equipment able to **forward** toward their destination packets that it did not originate

# IP Addressing

---

# IP Addressing

---

## Utilisation of IP address

identify a host/router that implements IP

usually, **one IP address identifies one (physical) interface on one endhost or router**

(physical) interface is access point to datalink layer

usually endhosts have a single interface

routers have more than one interface

## Encoding of 32 bits IP address

10001010 **00110000** 00011010 00000001  
138 . **48** . 26 . 1



# IP Addressing

## Utilisation of IP address

- identify a host/router that implements IP

- usually, one IP address identifies one (physical) interface on one endhost or router

- (physical) interface is access point to datalink layer

- usually endhosts have a single interface

- routers have more than one interface

## Encoding of 32 bits IP address

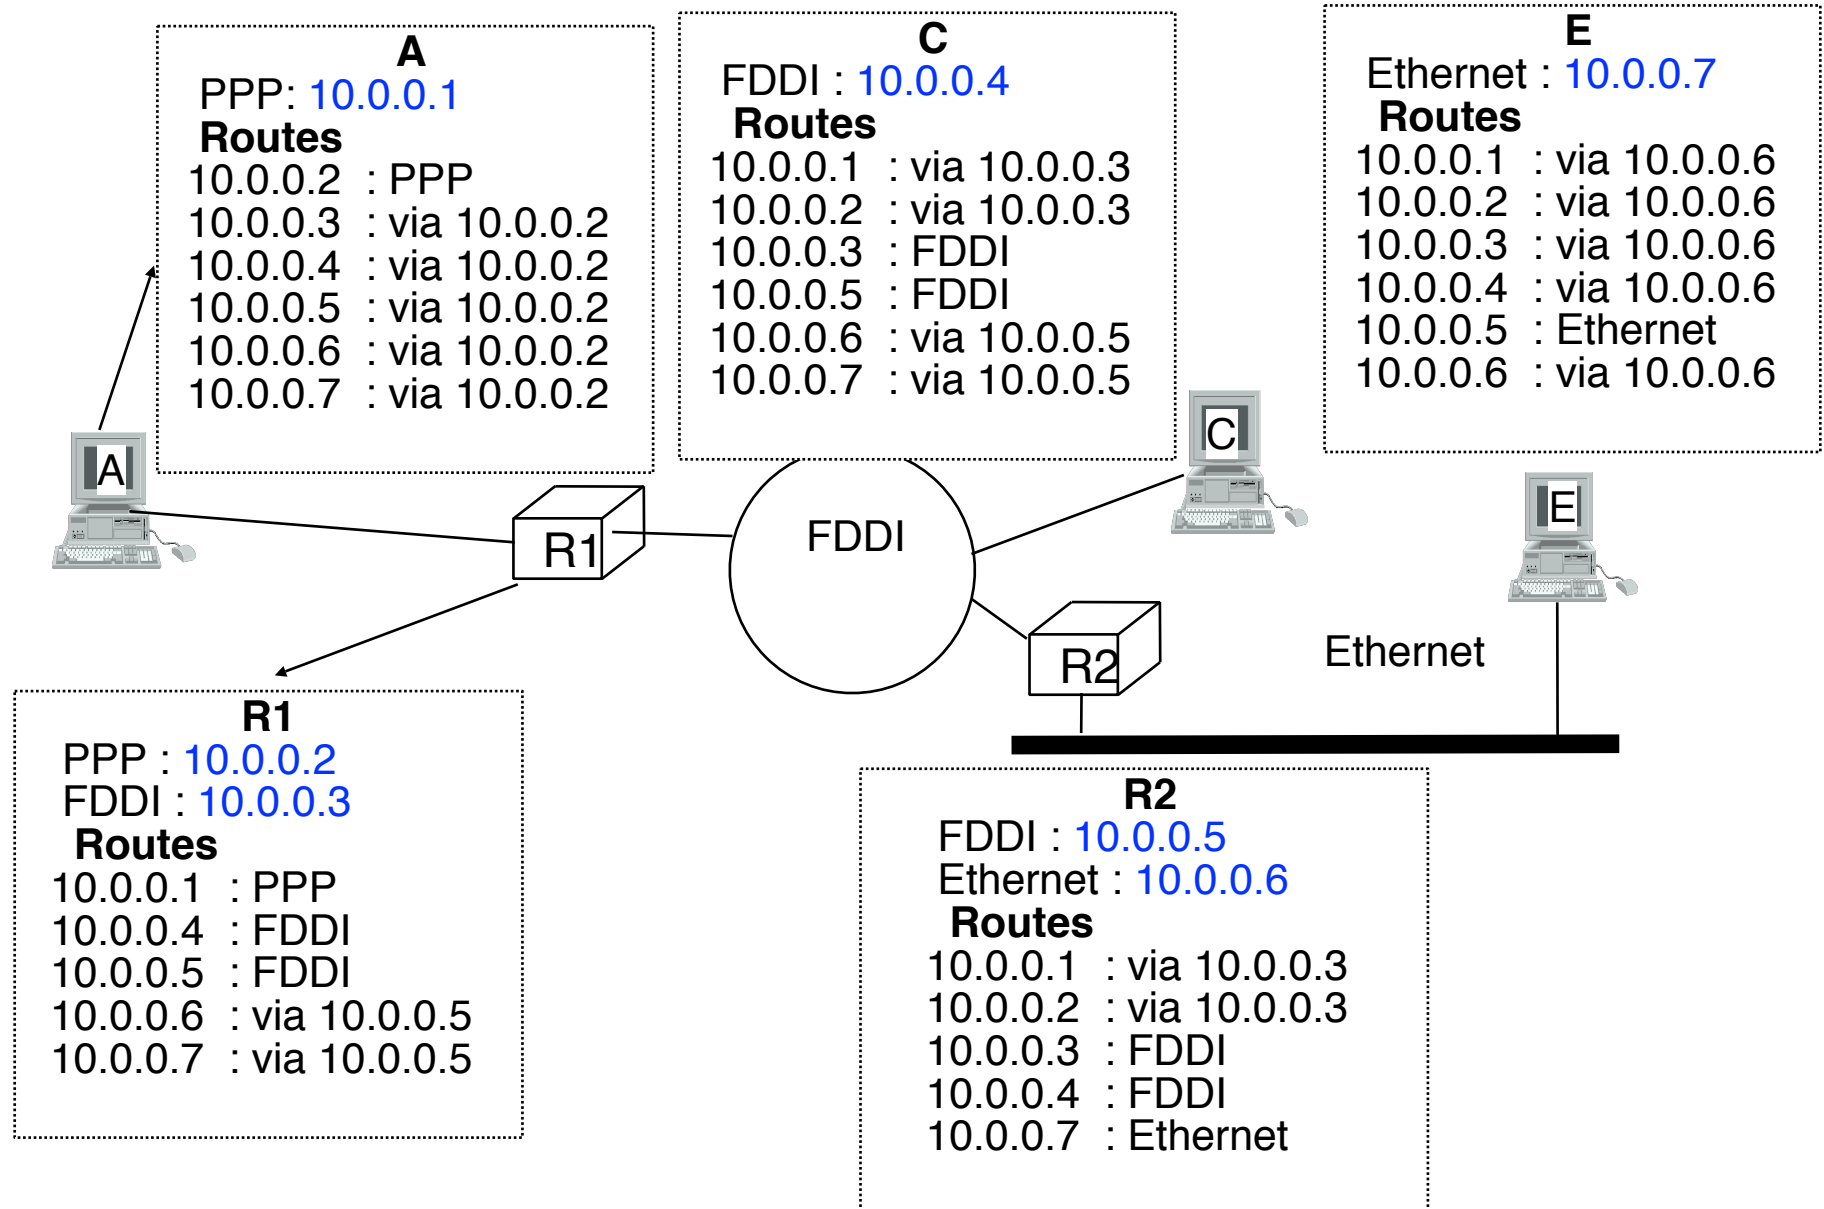
10001010 00110000 00011010 00000001  
138 . 48 . 26 . 1

## How to allocate IP addresses to hosts in a campus network

- Naive solution

- First come first served

# Naive IP addressing



# Hierarchical allocation of IP addresses

## Allocation of IP addresses

one address per interface

each address composed of two parts

1. subnetwork identifier  
M high order bits of IP address
2. equipment identifier inside the subnetworks  
32-M bits low order bits of IP address

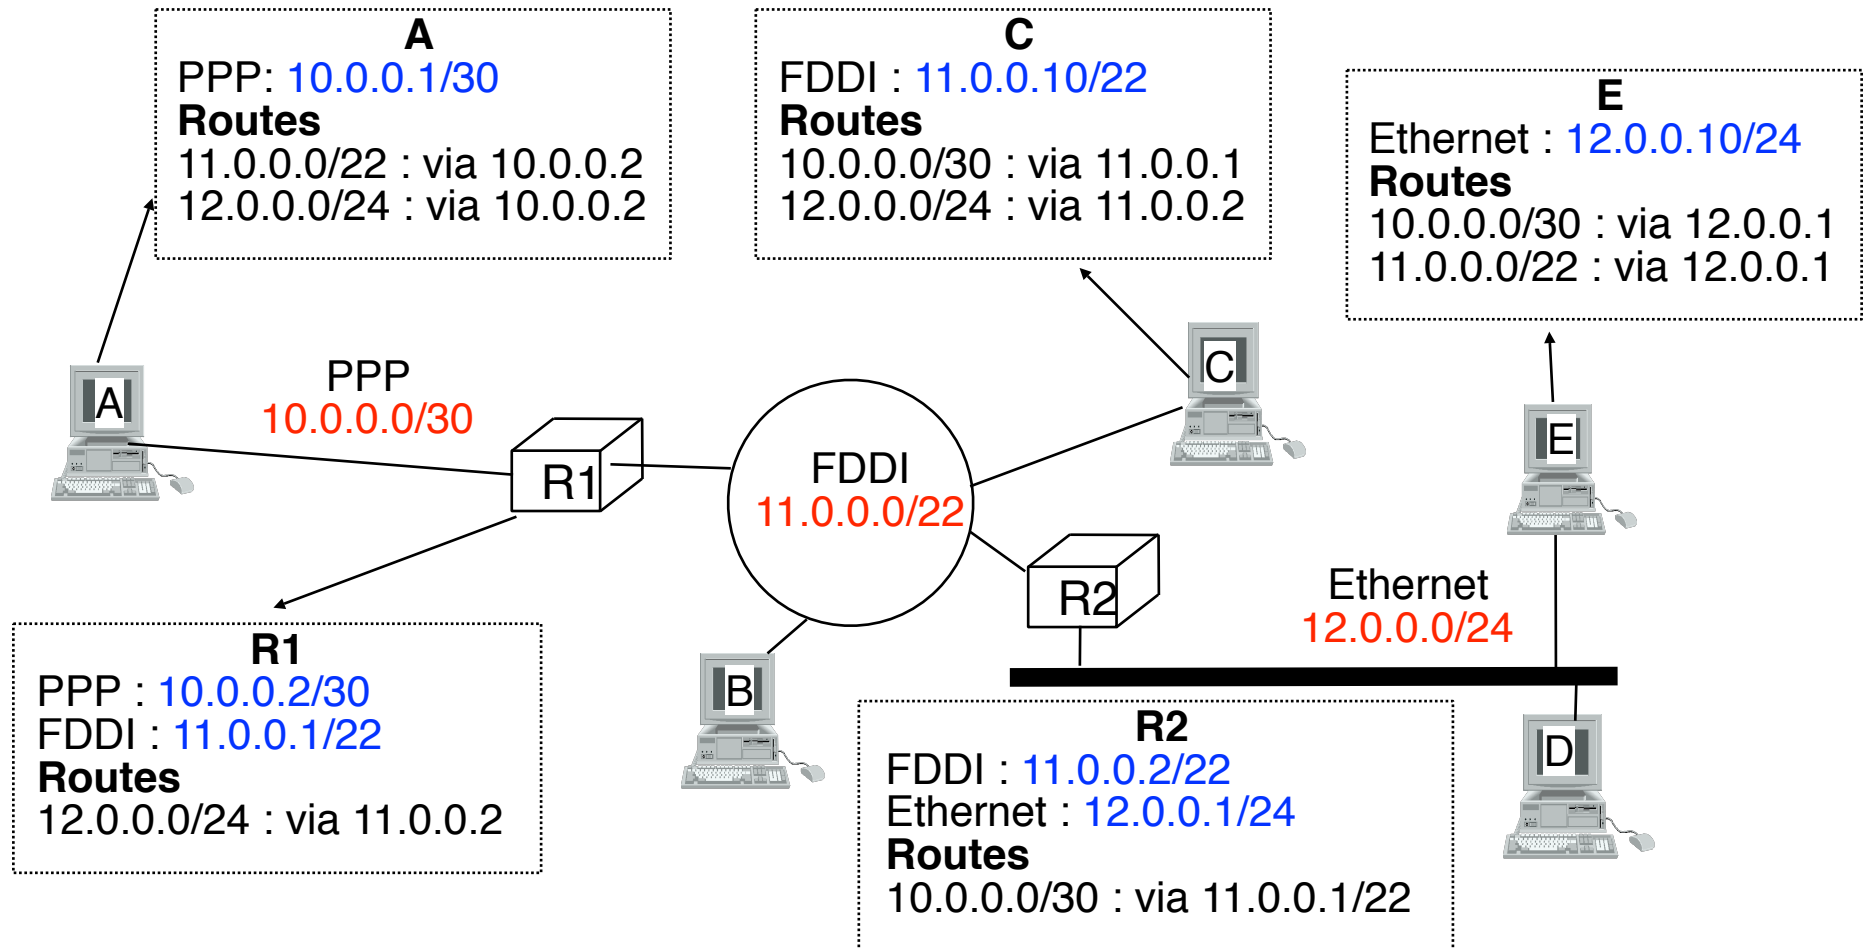
Example

10001010 00110000 00011010 00000001  
← subnetwork id                      host id

Notation 138.48.26.1/23 or 138.48.26.1 255.255.254.0

All hosts that belong to the same subnetwork can directly exchange frames through datalink layer

# IP addressing : examples



## Drawbacks of subnetworks

most subnetworks are not fully occupied

a campus network will need more IP addresses than the number of hosts attached to the network

# IP addresses

---

# IP addresses

---

Most addresses are allocated by IANA  
and the regional registries RIPE, ARIN, ...

# IP addresses

---

Most addresses are allocated by IANA  
and the regional registries RIPE, ARIN, ...

But some addresses play a special role

127.0.0.1

Loopback address on each host

Allows to reach servers on the local host

10.0.0.0/8, 172.16.0.0/12 and 192.168.0.0/16

used for private networks (not directly attached to Internet)

218.0.0.0/8 - 223.0.0.0/8 and 240.0.0.0/8 - 255.0.0.0/8

reserved for further utilization

224.0.0.0/8 - 239.0.0.0/8

used by IP multicast

255.255.255.255

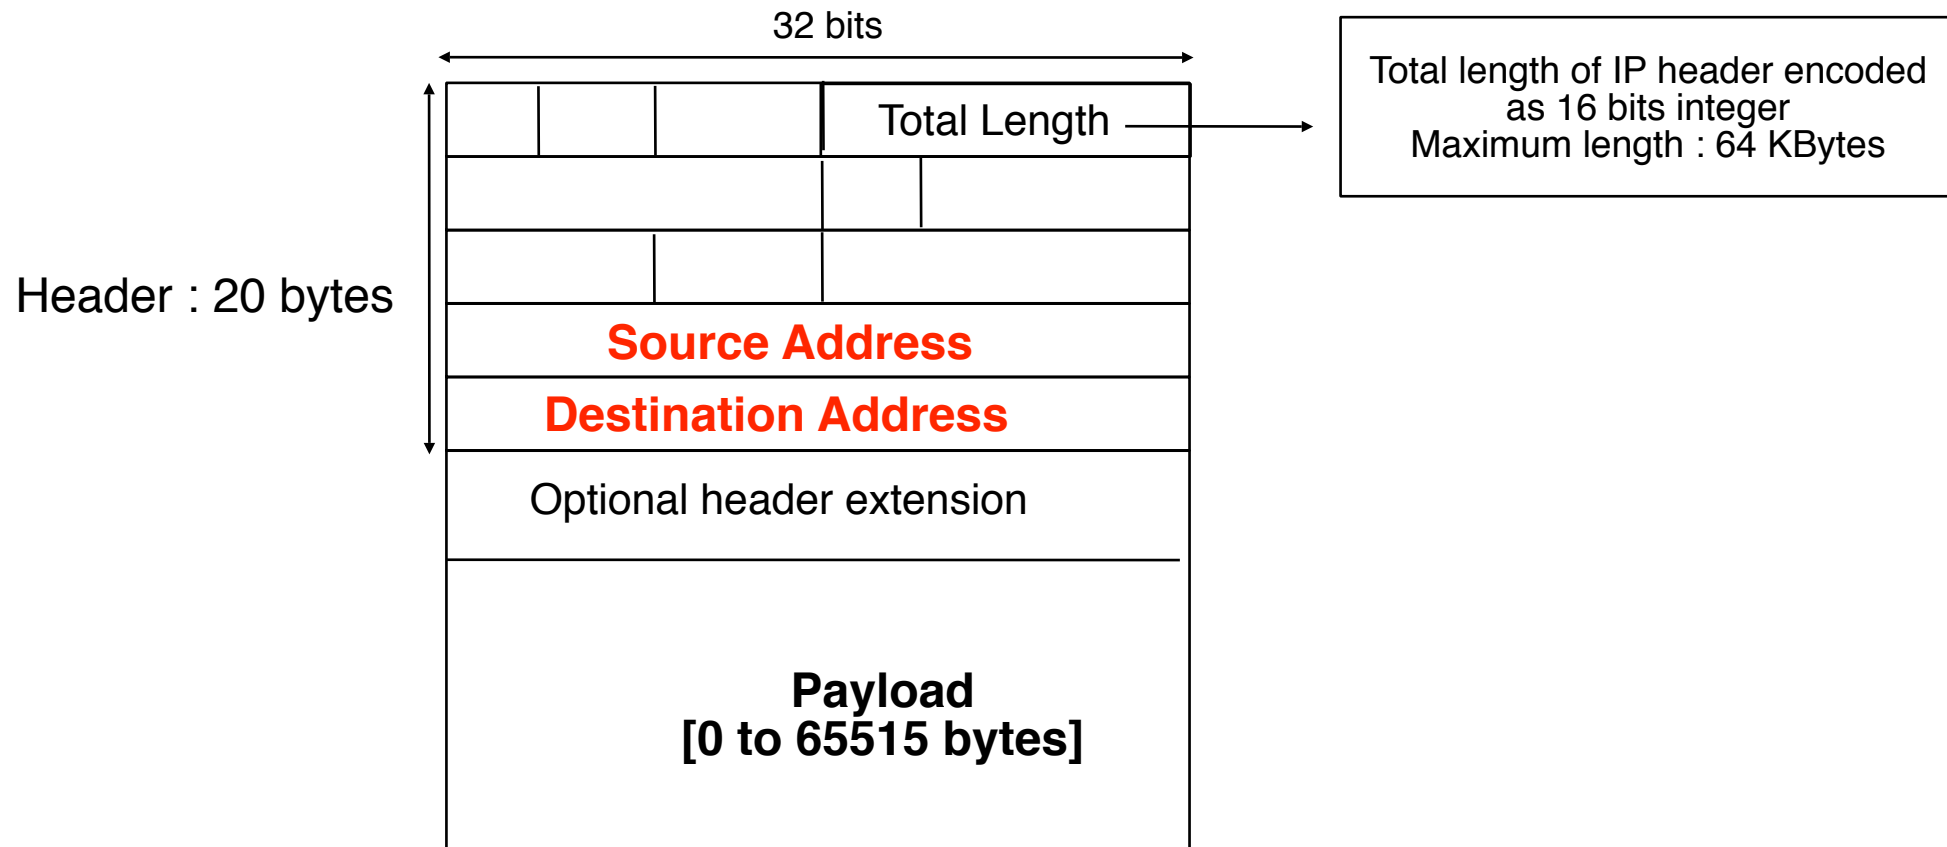
broadcast address

0.0.0.0

used when a host is booting and does not yet know its address

# IP Packets

## IP packet format



How can we transmit a 64 KBytes packet ?




# Transmission of long IP packets

## Principle


Each host and each router can fragment packets

Each **fragment** is a **complete IP packet** that contains source and destination IP addresses


Only the destination host performs reassembly



# Transmission of long IP packets (2)

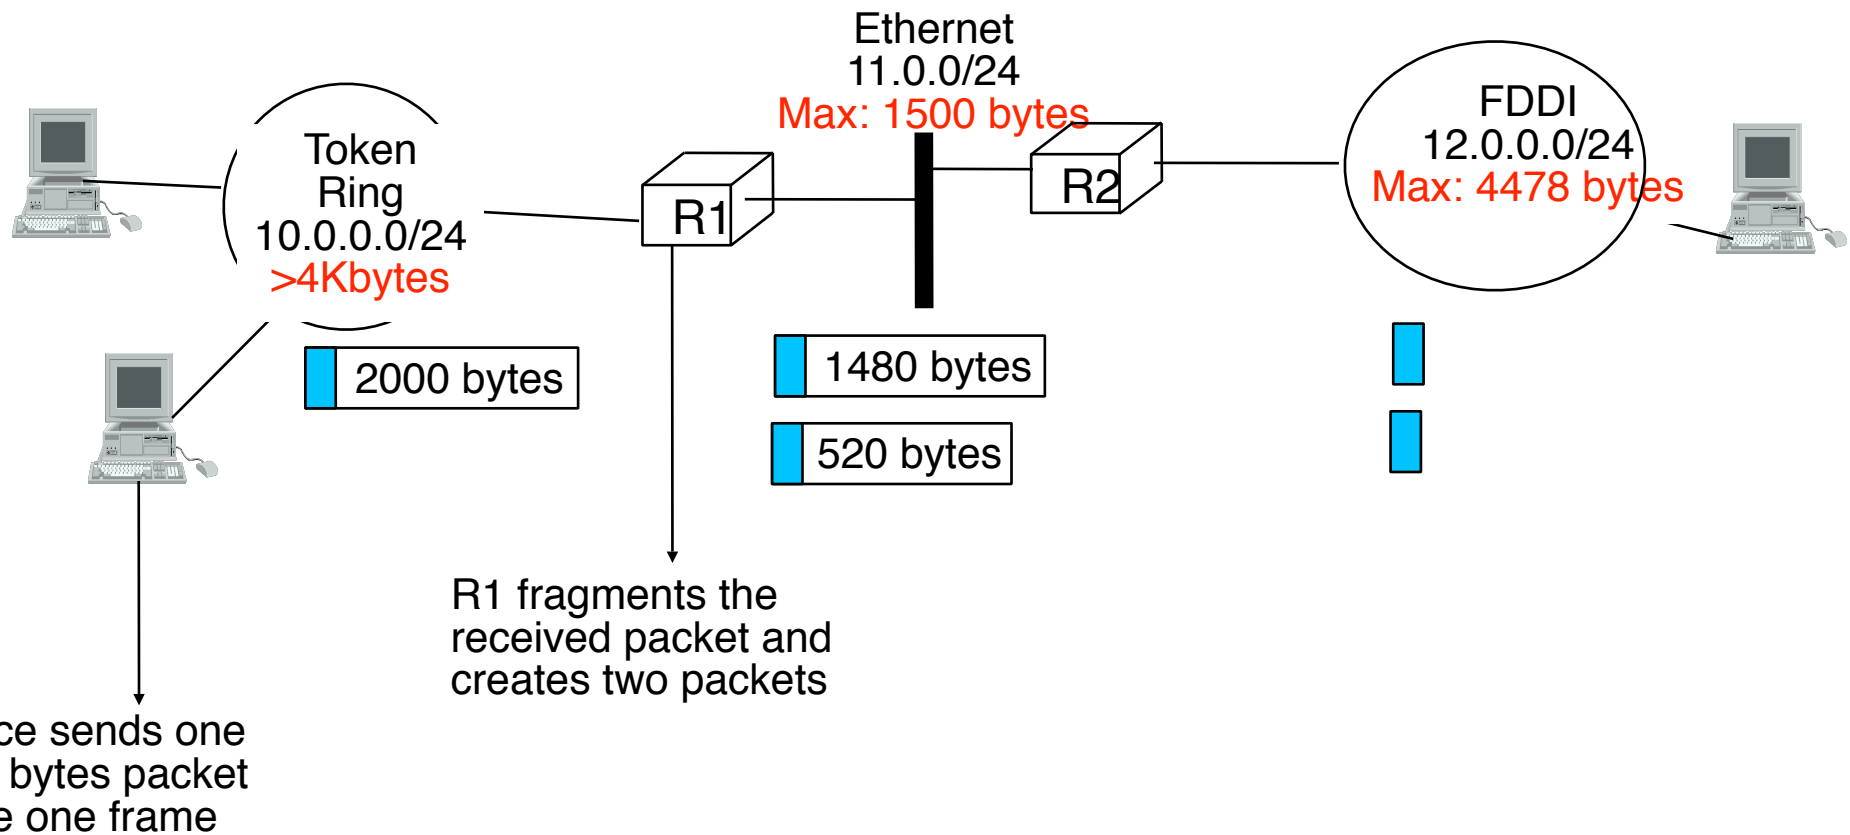


# Transmission of long IP packets (2)



Source sends one  
2000 bytes packet  
inside one frame

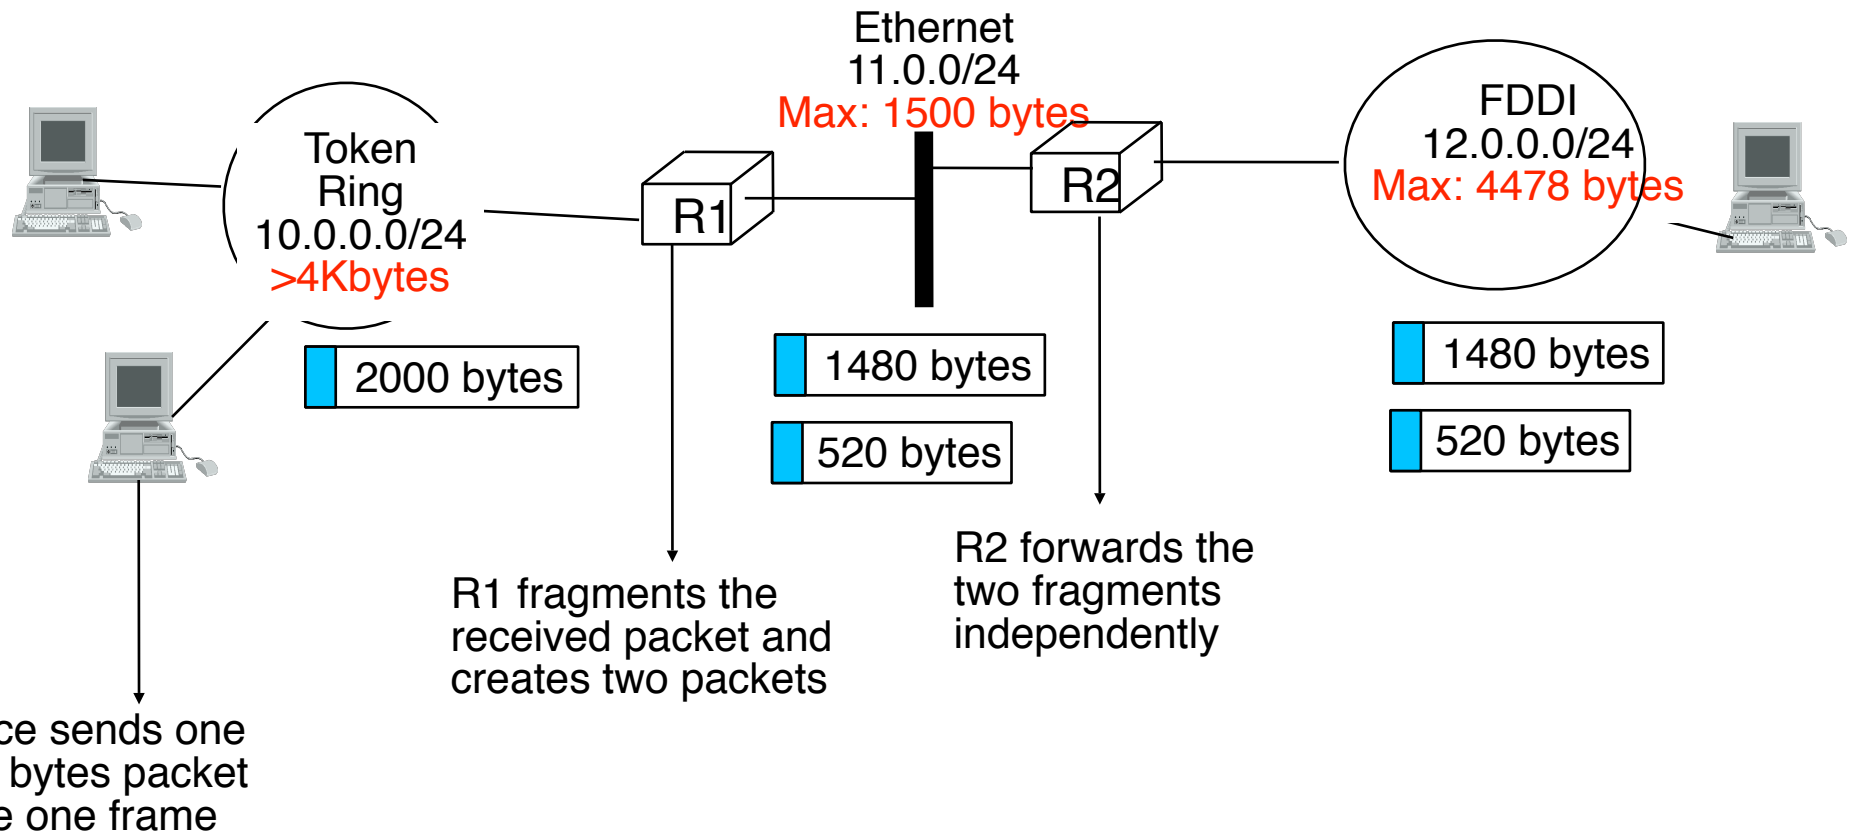
# Transmission of long IP packets (2)



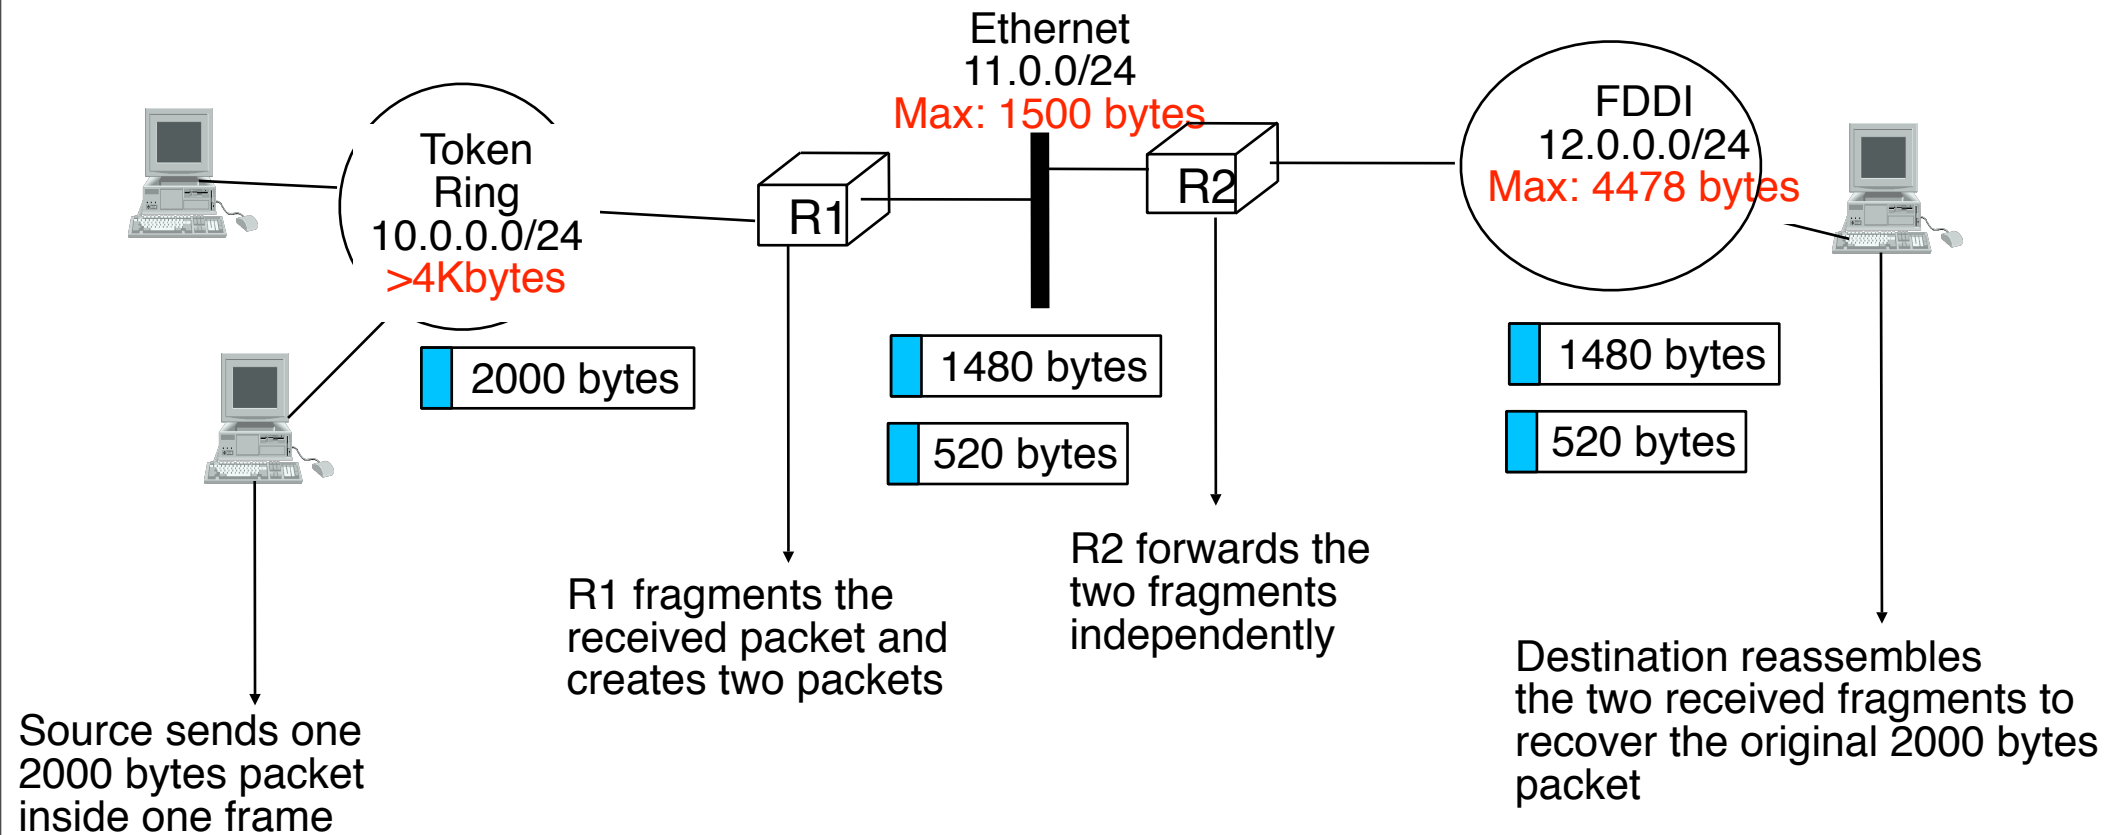
Source sends one 2000 bytes packet inside one frame

R1 fragments the received packet and creates two packets

# Transmission of long IP packets (2)



# Transmission of long IP packets (2)




# How to deal with limited MTU links ?

## IP fragmentation

Fragment the payload of IP packet

Each fragment must be numbered to recover from misordering

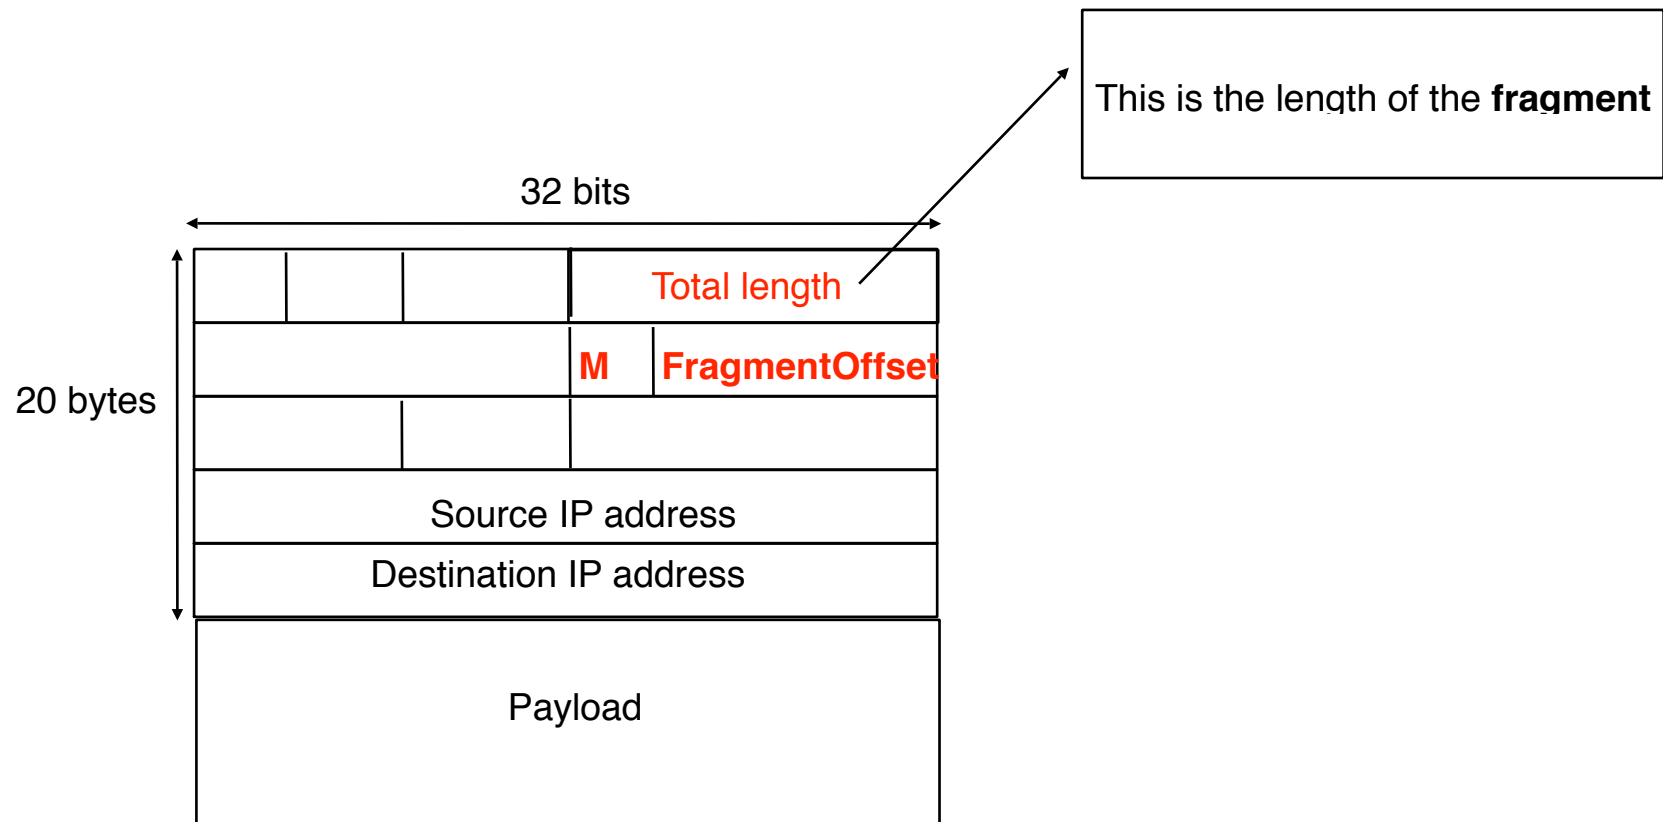


# How to deal with limited MTU links ?

## IP fragmentation

Fragment the payload of IP packet

Each fragment must be numbered to recover from misordering



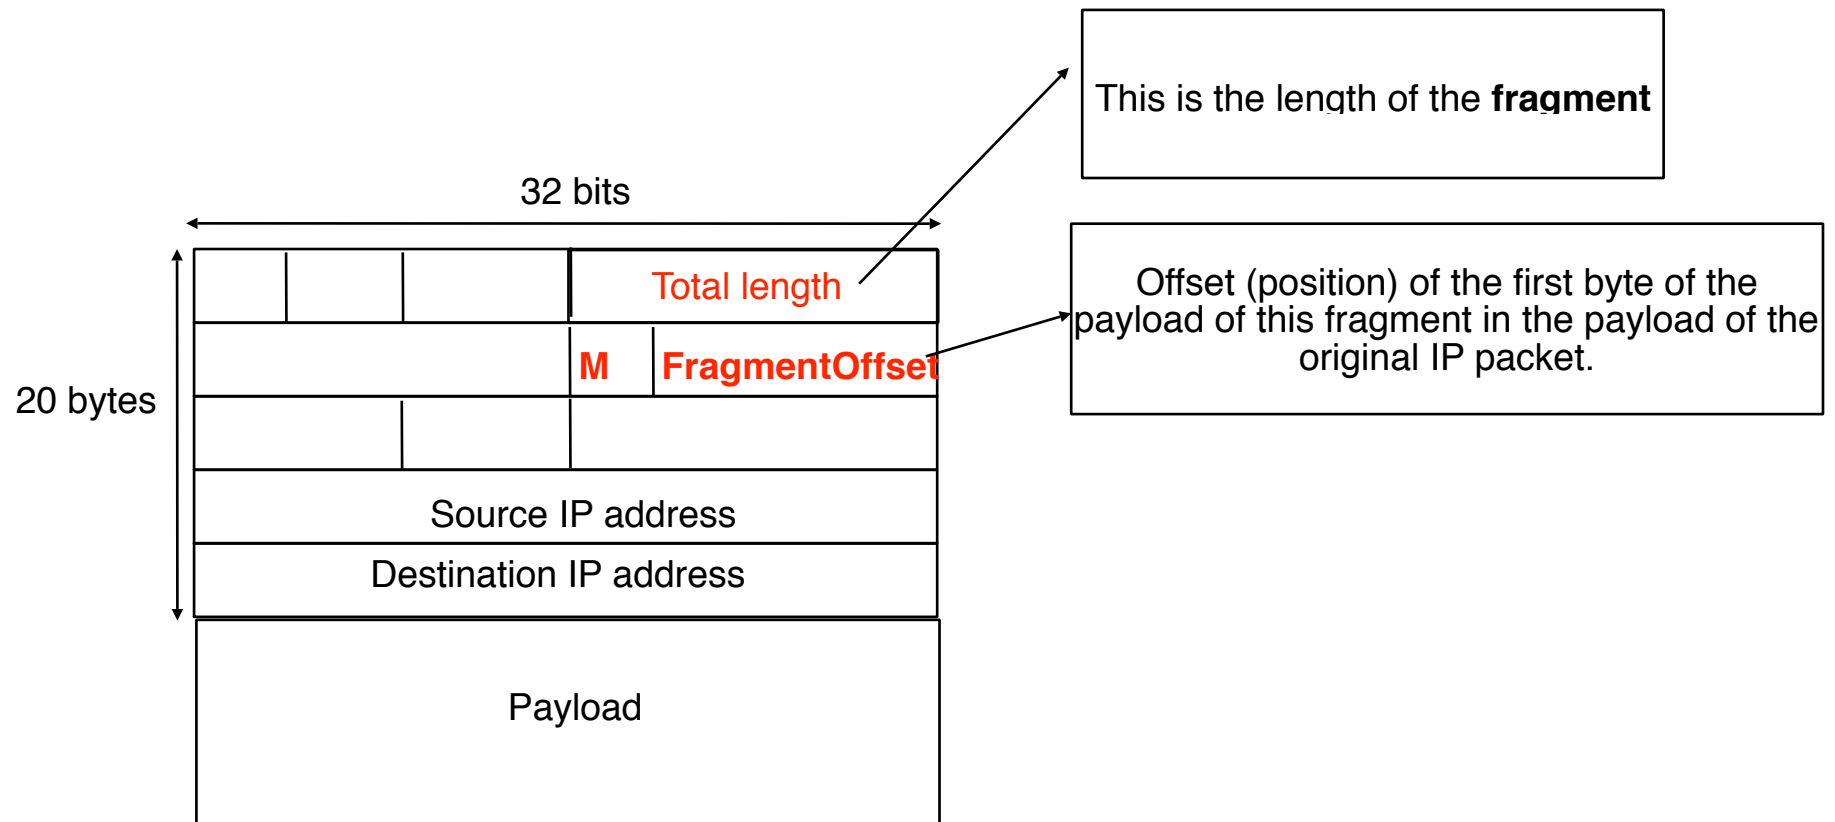


# How to deal with limited MTU links ?

## IP fragmentation

Fragment the payload of IP packet

Each fragment must be numbered to recover from misordering

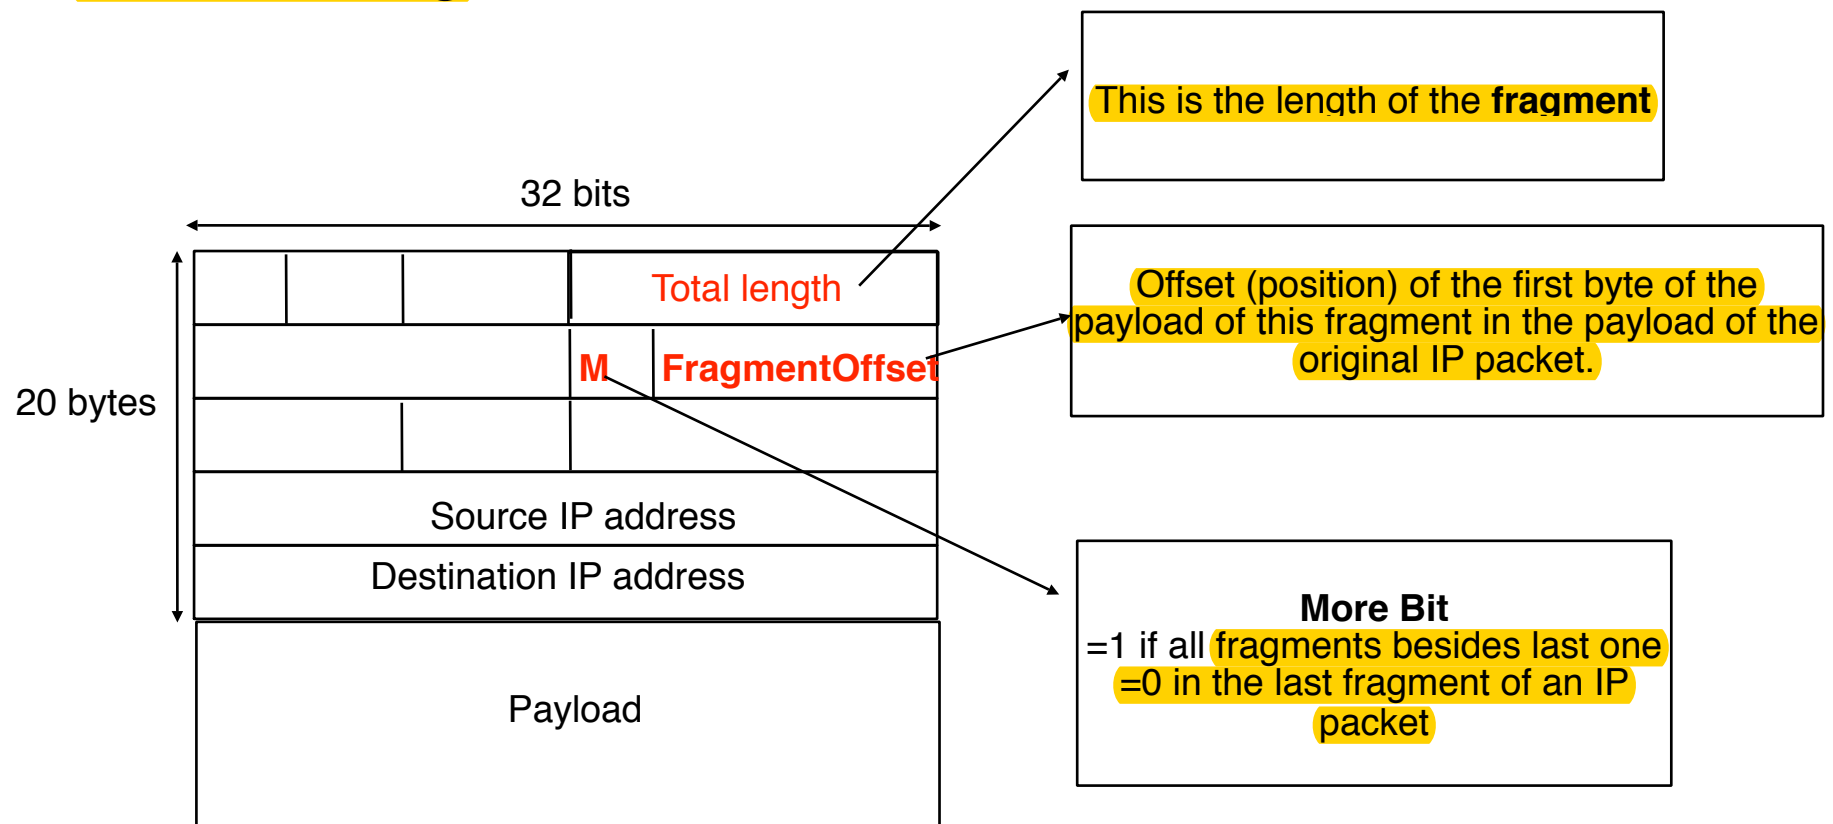


# How to deal with limited MTU links ?



## IP fragmentation

Fragment the payload of IP packet

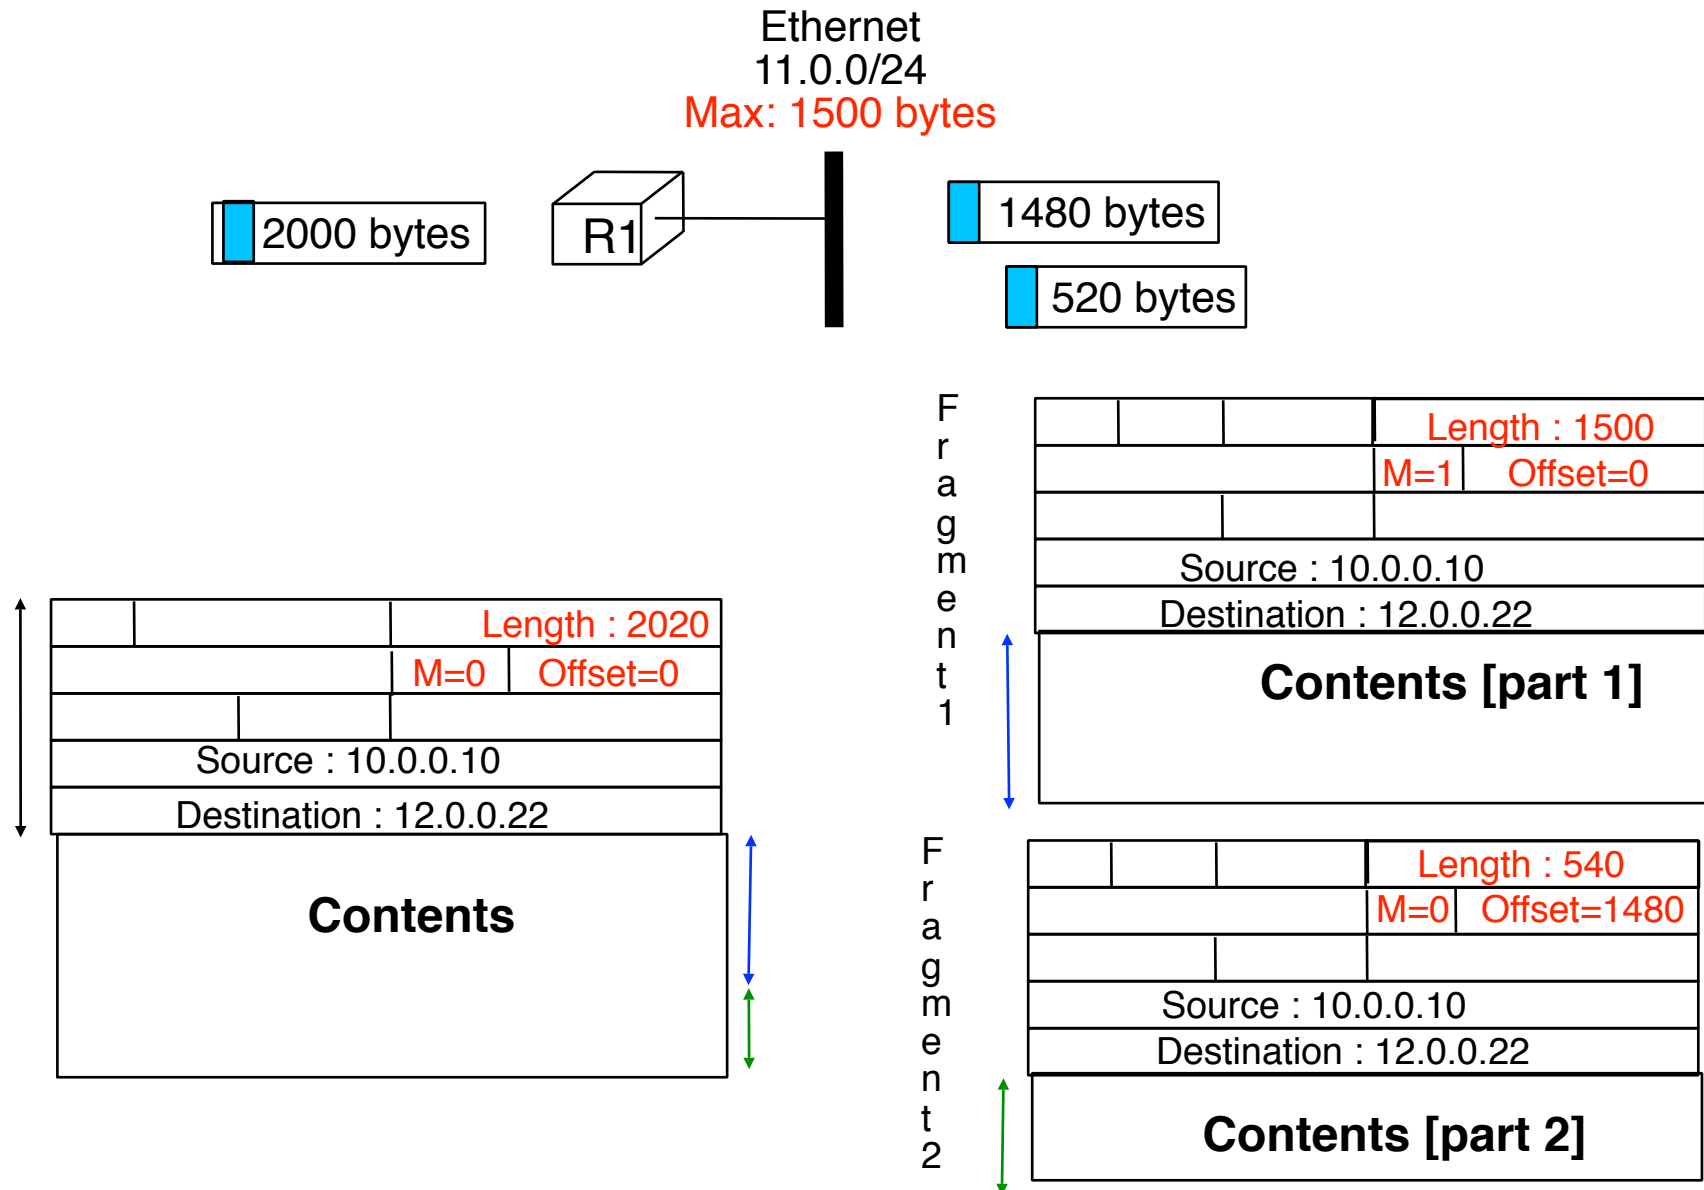
Each fragment must be numbered to recover from misordering



# Fragmentation : example



# Fragmentation : example



# Reassembly

---

## Issues

When does the destination has received all fragments ?

- Last fragment contains bit More=0

- How to handle lost fragments ?

  - the IP packet will not be reassembled by destination and received fragments of this packet will be discarded

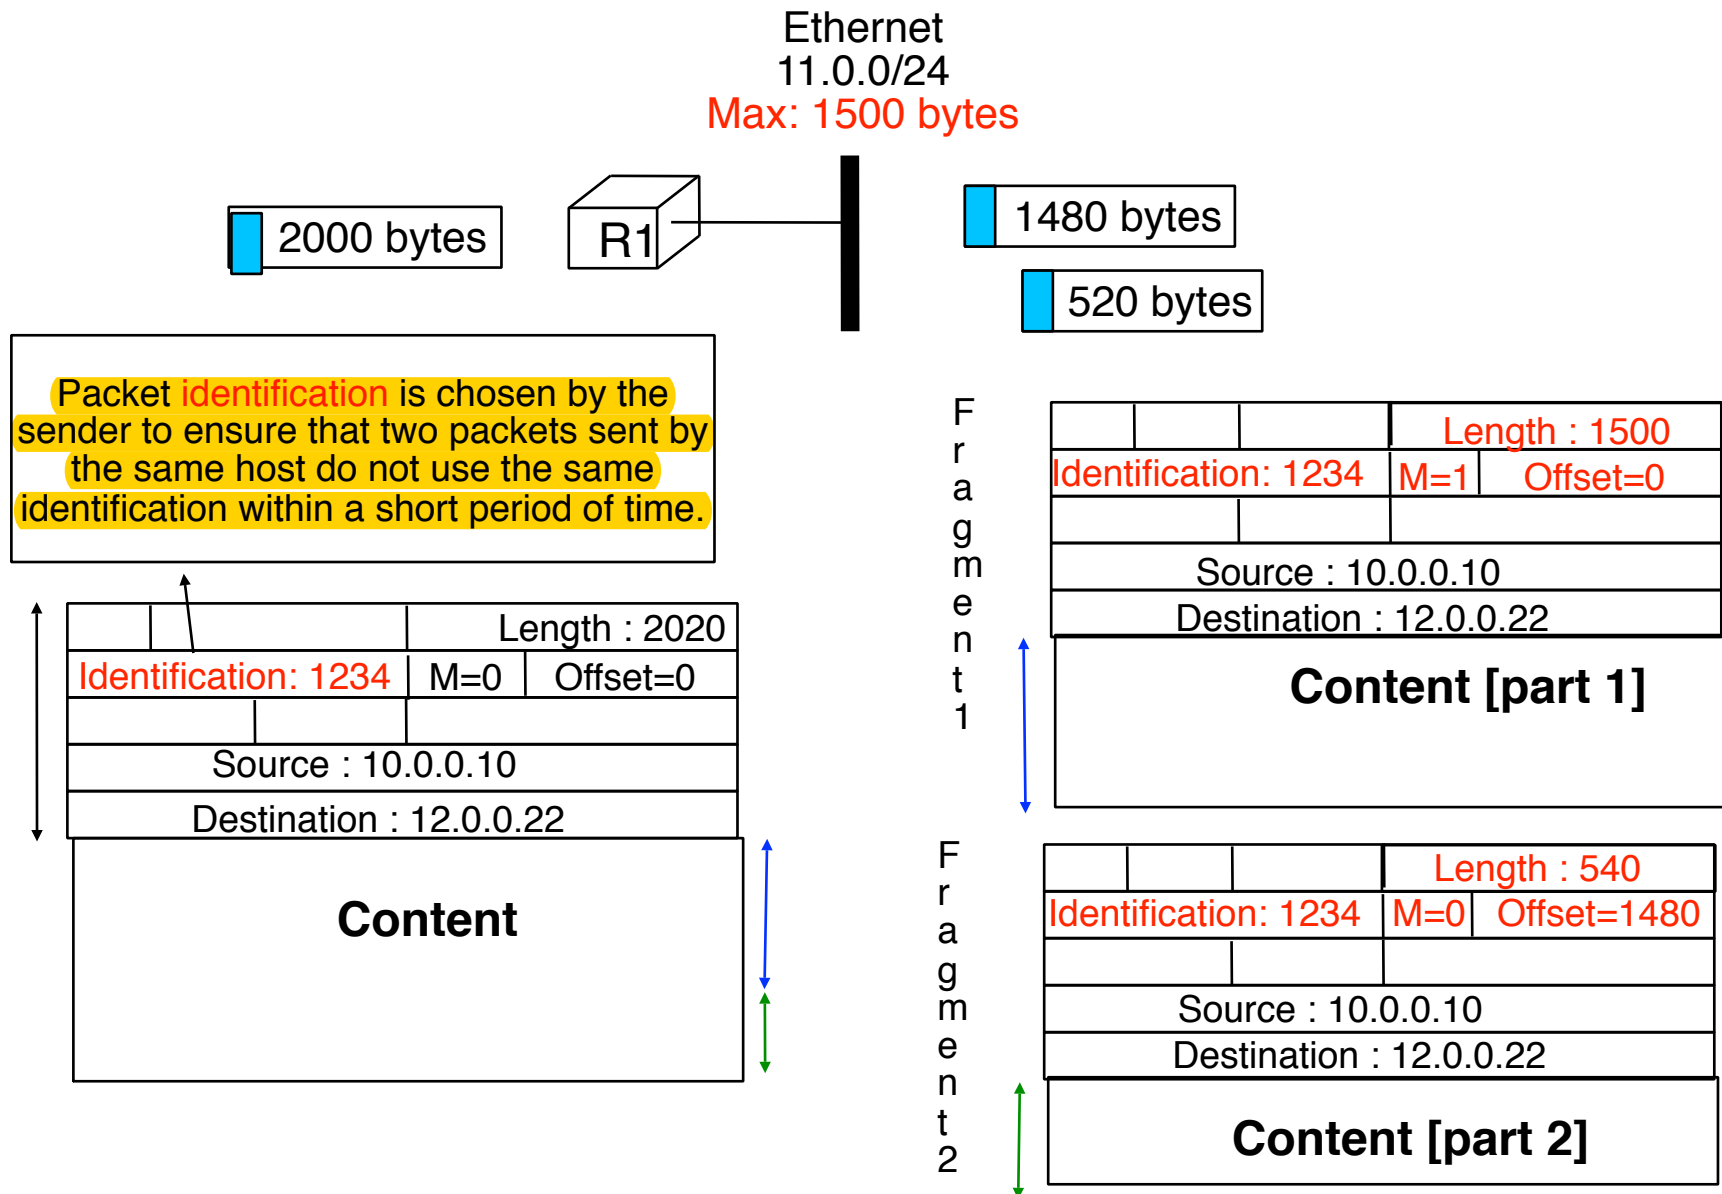
## How to deal with misordering

- Offset field allows to reorder fragments from same packet

- But misordering can cause fragments from multiple packets to be mixed

  - Each fragment must contain an identification of the original packet from which is was created

# Packets and fragments identification



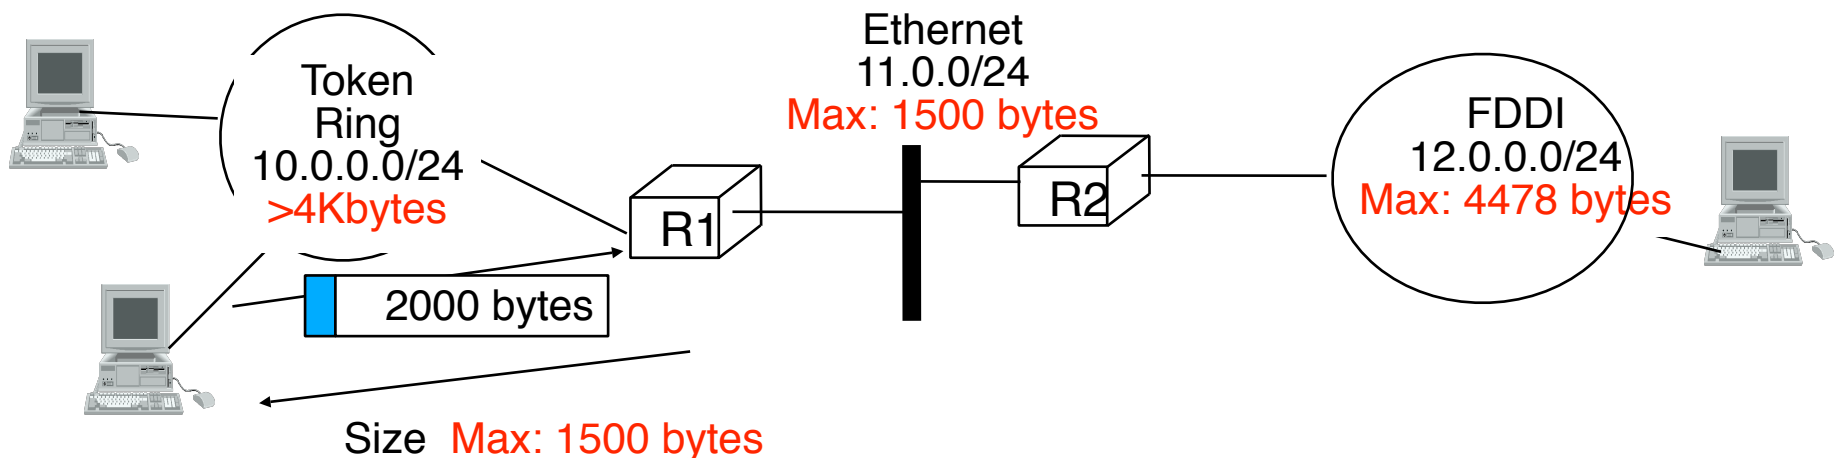
# How to avoid fragmentation ?

## Problem

How can a host determine the maximum packet size that he can use to reach a destination ?

## Solution

Instead of performing fragmentation, the router could indicate the maximum packet size that it supports



Knowing this maximum packet size, the endhost can send correctly sized packets

# Transmission errors

---

## How should IP react to transmission errors ?

### Transmission error inside packet content

some applications may continue to work despite this error

**IP : no detection of transmission errors in packet payload**

### Transmission error inside packet header

could cause more problems

imagine that the transmission error changes the source or destination

IP address

**IP uses a checksum to detect transmission errors in header**

16 bits checksum (same as TCP/UDP) computed only on header  
each router and each end host verifies the checksum of all packets  
that it receives. A packet with an errored header is immediately  
discarded



# Transient and permanent loops

---

# Transient and permanent loops

---

## Problem

Loops can occur in an IP network

permanent loops due to configuration errors

transient loops while routing tables are being updated

# Transient and permanent loops

---

## Problem

- Loops can occur in an IP network

  - permanent loops due to configuration errors

  - transient loops while routing tables are being updated

## Solution

- Each packet contains a **Time-to-Live (TTL)** that indicates the maximum number of intermediate routers that the packet can cross

  - many hosts set the initial TTL of their packets to 32 or 64

- each router checks the TTL of all packets

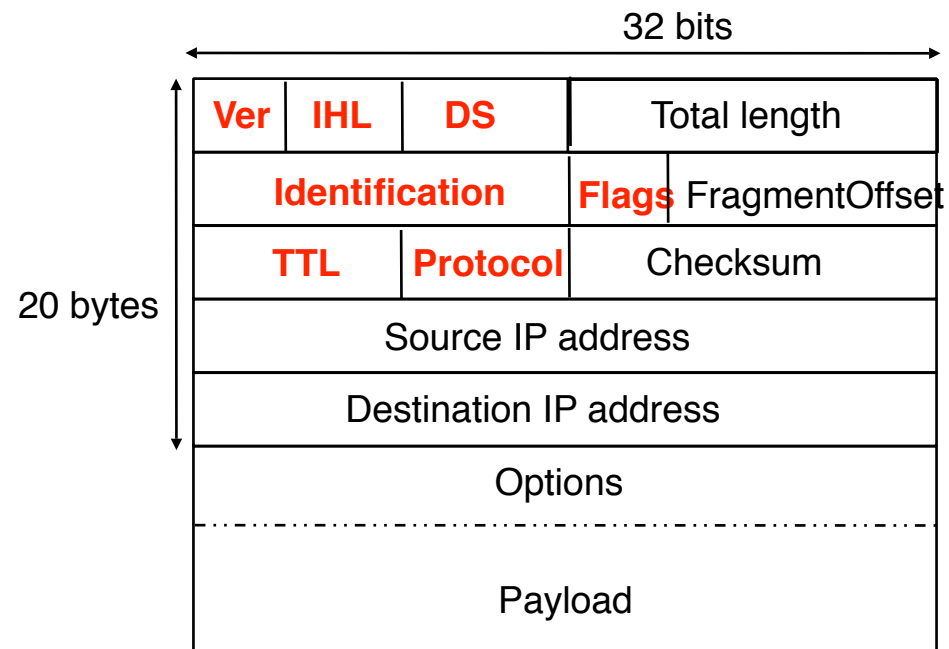
  - If  $TTL=1$ , packet is discarded and source is notified

  - If  $TTL>1$ , packet is forwarded and TTL is decremented by at least 1

    - routers thus must recompute checksum of all forwarded packets

- Utilisation of TTL is a means to bound the lifetime of packets inside the Internet**

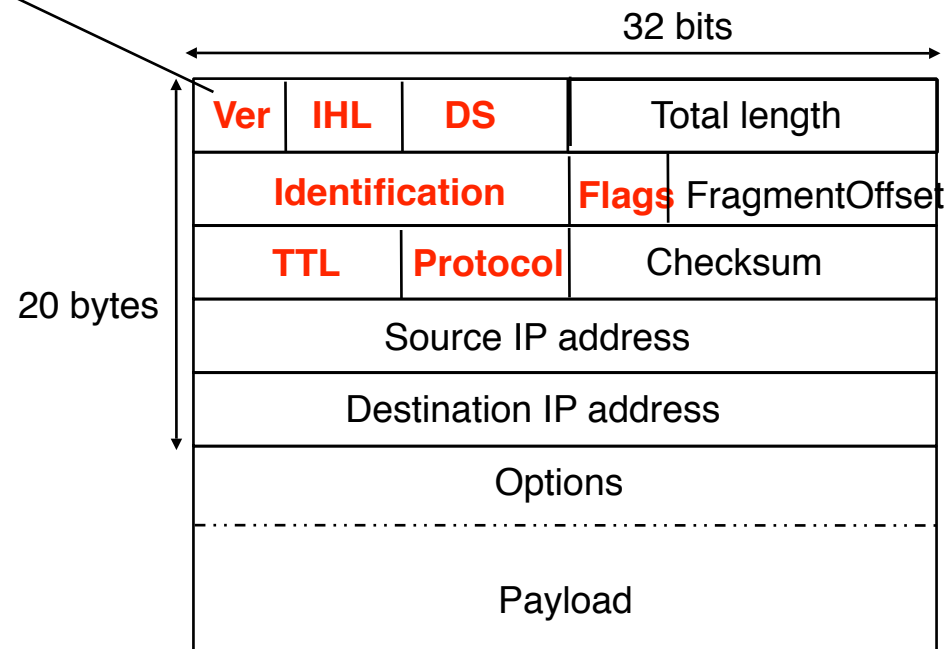
# IP header format



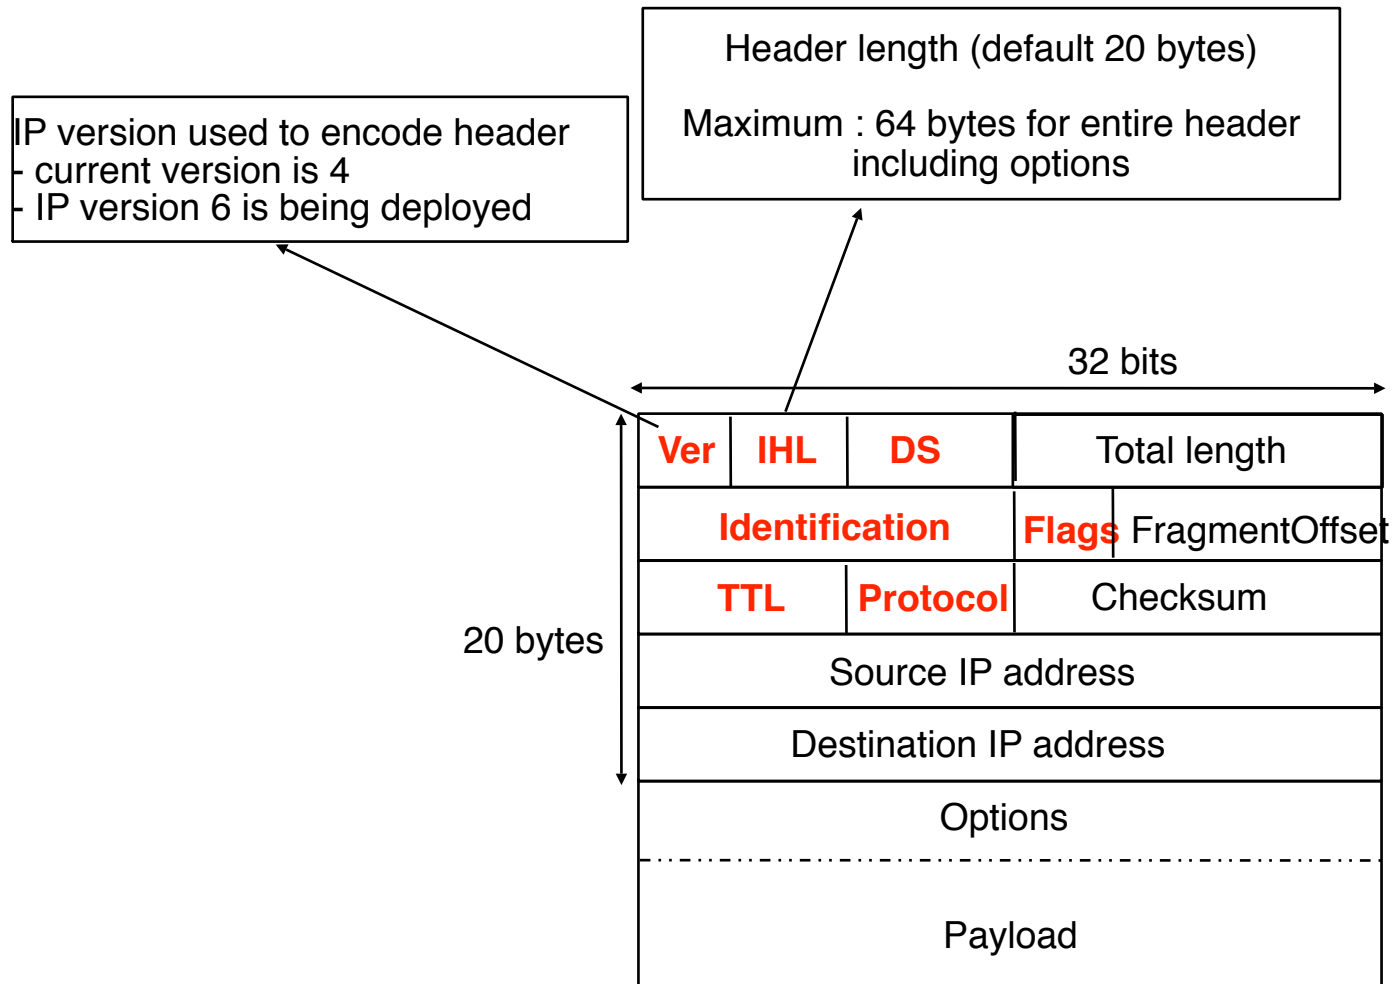
# IP header format

IP version used to encode header

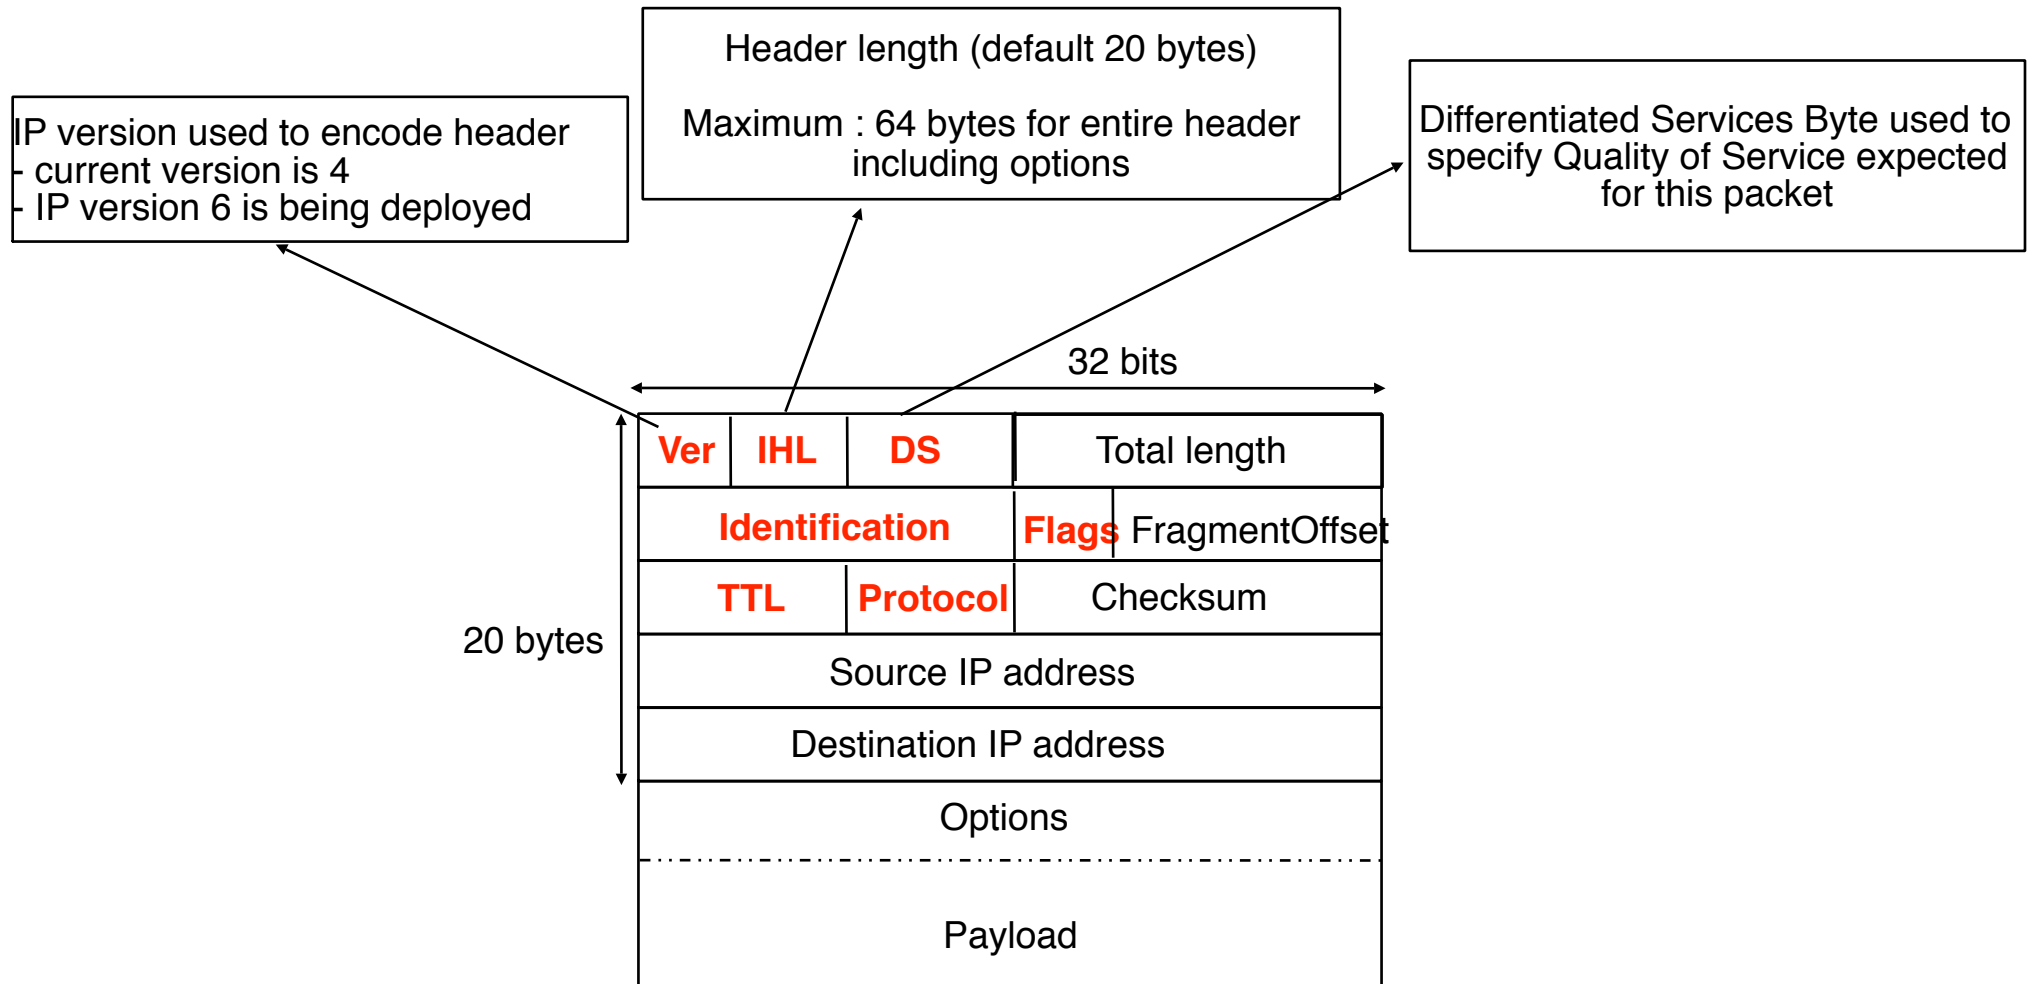
- current version is 4
- IP version 6 is being deployed



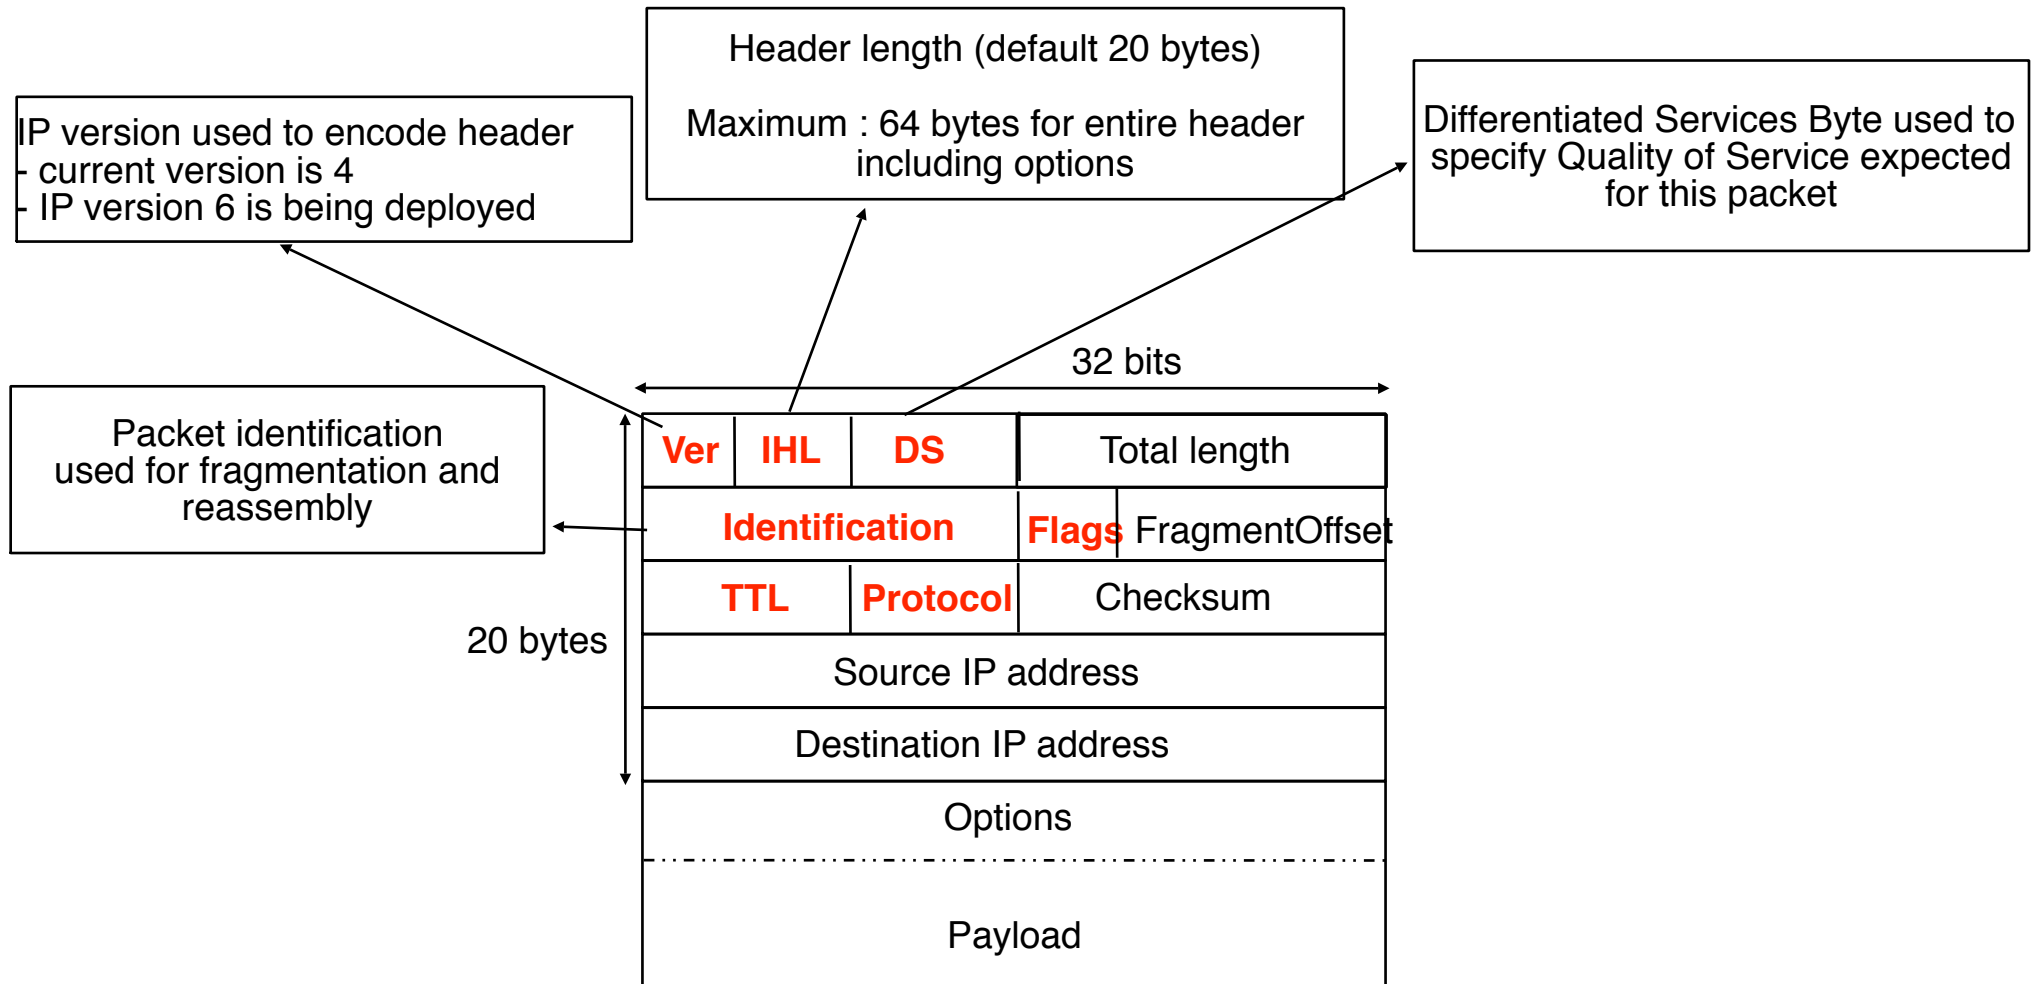
# IP header format



# IP header format

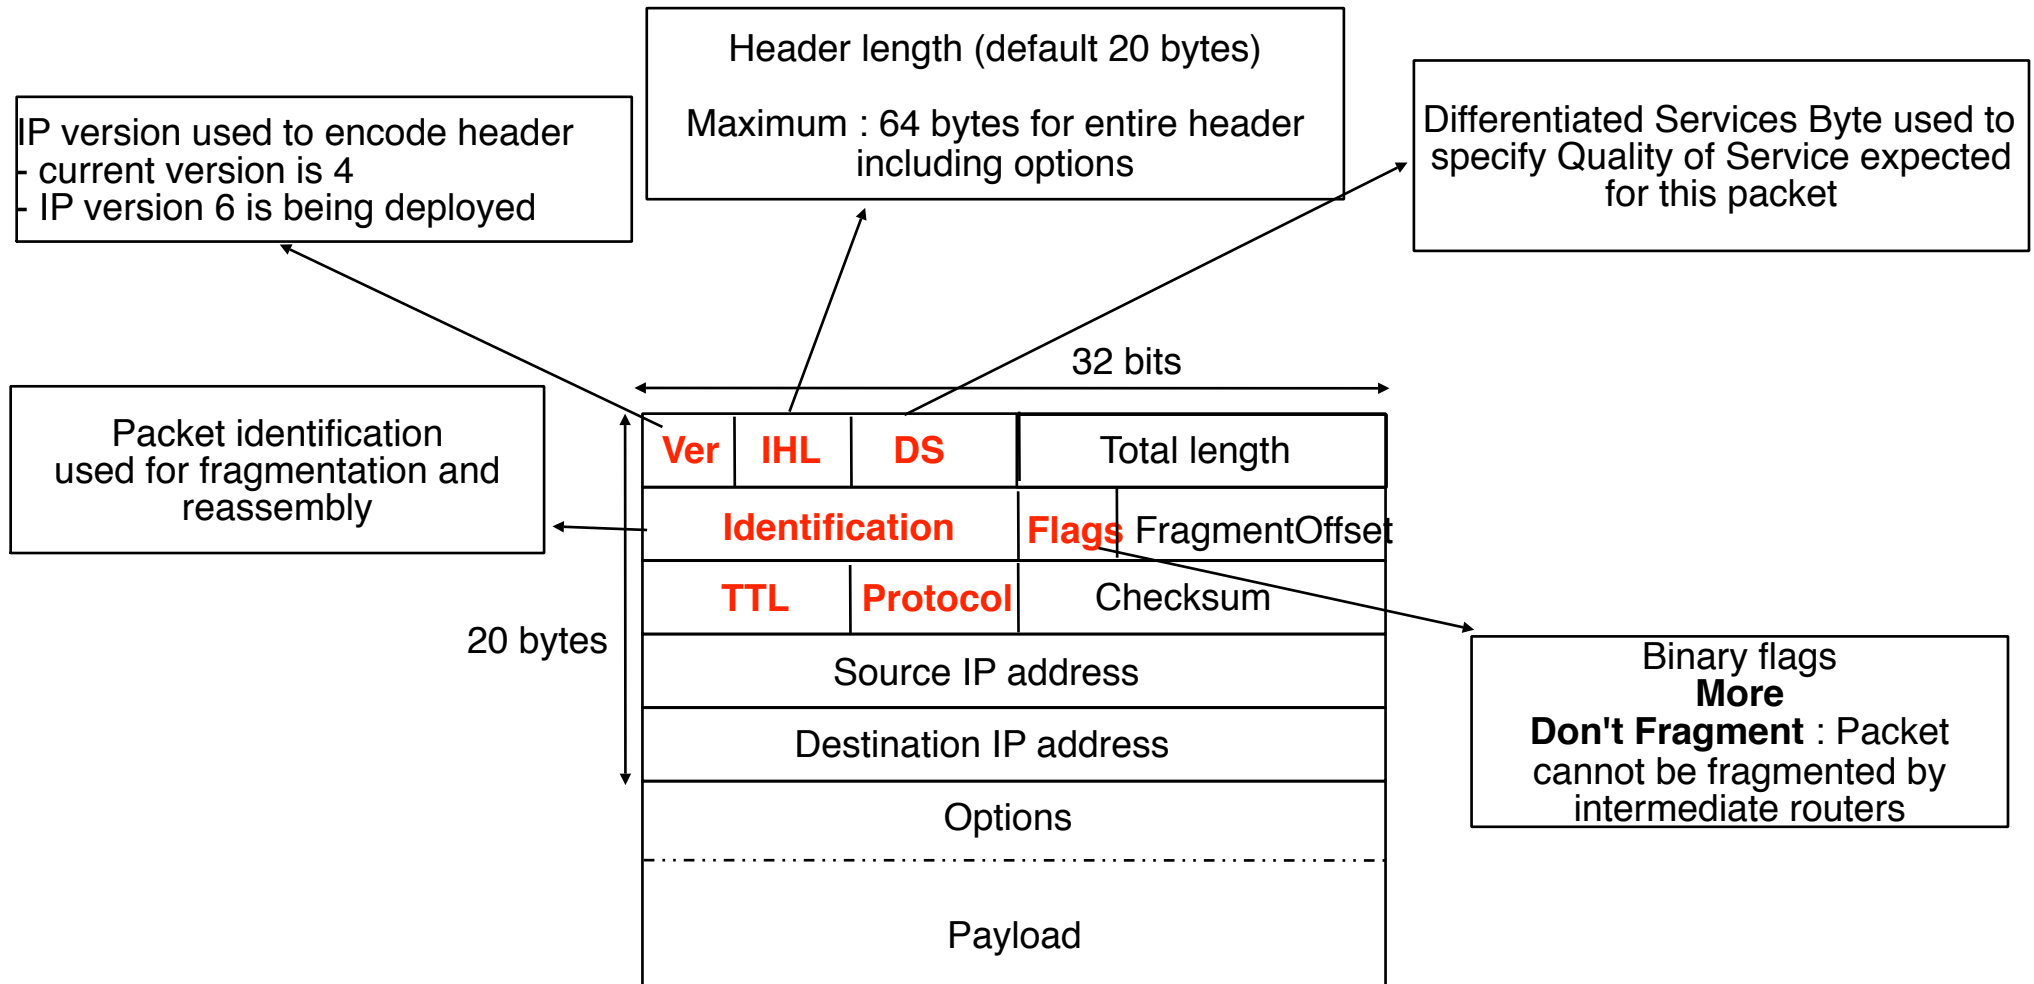


# IP header format

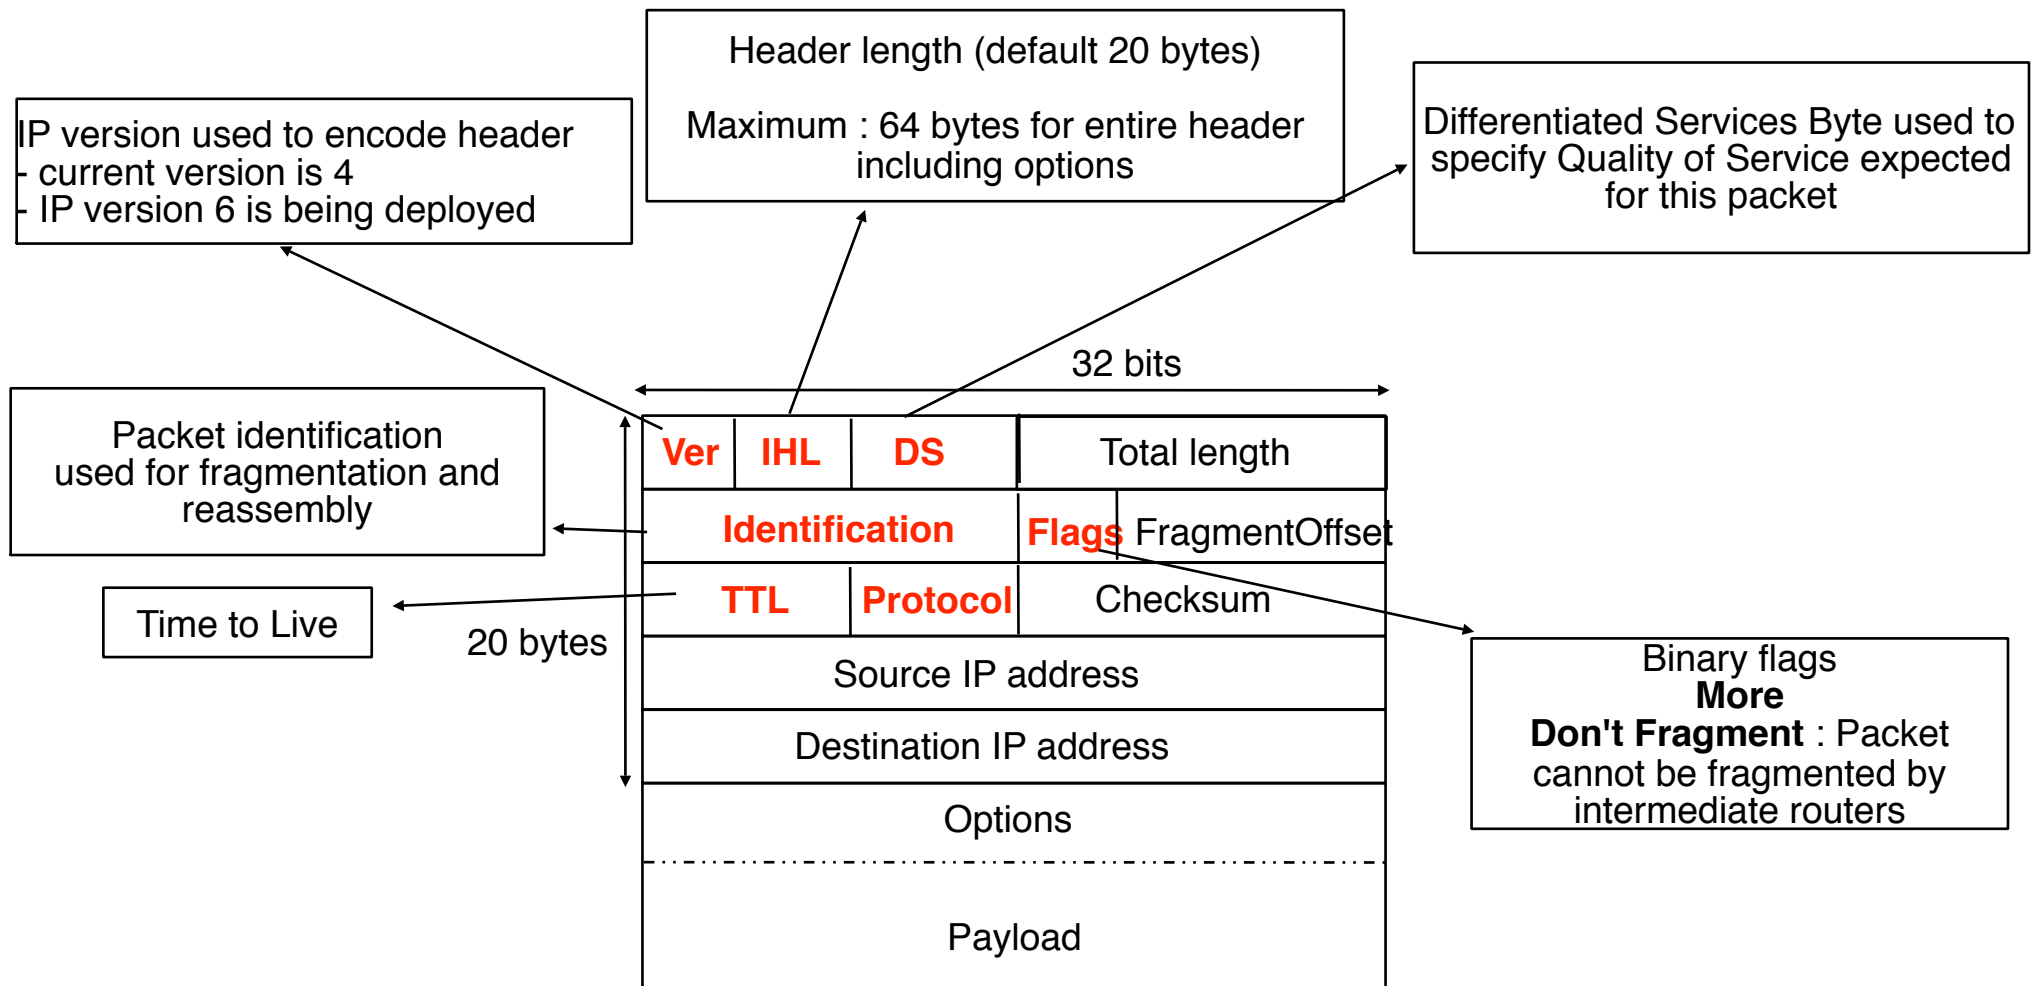




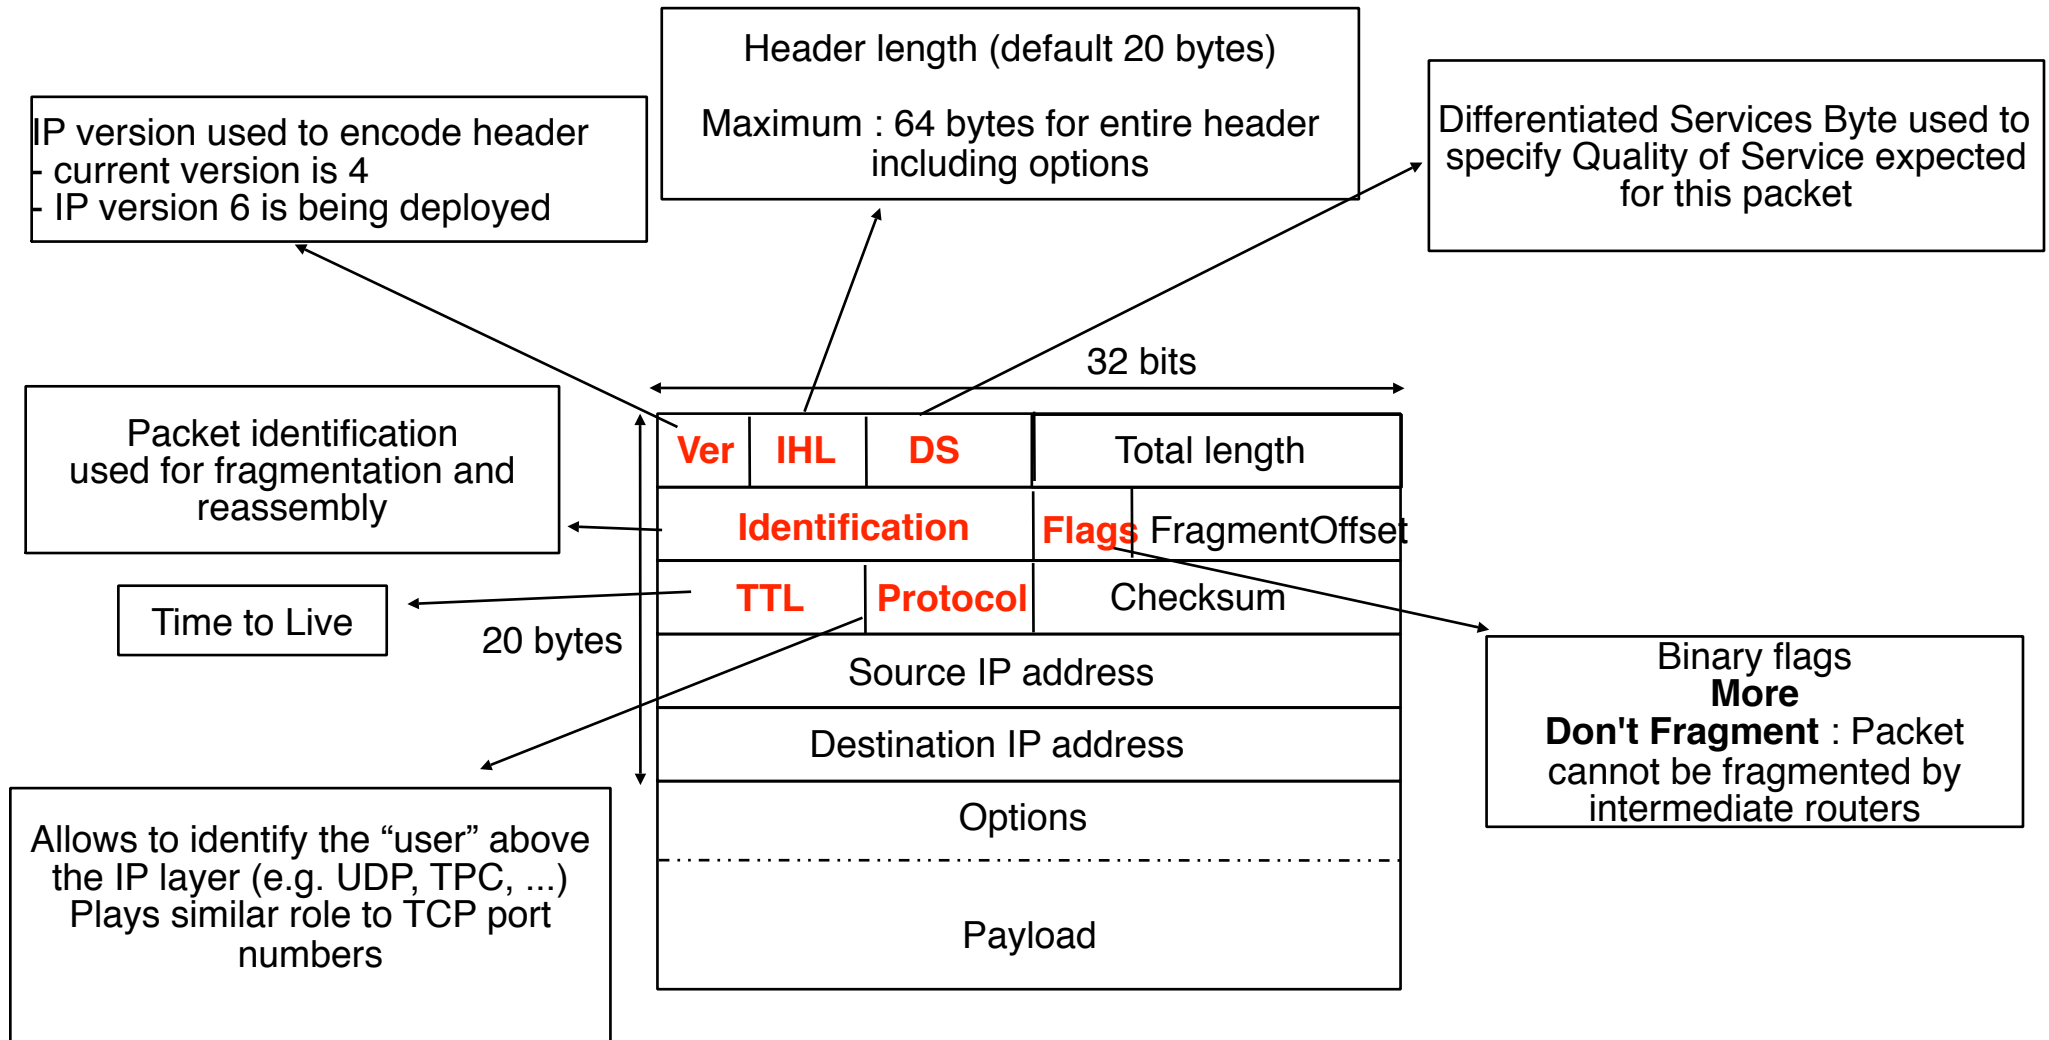
# IP header format



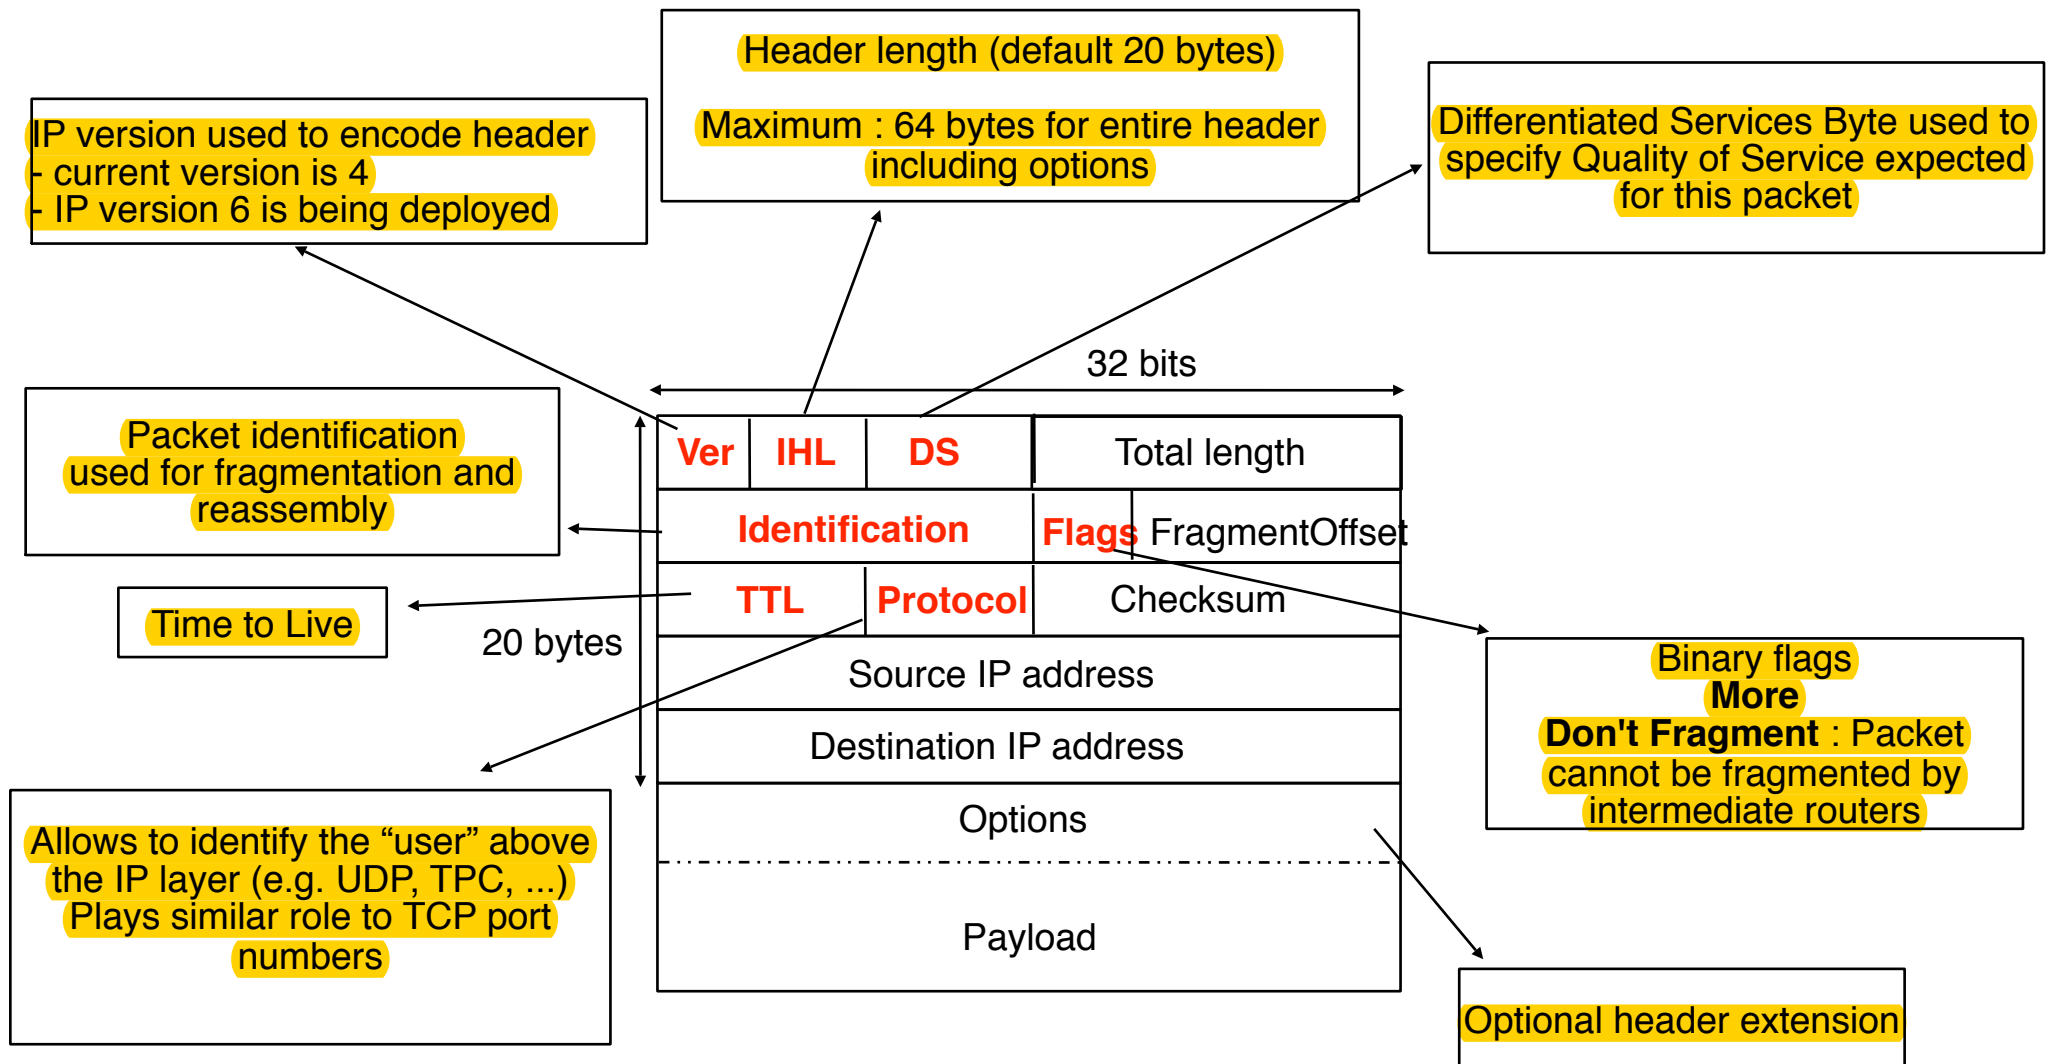
# IP header format



# IP header format



# IP header format



# IP Options

---

## Sample IP header options

### Strict source route option

allows the source to list IP addresses of all intermediate routers to reach destination between source and destination

### Loose source route option

allows the source to list IP addresses of some intermediate routers to reach destination between source and destination

### Record route option

allows each router to insert its IP address in the header  
rarely used because limited header length

### Router alert


allows the source to indicate to routers that there is something special to be done when processing this packet

Constraint : maximum header size with option 64 bytes

# IP Source Routing

---


Strict  
source routing



# IP Source Routing

Source : S  
Destination : D  
Path : RA,RB,RD

Strict  
source routing

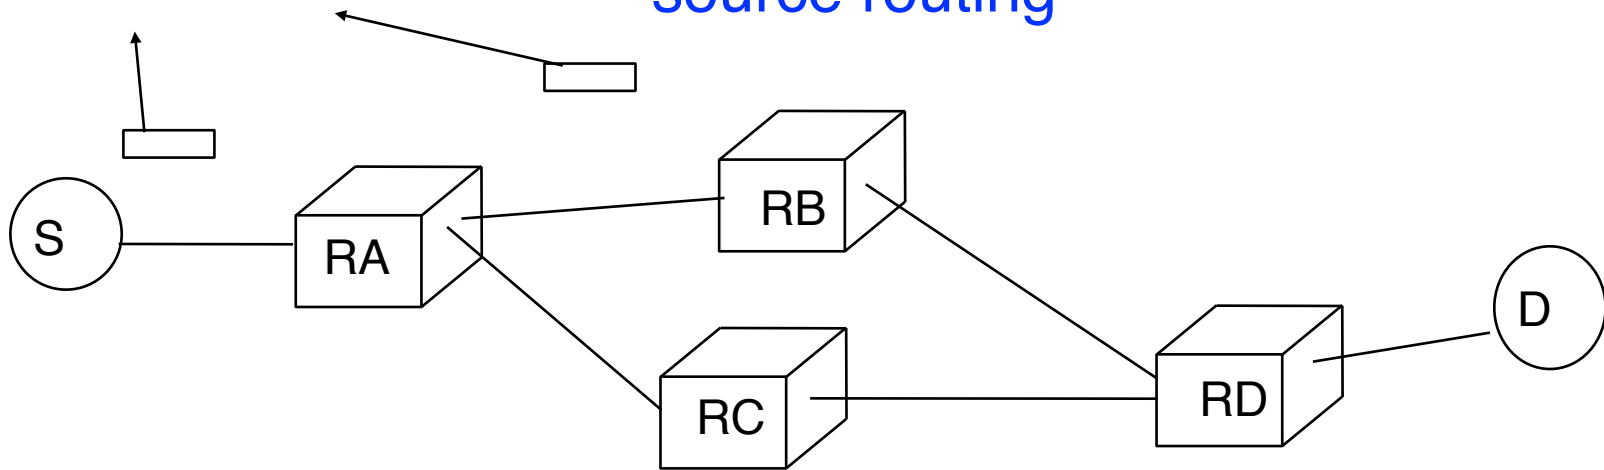


# IP Source Routing

Source : S  
Destination : D  
Path : RA,RB,RD

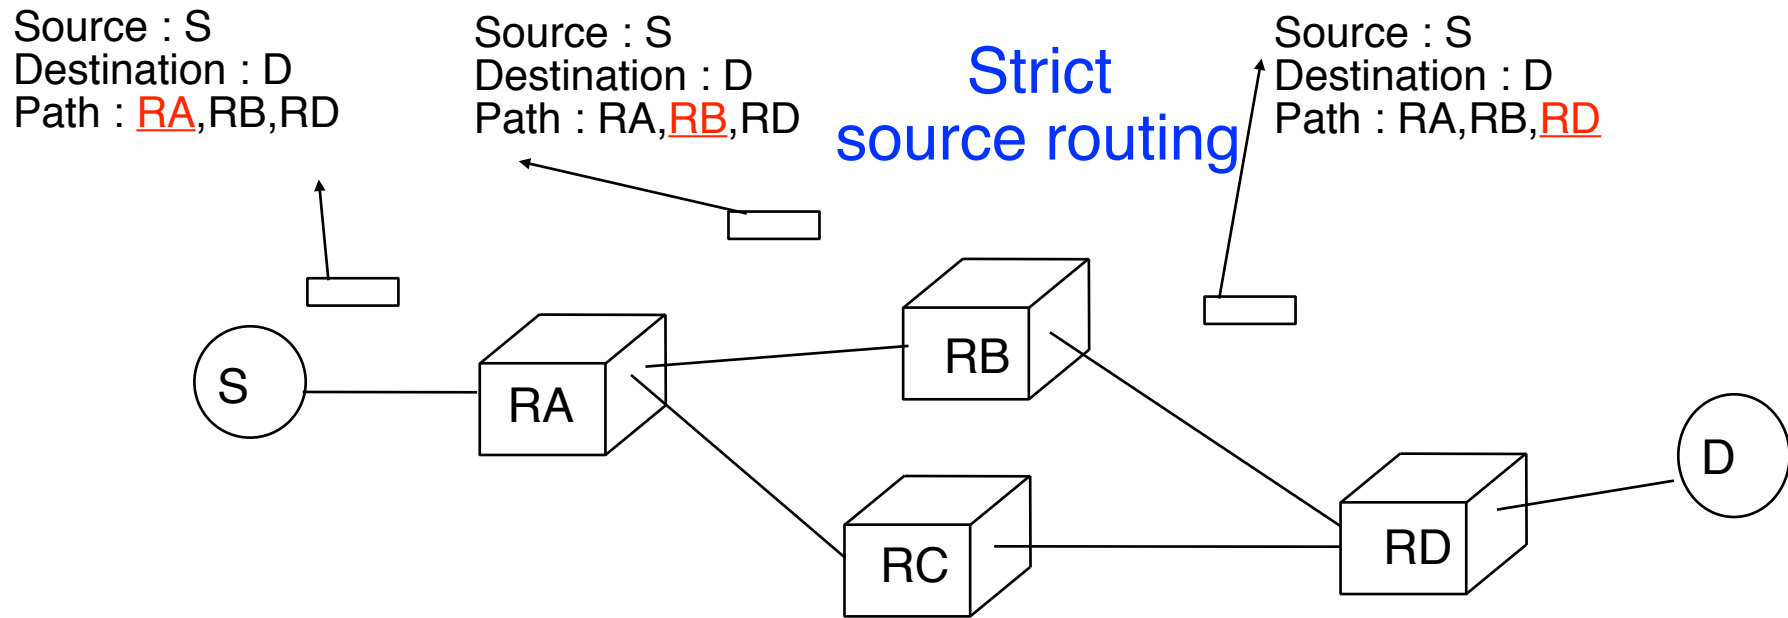
Source : S  
Destination : D  
Path : RA,RB,RD

Strict  
source routing

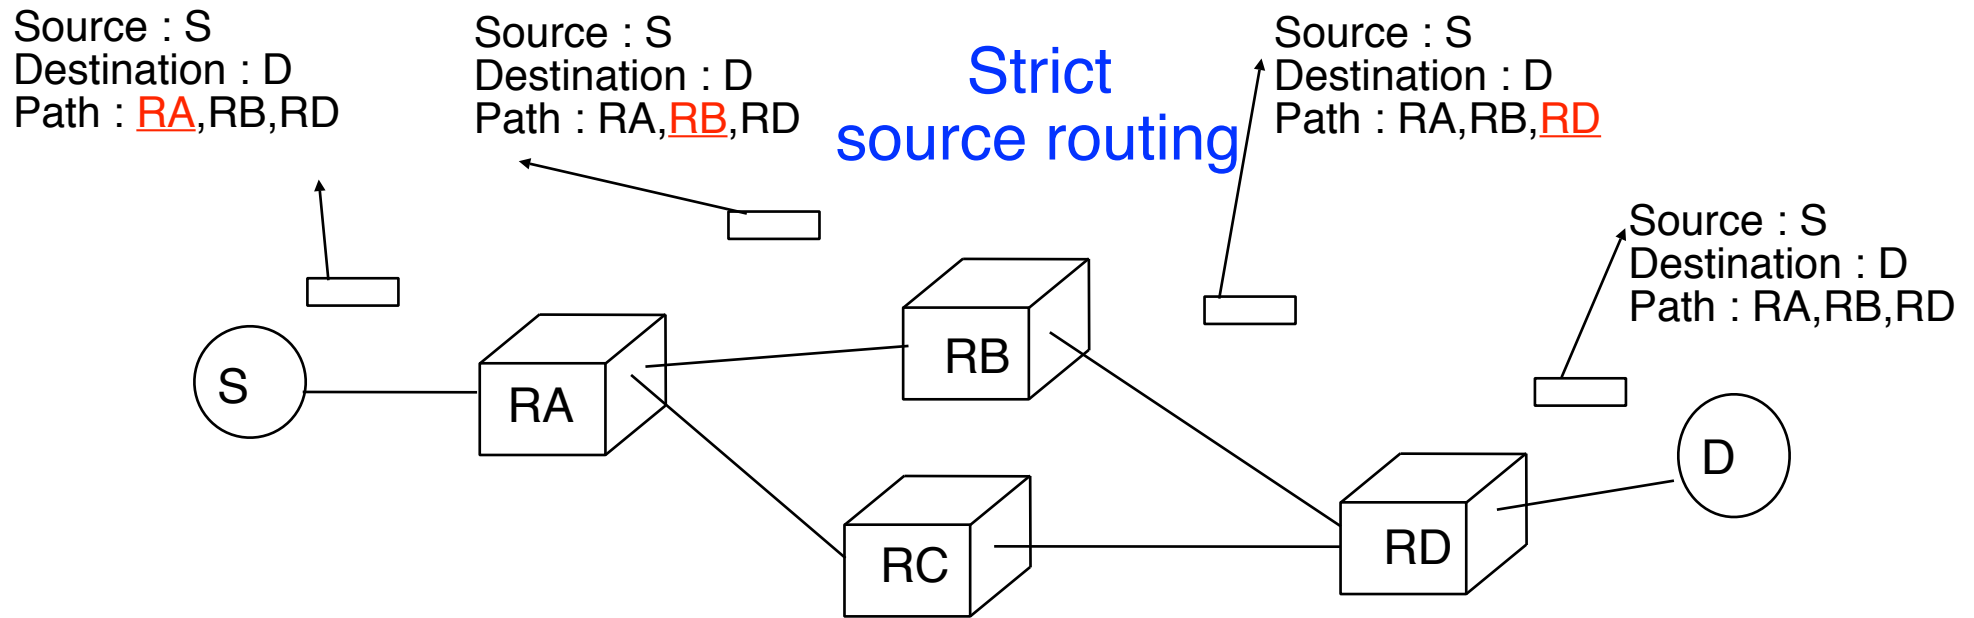




# IP Source Routing



# IP Source Routing



# IP Source Routing


Source : S  
Destination : D  
Path : RA,RB,RD

Source : S  
Destination : D  
Path : RA,RB,RD


Strict  
source routing

Source : S  
Destination : D  
Path : RA,RB,RD

Source : S  
Destination : D  
Path : RA,RB,RD



Loose  
source routing



# IP Source Routing


Source : S  
Destination : D  
Path : RA,RB,RD

Source : S  
Destination : D  
Path : RA,RB,RD

**Strict  
source routing**


Source : S  
Destination : D  
Path : RA,RB,RD

Source : S  
Destination : D  
Path : RA,RB,RD

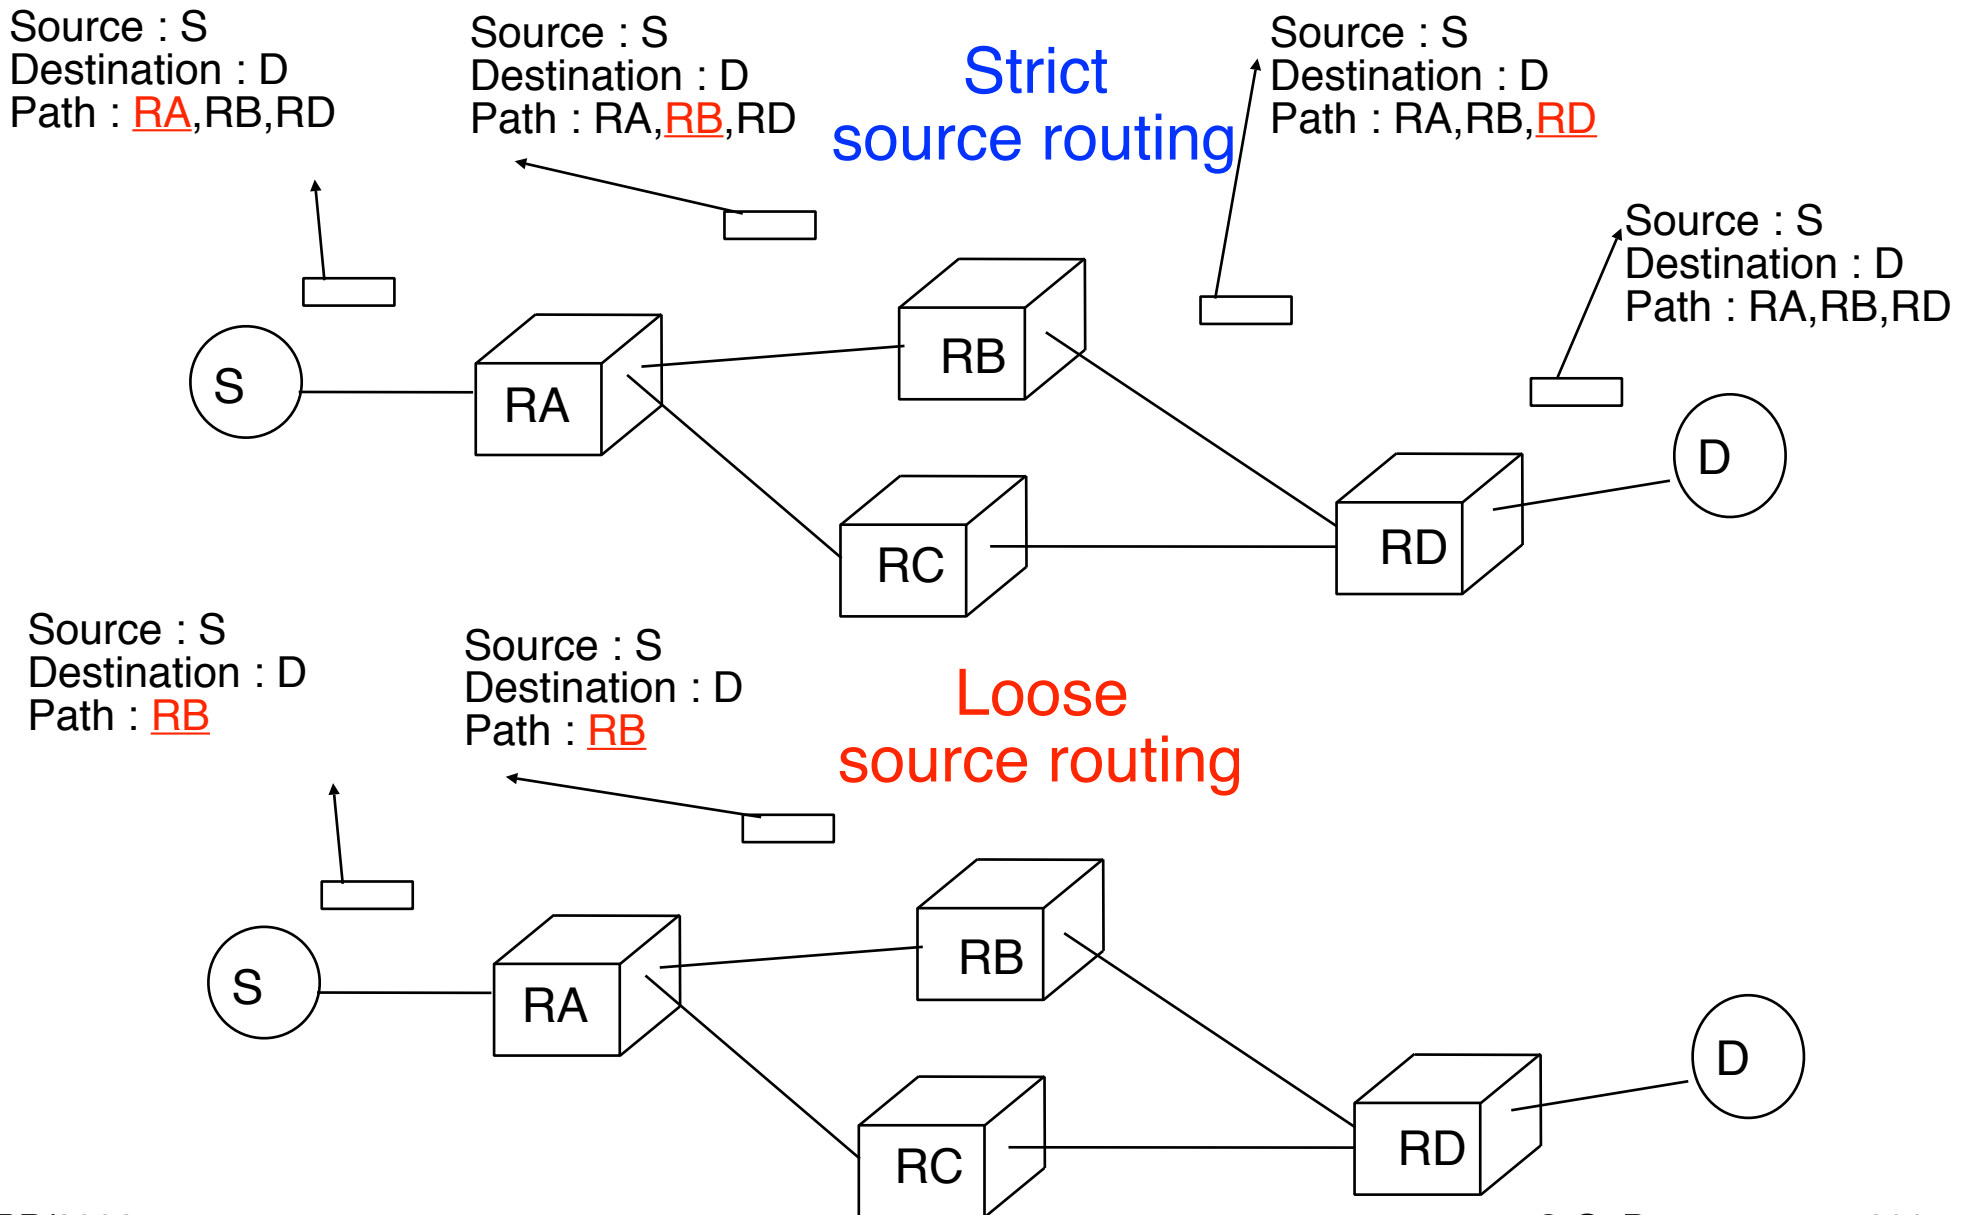


Source : S  
Destination : D  
Path : RB

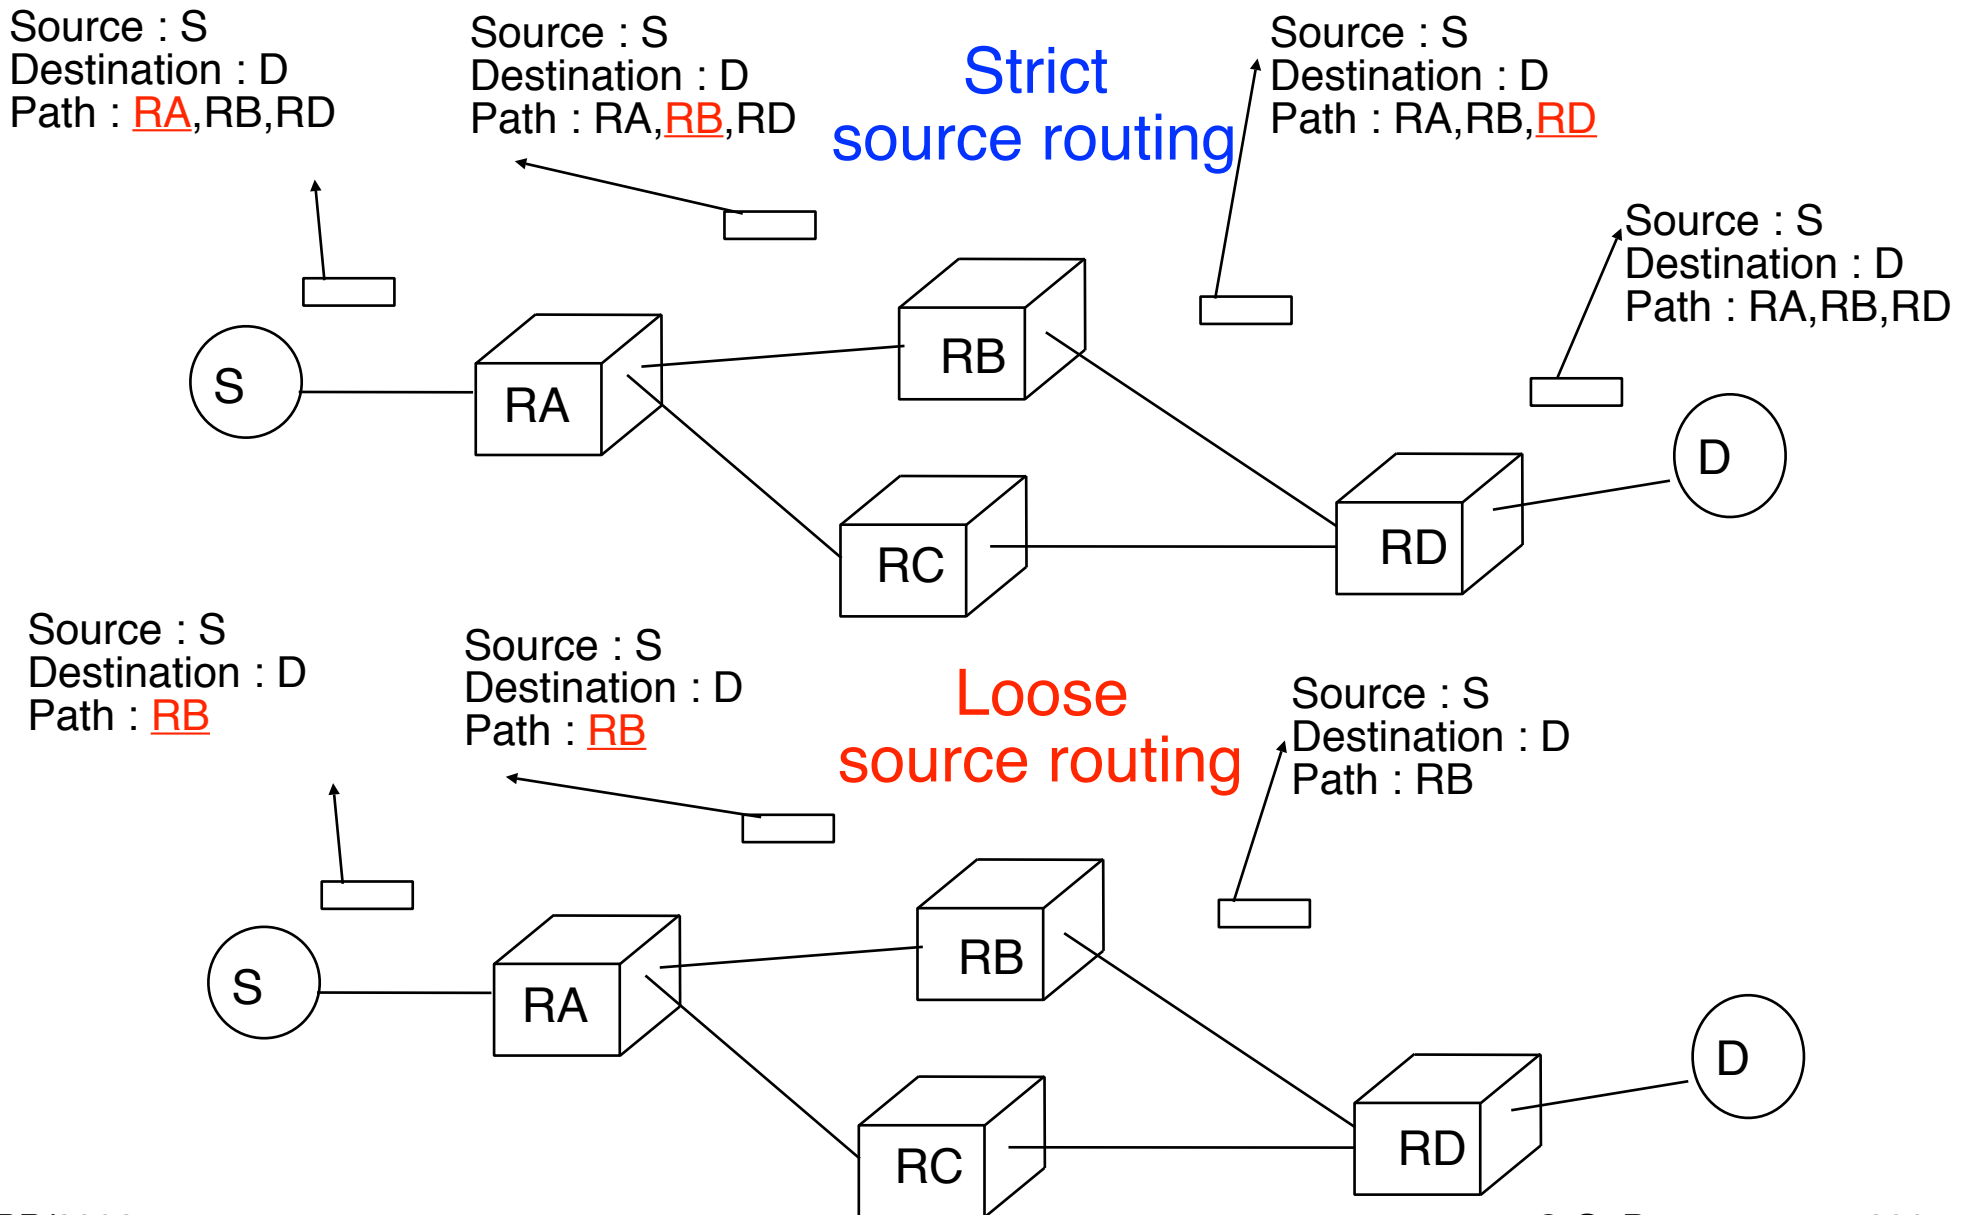
**Loose  
source routing**



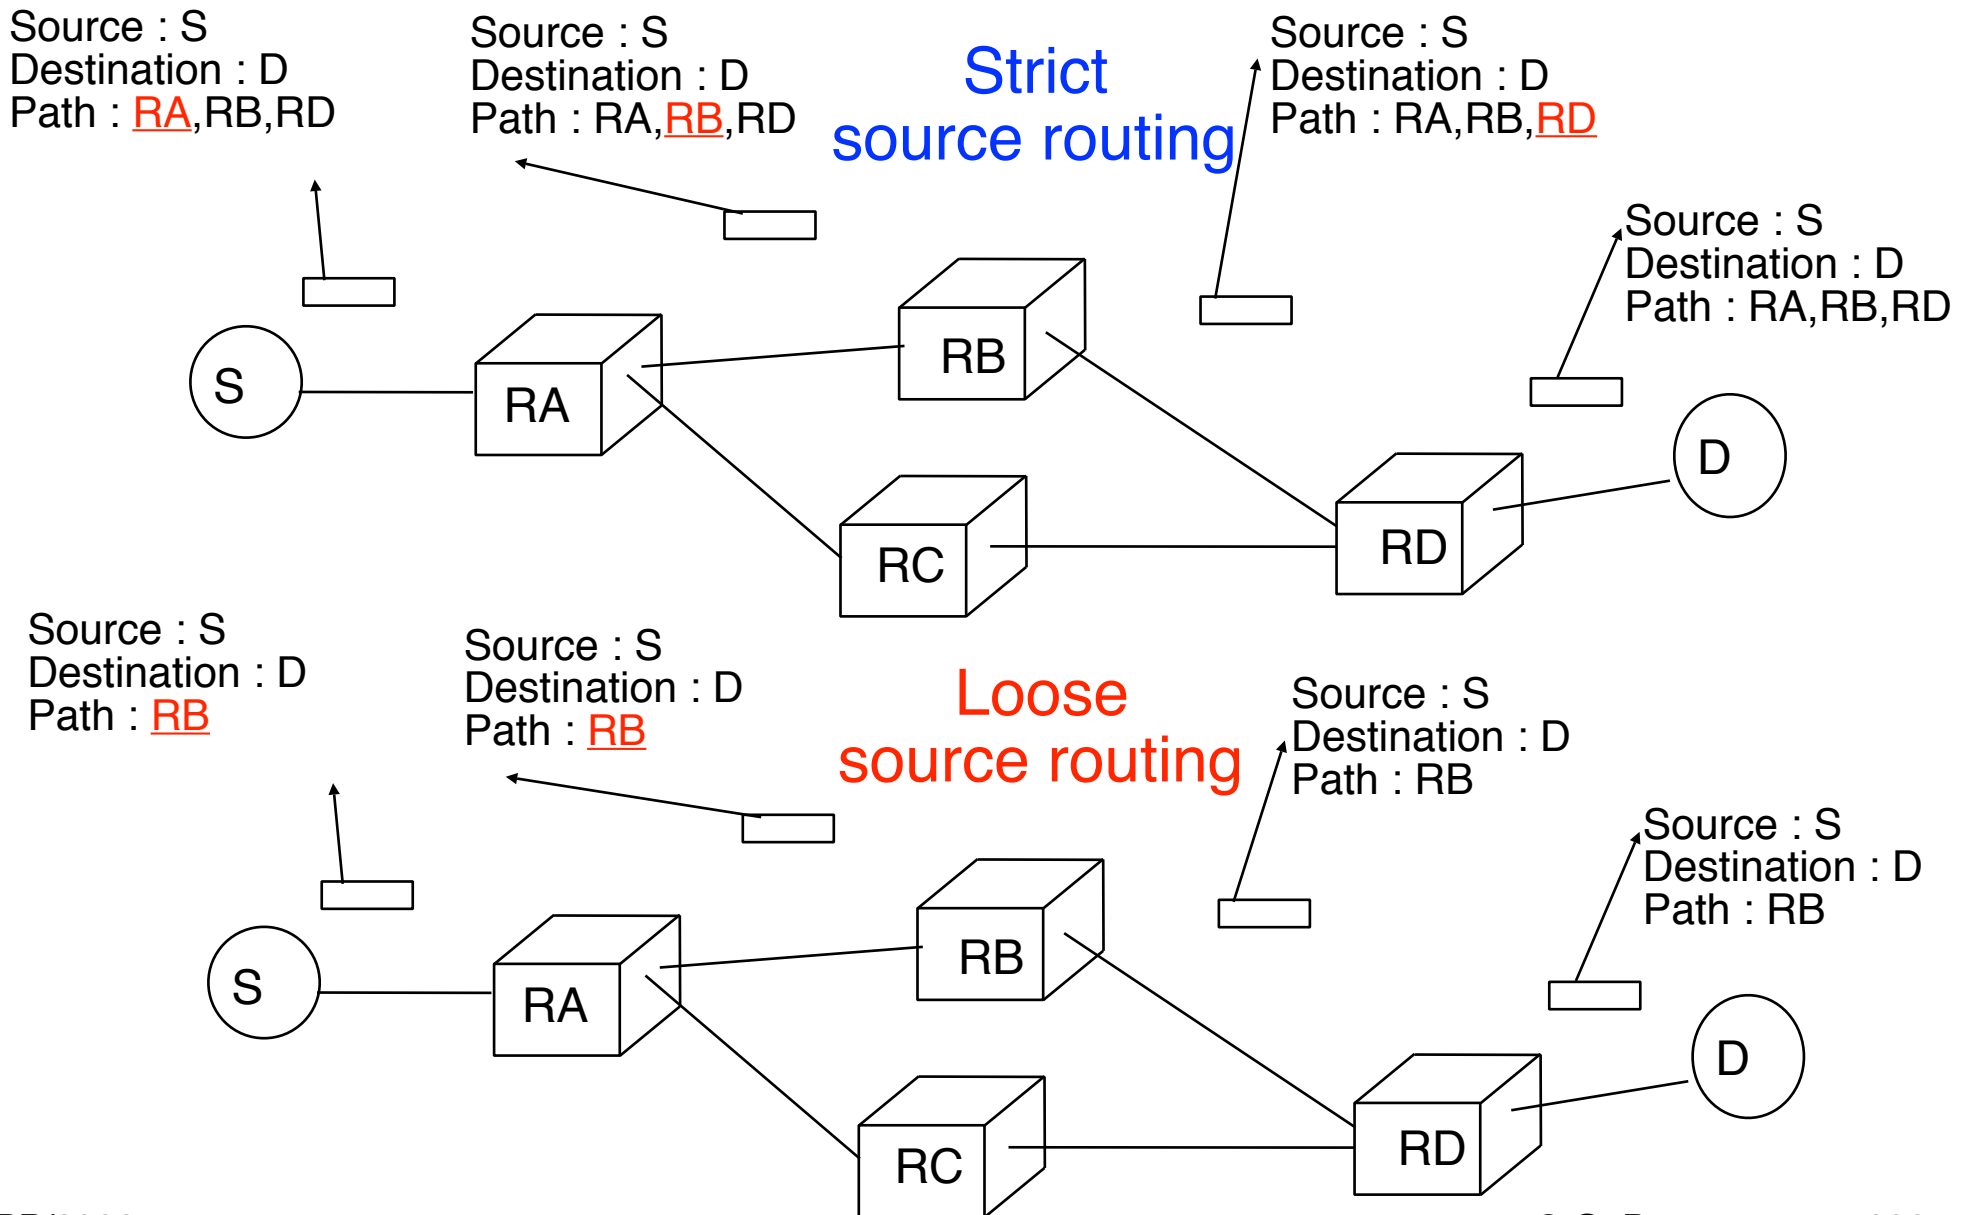
# IP Source Routing



# IP Source Routing



# IP Source Routing



# Operation of an IP endhost

---

## Required information on an IP endhost

### IP addresses of its interfaces

For each address, the subnet mask allows the endhost to determine the addresses that are directly reachable through the interface

### (small) routing table

#### Directly connected subnets

From the subnet mask of its own IP addresses

#### Default router

Router used to reach any unknown address

By convention, default route is 0.0.0.0/0

#### Other subnets known by endhost

Could be manually configured or learned through routing protocols  
are special packets (see later)



# IP address configuration

---

## How does a host know its IP address

### Manual configuration

- Used in many small networks

### Server-based autoconfiguration RARP

#### DHCP

- Dynamic Host Configuration Protocol

- Principle

- When it attaches to a subnet, endhost broadcasts a request to find DHCP server

- DHCP server replies and endhost can contact it to obtain IP address

- DHCP server allocates an IP address for some time period and can also provide additional information (subnet, default router, DNS resolver, ...)


- DHCP servers can be configured to always provide the same IP address to a given endhost or not

- Endhost reconfirms its allocation regularly

### Serverless autoconfiguration

- Used by IPv6

# Example



## Local addresses

11.0.0.1 ; 138.48.0.1 ; 192.168.1.1 ; 10.0.0.1

## Routing table

138.48.0.0/16 [North]  
11.0.0.0/8 [West]  
192.168.1.0/24 [East]  
10.0.0.0/8 [South]  
4.0.0.0/8 via 11.0.1.1 [West]  
4.10.11.0/24 via 10.0.1.1 [South]  
12.0.0.0/8 via 138.48.1.2 [North]  
0.0.0.0/0 via 138.48.1.1 [North]

# Operation of an IP router

---

## Required information on an IP router

### IP addresses of its interfaces

For each address, the subnet mask allows the endhost to determine the addresses that are directly reachable through the interface

### Routing table

#### Directly connected subnets

From the subnet mask of its own IP addresses

#### Other known subnets

Usually learned via routing protocols, sometimes manually configured

#### Default router

Router used to reach any unknown address

By convention, default route is 0.0.0.0/0

# Operation of an IP router (2)

---


## Operations performed for each packet

1. Check whether the packet's destination address is one of the router's addresses  
If yes, packet reached destination
2. Query Forwarding Information Base that contains  
list of directly connected networks with masks  
list of reachable networks and intermediate router
3. Lookup the most **specific route** in FIB  
For each route A.B.C.D/M via Rx  
compare M higher order bits of destination address with  
M higher order bits of routes to find longest match  
forward packet along this route

# Forwarding Information Base Lookup

How to find most specific route ?  
similar to longest prefix match in a text  
Trie

| Subnet        | Prefix | Next-hop |
|---------------|--------|----------|
| 138.48.0.0    | 16     | A        |
| 139.165.0.0   | 16     | B        |
| 139.165.16.0  | 24     | C        |
| 138.48.232.0  | 24     | E        |
| 138.48.32.200 | 32     | G        |
| 0.0.0.0       | 0      | F        |



Cost of lookup  
f(average length of prefixes)  
comparisons  
memory accesses  
caches for most frequently used routes

# Handling IP packets in error

---

## Problem

What should a router/host do when it receives an errored packet

### Example

- Packet whose destination is not the current endhost

- Packet containing a header with invalid syntax

- Packet received with TTL=1

- Packet destined to protocol not supported by host

## Solutions

- Ignore and discard the errored packet

- Send a message to the packet's source to warn it about the problem

  - ICMP : Internet Control Message Protocol

  - ICMP messages are sent inside IP packets by routers (mainly) and hosts

    - To avoid performance problems, most hosts/routers limit the amount of ICMP messages that they send

# Sample ICMP messages

---

## Routing error

### Destination unreachable

- Final destination of packet cannot be reached

  - Network unreachable for entire subnet

  - Host unreachable for an individual host

  - Protocol/Port unreachable for protocol/port on a reachable host

### Redirect

- The packet was sent to an incorrect first-hop router and should have been instead sent to another first-hop router

## Error in the IP header

### Parameter Problem

- Incorrect format of IP packet

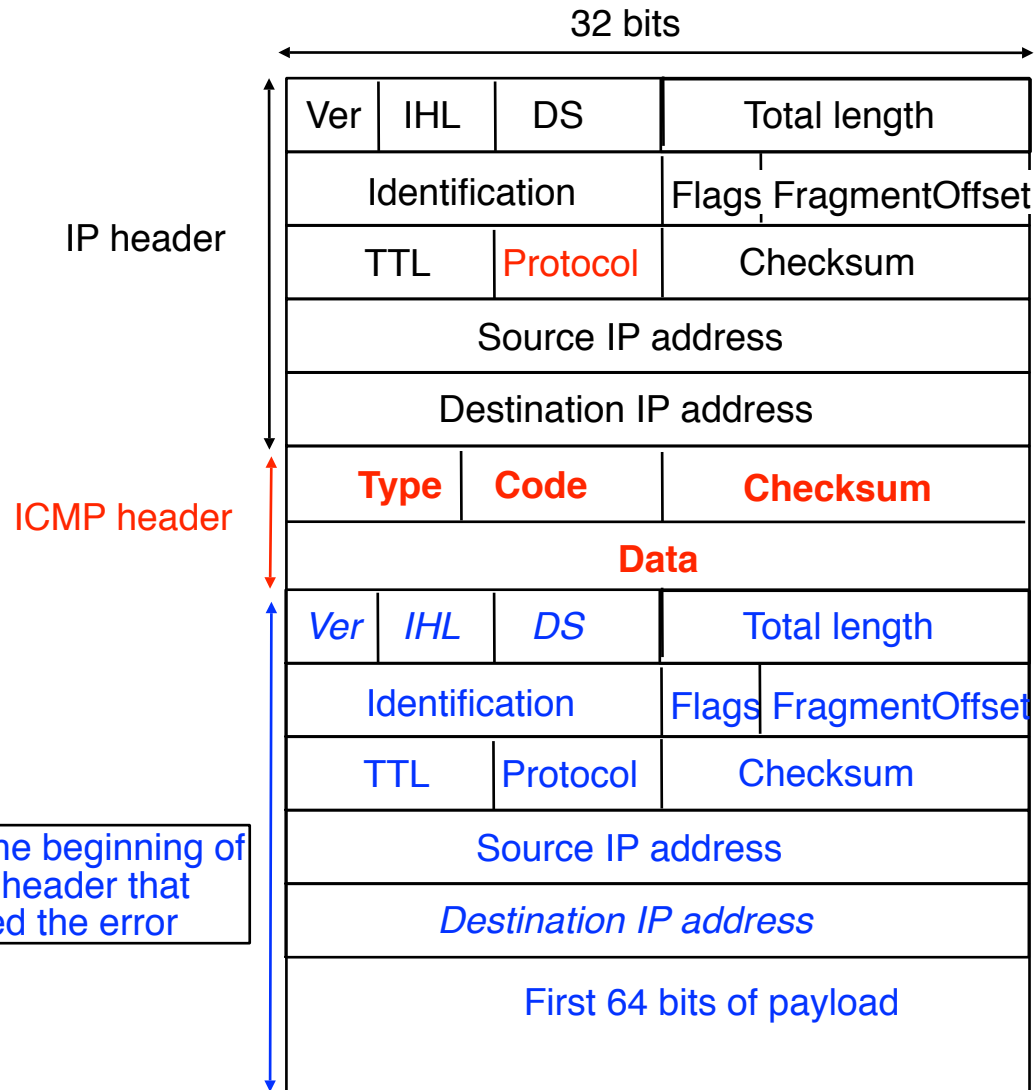
### TTL Exceeded

- Router received packet with TTL=1

### Fragmentation

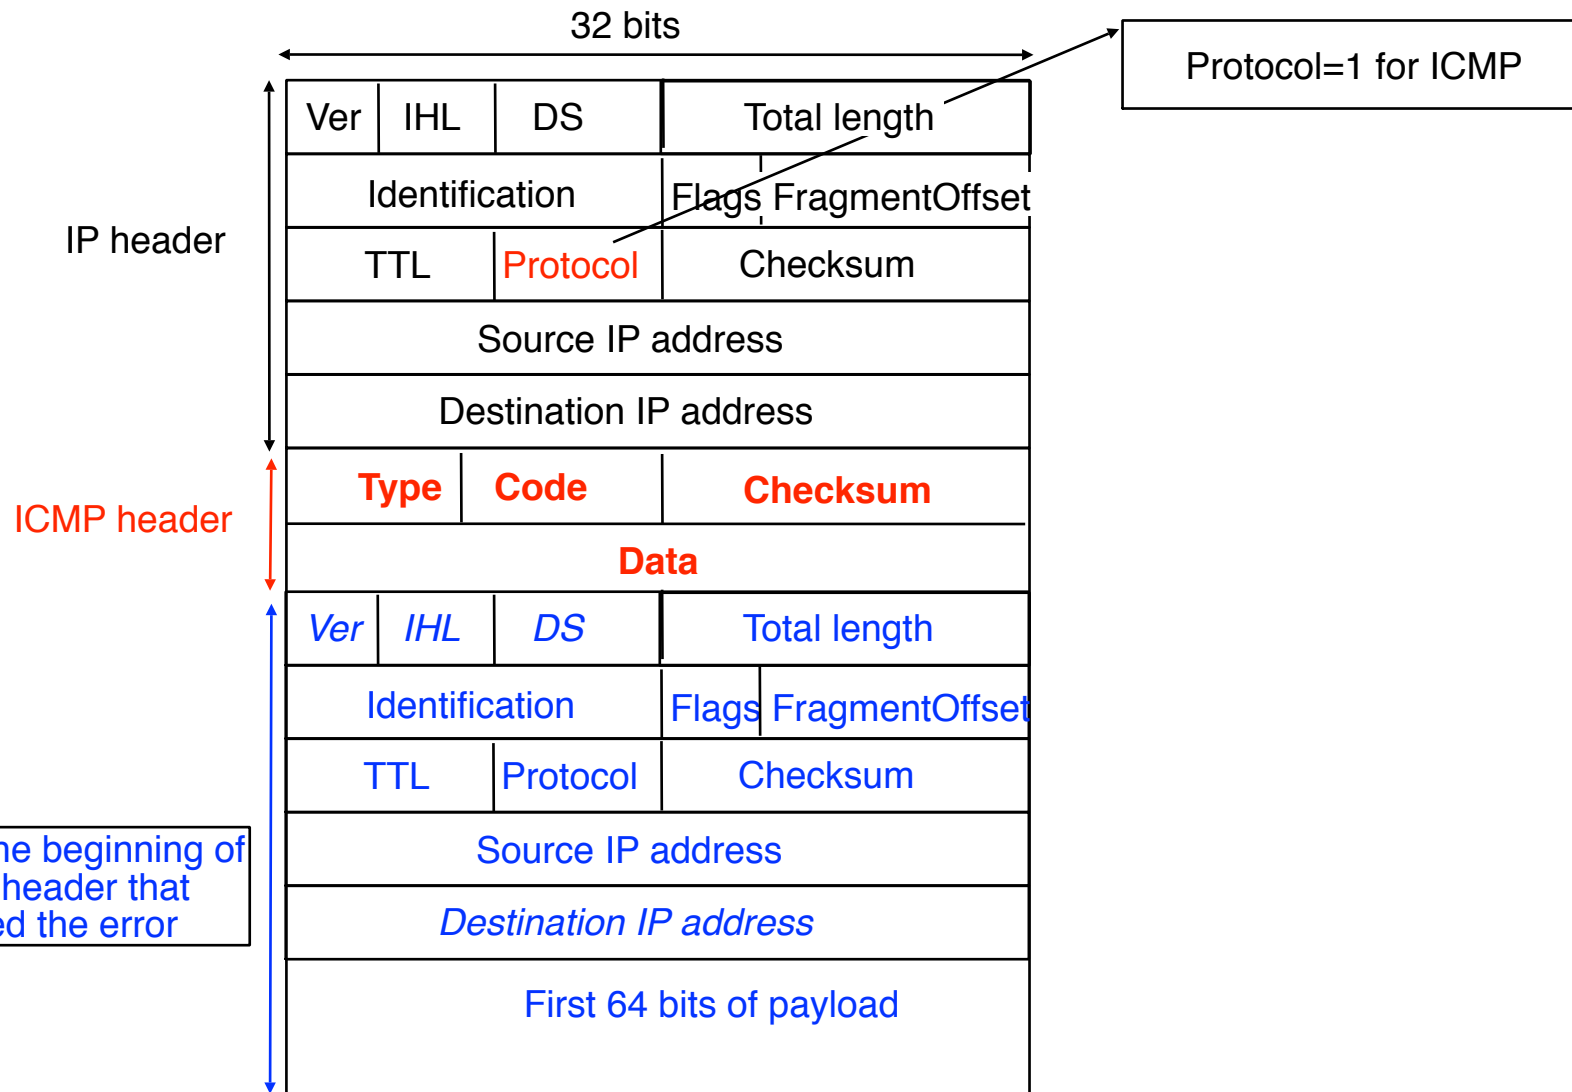
- the packet should have been fragmented, but its DF flag was true

# ICMP messages

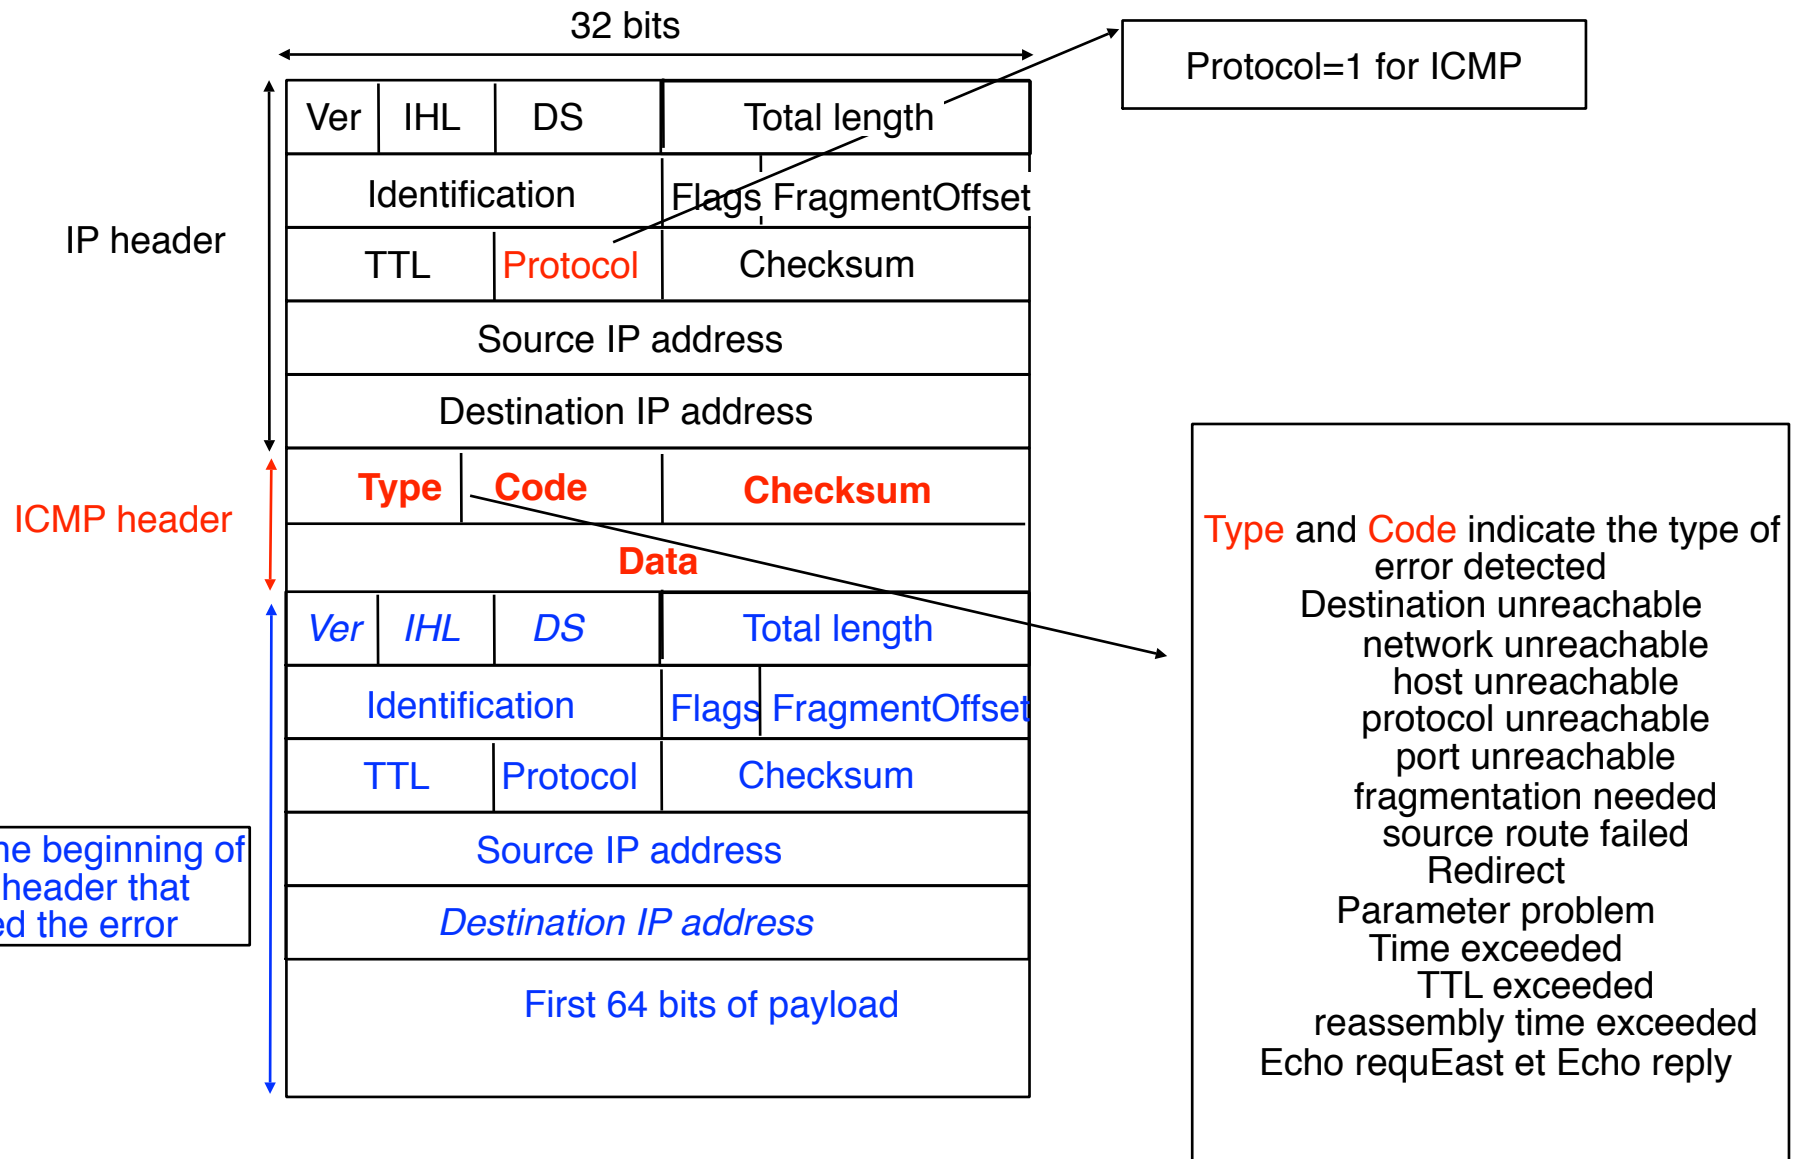




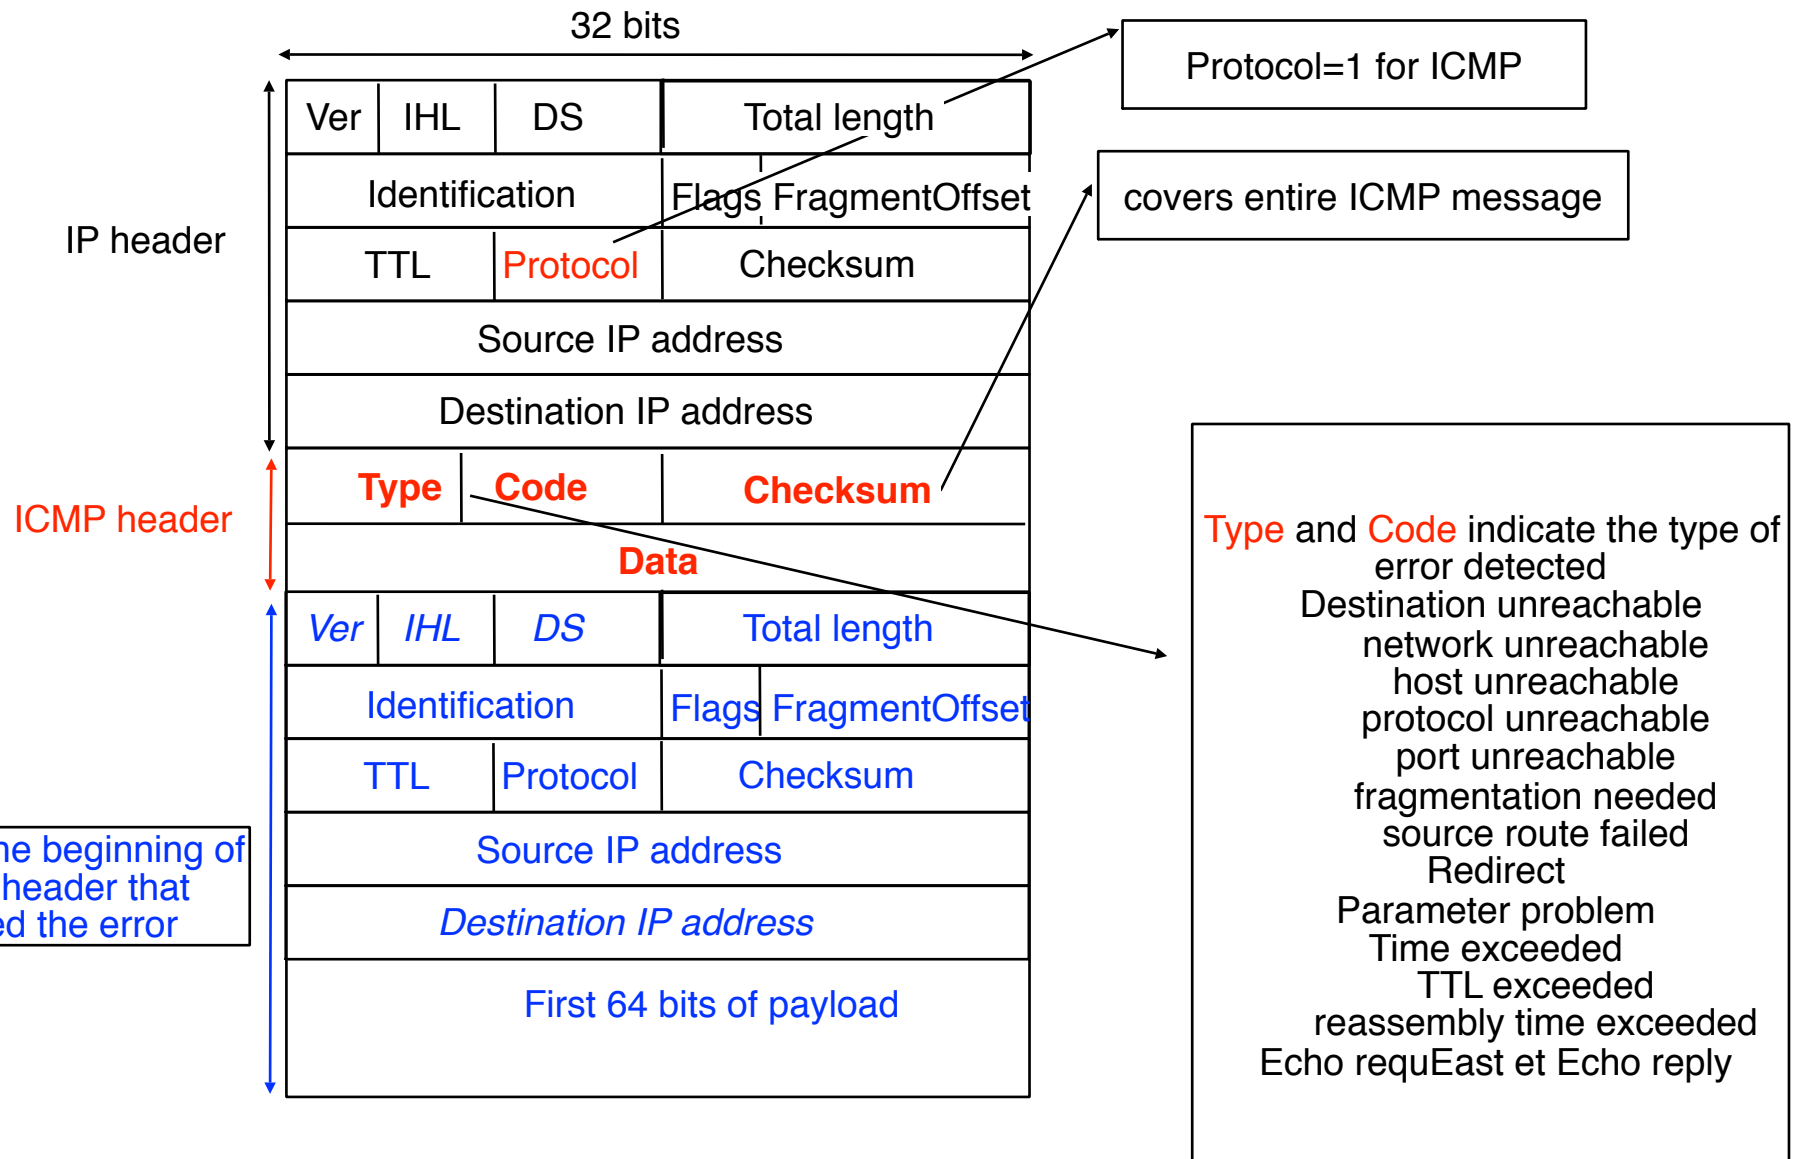
# ICMP messages



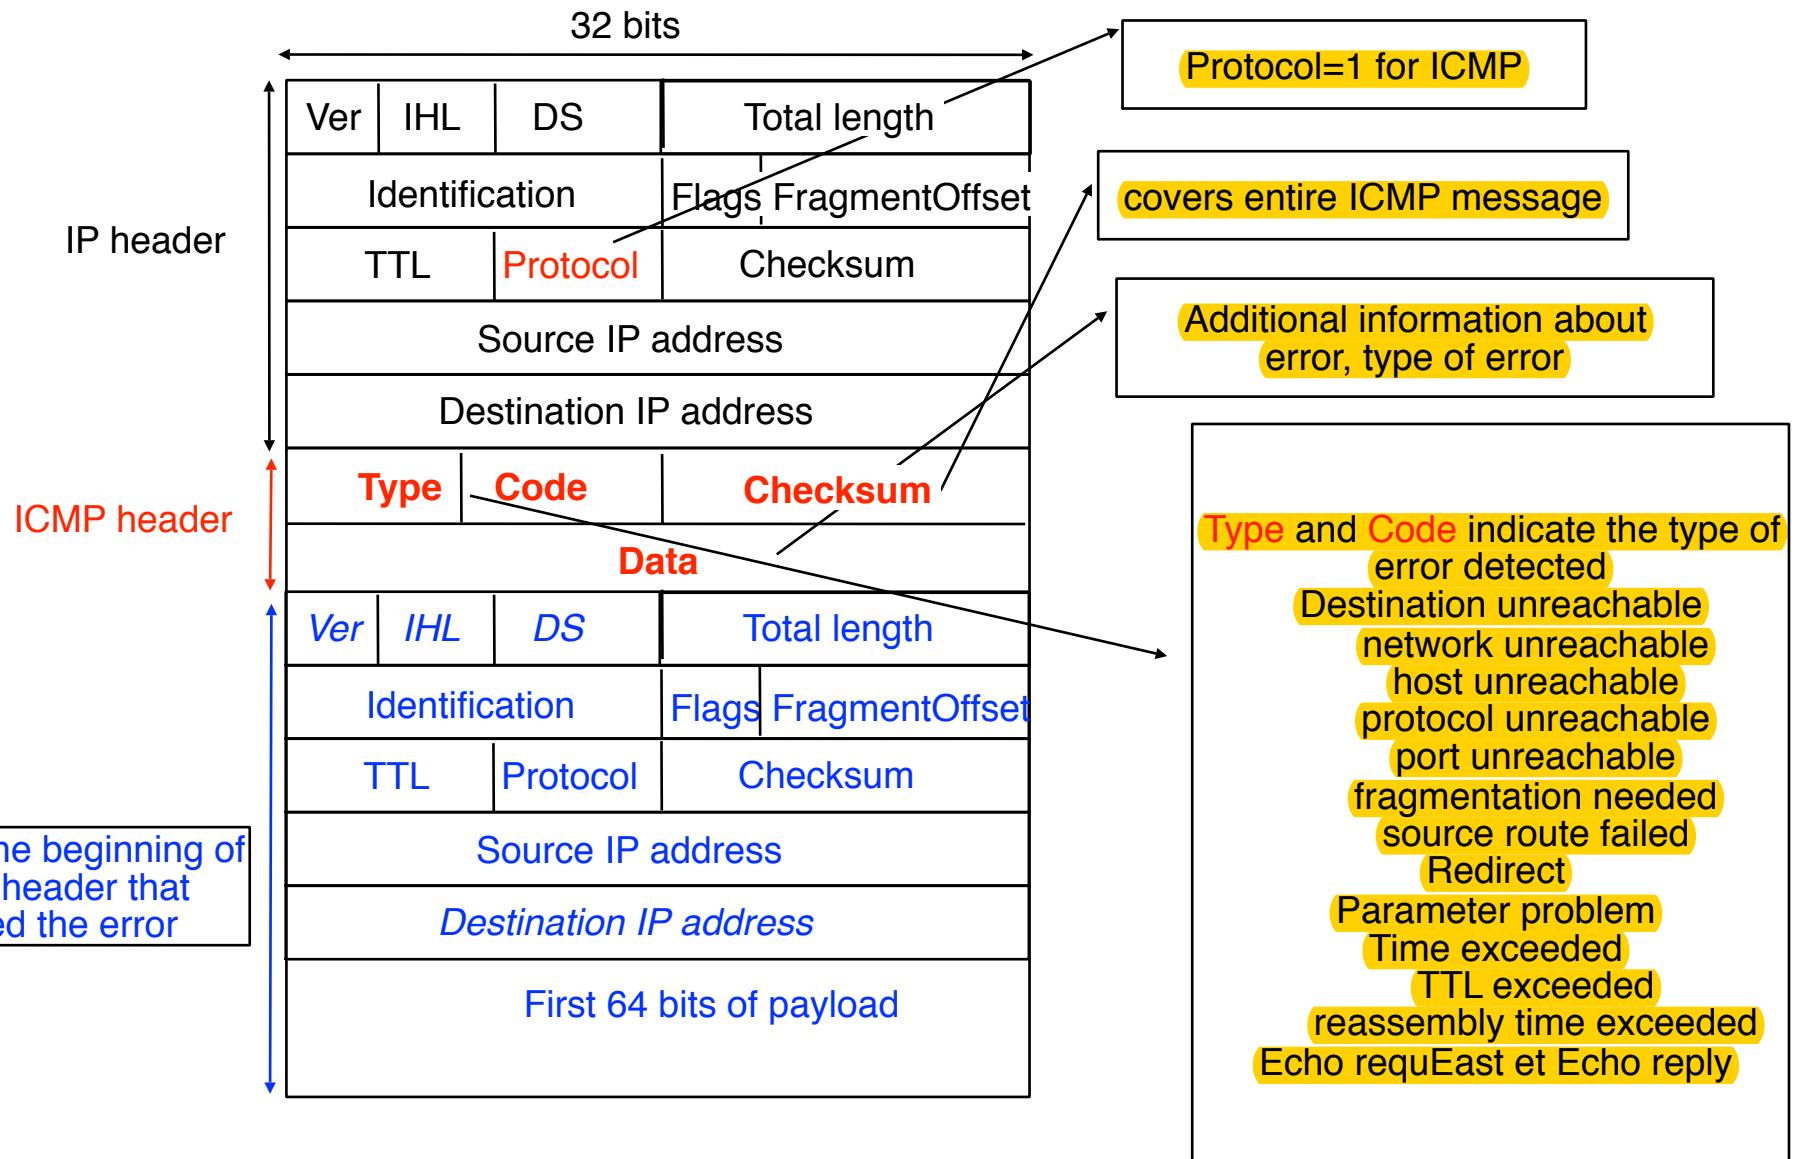
# ICMP messages



# ICMP messages



# ICMP messages



# Usage of ICMP messages

---

## Examples

### destination unreachable

the router sending this message did not have a route to reach the destination

### time exceeded

the router sending the message received an IP packet with TTL=0

used by `traceroute`

### redirect

to reach destination, another router must be used and ICMP message provides address of this router

### echo request / echo reply

used by `ping`

### fragmentation impossible

the packet should have been fragmented by the router sending the ICMP message by this packet had "Don't Fragment" set to true

# Middleboxes

---

The original TCP/IP architecture only defined hosts and routers

Today's networks contain devices that  
process  
analyse  
and possibly modify IP packets

## Examples

- Firewall

- Network Address Translator

- Traffic shaper


- Deep Packet Inspection

- Intrusion Detection System

- Load balancer

# Firewall

## Problem



How to control the packets entering the enterprise network ?

- only allow external access to some servers

- allow internal clients to use web

- allow smtp (inbound and outbound) only on some dedicated servers

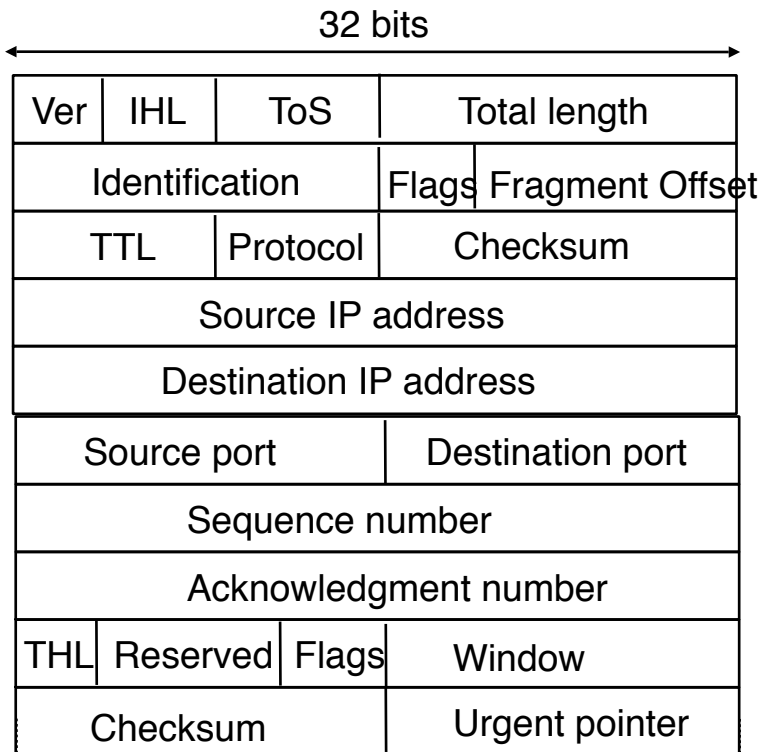
...

# Firewalls (2)

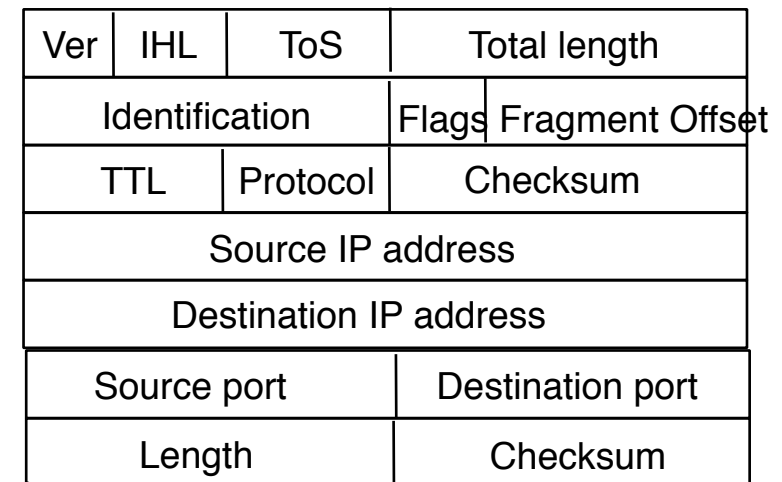
## Principle

Firewall analyses all packet headers  
rules specify which packets should be accepted  
and rejected

TCP



32 bits



UDP



# Firewalls : example

---

## Wifi UCLouvain

### Outbound

Standard IPsec VPN...

SSH: TCP/22

HTTP: TCP/80, HTTPS: TCP/443

IMAP2+4: TCP/143, IMAP3: TCP/220, IMAPS: TCP/993, POP: TCP/110, POP3S: TCP/995, SMTPS: TCP/465, SMTP submit with STARTTLS: TCP/587

Passive (S)FTP: TCP/21

RDP: TCP/3389

IPv6 Tunnel Broker service: IP protocol 41

### Inbound

OpenVPN 2.0: UDP/1194, IPsec NAT-Traversal UDP/4500, PPTP VPN: IP protocol 47 (GRE), Standard

IPsec VPN: IP protocols 50 (ESP) and 51 (AH)

IPv6 Tunnel Broker service: IP protocol 41

# Network Address Translator

## Problem


Limited number of public IPv4 addresses

## Solution


Use private addresses inside enterprise and home networks

Use one or a few public addresses

translate packets sent to public Internet



# Simple enterprise NAT



Operation  
Mapping table


Internal address <-> public address

Packet arrival from internal network

Packet arrival from public network

How long should mapping remain ?

# Single address NAT



## Single address NAT

NAT translates IP addresses and TCP/UDP port numbers

| Private address | Protocol | Port inside | public address  | Port outside |
|-----------------|----------|-------------|-----------------|--------------|
| 192.168.10.10   | UDP      | 2340        | 130.104.228.200 | 4567         |
| 192.168.10.10   | TCP      | 512         | 130.104.228.200 | 520          |
| 192.168.10.11   | TCP      | 1024        | 130.104.228.200 | 2048         |

# Network layer

---

Basics

Routing

IP : Internet Protocol

Routing in IP networks

→ Internet routing organisation

Intradomain routing : RIP

Intradomain routing : OSPF

Interdomain routing : BGP

# Internet organisation

---

Internet is an internetwork with a large number of Autonomous Systems (AS)

an AS is a set of routers that are managed by the same administrative entity

Examples : BELNET, UUNET, SKYNET, ...  
about 20000 ASes in 2007

Autonomous Systems are interconnected to allow the transmission of IP packets from any source to any destination

On the Internet, most packets need to travel through several transit Autonomous Systems

# Organisation of the Internet

---

Internet is composed of about 30.000  
autonomous routing domains

A domain is a set of routers, links, hosts and local  
area networks under the same administrative control

A domain can be very large...

AS568: SUMNET-AS DISO-UNRRA contains 73154560 IP addresses

A domain can be very small...

AS2111: IST-ATRIUM TE Experiment a single PC running Linux...

Domains are interconnected in various ways


The interconnection of all domains should in theory allow  
packets to be sent anywhere

Usually a packet will need to cross a few ASes to reach its  
destination

# Types of domains

## Transit domain

A **transit domain** allows external domains to use its own infrastructure to send packets to other domains



## Examples

UUNet, OpenTransit, GEANT, Internet2, RENATER, EQUANT, BT, Telia, Level3,...



# Types of domains (2)


## Stub domain

A stub domain does not allow external domains to use its infrastructure to send packets to other domains

A stub is connected to at least one transit domain

Single-homed stub : connected to one transit domain

Dual-homed stub : connected to two transit domains



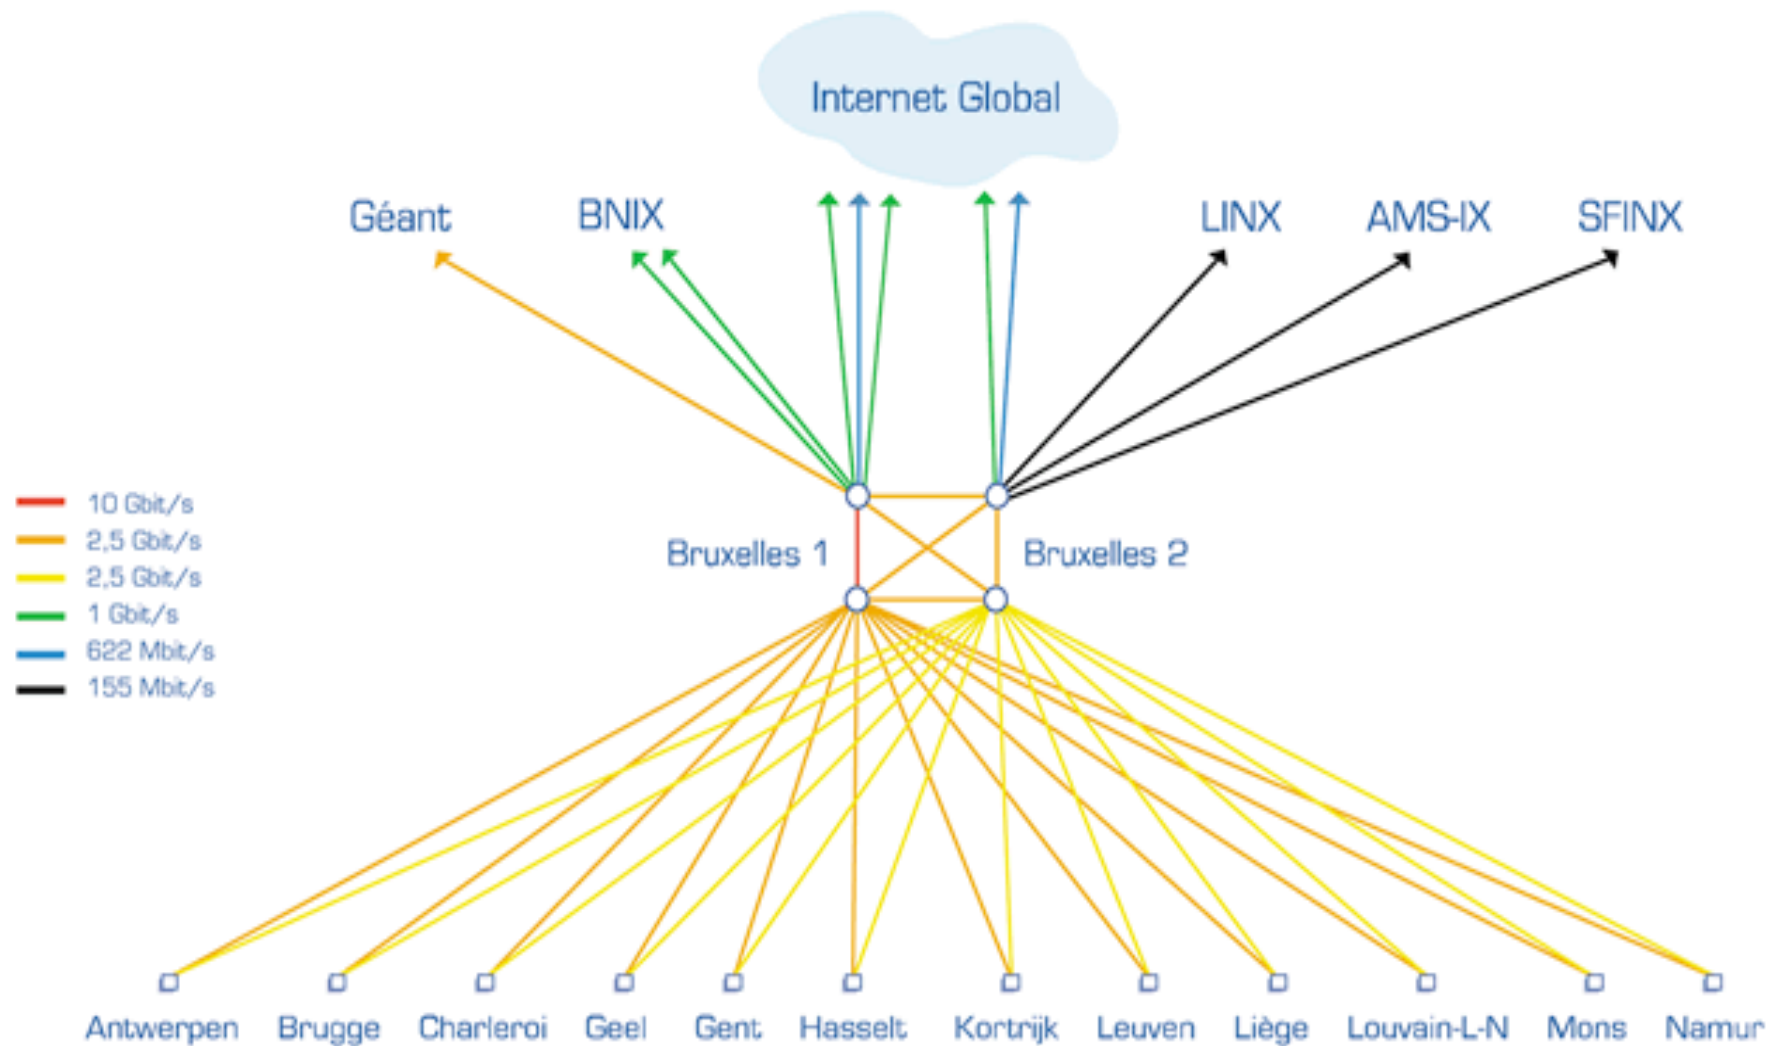
## Content-rich stub domain

Large web servers : Yahoo, Google, MSN, TF1, BBC,...

## Access-rich stub domain

ISPs providing Internet access via CATV, ADSL, ...

# Sample network : Belnet

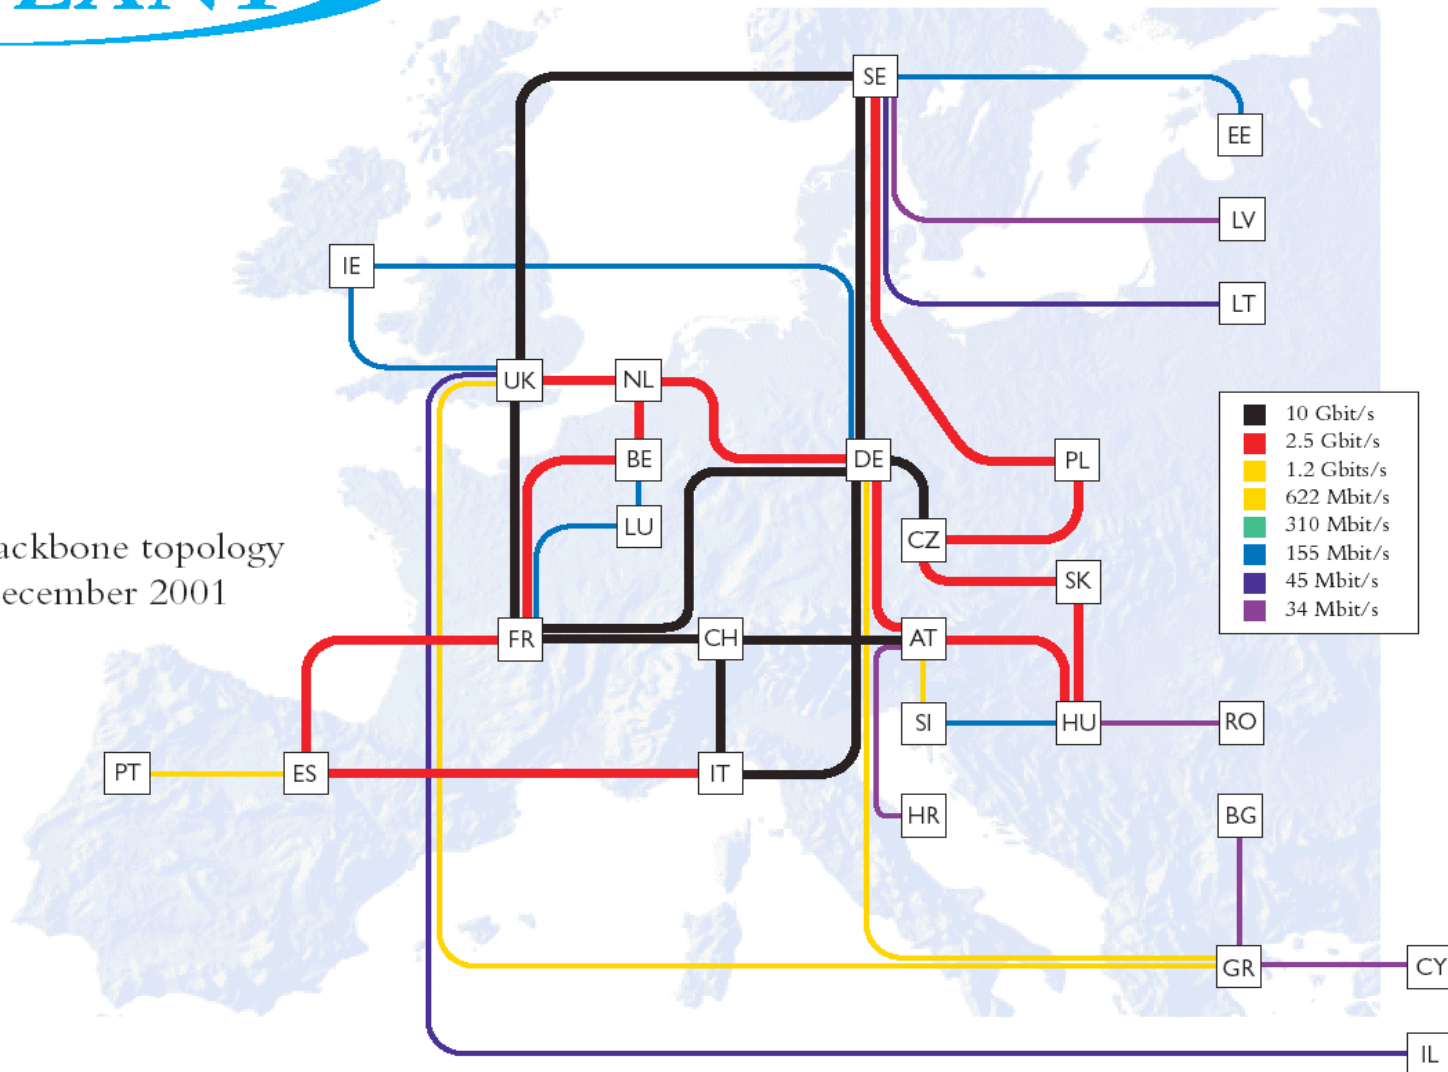


# Sample network : GEANT

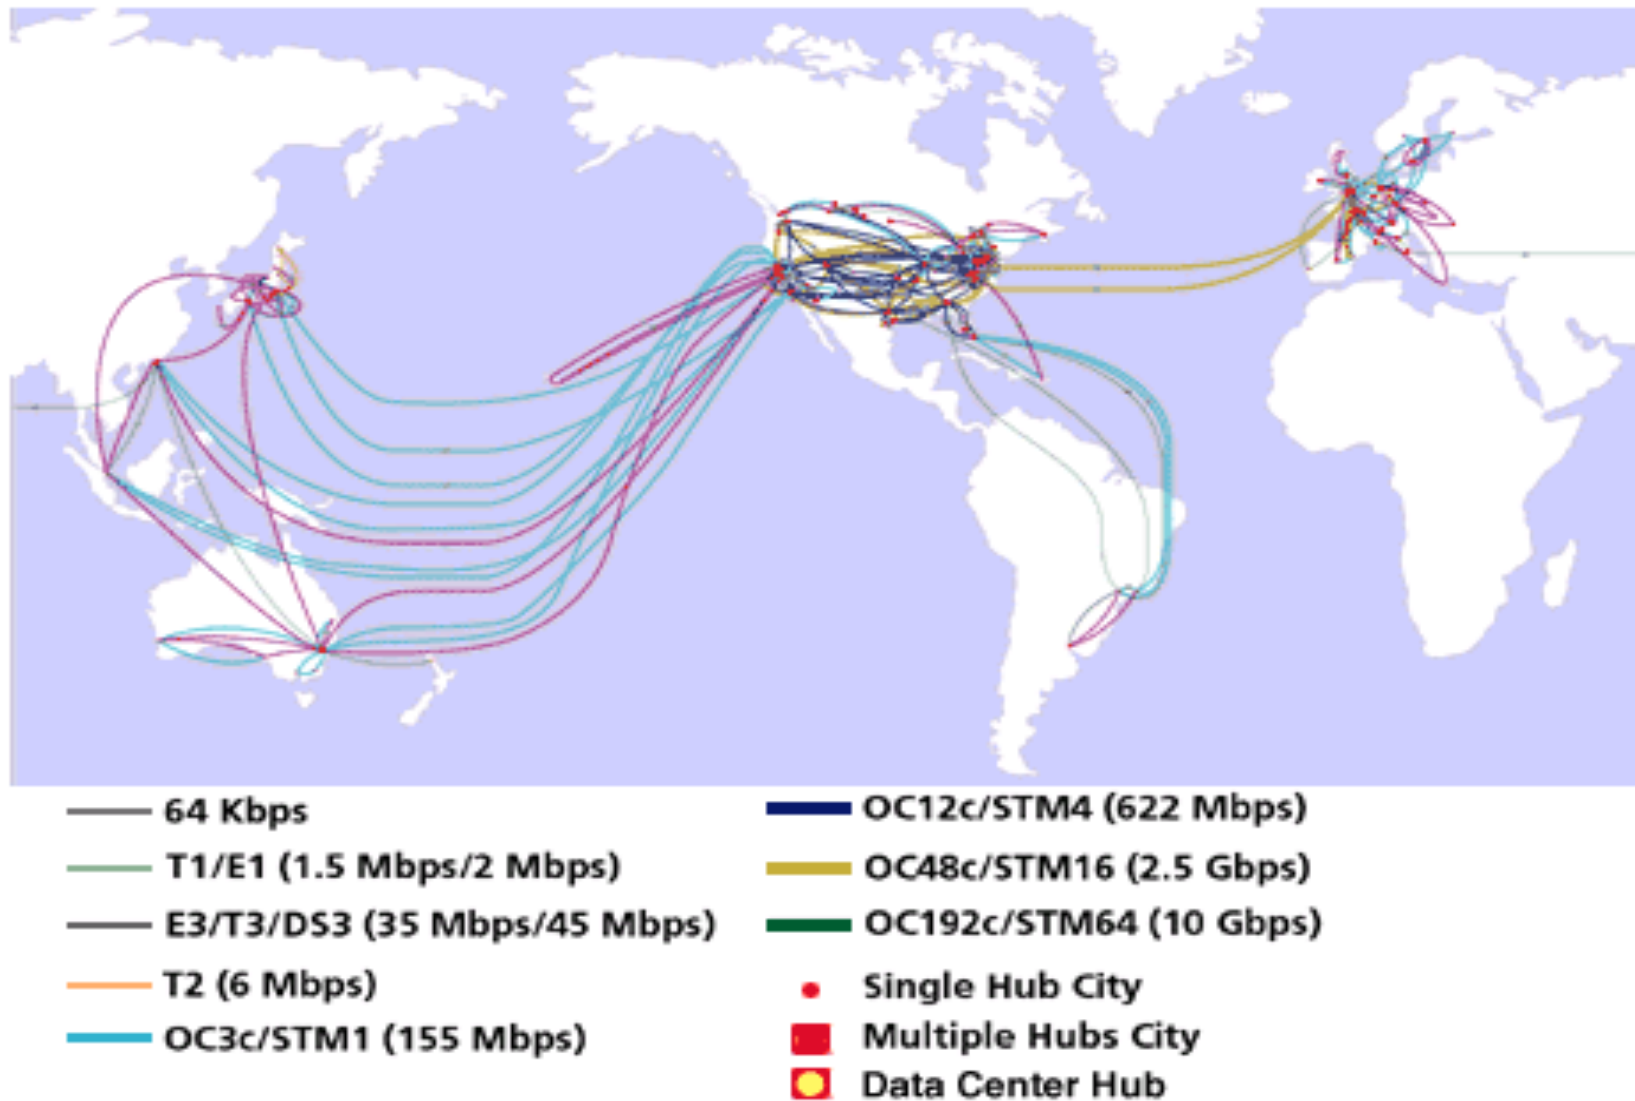


The Gigabit Research Network

Backbone topology  
December 2001



# A large worldwide network : UUNet

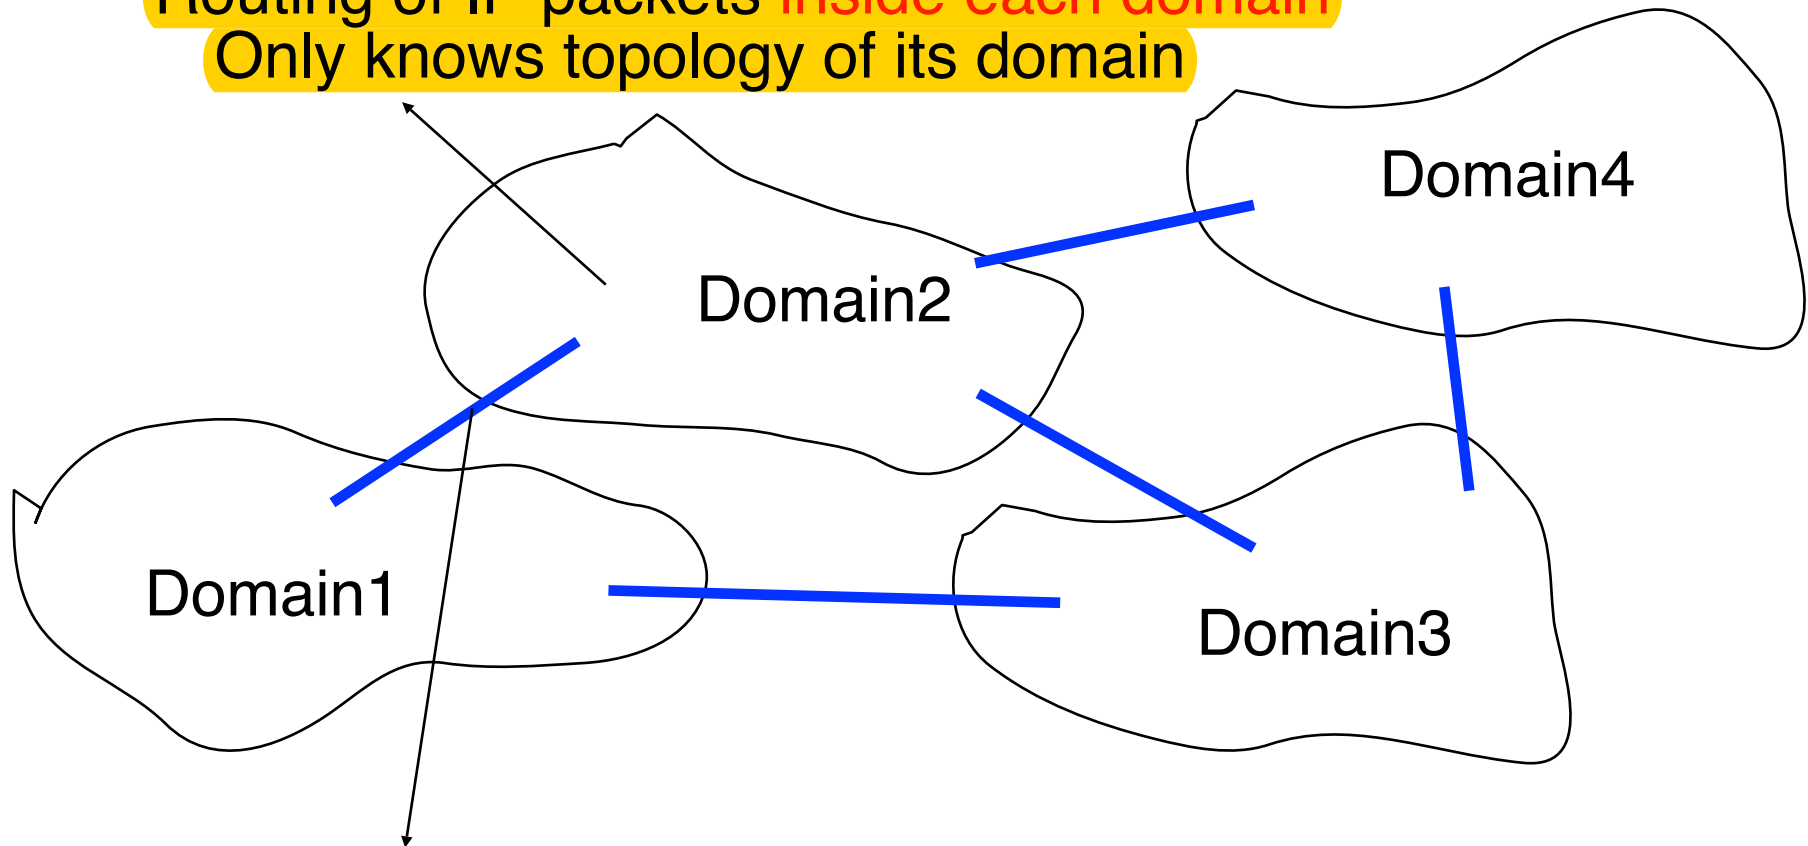


# Internet routing

## Interior Gateway Protocol (IGP)

Routing of IP packets **inside each domain**

Only knows topology of its domain



## Exterior Gateway Protocol (EGP)

Routing of IP packets **between domains**

Each domain is considered as a blackbox

# Intradomain routing

---

## Goal

- Allow routers to transmit IP packets along the best path towards their destination

  - best** usually means the shortest path

    - Shortest measured in seconds or as number of hops

  - sometimes **best** means the less loaded path

- Allow to find alternate routes in case of failures

## Behaviour

- All routers exchange routing information

  - Each domain router can obtain routing information for the whole domain

  - The network operator or the routing protocol selects the cost of each link

# Three types of Interior Gateway Protocols

---

## Static routing

Only useful in very small domains

## Distance vector routing

Routing Information Protocol (RIP)

Still widely used in small domains despite its limitations

## Link-state routing

Open Shortest Path First (OSPF)

Widely used in enterprise networks

## Intermediate System- Intermediate-System (IS-IS)

Widely used by ISPs

# Network layer

---

Basics

Routing

IP : Internet Protocol

Routing in IP networks

Internet routing organisation

→ Intradomain routing : RIP

Intradomain routing : OSPF

Interdomain routing : BGP



# RIP

## Routing Information Protocol

---

Simple routing protocol that relies on distance vectors

Defined in RFC2453

### Principle

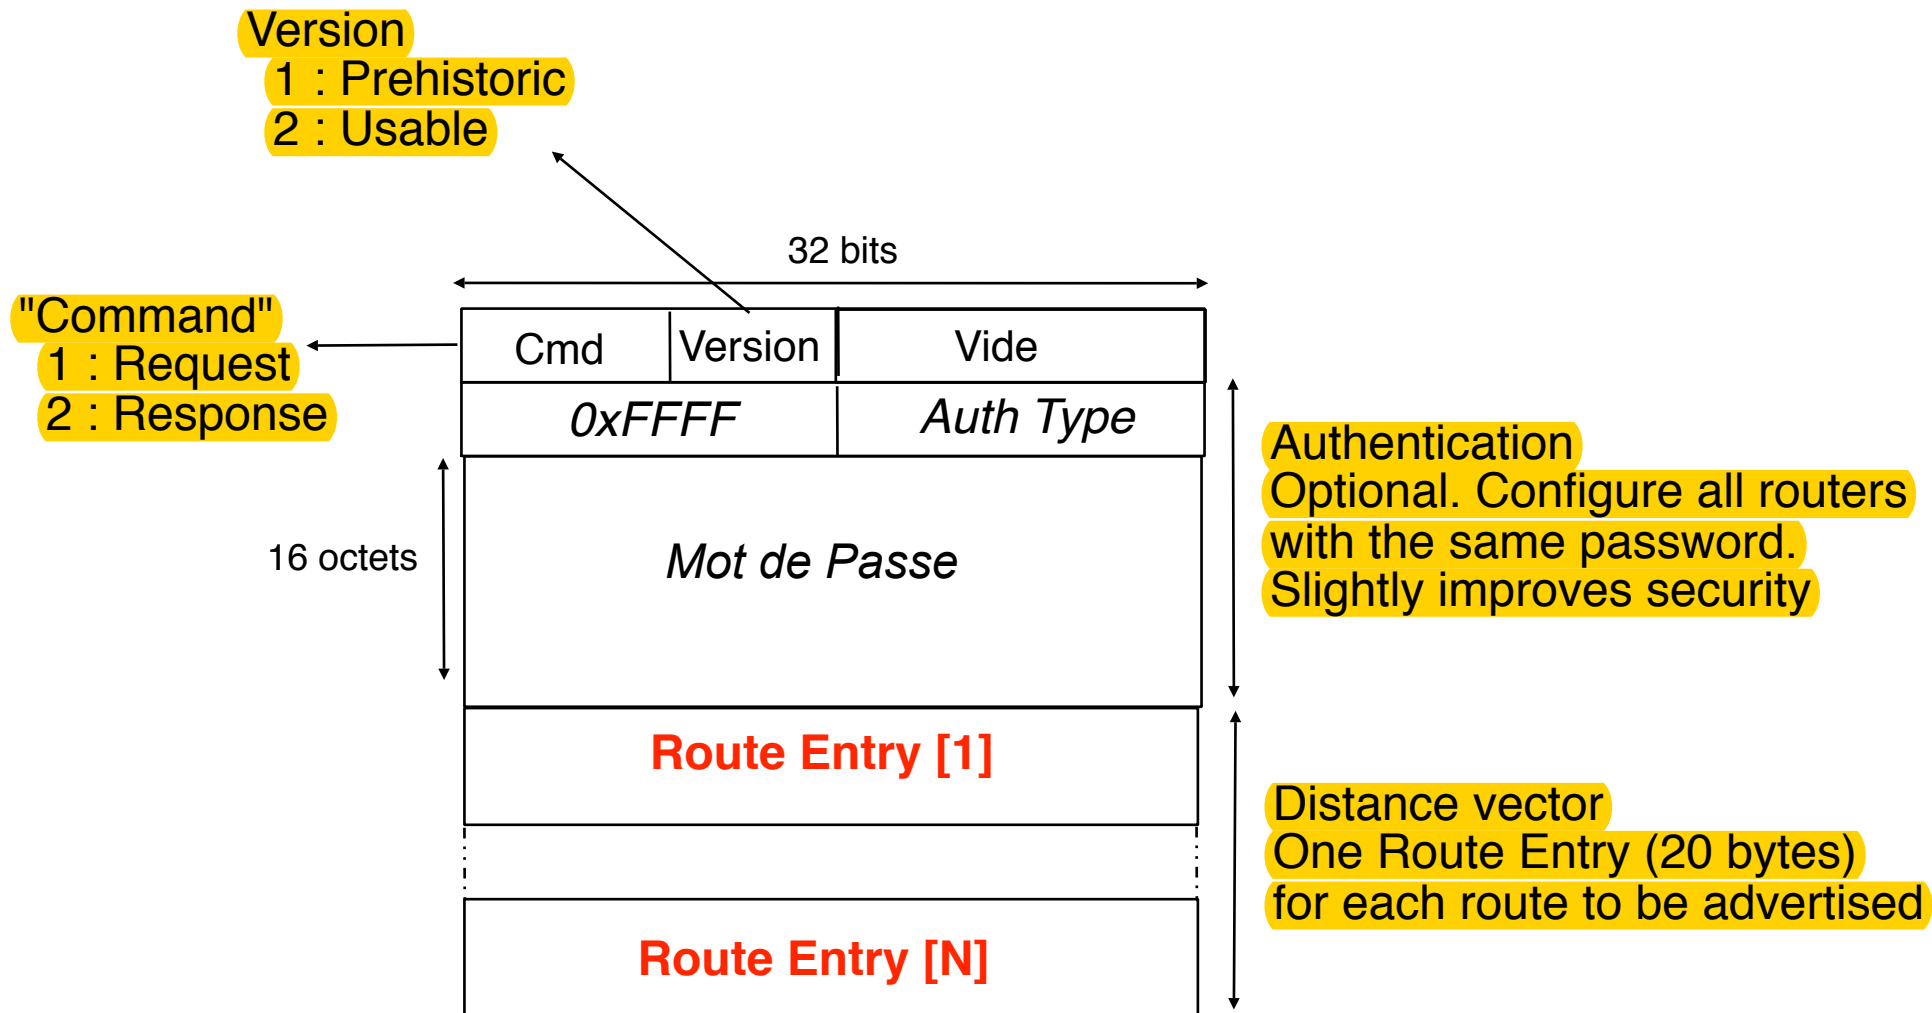
- Each router periodically sends its distance vectors
  - default period : 30 seconds
  - distance vector is sent in UDP message with TTL=1 to all routers in local subnets (via IP multicast )

- Optional extension : send a distance vector when the routing table changes

  - simple solution : send distance vector after each change
  - but some links flaps...

  - solution : send a distance vector if routing table changed and we did not send another vector within the last 5 seconds

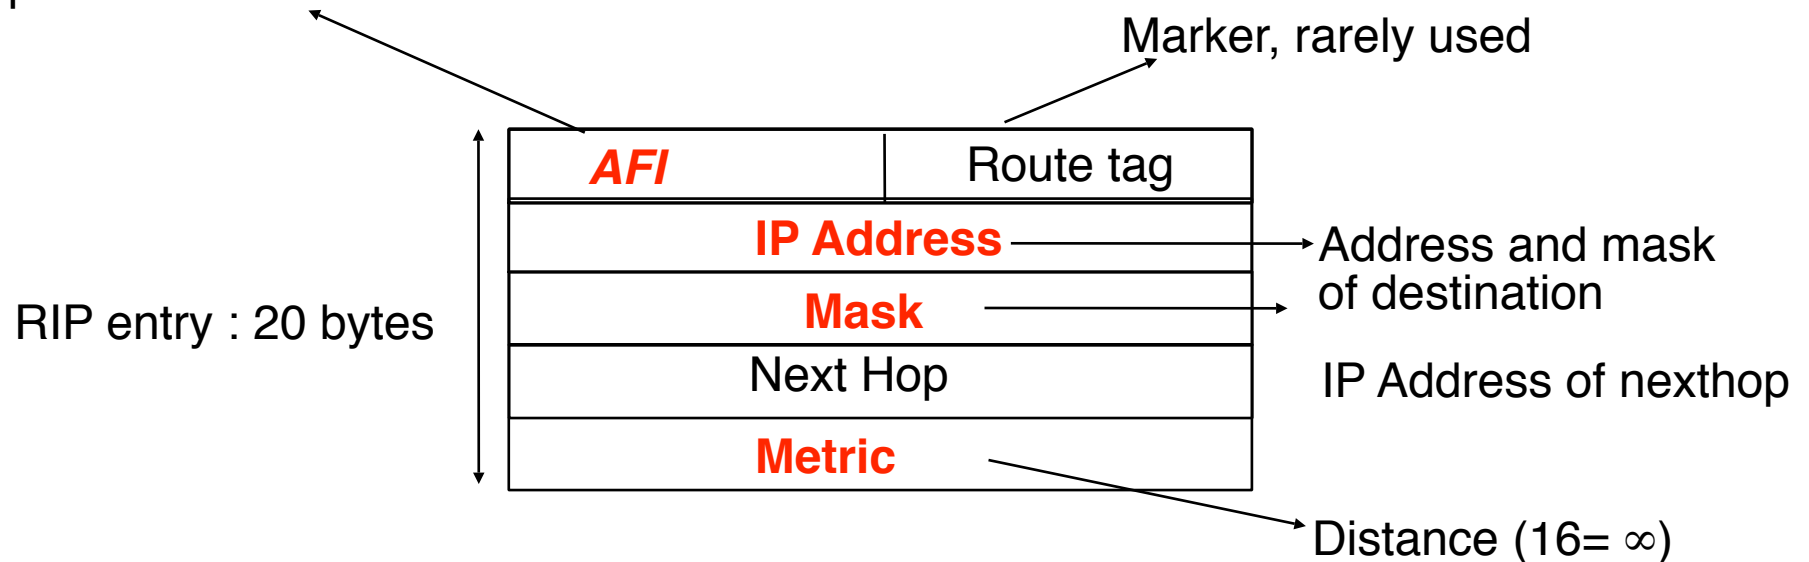
# RIP : message format



**RIP messages are sent by UDP  
port 520**

# RIP : Route Entries

AFI : Address Family Identifier  
type of addresses used  
2= Ipv4



## Default route

IP Address = 0.0.0.0, Mask = 0

Each RIP message can contain up to 25 route entries  
(24 with authentication)

If the routing table is larger than 25 entries, router will  
need to send several RIP messages

# RIP timers

---

## Operation

At each expiration of its 30-sec timer, each router sends its distance vector and restarts its timer

## Problem

After a power failure, all routers might restart at same time and have synchronised RIP timers

Each router will need to process bursts of RIP messages

## Solution

Add some randomness to the timers

Restart timer after  $\text{random}[27.5, 32.5]$  instead of 30 seconds

commonly used technique to avoid synchronisation problems in distributed protocols

# Network layer

---

Basics

Routing

IP : Internet Protocol

Routing in IP networks

Internet routing organisation

Intradomain routing : RIP

→ Intradomain routing : OSPF

Interdomain routing : BGP

# OSPF

---

## Standardised link state routing protocol

### Operation

#### Router startup

- HELLO packets to discover neighbours

#### Update of routing tables

##### Link state packets

- acknowledgements, sequence numbers, age

- periodic transmission

- transmission upon link changes

#### Database description

- provides the list of sequence numbers of all LSPs

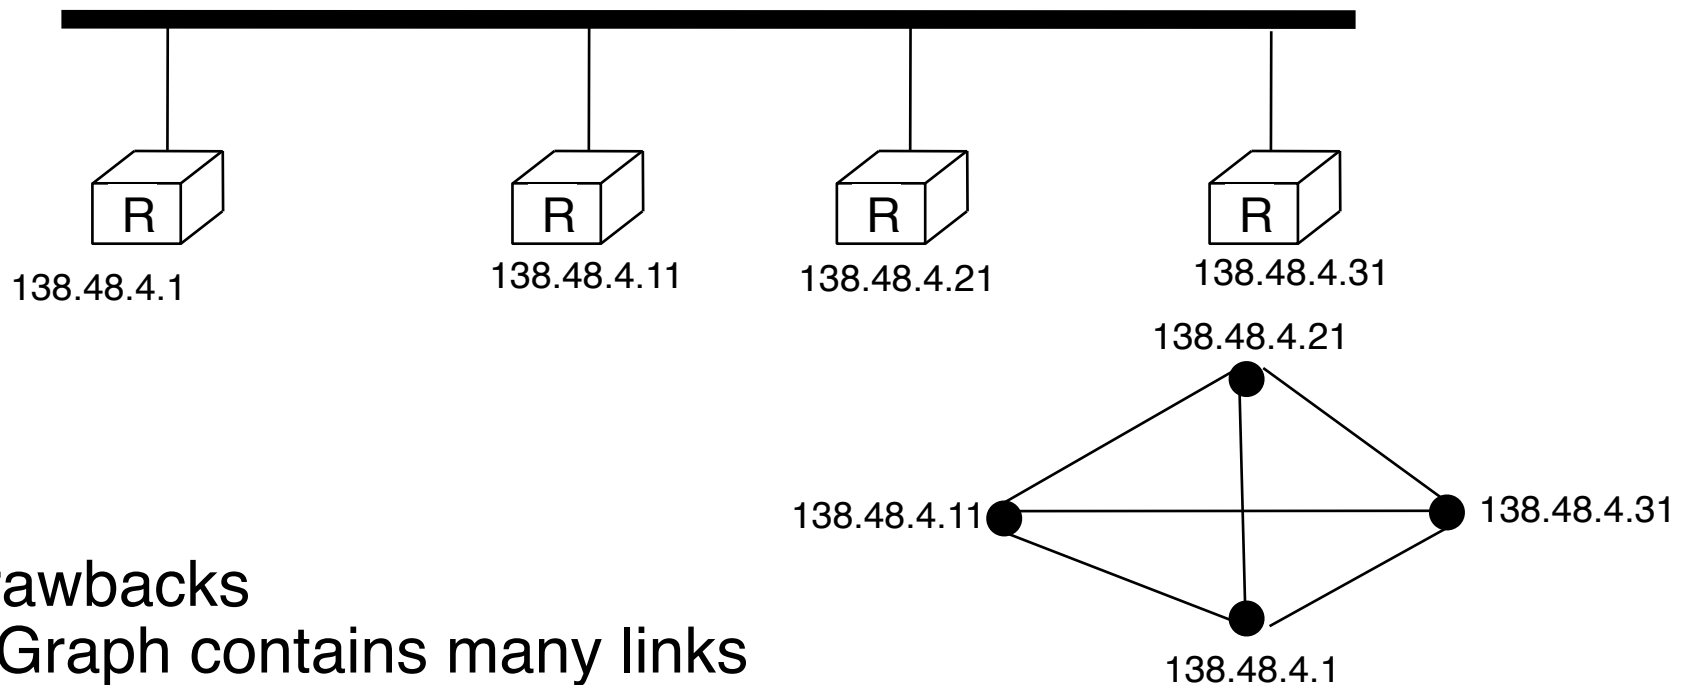
- stored by router

#### Link state Request

- used when a router boots to request link state packets from neighbours

# OSPF details

Routers are often attached to LANs  
How to describe a LAN full of routers as a graph



## Drawbacks

Graph contains many links

Routers need to exchange lots of HELLOs

Does not really describe the LAN

a failure of the LAN would cause a disconnection of all routers while the graph indicates a redundant topology

# OSPF details (2)

## Solution


represent the LAN as a star with one router acting as the LAN

Designated router

One router is elected in the LAN to originate link state packets for the LAN

Adjacent router

Maintain adjacencies with the designated router





# OSPF details (3)

---

# OSPF details (3)

---

OSPF in large networks  
avoid too large routing tables in OSPF routers

# OSPF details (3)

---

## OSPF in large networks

avoid too large routing tables in OSPF routers

## Solution

Divide network in **areas**

Backbone area : network backbone

all routers connected to two or more areas belong to the backbone area

All non-backbone areas must be attached to the backbone area

at least one router inside each area must be attached to the backbone

# OSPF details (3)

---

## OSPF in large networks

avoid too large routing tables in OSPF routers

## Solution

Divide network in **areas**

Backbone area : network backbone

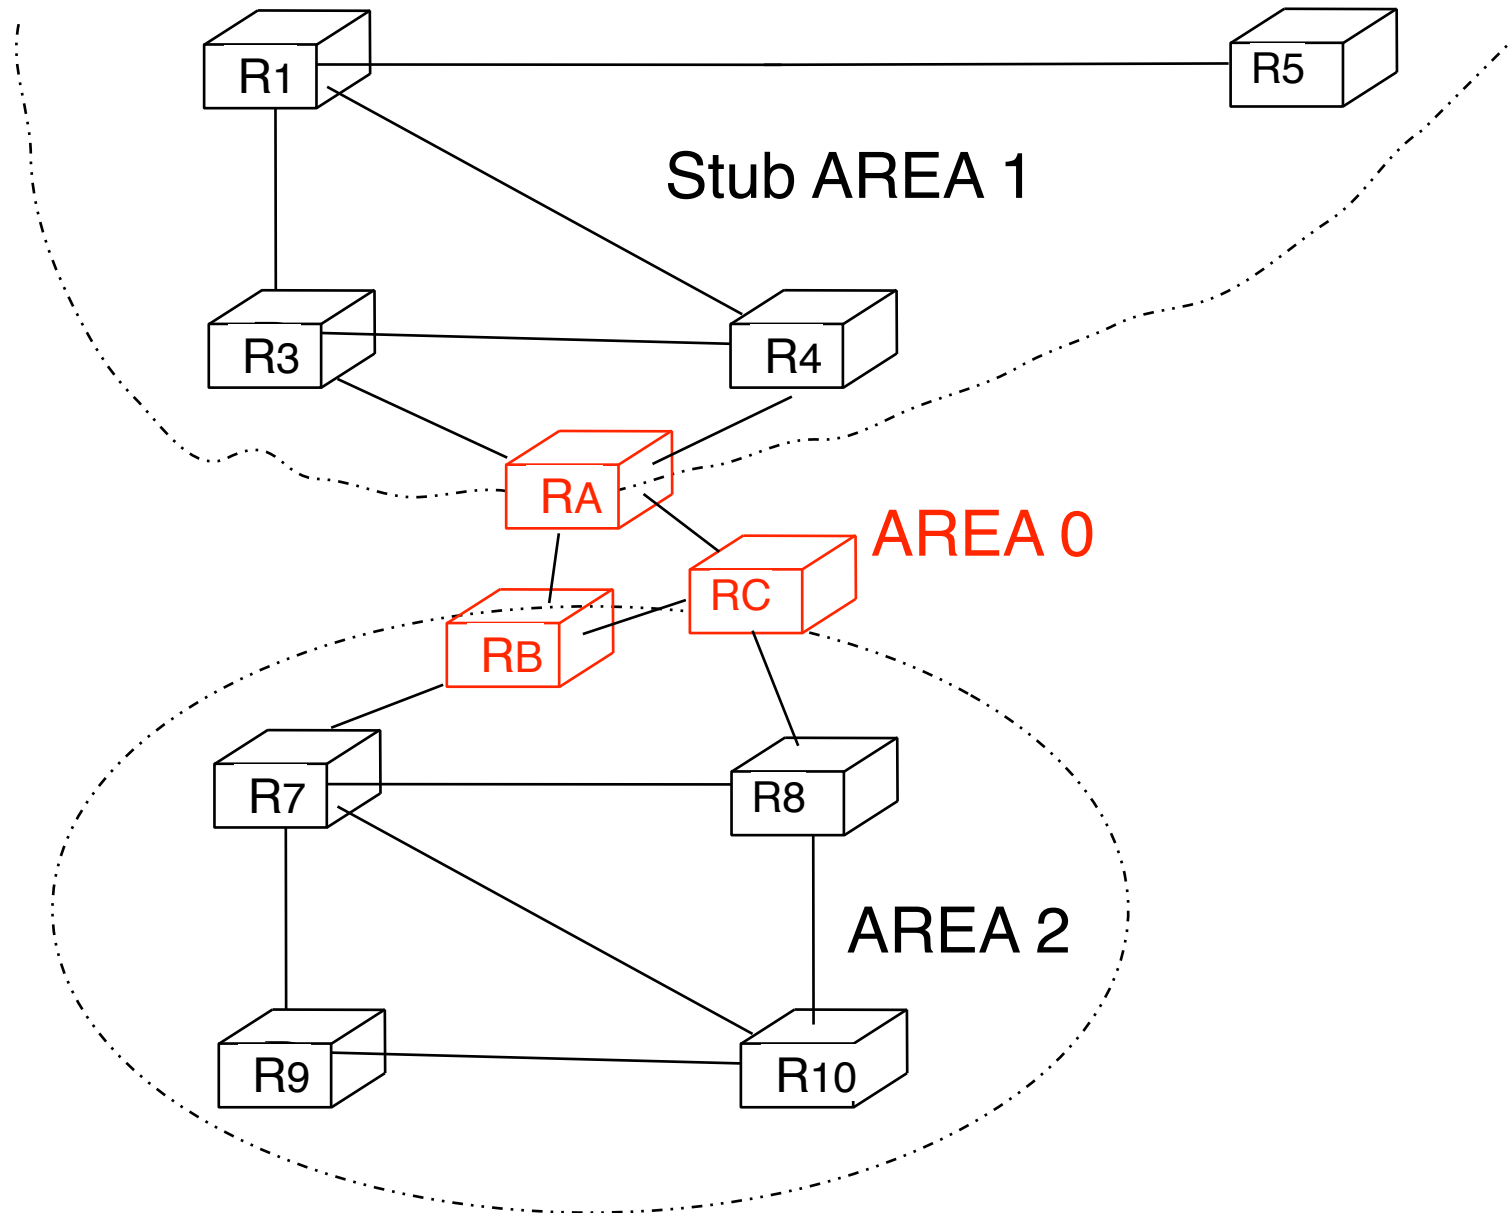
all routers connected to two or more areas belong to the backbone area

All non-backbone areas must be attached to the backbone area

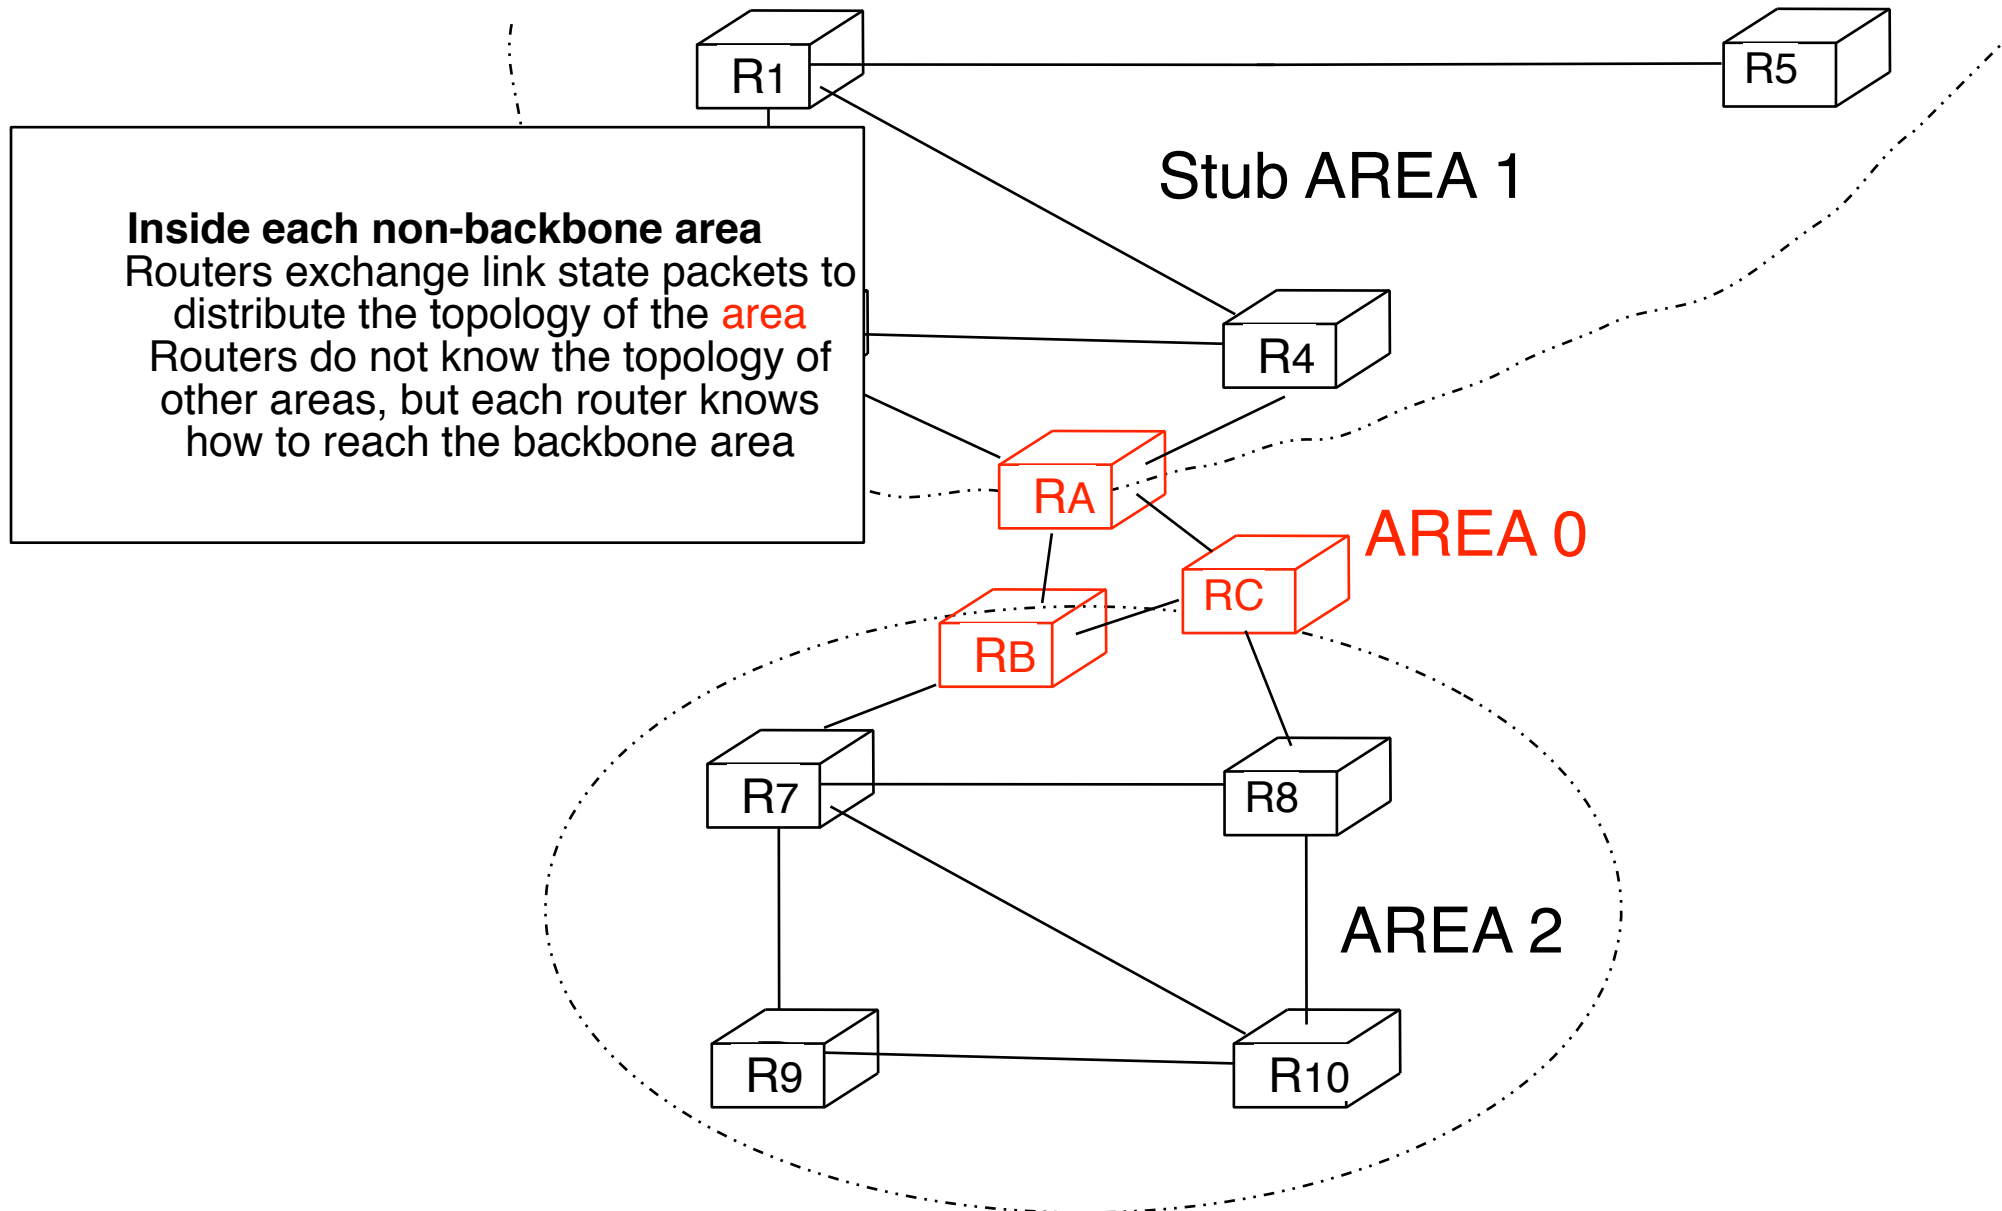
at least one router inside each area must be attached to the backbone

**OSPF routing must allow any router to send packets to any other router**

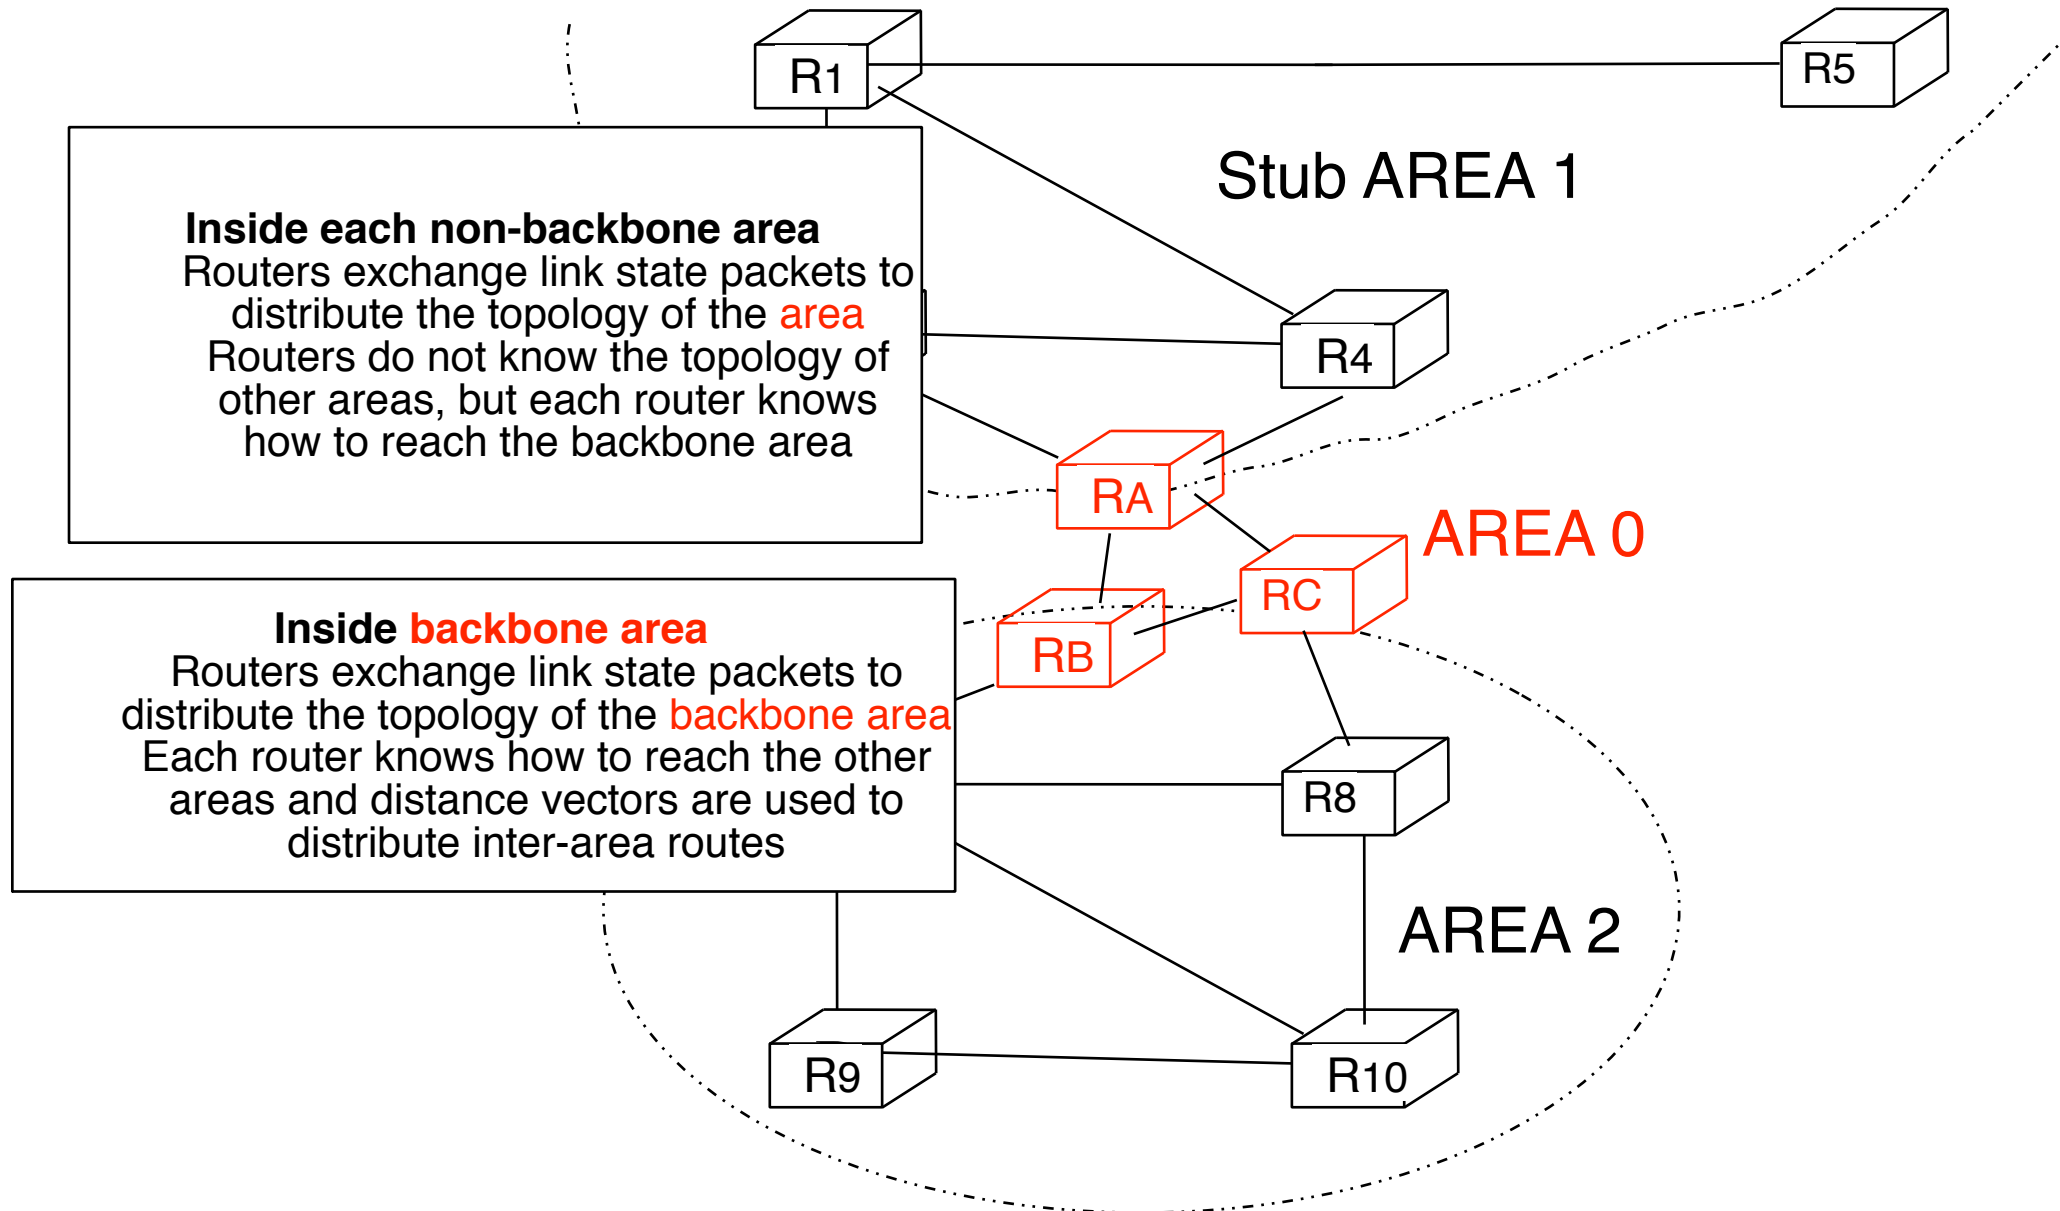
# OSPF details (4)



# OSPF details (4)



# OSPF details (4)



# OSPF areas : Example

Routes learned by R4

192.168.1.0/24, distance 3 (via RA)  
192.168.10.0/24, distance 3 (via RA)

Routes chosen par RA

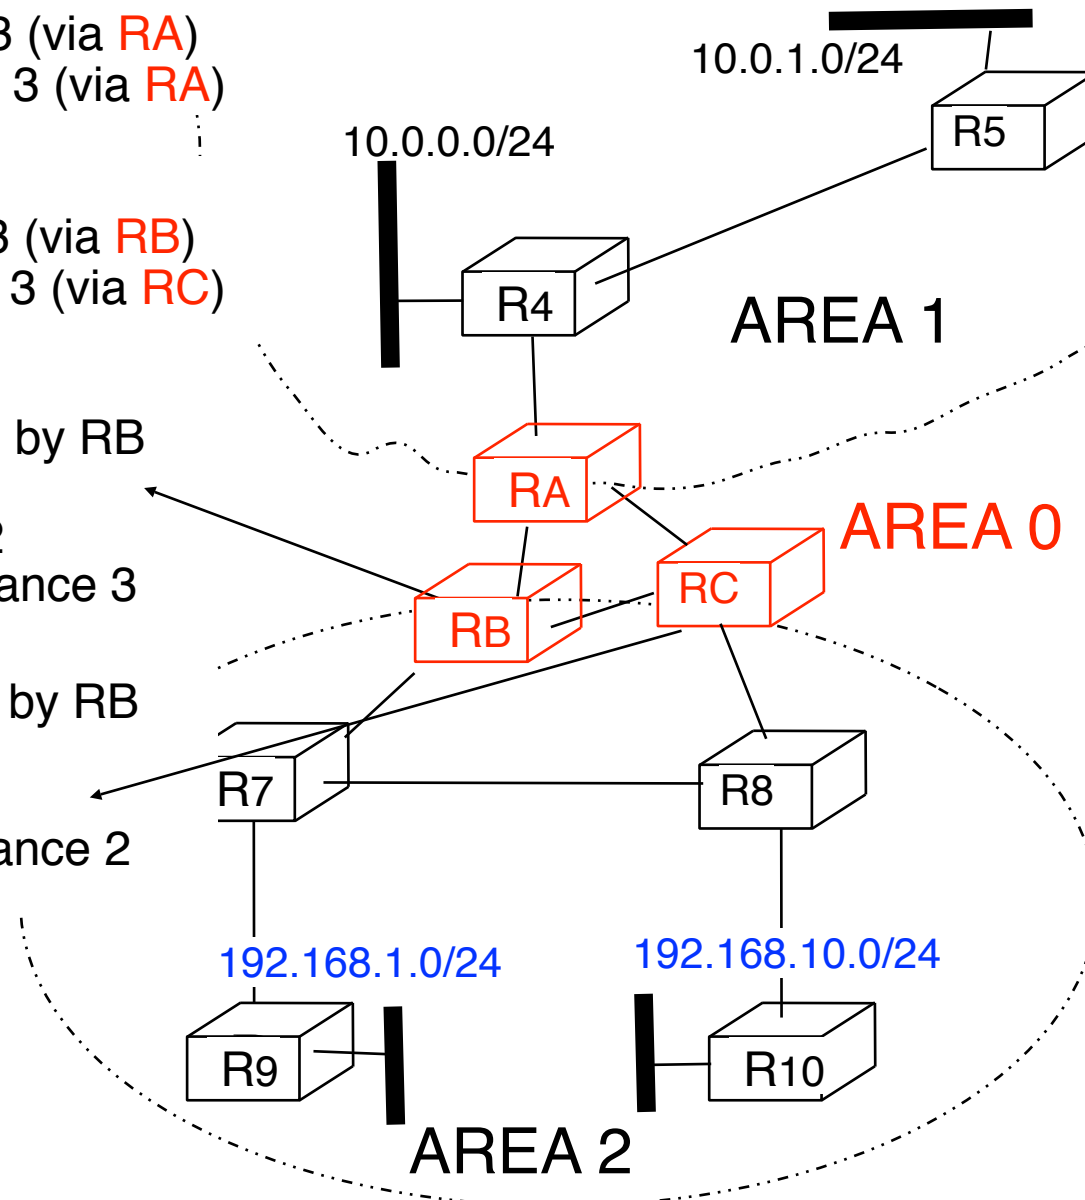
192.168.1.0/24, distance 3 (via RB)  
192.168.10.0/24, distance 3 (via RC)

Distance vectors advertised by RB  
in backbone area

192.168.1.0/24, distance 2  
and 192.168.10.0/24, distance 3

Distance vectors advertised by RB  
in backbone area

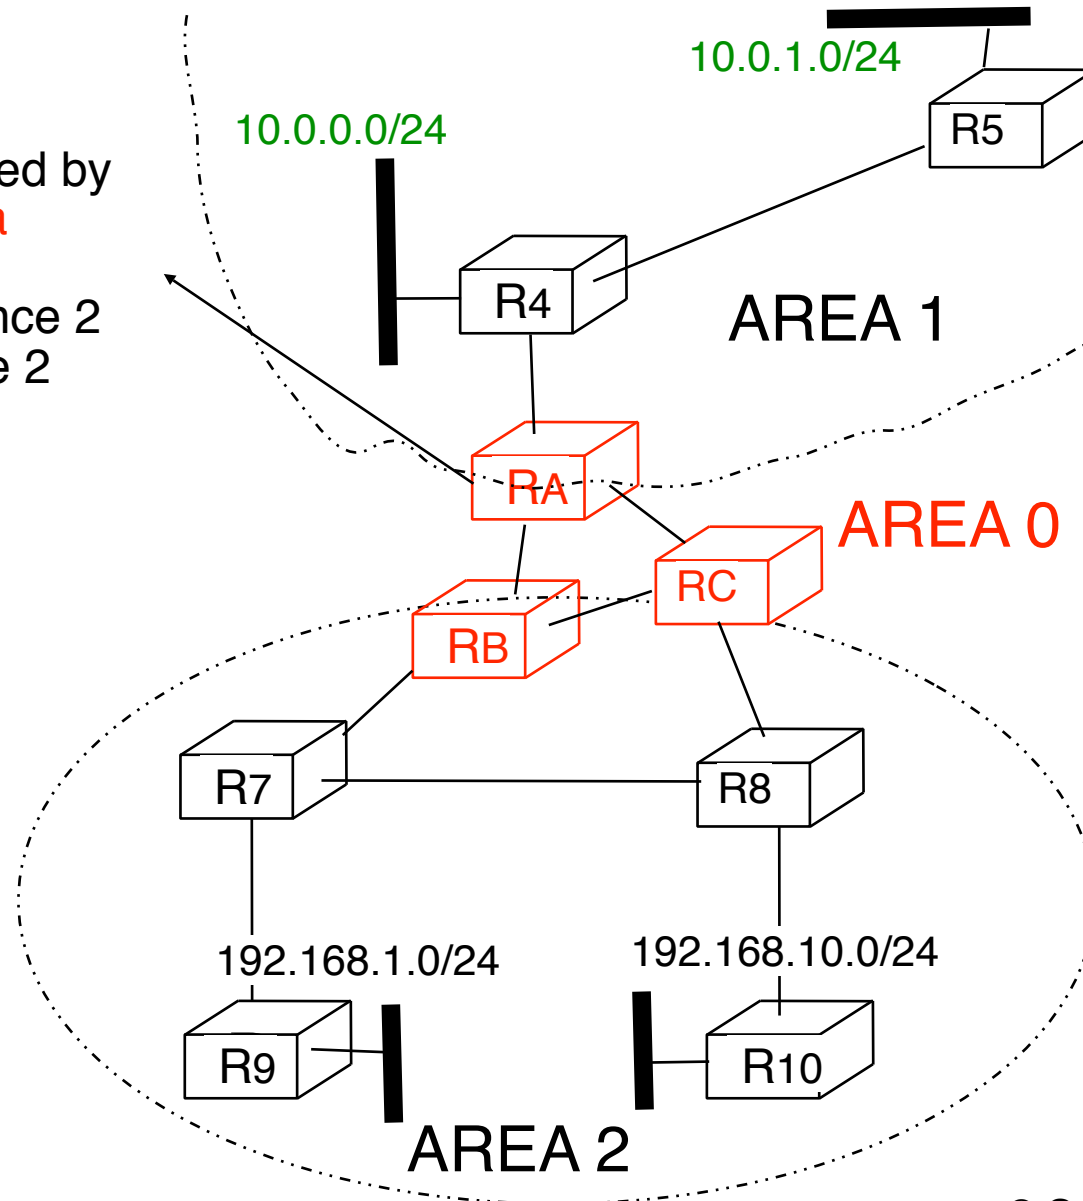
192.168.1.0/24, distance 3  
and 192.168.10.0/24, distance 2





# Areas OSPF : Example (2)

Distance vector advertised by  
**RA** in the **backbone area**  
**10.0.0.0/24**, distance 1  
and **10.0.1.0/24**, distance 2  
or **10.0.0.0/23**, distance 2



# Network layer

---

Basics

Routing

IP : Internet Protocol

Routing in IP networks

Internet routing organisation

Intradomain routing : RIP

Intradomain routing : OSPF

→ Interdomain routing : BGP

# Interdomain routing

---

## Goals

Allow to transmit IP packets along the **best path** towards their destination through several transit domains while taking into account the **routing policies** of each domain without knowing the detailed topology of those domains

From an interdomain viewpoint, **best path** often means *cheapest path*


**Each domain** is free to specify inside its **routing policy** the domains for which it agrees to provide a transit service and the method it uses to select the best path to reach each destination

# Types of interdomain links

## Two types of interdomain links


### Private link

Usually a leased line between two routers belonging to the two connected domains



### Connection via a public interconnection point

Usually Gigabit or higher Ethernet switch that interconnects routers belonging to different domains



# Routing policies

---

In theory BGP allows each domain to define its own routing policy...

In practice there are two common policies

**customer-provider peering**


**Customer c** buys Internet connectivity from **provider P**

**shared-cost peering**

**Domains x** and **y** agree to exchange packets by using a direct link or through an interconnection point

# Customer-provider peering

---



## Principle


Customer sends to its provider its internal routes and the routes learned from its own customers

Provider will advertise those routes to the entire Internet to allow anyone to reach the Customer

Provider sends to its customers all known routes

Customer will be able to reach anyone on the Internet

# Shared-cost peering



## Principle

PeerX sends to PeerY its internal routes and the routes learned from its own customers

PeerY will use shared link to reach PeerX and PeerX's customers

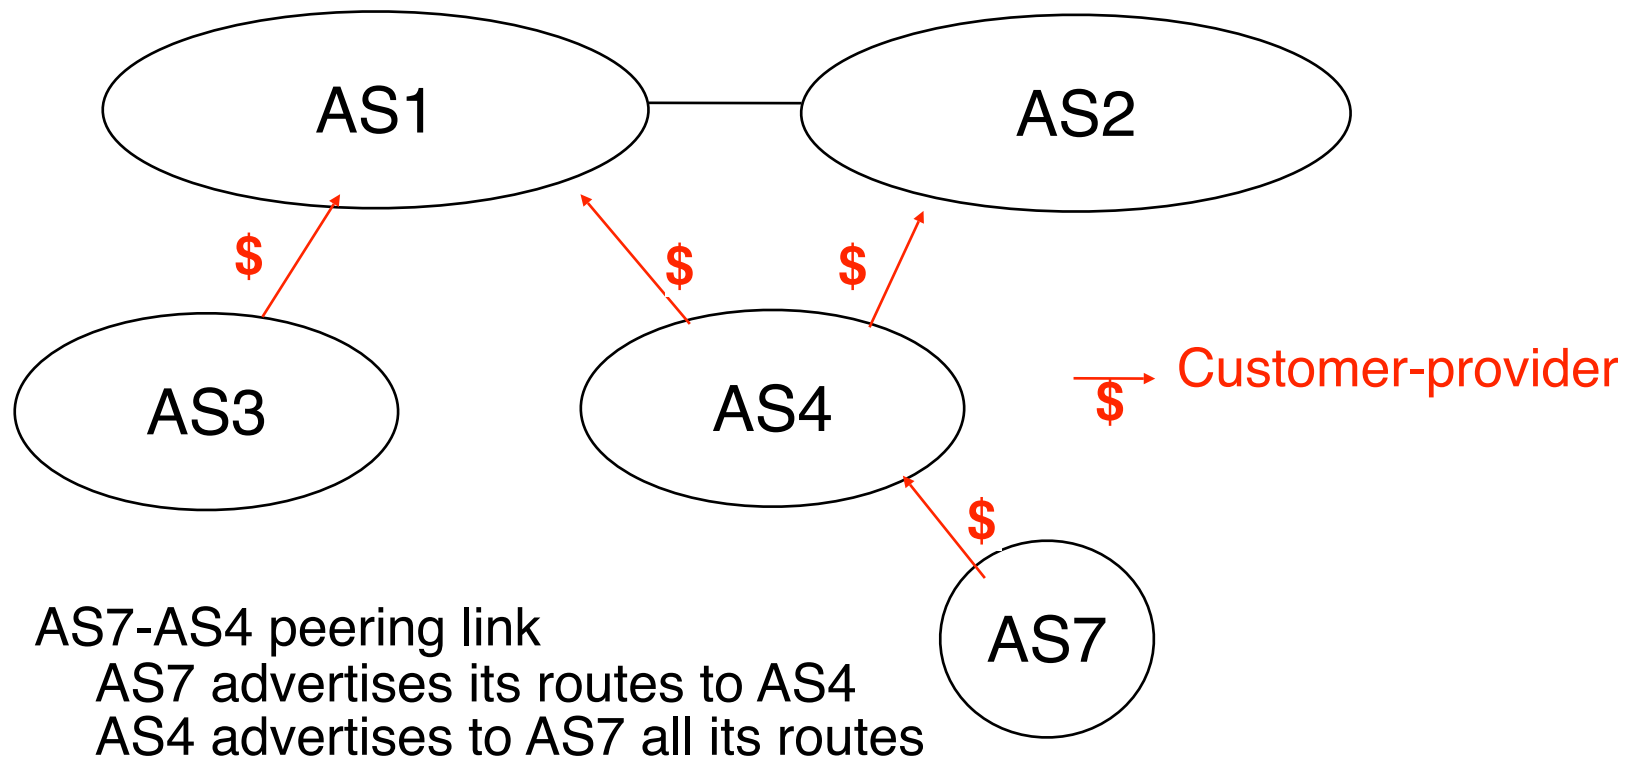
PeerX's providers are not reachable via the shared link

PeerY sends to PeerX its internal routes and the routes learned from its own customers

PeerX will use shared link to reach PeerY and PeerY's customers

PeerY's providers are not reachable via the shared link


# Customer-provider peering : example



AS4-AS2 peering link  
AS4 advertises its own routes et those of its customers (AS7)  
AS2 advertises to AS2 all known routes



# Shared-cost peering : example



## AS3-AS4 peering link

AS3 advertises its own routes

AS4 advertises its own routes and those received from its clients (AS7)

## AS1-AS2 peering link

AS1 advertises its own routes and those received from its clients (AS3 and AS4)

AS1 advertises its own routes and those received from its clients (AS4)

# Routing policies

---

A domain specifies its routing policy by defining on each BGP router two sets of filters for each peer

## Import filter

Specifies which routes can be accepted by the router among all the received routes from a given peer

## Export filter

Specifies which routes can be advertised by the router to a given peer

Filters can be defined in RPSL

Routing Policy Specification Language

defined in RFC2622 and examples in RFC2650

See also <http://www.ripe.net/ripencc/pub-services/whois.html>

# RPSL

---

## Simple import policies

### Syntax

```
import: from AS# accept list_of_AS
```

### Examples

```
Import: from Belgacom accept Belgacom WIN
```

```
Import: from Provider accept ANY
```

## Simple export policies

### Syntax

```
Export: to AS# announce list_of_AS
```

### Example


```
Export: to Customer announce ANY
```

```
Export: to Peer announce Customer1 Customer2
```

# Routing policies

## Simple example with RPSL

---



### Import policy for AS4

Import: from AS3 accept AS3  
import: from AS7 accept AS7  
import: from AS1 accept ANY  
import: from AS2 accept ANY

### Export policy for AS4

export: to AS3 announce AS4 AS7  
export: to AS7 announce ANY  
export: to AS1 announce AS4 AS7  
export: to AS2 announce AS4 AS7

### Import policy for AS7


Import: from AS4 accept ANY

### Export policy for AS4

export: to AS4 announce AS7


# The organisation of the Internet

---



# The organisation of the Internet

---



## Tier-1 ISPs

Dozen of large ISPs  
interconnected by **shared-cost**

Provide transit service

Uunet, Level3, OpenTransit, ...

## Tier-2 ISPs

Regional or National ISPs

Customer of T1 ISP(s)

Provider of T3 ISP(s)

**shared-cost** with other T2 ISPs

France Telecom, BT, Belgacom

## Tier-3 ISPs

Smaller ISPs, Corporate

Networks, Content providers

Customers of T2 or T1 ISPs


**shared-cost** with other T3 ISPs

# The Border Gateway Protocol

## Principle

### Path vector protocol

BGP router advertises its best route to each destination




# The Border Gateway Protocol

## Principle

### Path vector protocol

BGP router advertises its best route to each destination



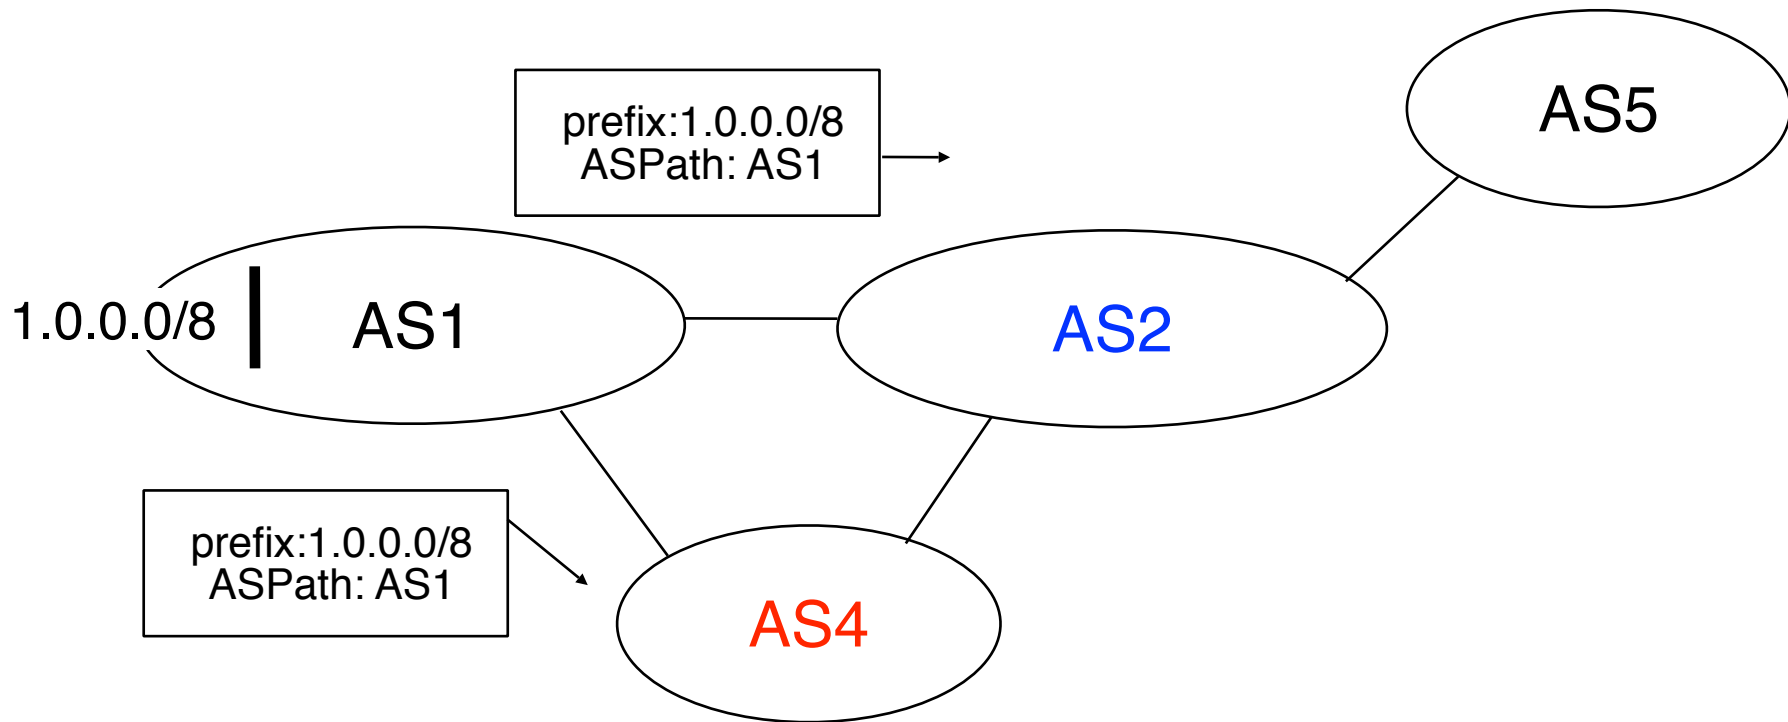


# The Border Gateway Protocol

## Principle

### Path vector protocol

BGP router advertises its best route to each destination

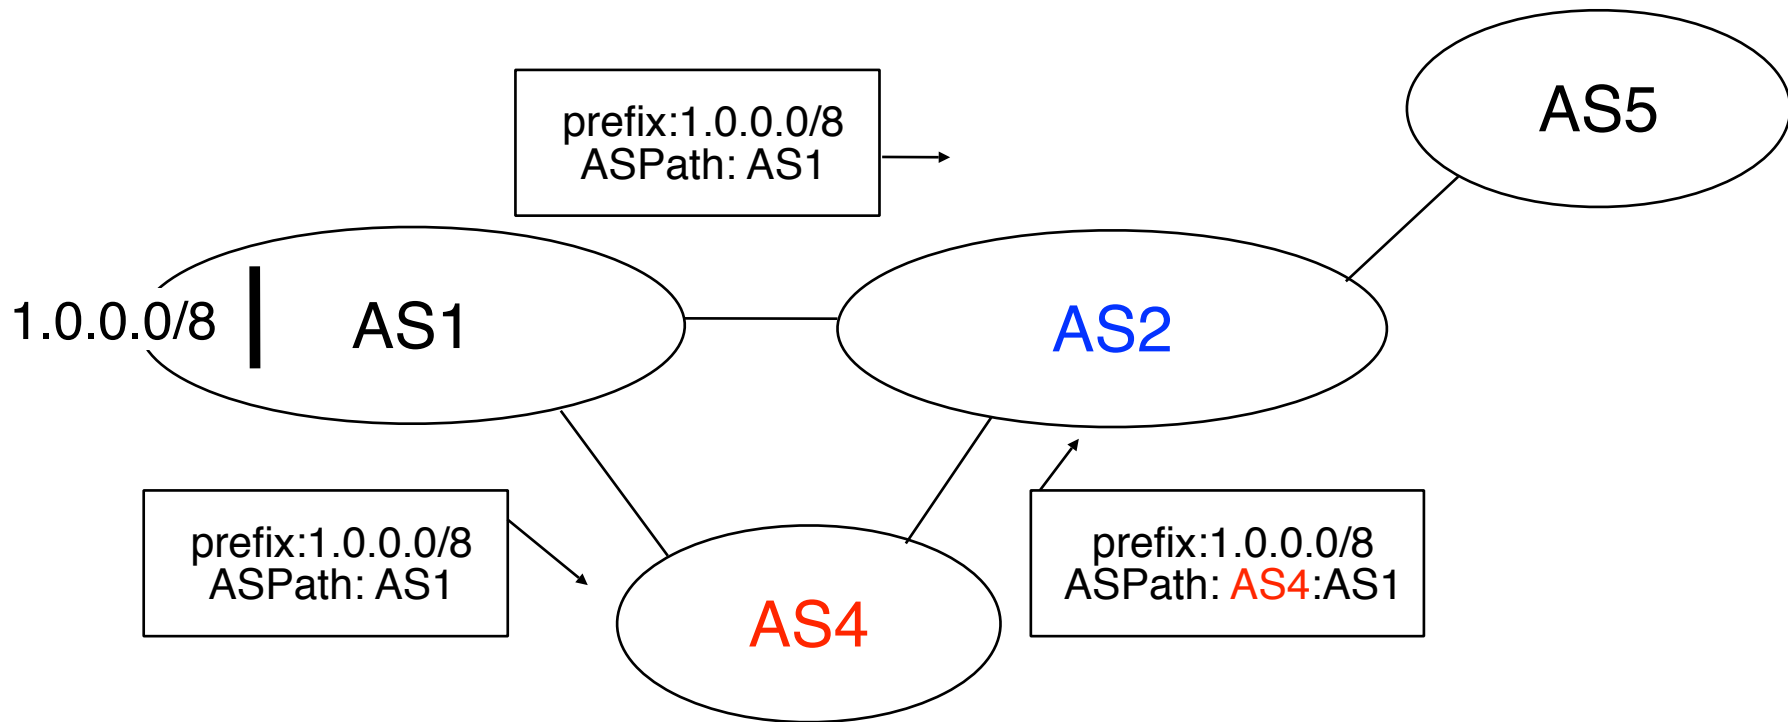


# The Border Gateway Protocol

## Principle

### Path vector protocol

BGP router advertises its best route to each destination

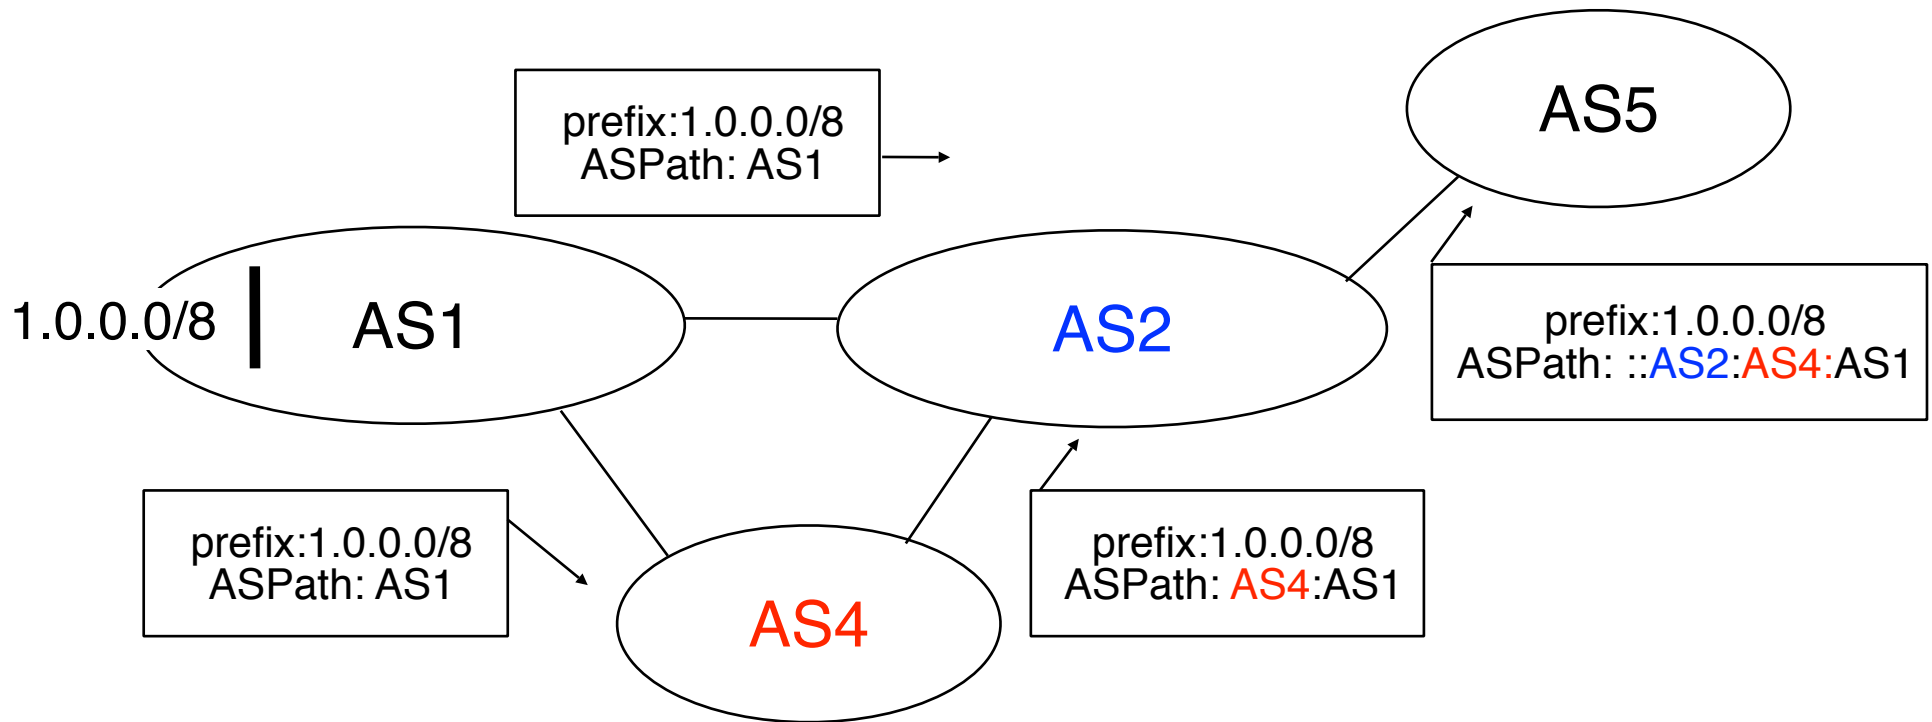


# The Border Gateway Protocol

## Principle

### Path vector protocol

BGP router advertises its best route to each destination




# The Border Gateway Protocol

## Principle

### Path vector protocol

BGP router advertises its best route to each destination



... with incremental updates

Advertisements are only sent when their content changes


# BGP : Principles of operation

---

## Principles

BGP relies on the

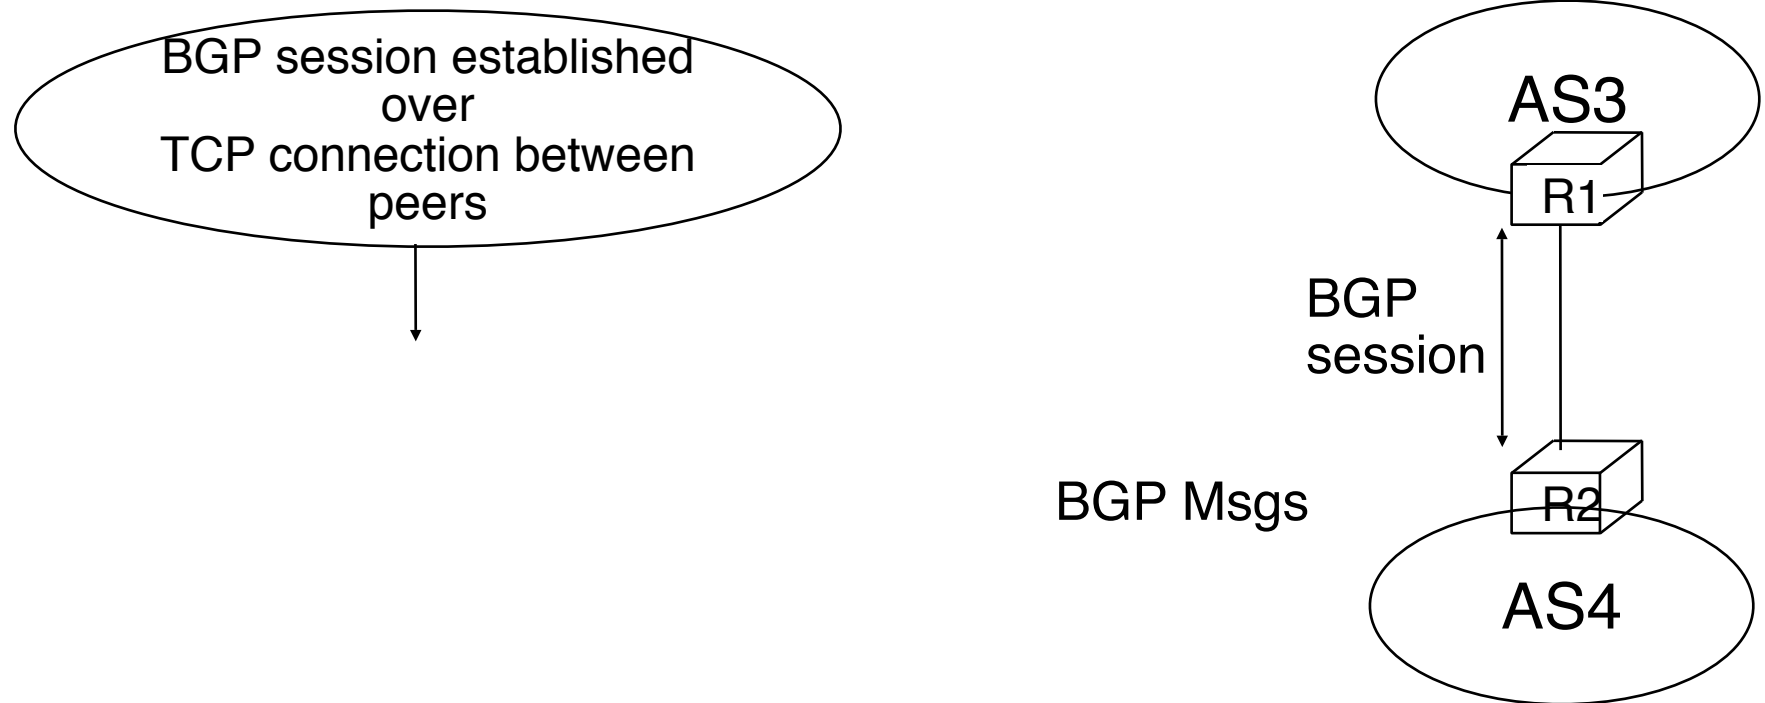
incremental exchange of path vectors



# BGP : Principles of operation

## Principles



BGP relies on the  
incremental exchange of path vectors



# BGP : Principles of operation

## Principles

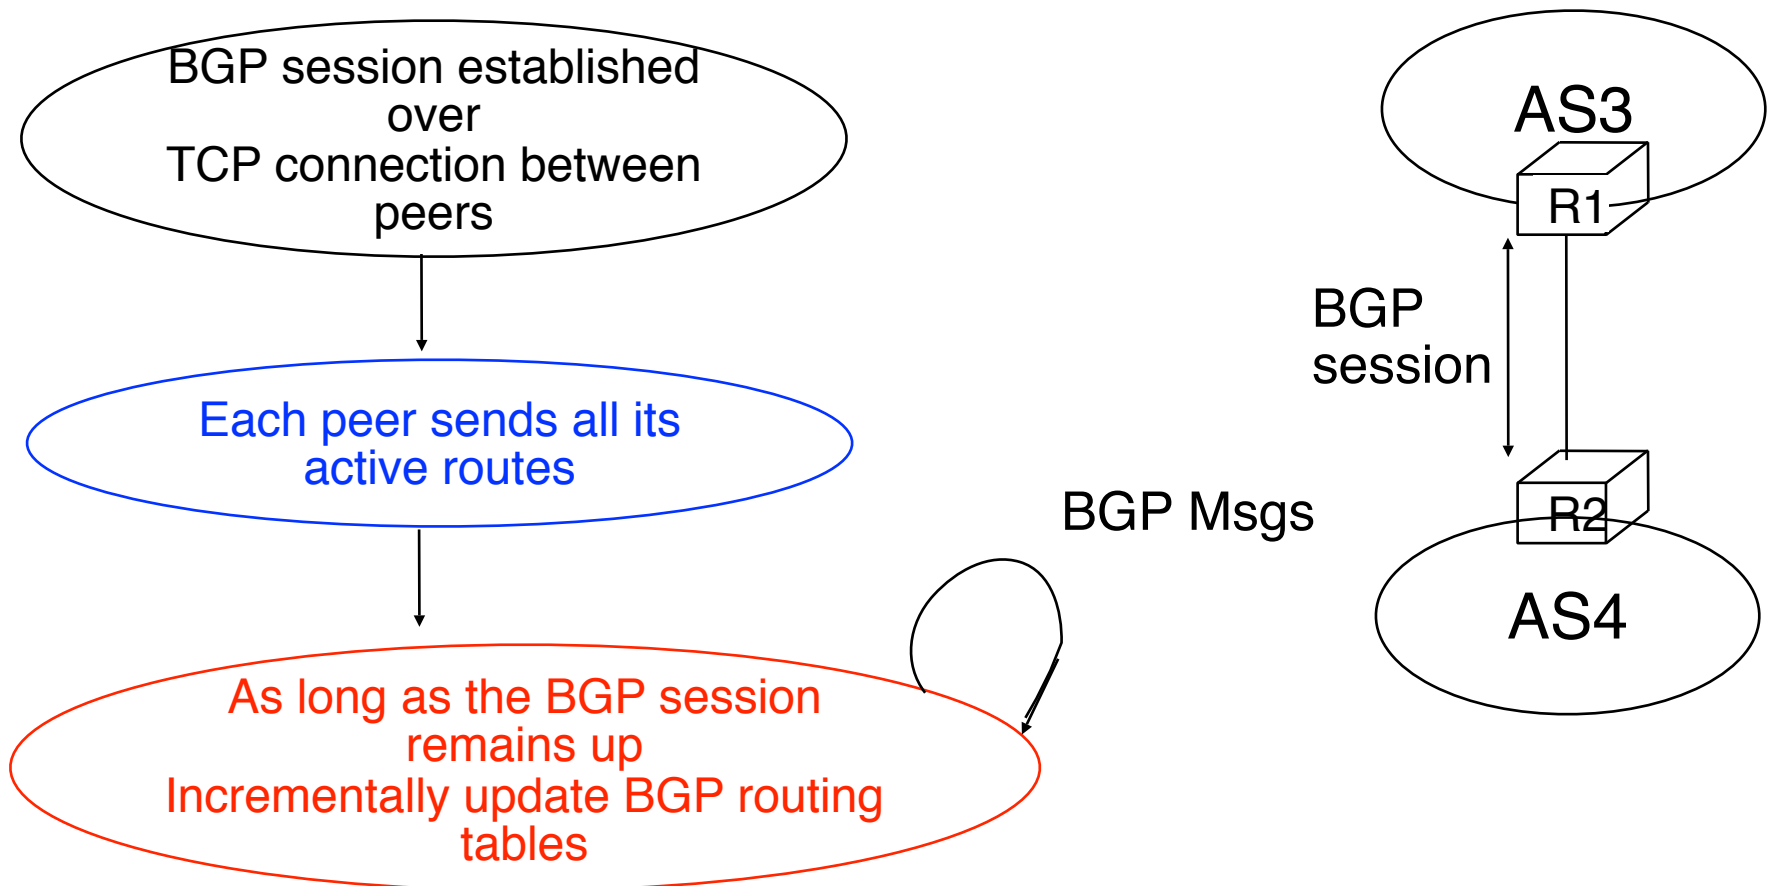
BGP relies on the  
incremental exchange of path vectors



# BGP : Principles of operation

## Principles

BGP relies on the  
**incremental exchange of path vectors**





# BGP : Principles of operation (2)

---

## Simplified model of BGP

2 types of BGP path vectors

### UPDATE

Used to announce a route towards one prefix

#### Content of UPDATE

- Destination address/prefix

- Interdomain path used to reach destination (AS-Path)

- Nexthop (address of the router advertising the route)

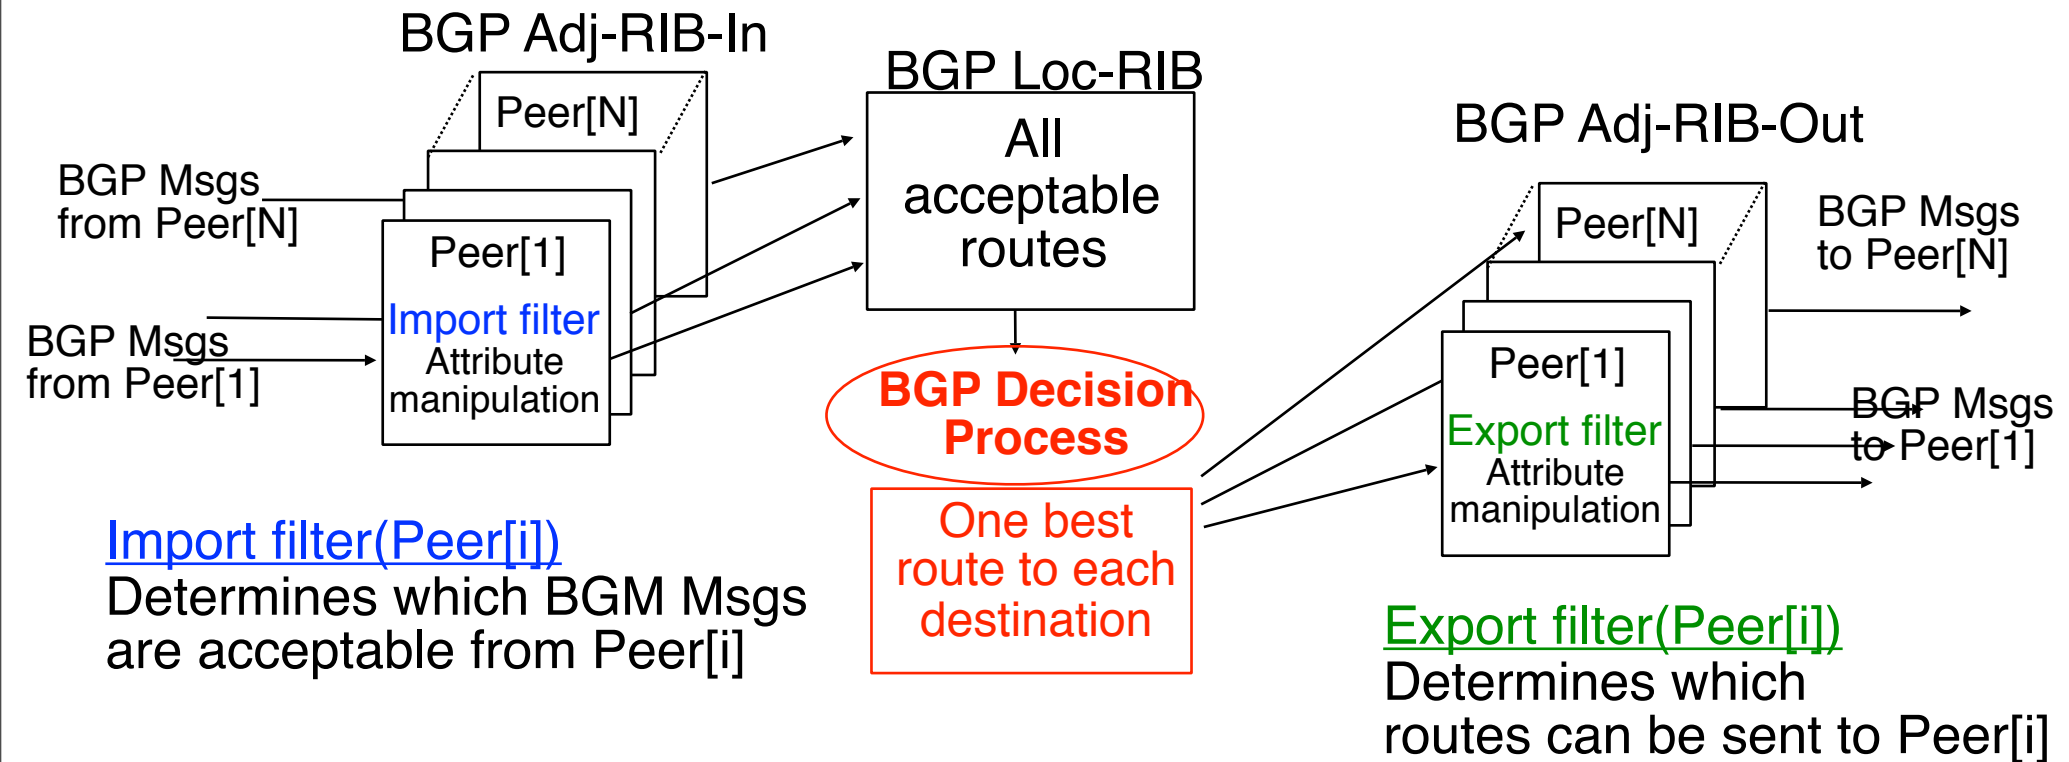
### WITHDRAW

Used to indicate that a previously announced route is not reachable anymore

#### Content of WITHDRAW

- Unreachable destination address/prefix

# Conceptual model of a BGP router



## BGP Routing Information Base

Contains all the acceptable routes learned from all Peers + internal routes

**BGP decision process** selects **the best route** towards each destination

# Where do the routes advertised by BGP routers come from ?

---

## Learned from another BGP router

Each BGP router advertises best route towards each destination

## Static route

Configured manually on the router

Ex : The BGP router at UCL advertises 130.104.0.0/16

Drawback

Requires manual configuration

Advantage

BGP advertisements are stable

## Learned from an intradomain routing protocol

BGP might try to aggregate the route before advertising it

Advantage :

BGP advertisements correspond to network status

Drawback

Routing instabilities inside a domain might propagate in

Internet

# BGP : Session Initialization

---

```
Initialize_BGP_Session(RemoteAS, RemoteIP)
{ /* Initialize and start BGP session */
/* Send BGP OPEN Message to RemoteIP on port 179*/
/* Follow BGP state machine */

/* advertise local routes and routes learned from peers*/
foreach (destination=d inside RIB)
{
    B=build_BGP_UPDATE(d);
    S=apply_export_filter(RemoteAS,B);
    if (S<>NULL)
        { /* send UPDATE message */
            send_UPDATE(S,RemoteAS, RemoteIP)
        }
}
/* entire RIB was sent */
/* new UPDATE will be sent only to reflect local or distant
   changes in routes */
...
}
```

# Events during a BGP session

---

## 1. Addition of a new route to RIB

- A new internal route was added on local router
  - static route added by configuration
  - Dynamic route learned from IGP

- Reception of UPDATE message announcing a new or modified route

## 2. Removal of a route from RIB

- Removal of an internal route
  - Static route is removed from router configuration
  - Intradomain route declared unreachable by IGP

- Reception of WITHDRAW message

## 3. Loss of BGP session

- All routes learned from this peer removed from RIB

# Export and Import filters

---

```
BGPMsg Apply_export_filter(RemoteAS, BGPMsg)
{ /* check if Remote AS already received route */
if (RemoteAS isin BGPMsg.ASPath)
    BGPMsg=NULL;
/* Many additional export policies can be configured : */
/* Accept or refuse the BGPMsg */
/* Modify selected attributes inside BGPMsg */
}
```

```
BGPMsg apply_import_filter(RemoteAS, BGPMsg)
{ /* check that we are not already inside ASPath */
if (MyAS isin BGPMsg.ASPath)
    BGPMsg=NULL;
/* Many additional import policies can be configured : */
/* Accept or refuse the BGPMsg */
/* Modify selected attributes inside BGPMsg */
}
```

# BGP : Processing of UPDATES

---

```
Recv_BGPMsg(Msg, RemoteAS)
{
    B=apply_import_filter(Msg,RemoteAS);
    if (B==NULL) /* Msg not acceptable */
        exit();
    if IsUPDATE(Msg)
    {
        Old_Route=BestRoute(Msg.prefix);
        Insert_in_RIB(Msg);
        Run_Decision_Process(RIB);
        if (BestRoute(Msg.prefix)<>Old_Route)
        { /* best route changed */
            B=build_BGP_Message(Msg.prefix);
            S=apply_export_filter(RemoteAS,B);
            if (S<>NULL) /* announce best route */
                send_UPDATE(S,RemoteAS);
            else if (Old_Route<>NULL)
                send_WITHDRAW(Msg.prefix);
        } ...
    }
```

# BGP : Processing of WITHDRAW

---


```
RecvMsg(Msg, RemoteAS)
...
if IsWithdraw(Msg)
{
    Old_Route=BestRoute(Msg.prefix);
    Remove_from_RIB(Msg);
    Run_Decision_Process(RIB);
    if (Best_Route(Msg.prefix) <> Old_Route)
    { /* best route changed */
        B=build_BGP_Message(d);
        S=apply_export_filter(RemoteAS,B);
        if (S<>NULL) /* still one best route */
            send_UPDATE(S,RemoteAS, RemoteIP);
        else if (Old_Route<>NULL) /* no best route anymore */
            send_WITHDRAW(Msg.prefix,RemoteAS,RemoteIP);
    }
}
```



# BGP and IP

## A first example


### Initial updates



# BGP and IP

## A first example

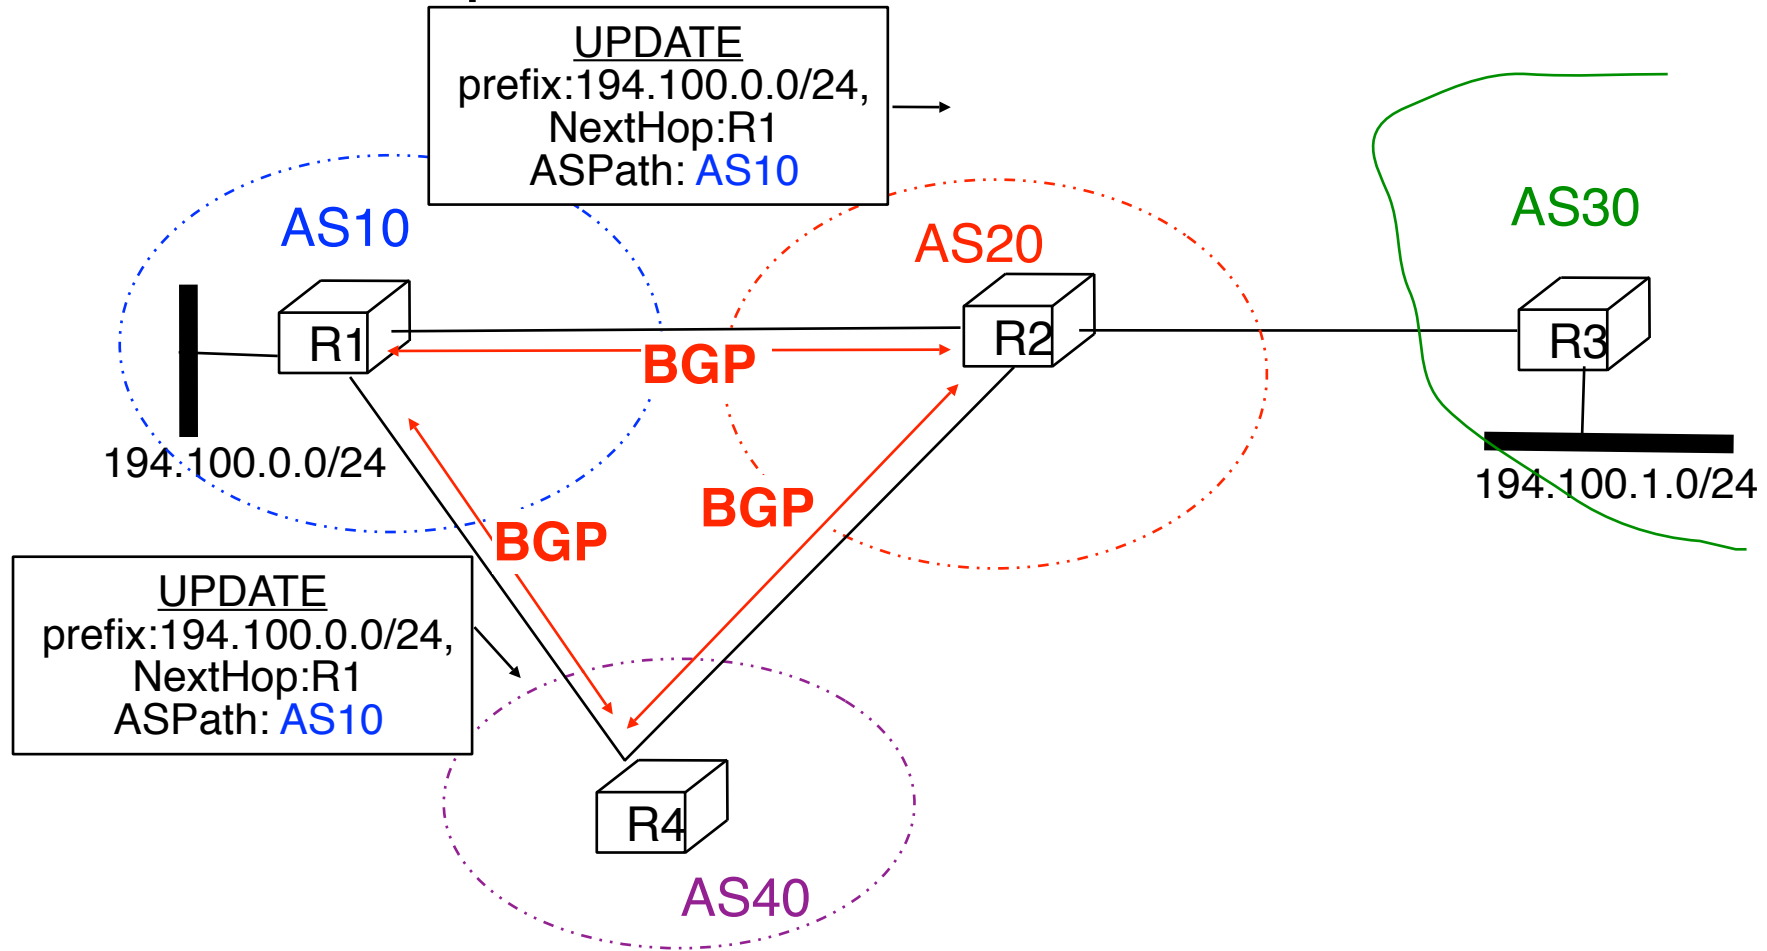
### Initial updates



# BGP and IP

## A first example

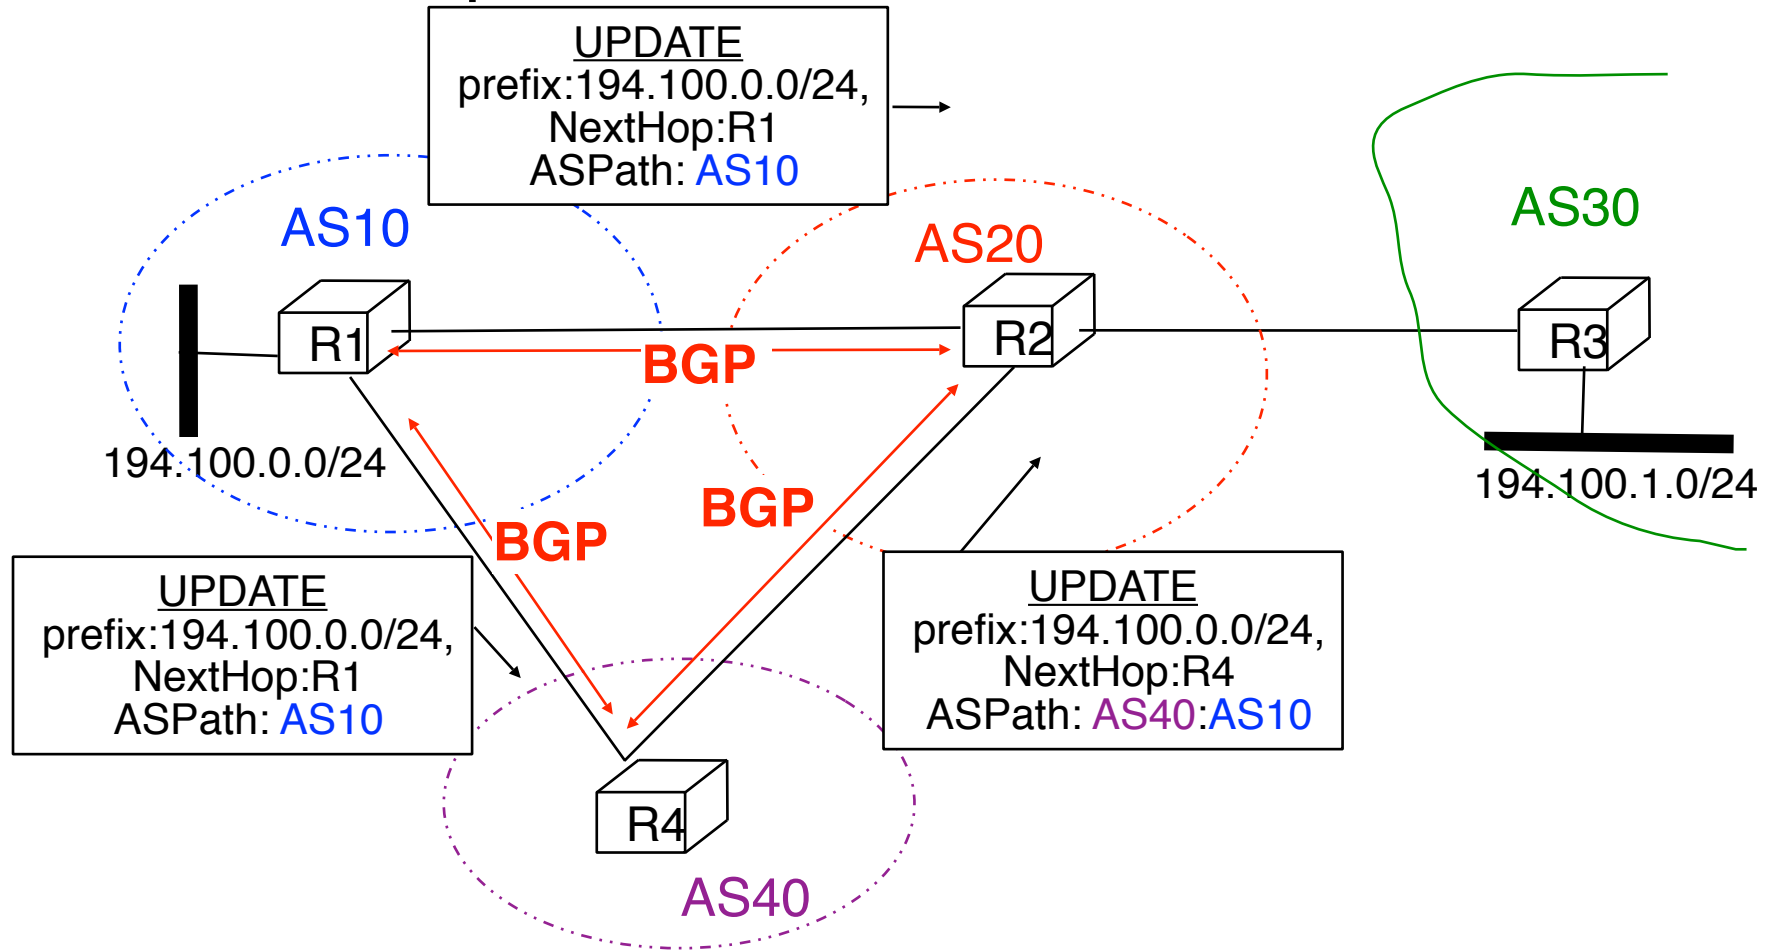
### Initial updates



# BGP and IP

## A first example

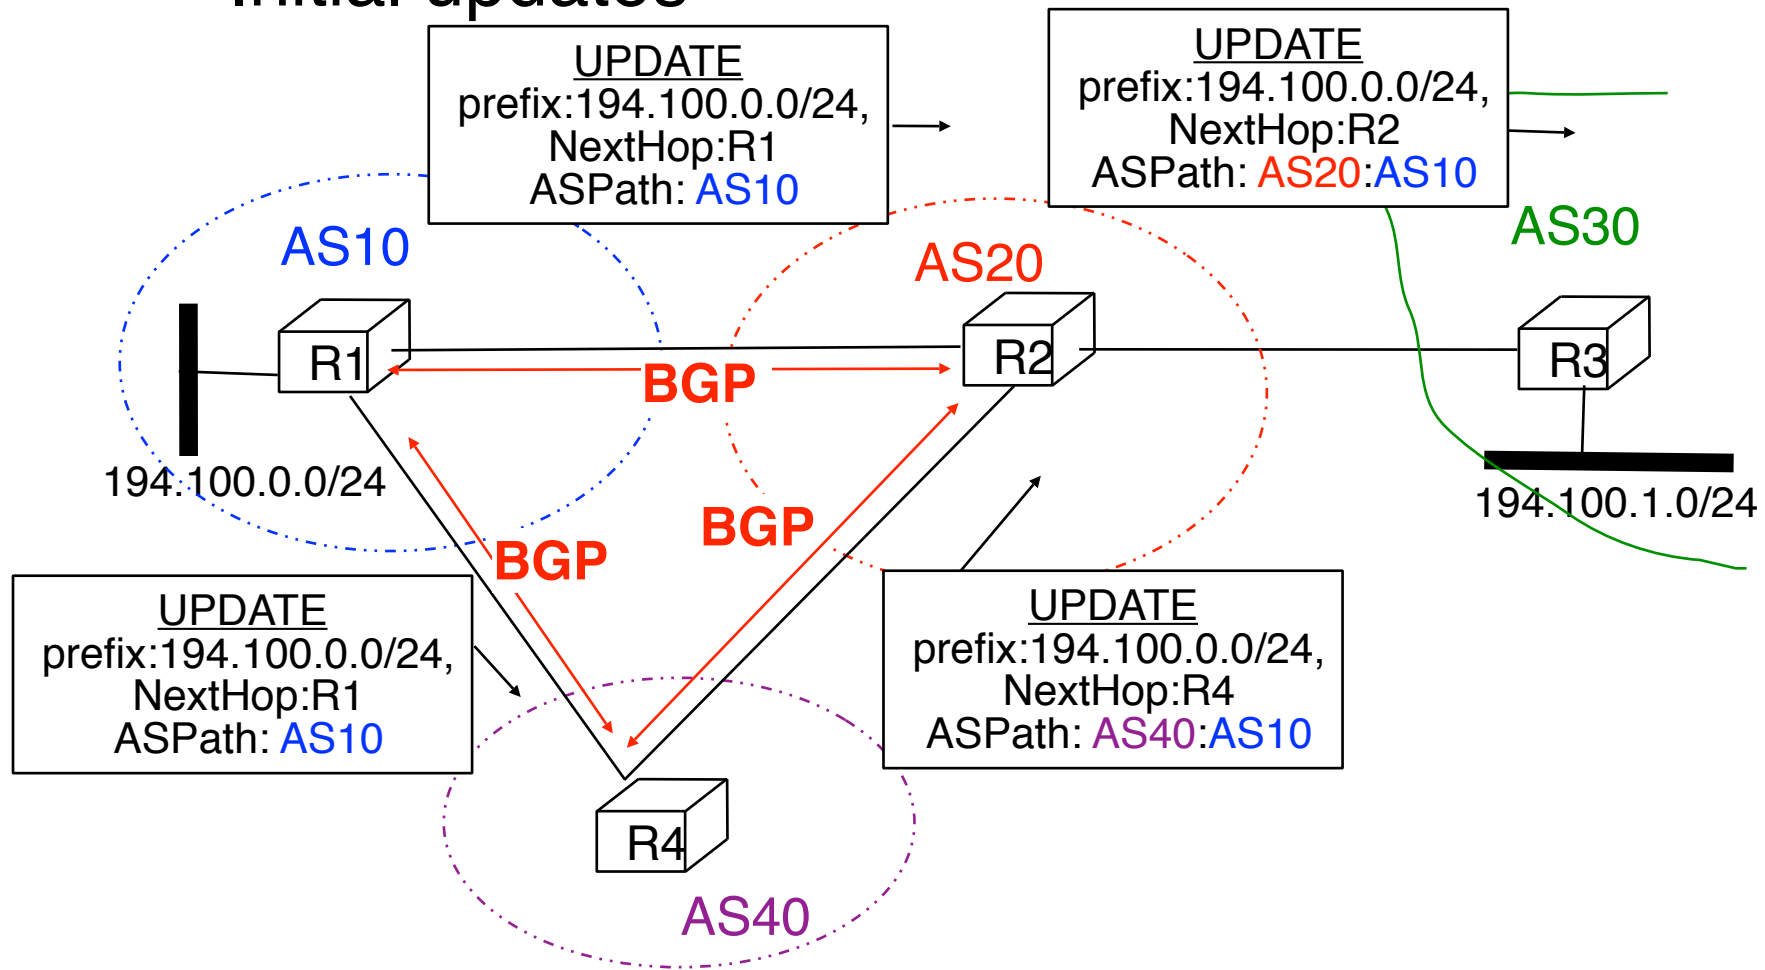
### Initial updates



# BGP and IP

## A first example

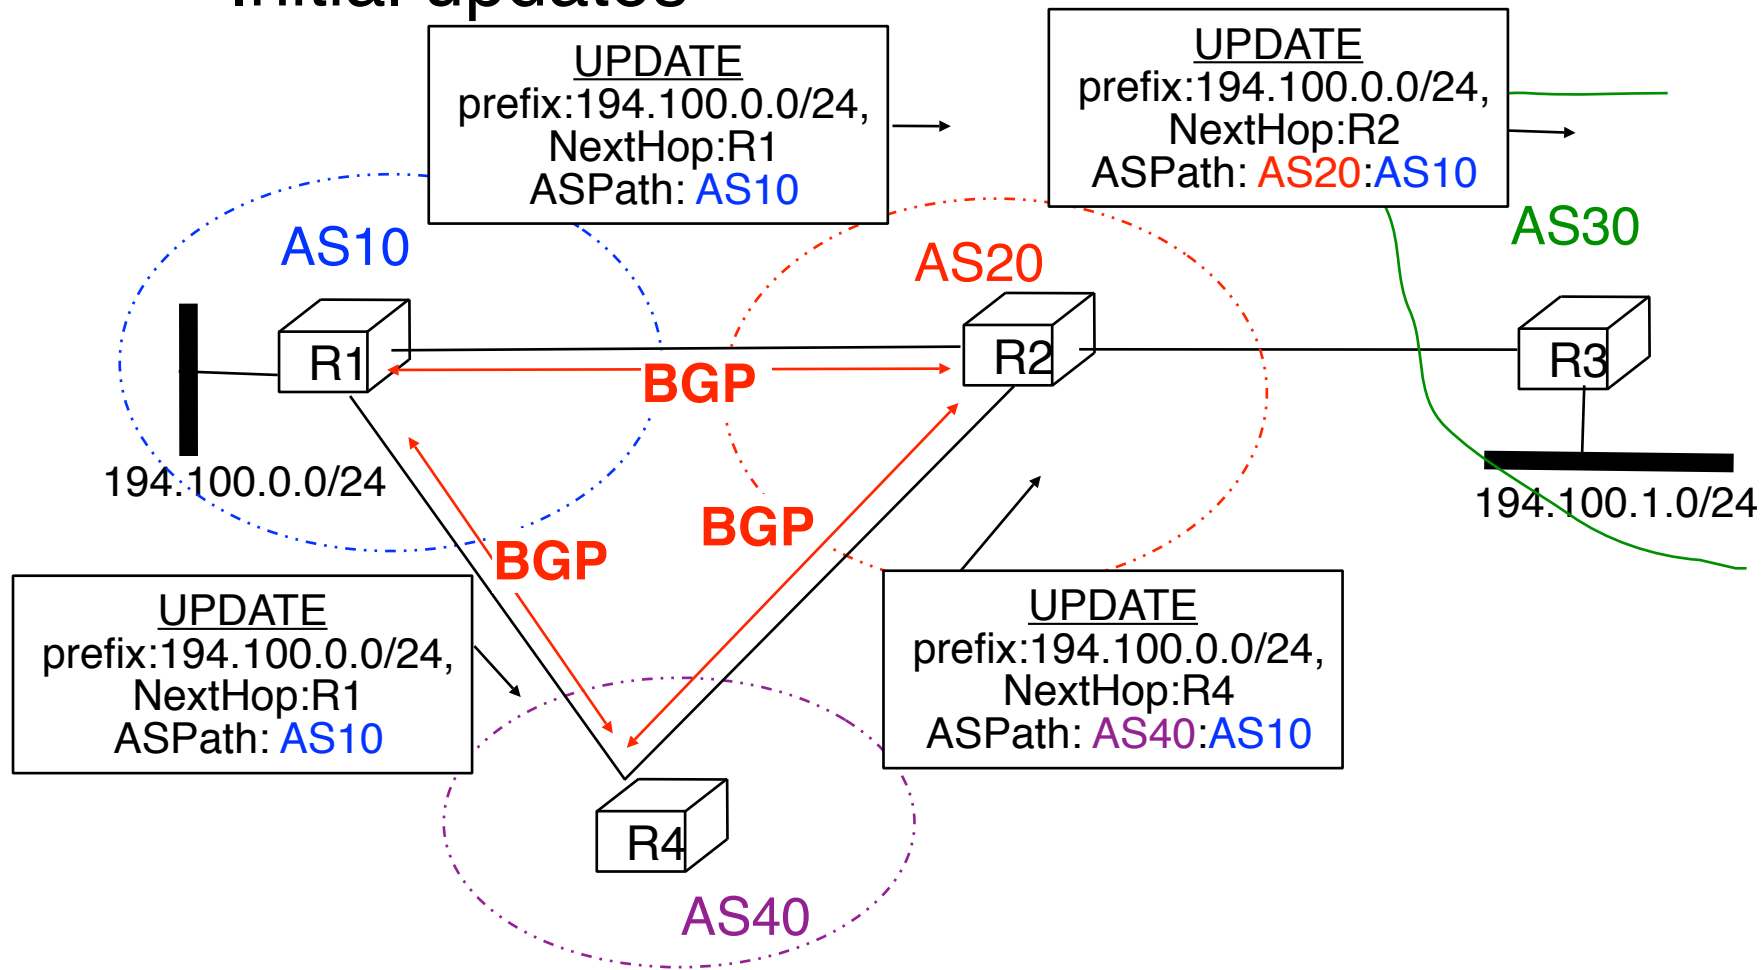
### Initial updates



# BGP and IP

## A first example


### Initial updates



What happens if link **AS10-AS20** goes down ?

# BGP and IP

## A second example




Main Path attributes of UPDATE message

NextHop : IP address of router used to reach destination

ASPath : Path followed by the route advertisement

# BGP and IP

## A second example



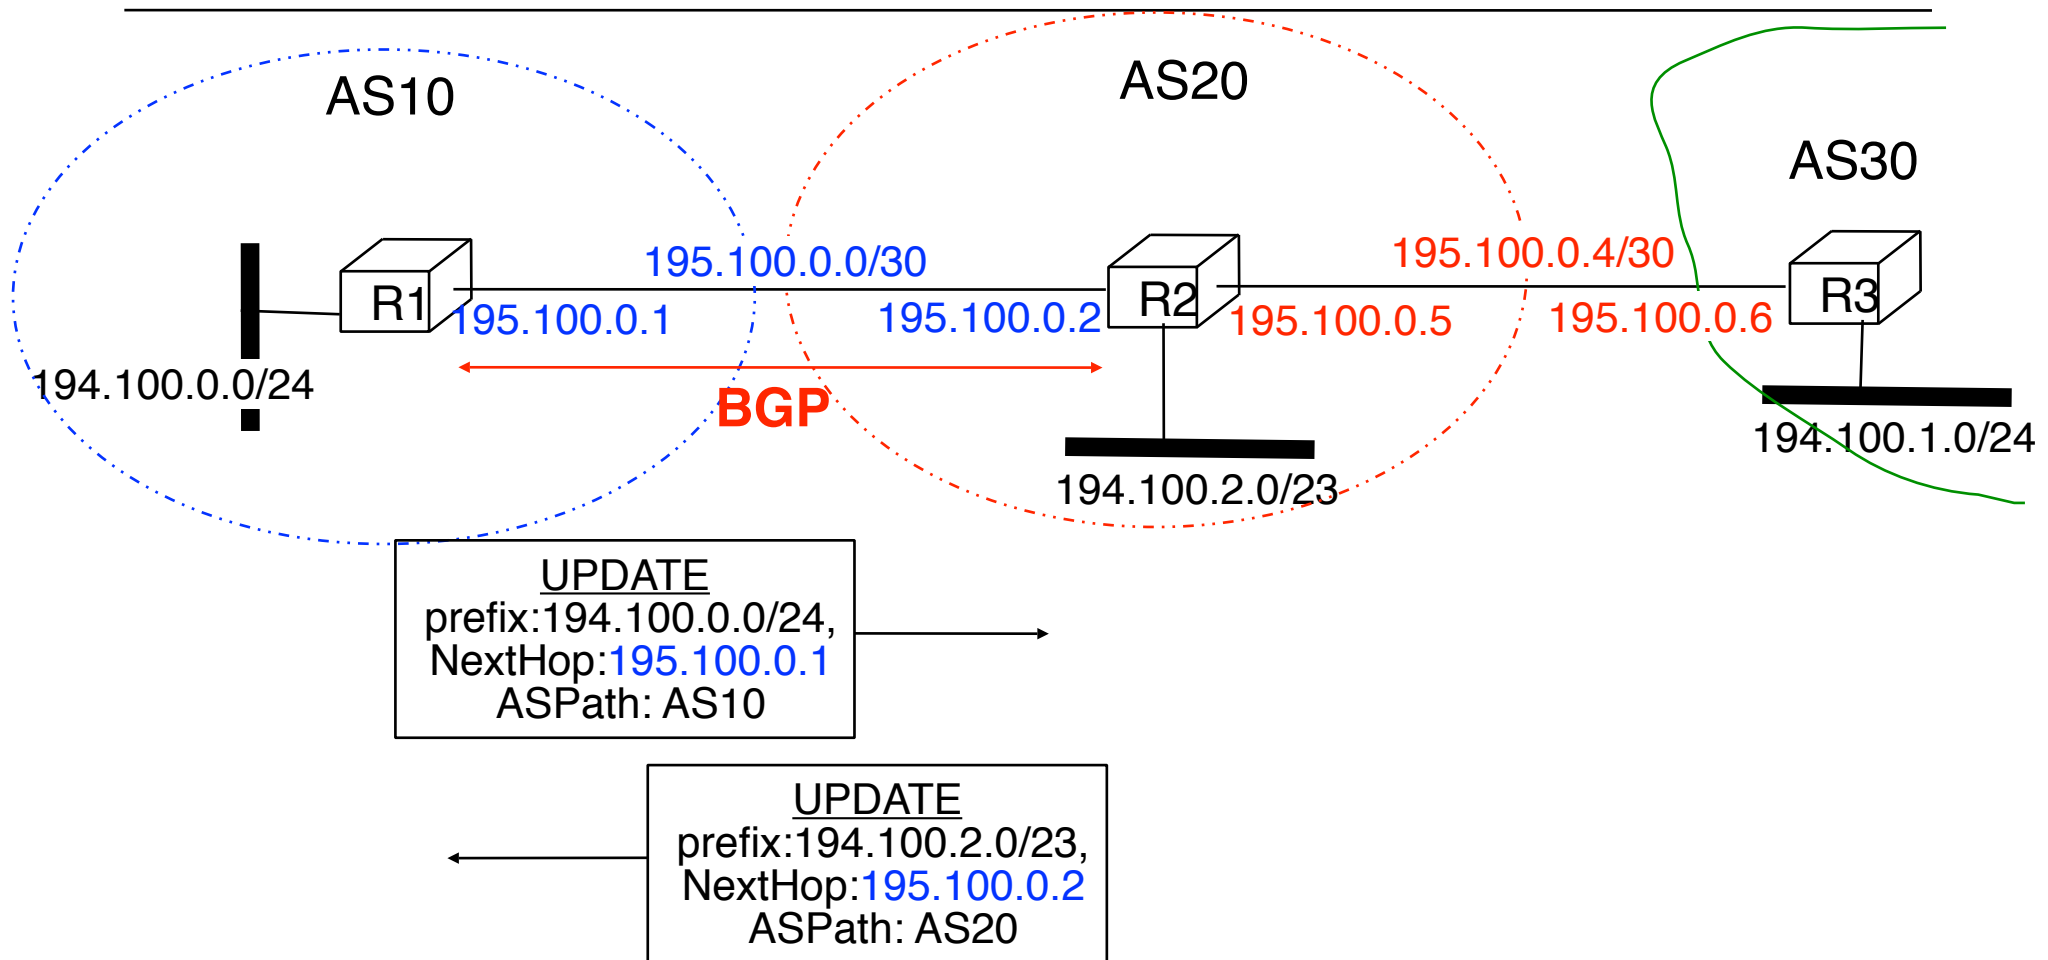
### Main Path attributes of UPDATE message

NextHop : IP address of router used to reach destination  
ASPath : Path followed by the route advertisement



# BGP and IP

## A second example




### Main Path attributes of UPDATE message

NextHop : IP address of router used to reach destination  
ASPath : Path followed by the route advertisement

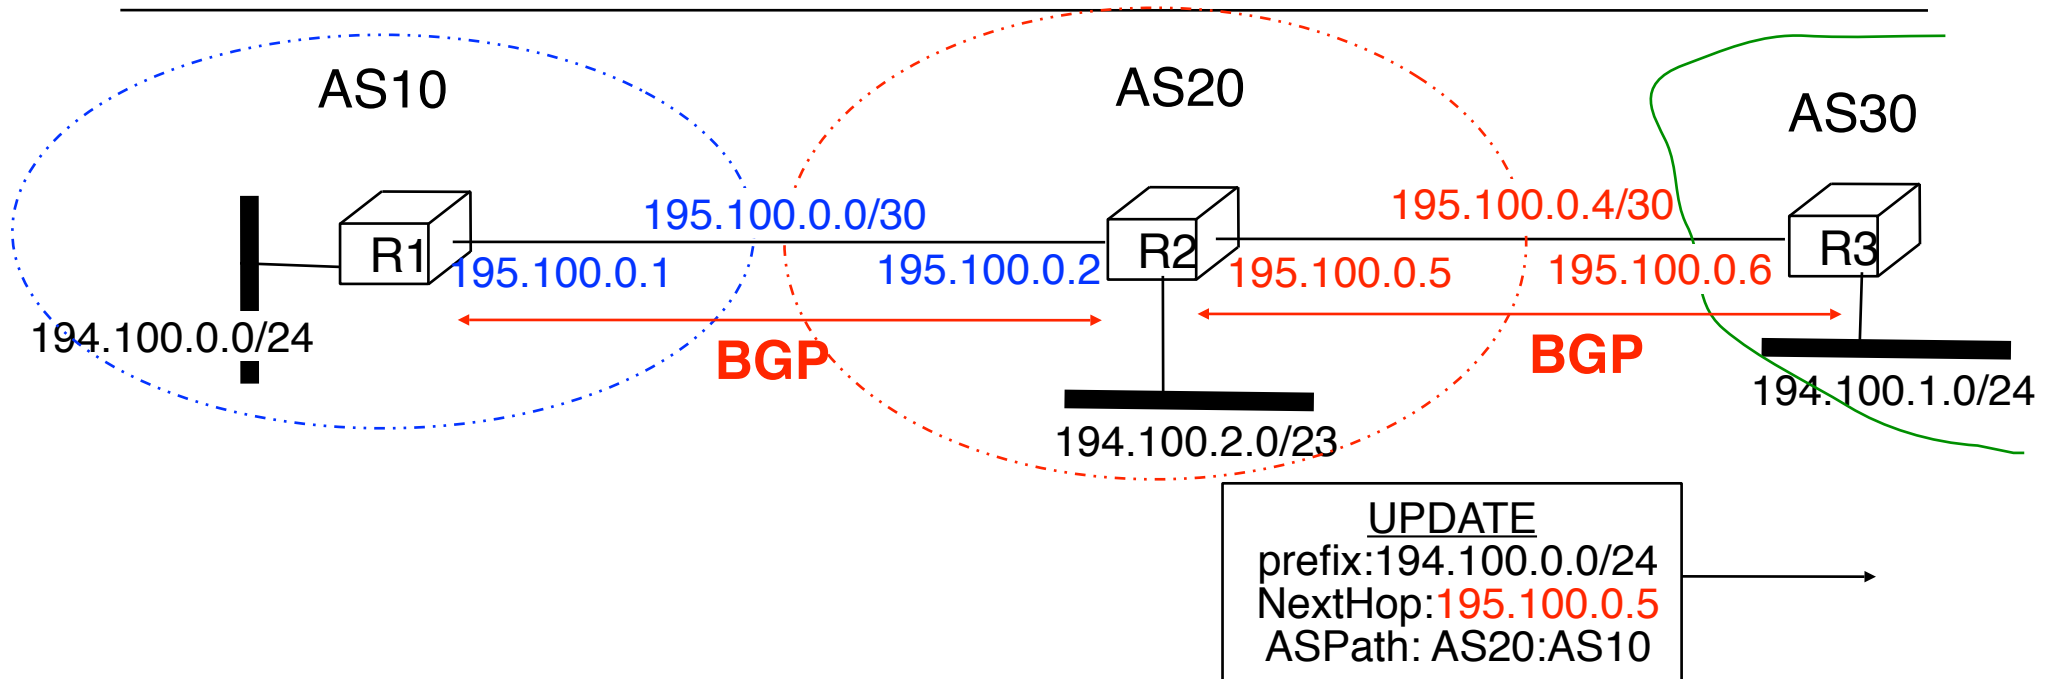
# BGP and IP

## A second example (2)



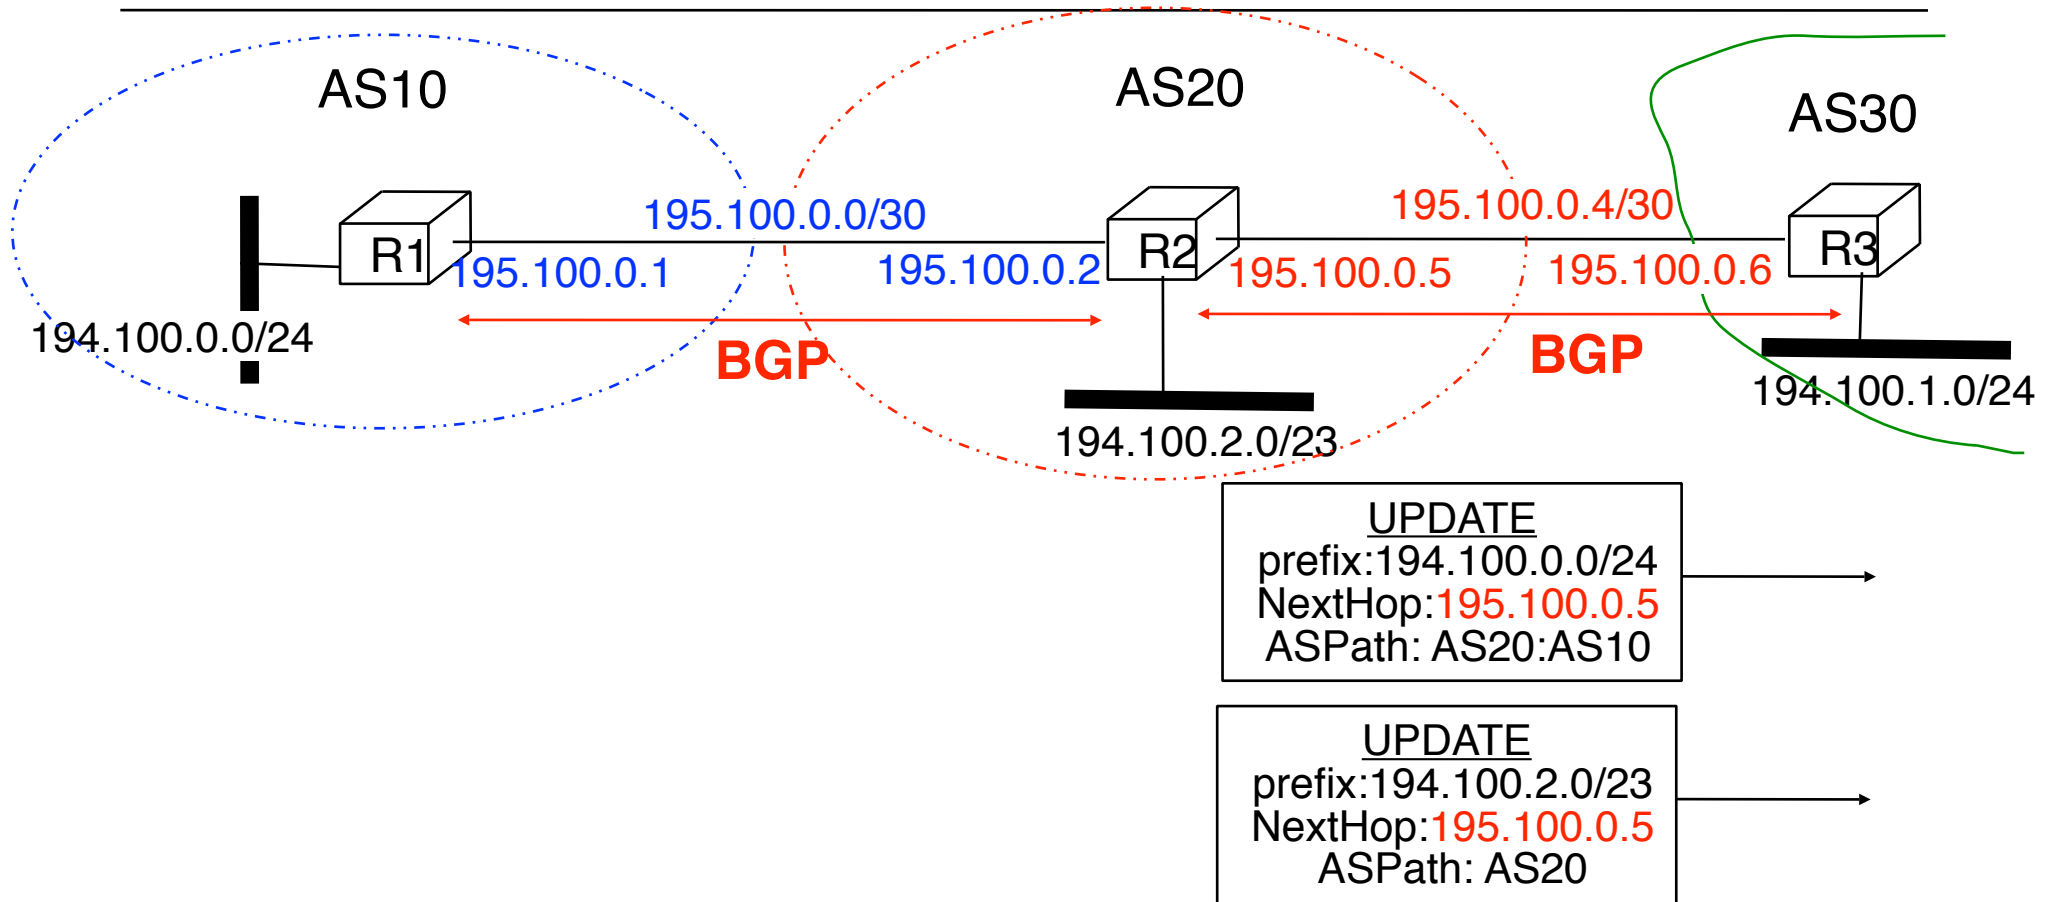
# BGP and IP

## A second example (2)



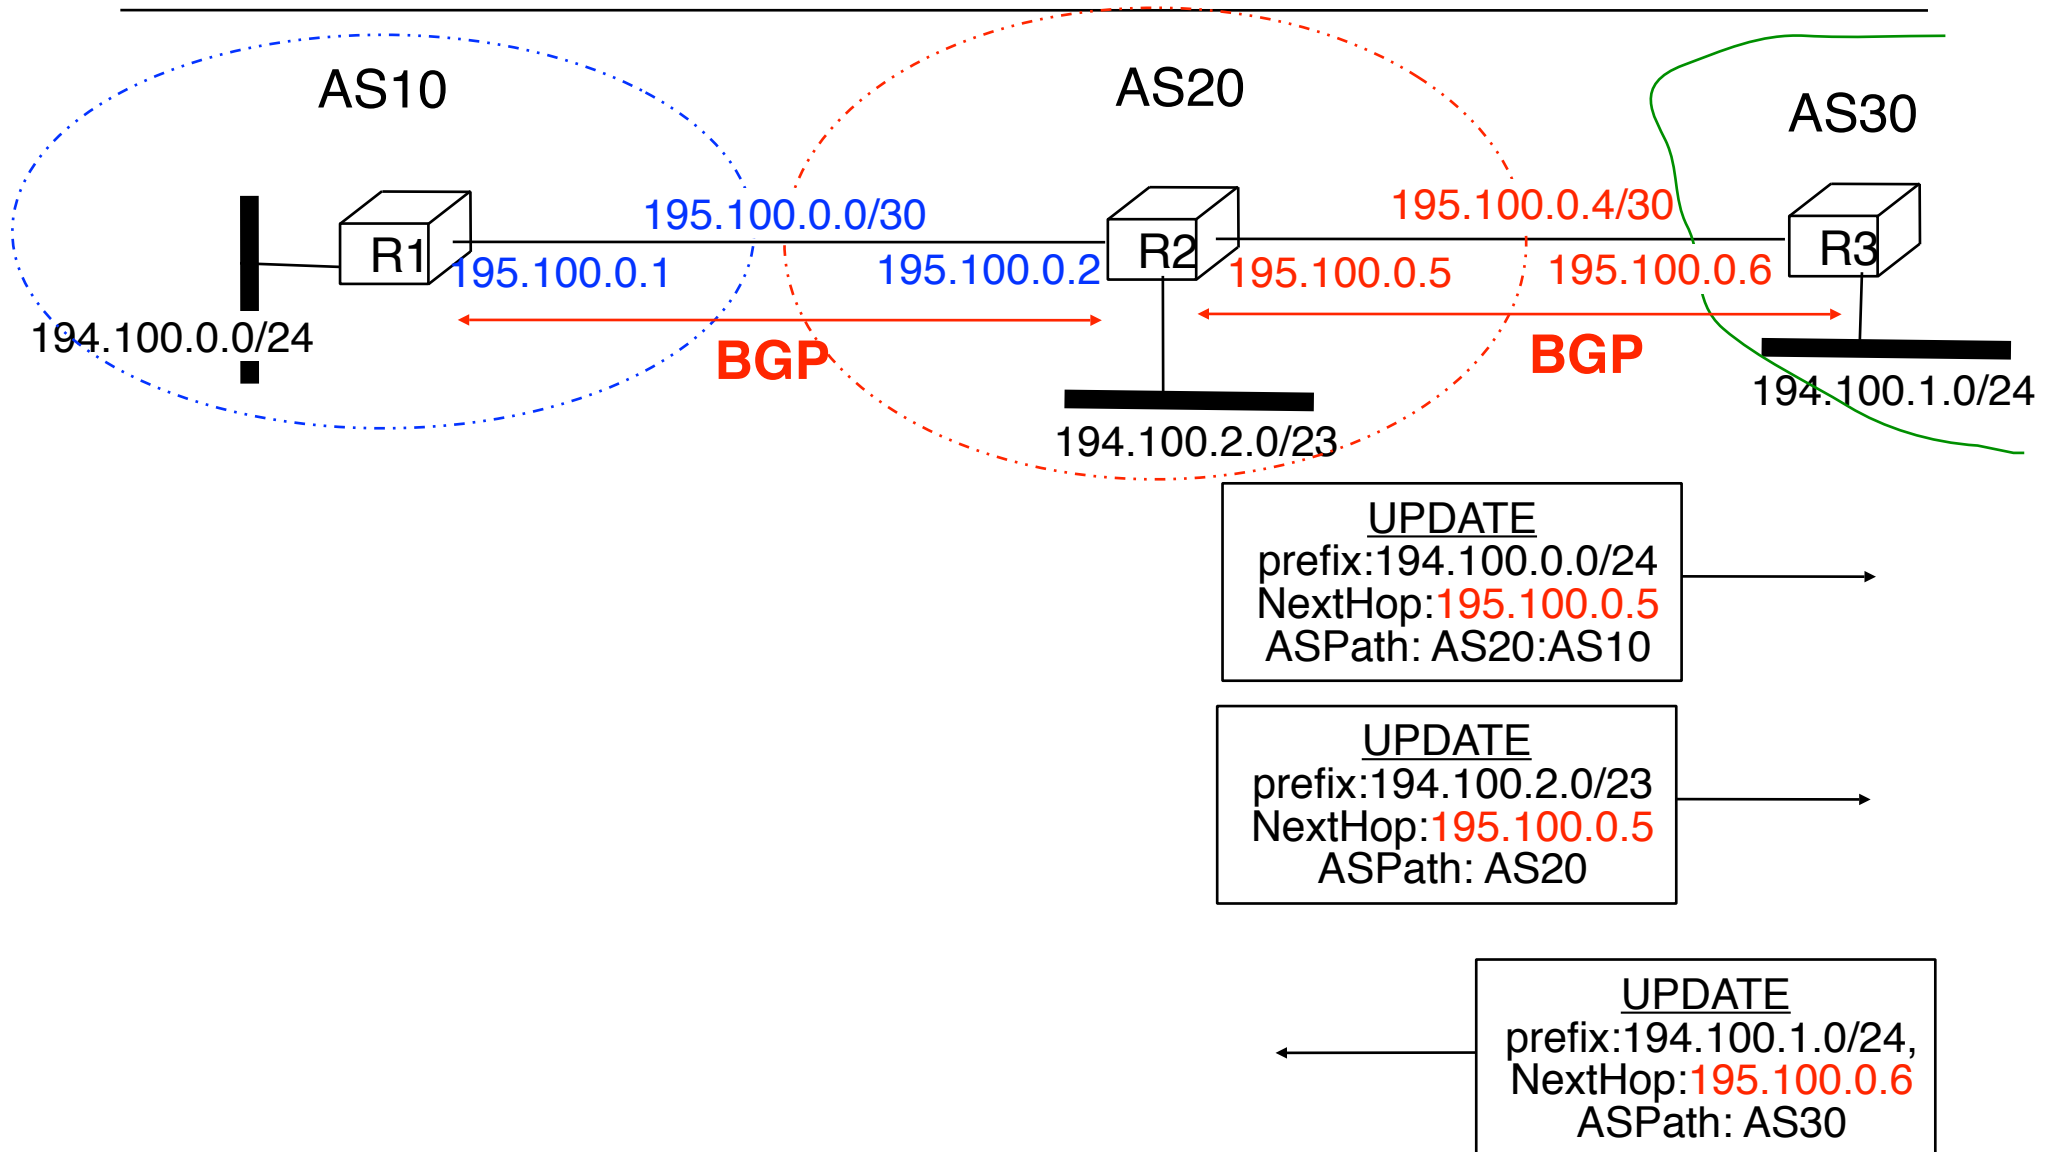
# BGP and IP

## A second example (2)



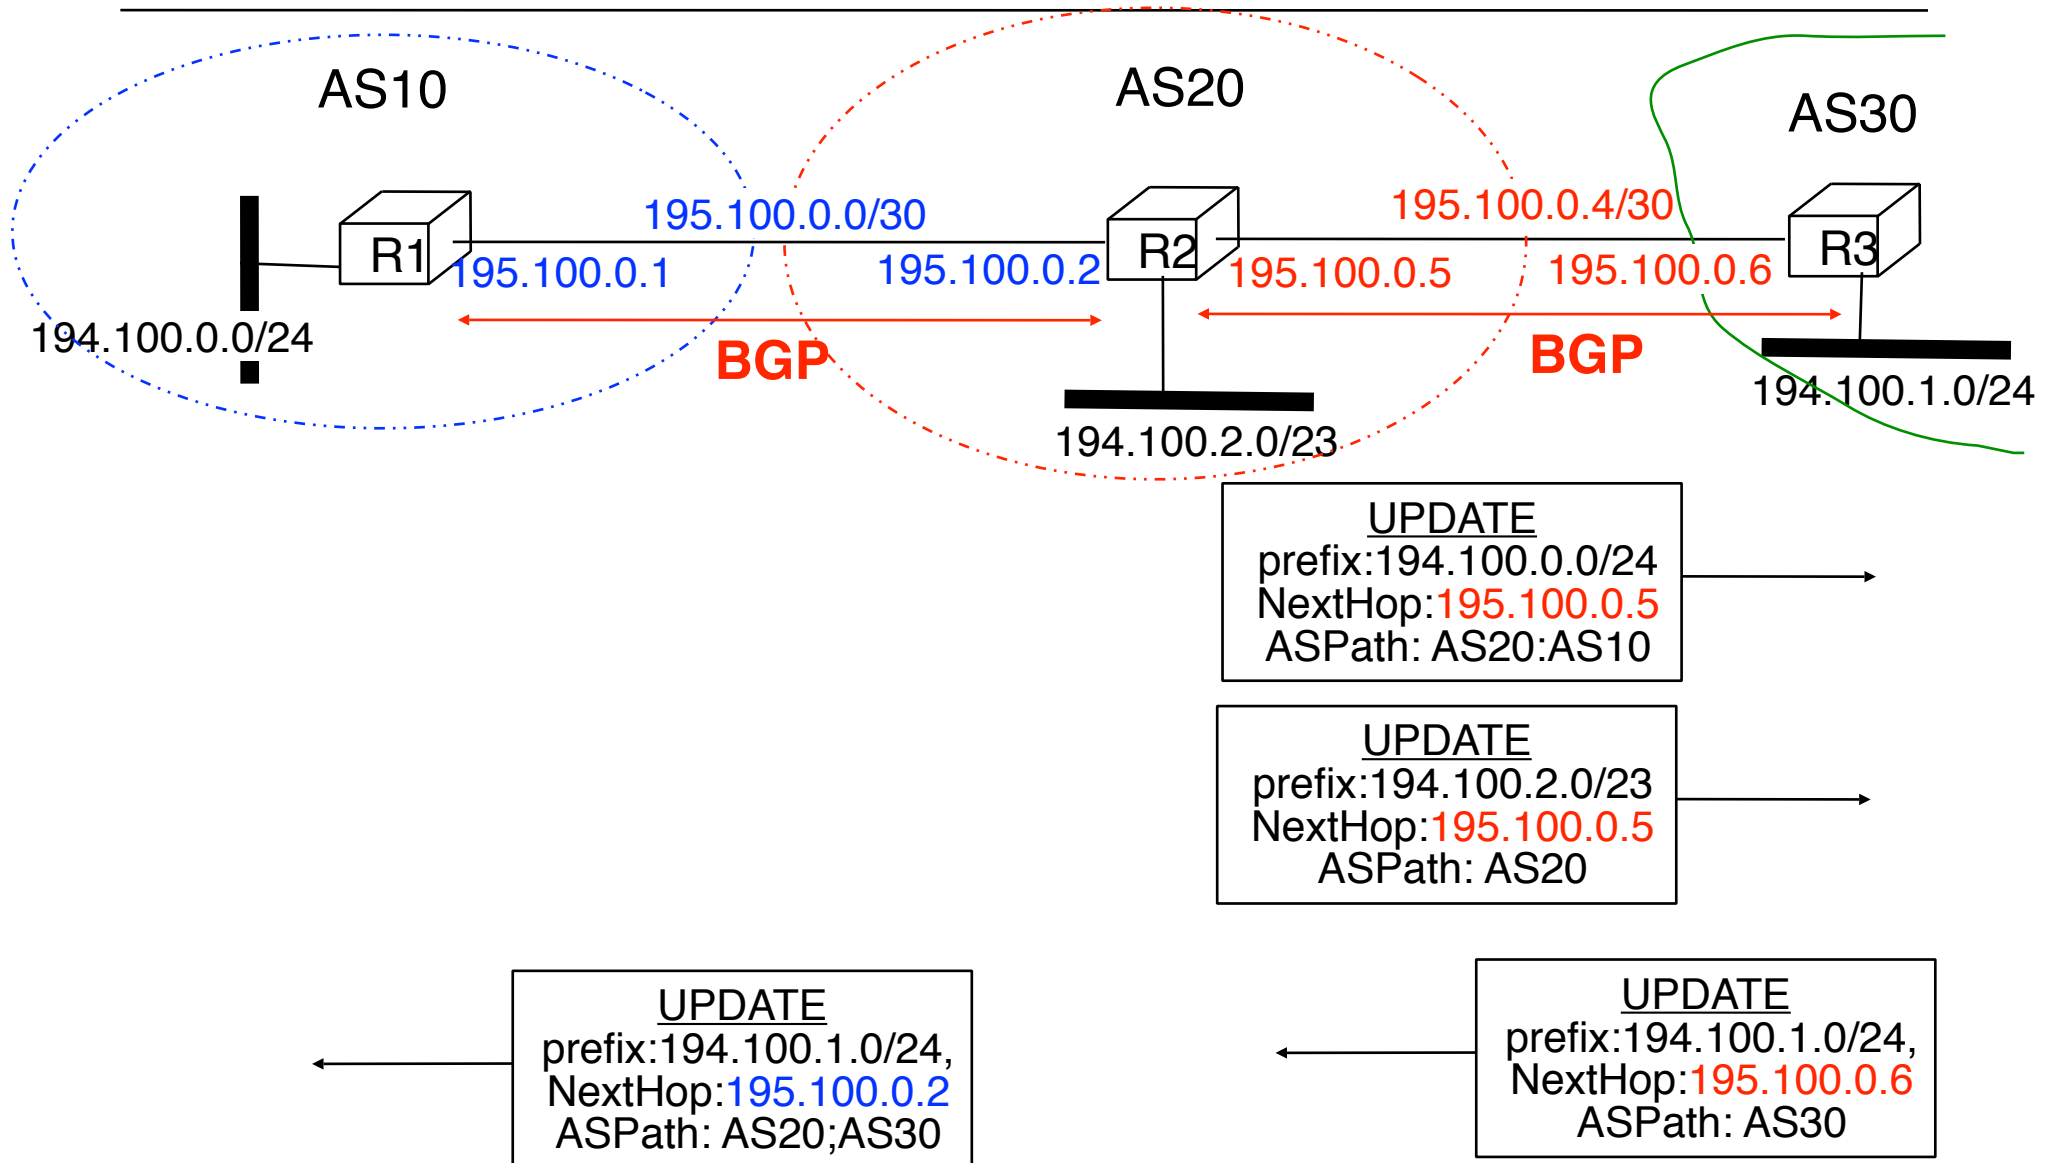
# BGP and IP

## A second example (2)



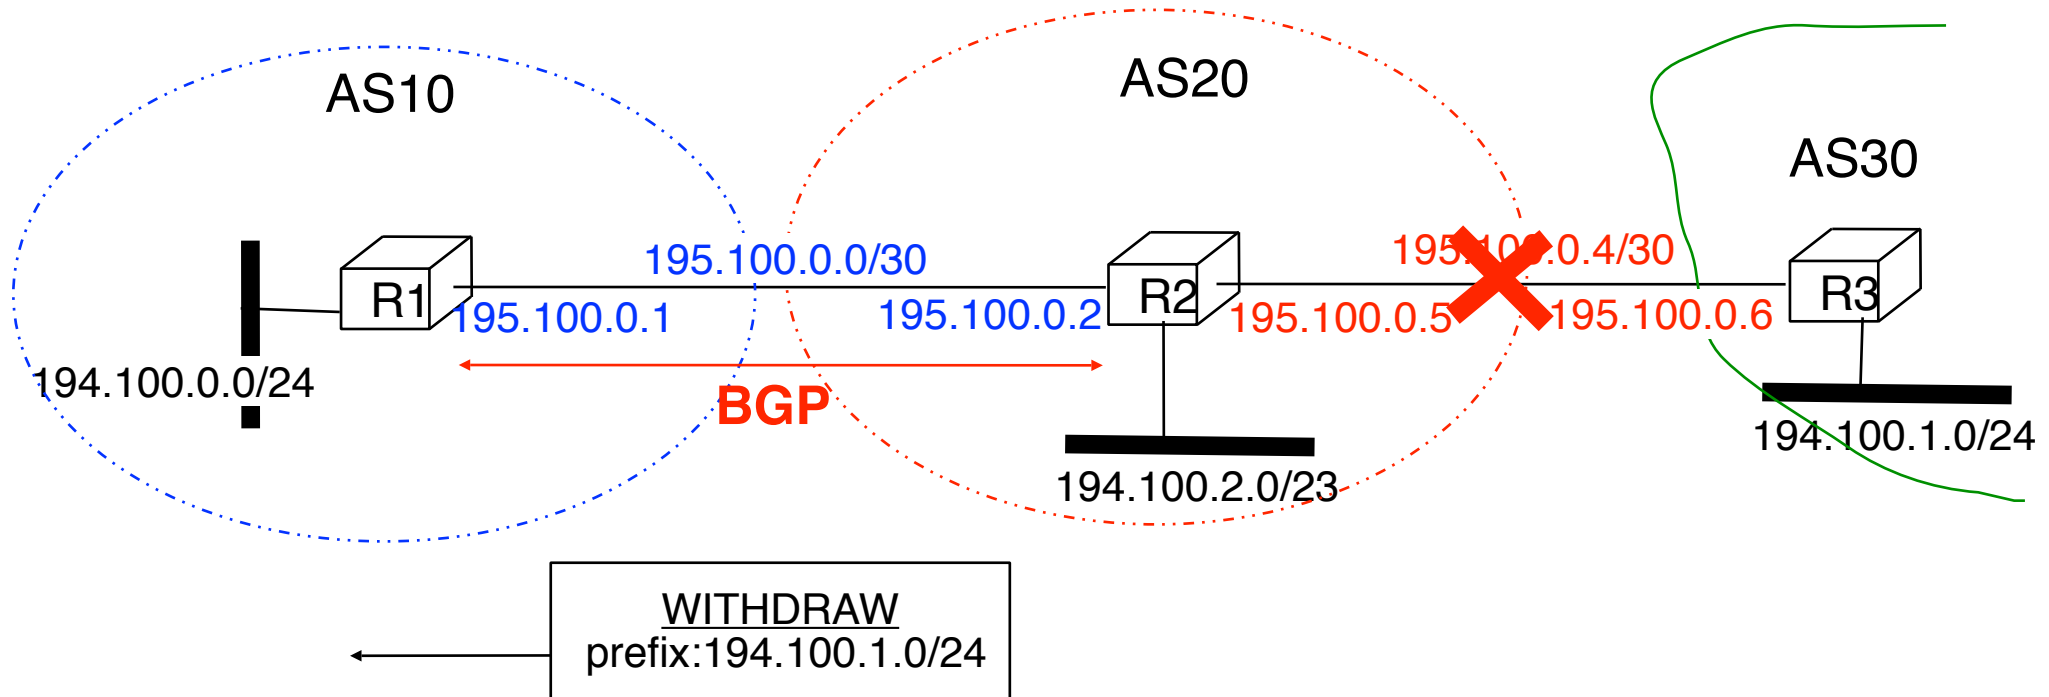
# BGP and IP

## A second example (2)




# BGP and IP

## A second example (3)




# How to prefer some routes over others ?

---



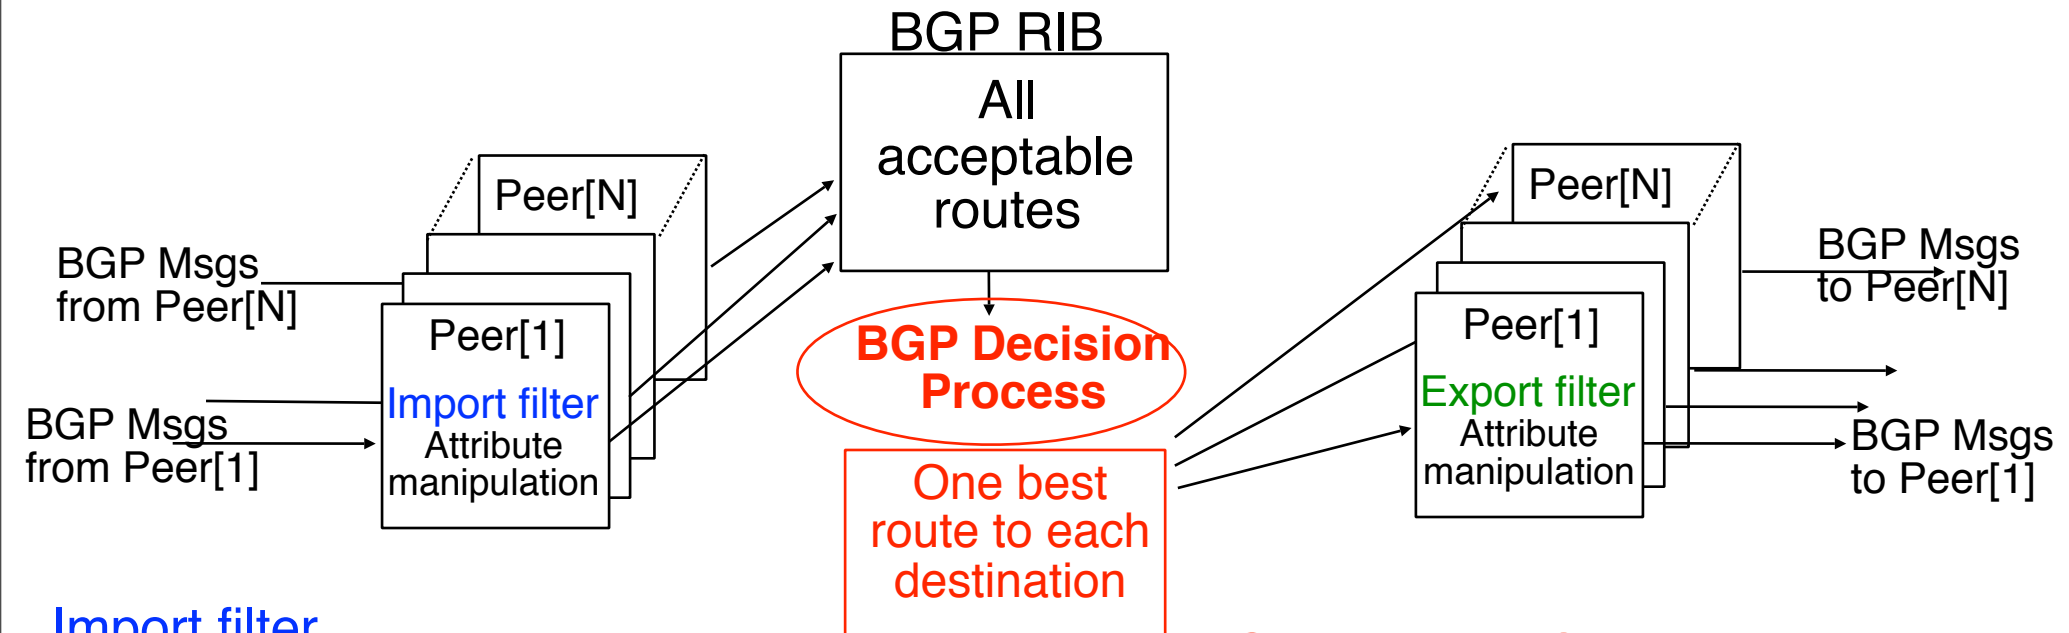
How to ensure that packets will flow on primary link ?



How to prefer cheap link over expensive link ?



# How to prefer some routes over others (2) ?



## Import filter


Selection of acceptable routes  
Addition of `local-pref` attribute  
inside received BGP Msg  
Normal quality route : `local-pref=100`  
Better than normal route : `local-pref=200`  
Worse than normal route : `local-pref=50`

## Simplified BGP Decision Process

Select routes with highest `local-pref`  
If there are several routes,  
choose routes with the  
shortest ASPath  
If there are still several routes  
tie-breaking rule

# How to prefer some routes over others (3) ?

---



## **RPSL-like policy for AS1**

aut-num: AS1

import: from AS2 RA at R1 set localpref=100;  
from AS2 RB at R1 set localpref=200;  
accept ANY

export: to AS2 RA at R1 announce AS1  
to AS2 RB at R1 announce AS1

## **RPSL-like policy for AS2**


aut-num: AS2

import: from AS1 R1 at RA set localpref=100;  
from AS1 R1 at RB set localpref=200;  
accept AS1

export: to AS1 R1 at RA announce ANY  
to AS2 R1 at RB announce ANY

# How to prefer some routes over others (4) ?

---



## **RPSL policy for AS1**

aut-num: AS1

import: from AS2 RA at R1 set localpref=100;

from AS4 R2 at R1 set localpref=200;

accept ANY

export: to AS2 RA at R1 announce AS1


to AS4 R2 at R1 announce AS1

AS1 will prefer to send packets over the cheap link  
But the flow of the packets destined to AS1 will depend  
on the routing policy of the other domains

# Limitations of local-pref

## In theory

Each domain is free to define its order of preference for the routes learned from external peers




How to reach 1.0.0.0/8 from AS3 and AS4 ?

# Limitations of local-pref (2)


---

AS1 sends its UPDATE messages ...



# Limitations of local-pref (2)

AS1 sends its UPDATE messages ...



**Preferred paths for AS3**

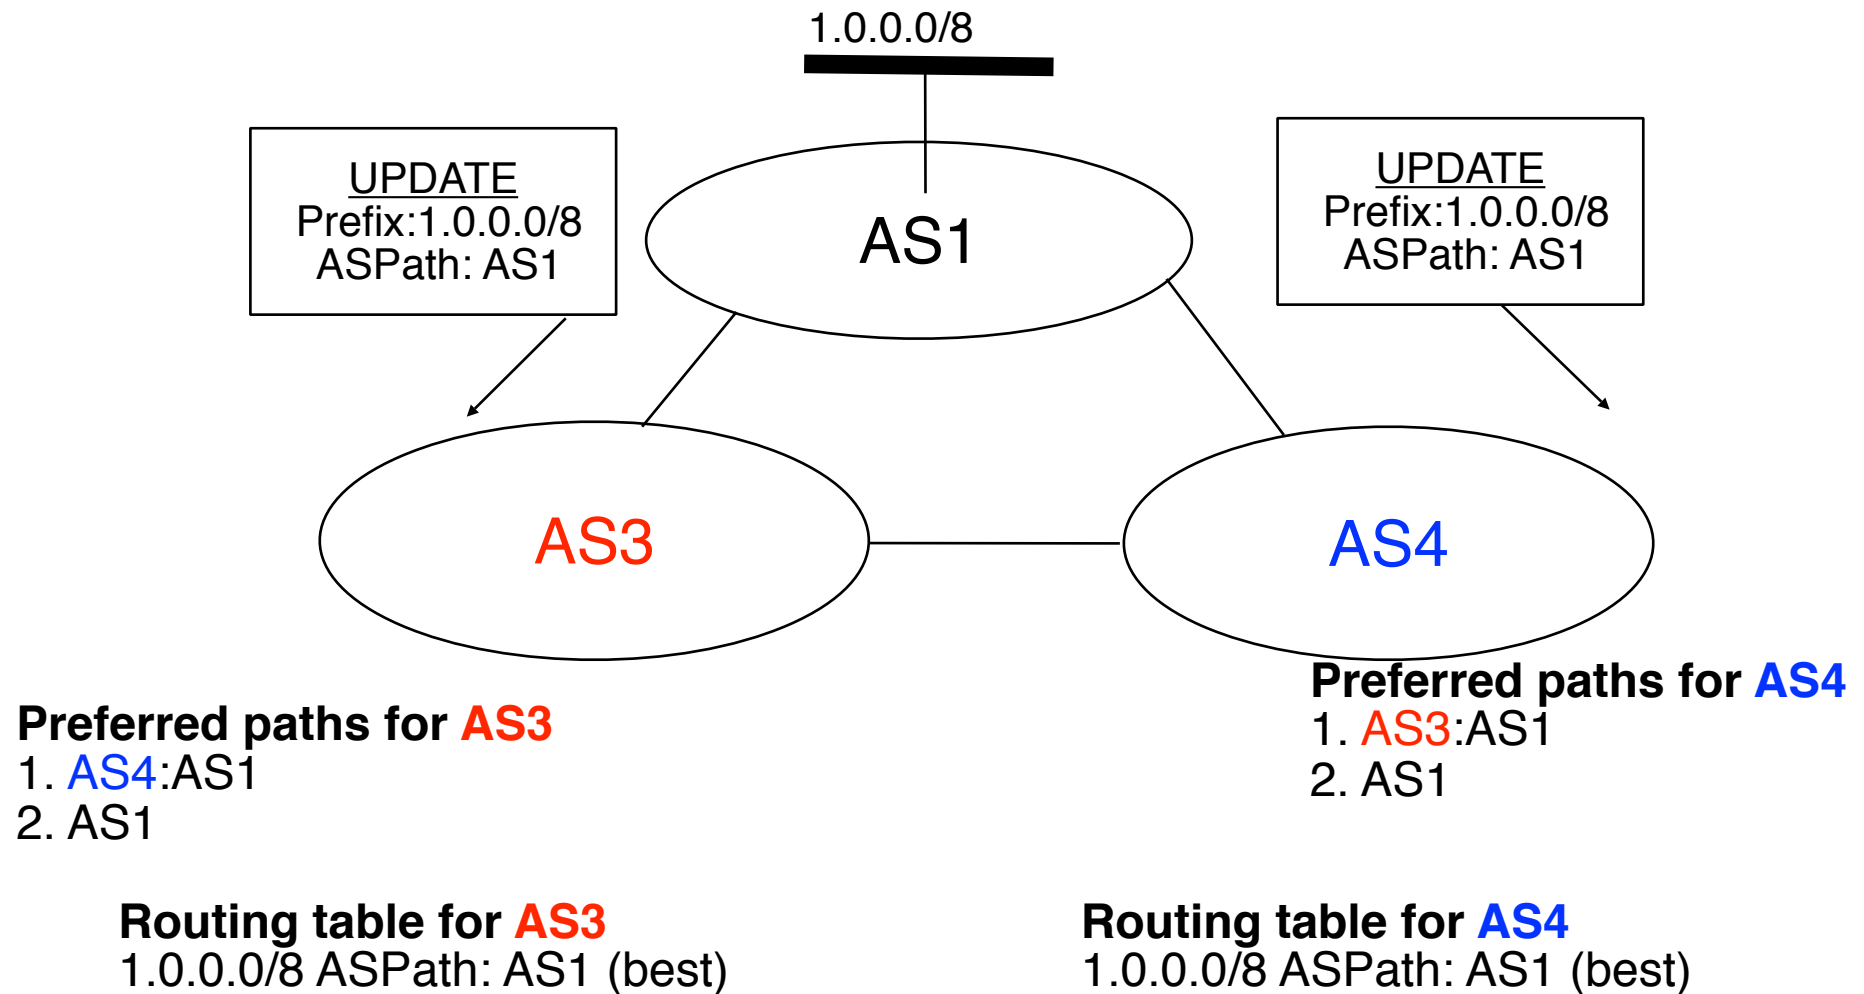
1. AS4:AS1
2. AS1

**Routing table for AS3**

1.0.0.0/8 ASPath: AS1 (best)

# Limitations of local-pref (2)

AS1 sends its UPDATE messages ...




# Limitations of local-pref (3)

---

First possibility

**AS3** sends its UPDATE first...



**Preferred paths for AS3**

1. **AS4**:AS1
2. AS1

**Routing table for AS3**

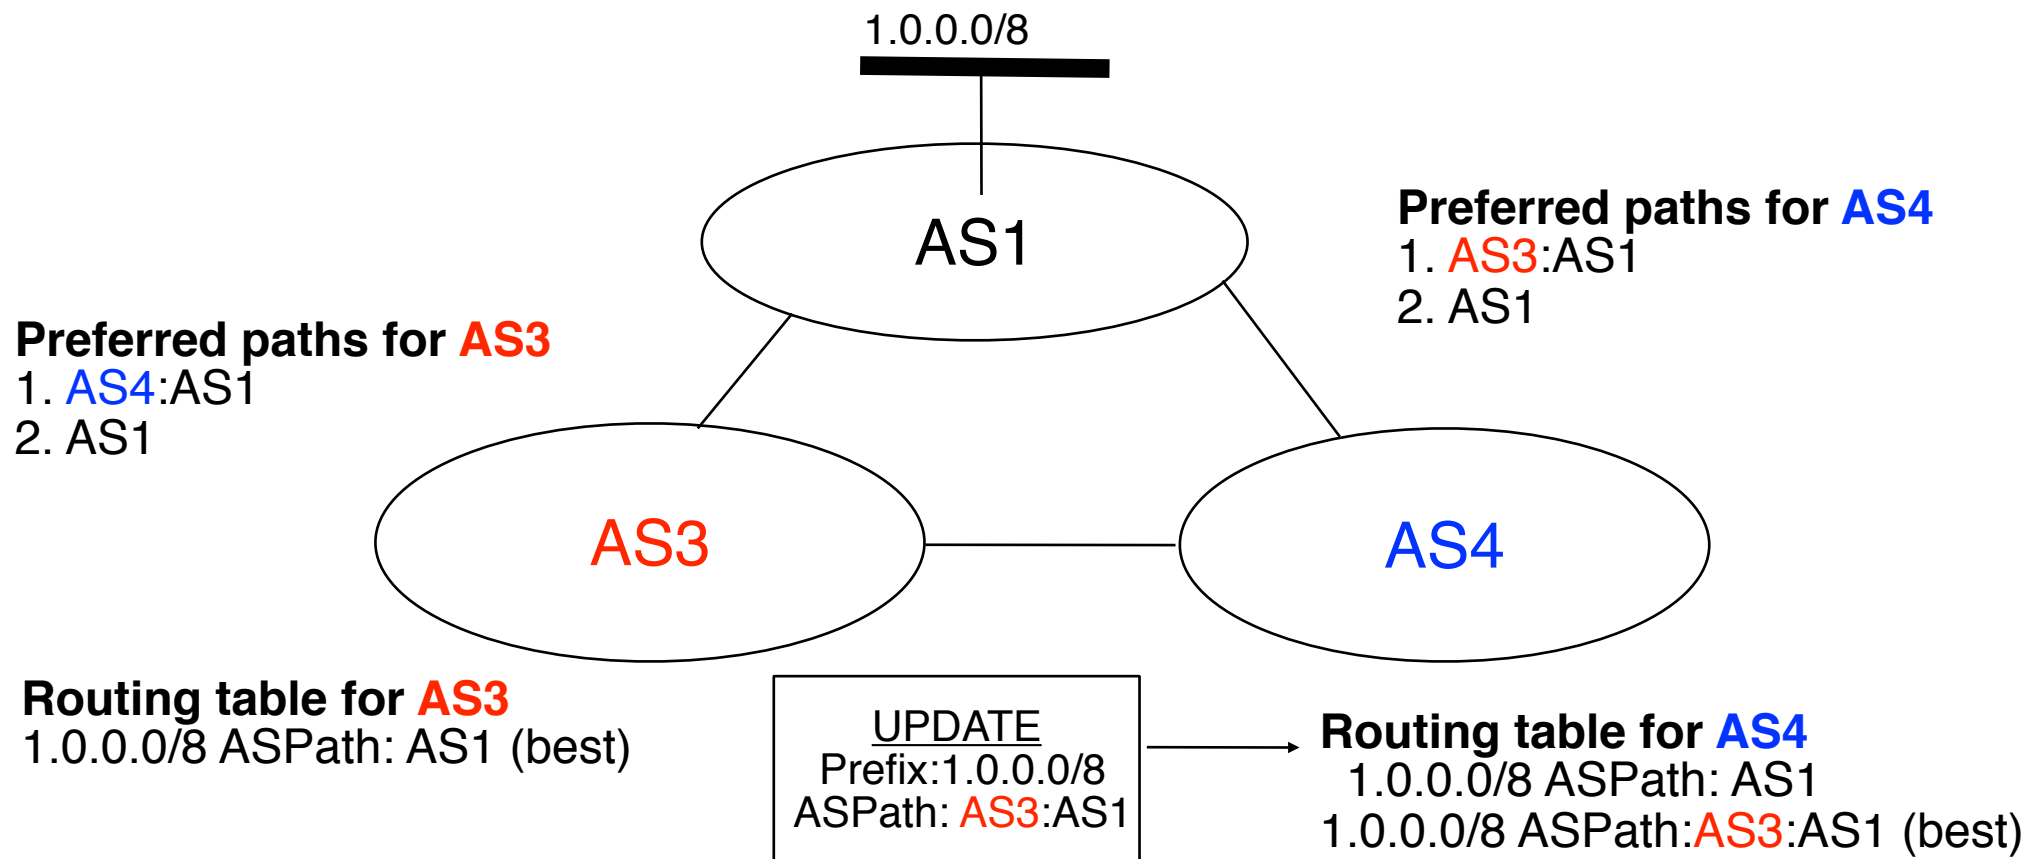
1.0.0.0/8 ASPath: AS1 (best)



# Limitations of local-pref (3)

First possibility

**AS3** sends its UPDATE first...




Stable route assignment

# Limitations of local-pref (4)

## Second possibility

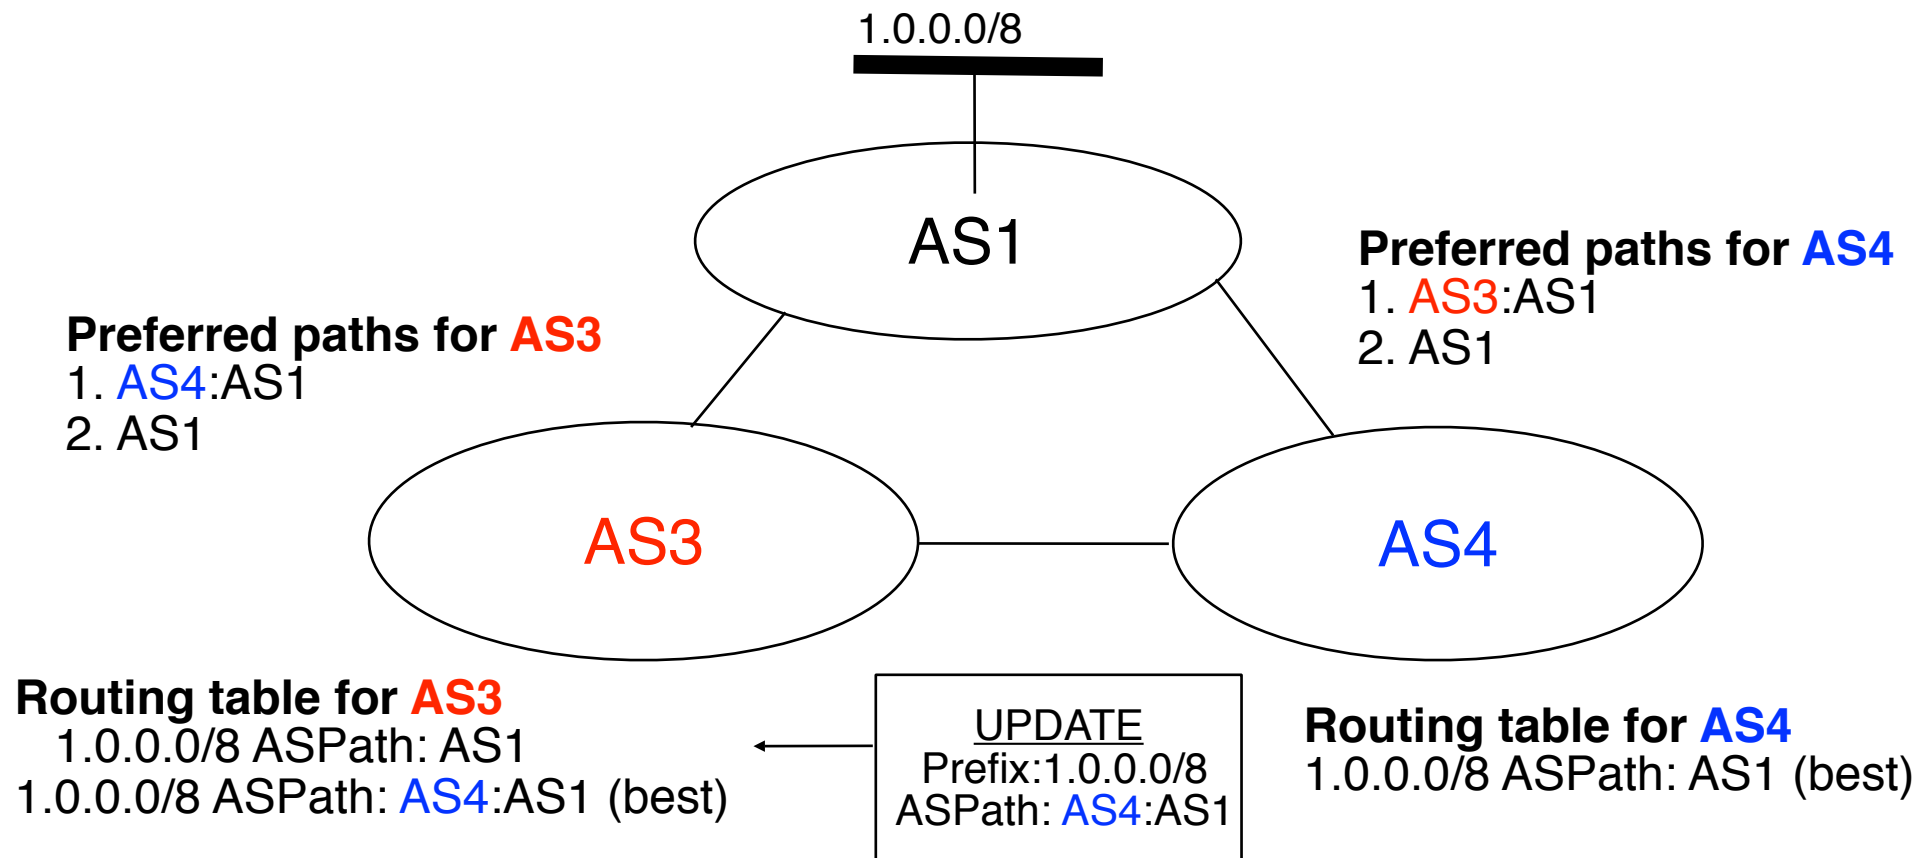
**AS4** sends its UPDATE first...



# Limitations of local-pref (4)

## Second possibility

**AS4** sends its UPDATE first...




Another (but different) stable route assignment

# Limitations of local-pref (5)

## Third possibility


**AS3** and **AS4** send their UPDATE together...



# Limitations of local-pref (5)

## Third possibility

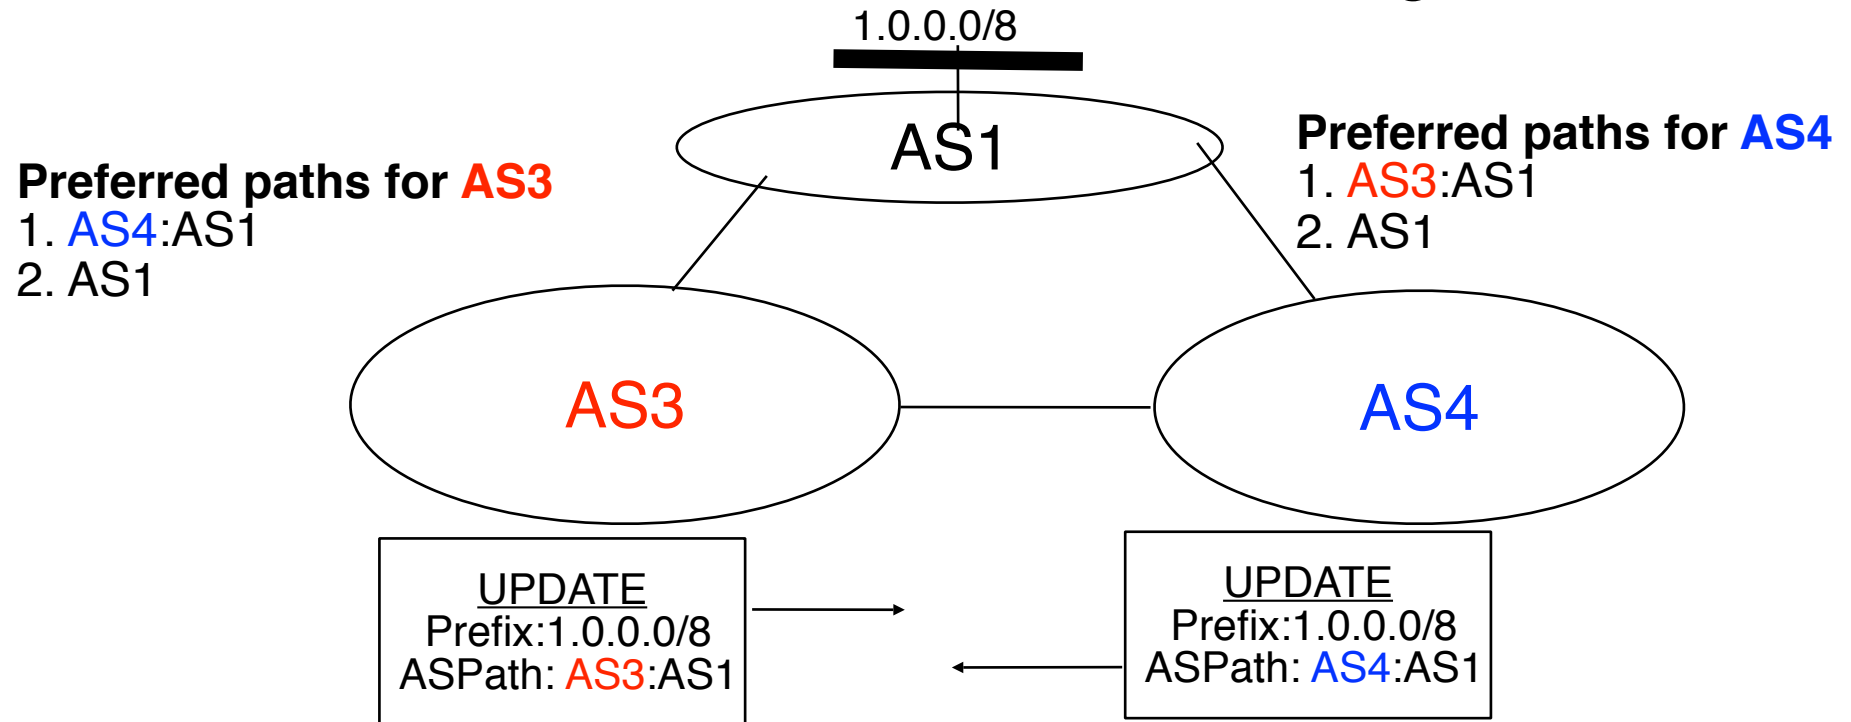
**AS3** and **AS4** send their UPDATE together...



# Limitations of local-pref (5)

## Third possibility

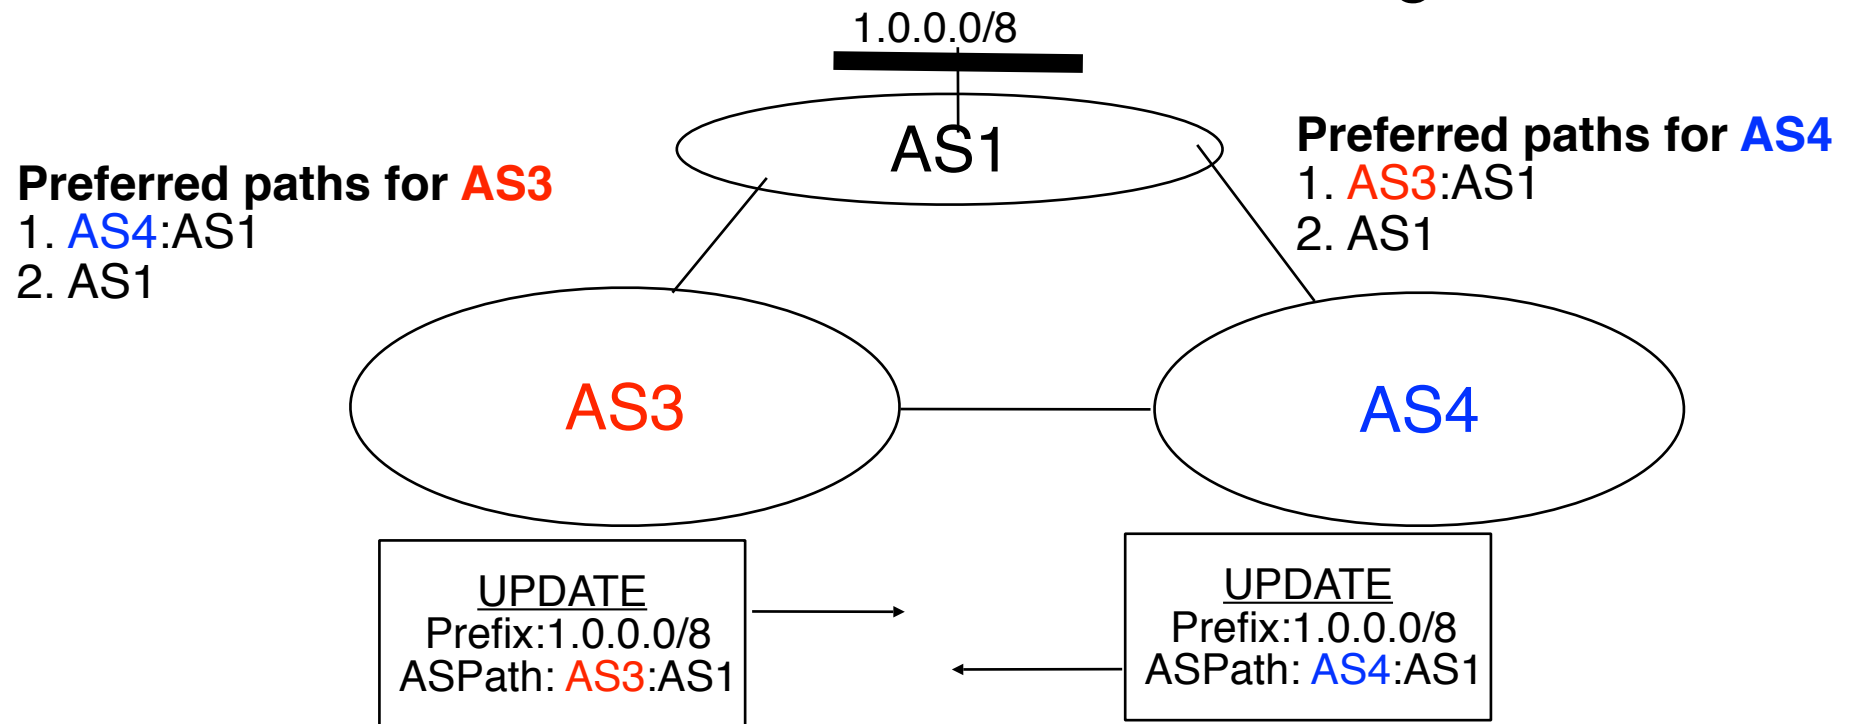
**AS3** and **AS4** send their UPDATE together...



# Limitations of local-pref (5)

## Third possibility

**AS3** and **AS4** send their UPDATE together...




**AS3** prefers the indirect path and will thus send withdraw since the chosen best path is via AS4

**AS4** prefers the indirect path and will thus send withdraw since the chosen best path is via AS3

# Limitations of local-pref (6)

Third possibility (cont.)

**AS3** and **AS4** send their UPDATE together...






# Limitations of local-pref (6)

## Third possibility (cont.)

**AS3** and **AS4** send their UPDATE together...




**AS3** learns that the indirect route is not available anymore  
AS3 will reannounce its direct route...

# Limitations of local-pref (6)

## Third possibility (cont.)

**AS3** and **AS4** send their UPDATE together...




**AS3** learns that the indirect route is not available anymore  
AS3 will reannounce its direct route...

**AS4** learns that the indirect route is not available anymore  
**AS4** will reannounce its direct route...

# More limitations of local-pref

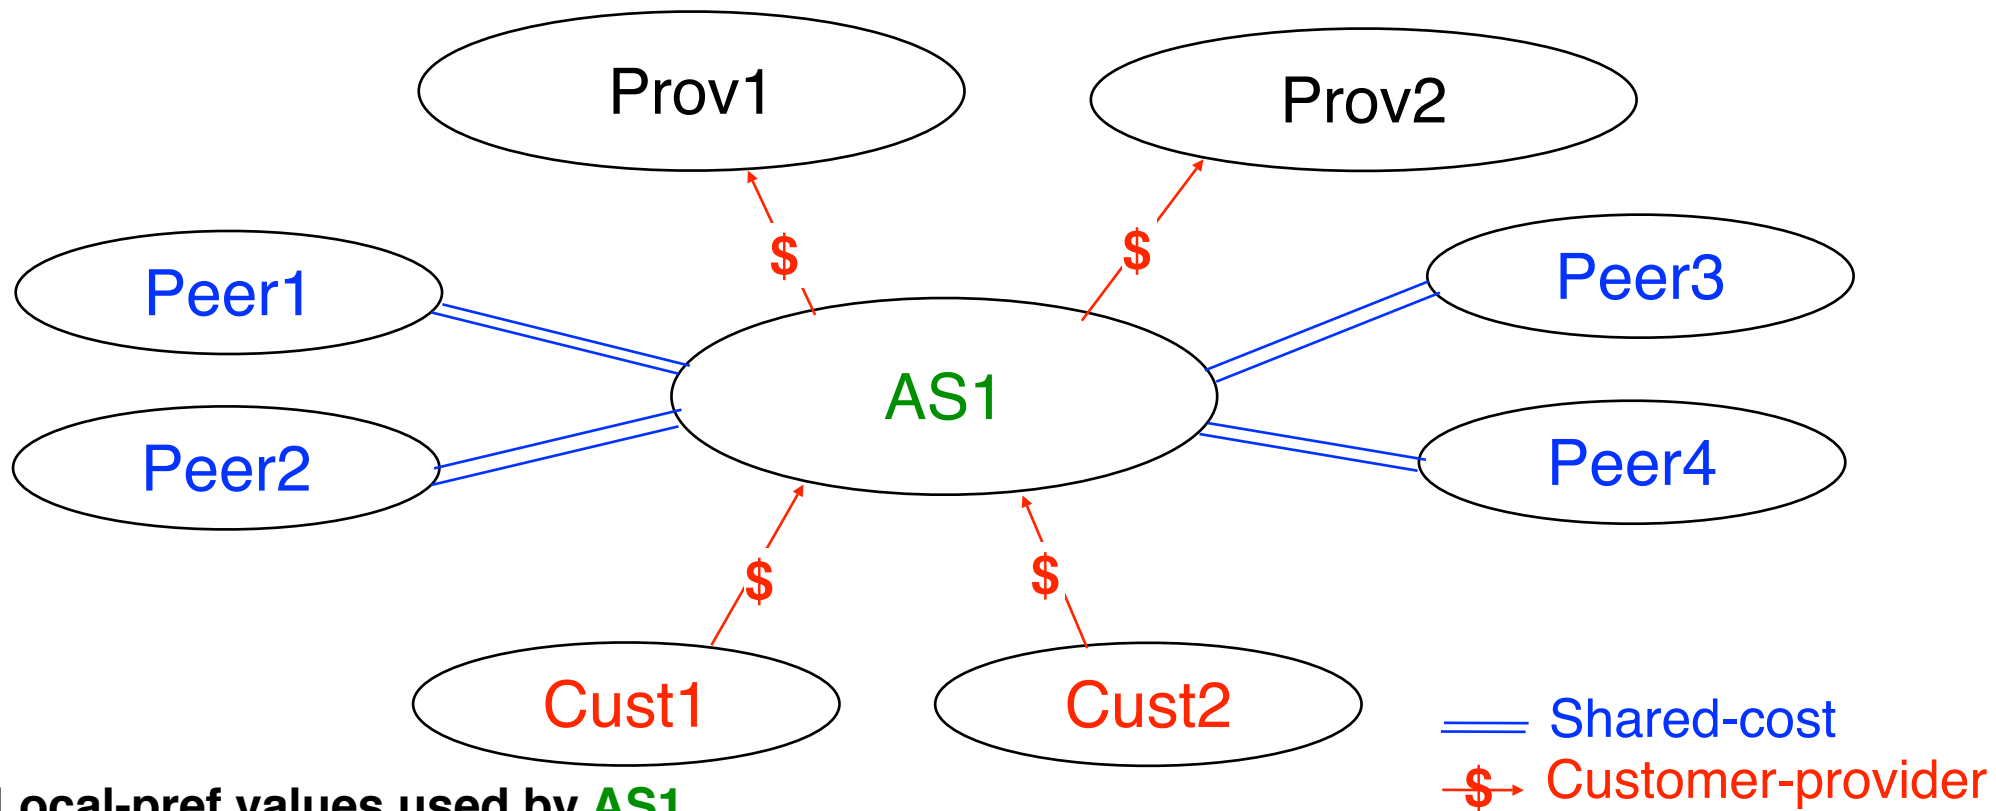
Unfortunately, interdomain routing may not converge at all in some cases...



How to reach a destination inside AS0 in this case ?

# local-pref and economical relationships

In practice, local-pref is often used to enforce economical relationships



## Local-pref values used by AS1


> 1000 for the routes received from a Customer

500 – 999 for the routes learned from a Peer

< 500 for the routes learned from a Provider

# Consequence of this utilisation of local-pref

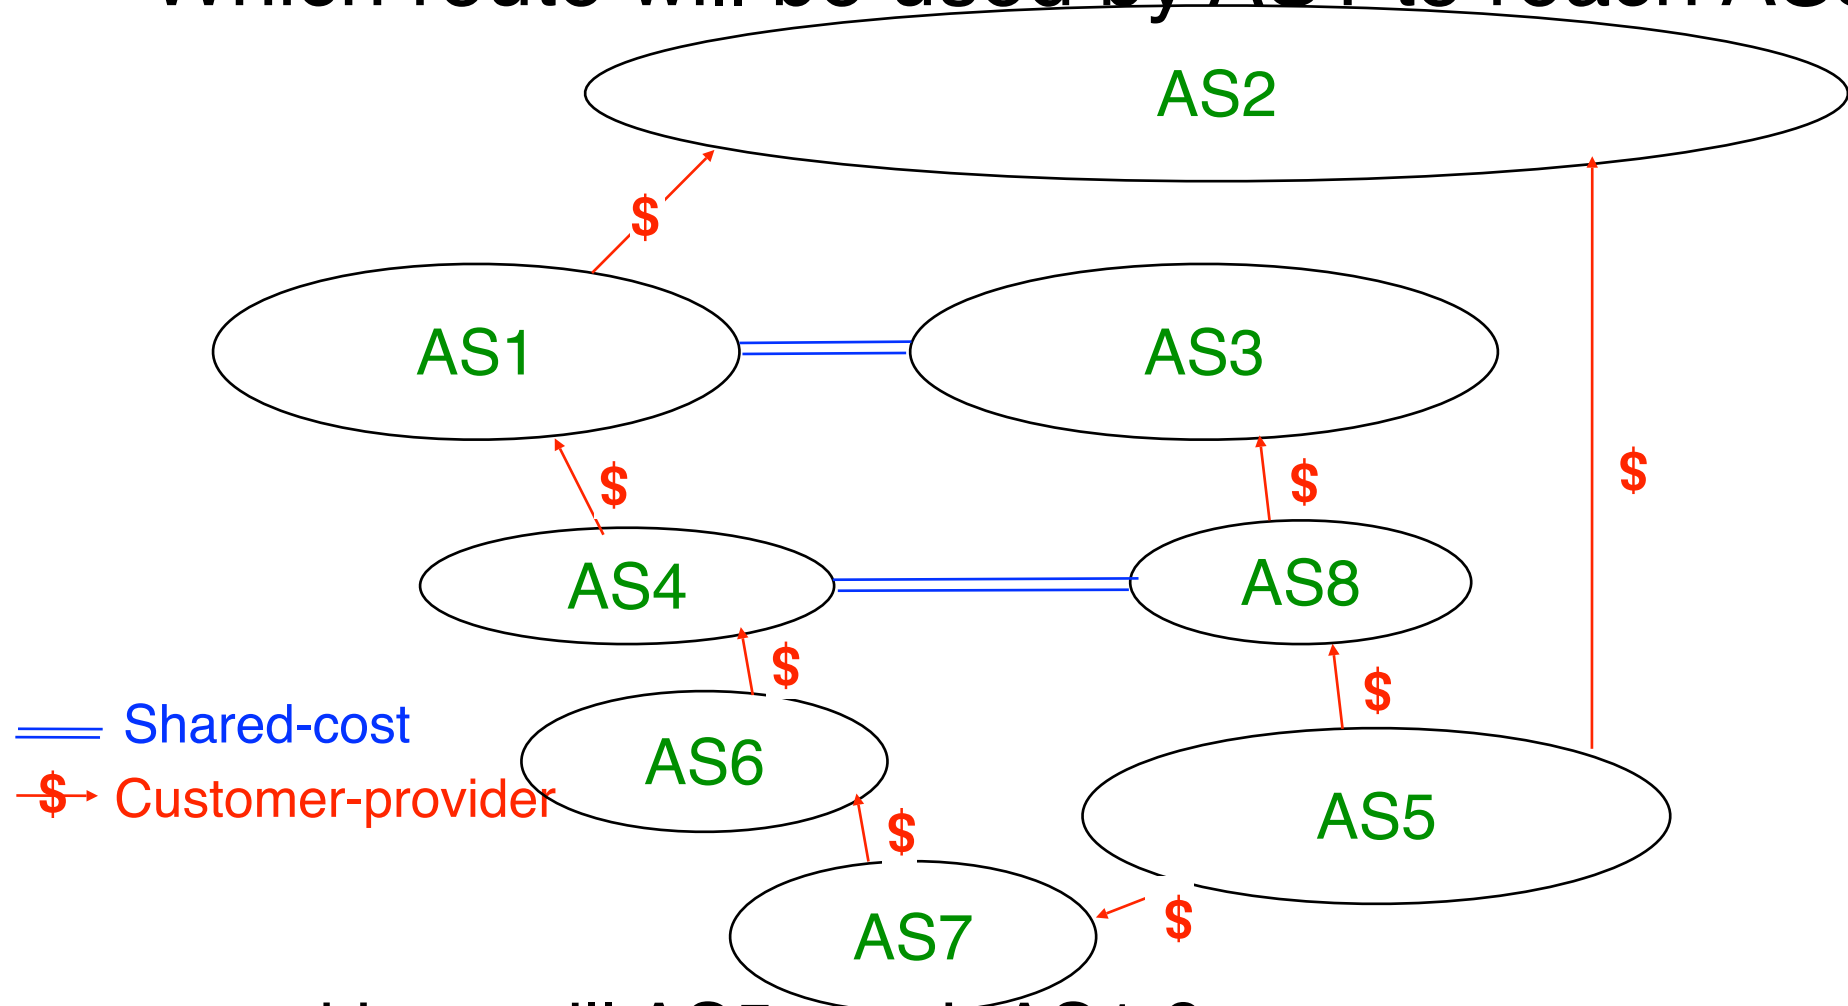
Which route will be used by AS1 to reach AS5 ?



and how will AS5 reach AS1 ?

# Consequence of this utilisation of local-pref


Which route will be used by AS1 to reach AS5 ?



and how will AS5 reach AS1 ?

# BGP and IP

## Second example

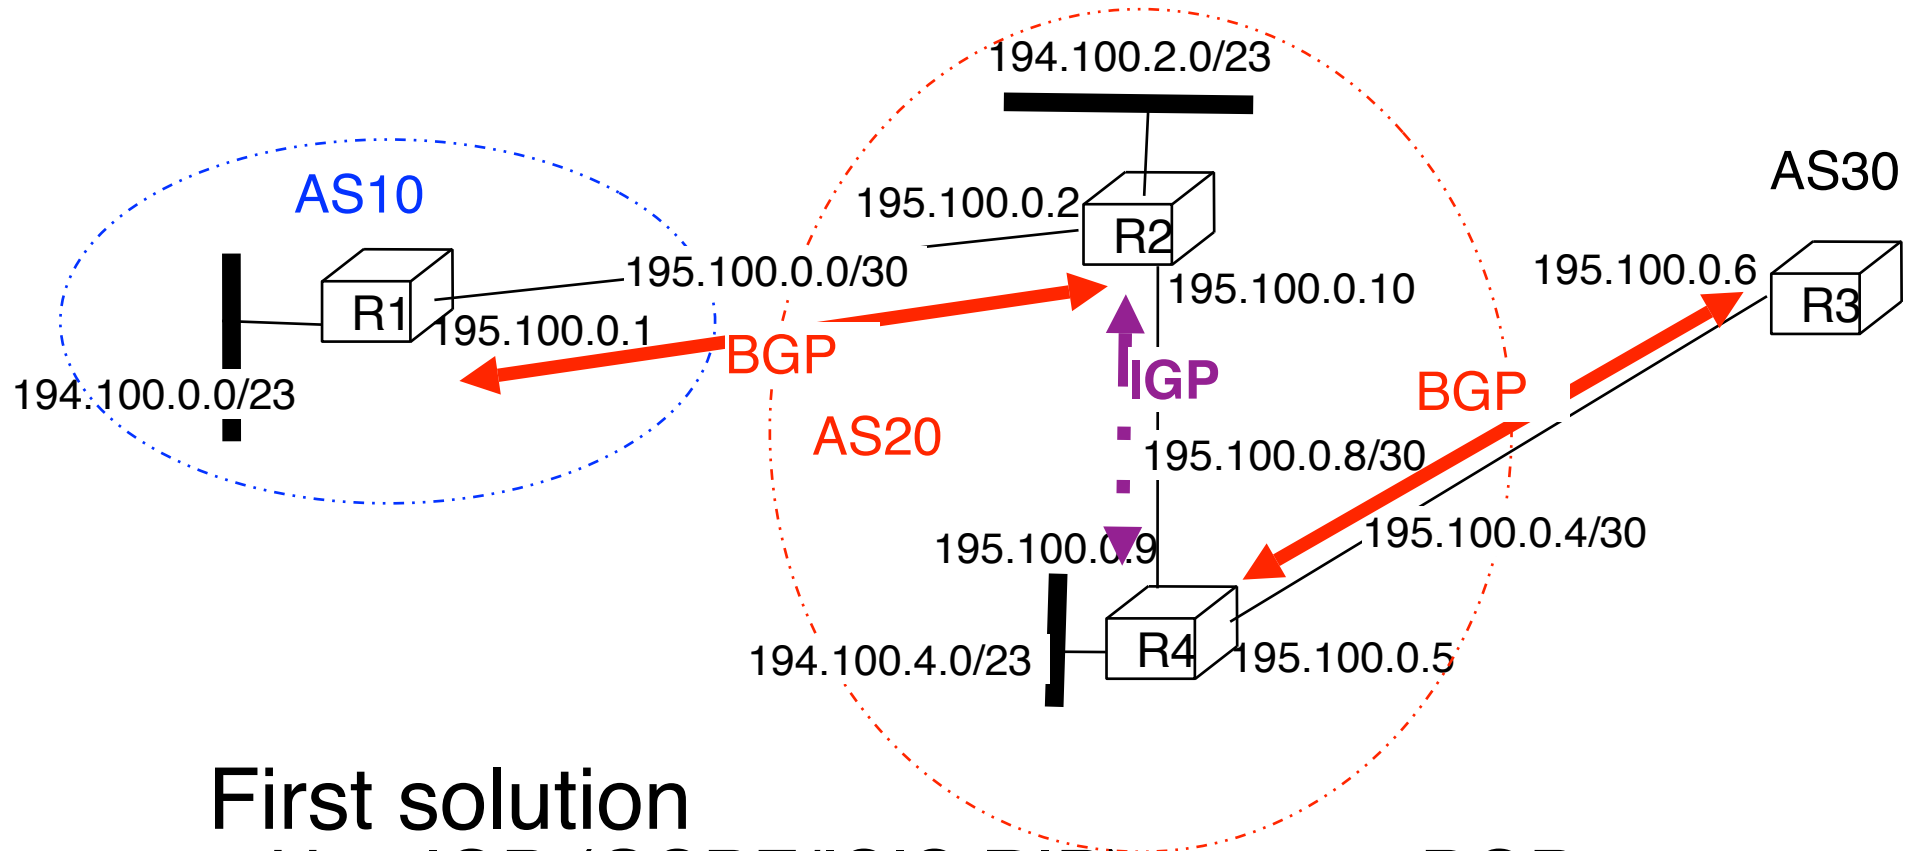


## Problem

How can R2 (resp. R4) advertise to R4 (resp. R2) the routes learned from AS10 (resp. AS30) ?

# BGP and IP

## Second example (2)



### First solution

Use IGP (OSPF/ISIS,RIP) to carry BGP routes

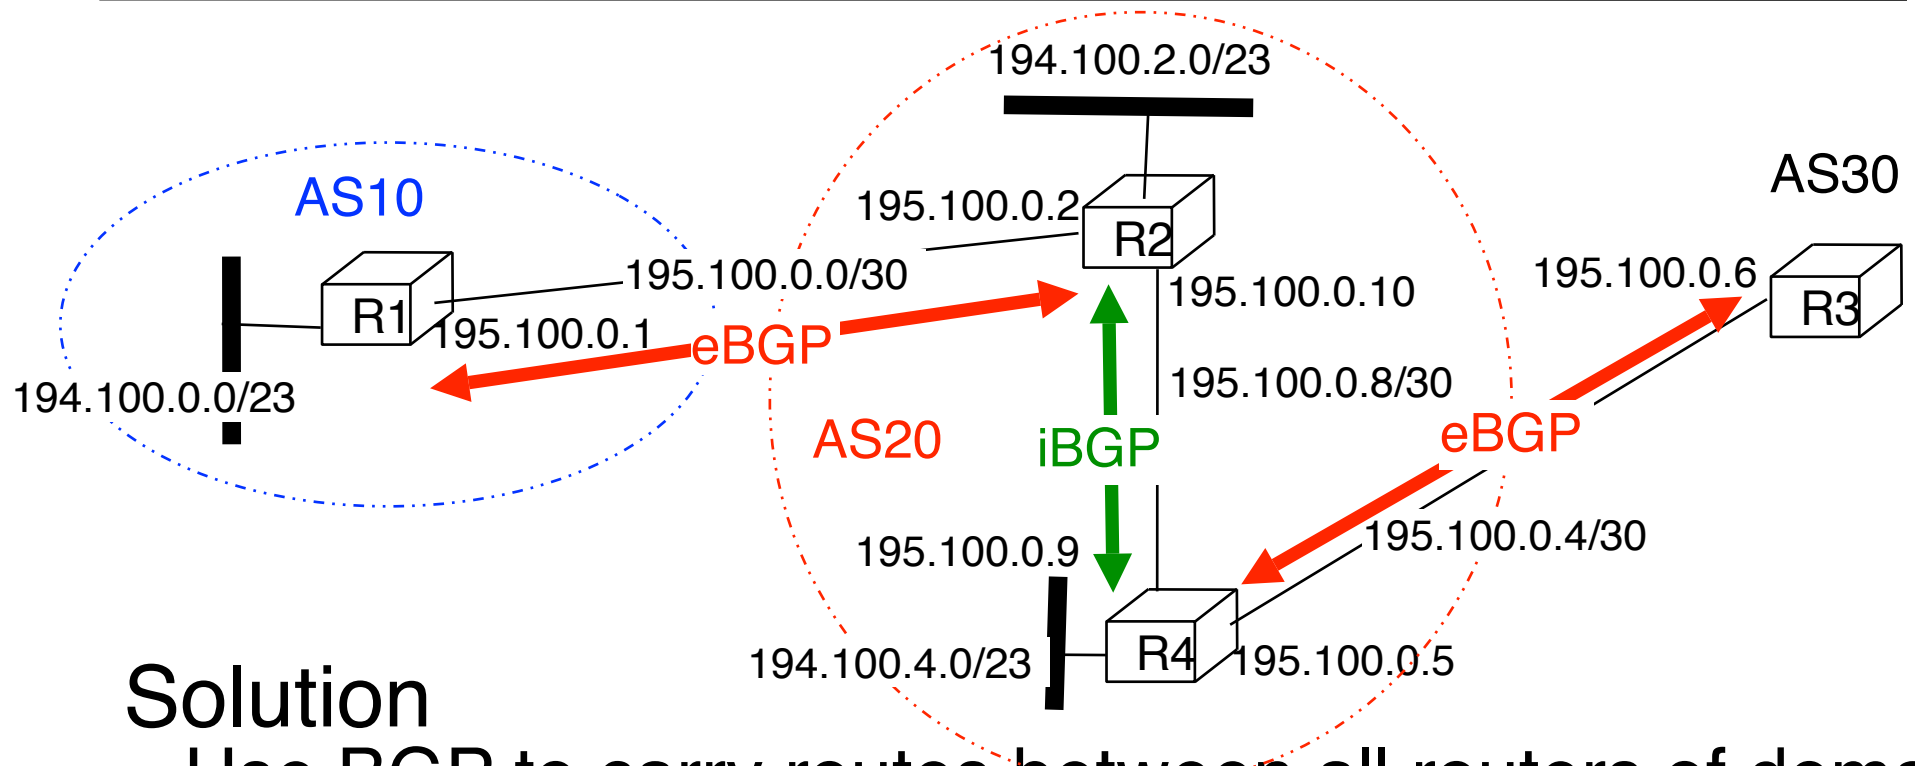
### Drawbacks

IGP may not be able to support so many routes

IGP does not carry BGP attributes like ASPath !



# iBGP and eBGP



## Solution

Use BGP to carry routes between all routers of domain

Two different types of BGP sessions

**eBGP** between routers belonging to different ASes

**iBGP** between each pair of routers belonging to the same AS

Each BGP router inside AS<sub>x</sub> maintains an **iBGP** session with all other BGP routers of AS<sub>x</sub> (full **iBGP** mesh)

Note that the iBGP sessions do not necessarily follow physical topology

# iBGP versus eBGP

---

## Differences between iBGP and eBGP

local-pref attribute is only carried inside messages sent over iBGP session

Over an eBGP session, a router only advertises its best route towards each destination

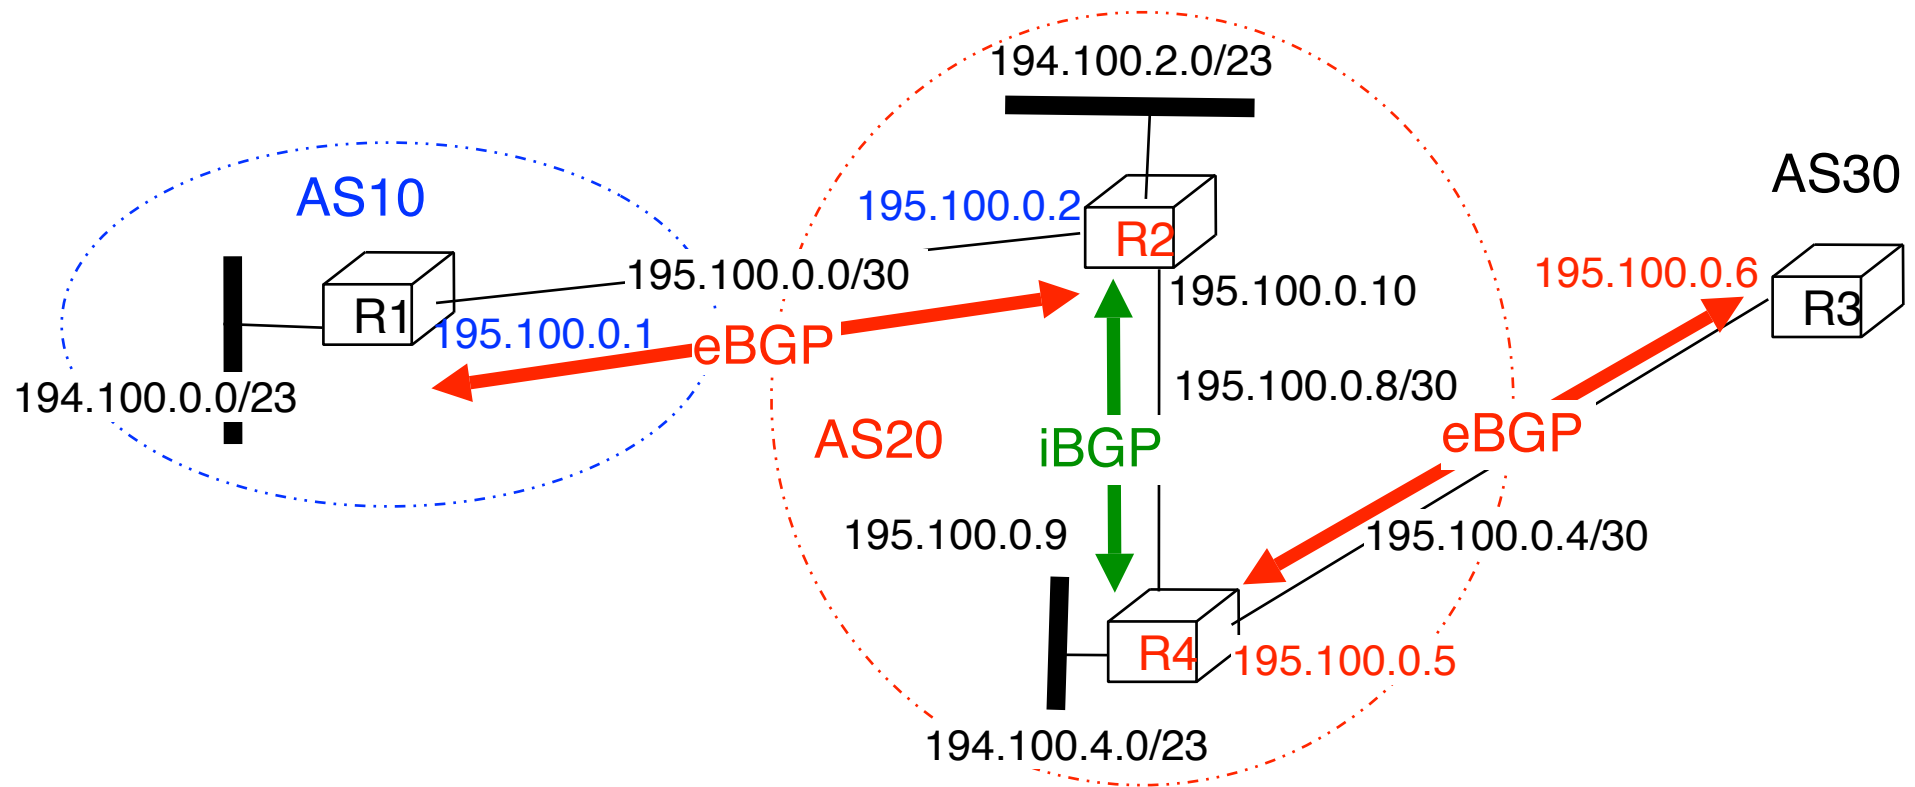
Usually, import and export filters are defined for each eBGP session

Over an iBGP session, a router advertises only its best routes learned over eBGP sessions

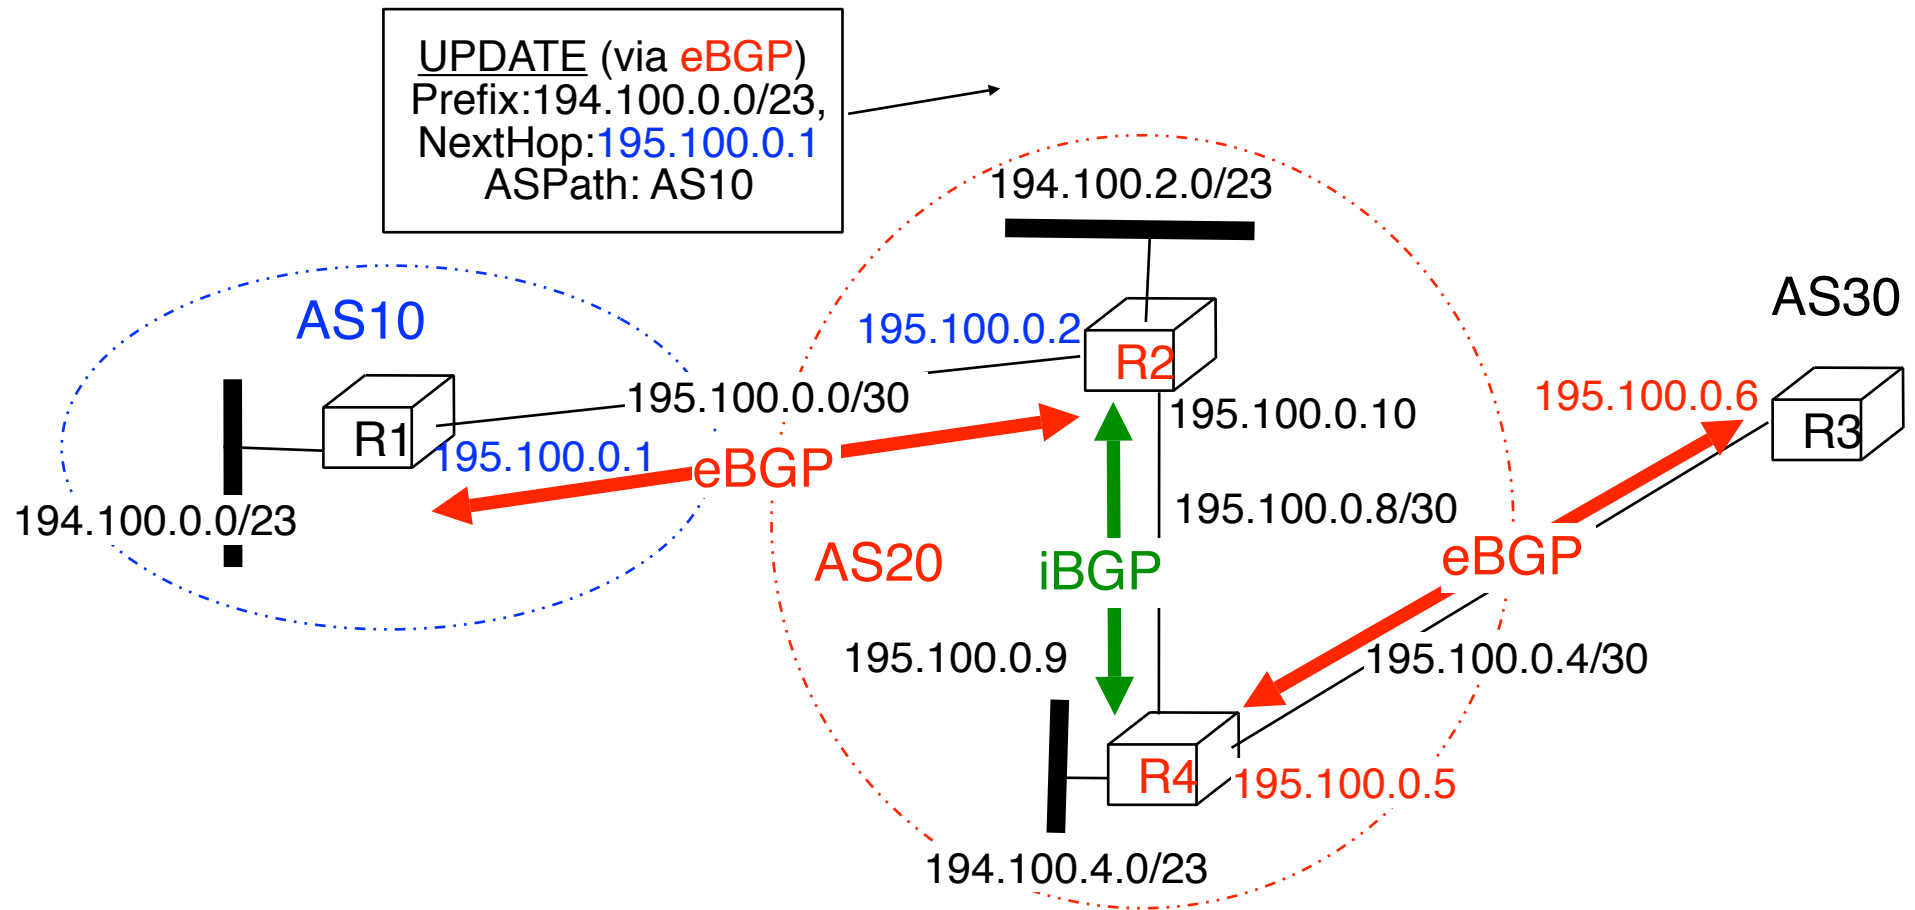
A route learned over an iBGP session is *never* advertised over another iBGP session

Usually, no filter is applied on iBGP sessions

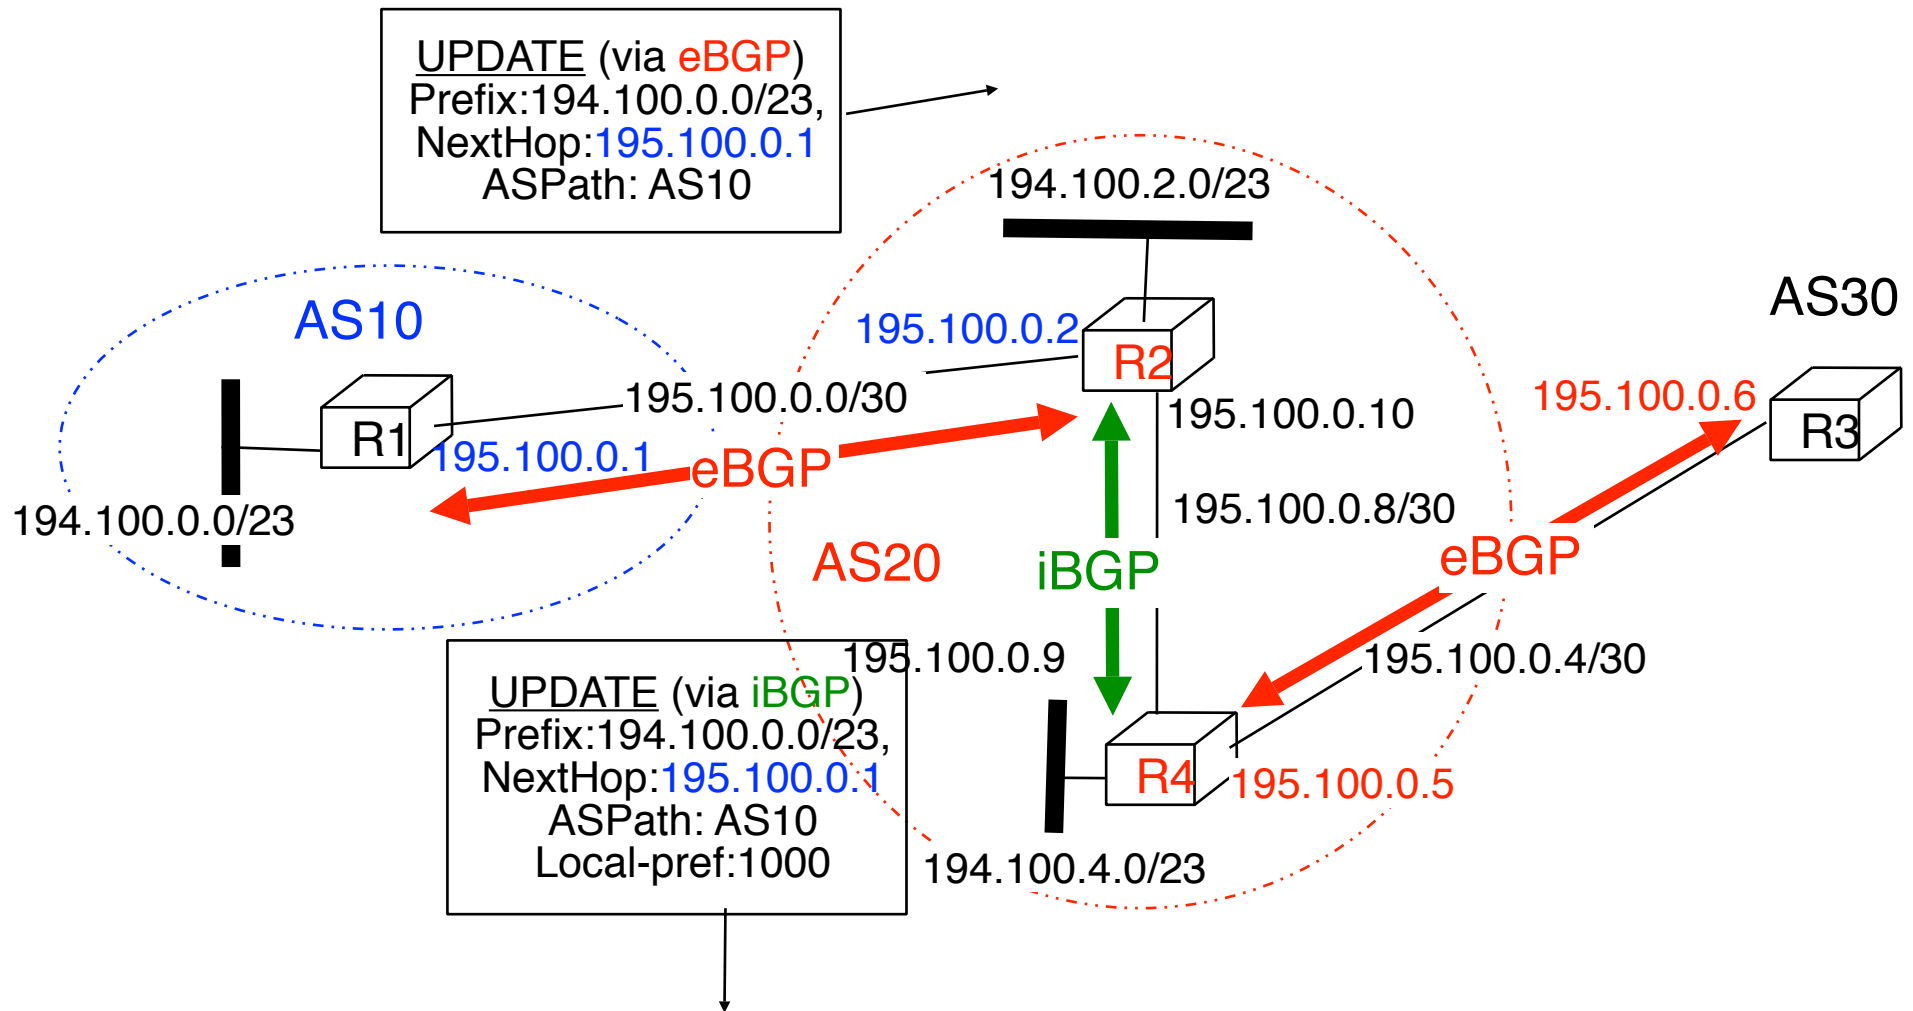
# iBGP and eBGP : Example



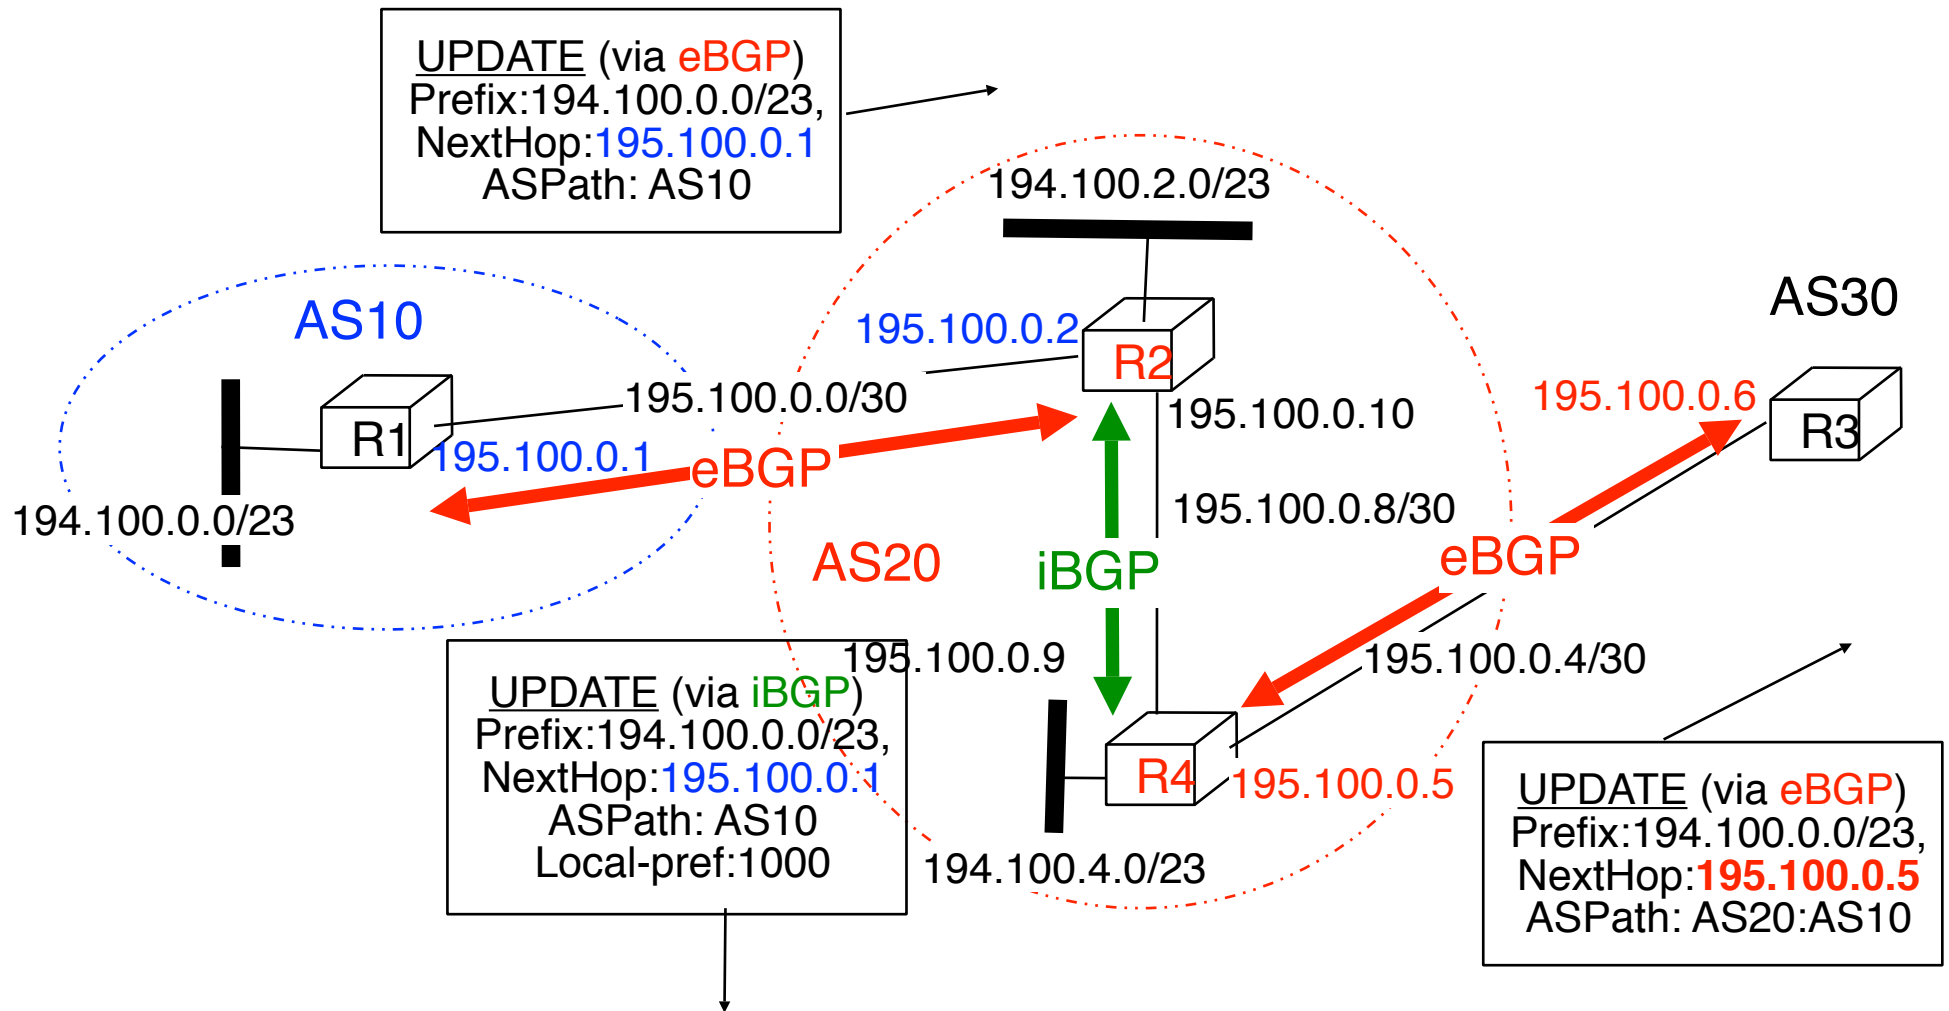
# iBGP and eBGP : Example



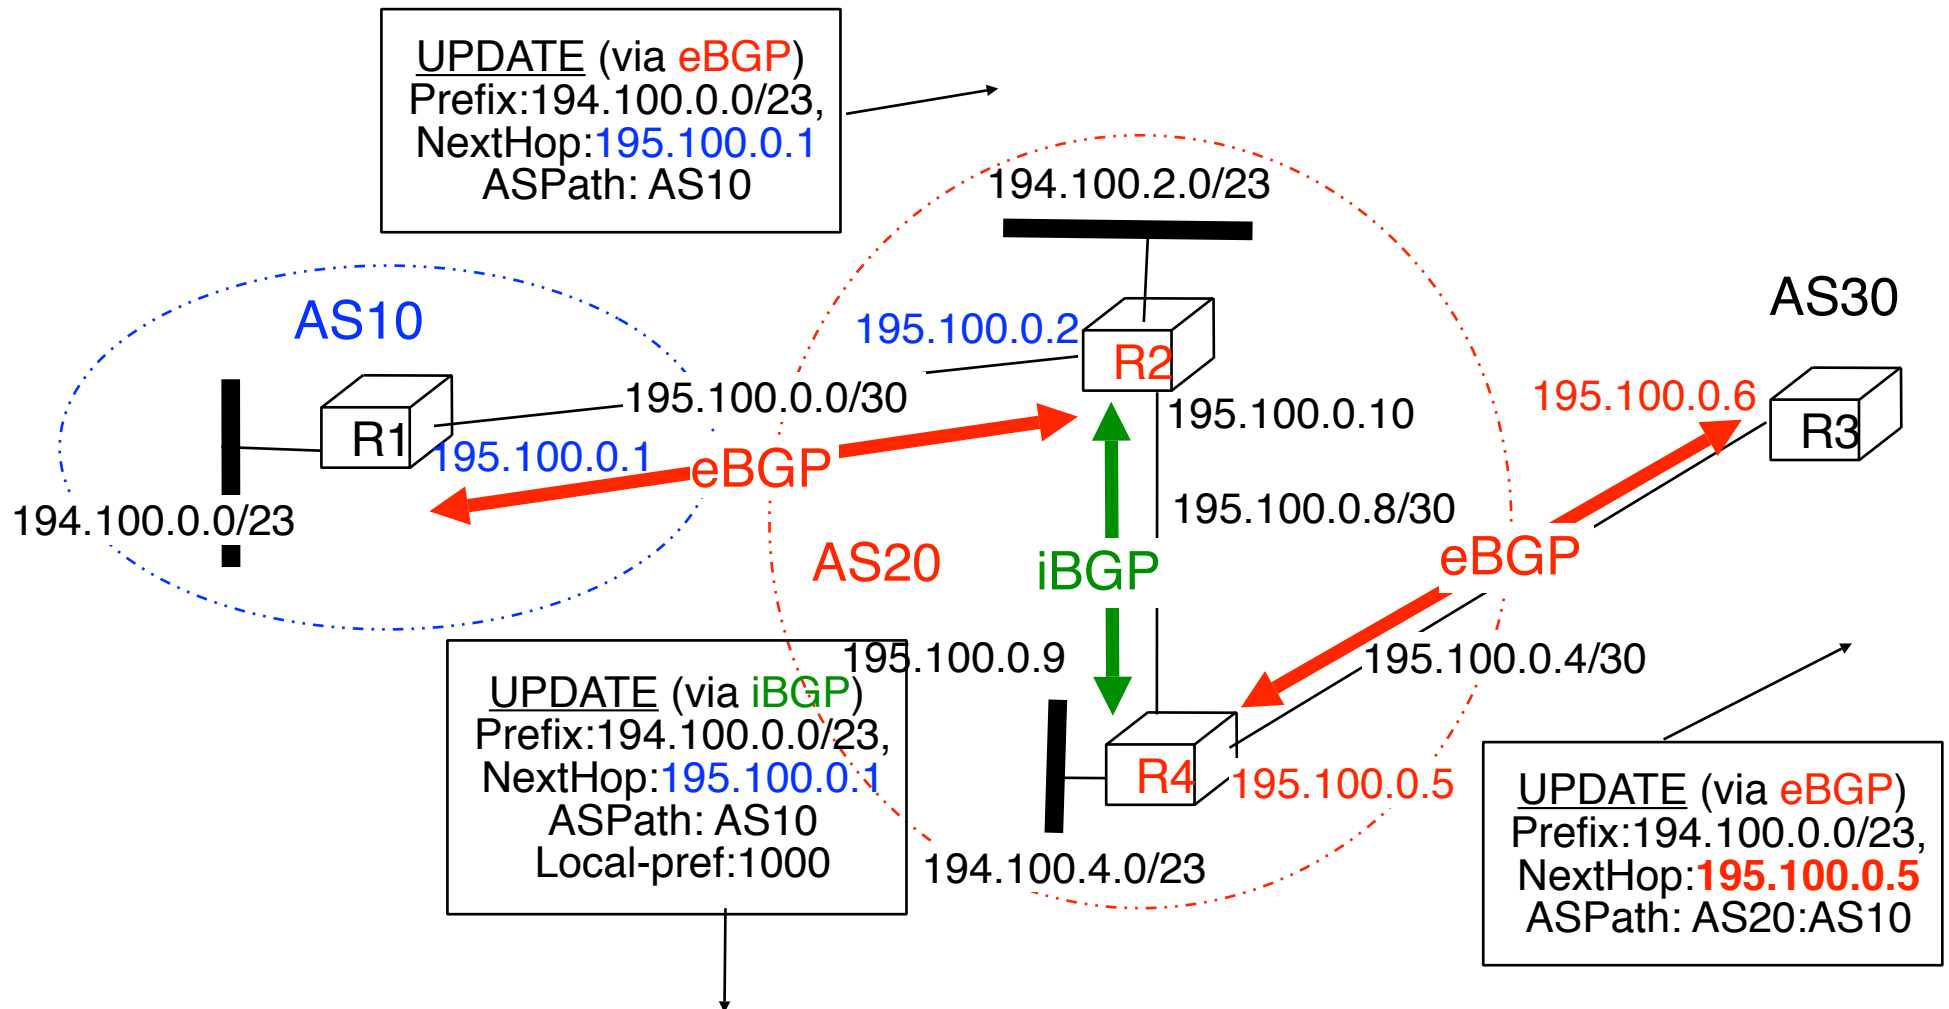
# iBGP and eBGP : Example



# iBGP and eBGP : Example

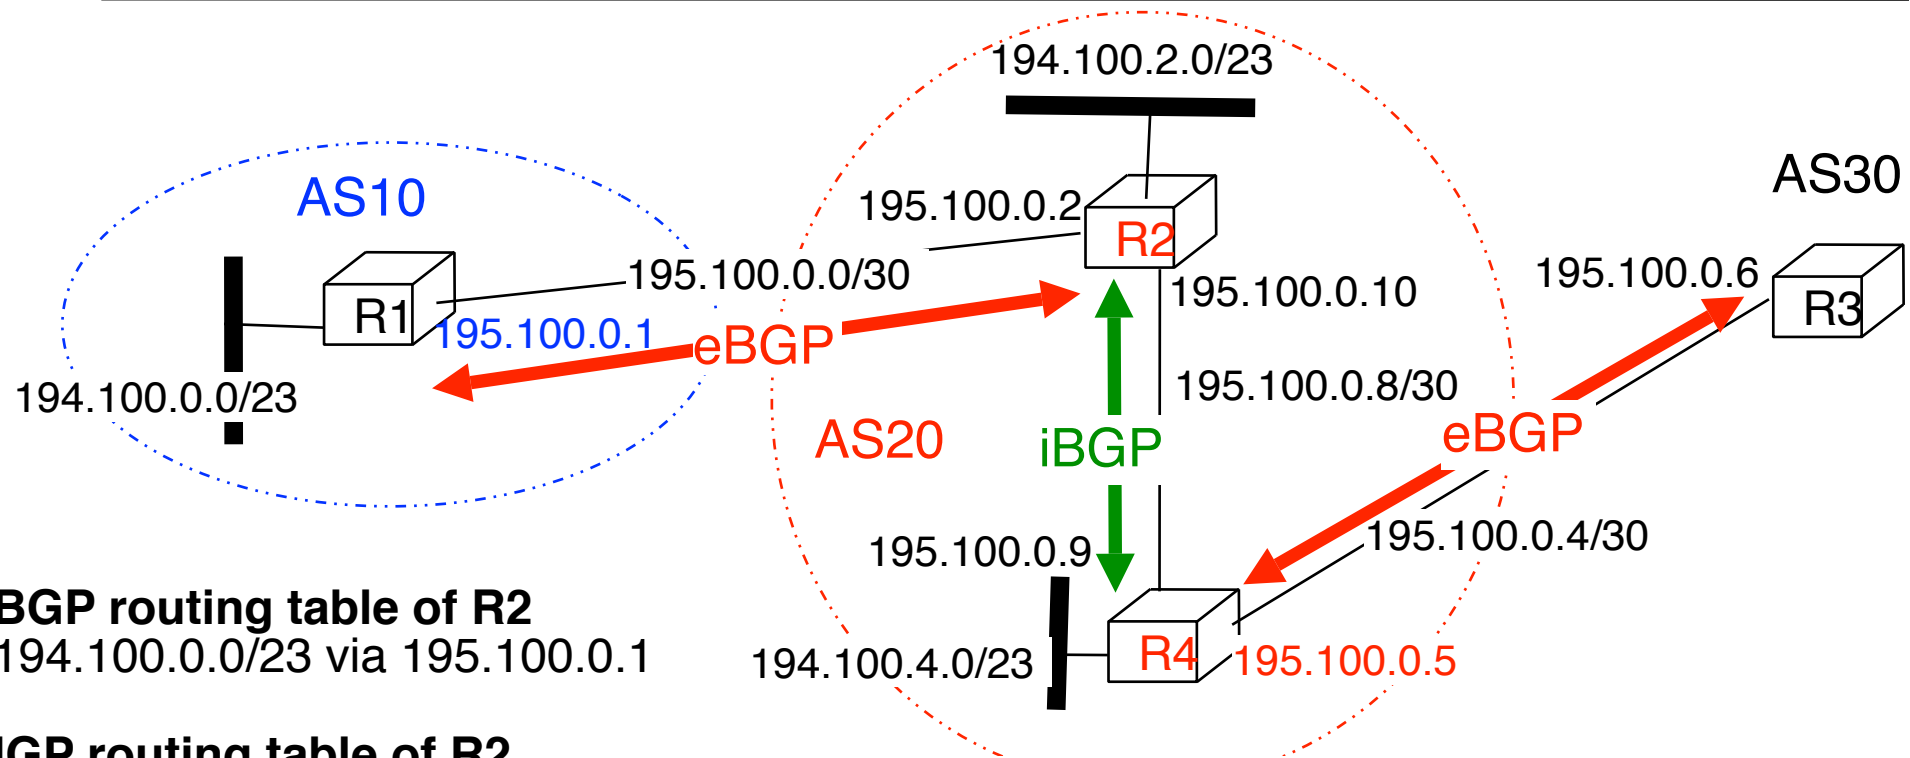


# iBGP and eBGP : Example



Note that the next-hop and the AS-Path of BGP update messages are only updated when sent over an eBGP session

# iBGP and eBGP Packet Forwarding



## BGP routing table of R2

194.100.0.0/23 via 195.100.0.1

## IGP routing table of R2

195.100.0.0/30 West  
195.100.0.4/30 via 195.100.0.9  
195.100.0.8/30 South  
194.100.0.4/23 via 195.100.0.9  
194.100.2.0/23 North

## BGP routing table of R4


194.100.0.0/23 via 195.100.0.1

## IGP routing table of R4

195.100.0.0/30 via 195.100.0.10  
195.100.0.4/30 East  
195.100.0.8/30 North  
194.100.2.0/23 via 195.100.0.10  
194.100.0.4/23 West



# iBGP and eBGP Packet Forwarding (2)



## BGP routing table of R4

194.100.0.0/23 via 195.100.0.1

## IGP routing table of R4

195.100.0.0/30 via 195.100.0.10

195.100.0.4/30 East

195.100.0.8/30 North

194.100.2.0/23 via 195.100.0.10

194.100.4.0/23 West

## Forwarding of R4

194.100.0.0/23 via 195.100.0.10

195.100.0.0/30 via 195.100.0.10

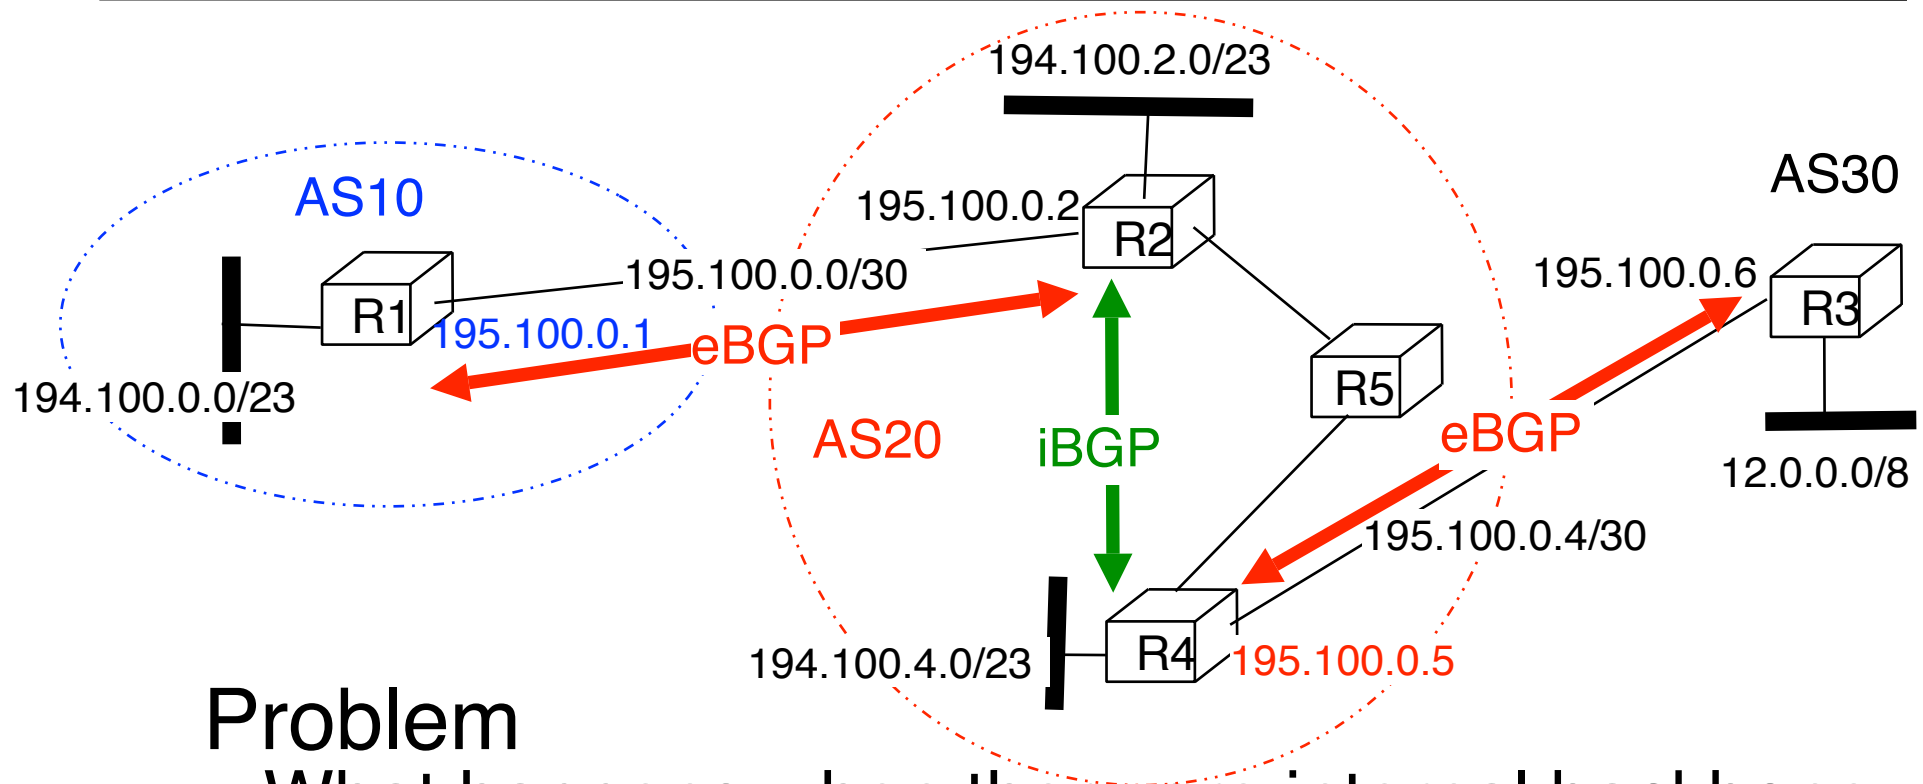
195.100.0.4/30 East

195.100.0.8/30 North

194.100.2.0/23 via 195.100.0.10

194.100.4.0/23 West

# Using non-BGP routers

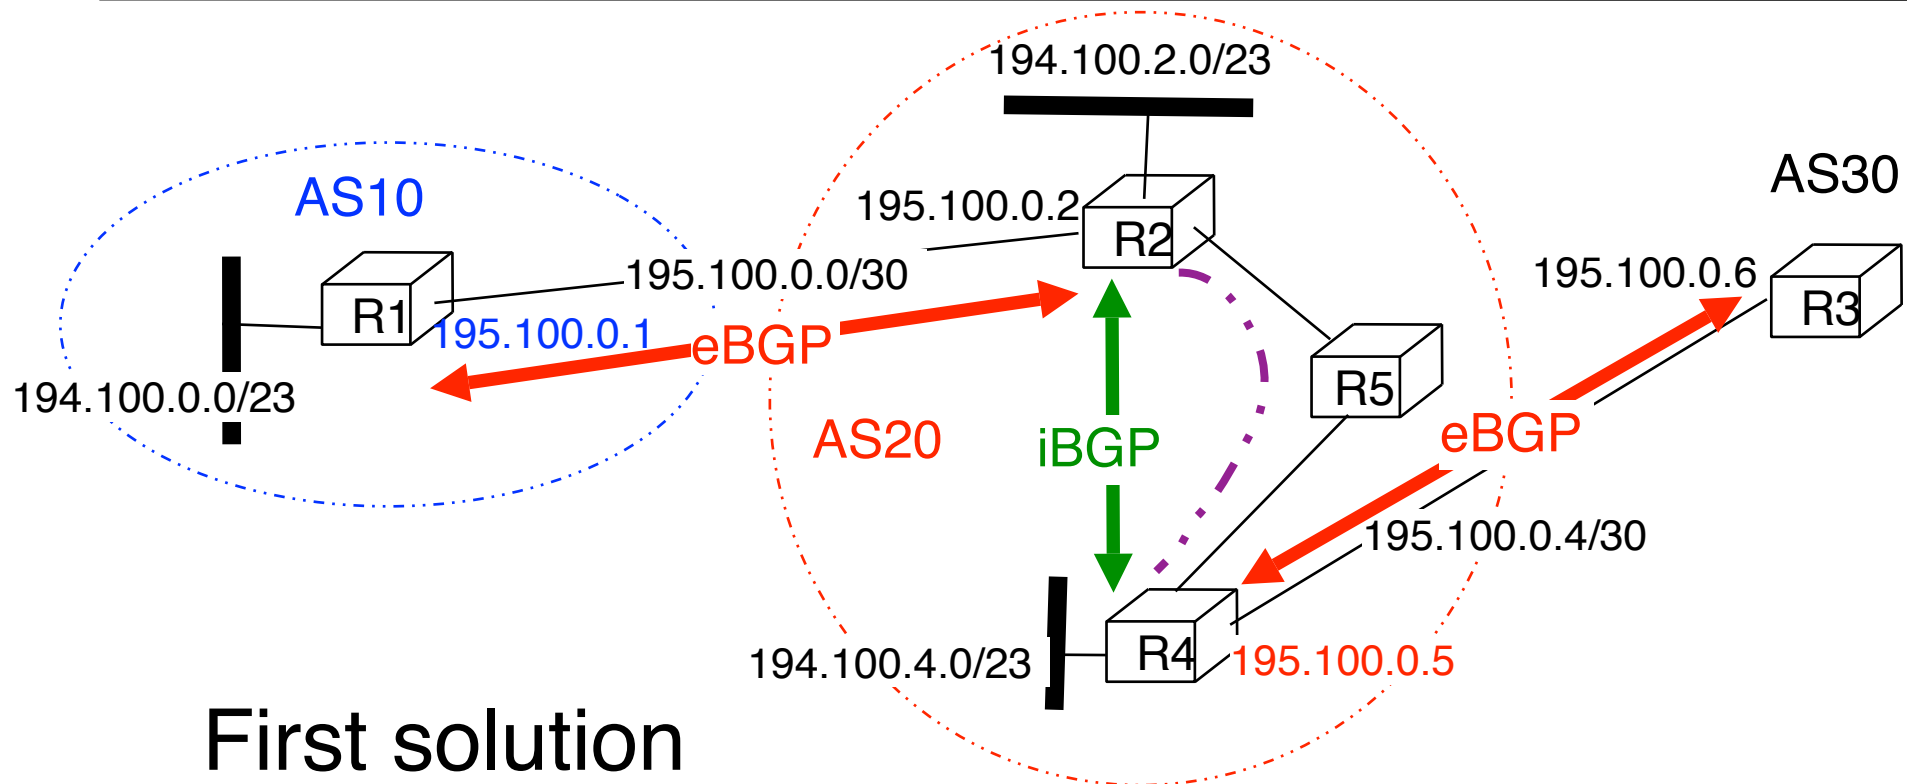


## Problem

What happens when there are internal backbone routers between BGP routers inside an AS ?

iBGP session between BGP routers is easily established when IGP is running since iBGP runs over TCP connection  
How to populate the routing table of the backbone routers to ensure that they will be able to route any IP packet ?

# Using non-BGP routers (2)



## First solution

Use tunnels between BGP routers to encapsulate interdomain packets

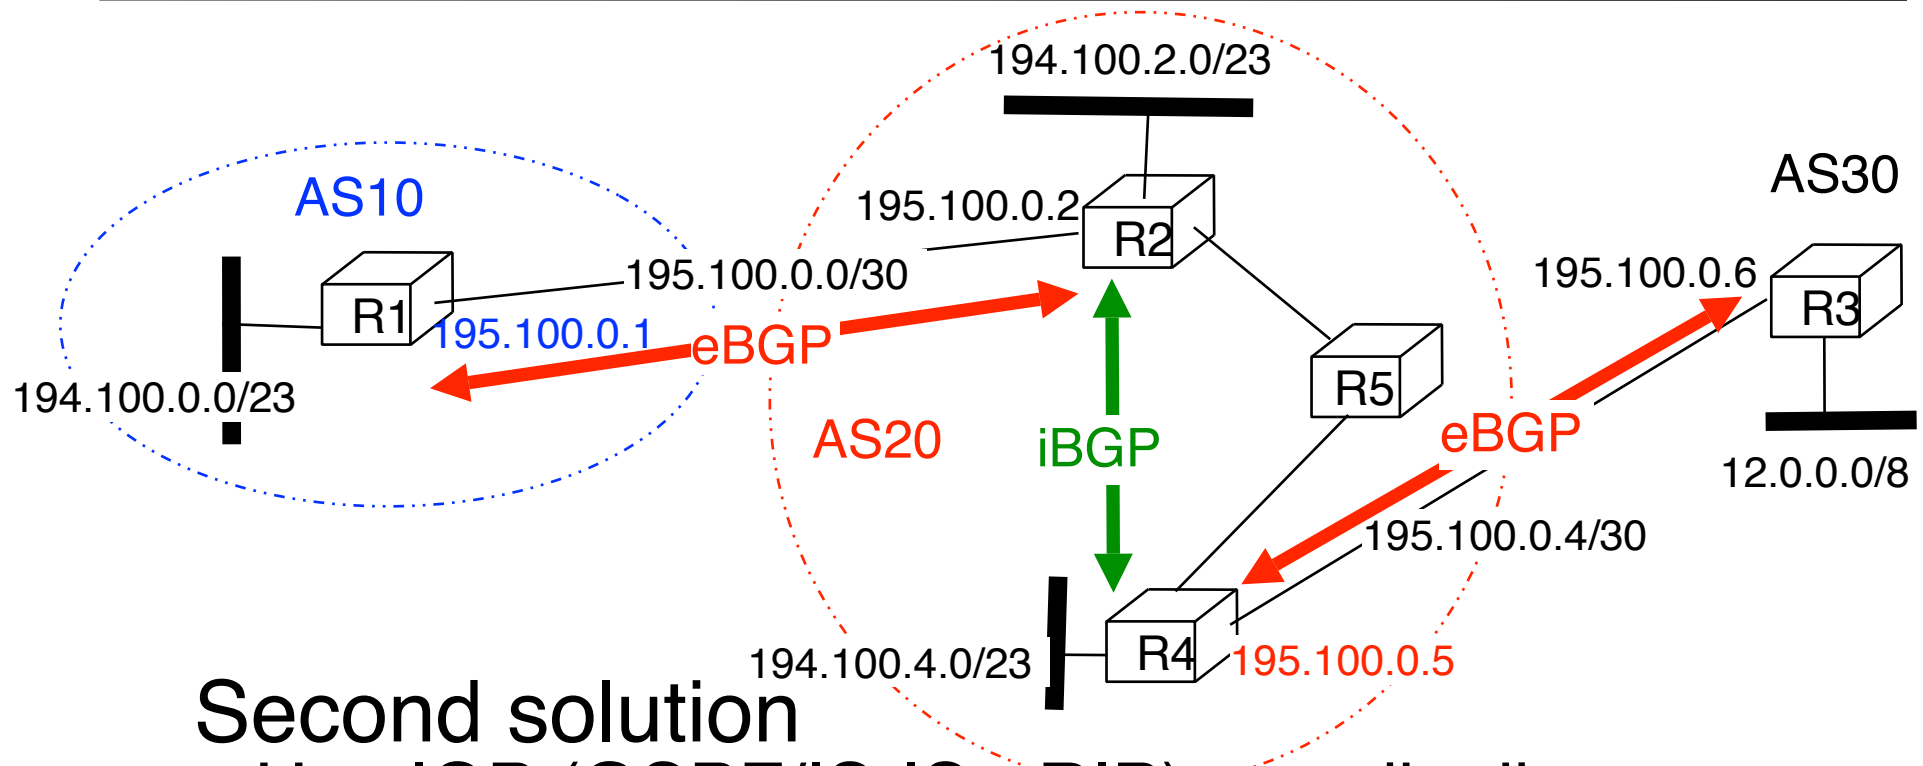
**GRE tunnel**

Needs static configuration and be careful with MTU issues

**MPLS tunnel**

Can be dynamically established in MPLS enabled backbone

# Using non-BGP routers (3)



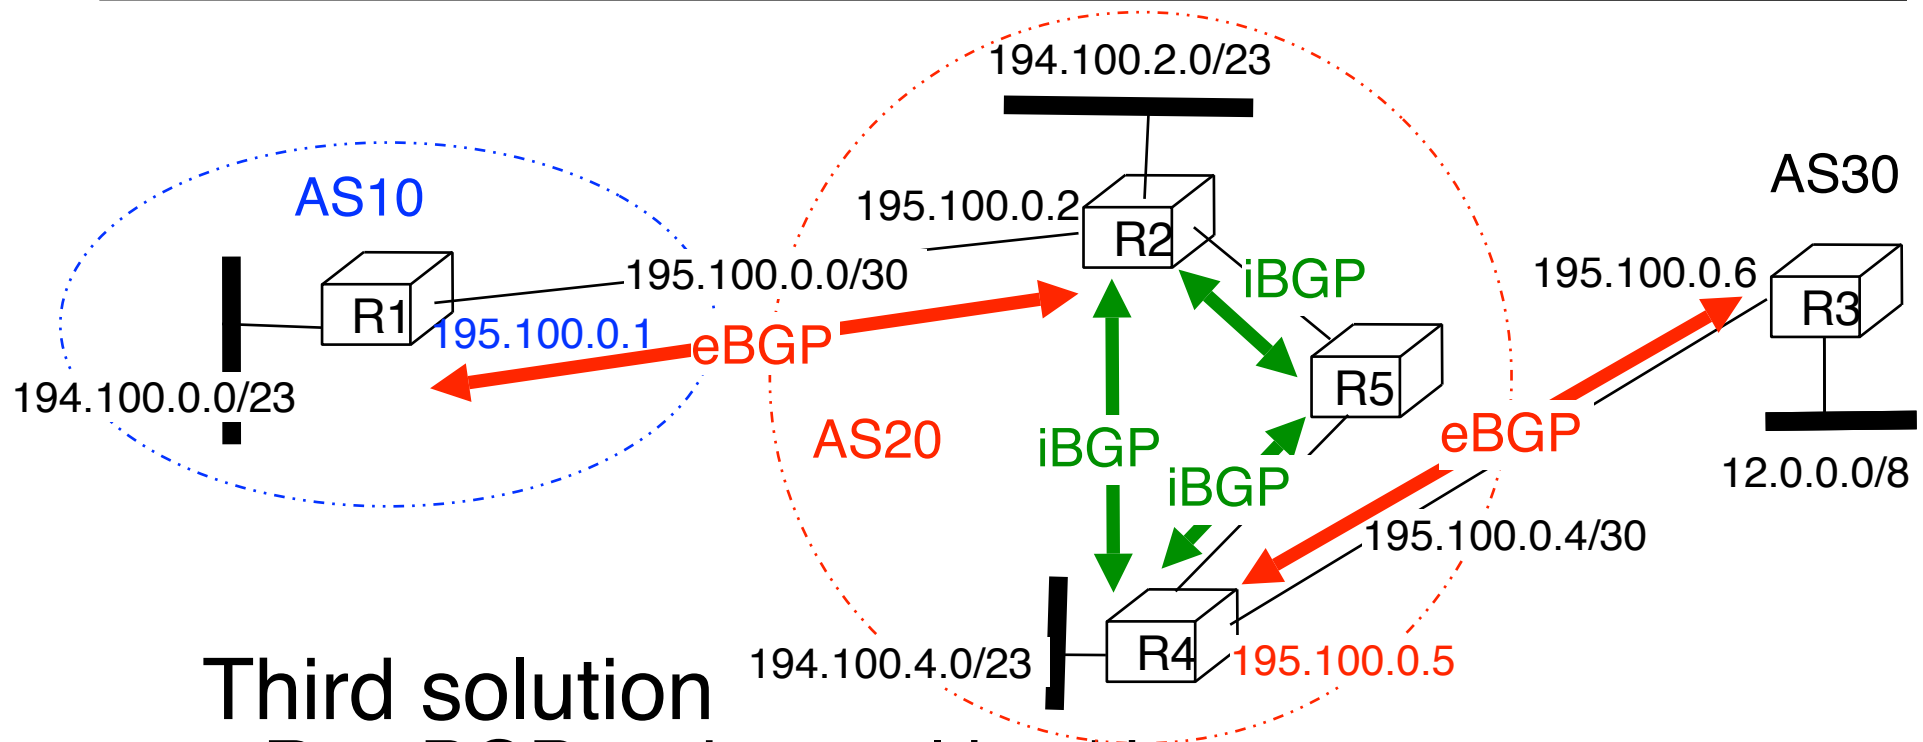
## Second solution

Use IGP (OSPF/IS-IS - RIP) to redistribute interdomain routes to internal backbone routers

### Drawbacks

Size of BGP tables may completely overload the IGP  
Make sure that BGP routes learned by R2 and injected inside IGP will not be re-injected inside BGP by R4 !

# Using non-BGP routers (4)



## Third solution

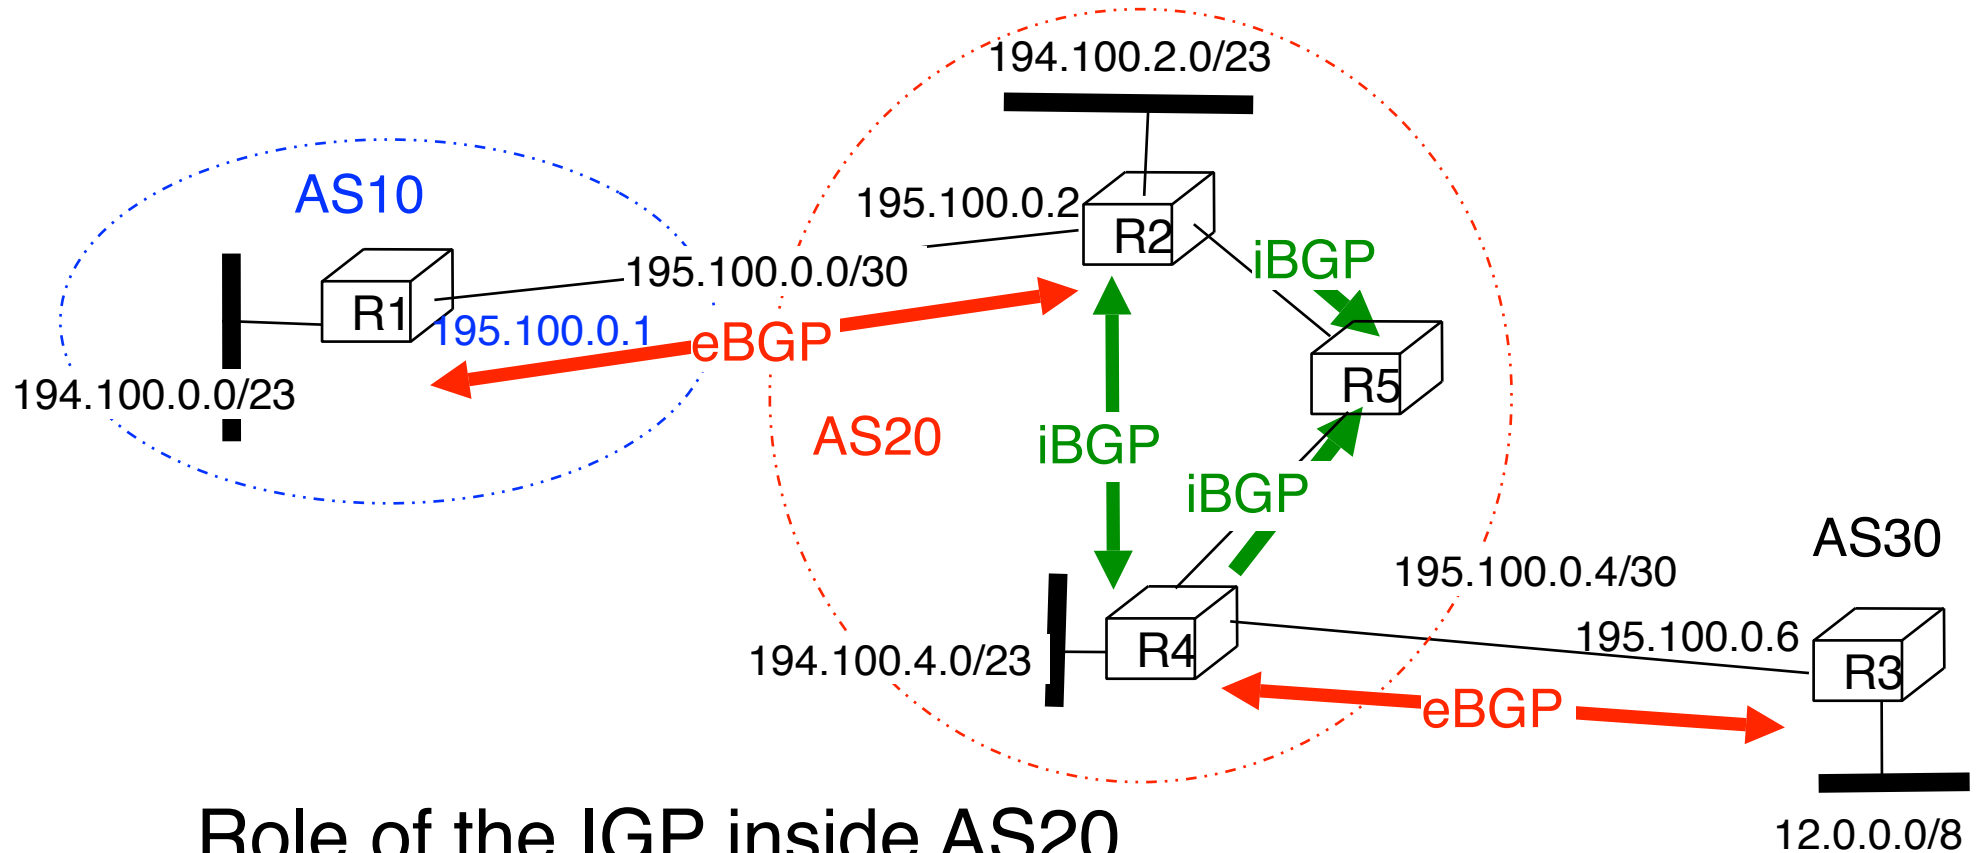
Run BGP on internal backbone routers

Internal backbone routers need to participate in iBGP full mesh

Internal backbone routers receive BGP routes via iBGP but never advertise any routes

Remember : a route learned over an iBGP session is never advertised over another iBGP session

# The roles of IGP and BGP



## Role of the IGP inside AS20

Distribute internal topology and internal addresses  
(R2-R4-R5)

## Role of BGP inside AS20


Distribute the routes towards external destinations

IGP must run to allow BGP routers to establish iBGP sessions


# The iBGP full mesh

## Drawback

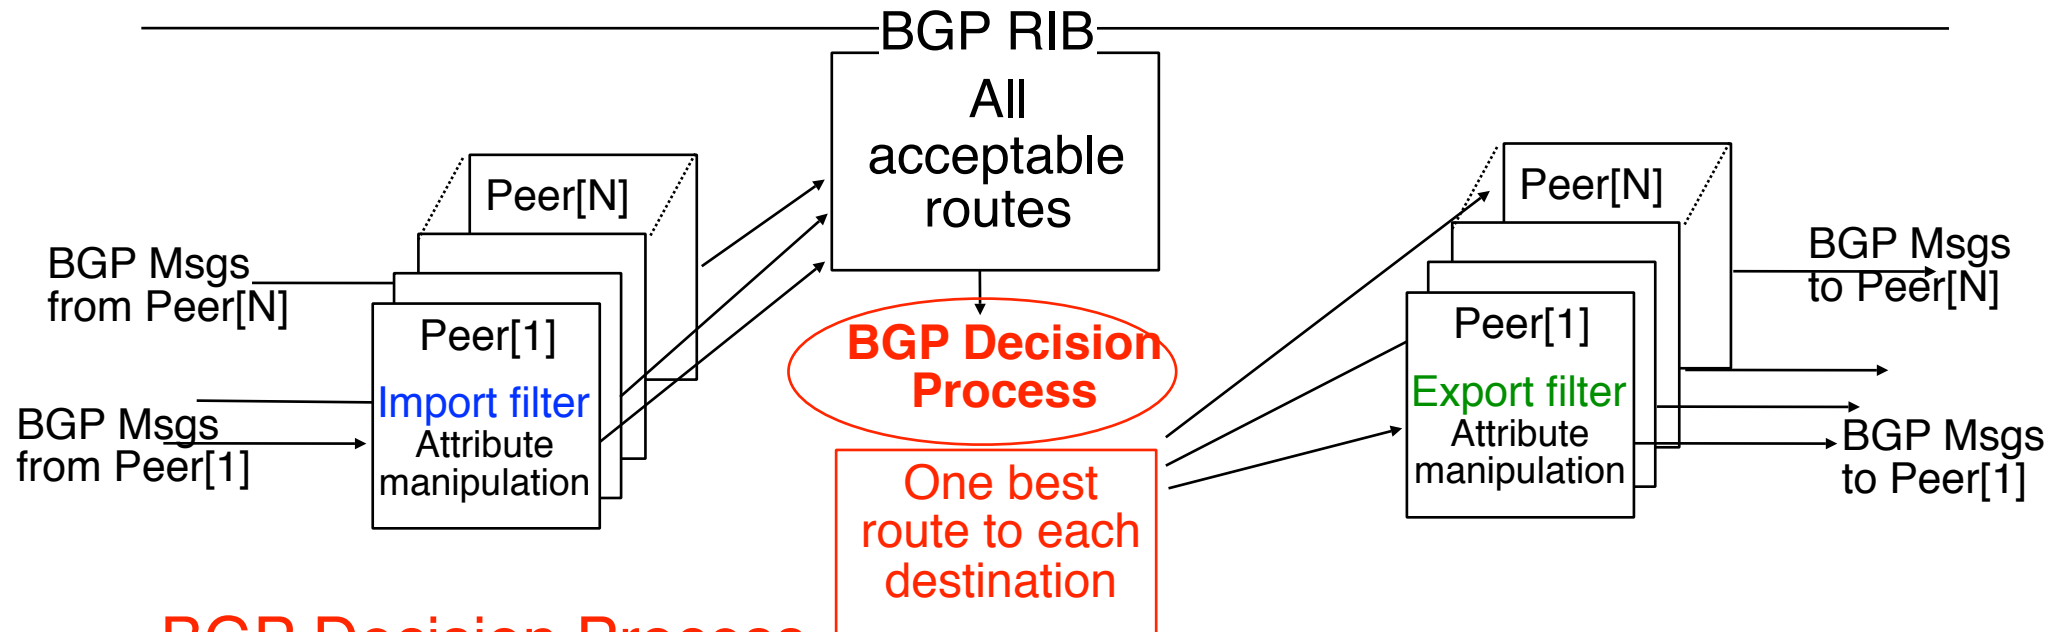
$N*(N-1)/2$  iBGP sessions for N routers



↔ iBGP session



# The BGP decision process



## BGP Decision Process

*Ignore routes with unreachable nexthop*

Prefer routes with highest local-pref

Prefer routes with shortest ASPath

Prefer routes with smallest MED

Prefer routes learned via eBGP over routes learned via iBGP

Prefer routes with closest next-hop

Tie breaking rules

Prefer Routes learned from router with lowest router id



# The shortest AS-Path step in the BGP decision process

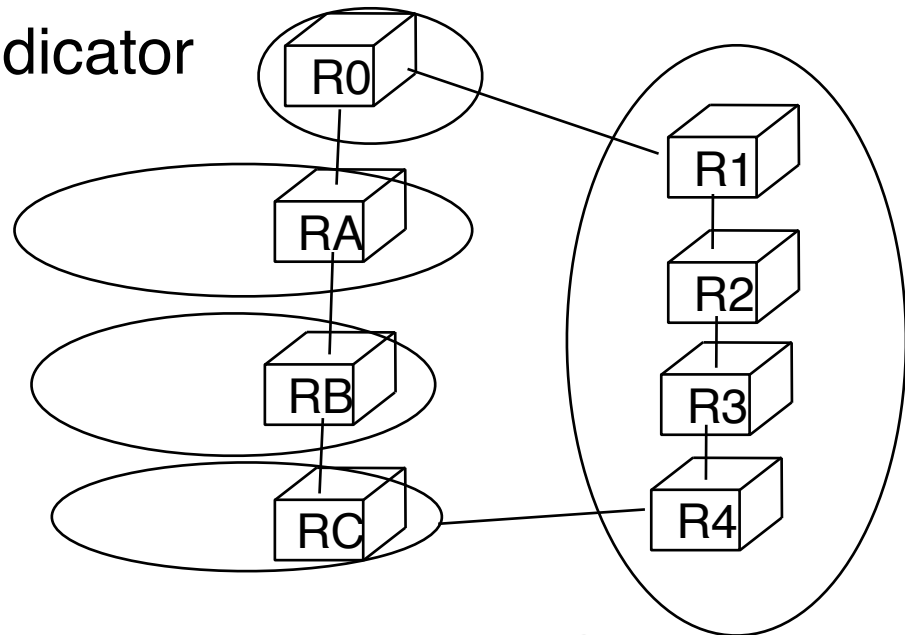
---

## Motivation

BGP does not contain a real “metric”

Use length of AS-Path as an indication of the quality of routes

Not always a good indicator



## Consequence

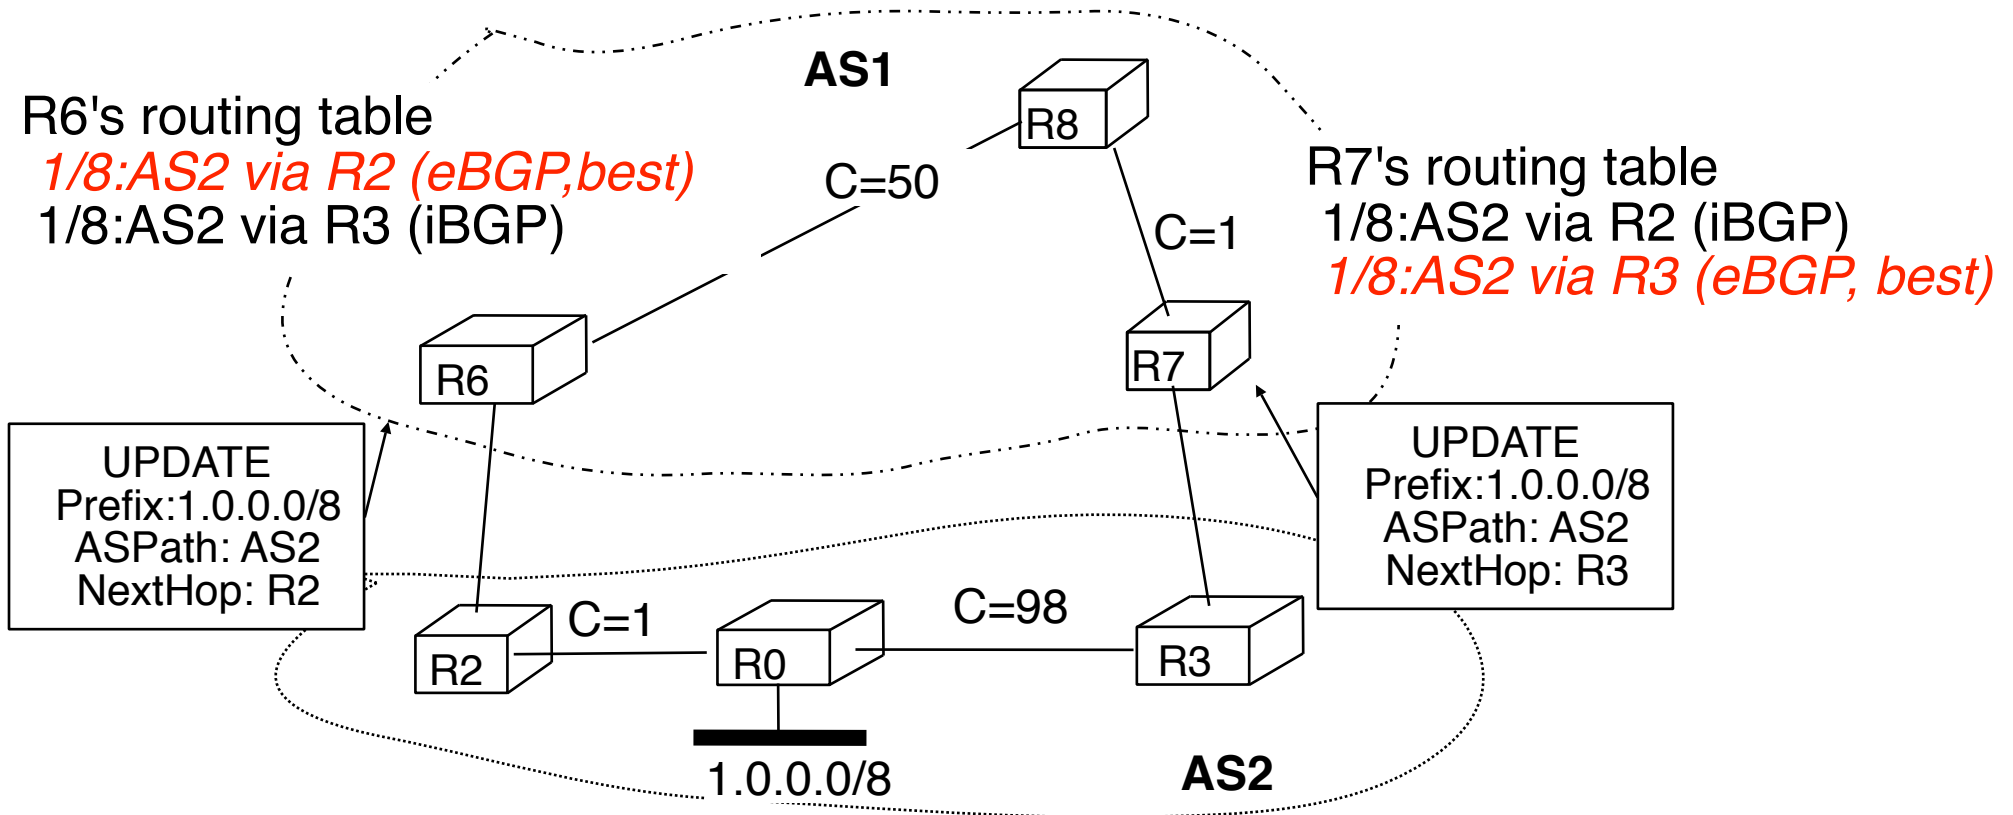
Internet paths tend to be short, 3-5 AS hops

Many paths converge at Tier-1 ISPs and those ISPs carry lots of traffic

# The prefer eBGP over iBGP step in the BGP decision process

## Motivation : hot potato routing

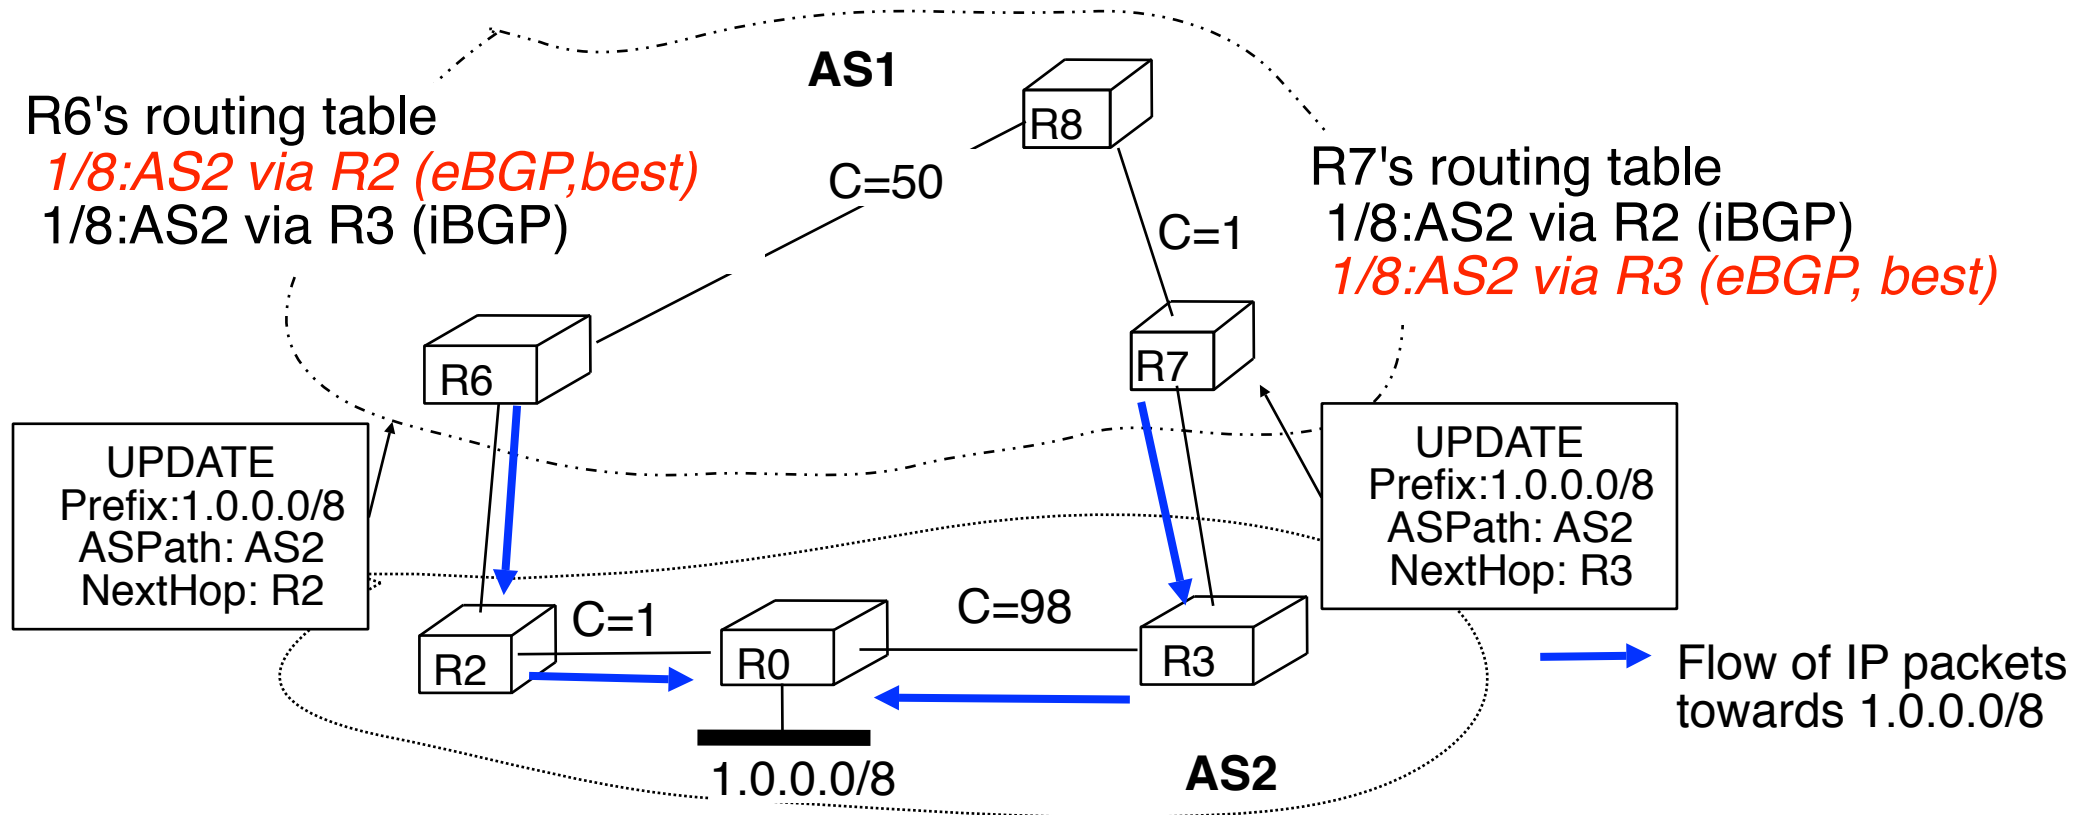
A router should try to get rid of packets sent to external domains as soon as possible



# The prefer eBGP over iBGP step in the BGP decision process

## Motivation : hot potato routing

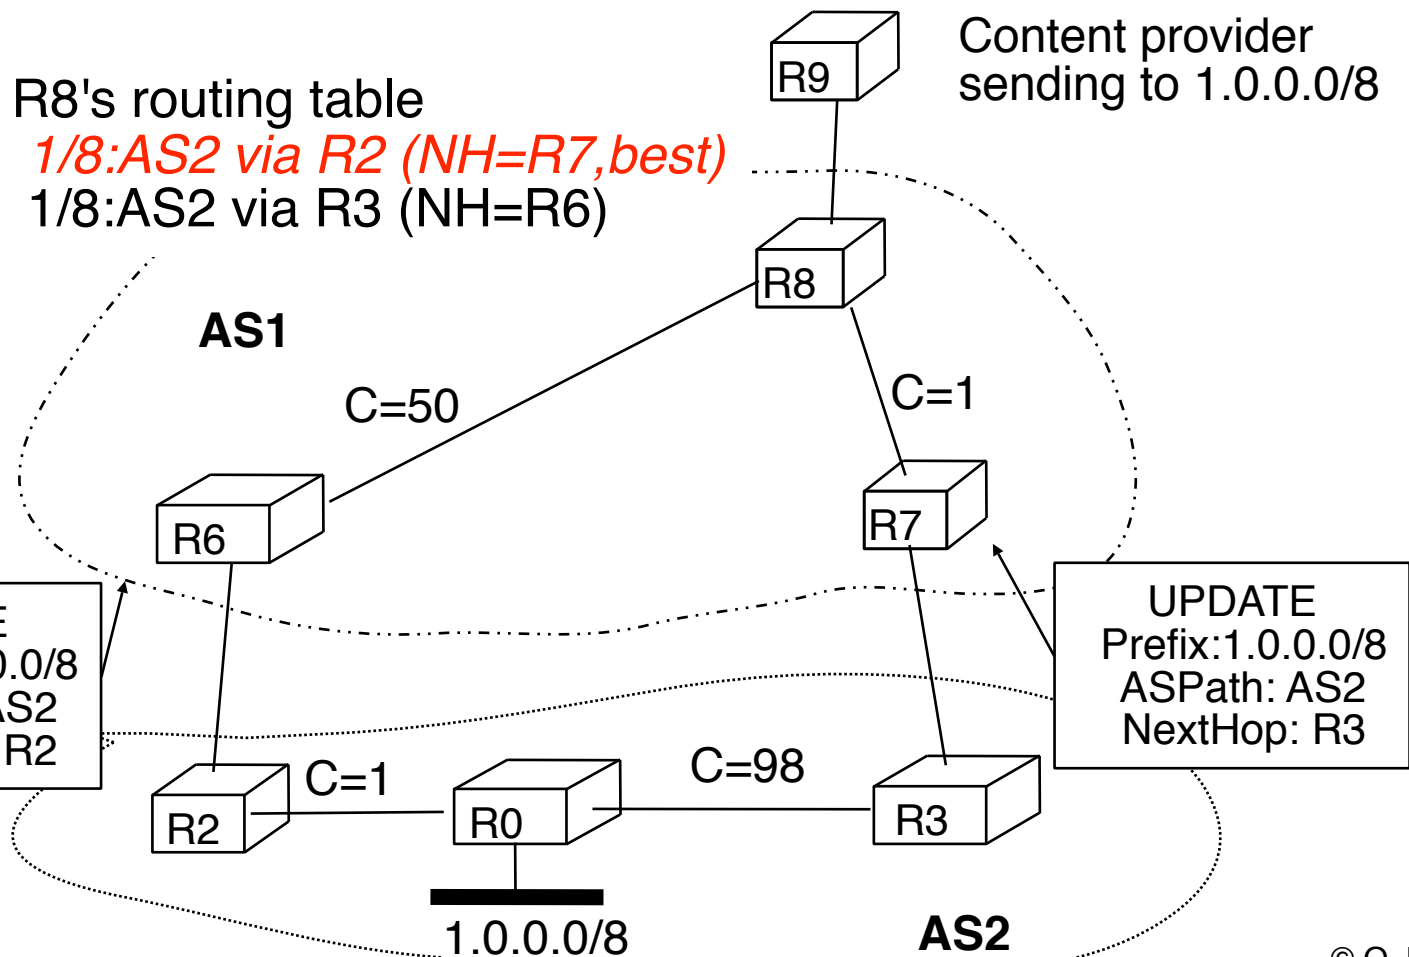
A router should try to get rid of packets sent to external domains as soon as possible



# The closest `nexthop` step in the BGP decision process

## Motivation : hot potato routing

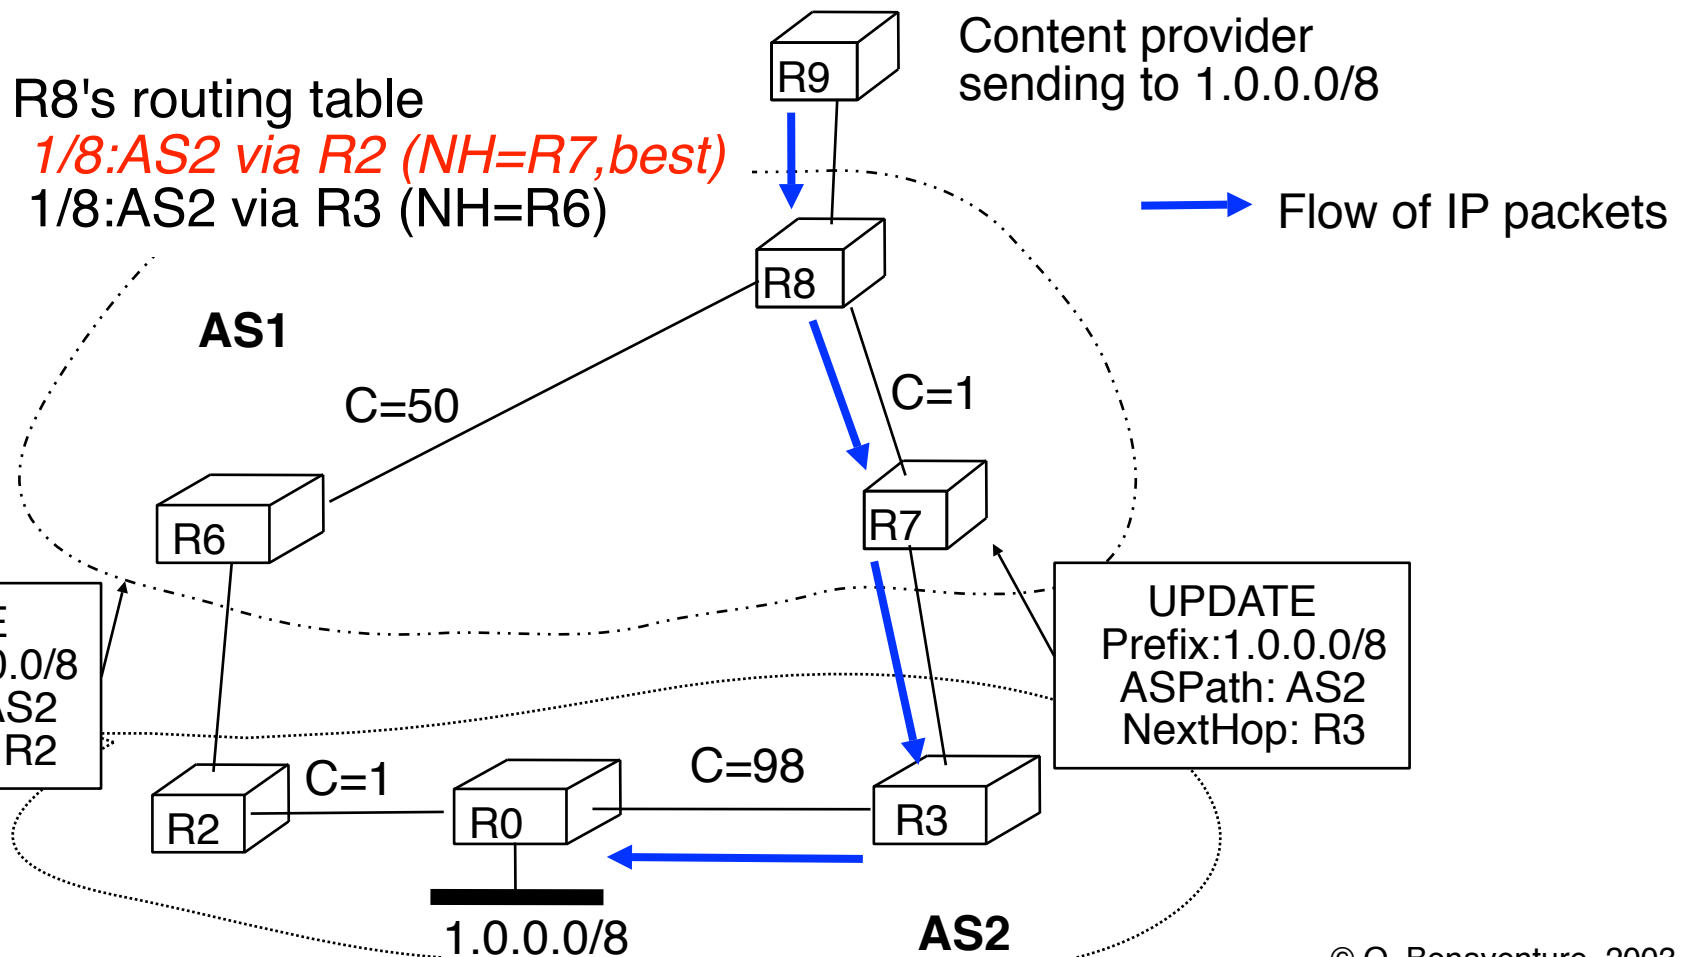
A router should try to get rid of packets sent to external domains as soon as possible



# The closest `nexthop` step in the BGP decision process

## Motivation : hot potato routing

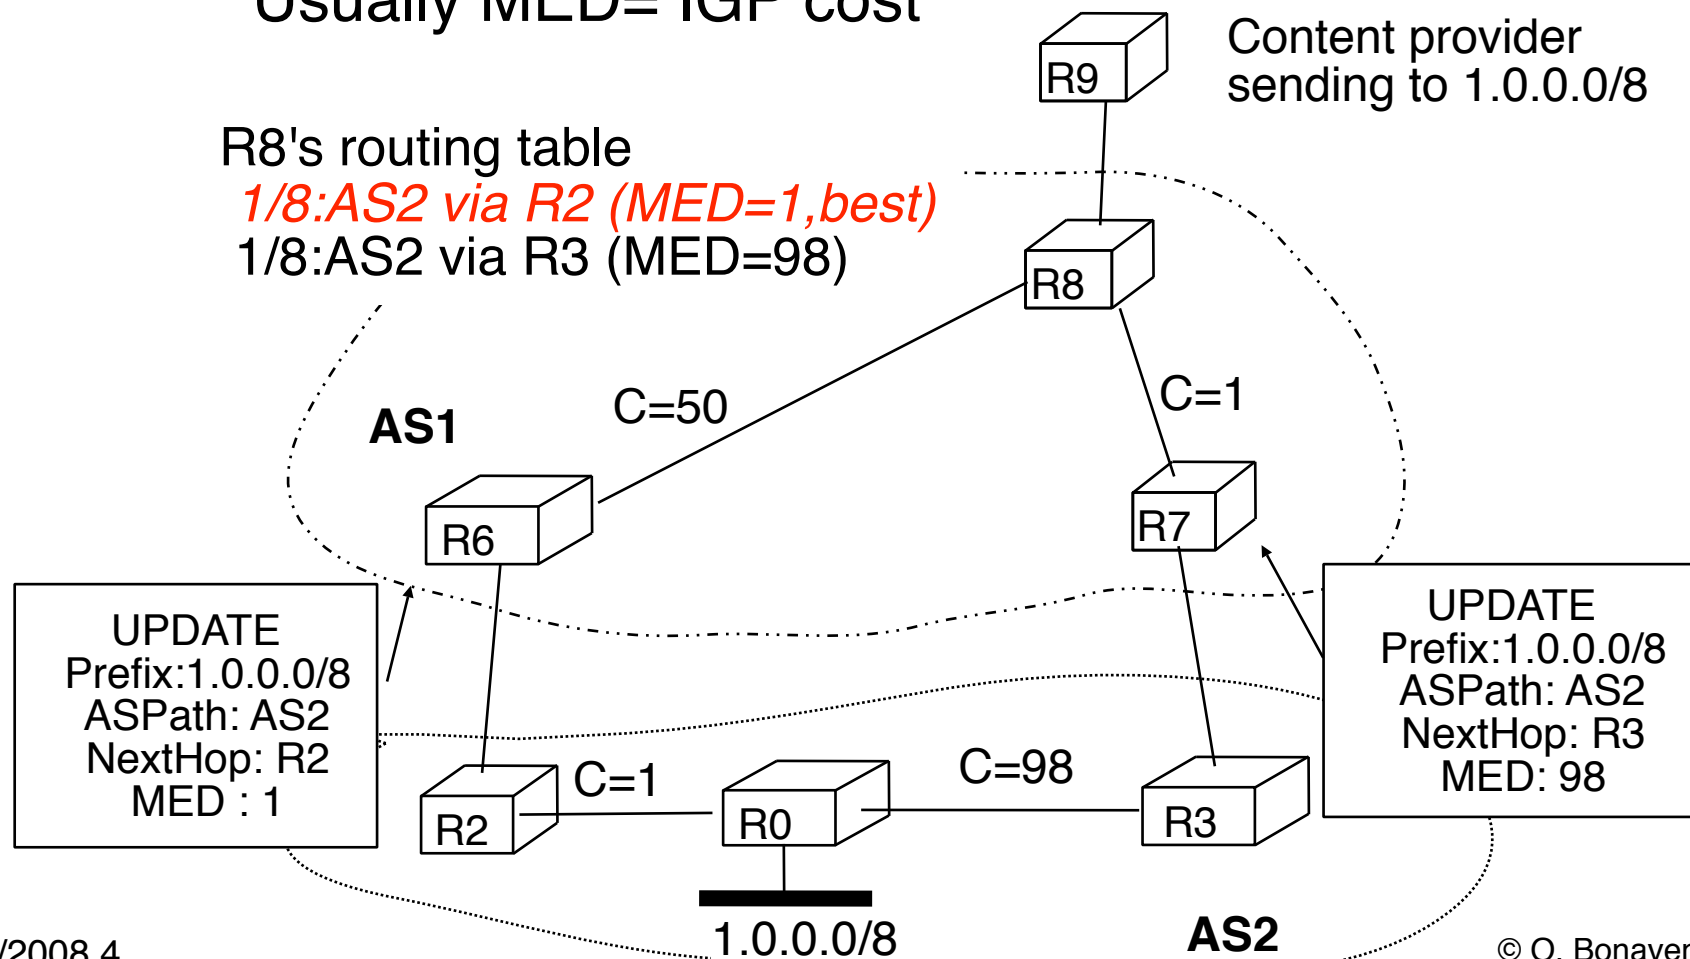
A router should try to get rid of packets sent to external domains as soon as possible



# The lowest MED step in the BGP decision process

## Motivation : cold potato routing

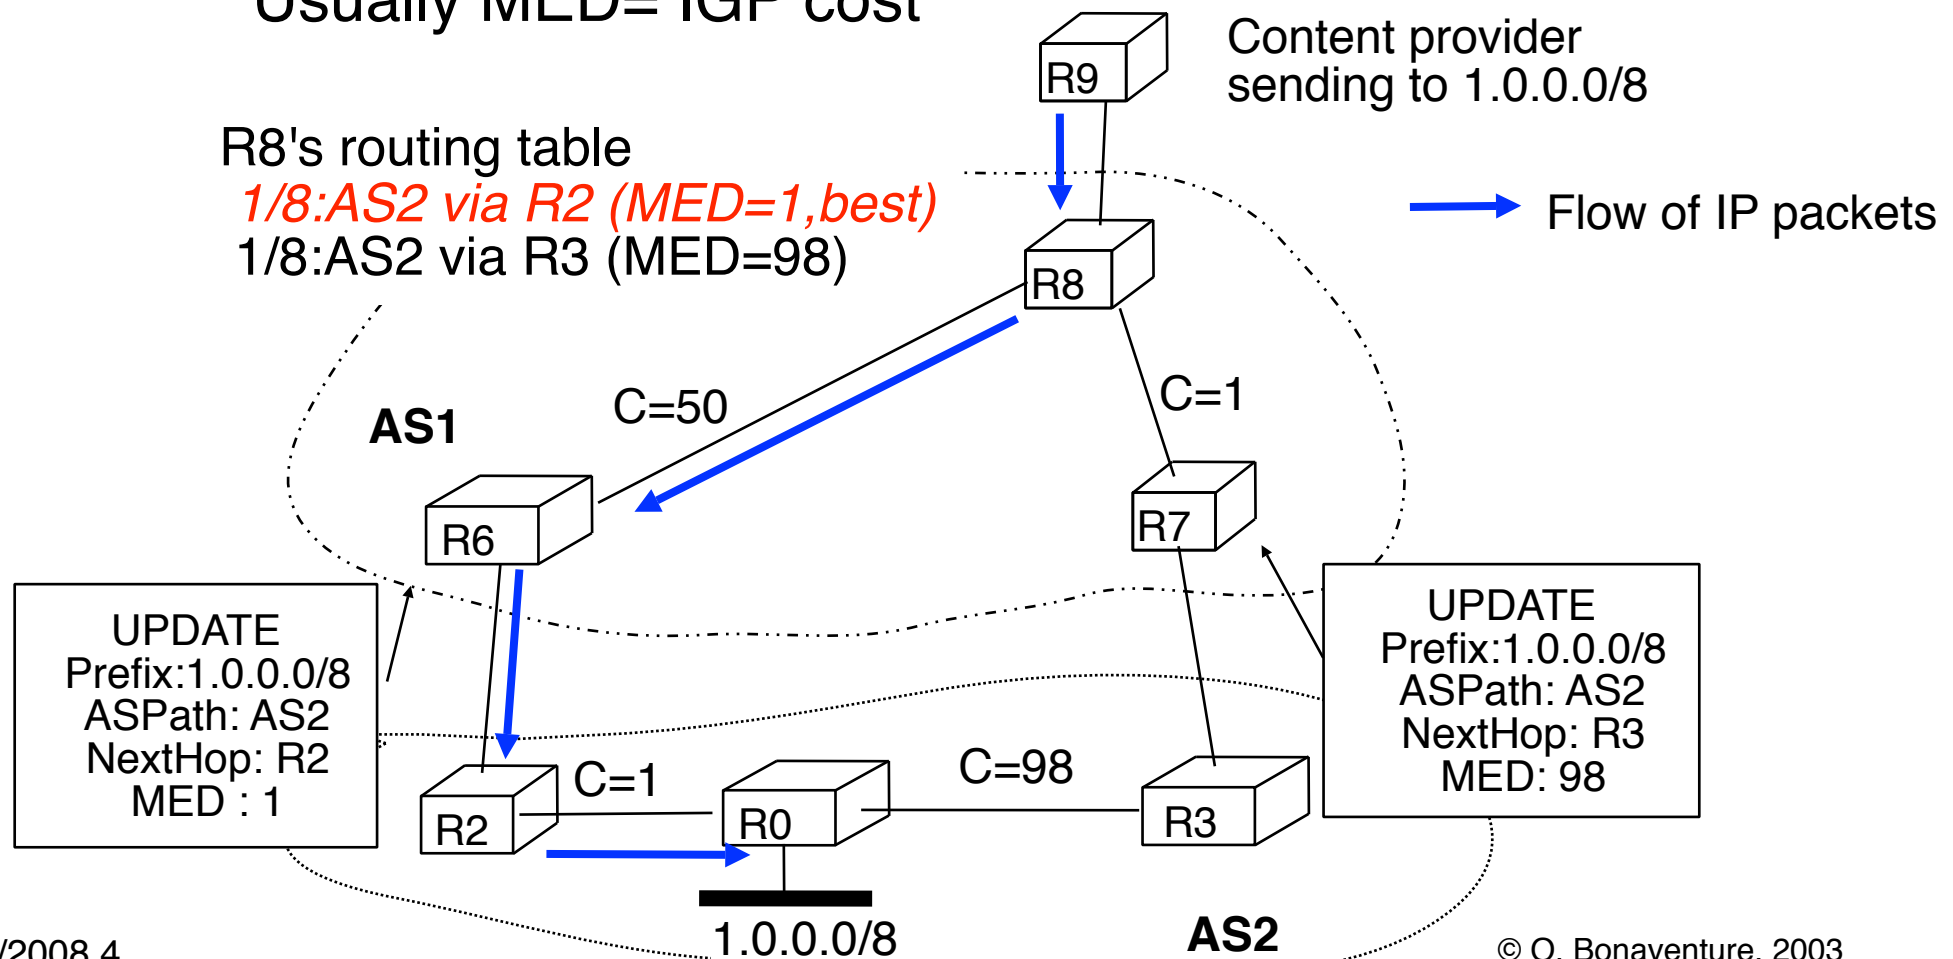
In a multi-connected AS, indicate which entry border router is closest to the advertised prefix  
Usually MED= IGP cost



# The lowest MED step in the BGP decision process

## Motivation : cold potato routing

In a multi-connected AS, indicate which entry border router is closest to the advertised prefix  
Usually MED= IGP cost

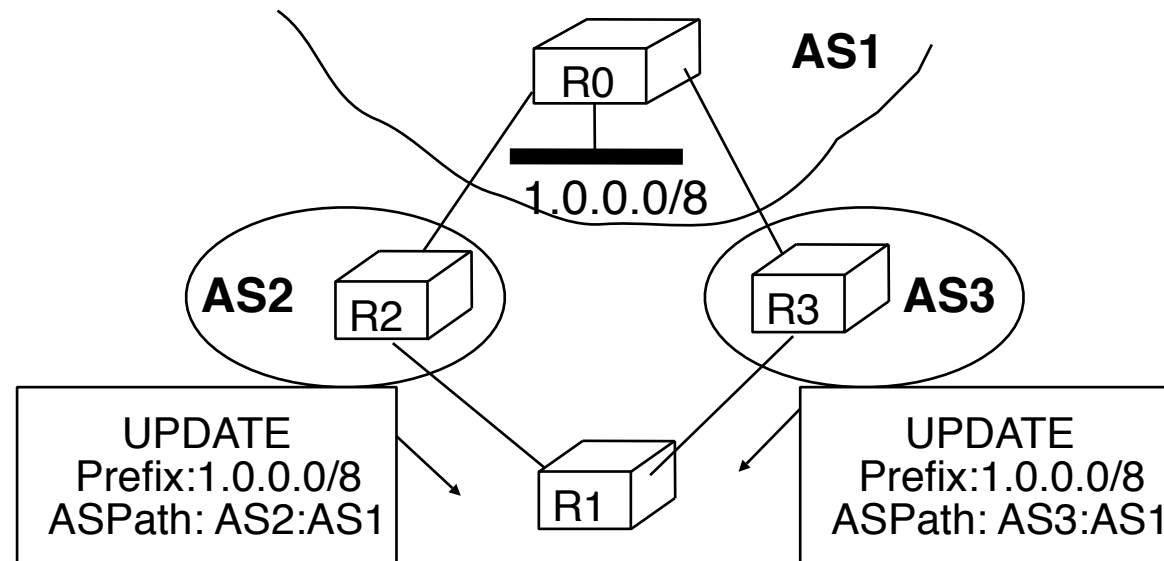


# The lowest router id step in the BGP decision process

## Motivation

A router must be able to determine *one* best route towards each destination prefix

A router may receive several routes with comparable attributes towards one destination



## Consequence

A router with a low IP address will be preferred



# Allocation of IP addresses

---

## How to allocate IP addresses

### First solution

Objective : **Ensure that IP addresses are unique**

Rule used by registries

Any organisation can be allocated a unique IP subnet on a FCFS basis

Size of the allocated subnet : three classes

Class A : subnet with 8 bits mask

Class B : subnet with 16 bits mask

Class C : subnet with 24 bits mask

### Drawbacks

Too rigid

Class A is too large for most networks and Class C too small  
**address waste !**

Difficult to aggregate prefixes

# Allocation of IP addresses (2)

---

## CIDR

### Goals

1. Ensure that IP addresses are unique
2. Allow BGP routers to advertise aggregated prefixes

### Rules used by registries

Only Internet Service Providers (and large companies) can obtain IP subnets

Size of allocated subnet is function of current and expected number of customers

An organisation willing to be connected to the Internet must obtain IP addresses from its ISP

### Advantage

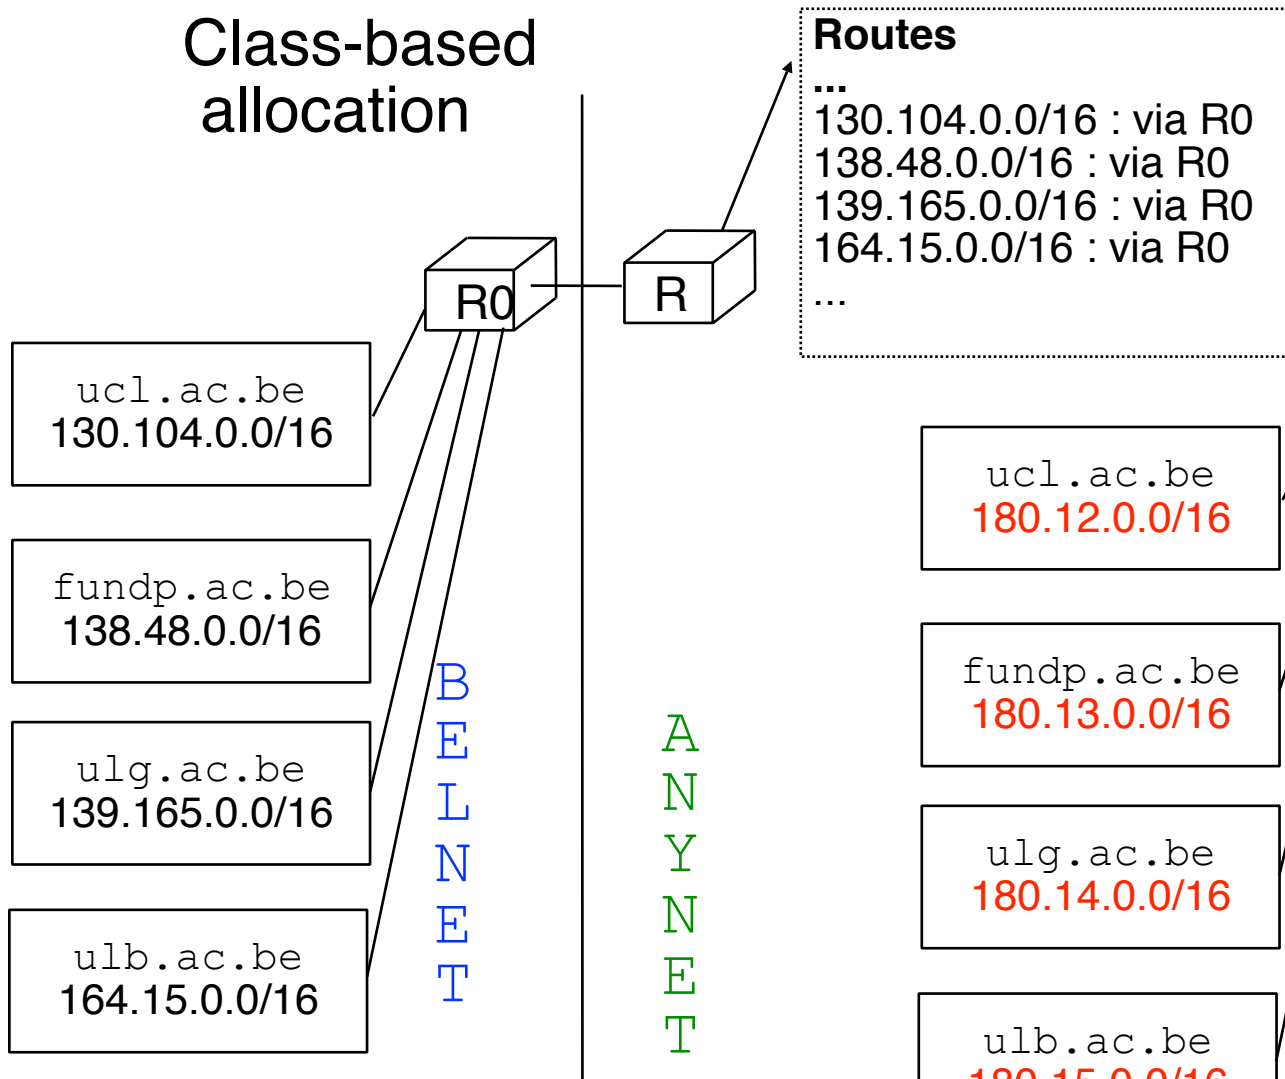
Improved aggregation of addresses

### Drawback

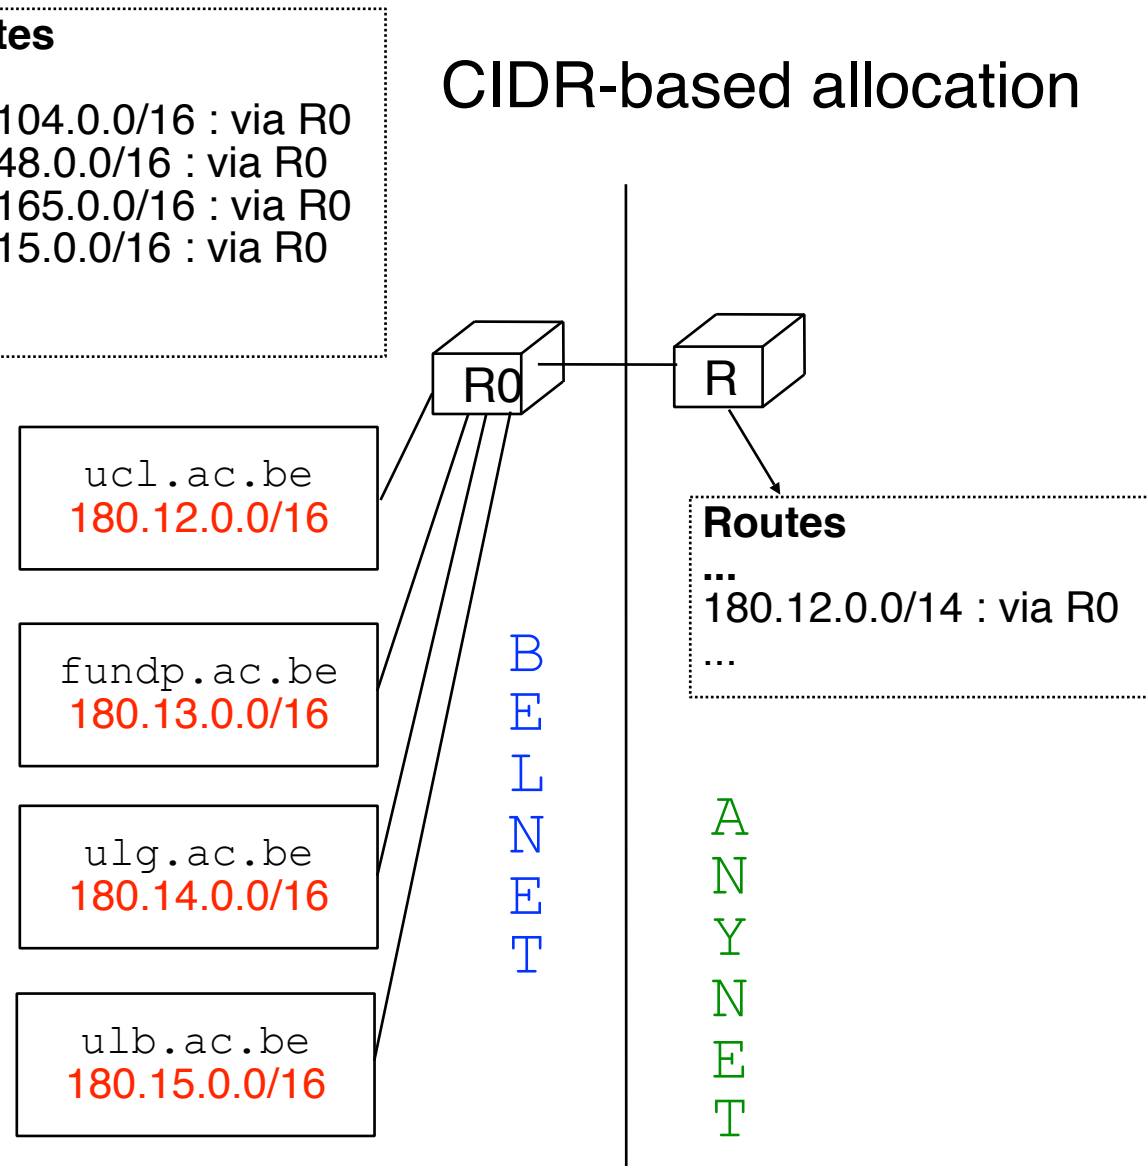
If a company switches from one provider to another, it will need to renumber its IP network - a real pain !

# Allocation of IP addresses (3)

## Class-based allocation




## CIDR-based allocation




# Internet evolution

## Number of Autonomous Systems




# Internet evolution (2)

## Size of the BGP routing tables Number of IPv4 prefixes in default-free routers




# Internet evolution (2)

## Size of the BGP routing tables Number of IPv4 prefixes in default-free routers




# Internet evolution (2)

## Size of the BGP routing tables Number of IPv4 prefixes in default-free routers



# Internet evolution (2)


## Size of the BGP routing tables Number of IPv4 prefixes in default-free routers





# Internet evolution (2)

## Size of the BGP routing tables Number of IPv4 prefixes in default-free routers



# Internet evolution (2)

## Size of the BGP routing tables Number of IPv4 prefixes in default-free routers

