

UCL

Université
catholique
de Louvain



Computer Networking : Principles, Protocols and Practice

Part 2 : Applications

Olivier Bonaventure
<http://inl.info.ucl.ac.be/>

The Application Layer

Contents

→ The client-server model

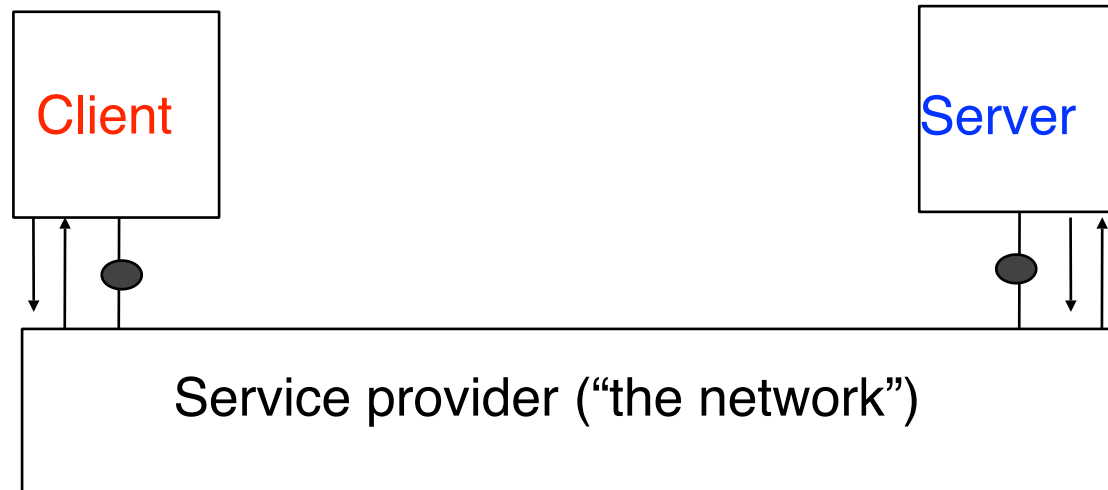
Name to address resolution

email

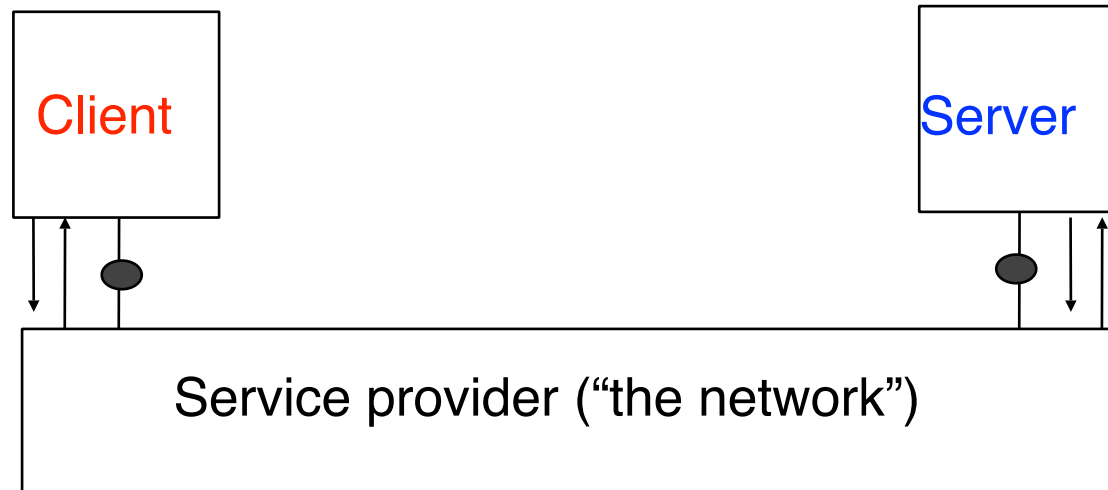
world wide web

peer-to-peer applications

The client server model



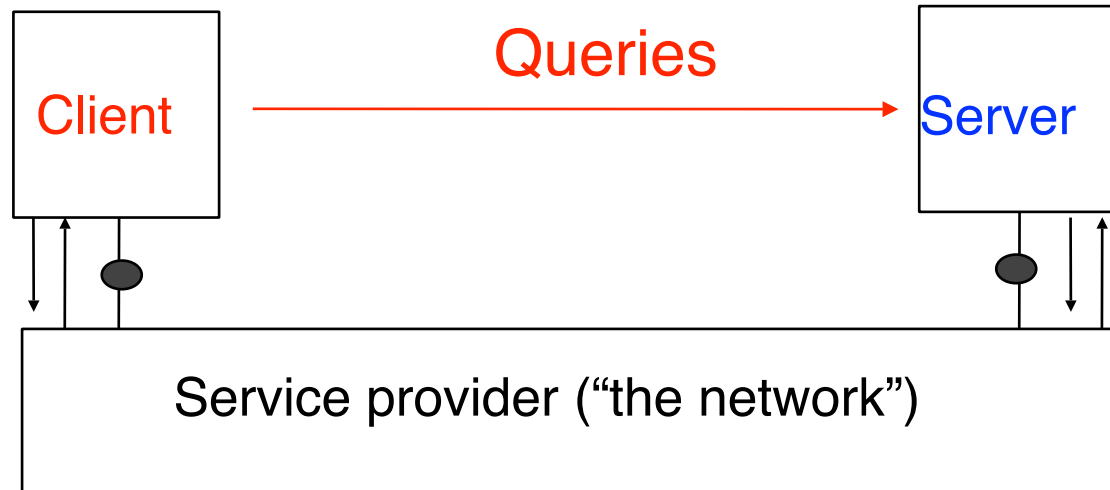
The client server model



Client

interacts with server through transport layer
sends queries or commands

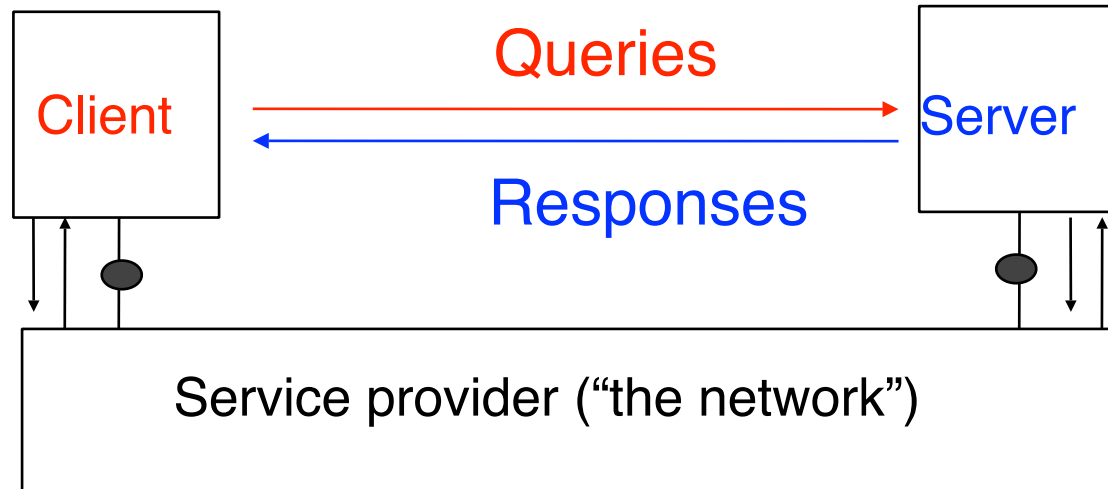
The client server model



Client

interacts with server through transport layer
sends queries or commands

The client server model



Client

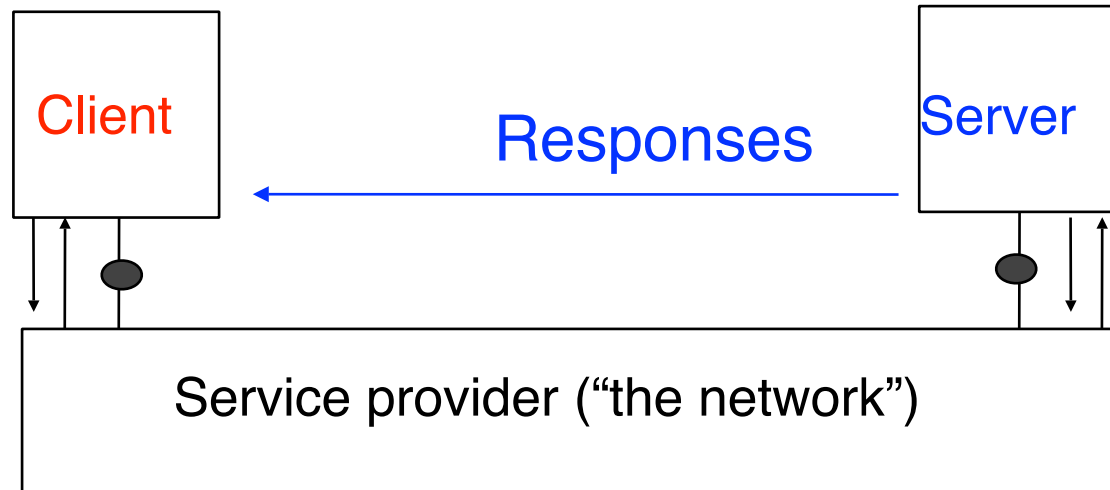
- interacts with server through transport layer
- sends queries or commands

Server

- Answers the queries received from clients
- Executes the commands from clients
- Many clients can use the same server

Example : email, www, ...

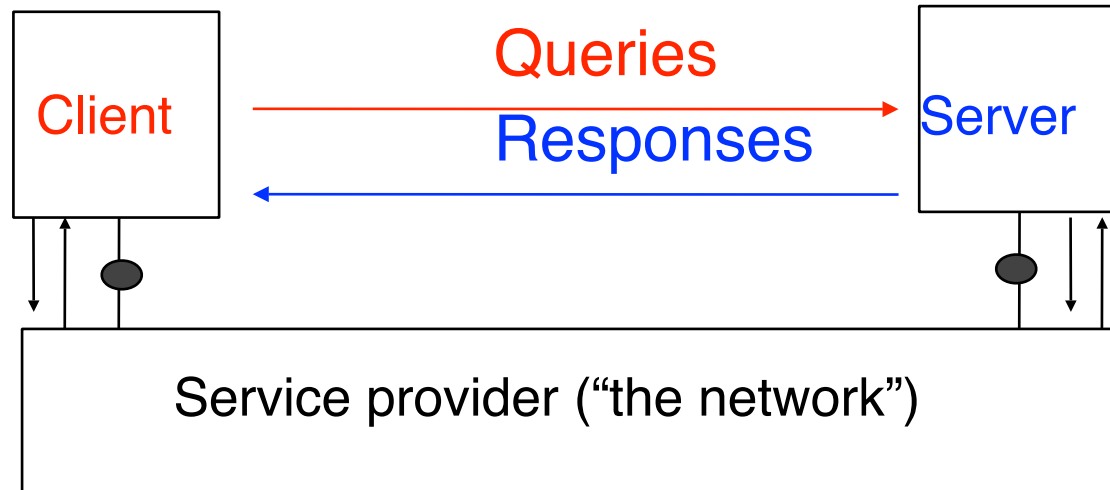
The client server model (2)



Client and servers interact with service provider
Both the client and the server must speak the
same language

Application-level protocol : set of syntactical and
semantical rules that define the messages exchanged
between the client and the server and their ordering

The client server model (2)



Client and servers interact with service provider
Both the client and the server must speak the
same language

Application-level protocol : set of syntactical and
semantical rules that define the messages exchanged
between the client and the server and their ordering

Transport service on the Internet

On the Internet, applications can use two different transport services

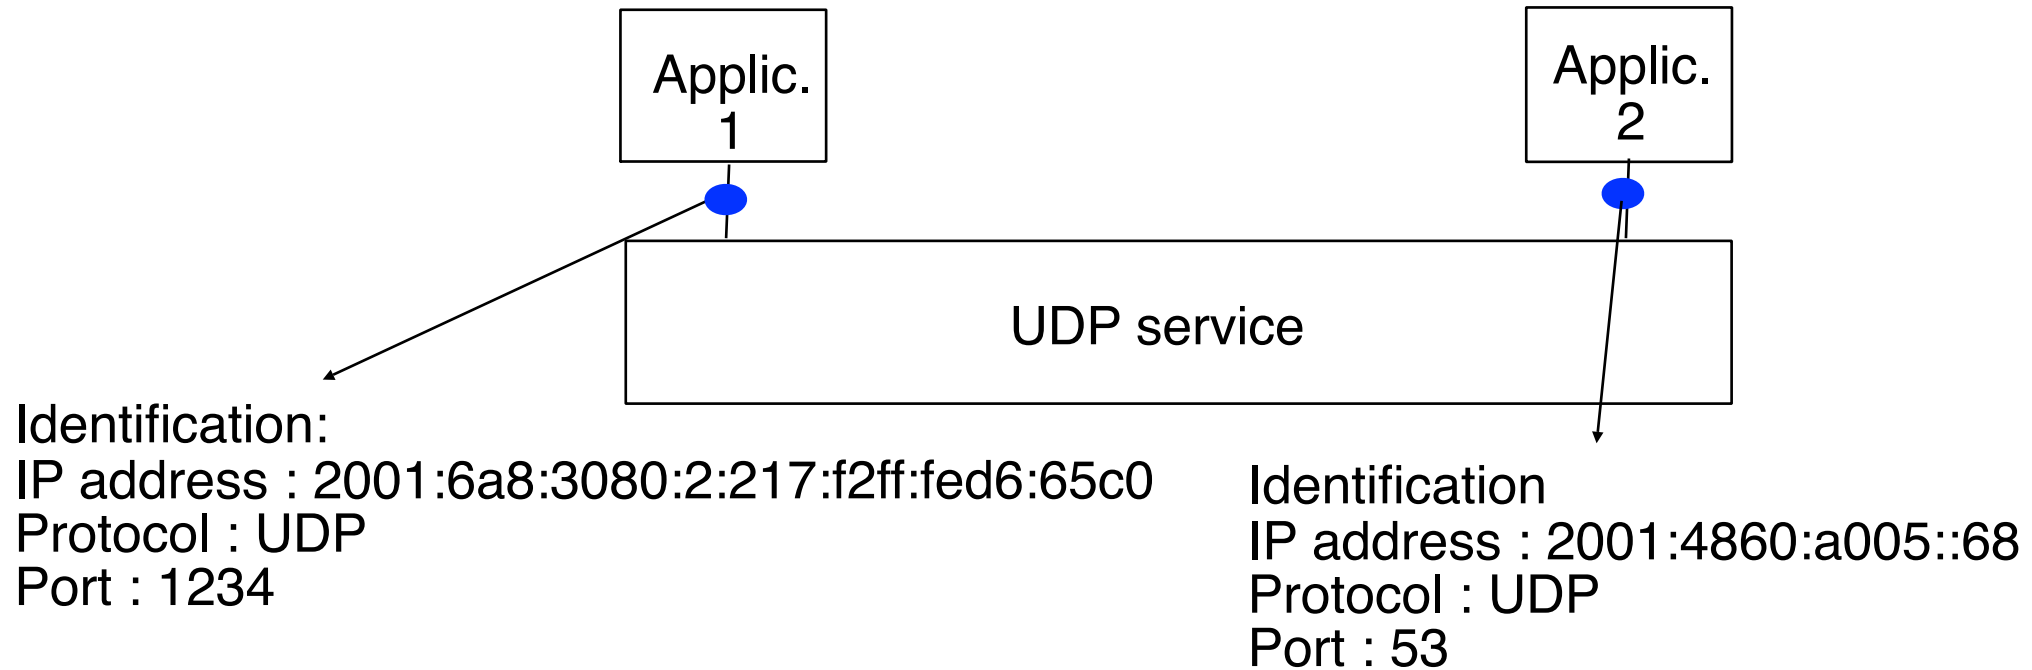
The service provided by the User Datagram Protocol (UDP)

unreliable connectionless service with error detection

The service provided by the Transmission Control Protocol (TCP)

reliable bytestream connection-oriented service

UDP service



Identification of an application

IP address + UDP + port number

Characteristics of UDP service

connectionless

unreliable

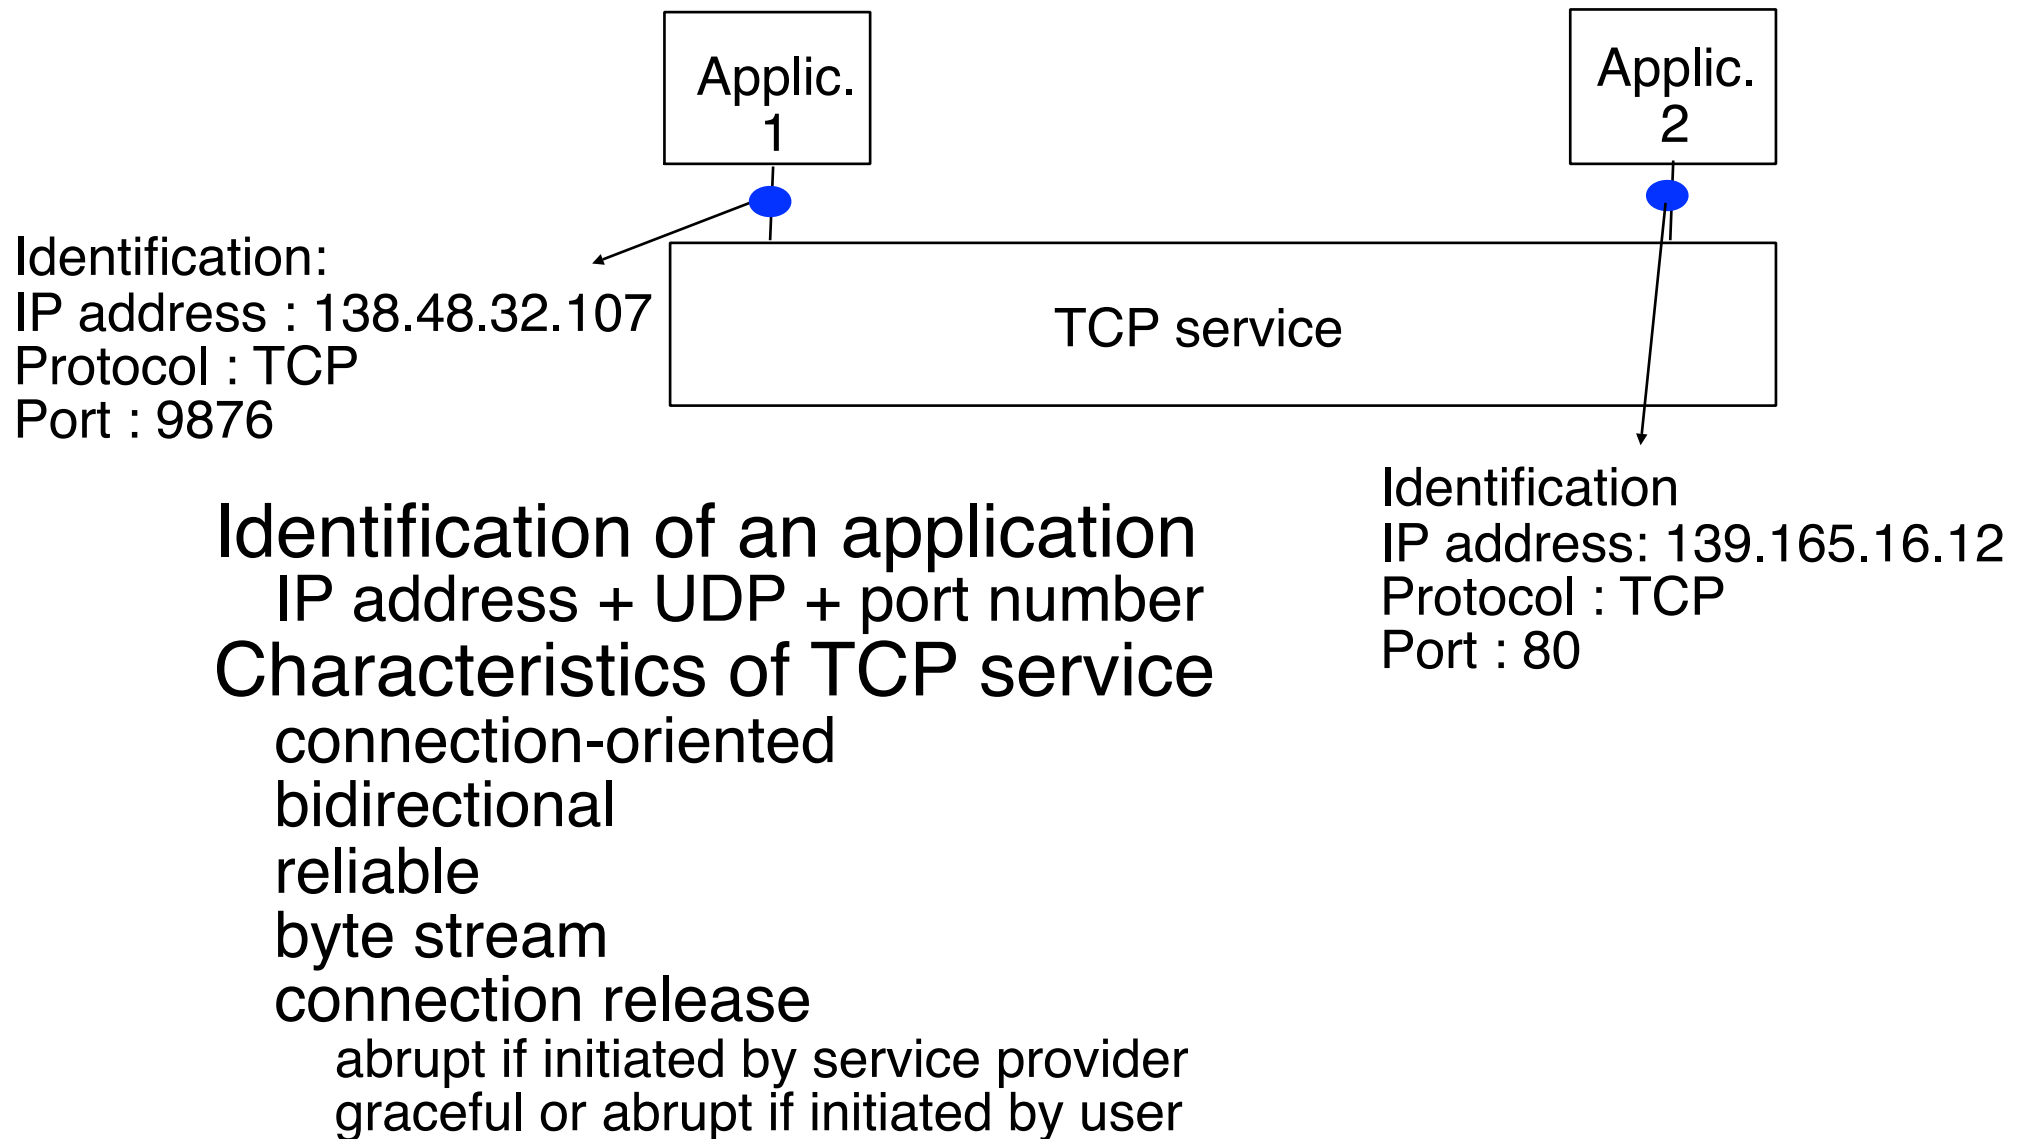
messages can be lost

transmission errors can be detected but not recovered

sequence is not preserved

Maximum size of messages : almost 64 Kbytes

TCP service



Internet applications

Contents

The client-server model

→ Name to address resolution

email

world wide web

peer-to-peer applications

Names and addresses

Names and addresses

Address of a server

IP Address of the host on which the server is running
port number (TCP or UDP)
usually well known port number

Names and addresses

Address of a server

IP Address of the host on which the server is running
port number (TCP or UDP)
usually well known port number

Drawback

Difficult to remember an IP address for a human

Names and addresses

Address of a server

IP Address of the host on which the server is running
port number (TCP or UDP)
usually well known port number

Drawback

Difficult to remember an IP address for a human

Idea

Replace IP address by a hostname
Easier for humans
but IP address is necessary to contact server

How to translate a hostname in an IP address ?

Names and addresses (2)

`hosts.txt` file
contains the name-address table
must be updated regularly

```
#  
# Internet host table  
#  
127.0.0.1      localhost  
138.48.32.99   babbage  
138.48.32.100  leibniz  
138.48.32.1    routeur  
138.48.32.92   corneille  
138.48.32.107  backus  
138.48.20.152  arzach  
138.48.32.137  almin01  
138.48.32.170  duke
```

cannot be used in a large network

Hostnames

Requirement

Host names should be unique

How to achieve this in a scalable manner ?

Introduce hierarchy

Each hostname is composed of two parts

- domain name (globally unique)

- hostname (unique within a given domain)

How to uniquely distribute domain names ?

Introduce hierarchy

- A small number of top-level domain names

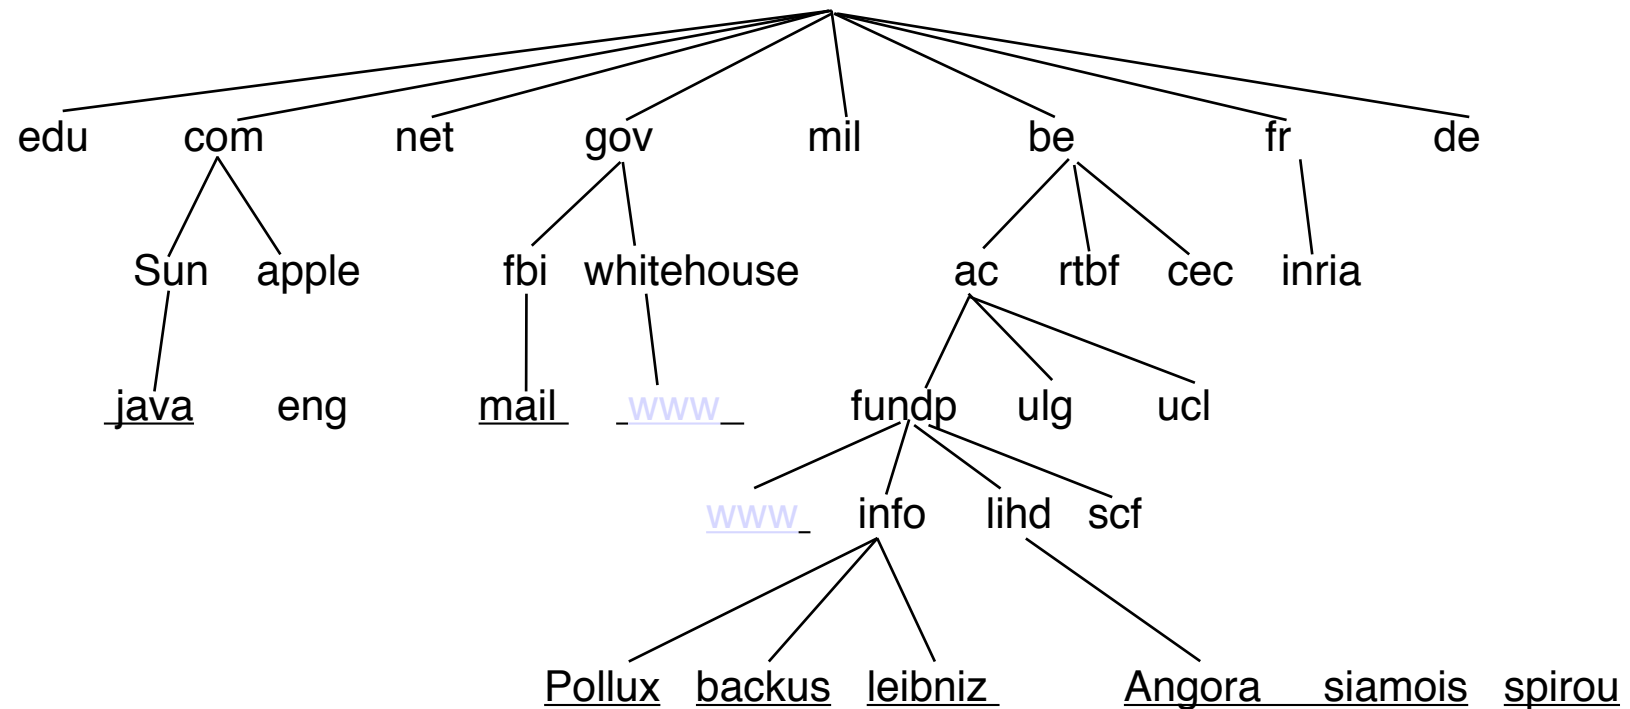
- Inside each top-level domain, allocate uniquely second level domain names

- Inside each seconde-level domain, allocate uniquely either third-level domain names or host names,

...

Host names and domain names

Tree of all host names



How to translate names into addresses ?

How to efficiently translate a host name ?

By using a centralised database


there are more than 1 billion host names today

By using a distributed database

DNS : Domain Name System

relies on the hierarchy of domain names

there is one server responsible for each domain and this server must be queried to translate host names inside this domain



How to translate names into addresses ?

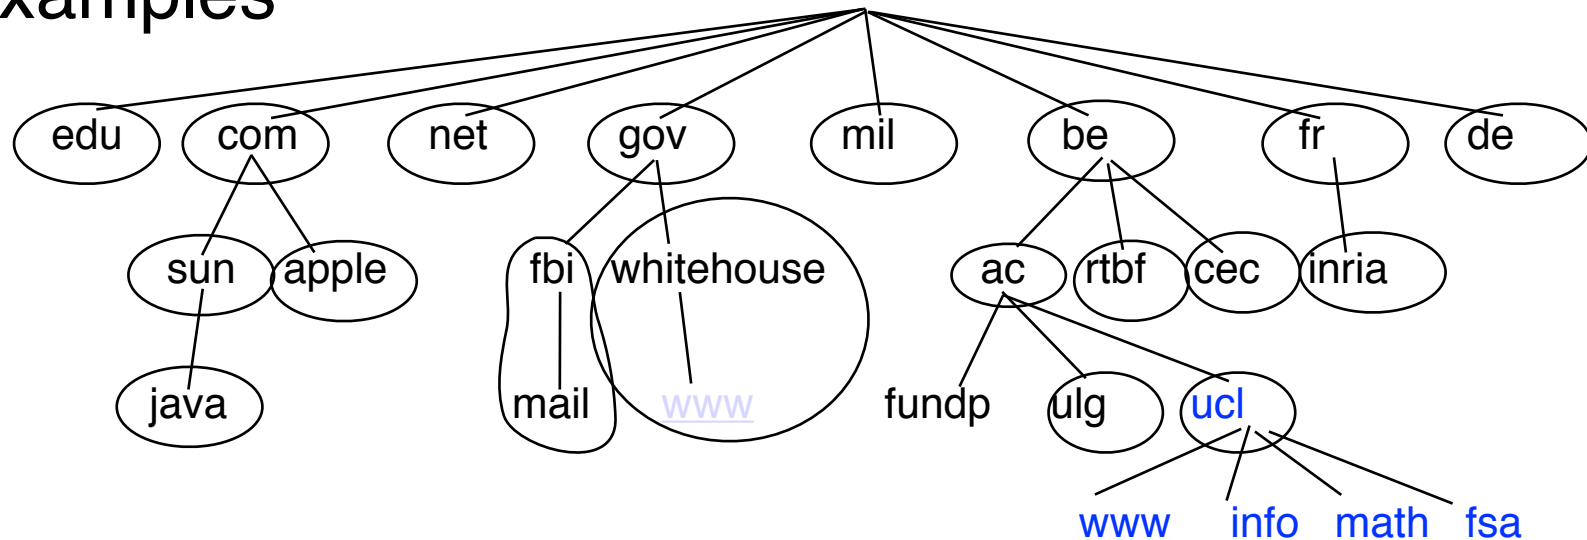
Domain Name Service (DNS)

Each DNS server is responsible for a domain and knows

The IP addresses of all host names in this domain

The IP addresses of the DNS servers responsible for subdomains

Examples



java.sun.com
www.ucl.ac.be

DNS resolver

To be able to translate name to addresses, a DNS implementation needs

- to know **actual** list of IP addresses of root servers
- to implement the DNS protocol and traverse the domain names hierarchy
- Difficult to do this on all endhosts

Solution

DNS resolver

- one resolver for a set of endhosts
- maintains up-to-date list of IP addresses of root servers
- implements DNS protocol

endhosts

- only need to be able to send DNS requests to resolver
- must know IP address of closest DNS resolvers

DNS : optimisations

Reduce risk of failures

- several root-servers

- server DNS servers authoritative for each domain

- each endhost can send queries to multiple resolvers

Improved performance

- avoid sending several times the same query

- cache memory on DNS resolvers containing

 - recent name-addresses translations

 - addresses of DNS servers recently contacted

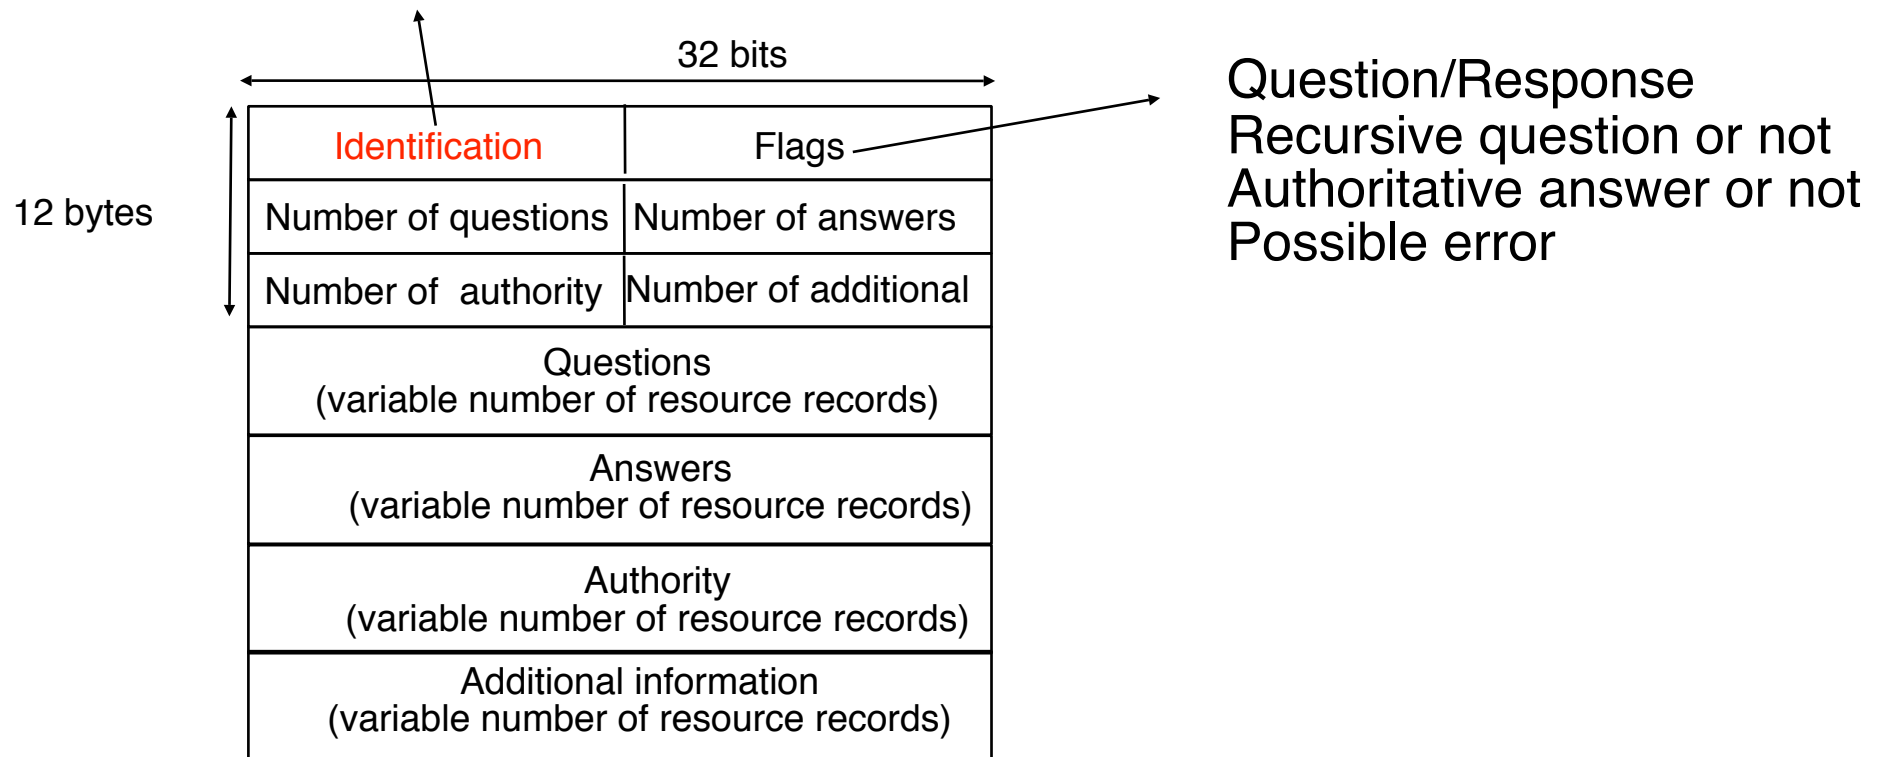
DNS protocol

- usually runs over UDP

- sometimes is also used over TCP

DNS : message format

Each DNS request contains a number that will be returned in the response by the server to allow the client to match the request.



DNS : resource records

Each DNS messages is composed of resource records (RR) encoded as **TLV**

< **Name**, **Value**, **Type**, **TTL** >

Types de RR

↙ Lifetime of the RR in server's cache

A (Address)

Name is a hostname and **Value** an **IPv4 address**

AAAA (Address)

Name is a hostname and **Value** an **IPv6 address**

NS (NameServer)

Name is a domain name and **Value** is the hostname of the DNS server responsible for this domain

MX (Mail Exchange)

Name is a domain name and **Value** is the name of the SMTP server that must be contacted to send emails to this domain

Type CNAME

Alias

Internet applications

Contents

The client-server model

Name to address resolution

→ email

world wide web

peer-to-peer applications

Email

Simplified model

Alice sends an email to Bob



Alice@a.net



Alice's
email server



b.net 's
email server

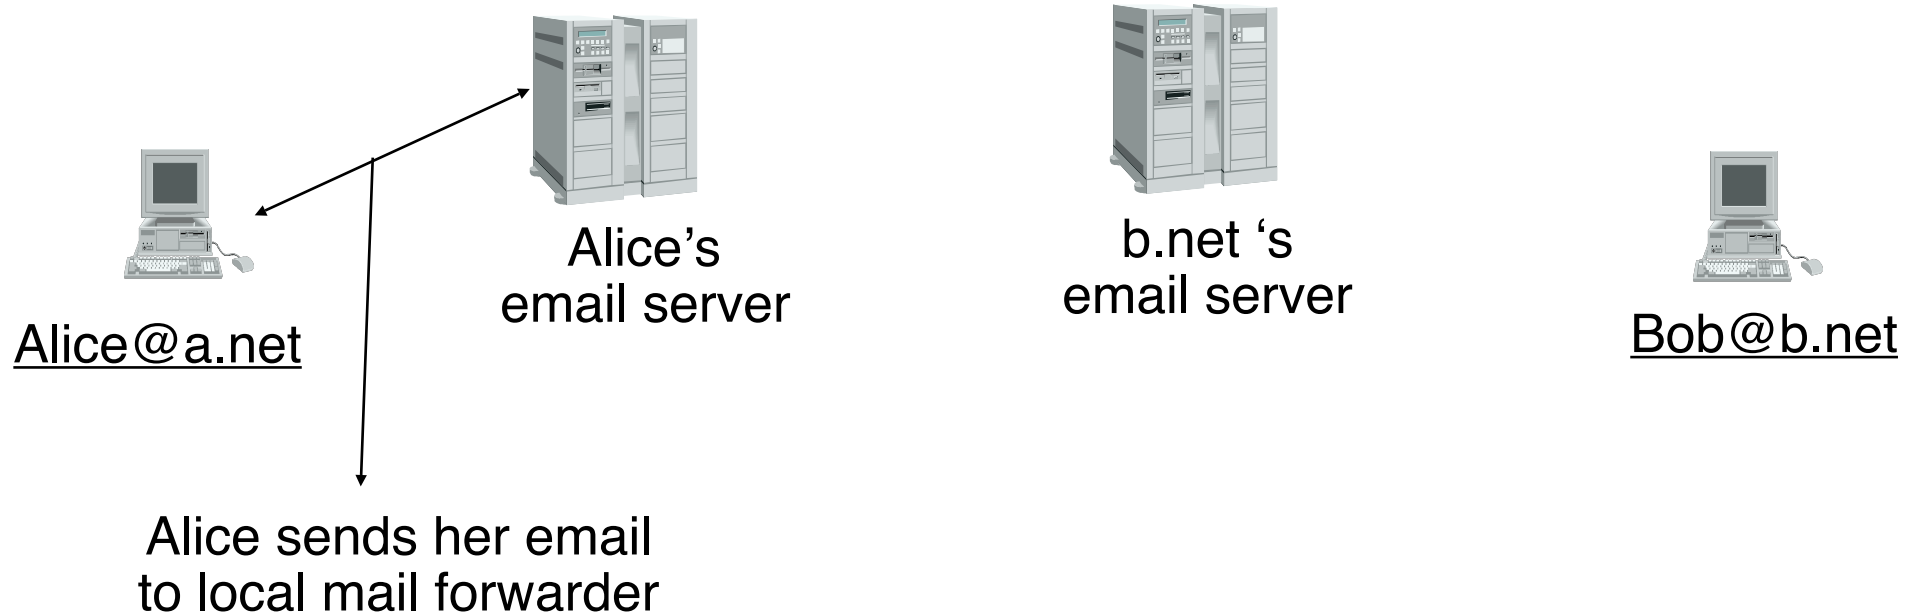


Bob@b.net

Email

Simplified model

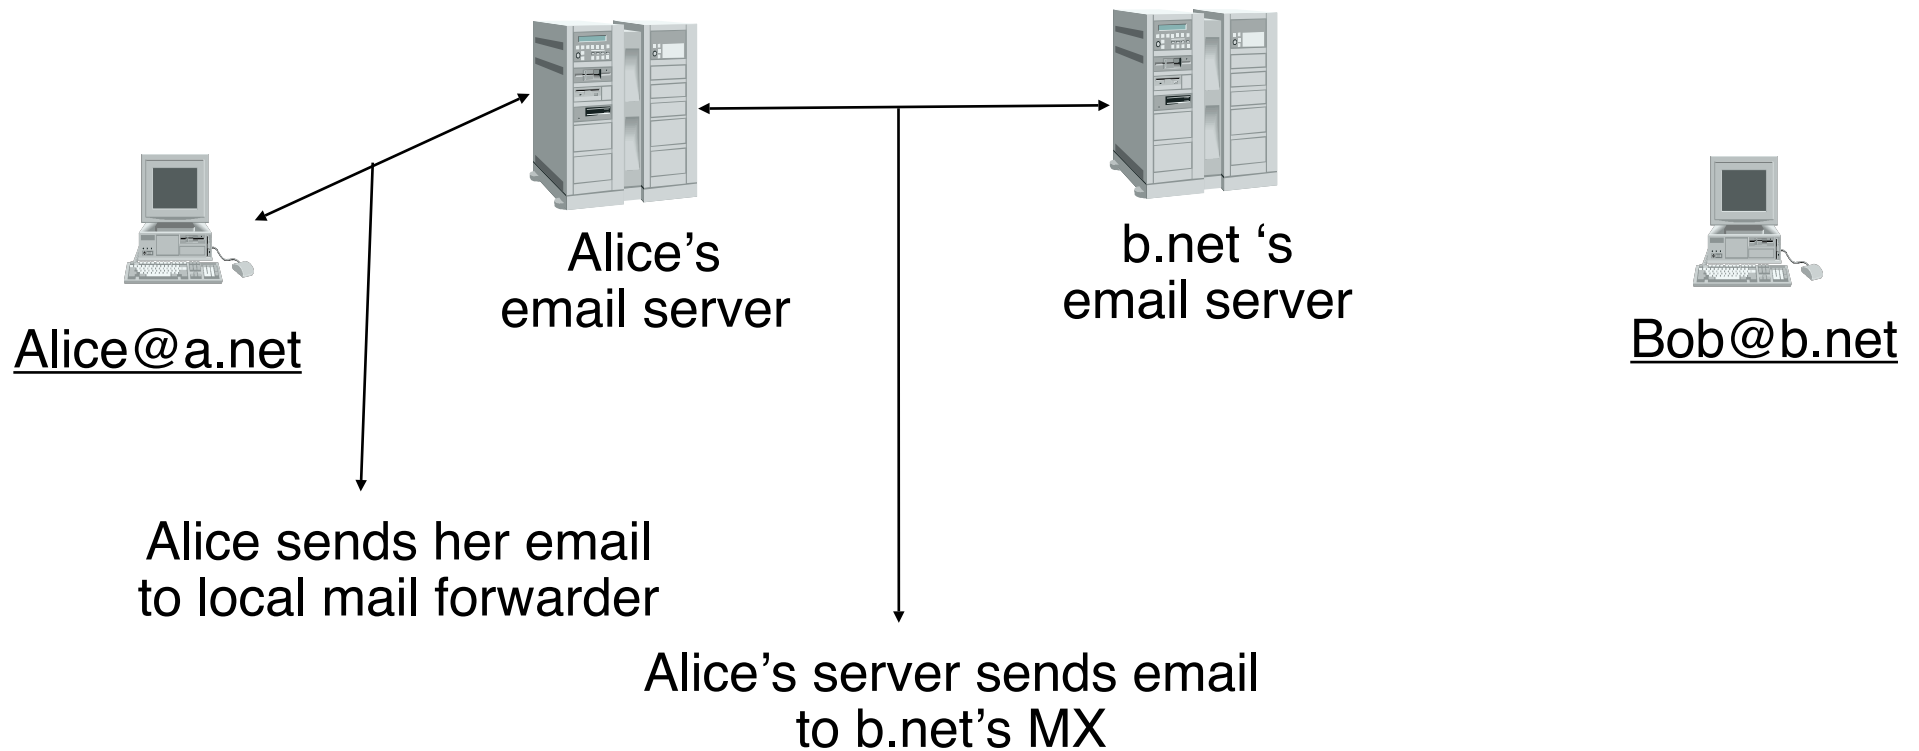
Alice sends an email to Bob



Email

Simplified model

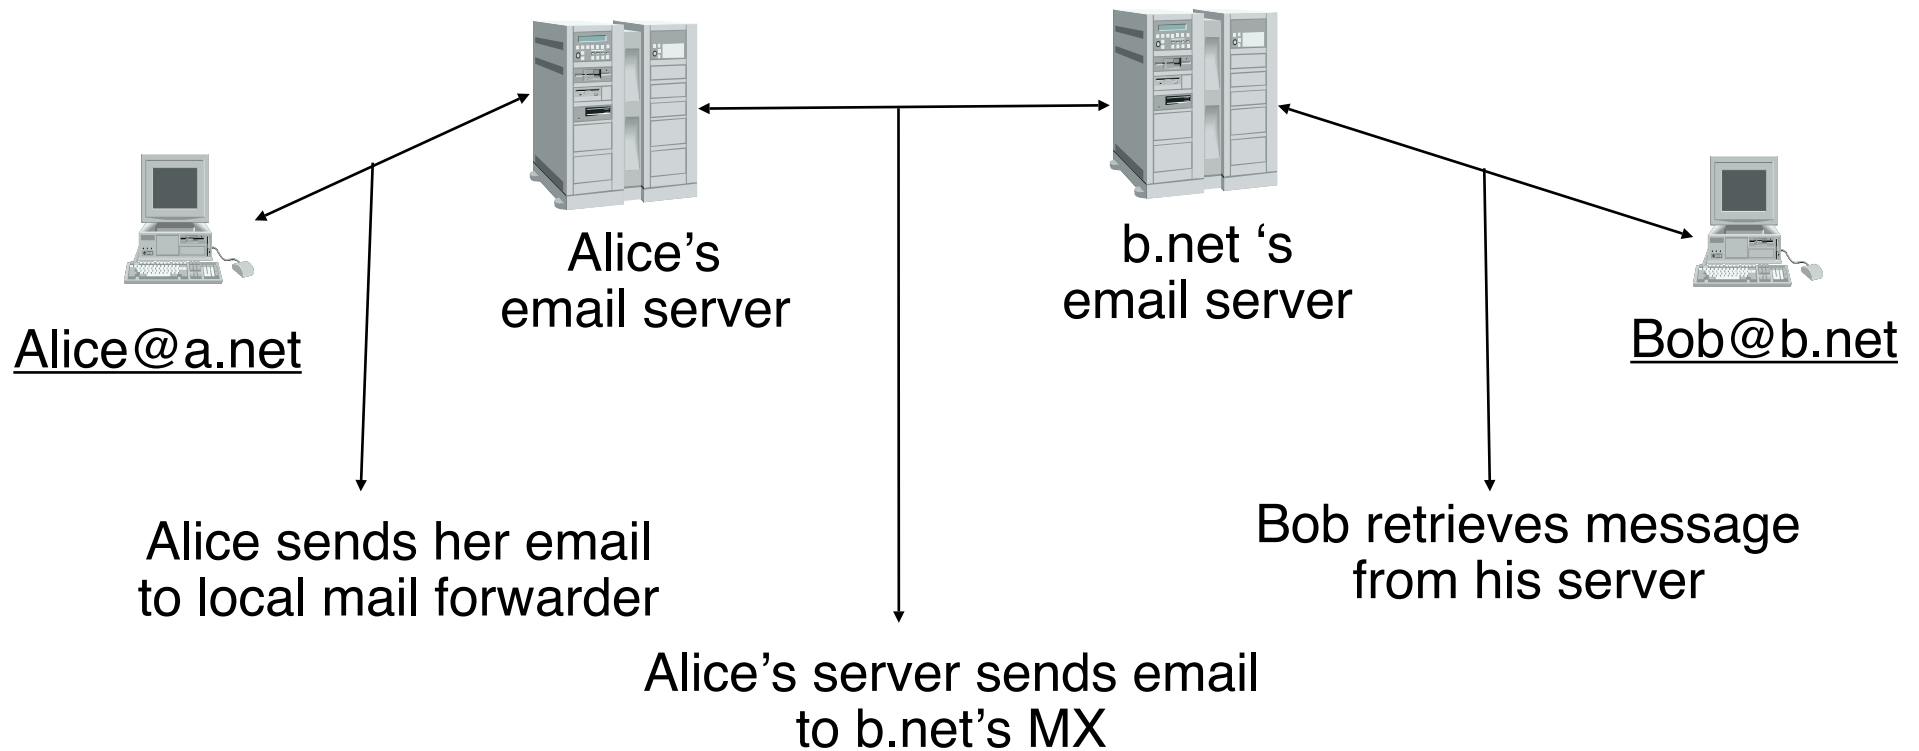
Alice sends an email to Bob



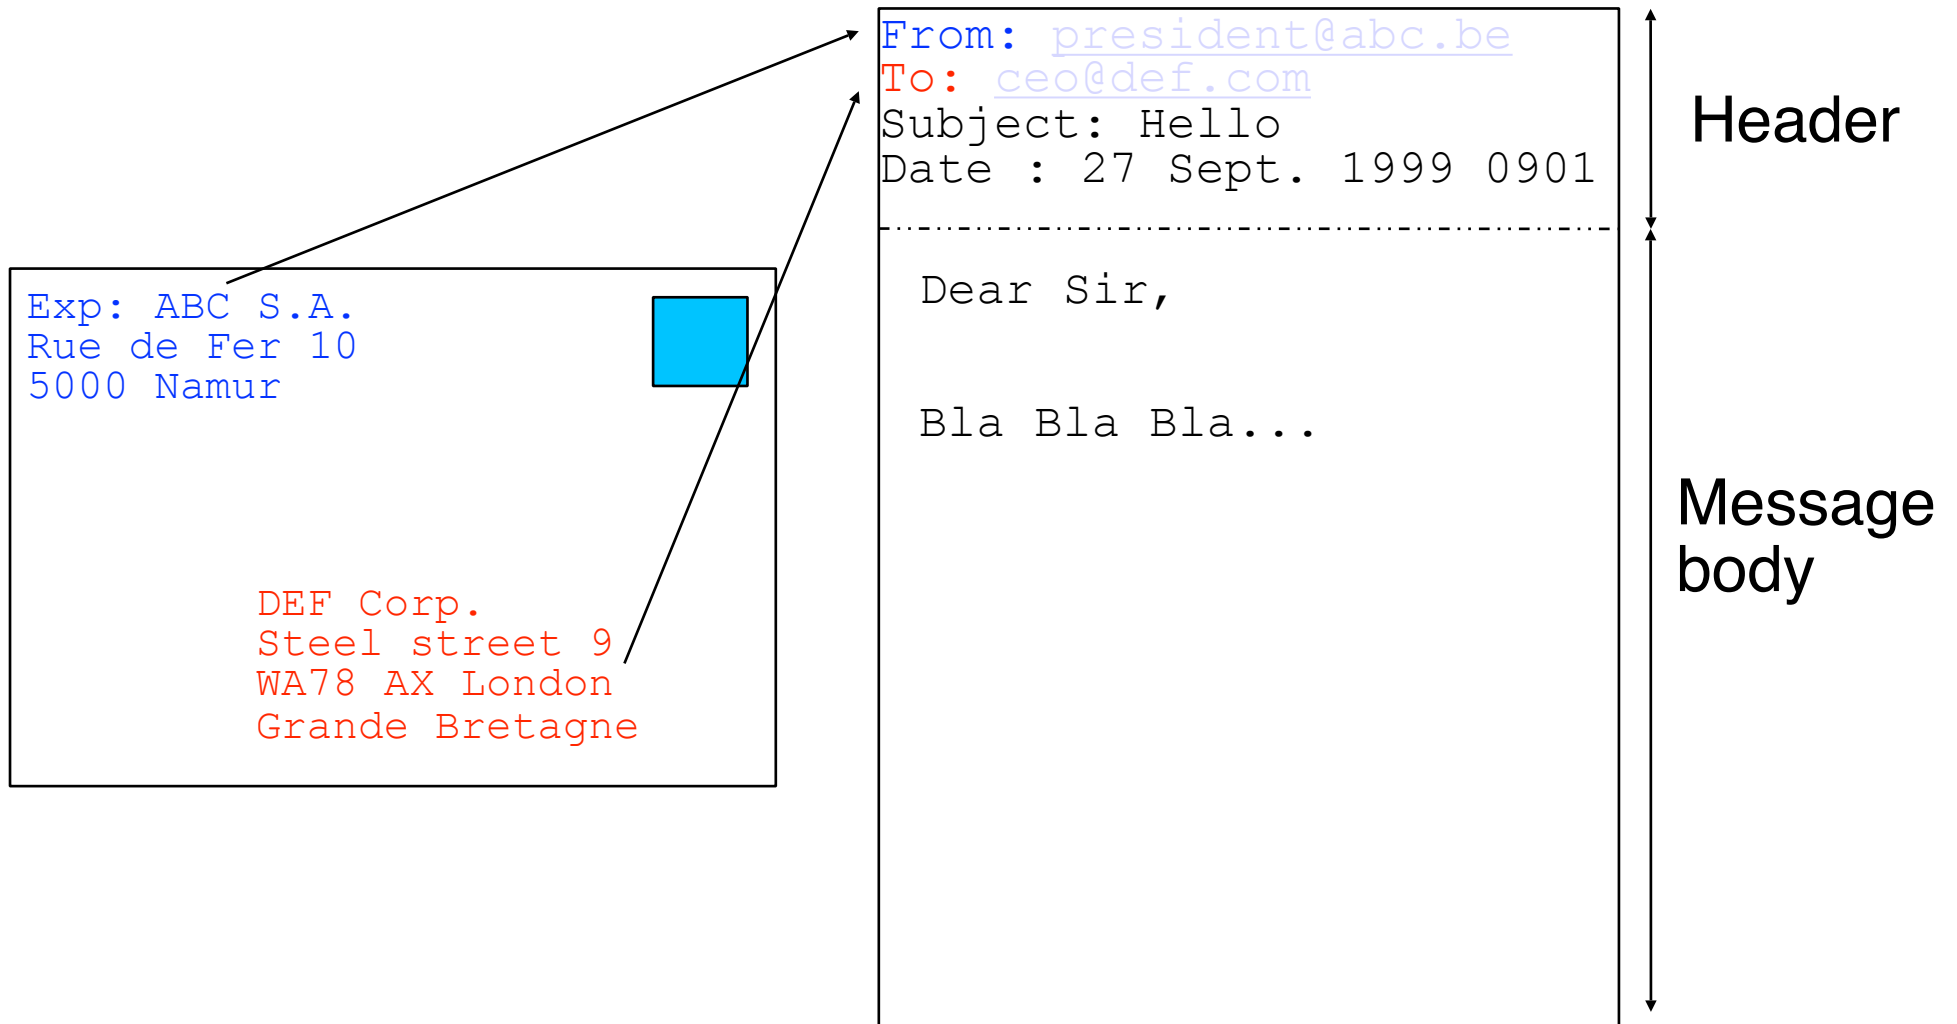
Email

Simplified model

Alice sends an email to Bob



Email message format



Message format (2)

Header format

Contains only US-ASCII (7bits) characters

At least three lines that end with <CRLF>

From: sender@domain

To: [recipient@domain](#)

Date: <creation date of message>

example : 26 Aug 199 1445 EDT

Optional fields

Subject: subject of message

cc: [copy@domain](#)

Message-ID: <[number@domain](#)>

Received: information on path followed by message

In-Reply-To: <message-ID>

Header ends with empty line (<CRLF>)

MIME

MIME

Internet email was designed for US-ASCII
How to transmit more complex messages ?

MIME

Internet email was designed for US-ASCII

How to transmit more complex messages ?

Multipurpose Internet Mail Extensions

Improved email message format

Constraints

- must remain compatible with old email servers

 - most of them only support US-ASCII and short lines

- must support non-English text

 - character set must be beyond 7bits US-ASCII

- must support various formats in a single message

 - message body, attachments, ...

- must allow to transmit audio, video, ...

 - need to identify the type of content

MIME

Internet email was designed for US-ASCII

How to transmit more complex messages ?

Multipurpose Internet Mail Extensions

Improved email message format

Constraints

- must remain compatible with old email servers

 - most of them only support US-ASCII and short lines

- must support non-English text

 - character set must be beyond 7bits US-ASCII

- must support various formats in a single message

 - message body, attachments, ...

- must allow to transmit audio, video, ...

 - need to identify the type of content

Solution

- add new optional fields in header

- add optional fields inside message body when

MIME (2)

New header fields

MIME-Version:

version of MIME used to encode message
current version : 1.0

Content-Description:

comment describing the content of the message

Content-Type:

type of information inside message

Content-Transfer-Encoding:

how the message has been encoded

Content-Id:

unique identifier for the content

MIME: Content-Type

Content-Type : type/encoding

type of content

text, image, video, application

multipart

encoding of content

text/plain , text/html

image/gif, image/jpeg

audio/basic

video/mpeg, video/quicktime

application/octet-stream, application/postscript

multipart/alternative

message contains several times the same information with different encodings

multipart/mixed

message contains several information of different types

example : text of message body and attachment

Character sets and content encoding

How to support rich character sets ?

Content-Type: text/plain; charset=us-ascii

ASCII 7bits, default

Content-Type: text/plain; charset=iso-8859-1

Character set suitable for Western European languages, defined by ISO, 8 bits per character

Content-Type: text/plain; charset=unicode

Universal character set, defined by ISO, 16 bits per character

How to encode non-text data ?

data must be encoded in US-ASCII 7 bits characters

Base64

uses ASCII caracteres A...Z, a...z, 0...9, "+" et "/"

A=0, B=1, C=2, ... +=62 et /=63

Each character is used to encode 6 bits

24 bits from initial message -> 4 ASCII characters

Special character "=" used for padding

Multipart/mixed

How to place different contents and encoding in a single message ?

We need a delimiter between the different content types placed inside message body

```
Date: Mon, 20 Sep 1999 16:33:16 +0200
From: Nathaniel Borenstein <nsb@bellcore.com>
To: Ned Freed <ned@innosoft.com>
Subject: Test
MIME-Version: 1.0
Content-Type: multipart/mixed; boundary="simple boundary"

preamble, to be ignored

--simple boundary
Content-Type: text/plain; charset=us-ascii

partie 1

--simple boundary
Content-Type: text/plain; charset=us-ascii

partie 2
--simple boundary
```



Email transmission

SMTP : Simple Mail Transfer protocol
uses TCP service

Address of SMTP server

IP address of server + TCP + port number: 25

RR of type MX can be used to find the SMTP server responsible for a given domain



SMTP

Client-server model

Server waits for email messages to relay/deliver

Client sends email messages through server

Application-level protocol

client opens TCP connection

Client sends commands composed of
command parameter <CRLF>

HELO

MAIL FROM:

RCPT TO:

DATA

QUIT

Server answers with one-line replies

numeric_code comment (text) <CRLF>

250 OK

221 closing

SMTP (2)

Three phases of SMTP

1. Establishment of an SMTP association

TCP connection established upon request from client

Server greetings

HELO command from client

2. Message transfer

MAIL FROM: <[user@domaine](#)>

RCPT TO: <[user@domaine](#)>

DATA

transmission of entire message including headers

one line containing only the dot “.” characters marks end of message

Other subsequent messages can be transmitted after

3. Release of the SMTP association

QUIT

Closing message from server

TCP connection is closed

Retrieval of email messages

In the old days

1. Destination is always connected to the Internet
email addresses are username@hostname
When an email arrives, it is stored in a file that belongs to the user, e.g. `/var/mail` on Unix

Today

Most networks have one or a few SMTP servers used to receive emails, but also detect spam, viruses, ...

Endusers retrieve their emails from this server

Post Office Protocol (POP)

Internet Mail Access Protocol (IMAP)

Webmail

POP

Goal

Allow authenticated users to retrieve email messages from server

Operation

POP uses TCP service

Address of POP server

Host address + TCP + port number : 110

Client send commands

command : one ASCII line ending with <CRLF>

USER, PASS, STAT, RETR, DELE, QUIT

server replies with

+OK if command was successful

email messages follow some +OK replies

-ERR in case of errors

POP (2)

Three phases of the protocol

1. Authorisation : checking the user credentials

USER <username>

PASS <password>

2. Transaction

retrieval and removal of messages

STAT

list headers of stored messages

RETR <n>

retrieval of the nth message

DELE <n>

the nth message is marked for deletion

3. Update

End of the retrieval phase

Messages marked for deletion are removed from server

TCP connection is closed

Internet applications

Contents

The client-server model

Name to address resolution

email


→ world wide web

peer-to-peer applications

World Wide Web

Goals

Allow browsers to browse hypertext documents stored on multiple servers



World Wide Web (2)

The five key elements of www

1. An addressing scheme that allows to identify any document stored on a server
URL : Uniform Resource Locator
2. An hypertext language that allows to easily write documents with hypertext links
HTML : HyperText Markup Language
3. An efficient and lightweight application-level protocol to exchange documents
HTTP : HyperText Transfer Protocol
4. Servers
5. Clients (browsers)

Uniform Resource Locator (URL)

Uniform Resource Locator (URL)

generic syntax : **<protocol>**://**<document>**

protocol used to retrieve document from server

http is the most common one but others are frequently used

document indicates the server and the location of the document

<user>:**<password>**@**<server>**:**<port>**/**<path>**

<user> : optional username

<password> : optional password

<machine> : hostname or IP address of the server that hosts the document

<port> : optional port number

<path> : document location on server

examples

<http://www.info.ucl.ac.be>

<http://alice:secret@inl.info.ucl.ac.be:80/index.html>

HTML

HyperText Markup Language

Language used to encode documents on the web

Keywords

```
<HTML>...</HTML>
<HEAD>...</HEAD>
<BODY>...</BODY>
<TITLE>...</TITLE>
<B>...</B>
<I>...</I>
<H1>...</H1>
<P>
<HR>
<UL>...</UL>
<OL>...</OL>
<IMG SRC="URL">
<A HREF="URL">text anchor</A>
```



HTML (2)

Example

<HTML>

<HEAD>

<TITLE>HTML test page</TITLE>

</HEAD>

<BODY>

<IMG SRC="<http://www.images.be/logo.gif>">

<H1>Web servers from UCL UCL</H1>

<HR>

<A HREF="<http://www.uclouvain.be>">UCL

<A HREF="<http://www.info.ucl.ac.be>">CSE Dept.

<A HREF="<http://www.math.ucl.ac.be>">Math

</BODY>

</HTML>

Header

Body

Image on remote server

First level title

External hypertext link

Information transfer www

HTTP 1.0 - non-persistent connection


Principle

relies on TPC service (default port : 80)

Client sends request, server sends reply

Client

Server




Information transfer [www](http://www.ietf.org/rfc/rfc2616.txt)

HTTP 1.0 - non-persistent connection

Principle

relies on TPC service (default port : 80)

Client sends request, server sends reply




Information transfer [www](http://www.ietf.org/rfc/rfc2616.txt)

HTTP 1.0 - non-persistent connection

Principle

relies on TPC service (default port : 80)

Client sends request, server sends reply




Information transfer [www](http://www.ietf.org/rfc/rfc2616.txt)

HTTP 1.0 - non-persistent connection

Principle

relies on TPC service (default port : 80)

Client sends request, server sends reply



HTTP

Client



Server



HTTP

Header contains additional information
about request sent by client

Method
GET
POST
...

Request

Method
Header

CRLF

MIME Document

Client



Server



HTTP

Header contains additional information about request sent by client

Method
GET
POST
...

Request




Client



Server



Response



Success or failure

Header contains information about server and optional parameters specific to response

HTTP

Header contains additional information about request sent by client

Method
GET
POST
...

Request




Client



Server



Response



Success or failure

Header contains information about server and optional parameters specific to response

HTTP is a stateless protocol, server does not maintain any state from one request to another

POP, FTP, SMTP are examples of stateful protocols in contrast

HTTP : Example

Client




Server


www.info.ucl.ac.be



HTTP : Example



HTTP : Example



HTTP : Methods

Methods

GET

method used to request a "document" stored on server

GET <document> HTTP/1.0

example

GET /index.html HTTP/1.0

POST

method used to send a "document" to a server

document is part of the request and encoded as a MIME document

HTTP : Request headers

Request headers

Allow to add information about the client or the request

Host: <name>

Name of the server where the document is stored

Authorization

allows to perform access control

If-Modified-Since: <date>

server will only send the requested document if the document is more recent than **date**

Referer: <url>

Information, indicates the **URL** visited by the client before this request

User-Agent: <agent>

information, indicates the browser used on the client

HTTP : Status line

Status line

Format : Version_HTTP Code Comment

Success/Failure

1xx : For information (unused)

2xx : Success

Example : HTTP/1.0 200 OK

3xx : Redirection

Request could not be handled on local server and should be sent to another server

Example :

HTTP/1.0 301 Moved permanently
attached MIME document will contain URL of document

4xx : Client-side error

examples

syntax error, unreachable URL, unauthorised, ...

5xx : Server-side error

examples :

internal error, method not implemented on server, ...

HTTP : Response headers

Header

Optional information about the server, the response or the document attached to the response

Date

date of the document attached to response

example : `Date: Wed, 05 Sep 2001 13:27:34 GMT`

Server

Name and version of http server used

example :

`Server: Apache/1.3.20 (Unix)ApacheJServ/1.1.2 PHP/4.0.6`

Content-*

MIME header of the attached document

example :

`Content-Length: 5891`

`Content-Type: text/html`

HTTP 1.1

HTTP 1.1

HTTP 1.0

a single TCP connection used to transmit
a single document (html file, image,...)
the establishment and release of the TCP connection
induce a significant overhead, in particular for small pages

HTTP 1.1

HTTP 1.0

a single TCP connection used to transmit
a single document (html file, image,...)
the establishment and release of the TCP connection
induce a significant overhead, in particular for small pages

HTTP 1.1

uses a single persistent TCP connection
This TCP connection can be used for several requests
and the corresponding responses
the cost of establishing and releasing the TCP connection
is amortised over multiple requests
Although HTTP 1.1 uses a single TCP connection for
multiple requests, HTTP 1.1 remains stateless

HTTP 1.1 : Persistent connection

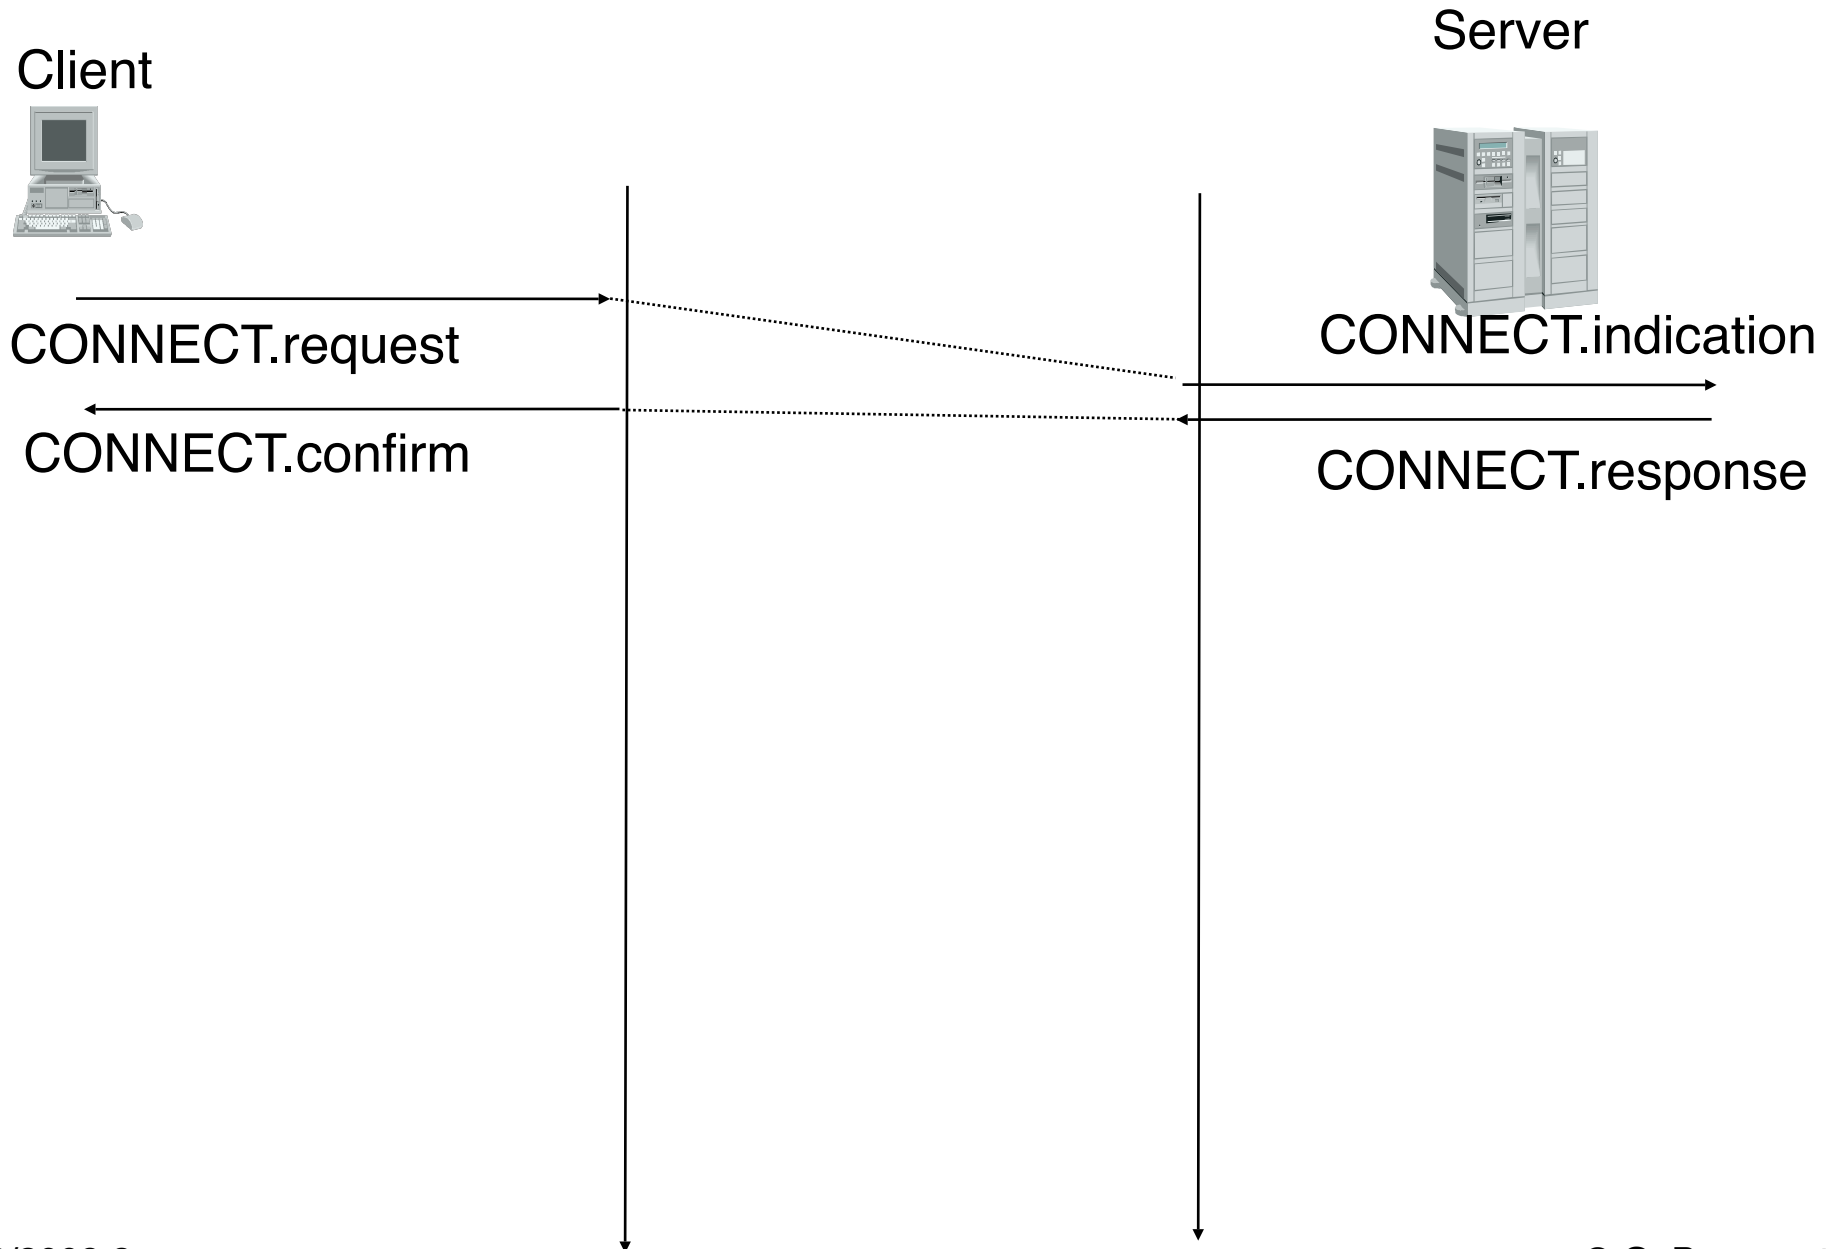
Client



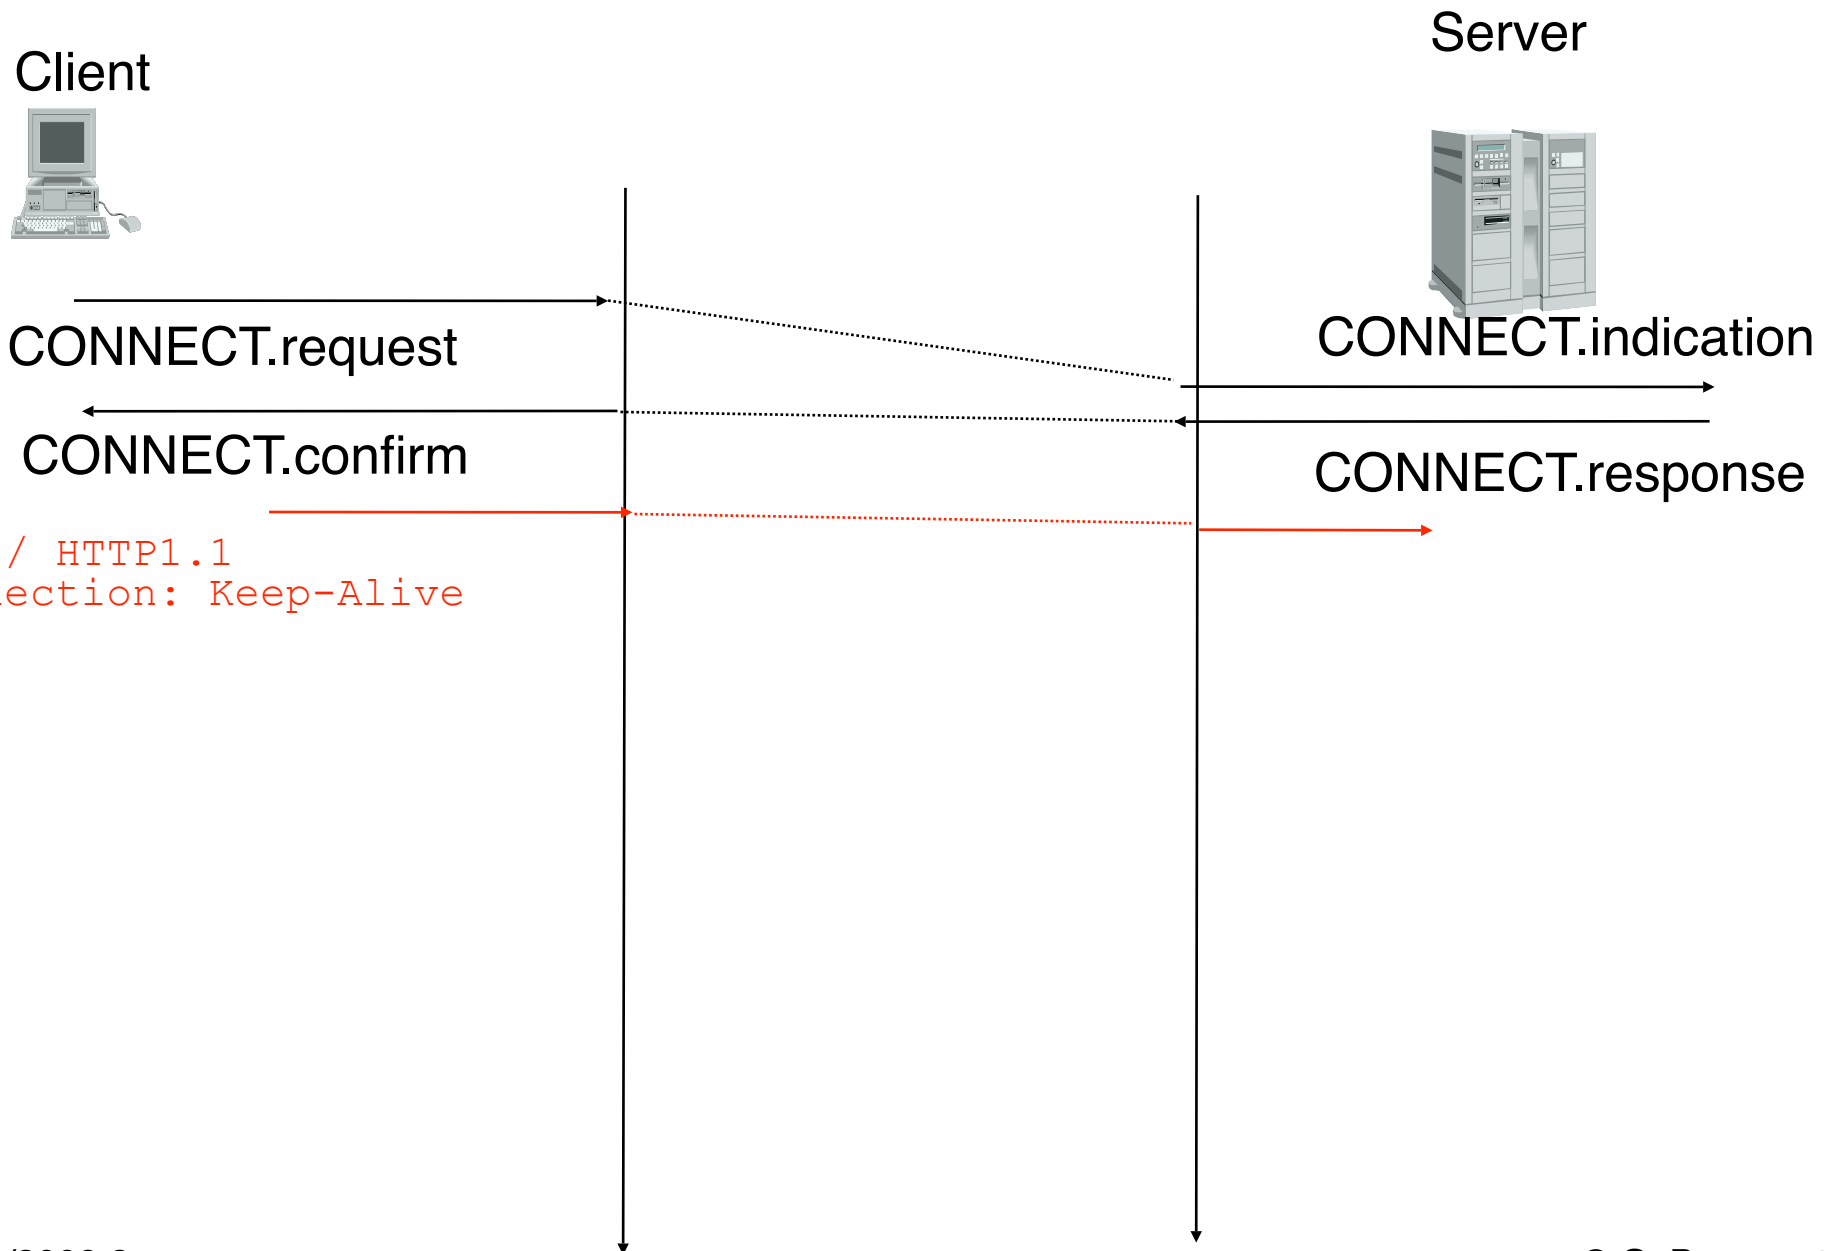
Server




HTTP 1.1 : Persistent connection




HTTP 1.1 : Persistent connection



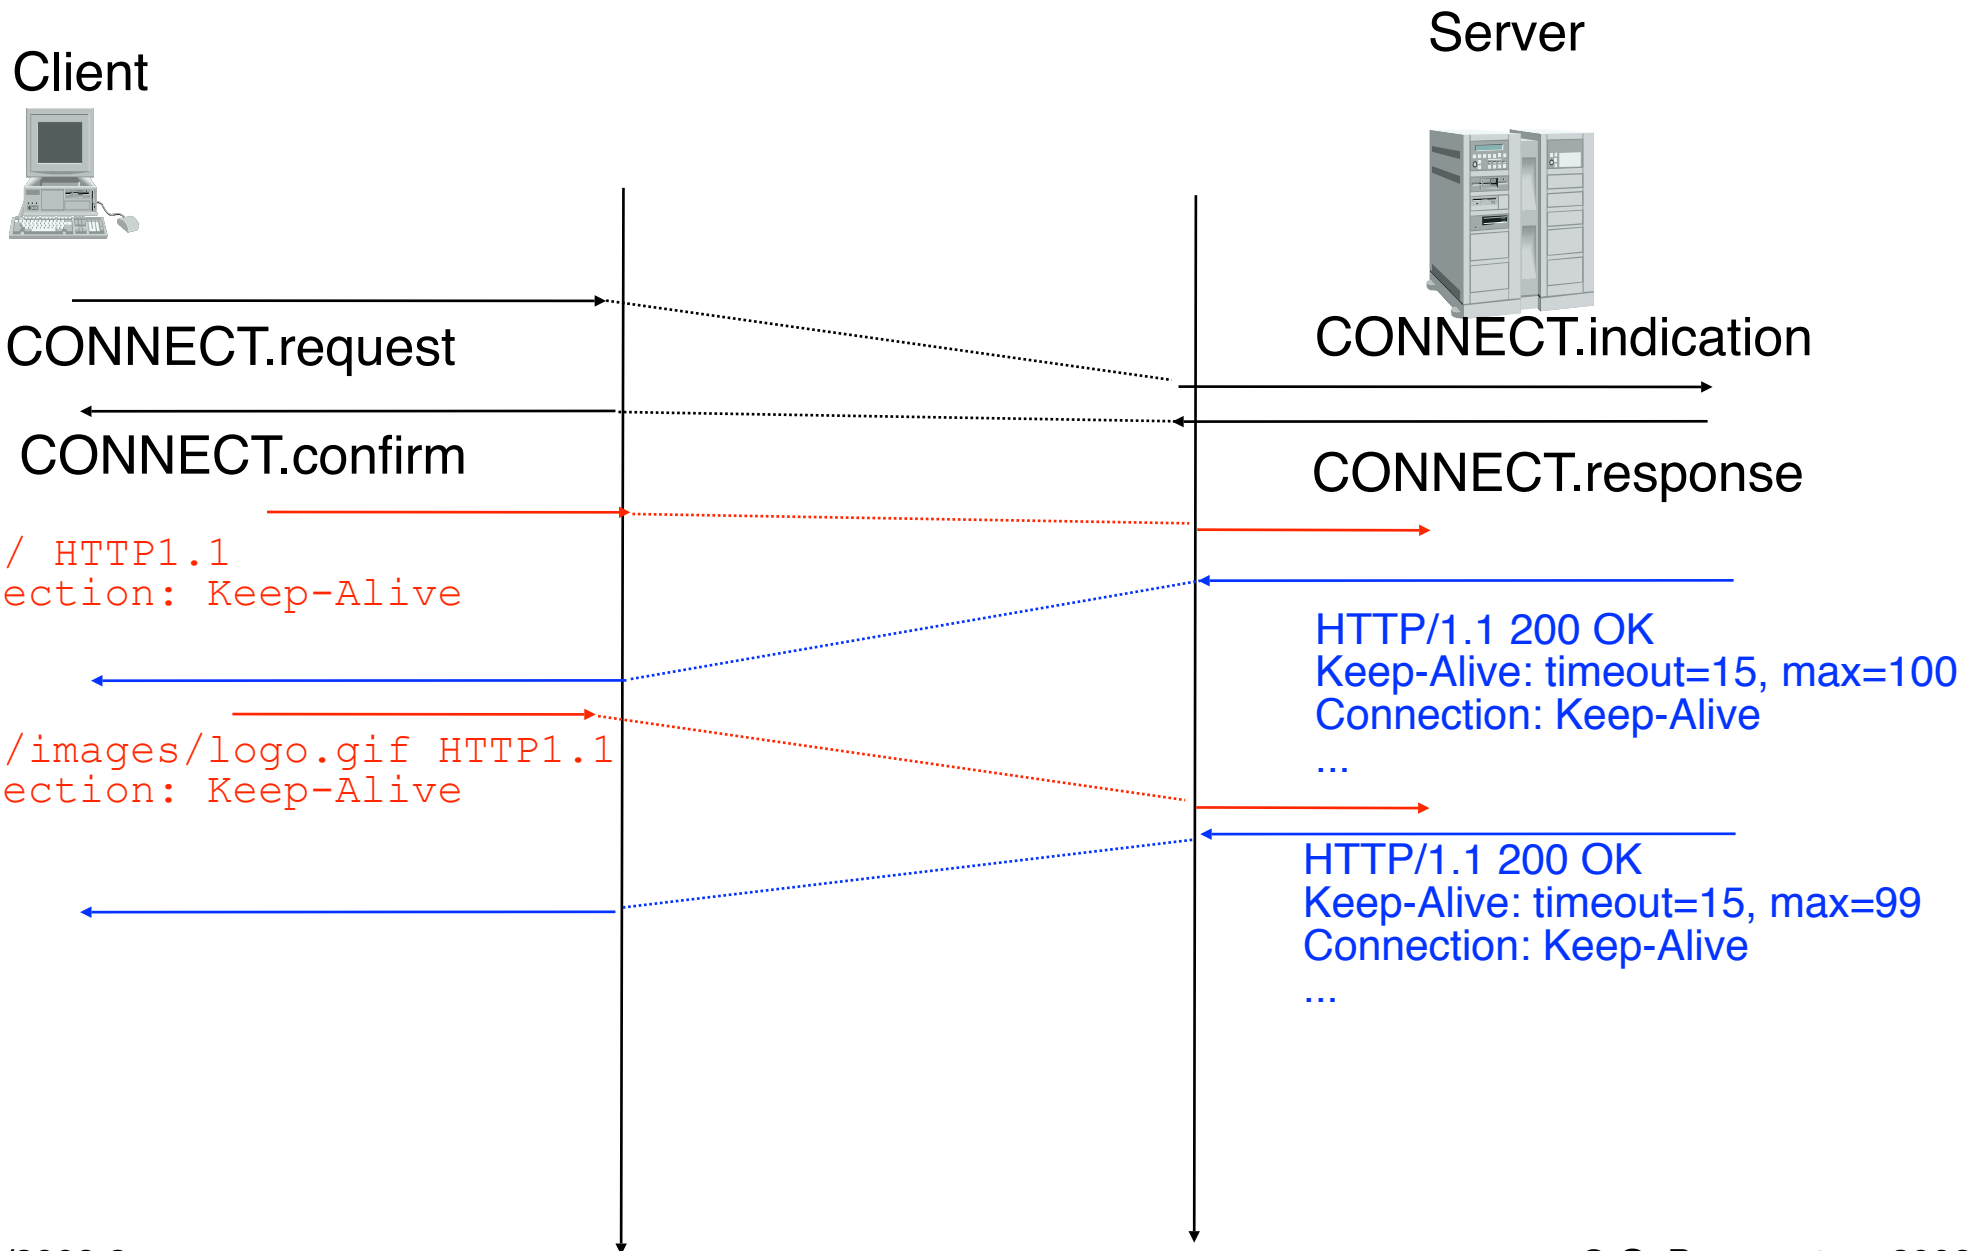
HTTP 1.1 : Persistent connection



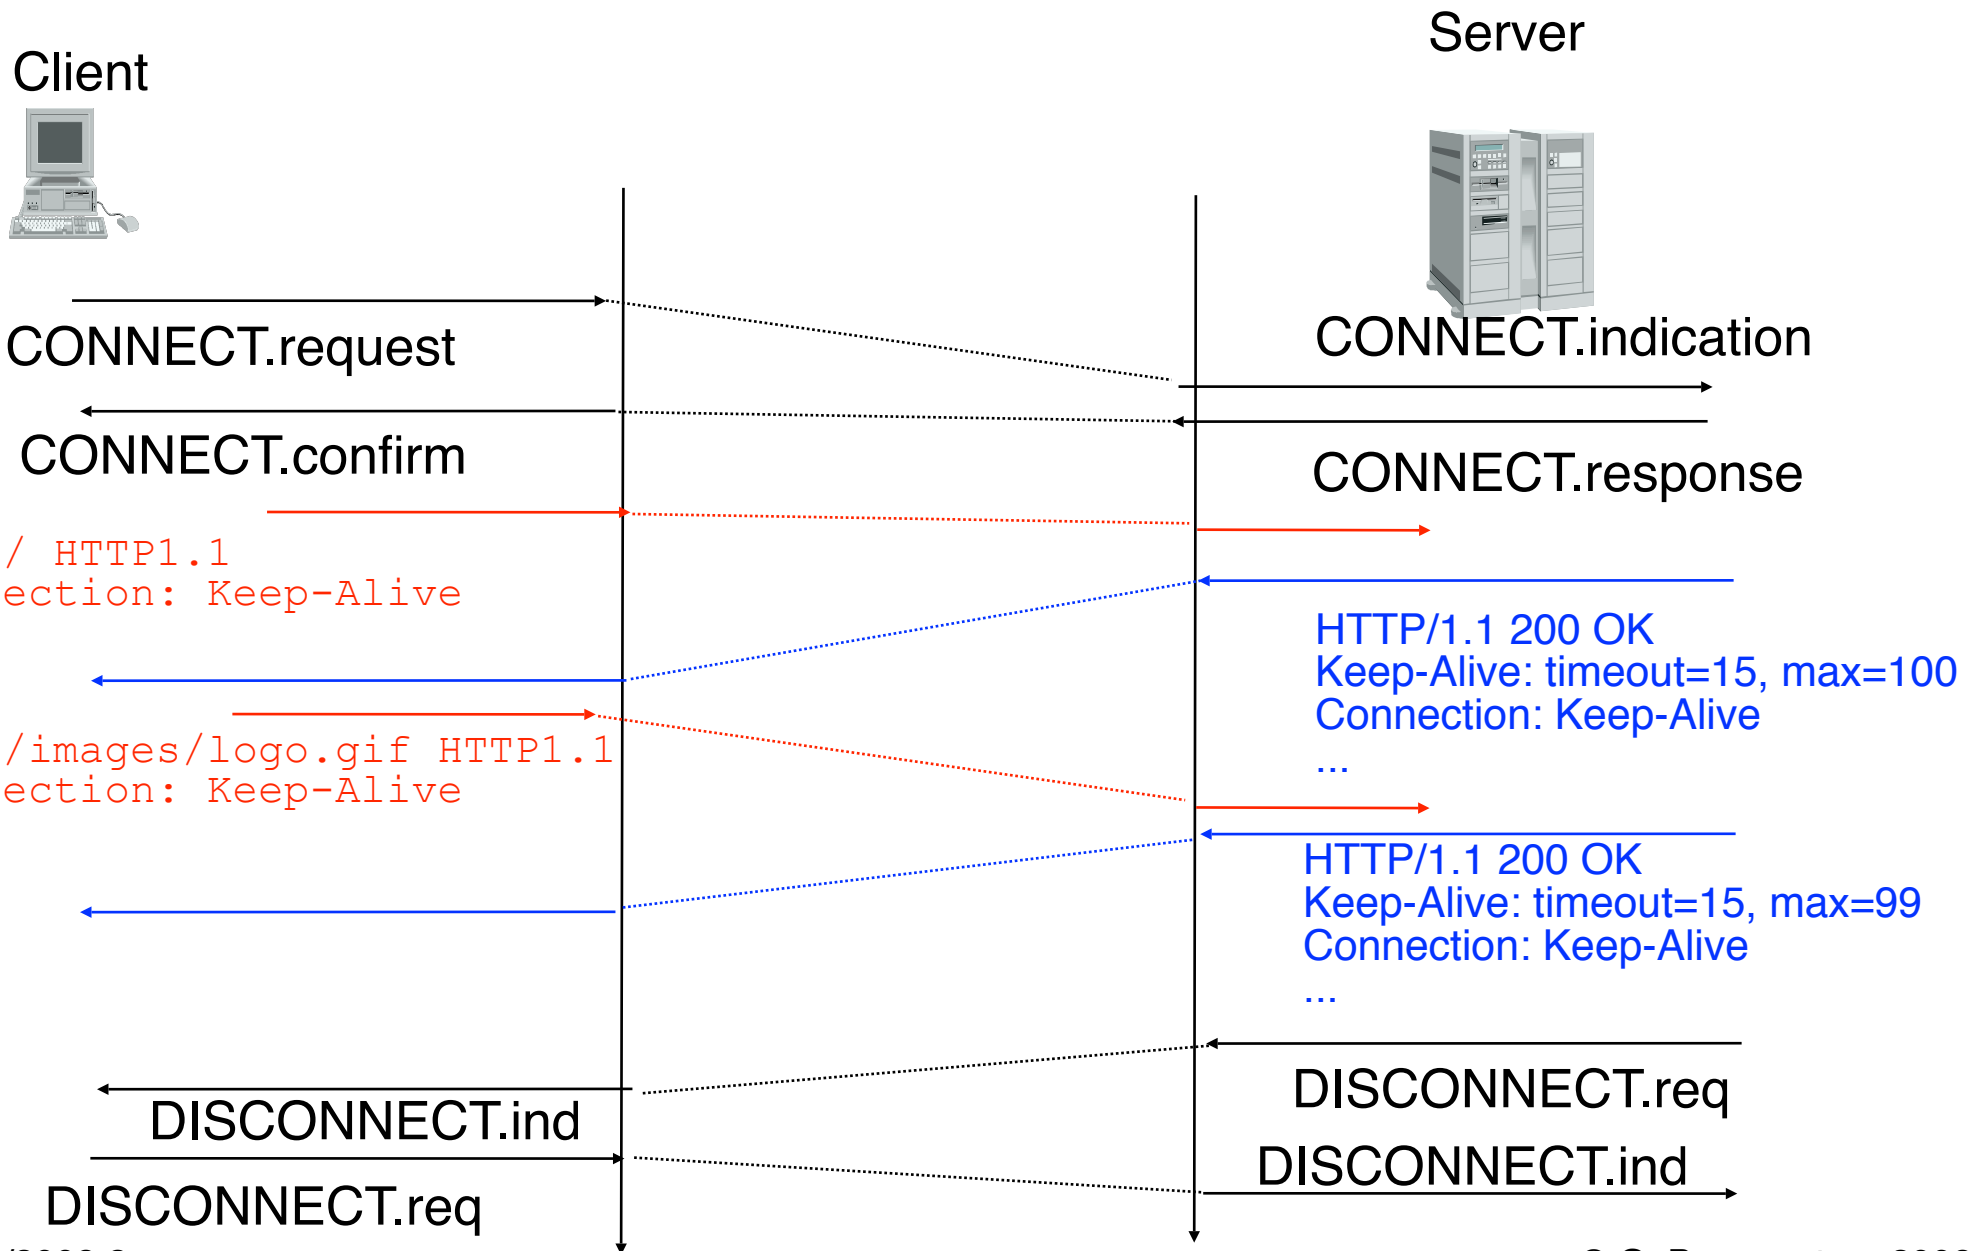
HTTP 1.1 : Persistent connection



HTTP 1.1 : Persistent connection



HTTP 1.1 : Persistent connection



Improving performance

Observation

Many pages are requested multiple times or from close endhosts

Solution

local cache on each client

`if-modified-since` header helps

one cache for multiple endhosts

Improving performance

Observation


Many pages are requested multiple times or from close endhosts

Solution

local cache on each client

`if-modified-since` header helps

one cache for multiple endhosts



HTTP Cookies

Example

Client




Server

Normal response



HTTP Cookies

Example



HTTP Cookies

Example

Client



Server

GET / HTTP1.1
...

HTTP/1.0 200 OK
Set-Cookie: machin
...

Normal response

Browser saves cookie

GET /doc HTTP1.1
Cookie: machin
...

HTTP/1.1 200 OK
...

*Response is function
of URL and cookie*

GET /images/t.gif HTTP1.1
Cookie: machi
...

*Browser sends cookie in all
requests sent to server*