



Systems Calls Input/Output (I)

Alain NINANE – RSSI UCL
23 Février 2016





System Calls

- Interfaces to the kernel
 - Written and used in C programs
 - Can be used from other languages programs (Fortran, ...)
 - (UNIX allows mixing of languages ... !)
- Limited set for each Unix kernel
- Kernel entry point is unique
 - System calls are indexes in a 'C' switch statement
 - MacOSX: `/usr/include/sys/syscall.h`
 - Linux: `/usr/include/sys/syscall.h`
`/usr/include/bits/syscall.h`
`/usr/include/asm/unistd.h`
 - Systems calls are a finite set
 - Unless rewriting/upgrading the kernel :-)
 - Number depends on the type of unix/linux kernel



System Calls

- I/O
 - open, creat, read, write, lseek, dup, ioctl, close, ...
- File Management
 - link, unlink, mknod, mount, umount, stat, fcntl, ...
- Protection
 - access, chmod, chown, getuid, setuid, umask, ...
- Processes
 - exec, fork, alarm, signal, chdir, getpid, brk, times
- Misc
 - acct, time



Portability issues (I)

- Because of UNIX diversity
 - Programs using syscalls directly
 - are no more “Standard C programs”
 - are UNIX C programs
 - non-portable across UNIX'es
 - Syscalls annoyances
 - Header files
 - Types of arguments
- Importance of man pages
 - man 2 syscall_name



Portability Issues (II)

- Header files ...
- E.g. read()
 - macosx
 - #include <unistd.h>
 - linux
 - #include <sys/types.h>
 - #include <sys/uio.h>
 - #include <unistd.h>



Portability Issues (III)

- Type of arguments ...
- E.g. signal()
 - macosx
 - typedef void (***sig_t**) (int);
 - **sig_t** signal(int sig, **sig_t** func);
 - linux
 - typedef void (***sighandler_t**)(int);
 - **sighandler_t** signal(int signum, **sighandler_t** handler);



Good Prog. Practices (I)

- Use Compilation Directives
 - CPP - C Pre Processor
 - CPP is called before real compilation to resolves directives
- Directives
 - Example in a C program
 - `#ifdef MYMACOSX`
 -
 - `#endif`
 - User defined
 - `cc -DMYMACOSX test1.c`
 - Pre defined by compiler
 - `cc test1.c`



Good Prog. Practices (II)

```
#ifdef __MACH__
```

```
    Code for Apple UNIX
```

```
#elif __alpha
```

```
    Code for Digital/Compaq/HP UNIX
```

```
#elif __linux__
```

```
    Code for Linux
```

```
#else
```

```
#error "Unsupported O/S"
```

```
#endif
```


CPP - Predefinitions (I)



<code>__MACH__</code>	Kernel Mach (Apple ...)
<code>__alpha</code>	Digital Alpha 64 bit CPU
<code>__linux__</code>	A Linux system
<code>__gnu_linux__</code>	Linux gcc compiler
<code>__i386</code>	An Intel like 386 proc.
<code>__x86_64</code>	An Intel 64 bit proc.



CPP - Predefinitions (II)

- How to get them ... ?
 - Experience ...
- Linux, MacOSX
 - `cpp -dM someFile.c ...`
- Other Unixes
 - `man cc`
 - `man gcc`
 - `man cpp`



Autoconfiguration

- Unix Source Code Distr./Install.
 - tar zcvf code.tar.gz
 - Uncompression process
 - unzip - untar
- Compilation
 - ./configure
 - generates a Makefile with correct cpp definitions
 - Compilation
 - make
 - make install



Standards

- Standards developed around UNIX
 - POSIX: IEEE 1003.1
- Primitive system data types
 - Hides implementation details
 - typedef's in `<sys/types.h>`
 - `ino_t` file serial number
 - `off_t` file size
 - `size_t` i/o size
 - `time_t` system time
 - `pid_t` process id



Errors handling (1 - errno)

- All syscalls returns -1 upon error

```
int fd;  
fd = open("myfile",O_RDWR);  
if( fd == -1 )  
{ fprintf(stderr,"open error !\n");  
  return -1;  
}
```

... continue reading ...

- Return code MUST always be tested
- errno: external global integer
 - Give details about the error
 - #include <errno.h>
 - #include <sys/errno.h>
 - errno NOT cleared by "next" syscall

Error handling (II - errno)



```
#include <errno.h>
extern int errno;
int fd;
if( (fd = open("myfile",O_RDWR) == -1 )
{  switch( errno )
    {      case EPERM:
            ...
            break;
            case ENOENT:
            ...
            break;
            default:
            ...
    }
    return -1;
}
```



Error handling (III - perror)

- `#include <errno.h>`
- `void perror(const char *s)`
 - `s` is not NULL
 - print on the standard error:
 - `s: clear_text_errno_message`
 - `s` is NULL
 - print on the standard error:
 - `clear_text_errno_message`

Error handling (IV - example)



```
#include      <errno.h>
int fd;

if( (fd = open("myfile",O_RDWR)) < 0 )
{
    perror("open myfile");    /* open failure */
    return -1;
}
... access granted to file ...
```

```
% ./a.out
open myfile: Permission denied
```




I/O System Calls (I)

- UNIX provides a uniform interface for performing I/O operations on resources
- Resources can be ...
 - files
 - terminals
 - pipes
 - tapes
 - network sockets
 - ... other devices ...



I/O System Calls (II)

- Accessing the resource
 - `open()`, `creat()`, `close()`, ...
- R/W data from/to the resource
 - `read()`, `readv()`, `write()`, `writew()`, ...
- Controlling resource
 - `lseek()`, `ioctl()`, ...
- Special
 - `select()`, `dup()`, ...
- *Network*
 - `socket()`, `bind()`, `accept()`, ...



I/O System Calls (III)

- Ressource identified by a “small integer”, the file descriptor
 - the file descriptor is a program abstraction to access the ressource
- I/O Systems Calls are non buffered
 - each system call implies a kernel and device operation



File Descriptor (I)

- Opened files designed by a file descriptor (fd)
- fd is a “small” integer used to perform I/O
- fd returned by
 - `open()`
 - `creat()`
 - `socket()`
- `int fd = open(...);`
- `read(fd, ...);`



File Descriptor (II)

- For each process, the kernel maintains a table of open files
- Illustration:
 - command `ls -l /proc/PID/fd` = list open files
- Demo
 - `./testfd`
 - `ls -l /proc/PID/fd`
 - `/usr/sbin/lsof -p PID`



File Descriptor (III)

- Observation
 - fd for “messages” file is 3
 - there exists fd 0, 1 and 2
 - connected to my “terminal”

Fd	Purpose	Initial Device
0	standard input	keyboard
1	standard output	terminal
2	standard error	terminal



File Descriptor (IV)

- Pre defined for each process
 - `/* Read from standard input */`
 - `int ret = read(0,buffer,len);`
 - `/* Write on the standard output */`
 - `int ret = write(1,buffer,len);`
 - `/* Write on the standard error */`
 - `int ret = write(2,buffer,len);`



Open() system call (I)

```
#include <sys/types.h>
```

```
#include <stat.h>
```

```
#include <fcntl.h>
```

```
int open(const char *path, int flags, mode_t mode);
```

```
int open(const char *path, int flags);
```

```
int creat(const char *path, mode_t mode);
```

- Opens a file for reading or writing
- Returns the fd or -1
- **path**: string designating the resource in the file system



Open() system call (II)

- **flags**

O_RDONLY	read only
O_WRONLY	write only
O_RDWR	read/write
O_CREAT	creat if ! exist
O_APPEND	pointer set to eof
O_TRUNC	if exists - truncate
O_EXCL	creat fails if exists



Open() system calls (III)

- **flags**

- lot of **non-portable** flags
- on linux ... man 2 open
 - **O_SYNC**
 - **O_NOFOLLOW** (free bsd)
 - **O_DIRECTORY** (linux)
 - **O_LARGEFILE**
 - Allows > 2 GB files on some 32 bit systems



Open() system call (IV)

- **mode**
 - Argument used with O_CREAT
 - UNIX file mode (or'ed with umask)
- **errors** (specified by errno)
 - EACCES
 - access not allowed
 - ENOENT
 - file does not exist or missing path component
 -



Open() system call (V)

```
int fd;  
int flags = O_CREAT|O_TRUNC|O_WRONLY;  
  
if( (fd = open("../file",flags)) == -1 )  
{  perror("open file");  
    exit(1);  
}
```

- **Discussion: why a write only file ?**



Creat() & other system call

- `int creat (const char *path, mode_t mode)`
 - simple replacement for open with `O_CREATE`
- **Directories**
 - `int mkdir(const char *path, mode_t mode);`
- **Special files**
 - `int mknod(const char *path, mode_t mode, dev_t dev);`



Read() system call (I)

```
#include <unistd.h>
```

```
ssize_t read(int fd, void *buf, size_t len);
```

fd: file descriptor of the opened file to read

buf: address of the pointer where input must be copied

len: maximum number of bytes to read

Returns:

ret : the number of bytes actually read.

0 : EOF is reached

-1 : an error occurred (see errno or perror())

Warning:

ret may be lower than requested len



Read() system call (II)

- `ret < len`
 - `ret = -1`
 - error condition (see `errno`)
 - `ret = 0`
 - no more data to read
 - `0 < ret && ret < len`
 - file size smaller than `len`
 - input is a terminal and user hit CTRL-D (EOF)
 - input is a raw device (e.g. a tape drive, a terminal)
 - input is a network socket



Read() system call

- Always be prepared to read fewer data than expected
- Even if input is a “standard” file
 - std file can easily be substituted by
 - a pipe
 - a network socket
 - ...

Handling of incomplete reads



```
#define LEN    8192
char buffer[LEN];

int remaining = LEN;
int current_p = 0;

while( (ret = read(fd,&buffer[current_p],remaining)) > 0 )
{
    remaining -= ret;
    current_p += ret;
}

if( ret )
{
    perror("read");
    exit(1);
}
else
{
    printf("eof reached !");
}
```



Write() system call

```
#include <unistd.h>
```

```
ssize_t write(int fd, void *buf, size_t len);
```

fd: file descriptor of the opened file to write

buf: address of the pointer where output must be read

len: maximum number of bytes to write

Returns:

ret : the number of bytes actually written.

0 : EOF is reached

-1 : an error occurred (see errno or perror())

Warning:

ret may be lower than requested len



Some words about files

- Files are seen as sequence of bytes
 - no records
 - physical blocks not visible
- Current file offset tight to file descriptor
- Blocks devices (raw disk i/o, tapes) are seen as sequence of blocks
 - managing unit is the block
 - whole blocks are read/written at once



Lseek() system call

To move the pointer into the file ...

```
#include <sys/types.h>
```

```
#include <unistd.h>
```

```
off_t lseek(int fd, int offset, off_t whence);
```

fd: file descriptor of the opened file

offset: number of bytes to be moved by the pointer

whence: starting point

SEEK_SET: beginning of file;

SEEK_CUR: from the current position;

SEEK_END: from the end of file;



ioctl() system call

- for all other i/o operations
 - e.g. move a tape to next file
 - operations described in the driver (4) part of man
 - e.g. man st

```
#include <sys/mtio.h>
```

```
int ioctl(int fd, MTIOCTOP, (struct mtop *)mt_cmd);
```

```
/* Structure for MTIOCTOP - mag tape op command: */
```

```
struct mtop
```

```
{ short mt_op; /* operations defined below */
```

```
  int mt_count; /* how many of them */
```

```
};
```

```
MTBSF      Backward space over mt_count filemarks.
```

```
MTERASE    Erase tape.
```

```
MTFSF      Forward space over mt_count filemarks.
```



Close() system call

- `int close(int fd)`
 - Dissociates the “small integer” fd from the file
 - Returns 0 or -1
 - All files are closed on program exit.