

Systems Calls - Procs

Alain NINANE - UCL
15 mars 2016





Process - A definition (I)

- **A process is a program being executed**
 - **program: a static object**
 - An executable file on the file system
 - A file with the mode 'x' (executable) bit set
 - Executable program file
 - machine readable code (a compiled program)
 - human readable code (a program for an interpreter)
 - the interpreter is itself a compiled program
 - **process: a dynamic object**
 - transient in memory during execution

Process/Program - Examples



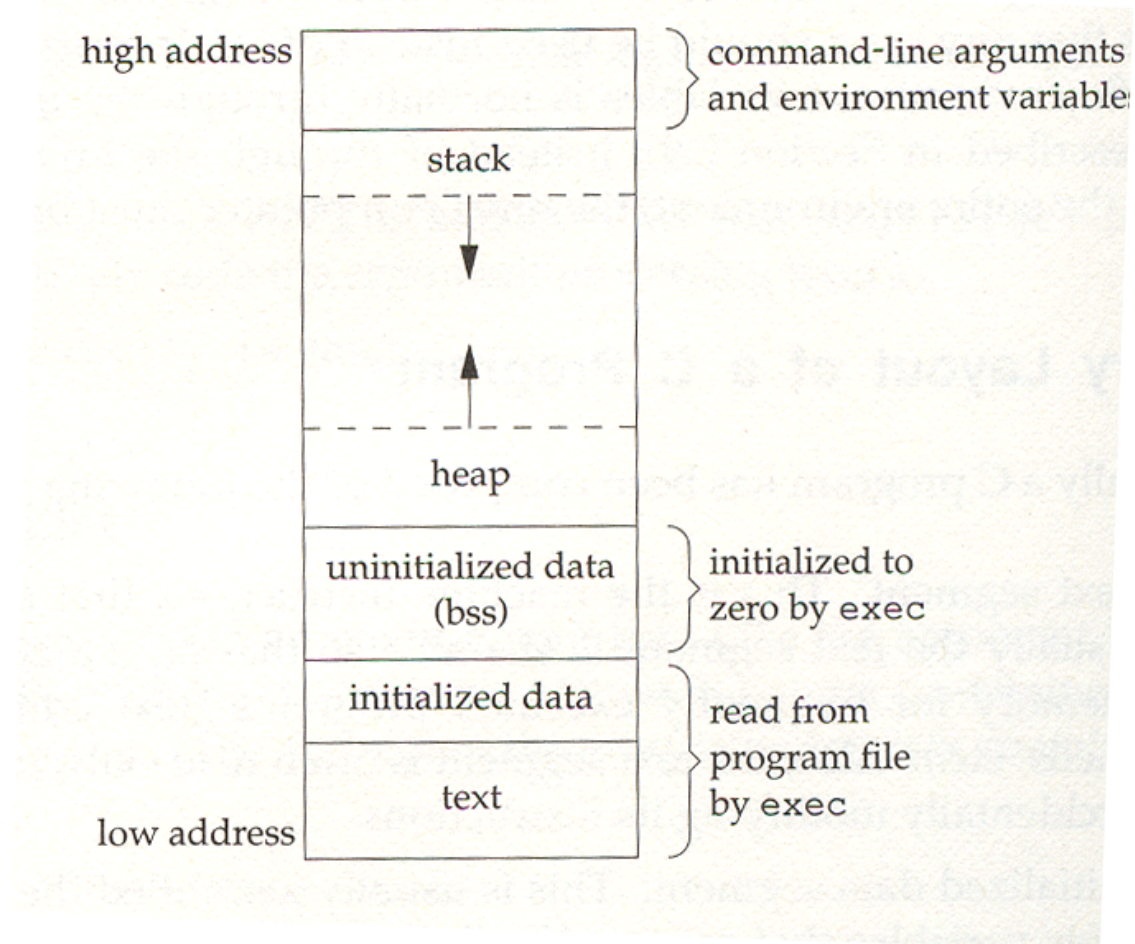
- **Compiled program**
 - Needs compilation: `cc -o test_01 test_01.c`
 - Demonstration
 - `./test_01 &`
 - `ps tree -G -p | more -c`
 - `ps axjf | grep 500`
- **Interpreted program**
 - Needs change mode: `chmod +x test_02.sh test_03.pl`
 - shell
 - `./test_02.sh &`
 - `ps axjf | grep 500` # 500 is my own uid on the test system !!
 - perl
 - `./test_03.pl &`
 - `ps axjf | grep 500`
 - Note: magic characters (`/usr/share/file/magic`)



Process - A definition (II)

- A process is a program being executed
- **A process contains in memory data**
 - **Text:** the instructions to be executed
 - `size a.out`
 - `nm a.out`
 - `objdump -d a.out`
 - **Global data:** initialized or not
 - **Local data:** on the stack
 - **Heap:** free memory for `malloc()`
 - **Execution environment:**
 - Program: counter, status, ...
 - System: file table, working dir., control. terminal, user, ...
 - **Environment variables:** system & user defined values
 - Environment variables is a set of pair variables/values that can be accessed by the process **and its childrens**

Process - In Memory Data





Environment - variables (I)

- A process holds environment variables
 - “System V family shell”
 - `printenv`, `export`
 - “Berkeley family shell”
 - `printenv`
- Settings users’s defined environment variables
 - “System V family shell”
 - `export MYENVDATA="/tmp"`
 - “Berkeley family shell”
 - `setenv MYENVDATA /tmp`



Environment - variables (II)

- Environment is a set of pairs
 - name=value
 - Example:
 - PATH=./usr/bin:/bin
 - USER=nina
 - HOME=/disks/home/n/nina
 - TERM=vt100
 - Access from a user's program
 - **extern char **environ;** /* or better !!!! */
 - char *getenv(const char *name);

Environment - PATH explained



- The PATH variable
 - Directory search list for executables programs
 - Example:
 - `/usr/kerberos/bin:/usr/local/bin:/bin:/usr/bin:/home/nina/bin`
 - a standard search path
 - `./:/home/nina/bin:/usr/local/bin:/bin/....`
 - a path with priority to owner & local directories
 - `./:/usr/ucb/bin:/usr/bin:....`
 - a specialized path on a Sun Solaris system (System V) with priority to compatibility Berkeley commands
 - example:
 - BSD: `ps aux`
 - System V: `ps -eaf`

strace demonstration



- `strace sh -c ls 2>&1 | more -cf`
 - Illustrates
 - system calls:
 - `execve()`, `mmap()`, `open()`, `read()`, `fstat()`, `lseek()`
 - shared libraries
 - `/etc/ld.so.cache`
 - `/lib64/libdl.so.2`
 - observation:
 - `mmap` **NOT** unmapped by `close`
 - search path
 - through `stat()` system call



Process - A definition (III)

- A process is a program being executed
- A process holds in memory data
- **A process:**
 - **is identified by a number**, unique in the UNIX system
 - the **pid** : process identifier
 - pids are positive 16 bits integer
 - maximum of ~32.000 processes
 - **executes on behalf of a user**
 - the **uid** : identifies the user executing the process



Process - Starting a process

- Two point of view
 - at the command level:
 - by means of a command interpreter (e.g. shell)
 - from within a program:
 - by means of system calls
- One point of view
 - The **ONLY** way to create a new process is when an **existing** process calls the **fork()** syscall.



fork - Process creation syscall

- `#include <sys/types.h>`
- `#include <unistd.h>`
- `pid_t fork(void)`
 - **The existing process is duplicated**
 - Processes have:
 - same copies of in-memory data, *with some exceptions ...*
 - Returns:
 - the original (**parent**) process gets child's pid
 - the forked (**child**) process gets 0
 - Errors:
 - returns -1 (set errno)

fork - The simplest example



- `#include <sys/types.h>`
- `#include <unistd.h>`
- `main()`
- `{ pid_t ret = fork();`
 - `if(ret > 0)`
 - `{ /* ret > 0 -> parent code */`
 - `myPid = getpid(); /* ret != myPid !!! */`
 - `}`
 - `else if(ret == 0)`
 - `{ /* ret == 0 -> child code */`
 - `myPid = getpid();`
 - `}`
 - `else`
 - `{ /* ret < 0 -> parent in error */`
 - `}`



fork - shared/unshared data

- Both processes **share**
 - same code
 - same values in variables (except ret value)
 - same open files, file pointers and file descriptors
 - same current working directory
 - same controlling terminal/window
- Both processes **do not share**
 - execution time counters (cleared for the child)
 - semaphore values
 - file locks
 - pending alarm or signal (cleared for the child)



fork - special cases

- Experience:
 - `./test_04_[init|zombie] &`
 - `ps -U nina lx`
- **the parent terminates before the child**
 - the child is attached to process “init” (PPID = 1)
- **the child terminates before the parent**
 - the child becomes a zombie (Z status)
 - a terminated child still exists until the parent:
 - terminates
 - check the child `exit()` value e.g. with `wait()` syscall
 - a zombie process doesn't consume any resources
 - except entries in the process table



fork - synchronization

- A parent process must sometimes wait until one of its children terminates:
 - synchronization
 - **system calls wait() and exit() [or _exit()]**
- General principe (**demo test_05**)
 - if(fork())
 - /* parent process */
 - wait()
 - else
 - /* child process */
 - exit(value)
 - father continues ...



wait - synchronization

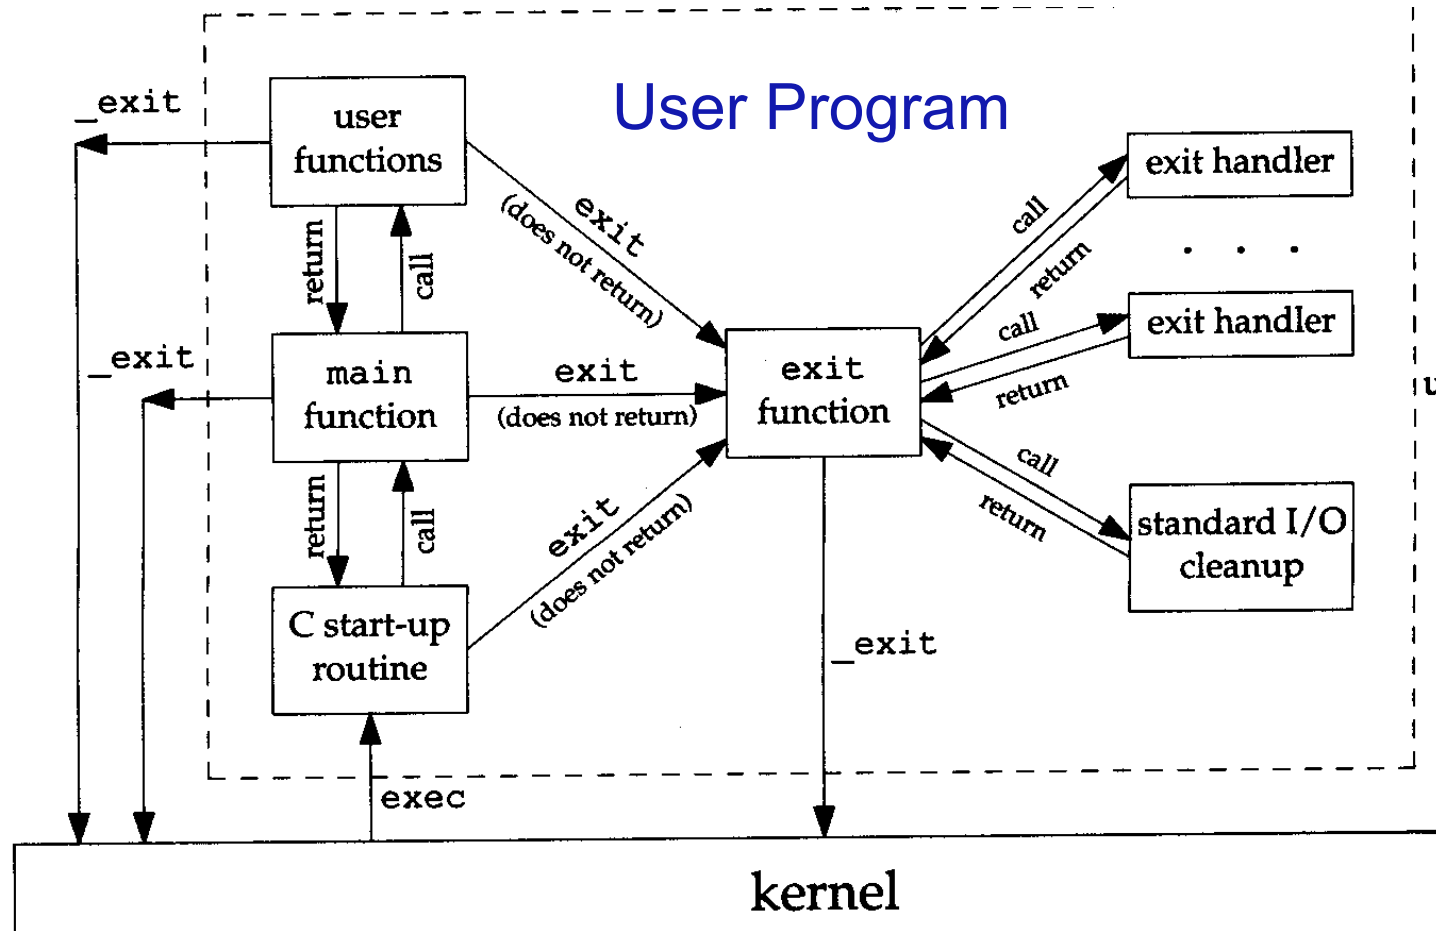
- `#include <sys/types.h>`
- `#include <sys/wait.h>`
 - `pid_t wait(int *status);`
 - `pid_t waitpid(pid_t pid, int *status, int options);`
- Wait for a child to terminates
 - suspend execution of the “current” process until one of the childrens terminates
 - **returns pid of waited child** or -1
 - can wait for a specific child
 - options & status: read the manual
 - can contains `exit_value`
- A process can wait for several processes
 - Implies several `wait()` or `waitpid()` calls !



exit - exit() vs _exit()

- `_exit(status)` is the “real” system call
 - terminates the current process immediately
 - opened fd's are closed
 - children are “moved” to another parent (init - pid 1)
 - parent's process receive a SIGCHLD signal
 - status is returned to the parent process
- `exit(status)` is a function
 - executes terminations functions
 - such as flushing streams
 - call void (*functions)(void) added to the process with `atexit()`
 - call `_exit(status)`
- demonstration: **test_06** and **test_06_exit**

exit - execution flow





exit - programming technique

- `exit()` should be used in place of `_exit()`
 - calls exit handlers
 - flush stdio streams
- Convention ...
 - program terminates normally: `exit(0)`
 - program terminates badly: `exit(something)`
 - something is between 1 and 255 (8 bits)
 - -1 --> 255
 - exit value can be checked by parent process
- Demo: check return value of `test_06[_exit]`
 - `./test_06[_exit]`
 - `echo $?`



fork - starting processes

- So far so good ...
- `fork()` starts a new process which is ...
 - just a copy of the parent process ...
- so we just have
 - a process tree (arborescence) of clones
- so we need something more
 - `exec()` family of functions



exec - executes a program

- executes - not start !!
- process is already started by fork
- **exec()** will overlay the current process with a new one taken from a new image file
- the process continue at the main() function of the image file
- since the calling image is lost, **exec()** should never returns

exec - the simplest example



```
#include <stdio.h>
main()
{  if( fork() == 0 )
    {  execl("bin/ls","ls","-l","/etc",(char*)NULL);
        /* SHOULD NEVER GETS HERE */
        exit(1);
    }
    /* father process waits */
    wait(&status);
    exit(status);
}
```

exec - a family of function



- **exec()** is a family of functions and **1** system call

```
#include <unistd.h>
```

```
int execl(const char *pathname, const char *arg0, ... /* (char *) 0 */ );
```

```
int execv(const char *pathname, char *const argv[]);
```

```
int execl(const char *pathname, const char *arg0, ...  
          /* (char *) 0, char *const envp[] */ );
```

```
int execve(const char *pathname, char *const argv[], char *const envp[]);
```

```
int execlp(const char *filename, const char *arg0, ... /* (char *) 0 */ );
```

```
int execvp(const char *filename, char *const argv[]);
```

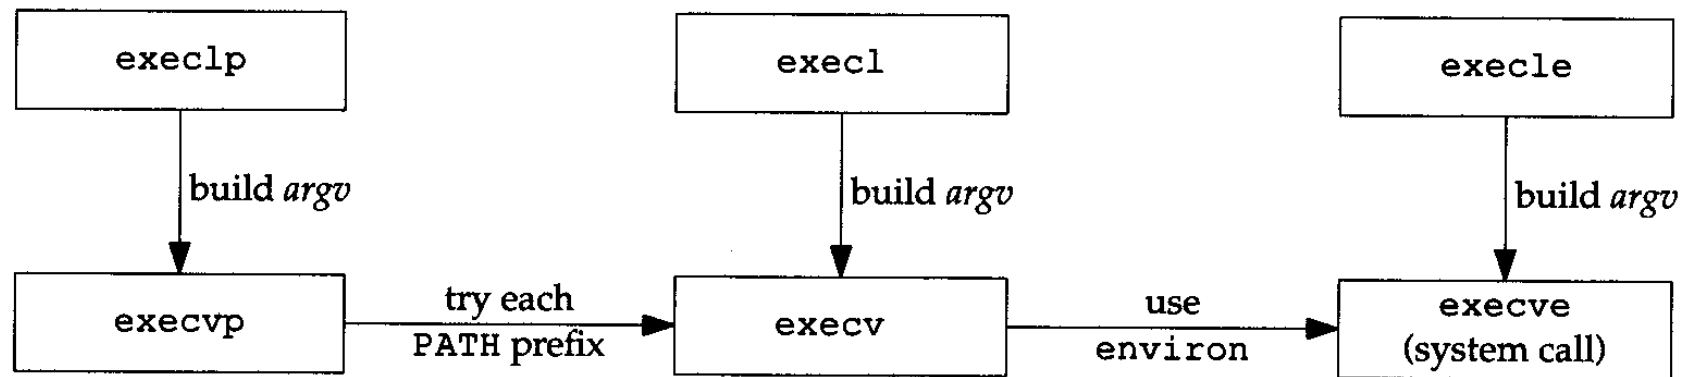
All six return: -1 on error, no return on success



exec - a family of functions

- functions differs by:
 - enabling or not the search PATH
 - execl, execv: pathname must be exact
 - execlp, execvp: search pathname in PATH
 - passing or not the environment
 - execl, execve: environment variables passed as an array
 - extern char **environ;
 - passing arguments “inline” or through array
 - execl, execlp: args as inline
 - execv, execvp: args as array
- 1 system call: **execve()**

exec - 6 functions pictures



exec - i/o redirection



- Example - i/o redirection in child process */

```
main()
{  if ( fork() == 0 )
    {  close(1);
        open("myoutput", O_CREAT|O_RDWR,0644);
        execl("/bin/date","date",(char*)NULL);
    }
    wait();
    ...
    exit(0);
}
```



exec - pass commands to shell

- Example - pass command to a new shell */

```
main()
{  if ( fork() == 0 )
    {      execl("/bin/sh","sh",
                "-c","ls /etc | wc > test_08.out",(char*)NULL);
    }
    wait();
    exit(0);
}
```

- Can be replaced by the system() function
 - system("ls /etc | wc > test_08.out");